Using Visualizations for Music Discovery

ISMIR 2009

Justin Donaldson
Paul Lamere

Using Visualizations for Music Discovery

ISMIR 2009

Justin Donaldson
Paul Lamere

Or ...

85 ways to visualize your music

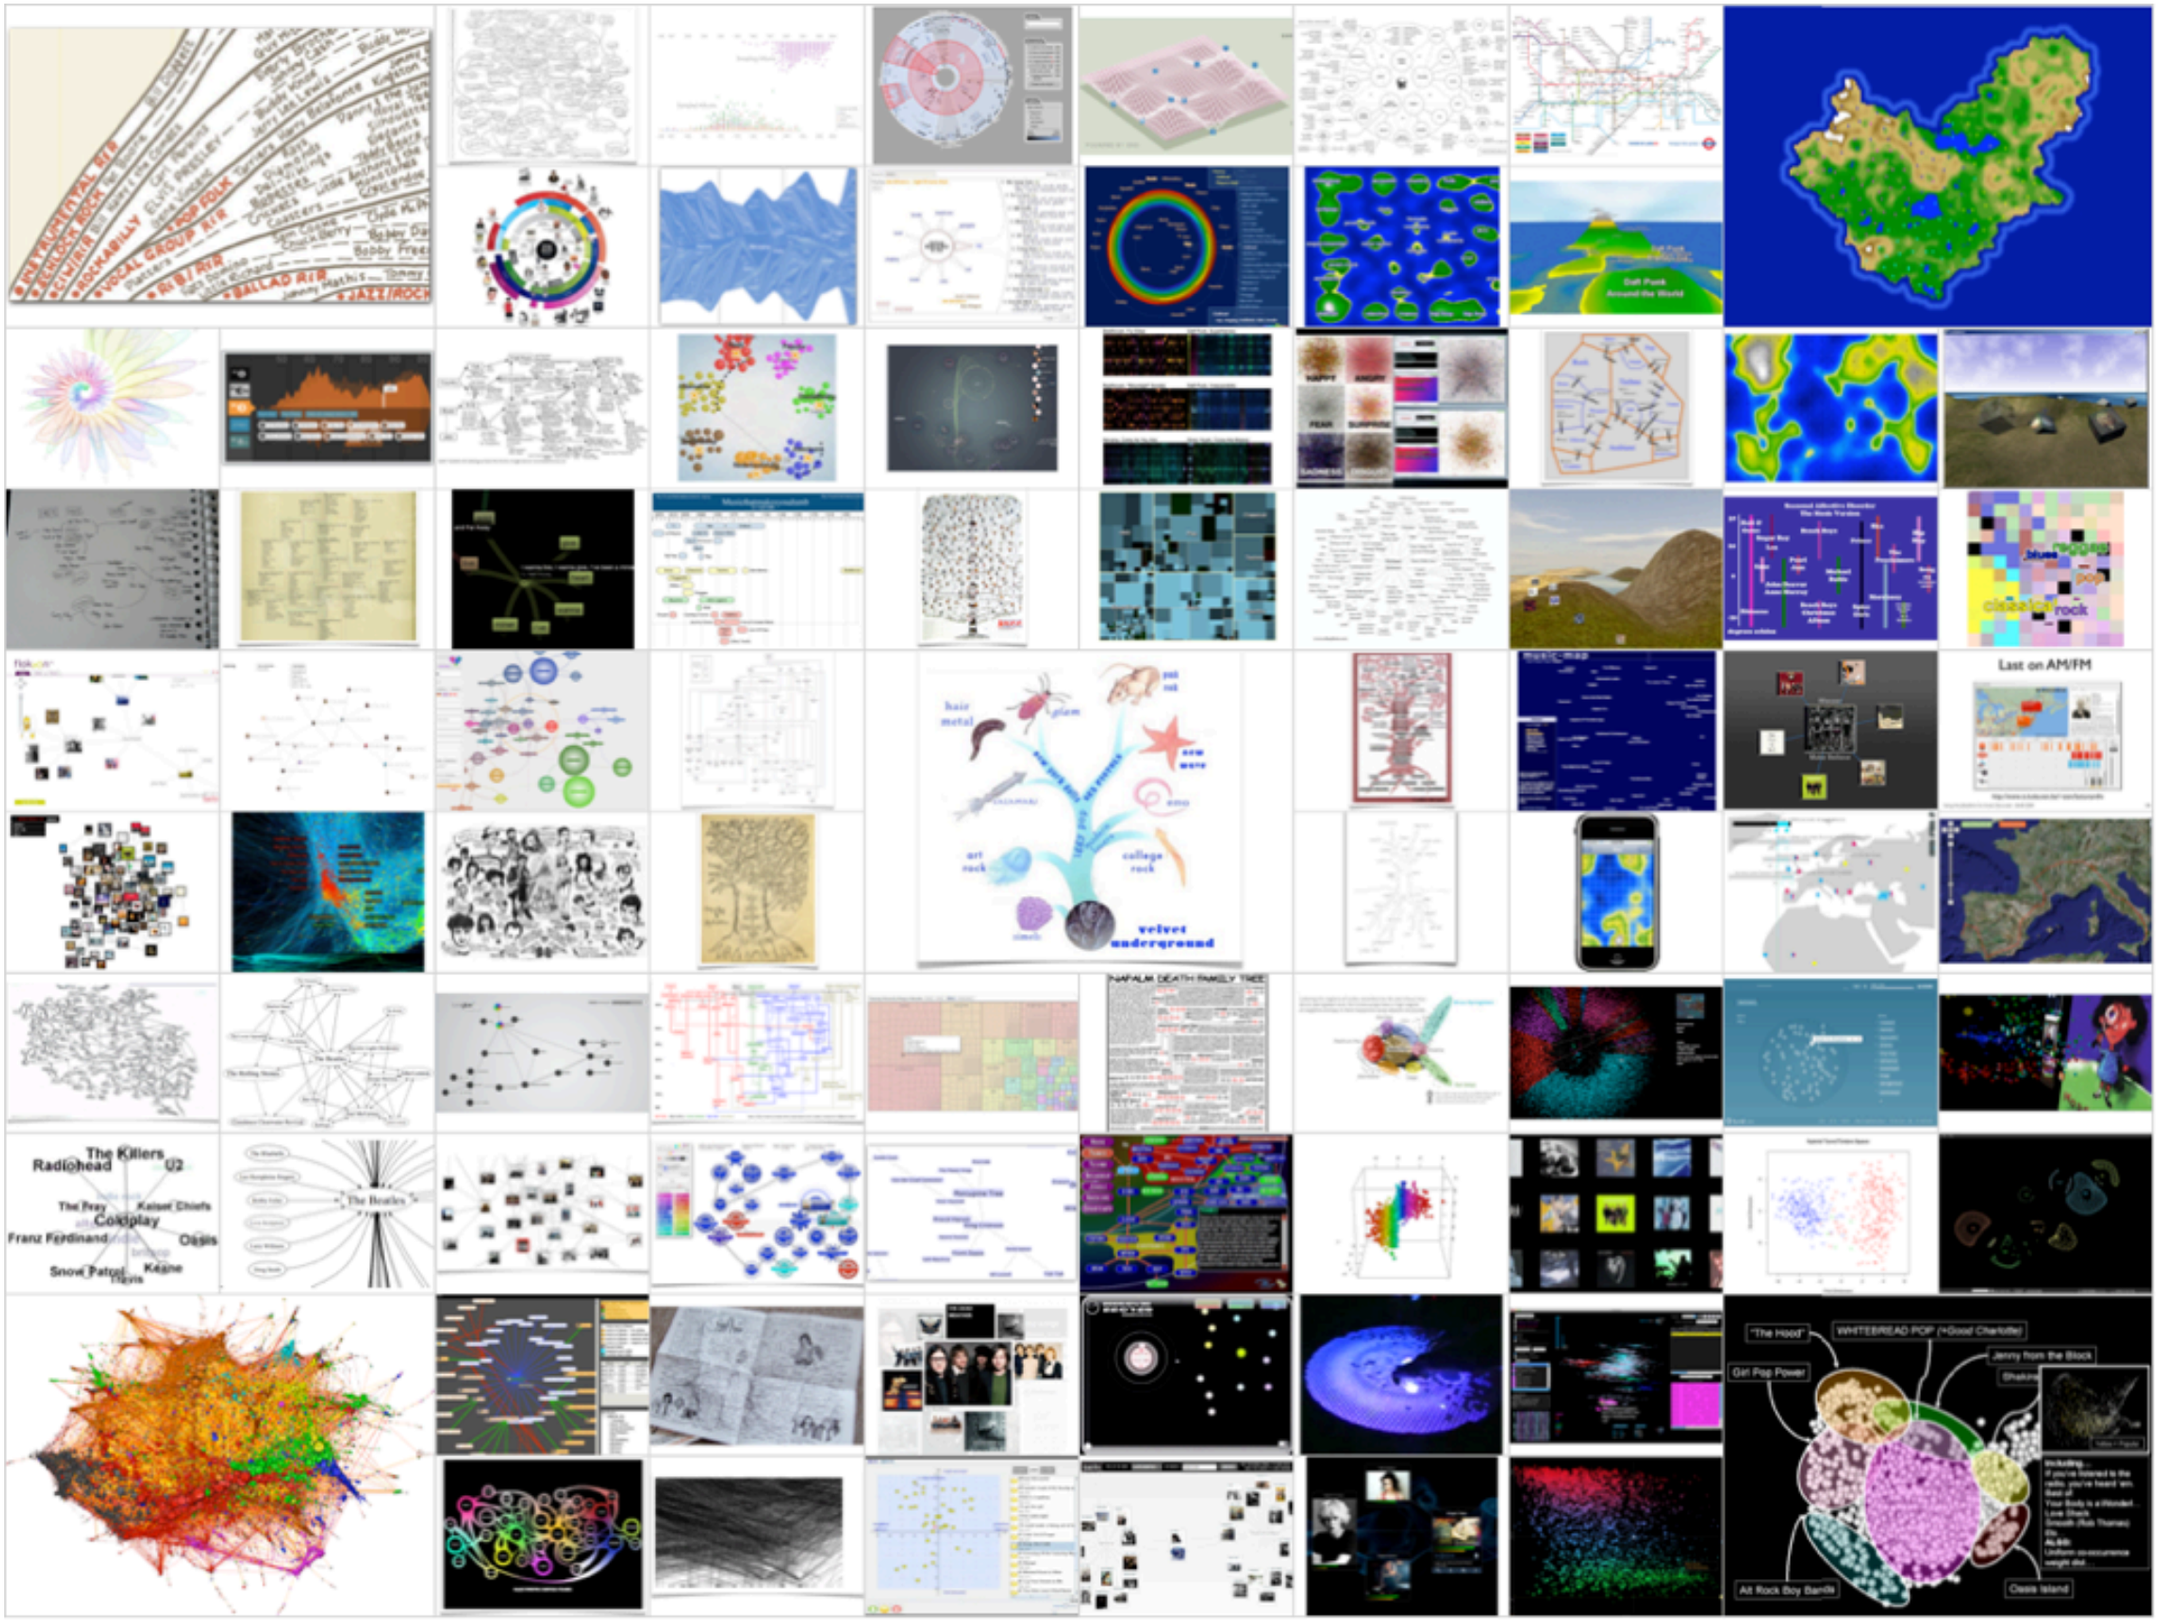

Speakers

Justin Donaldson is a PhD candidate at Indiana University School of Informatics, as well as a regular research intern at Strands, Inc. Justin is interested with the analyses and visualizations of social sources of data, such as those that are generated from playlists, blogs, and bookmarks.

Paul Lamere is the Director of Developer Community at The Echo Nest, a research-focused music intelligence startup that provides music information services to developers and partners through a data mining and machine listening platform. Paul is especially interested in hybrid music recommenders and using visualizations to aid music discovery. Paul also authors 'Music Machinery' a blog focusing on music discovery and recommendation.

Outline of the talk

- Why visualize
- Some history
- Core concepts
- Survey of Music Visualizations
- Tools / Resources
- Conclusions

Music Discovery Challenge

- Millions of songs
- We need tools
  - Recommendation
  - Playlisting
  - **Visualization**

Why Visualize?

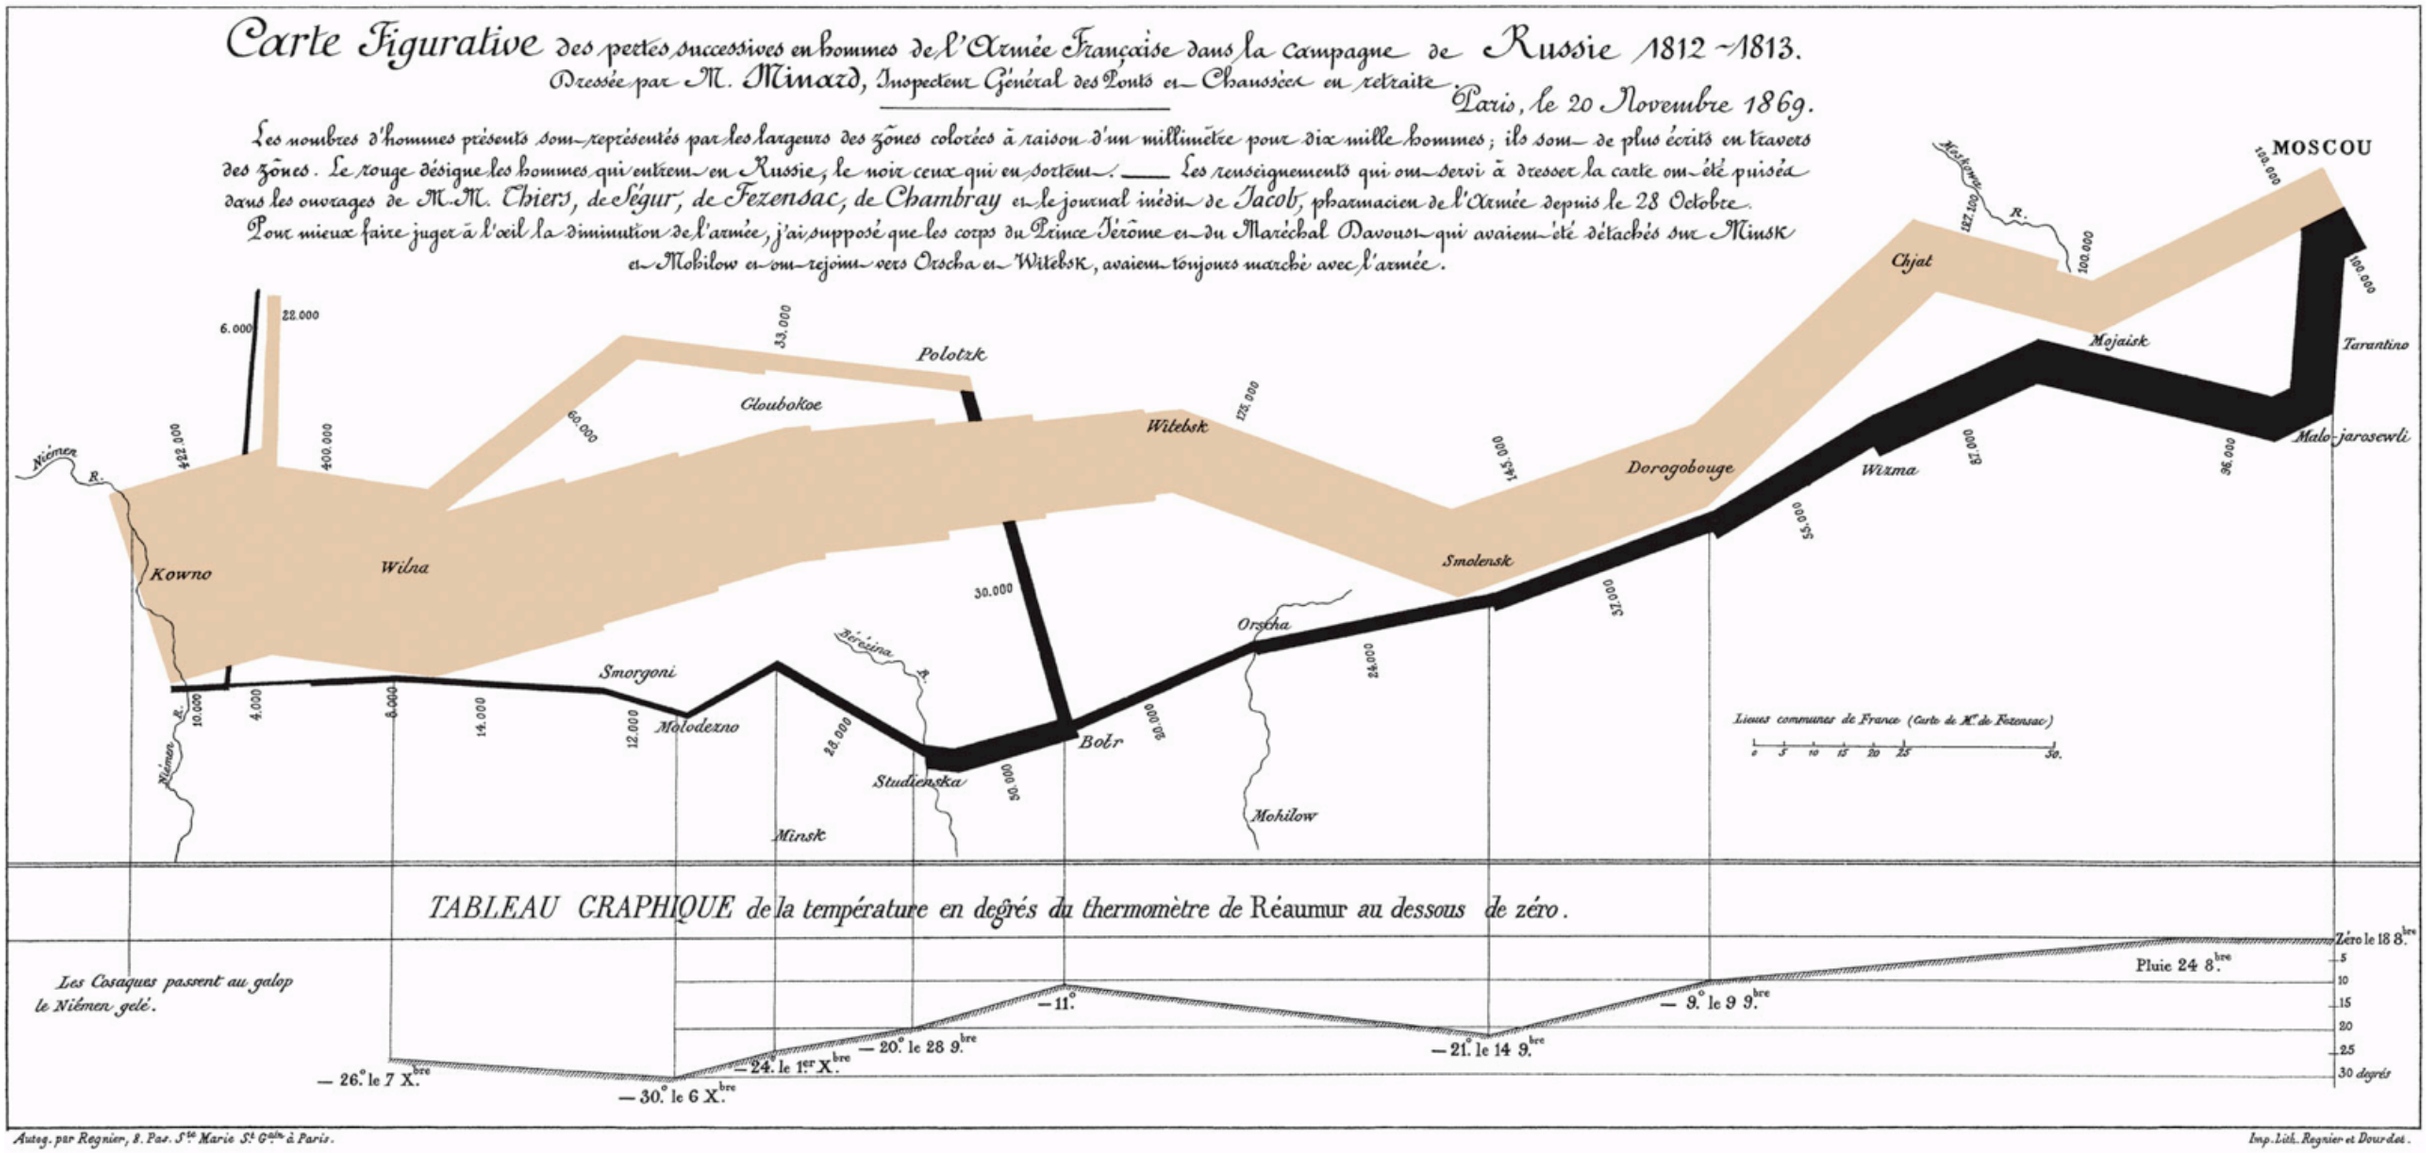

Visualizations can tell a story

Why Visualize?

Visualizations can engage our pattern-matching brain

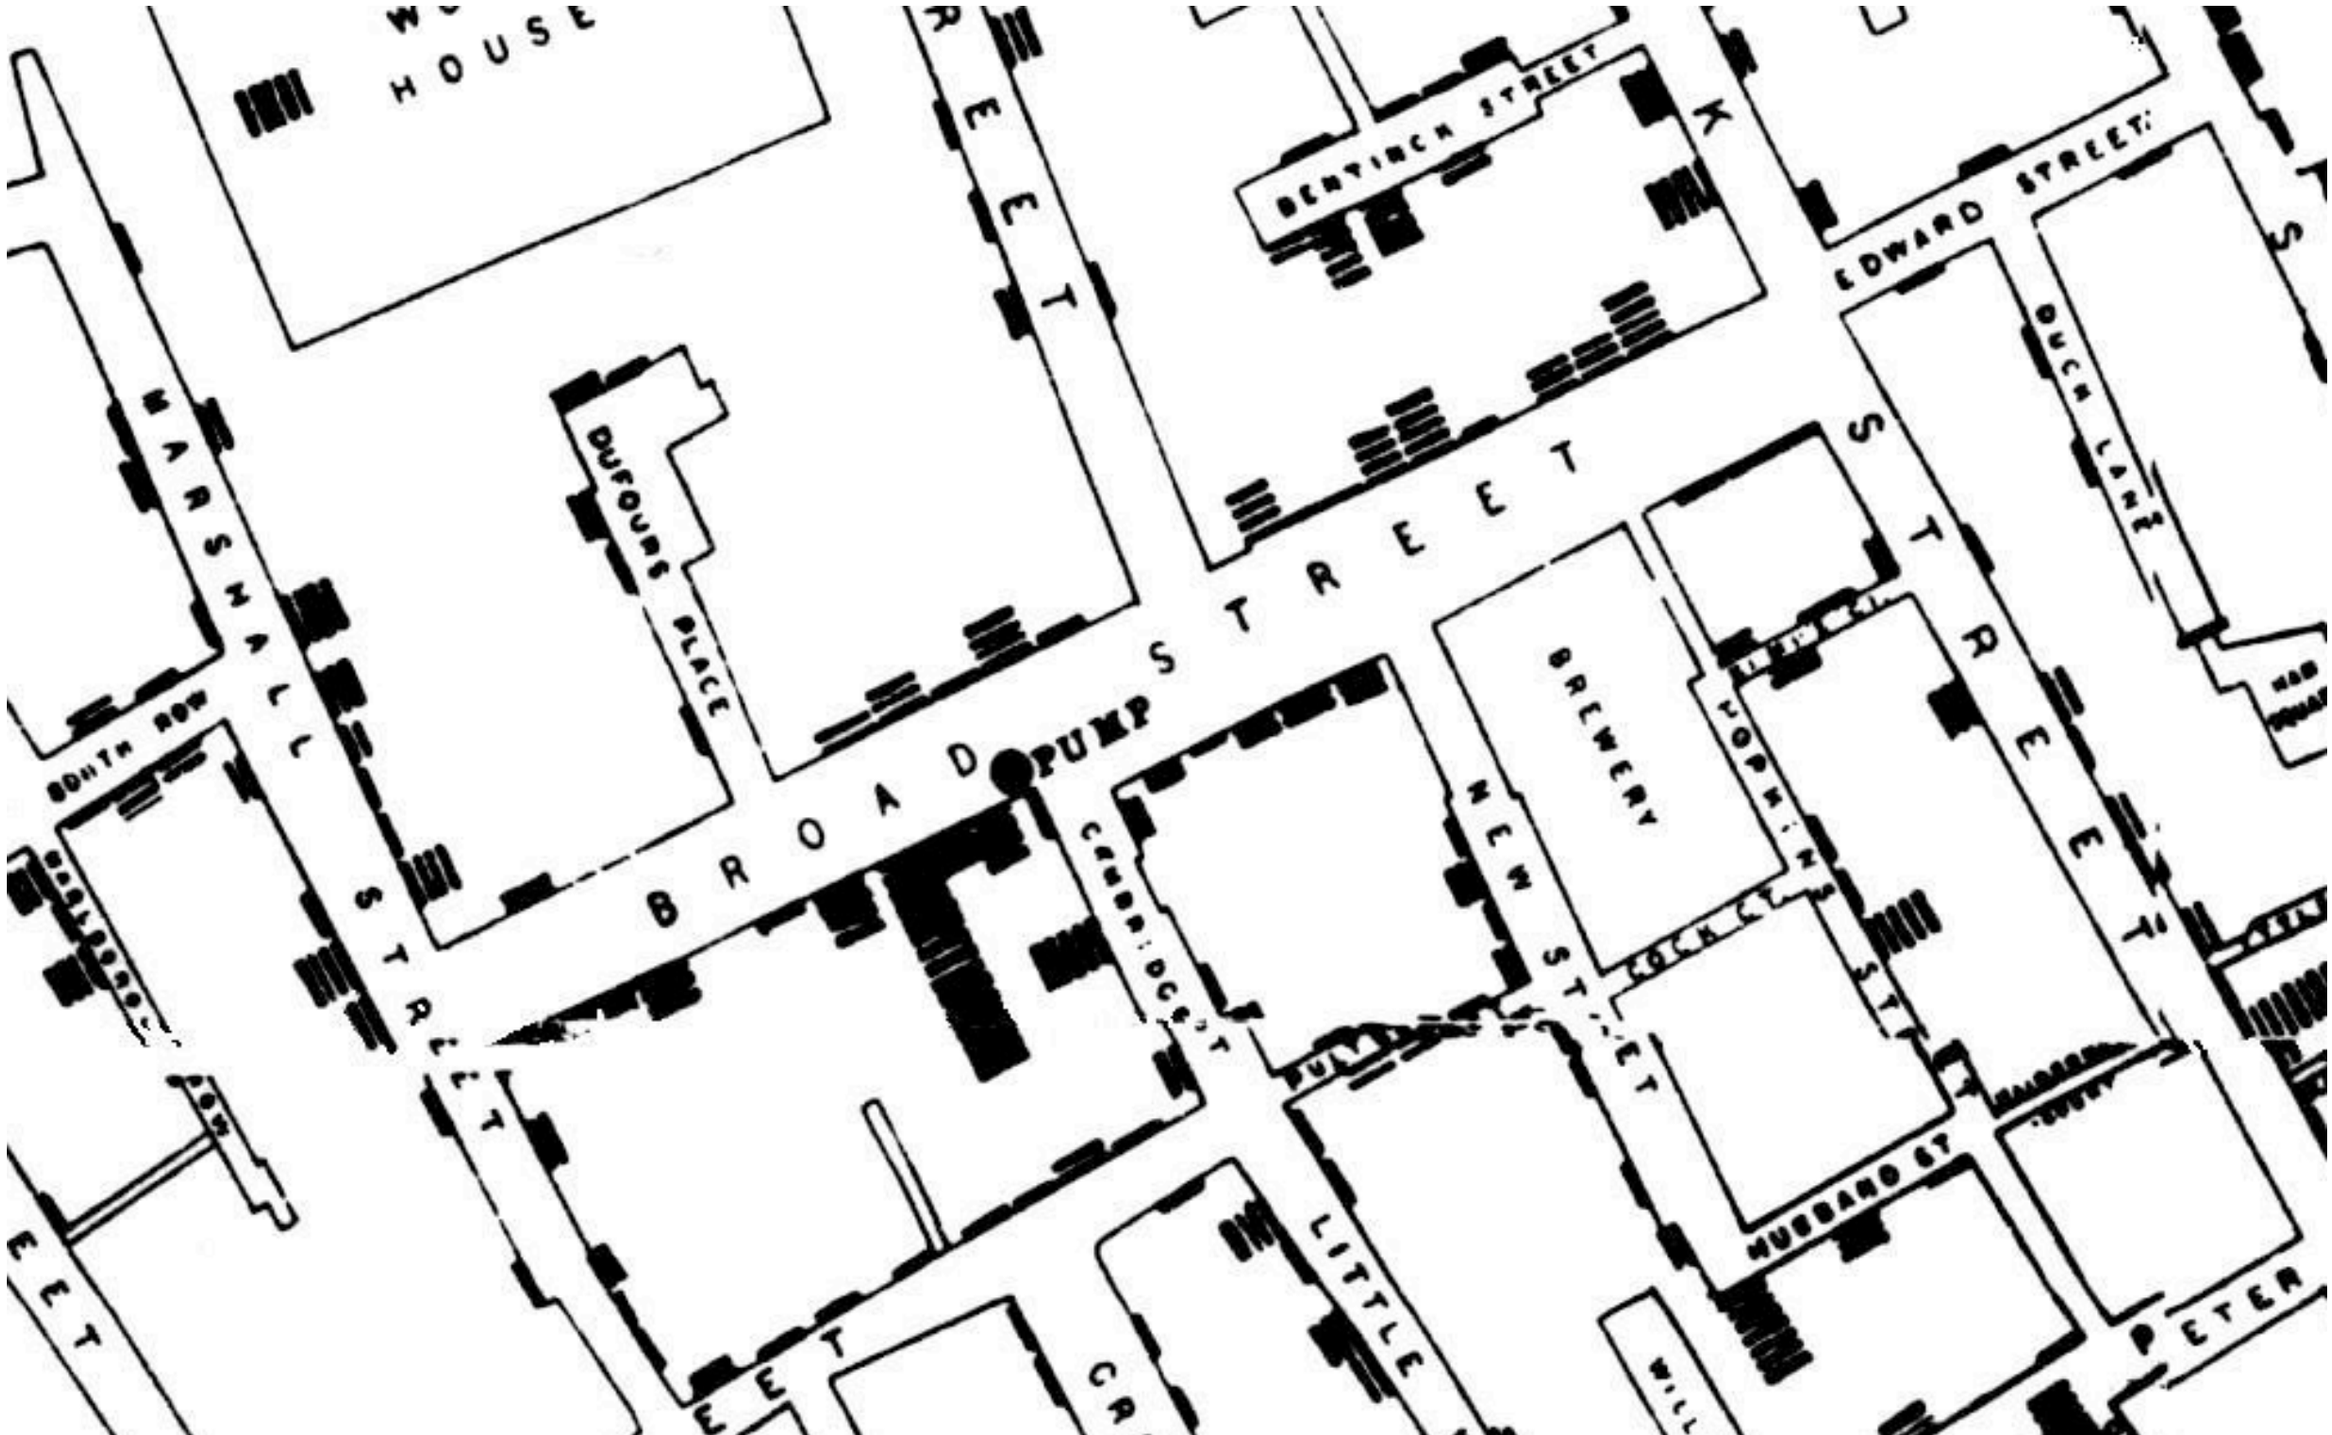

Why Visualize?

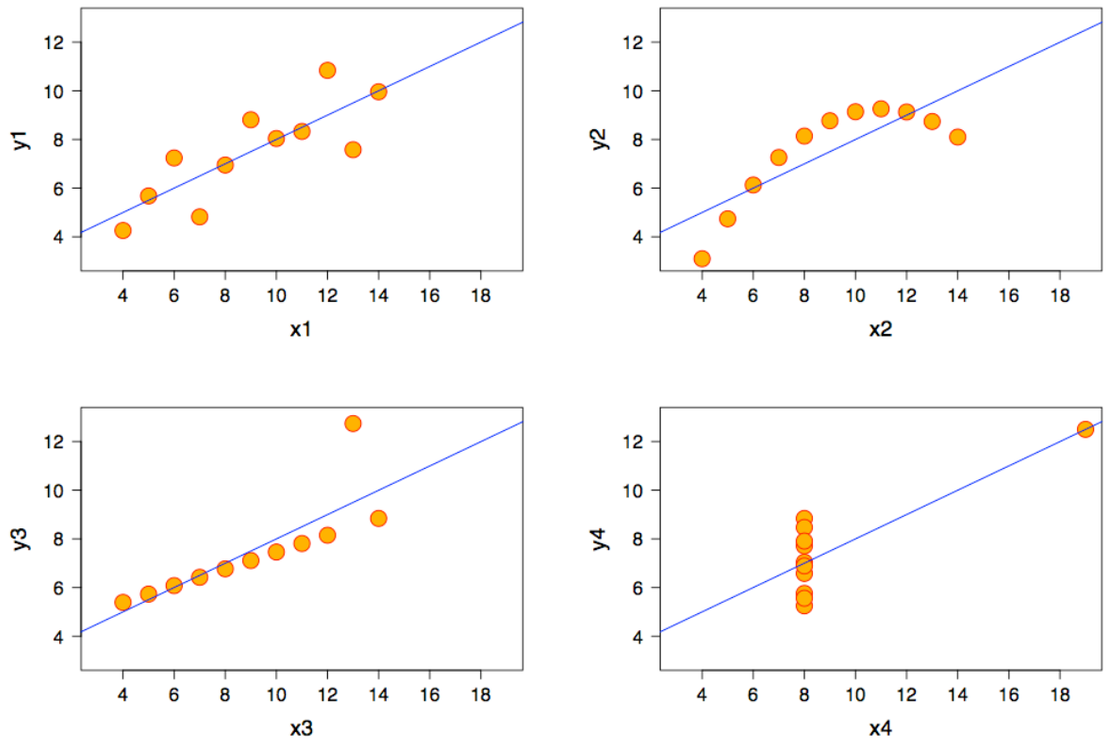

property	value
x mean	9
x var	10
y mean	7.5
y var	3.75
x y cor.	0.816
linear reg.	y = 3 + .5x

- Anscombe's "Quartet"
- Each distribution has equivalent summary statistics
- Means, correlations, variance: Sometimes one number summaries are deceiving.

http://en.wikipedia.org/wiki/Anscombe%27s_quartet

Semantic Pattern Mappings

Graphical Code		Semantics
Small shapes defined by closed contour, texture, color, shaded solid.		Object, idea, entity, node.
Spatially ordered graphical objects.		Related information or a sequence. In a sequence the left-to-right ordering convention borrows from the western convention for written language.
Graphical objects in proximity.		Similar concepts, related information.
Graphical objects having the same shape, color, or texture.		Similar concepts, related information.
Size of graphical object Height of graphical object.		Magnitude, quantity, importance.

Colin Ware - Visual Thinking for Design

Semantic Pattern Mappings

Graphical Code		Semantics
Shapes connected by contour.		Related entities, path between entities.
Thickness of connecting contour.		Strength of relationship.
Color and texture of connecting contour.		Type of relationship.
Shapes enclosed by a contour, or a common texture, or a common color.		Contained entities. Related entities.
Nested regions, partitioned regions.		Hierarchical concepts.
Attached shapes.		Parts of a conceptual structure.

Colin Ware - Visual Thinking for Design

Basic Popout Channels

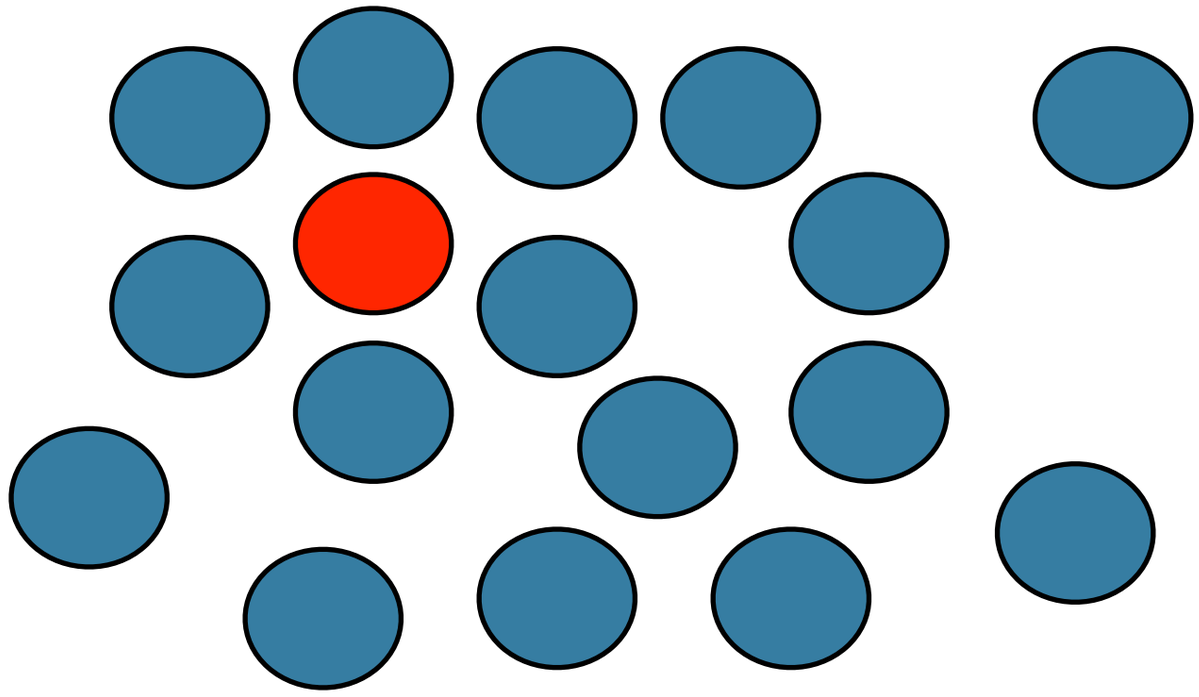

Color

Colin Ware - Visual Thinking for Design

Basic Popout Channels

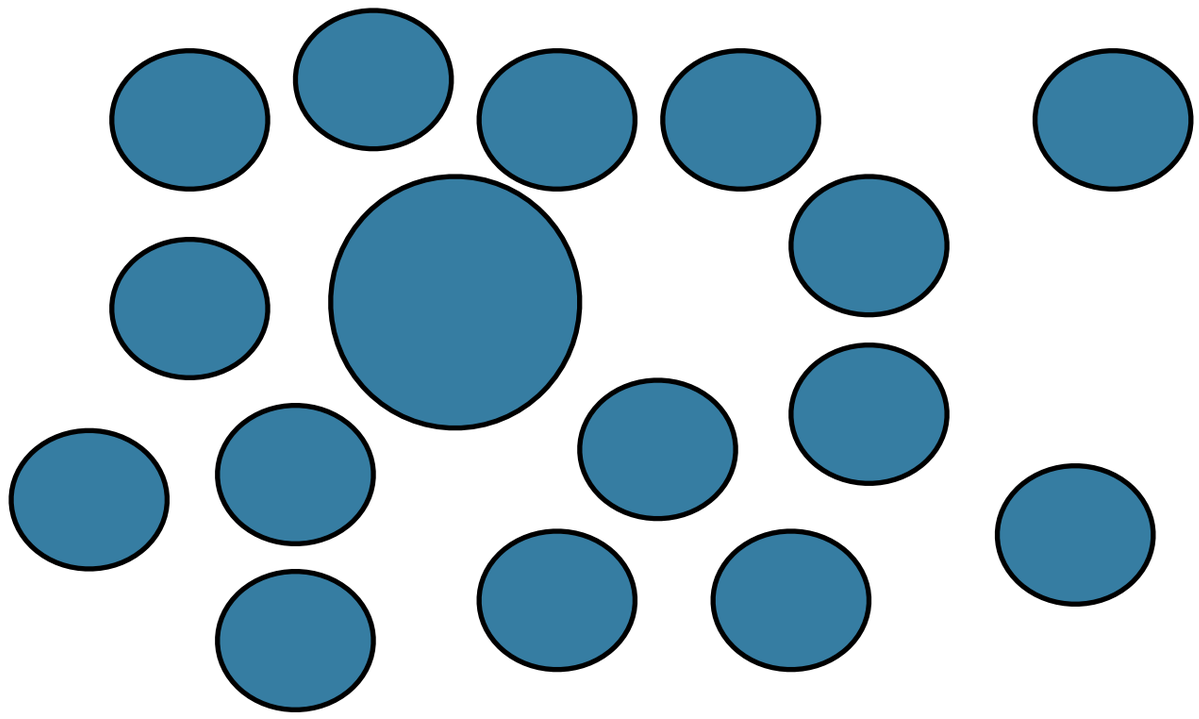

Size

Colin Ware - Visual Thinking for Design

Basic Popout Channels

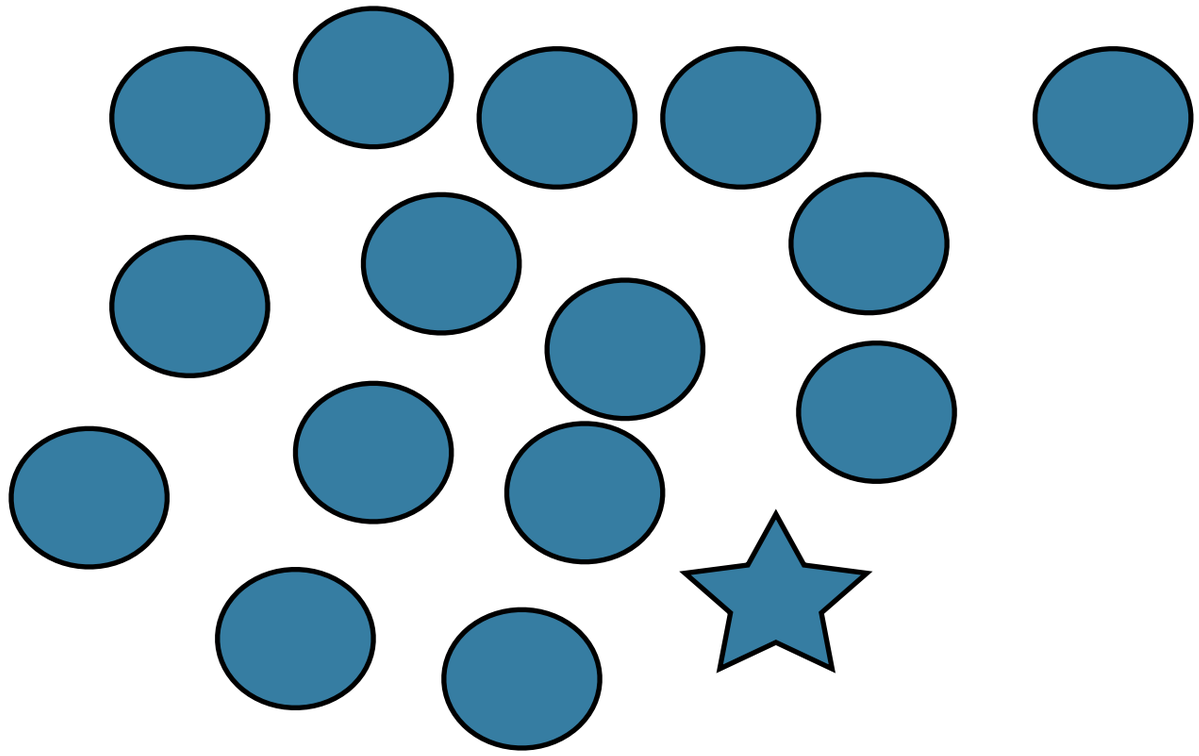

Shape

Colin Ware - Visual Thinking for Design

Basic Popout Channels

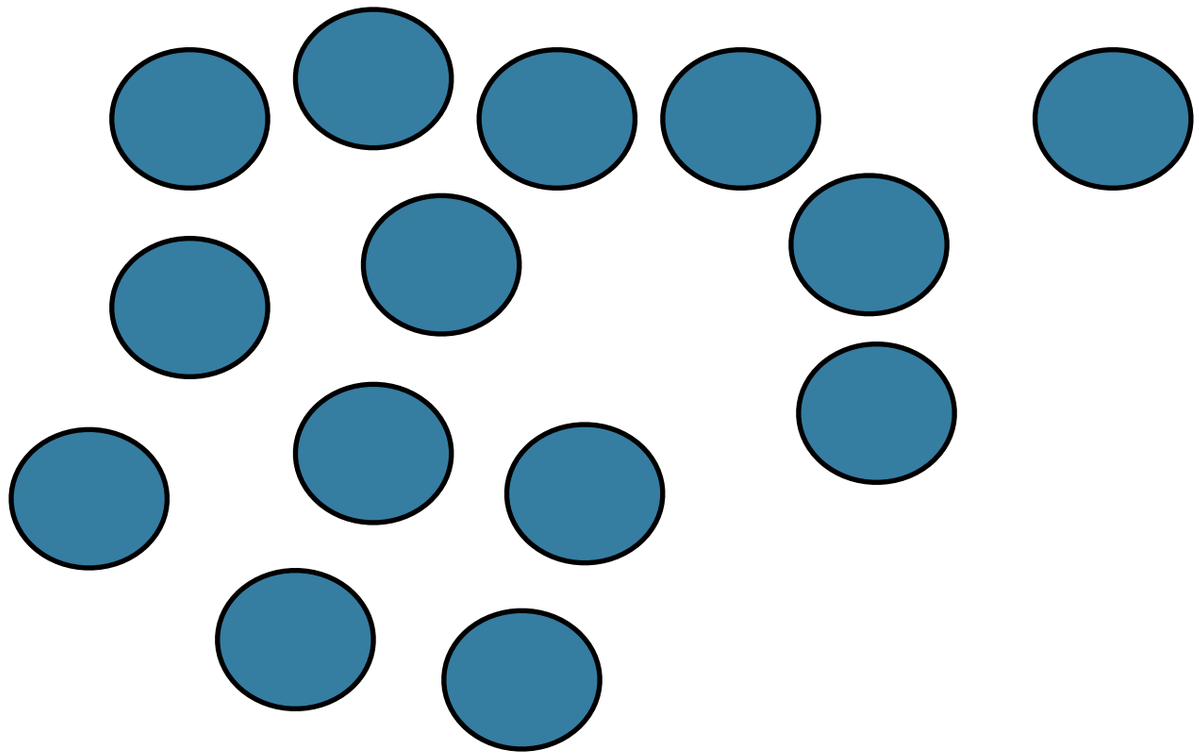

Motion

Colin Ware - Visual Thinking for Design

Basic Popout Channels

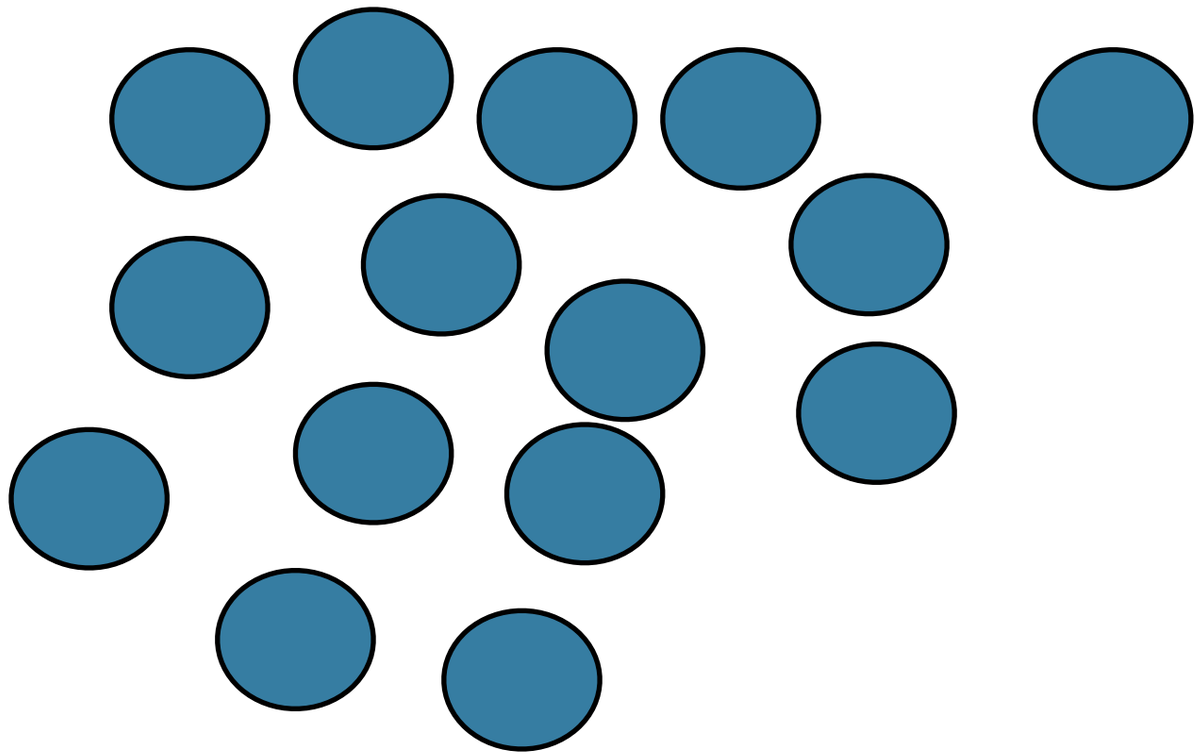

Motion

Colin Ware - Visual Thinking for Design

Basic Popout Channels

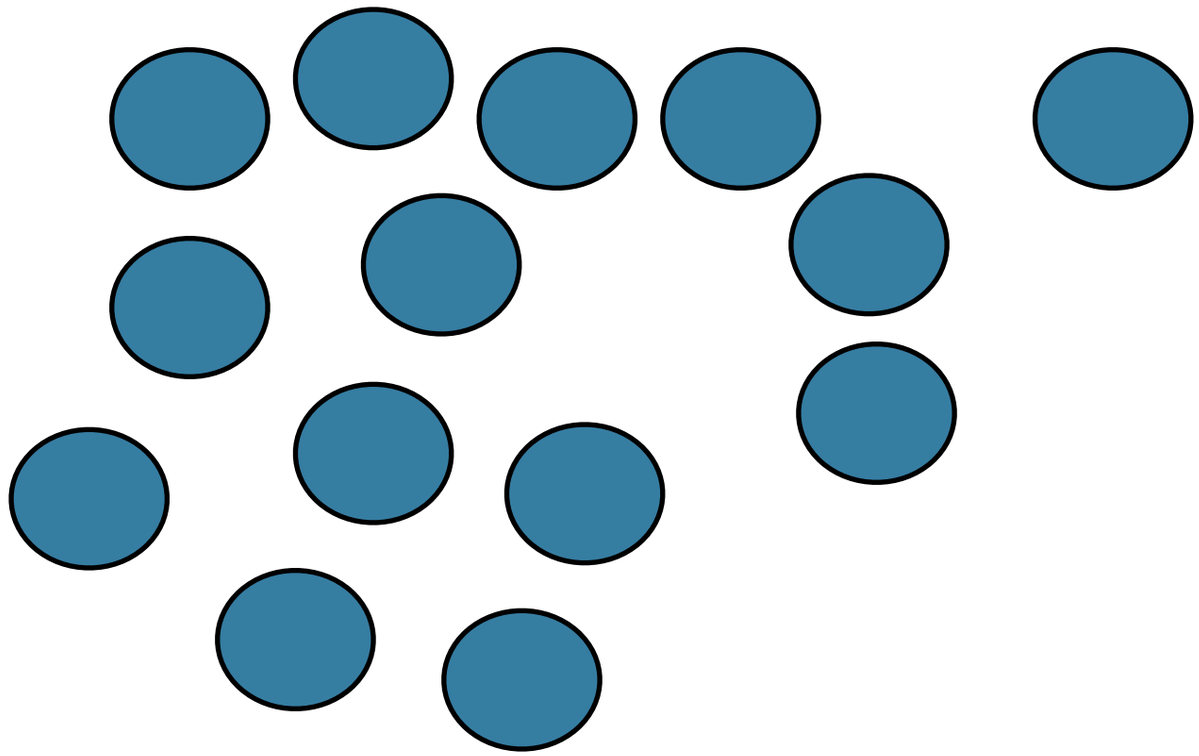

Motion

Colin Ware - Visual Thinking for Design

Basic Popout Channels

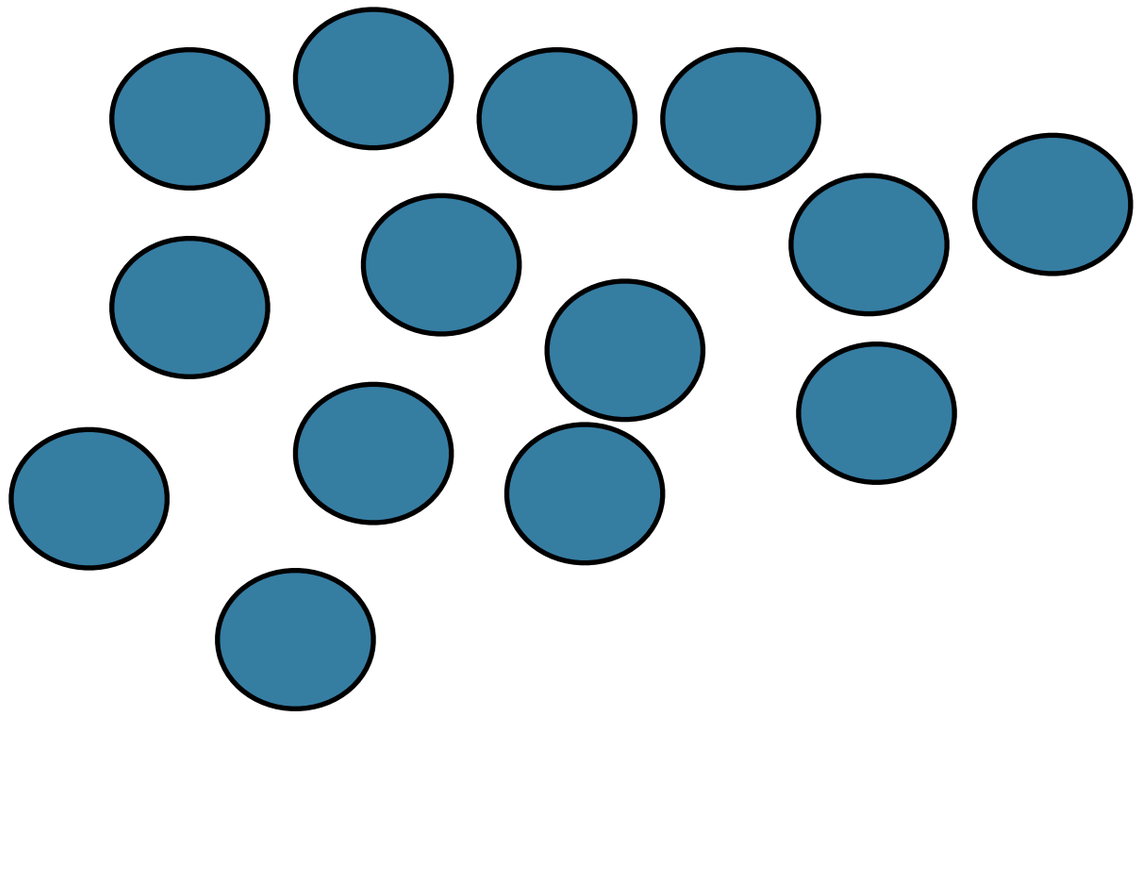

Spatial Grouping

Colin Ware - Visual Thinking for Design

Overloading the channels

Flickr photo by Stéfan

Visualization of Visualizations

History/Geneaology

Cuban Music Genealogy

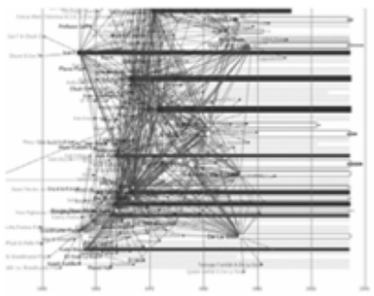

musician map

Garofalos' genealogy

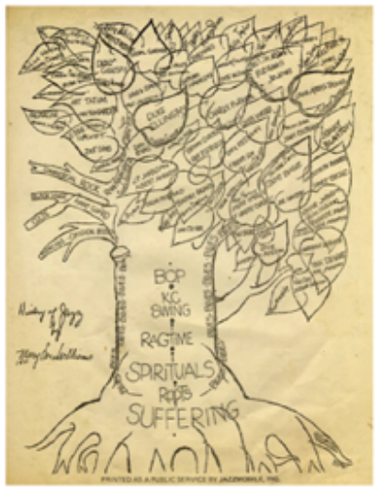

History of Jazz

Michael Jackson Samples

ishkur's guide

Authoritative/Researcher/Historian

Casual/Consumer

(Not many here...)

(dozens more here)

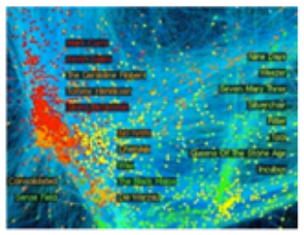

World of Music

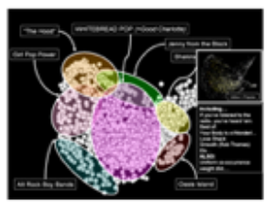

ZMDS

SITM

ESOM MusicMiner

London Tube Map

Last.fm Artist Map

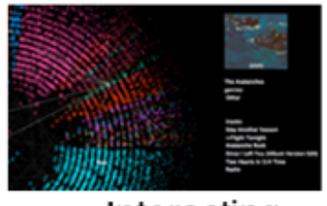

Interacting...

Acoustic/Social Similarity

Visualization Name

Types of visualizations

Tree

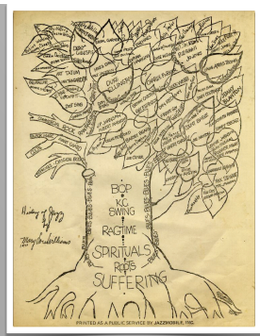

Map

Graph

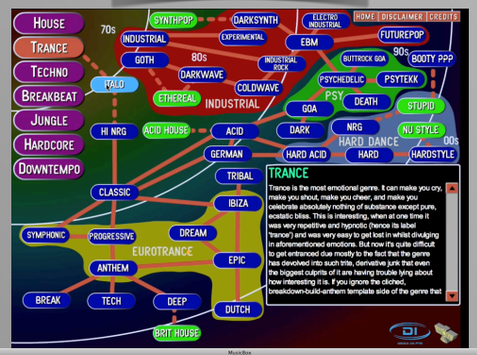

Time Series

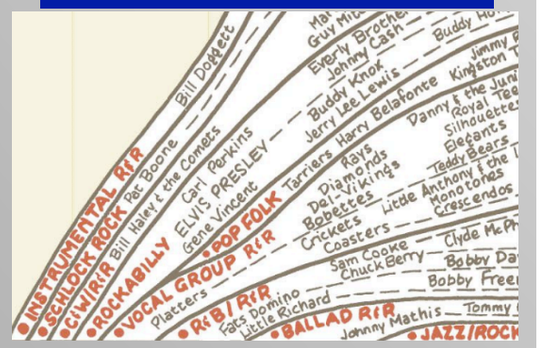

Scatter

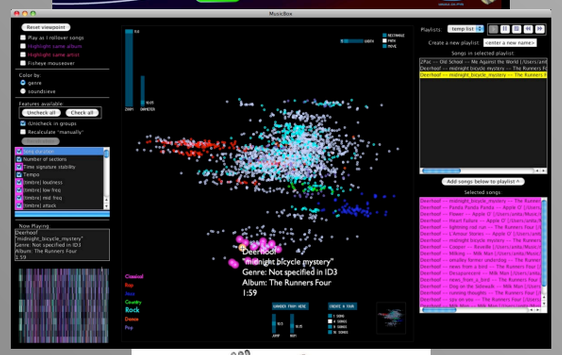

Hybrid

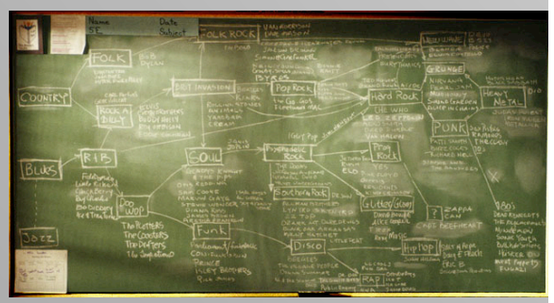

Connections

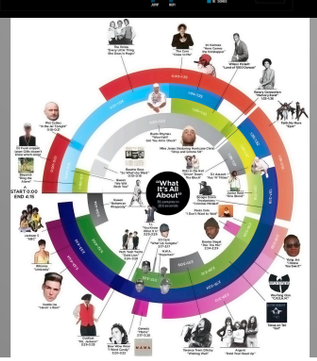

Interactive

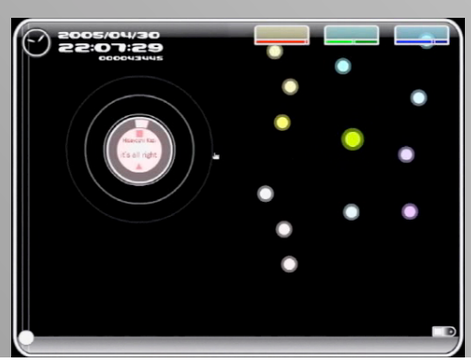

Objects to be discovered

	human	machine
Artists	23	37
Albums	0	6
Tracks	3	17
Samples	3	0
Playlists	0	0
Ordered	0	9
Concerts	0	0
Blogs	0	0
Music	0	0

	human	machine
Genre	18	14
Users	0	4
Radio	0	2
DJs	0	1
Games	0	0
Commu	0	0
Venues	0	0
Taste	2	2
Labels	1	0

26 Human rendered 55 Machine rendered visualizations

Features Used

Content	Beat, tempo, timbre, harmonic content, lyrics
Cultural	web and playlist co-occurrence, purchase/play history, social tags, genre, text aura, expert opinion, appearance, popularity
Relational	member of, collaborated on, supporting musician for, performed as, parent, sibling, married, same label, produced
Mood	happy, sad, angry, relaxed, autumnal
Location	city, country, lat/long,
Time	the hits from the 60s, 70s 80s and today, recent play history
Misc	IQ, Temperature, Name morphology

More features

A scene from High Fidelity

More features

A scene from High Fidelity

Human rendered visualizations

Human rendered visualizations
Some common techniques

- The Flow
- The Map
- The Tree
- The Graph

Human rendered visualizations
Some common techniques

- History
- Influences
- Similarity
- Popularity

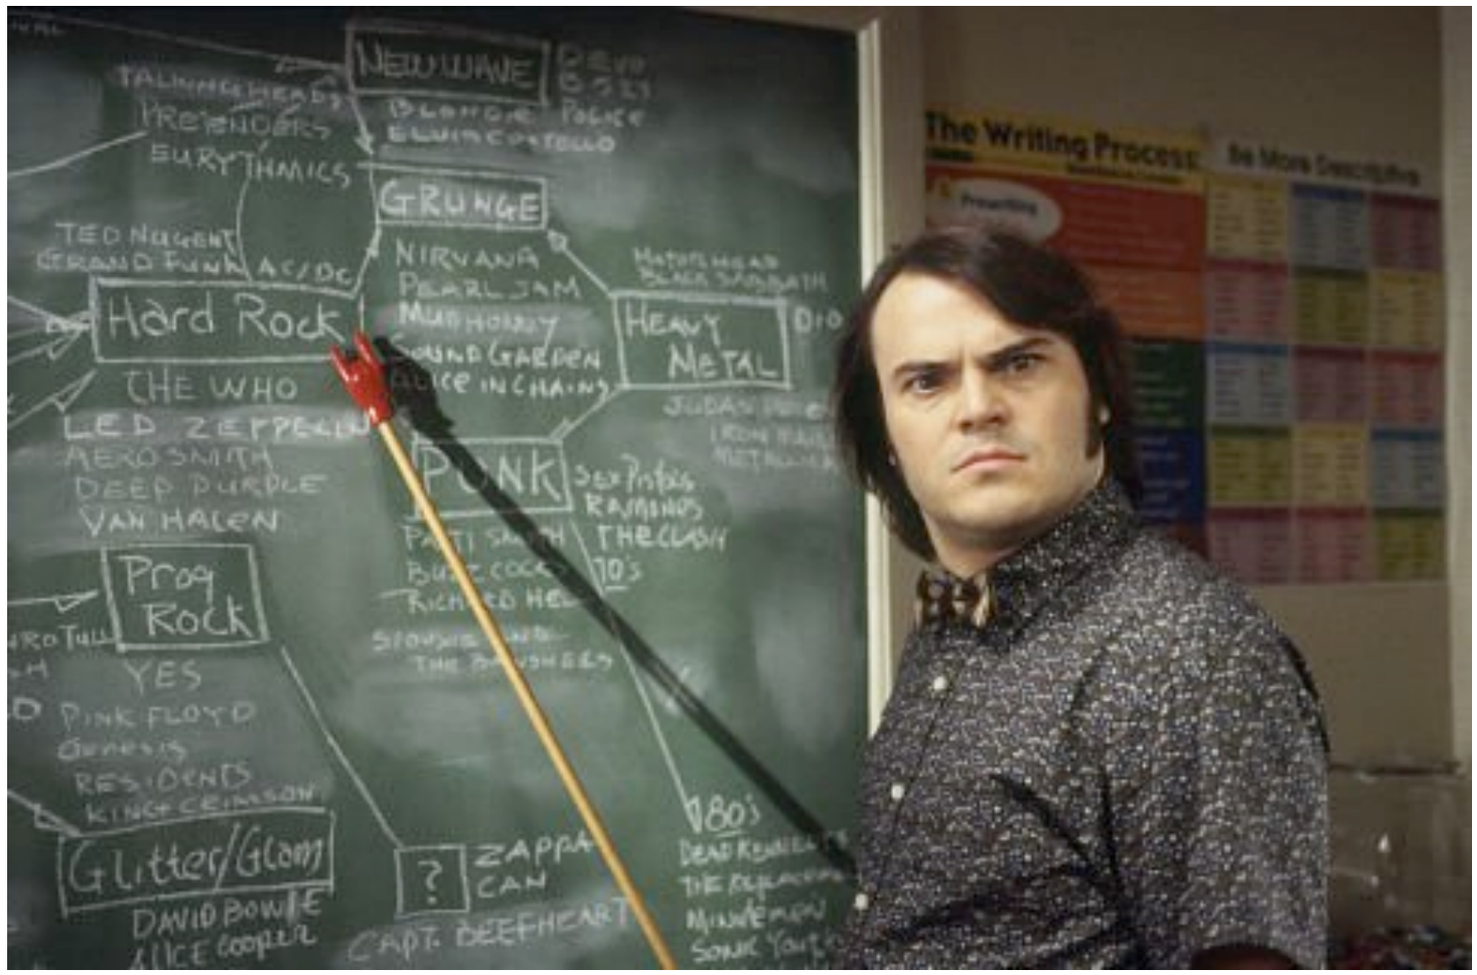

Human rendered visualizations
Some common techniques

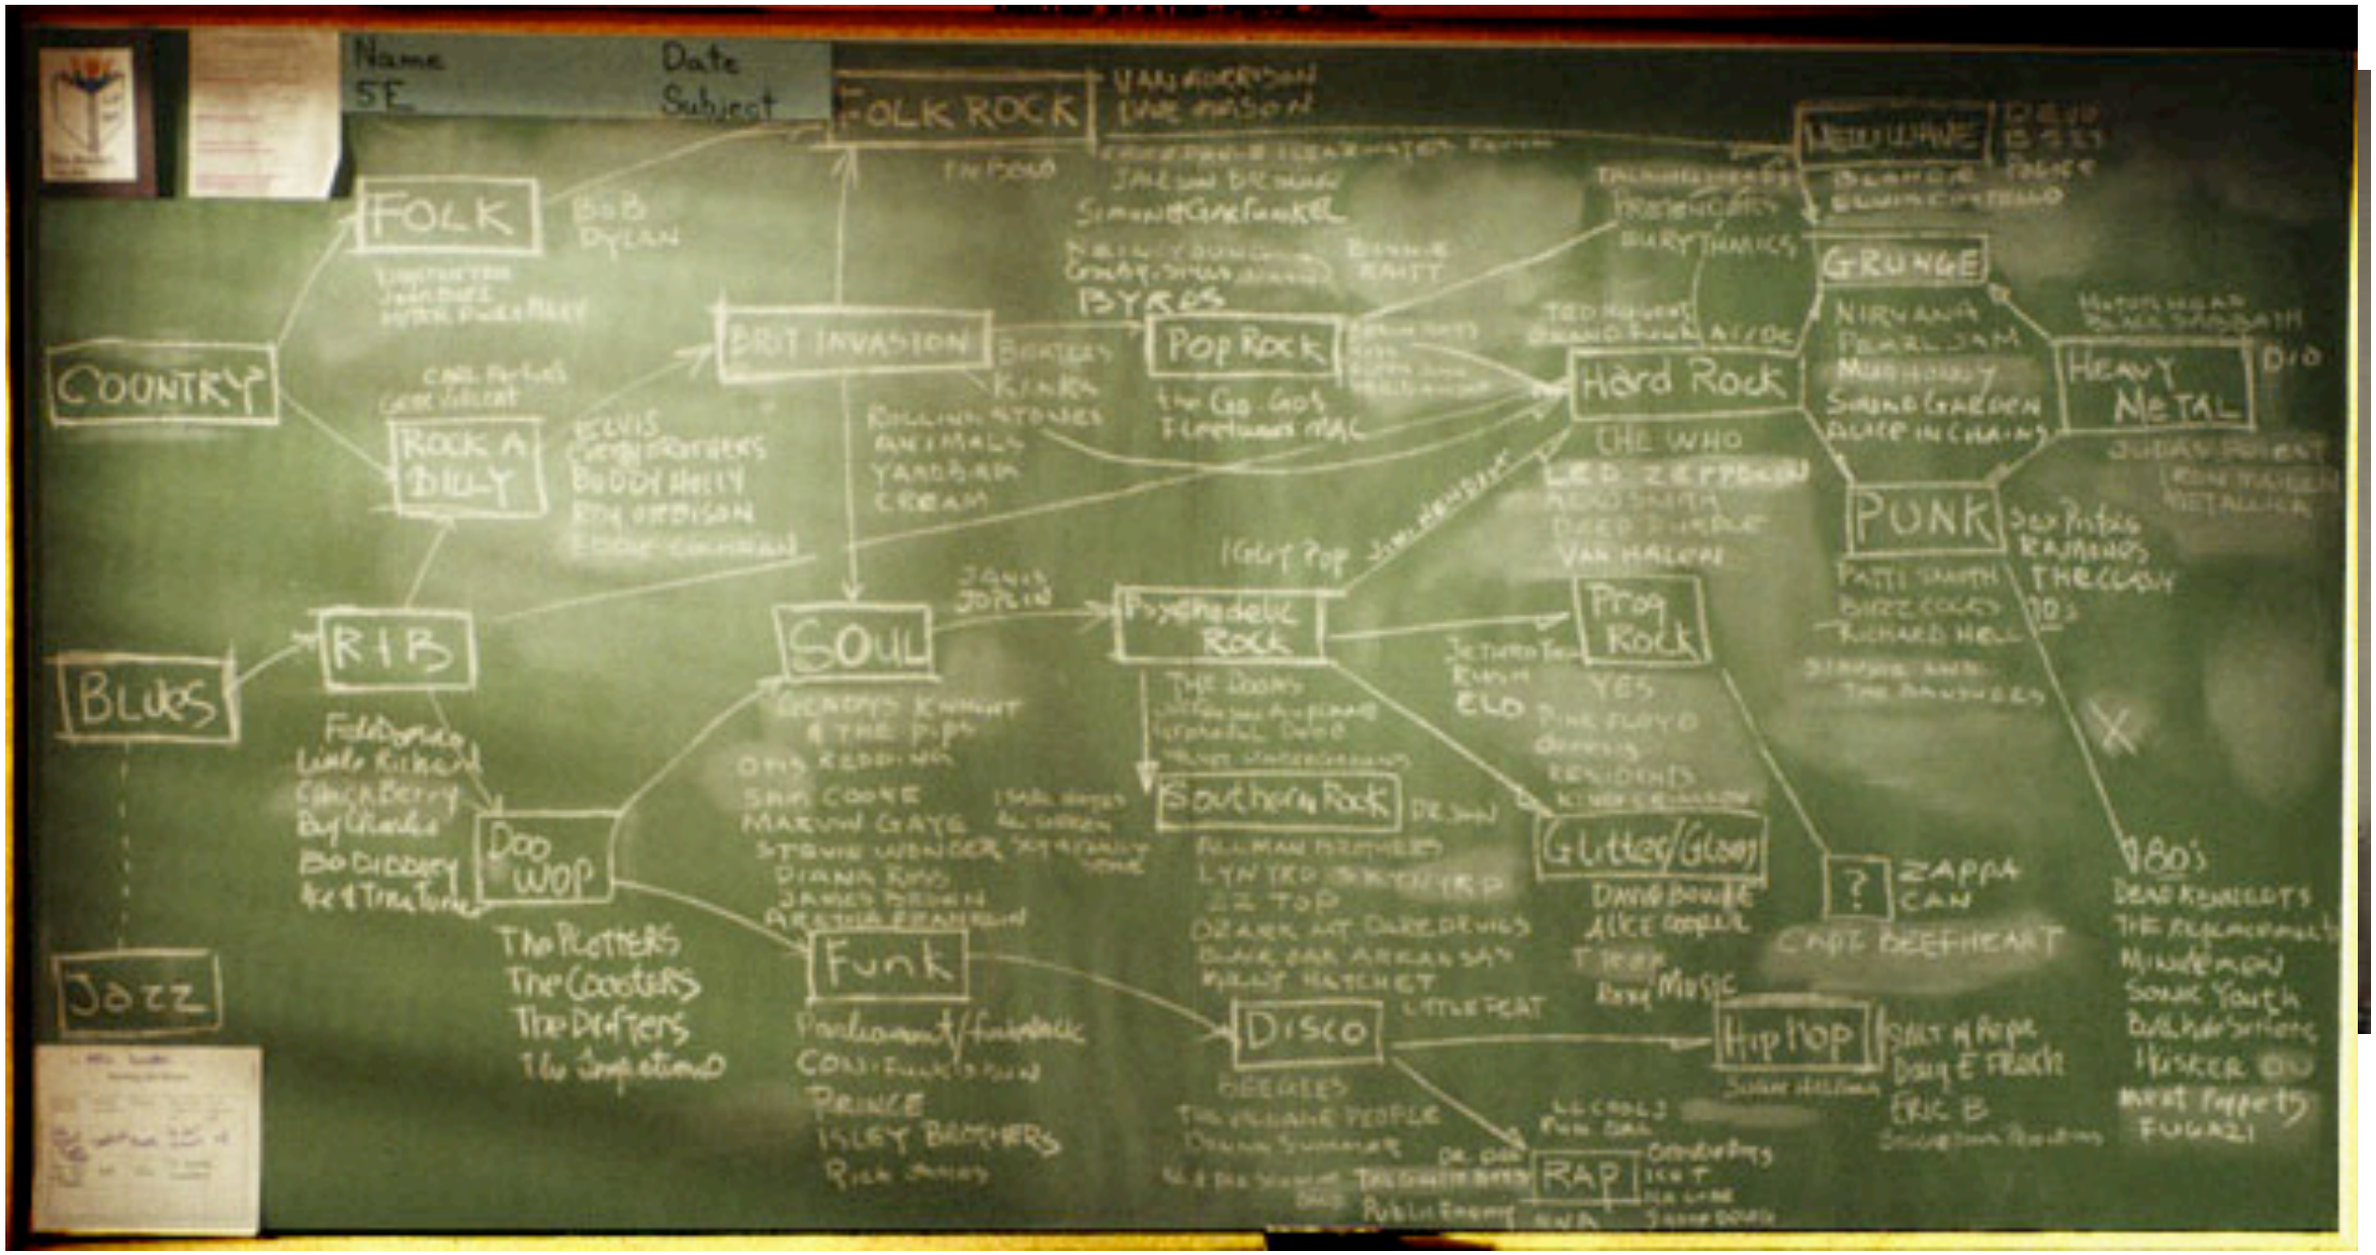

Human rendered visualizations
Some common techniques

Genealogy of Pop and Rock Music

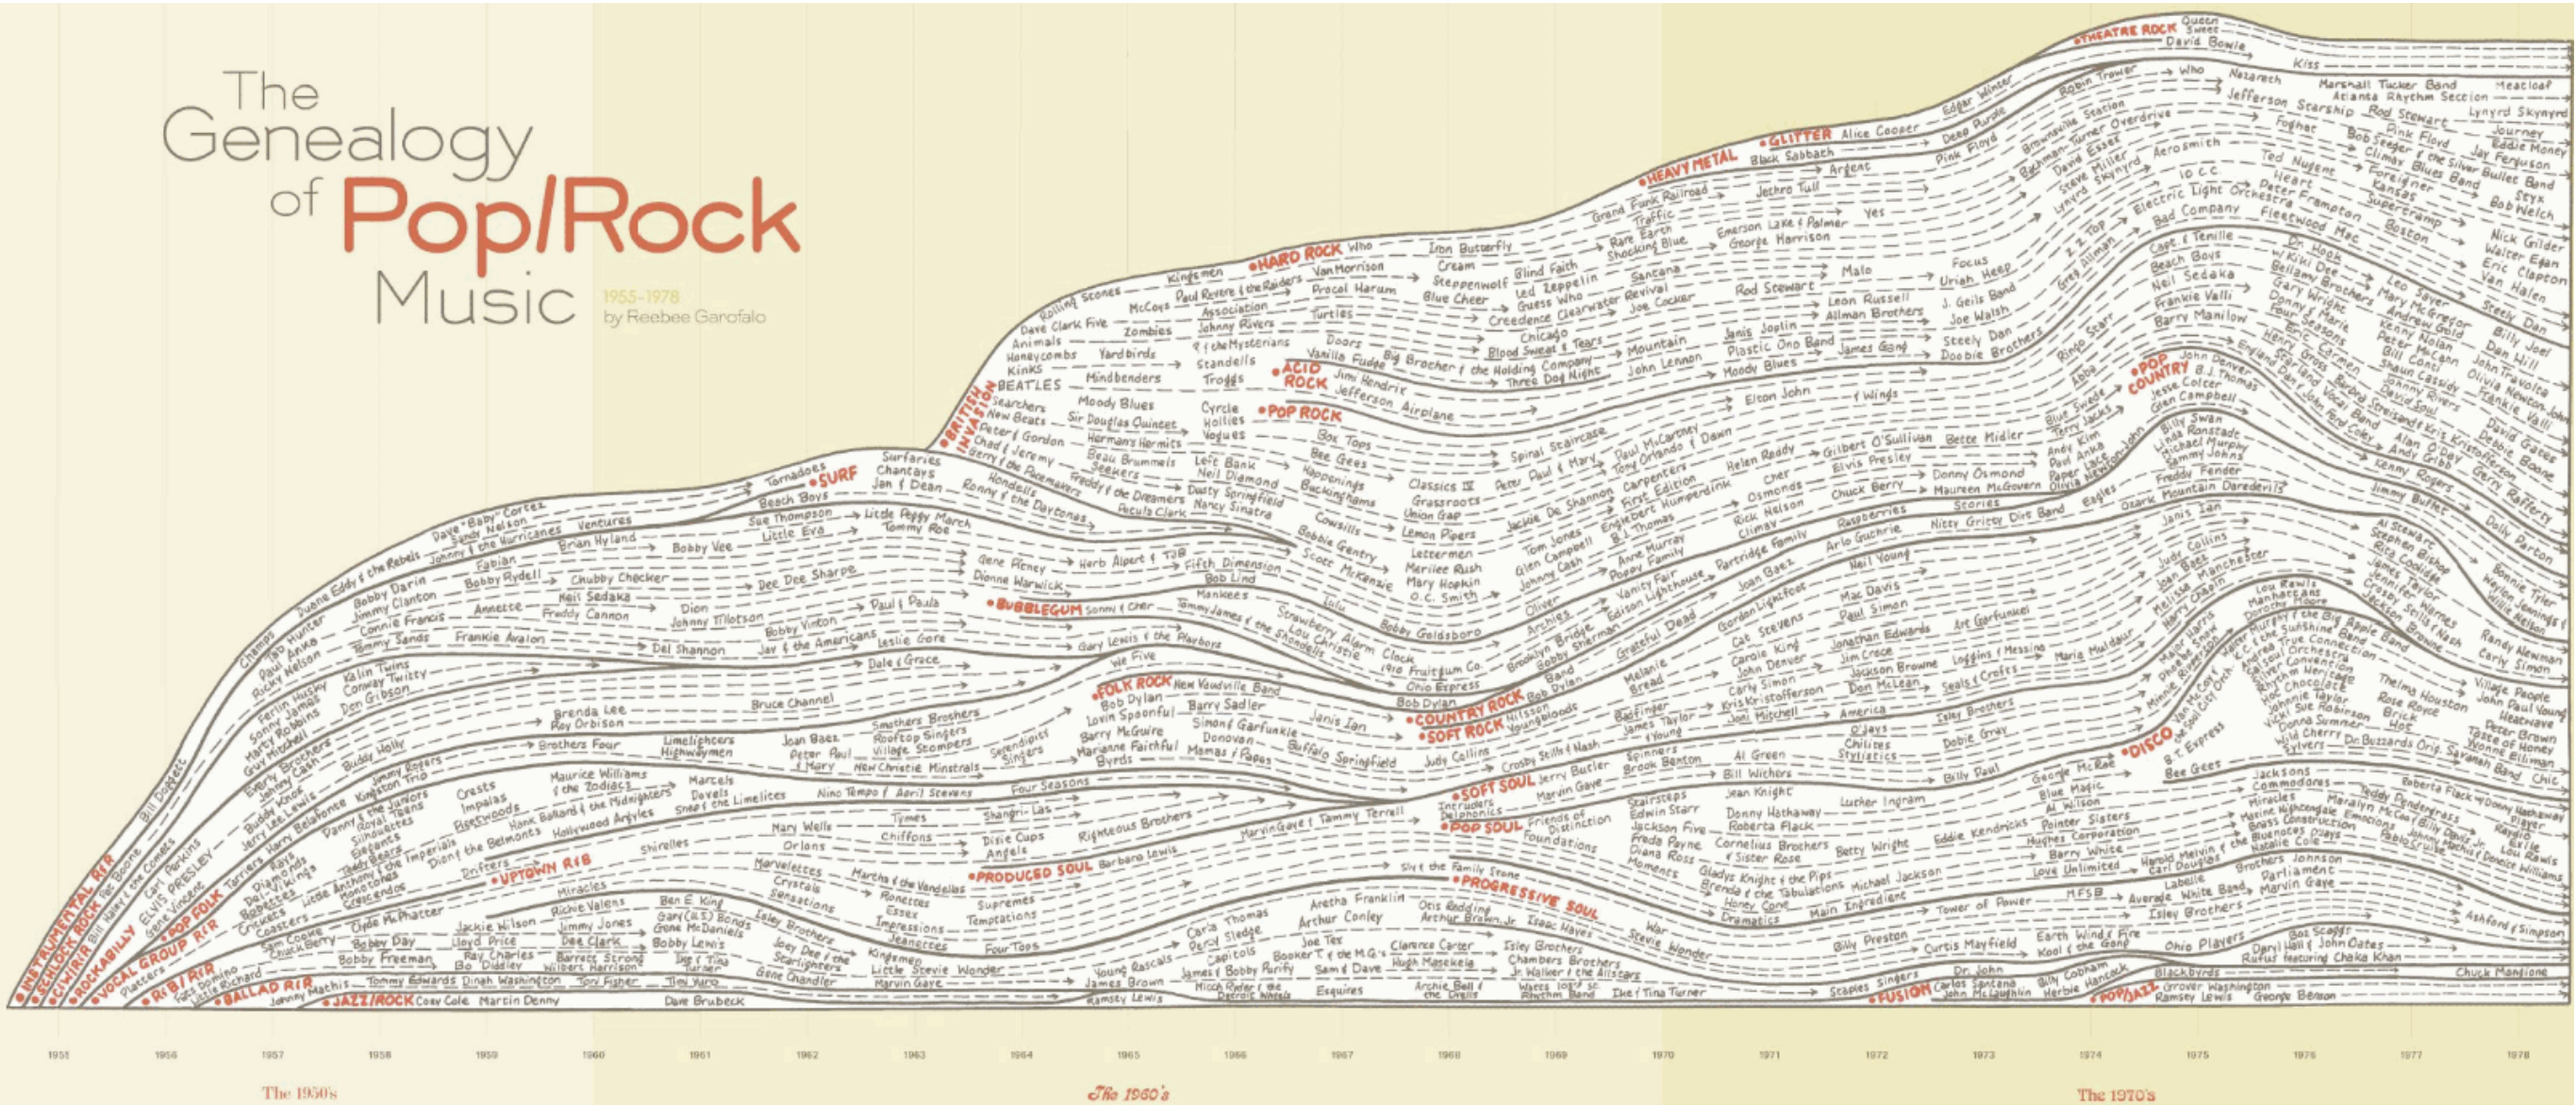

Genealogy of Pop and Rock Music - Reebee Garofalo

Genealogy of Pop and Rock Music

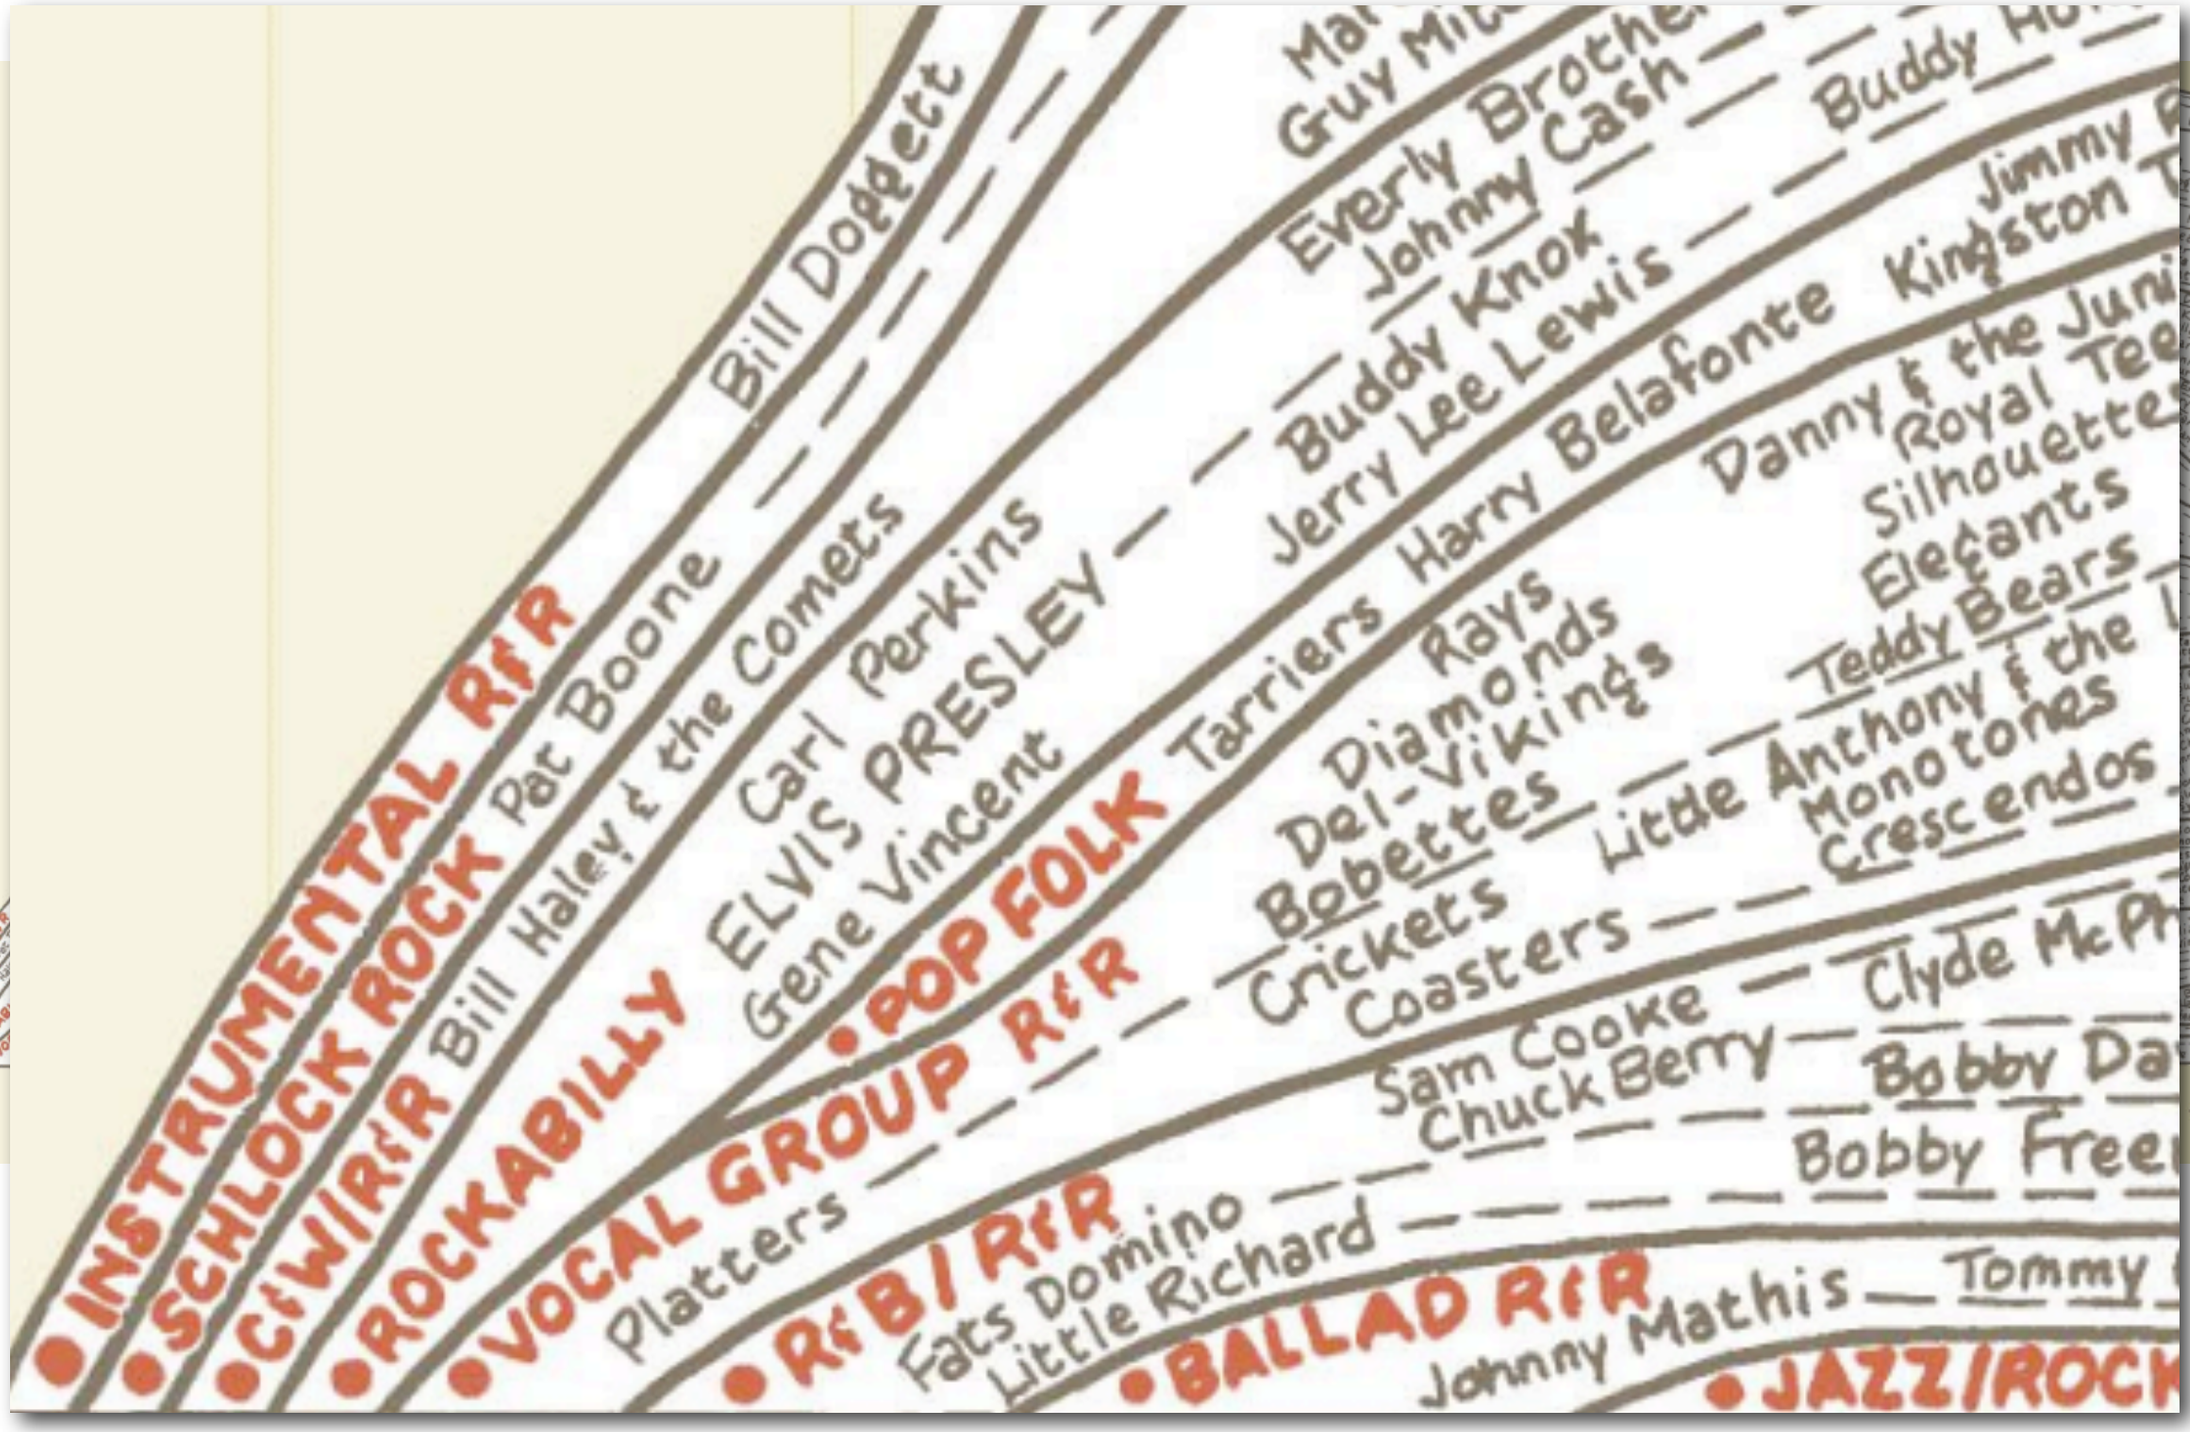

Genealogy of Pop and Rock Music - Reebee Garofalo

Genealogy of Pop and Rock Music

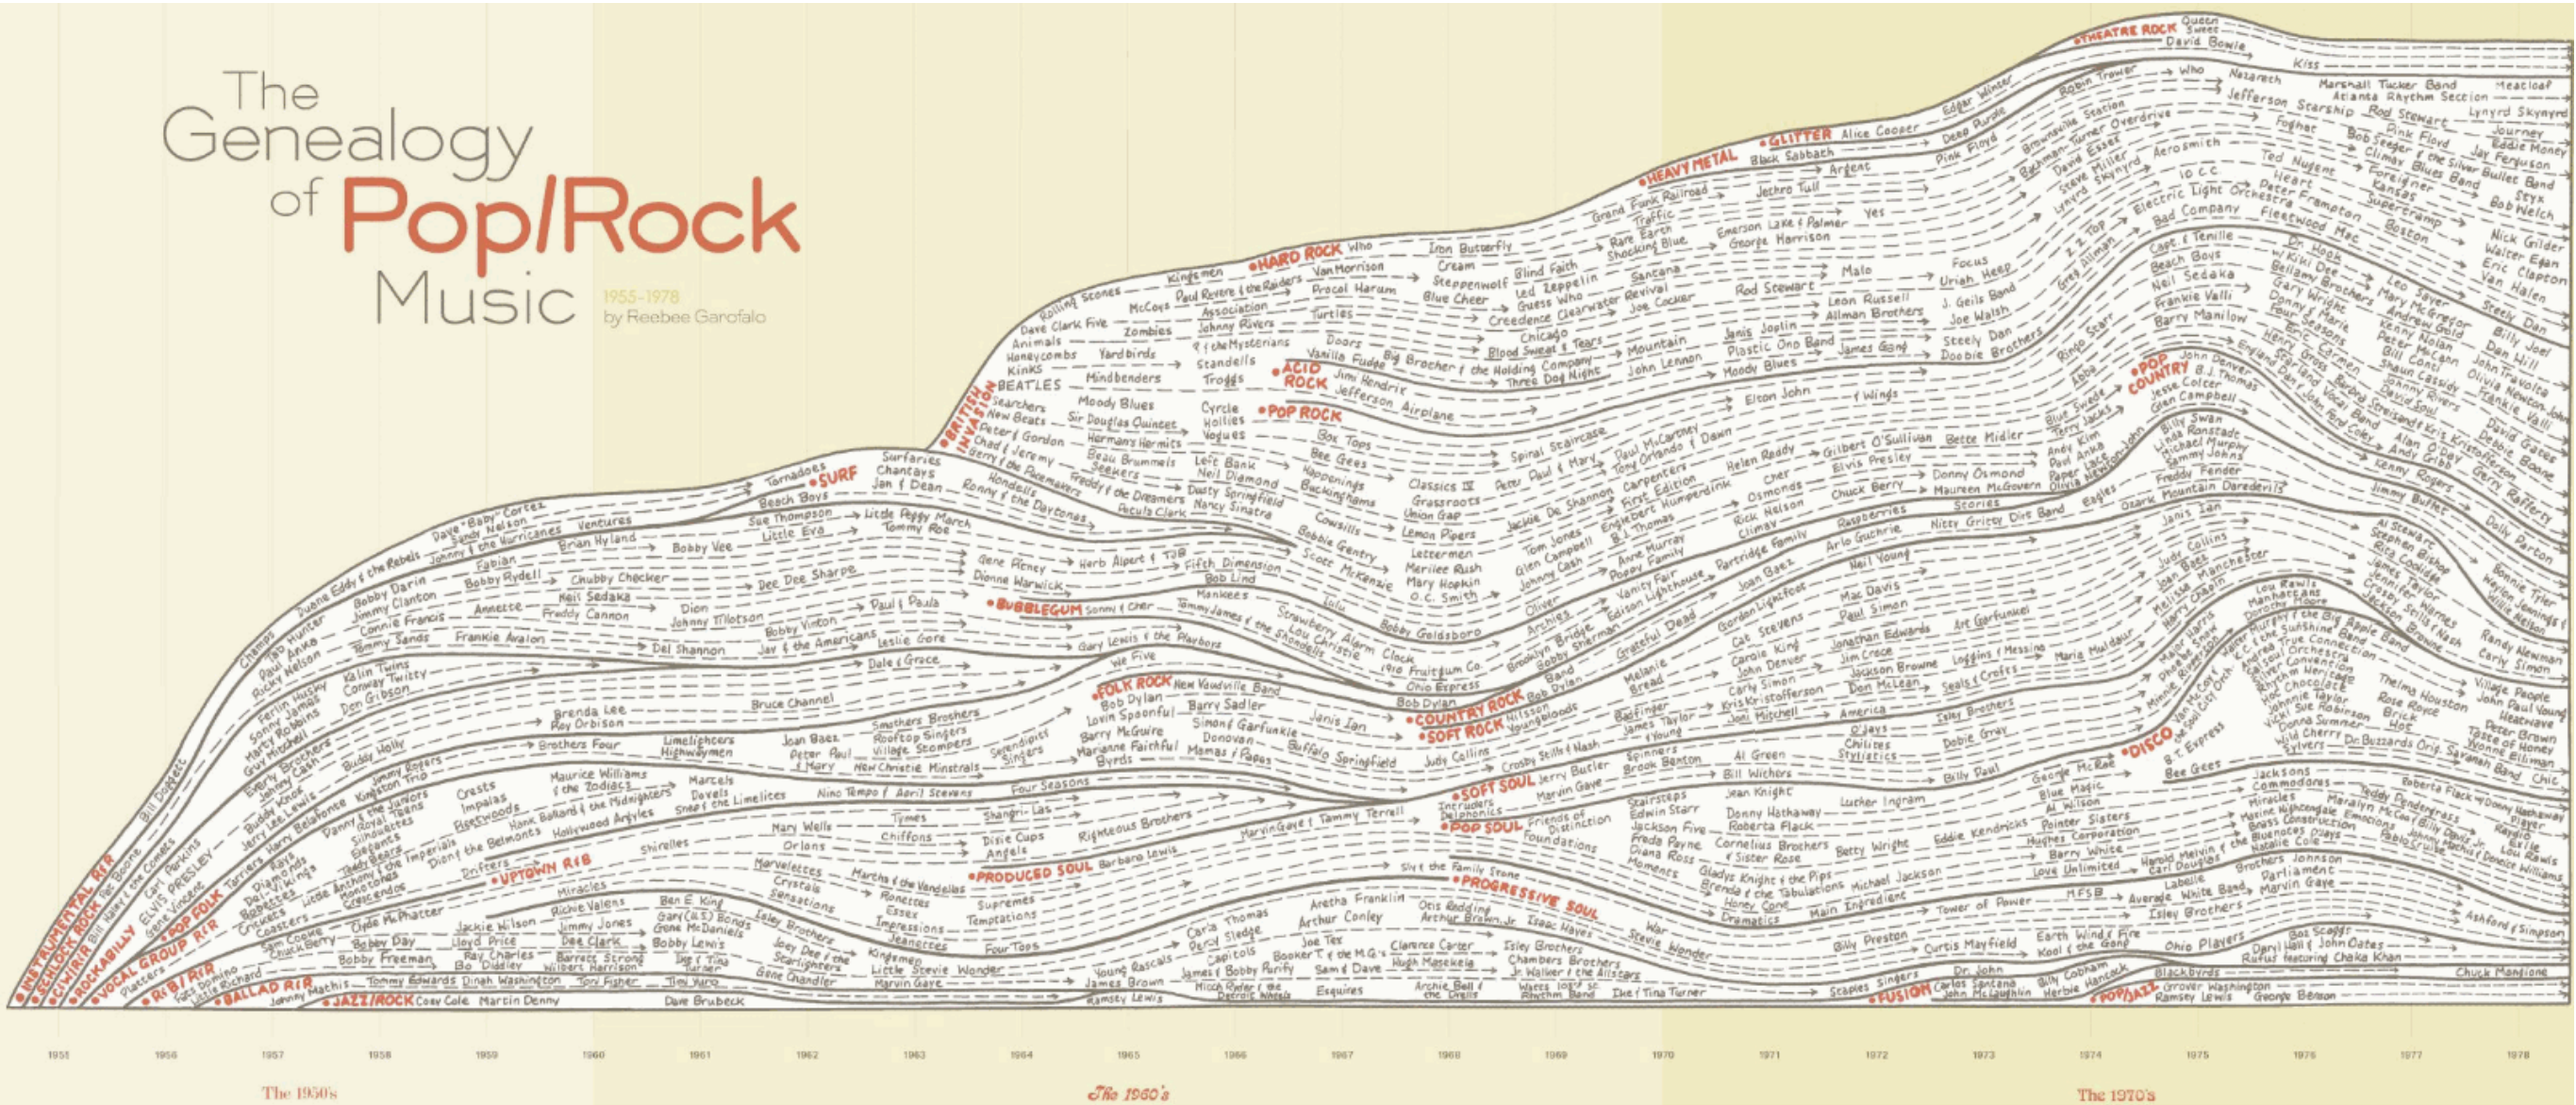

Genealogy of Pop and Rock Music - Reebee Garofalo

Ishkur's Guide to Electronic Dance Music

Ishkur's Guide to Electronic Dance Music

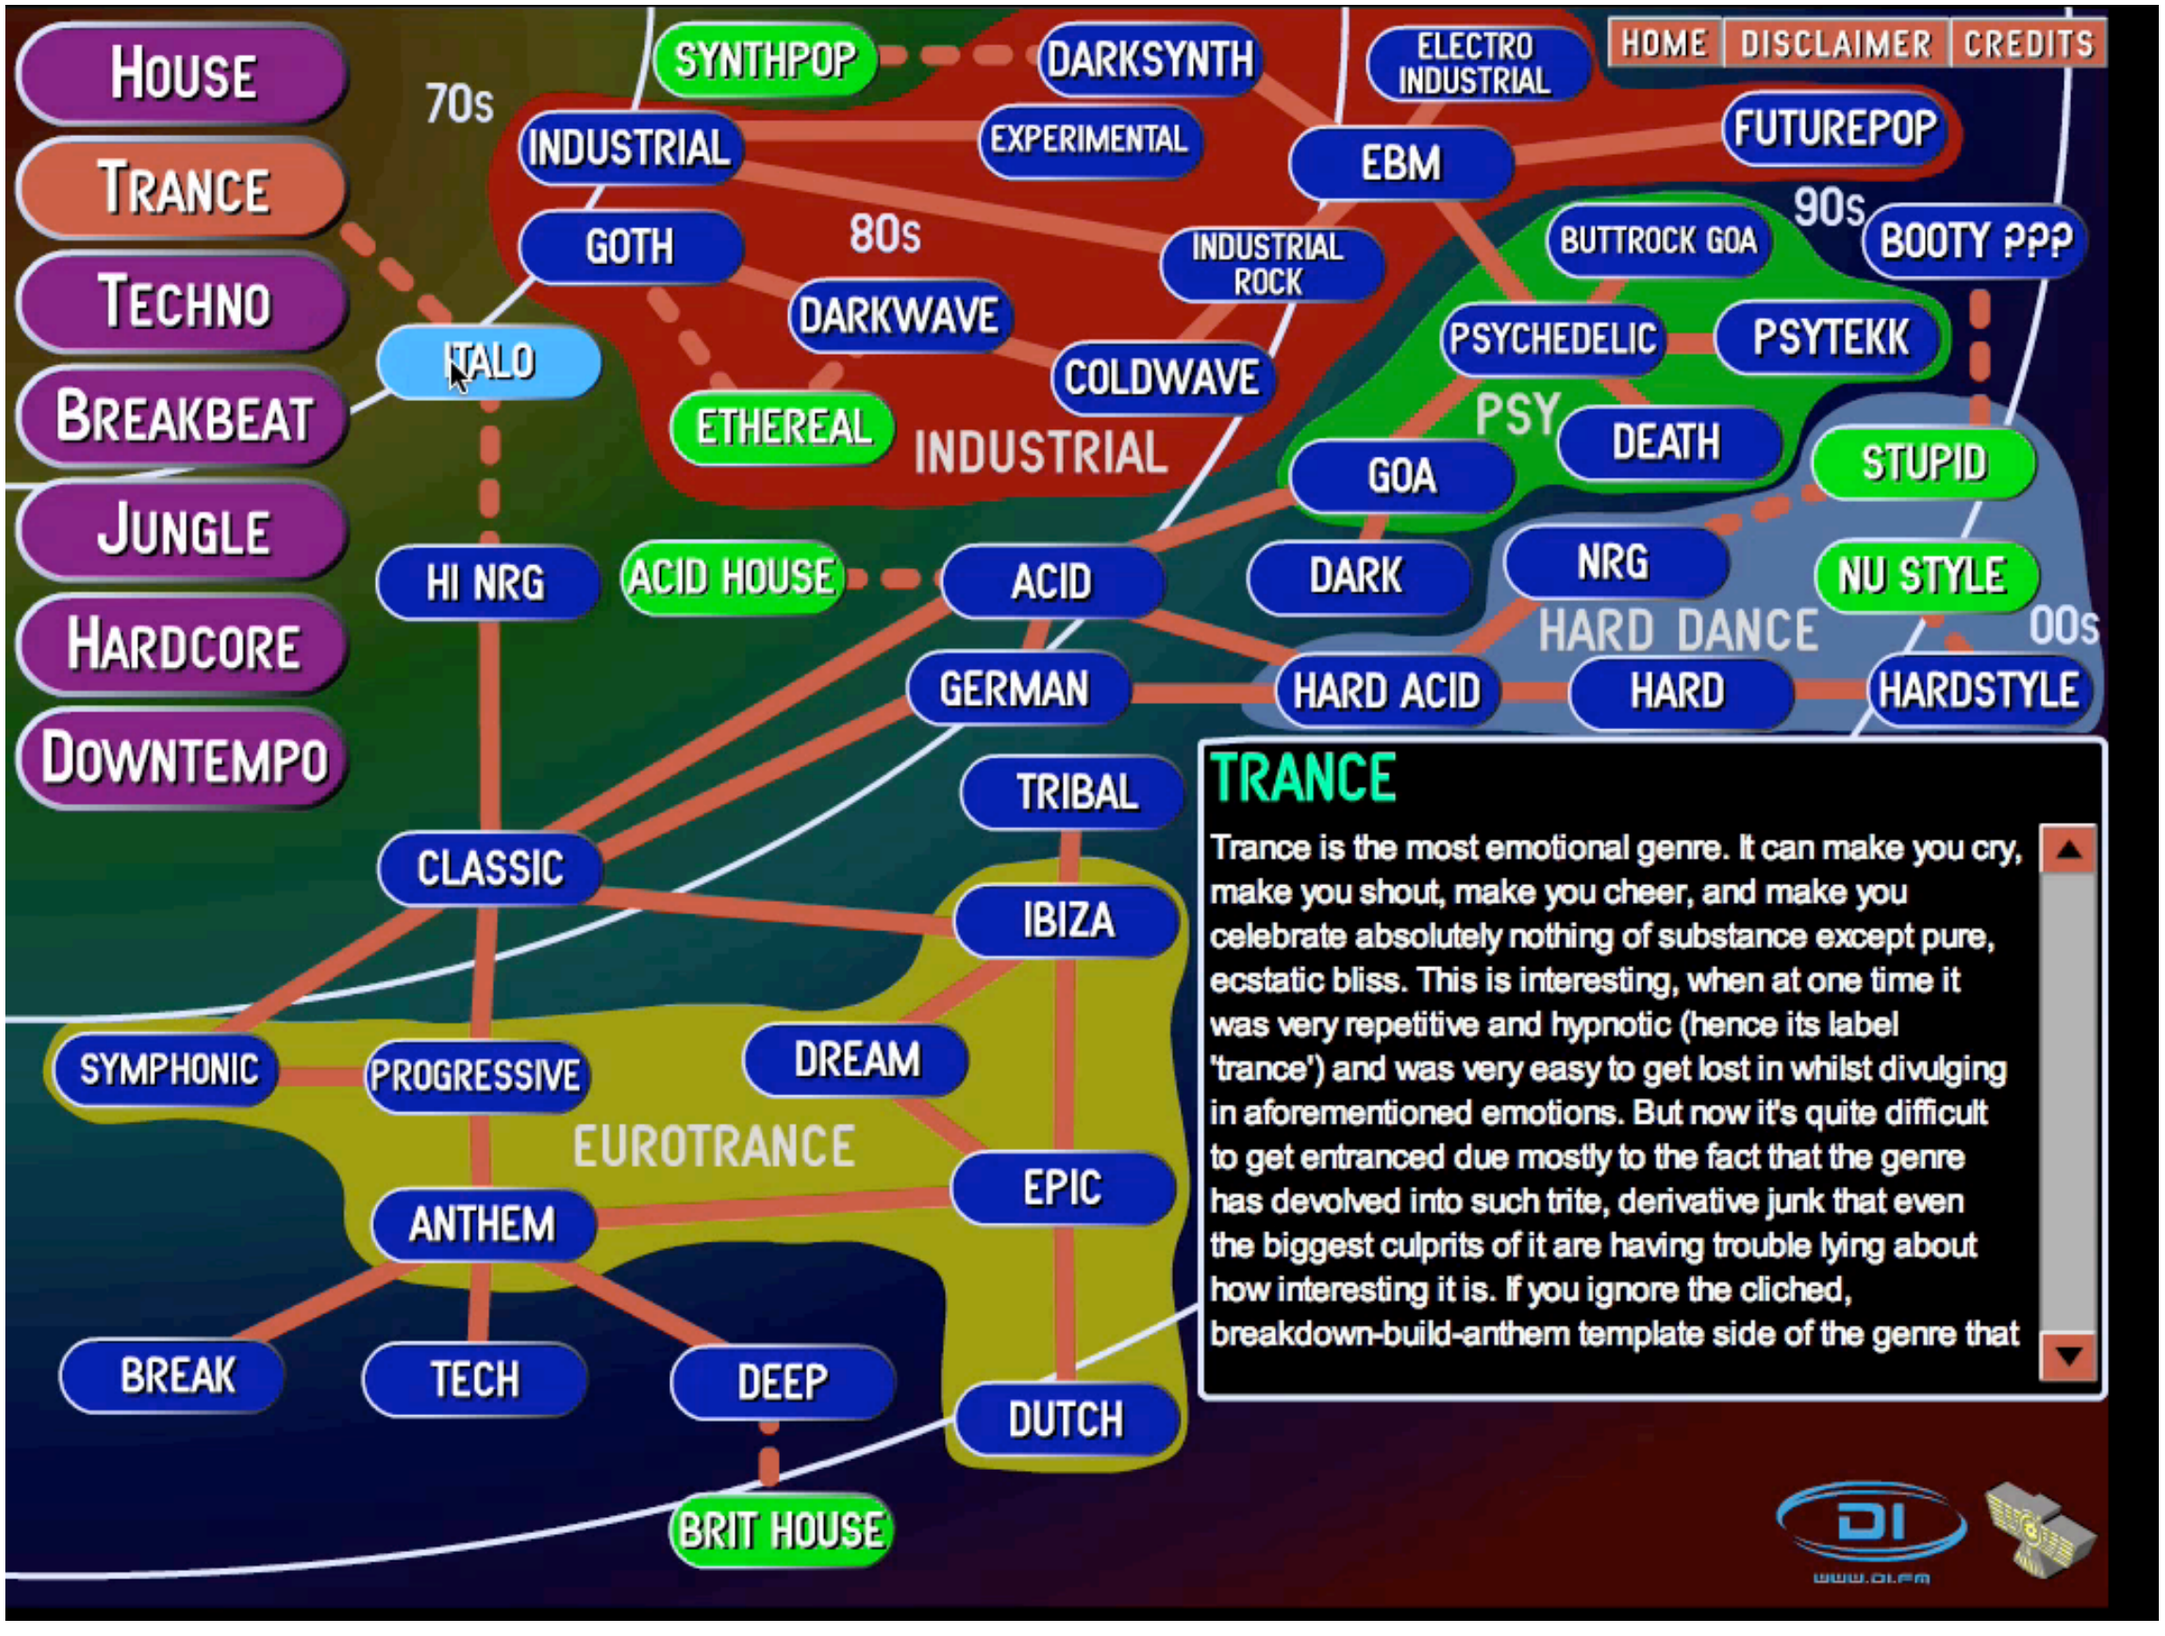

Music History as a London Tube Map

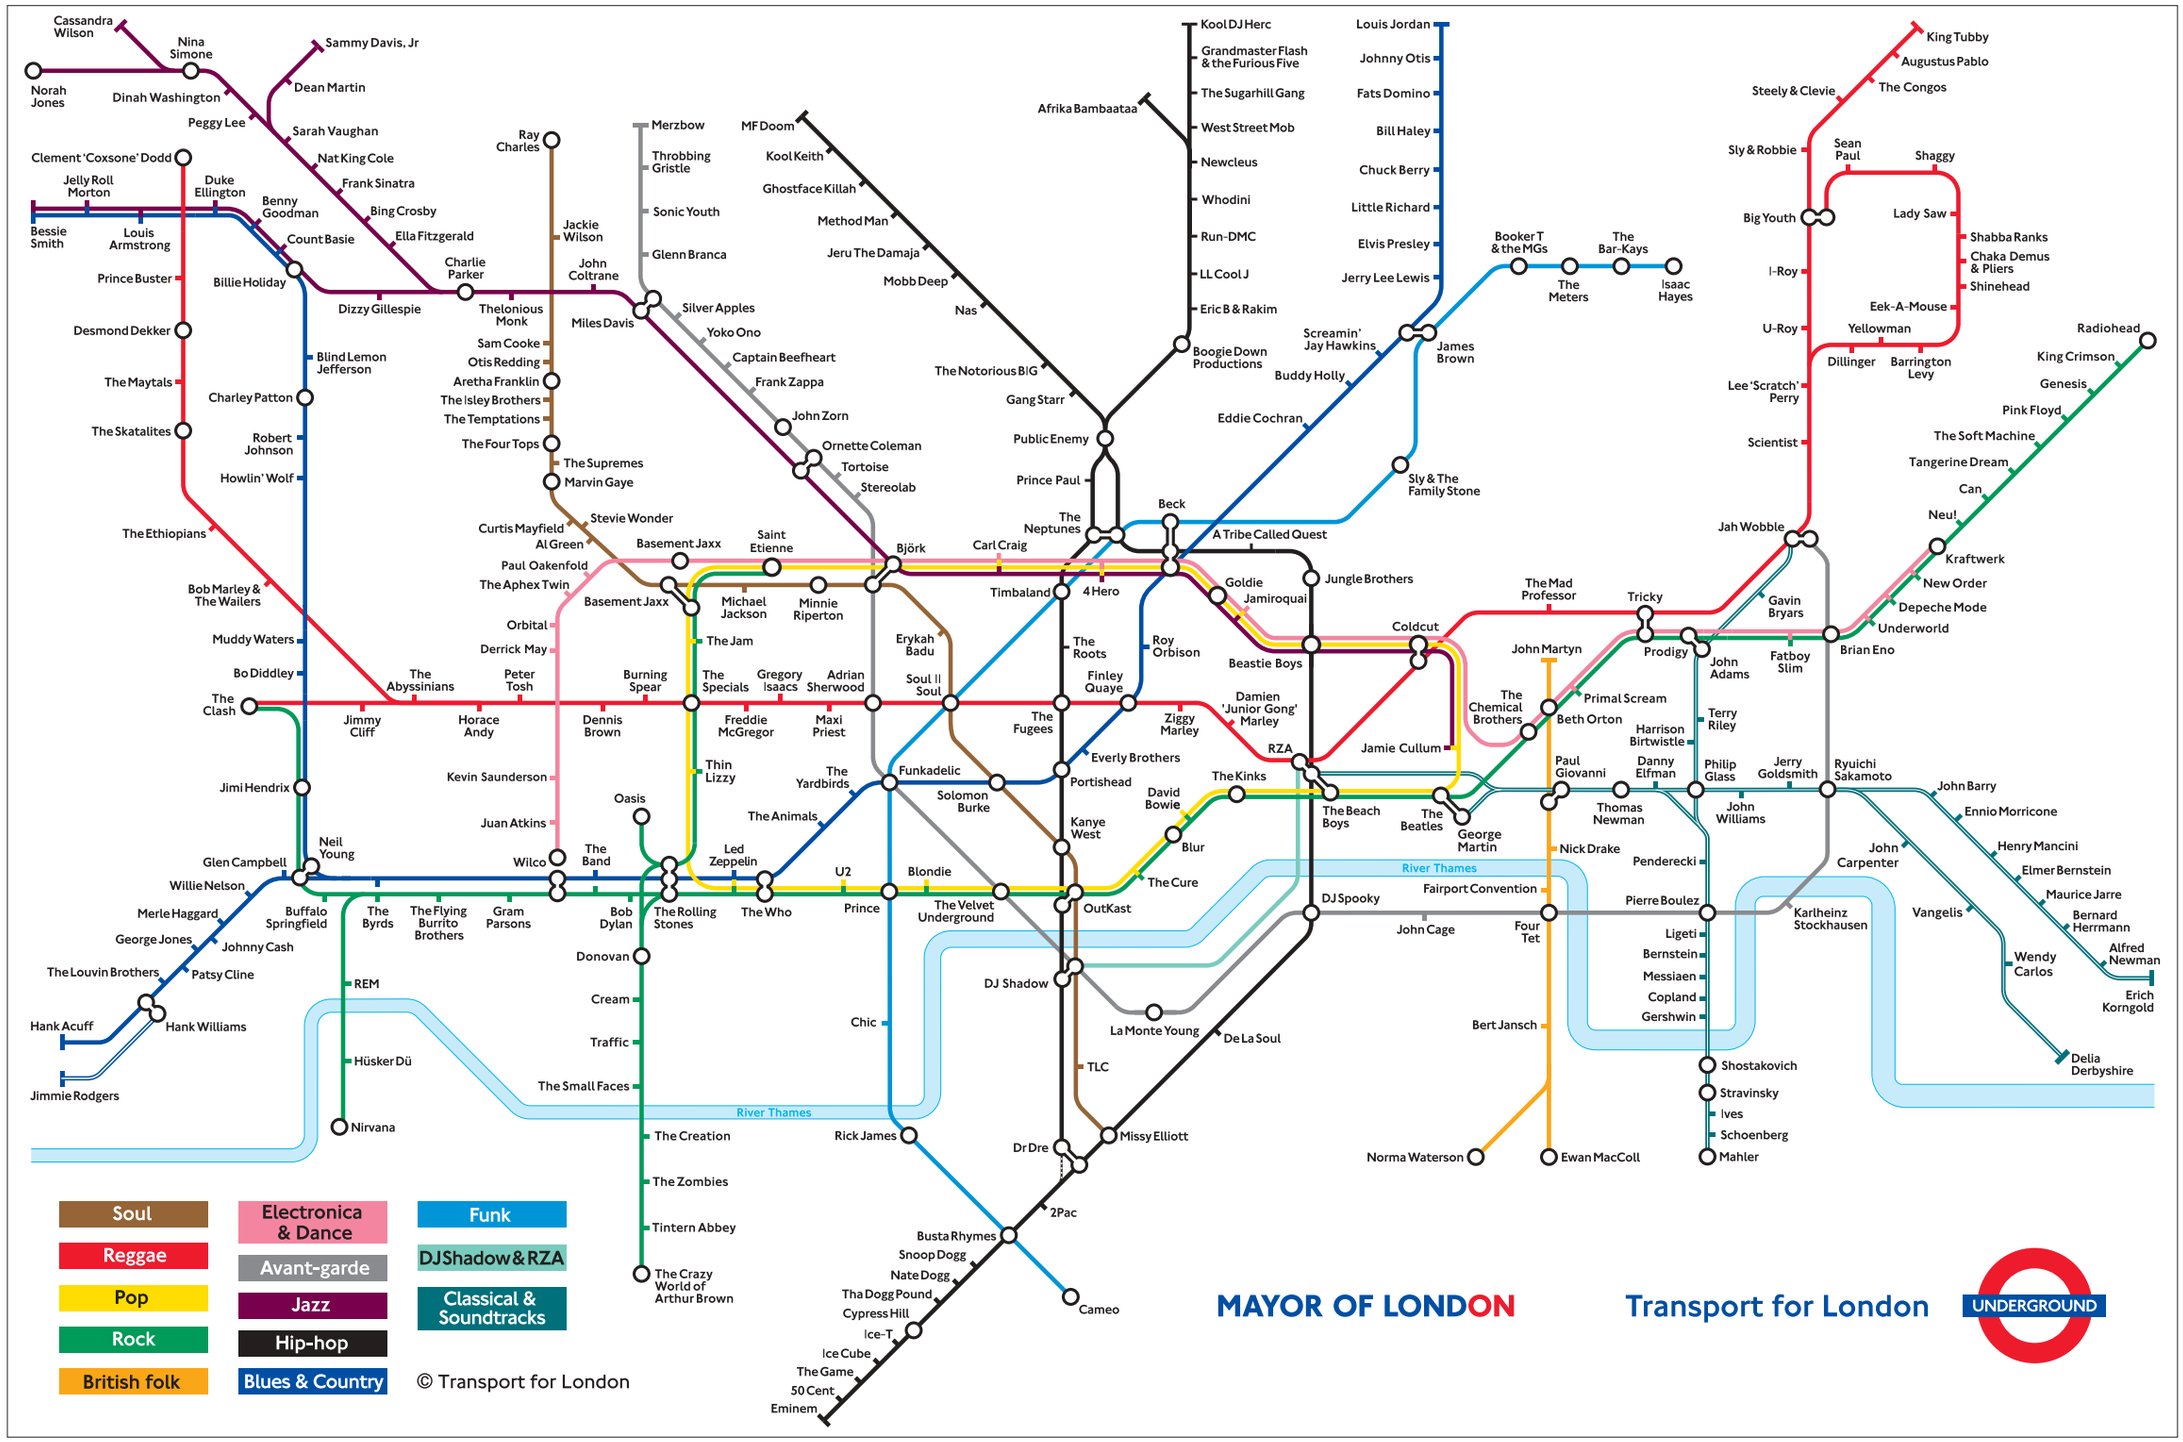

Dorian Lynskey - The Guardian

Music History as a London Tube Map

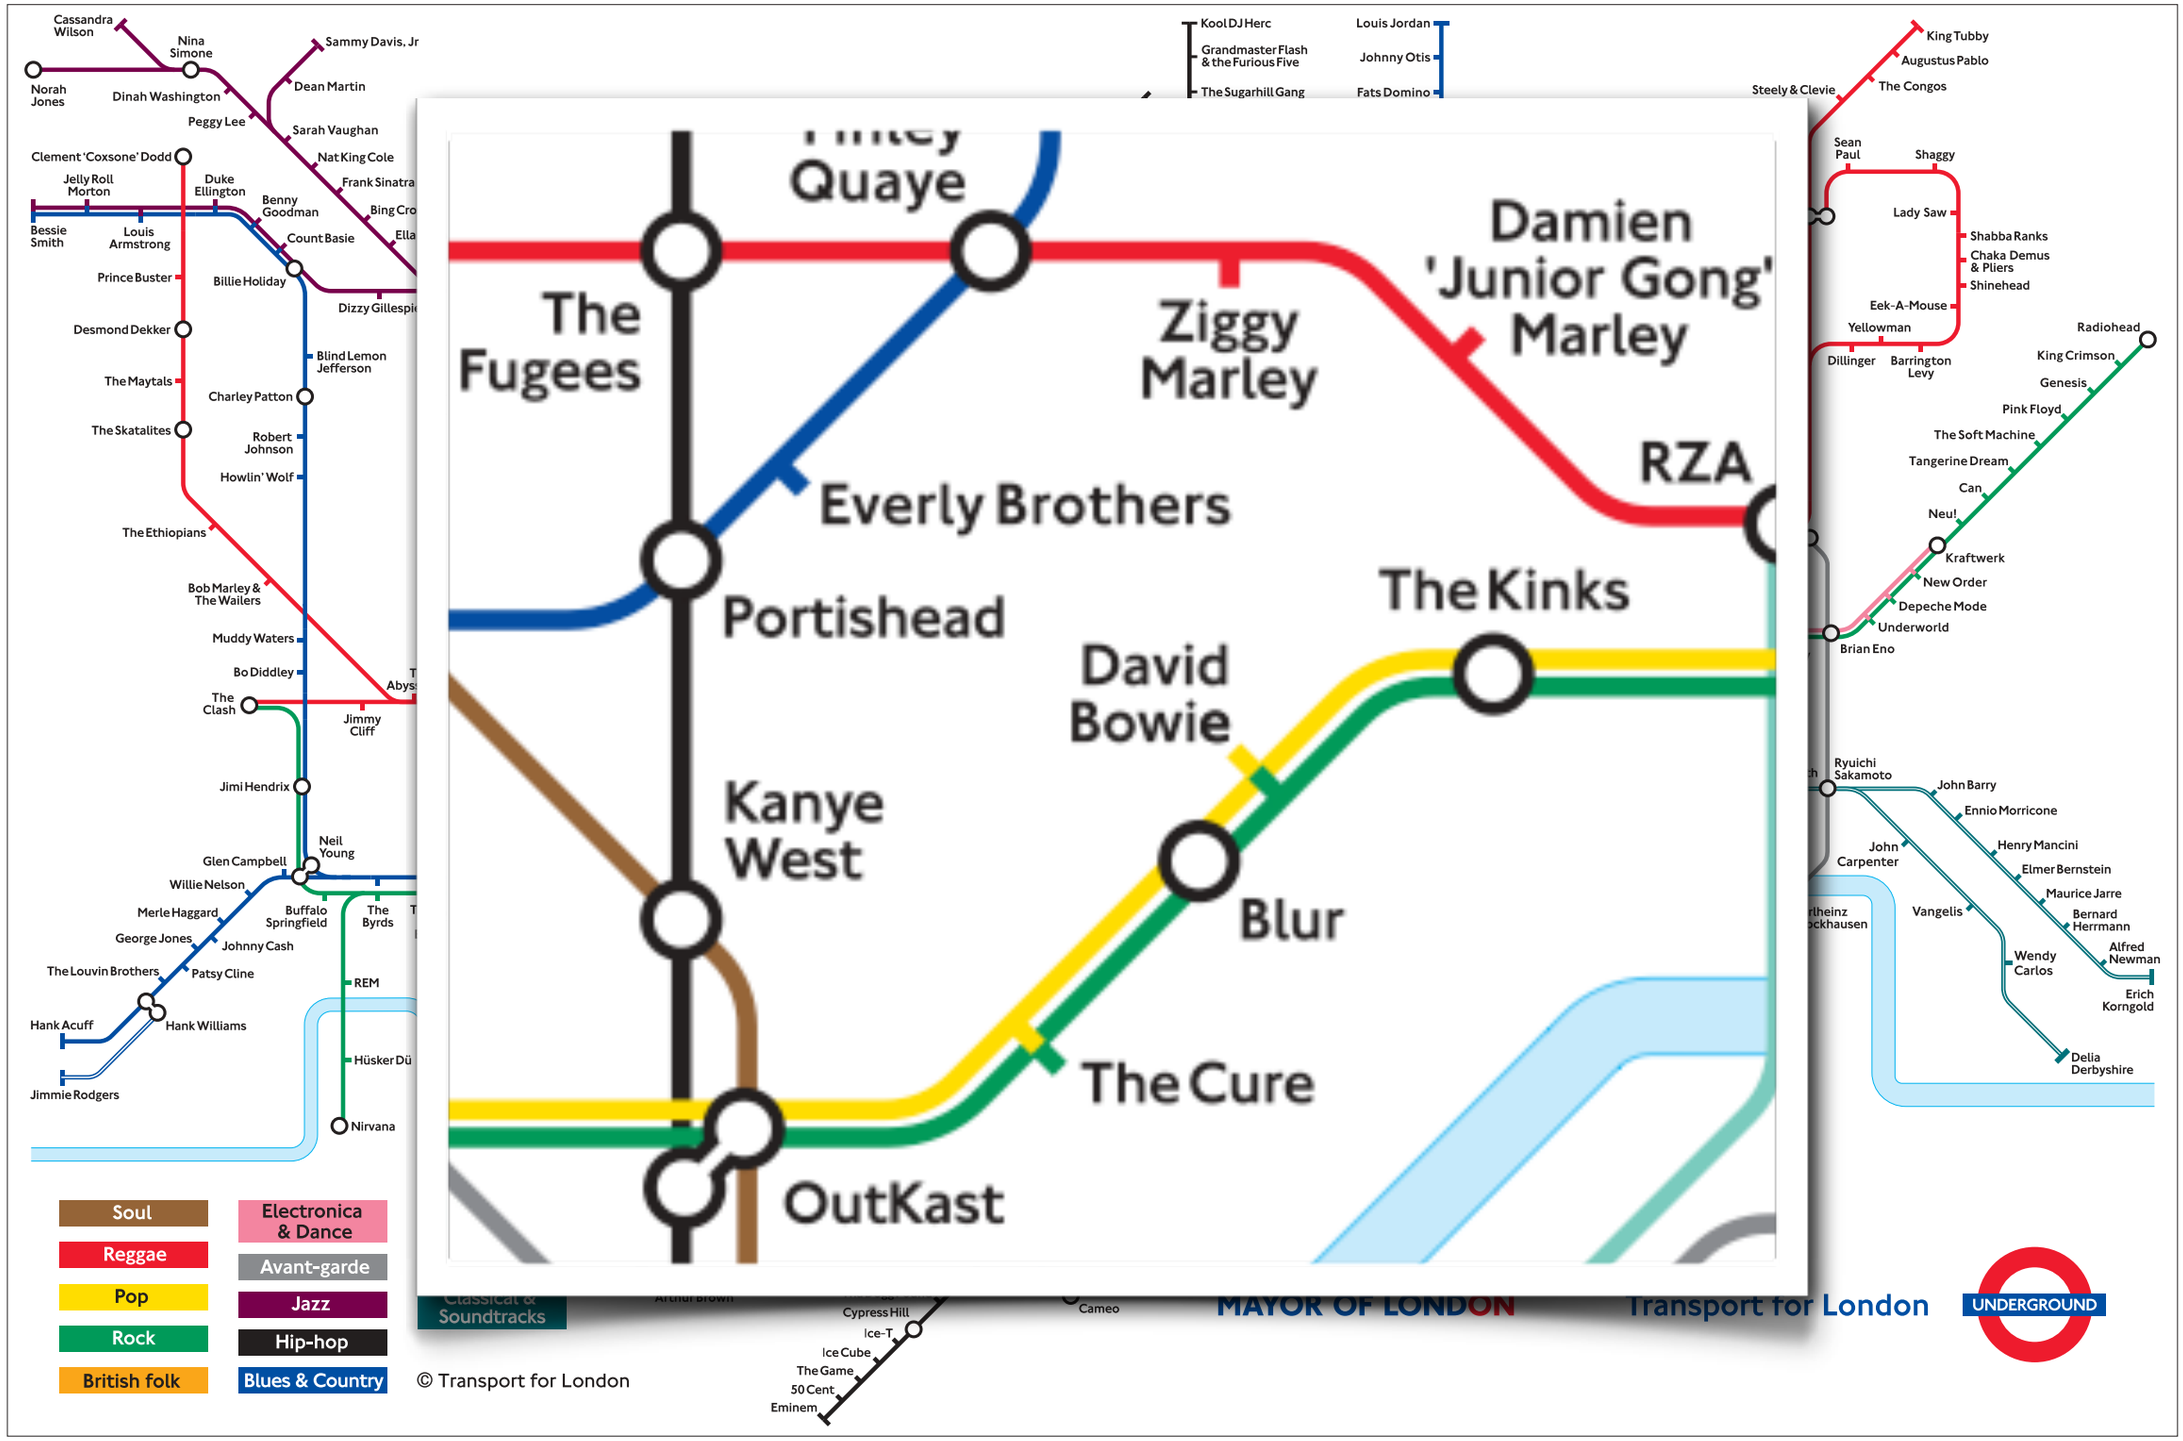

Dorian Lynskey - The Guardian

Visualizations as a graph

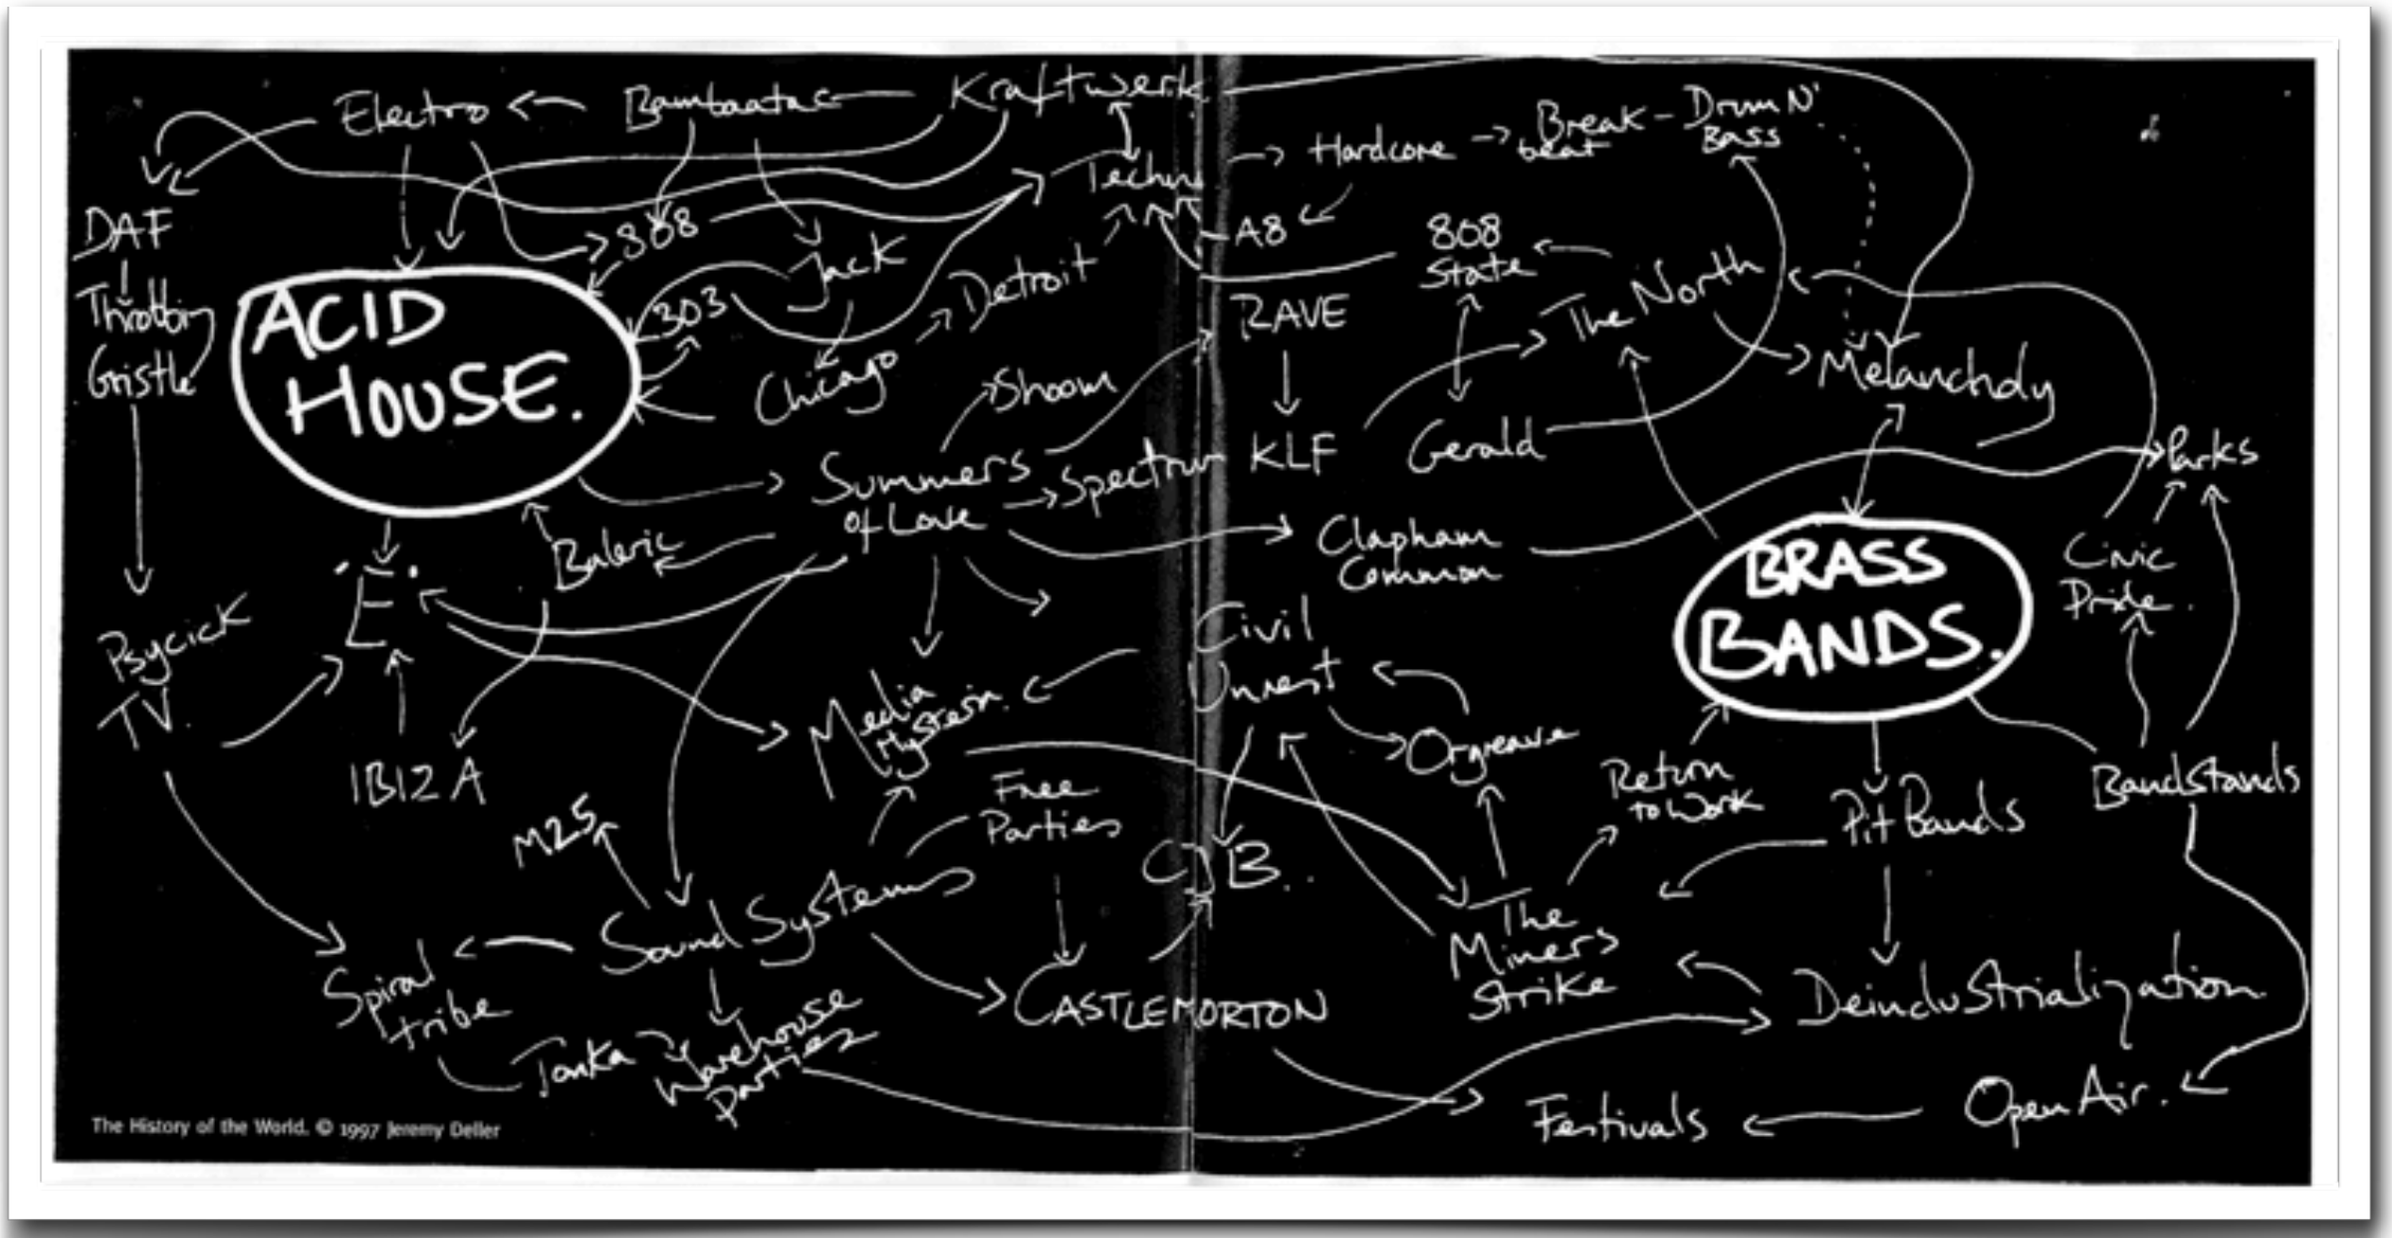

Visualizations as a graph

Visualizations as a graph

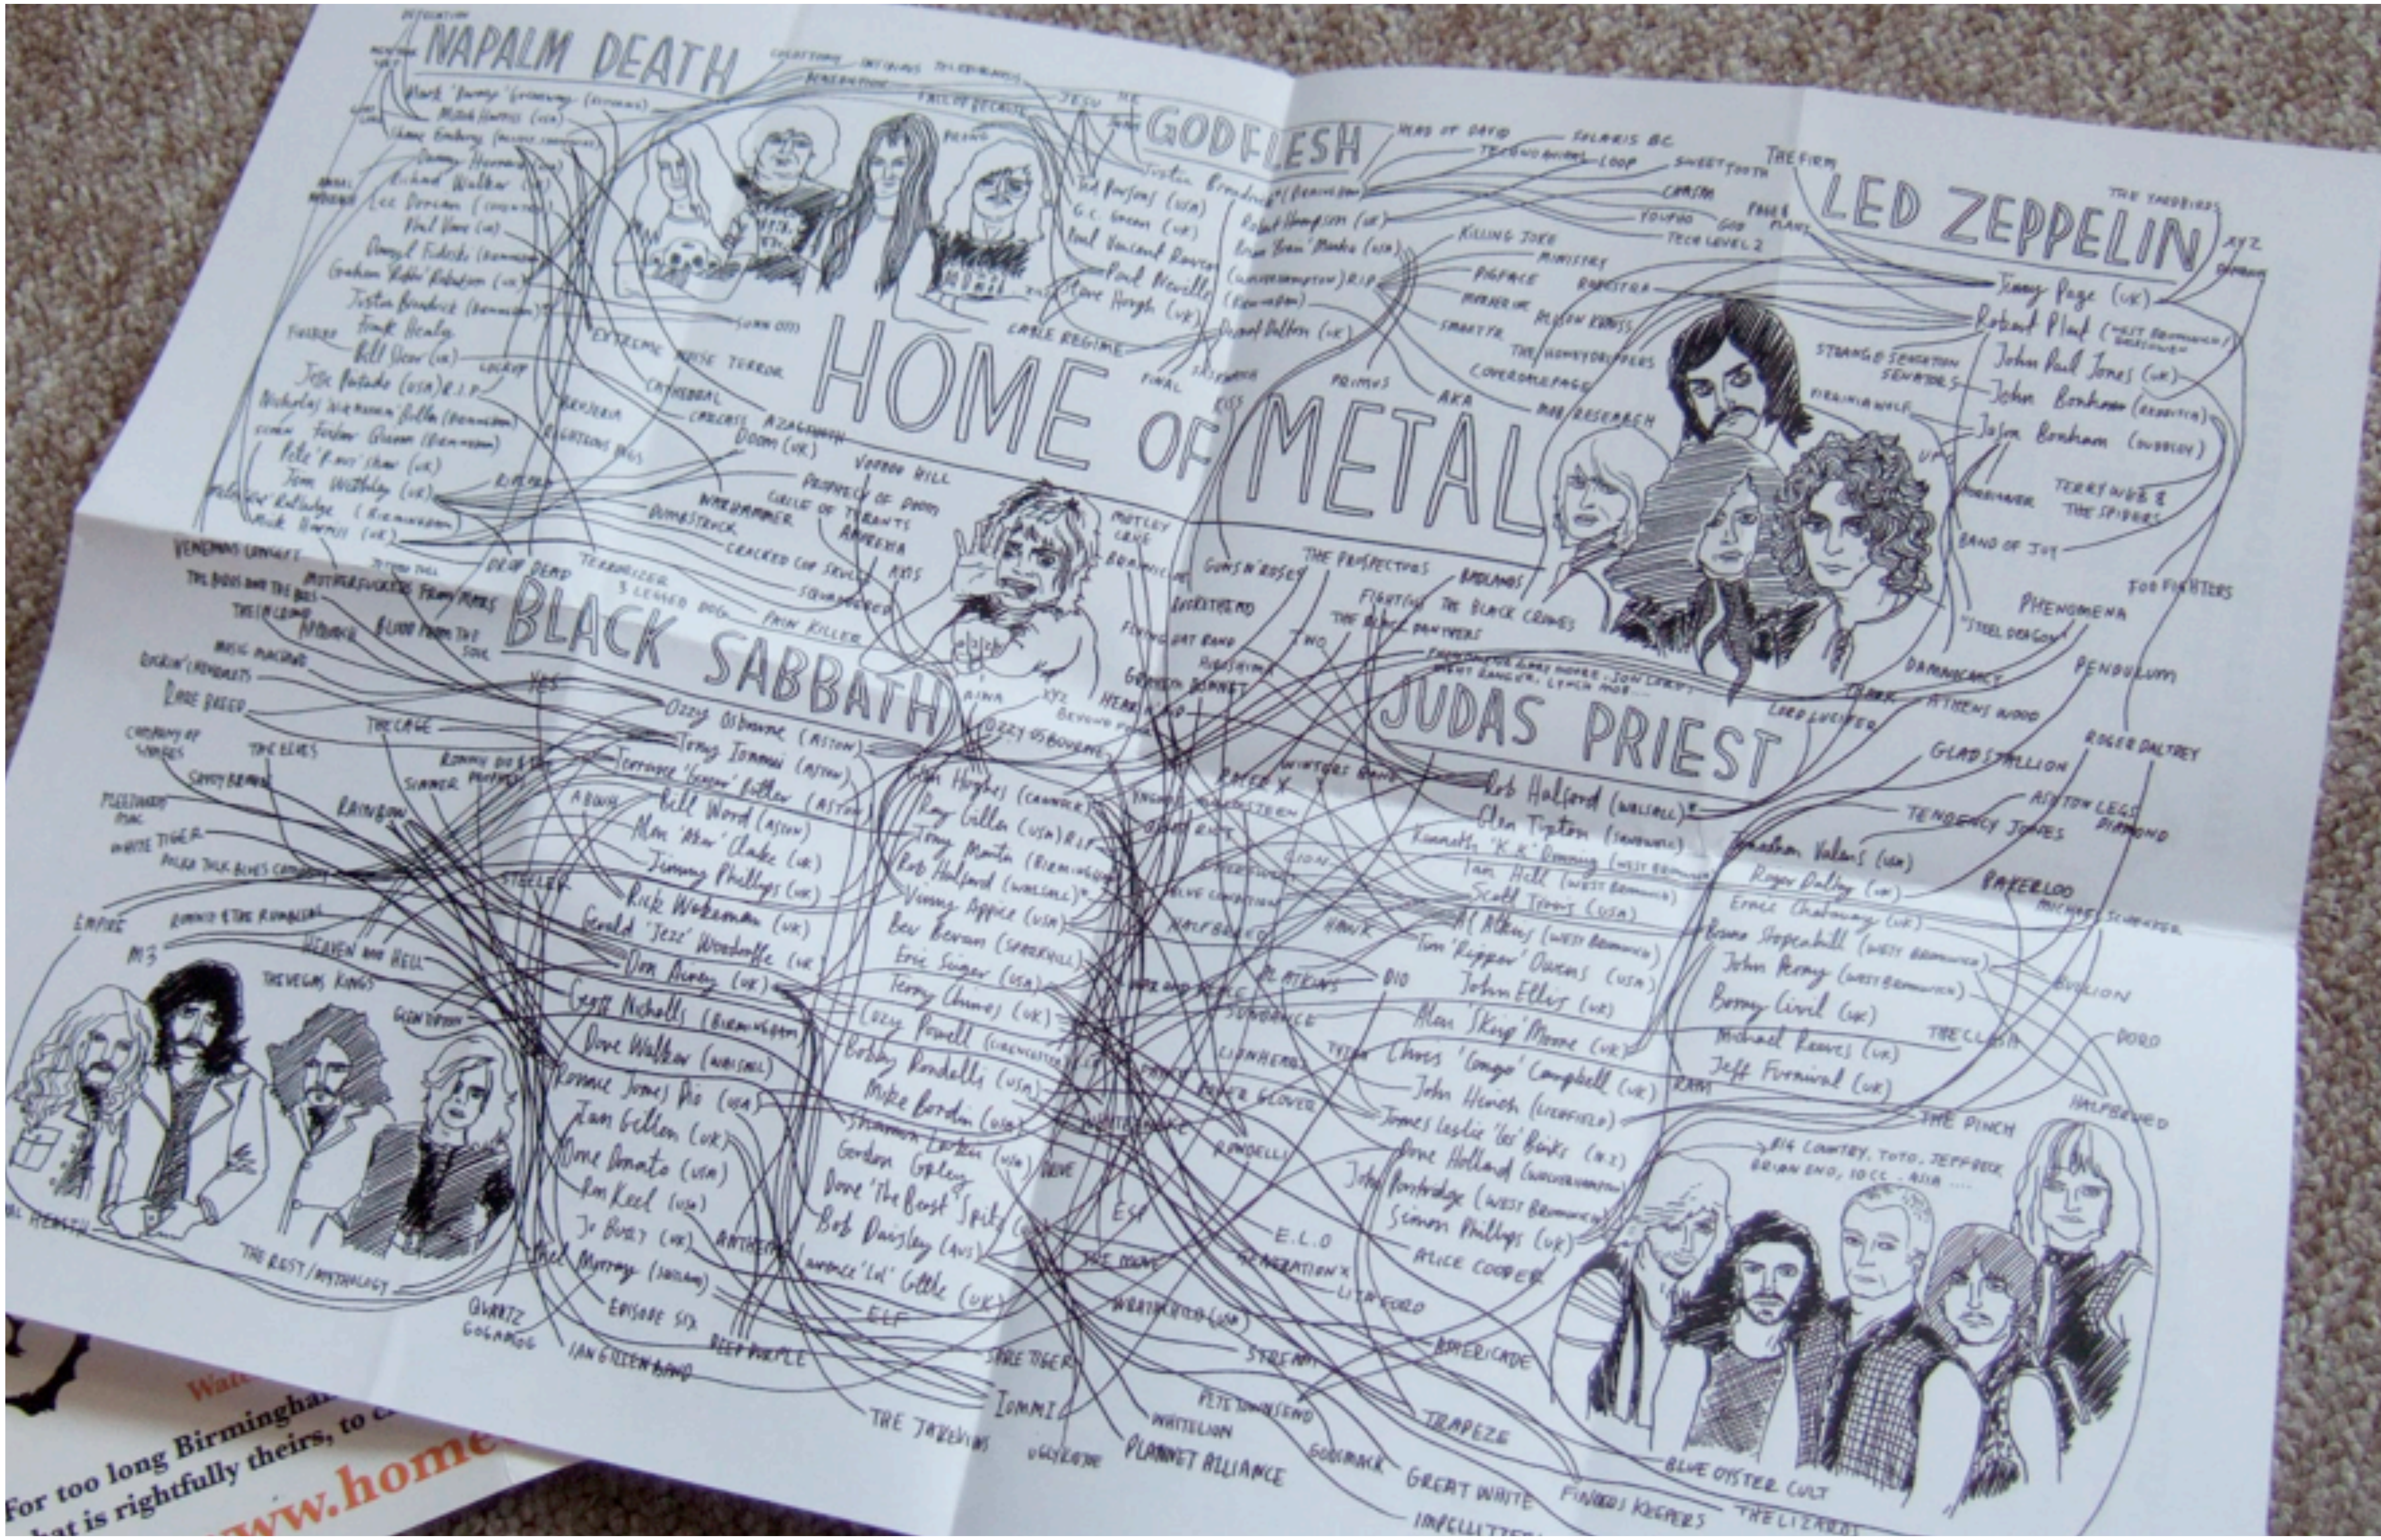

Visualizations as a graph

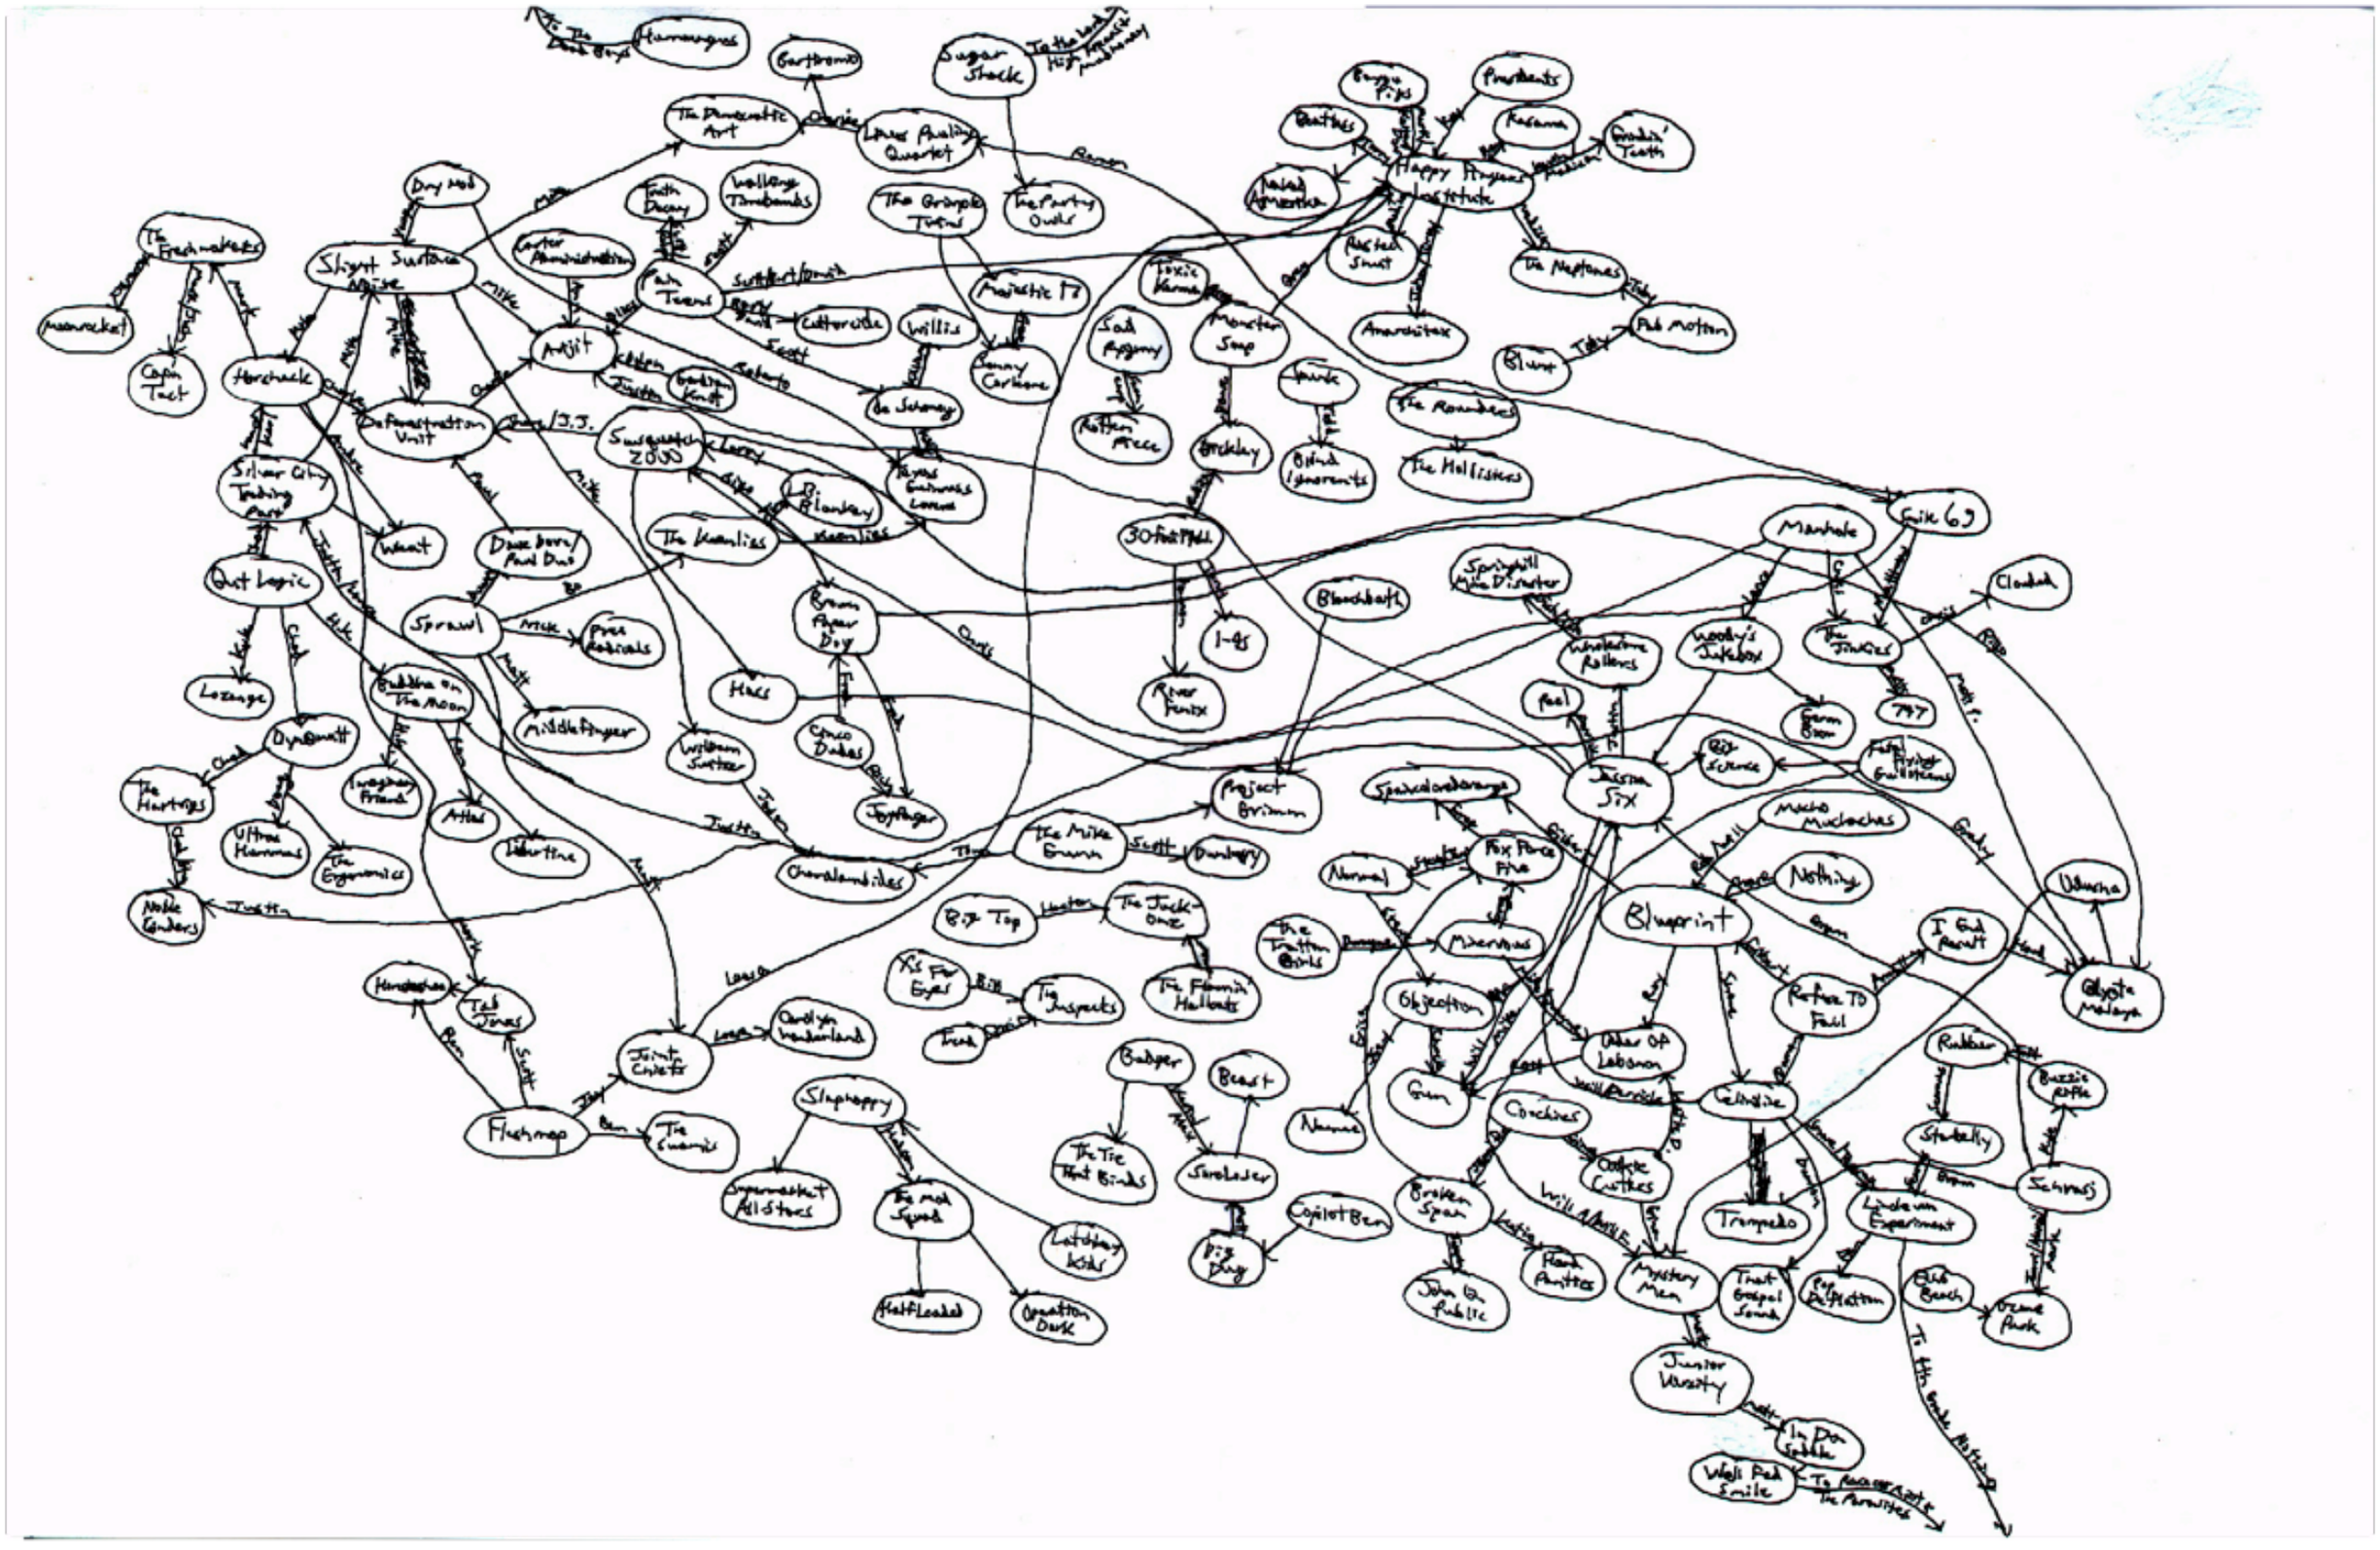

Visualizations as a graph

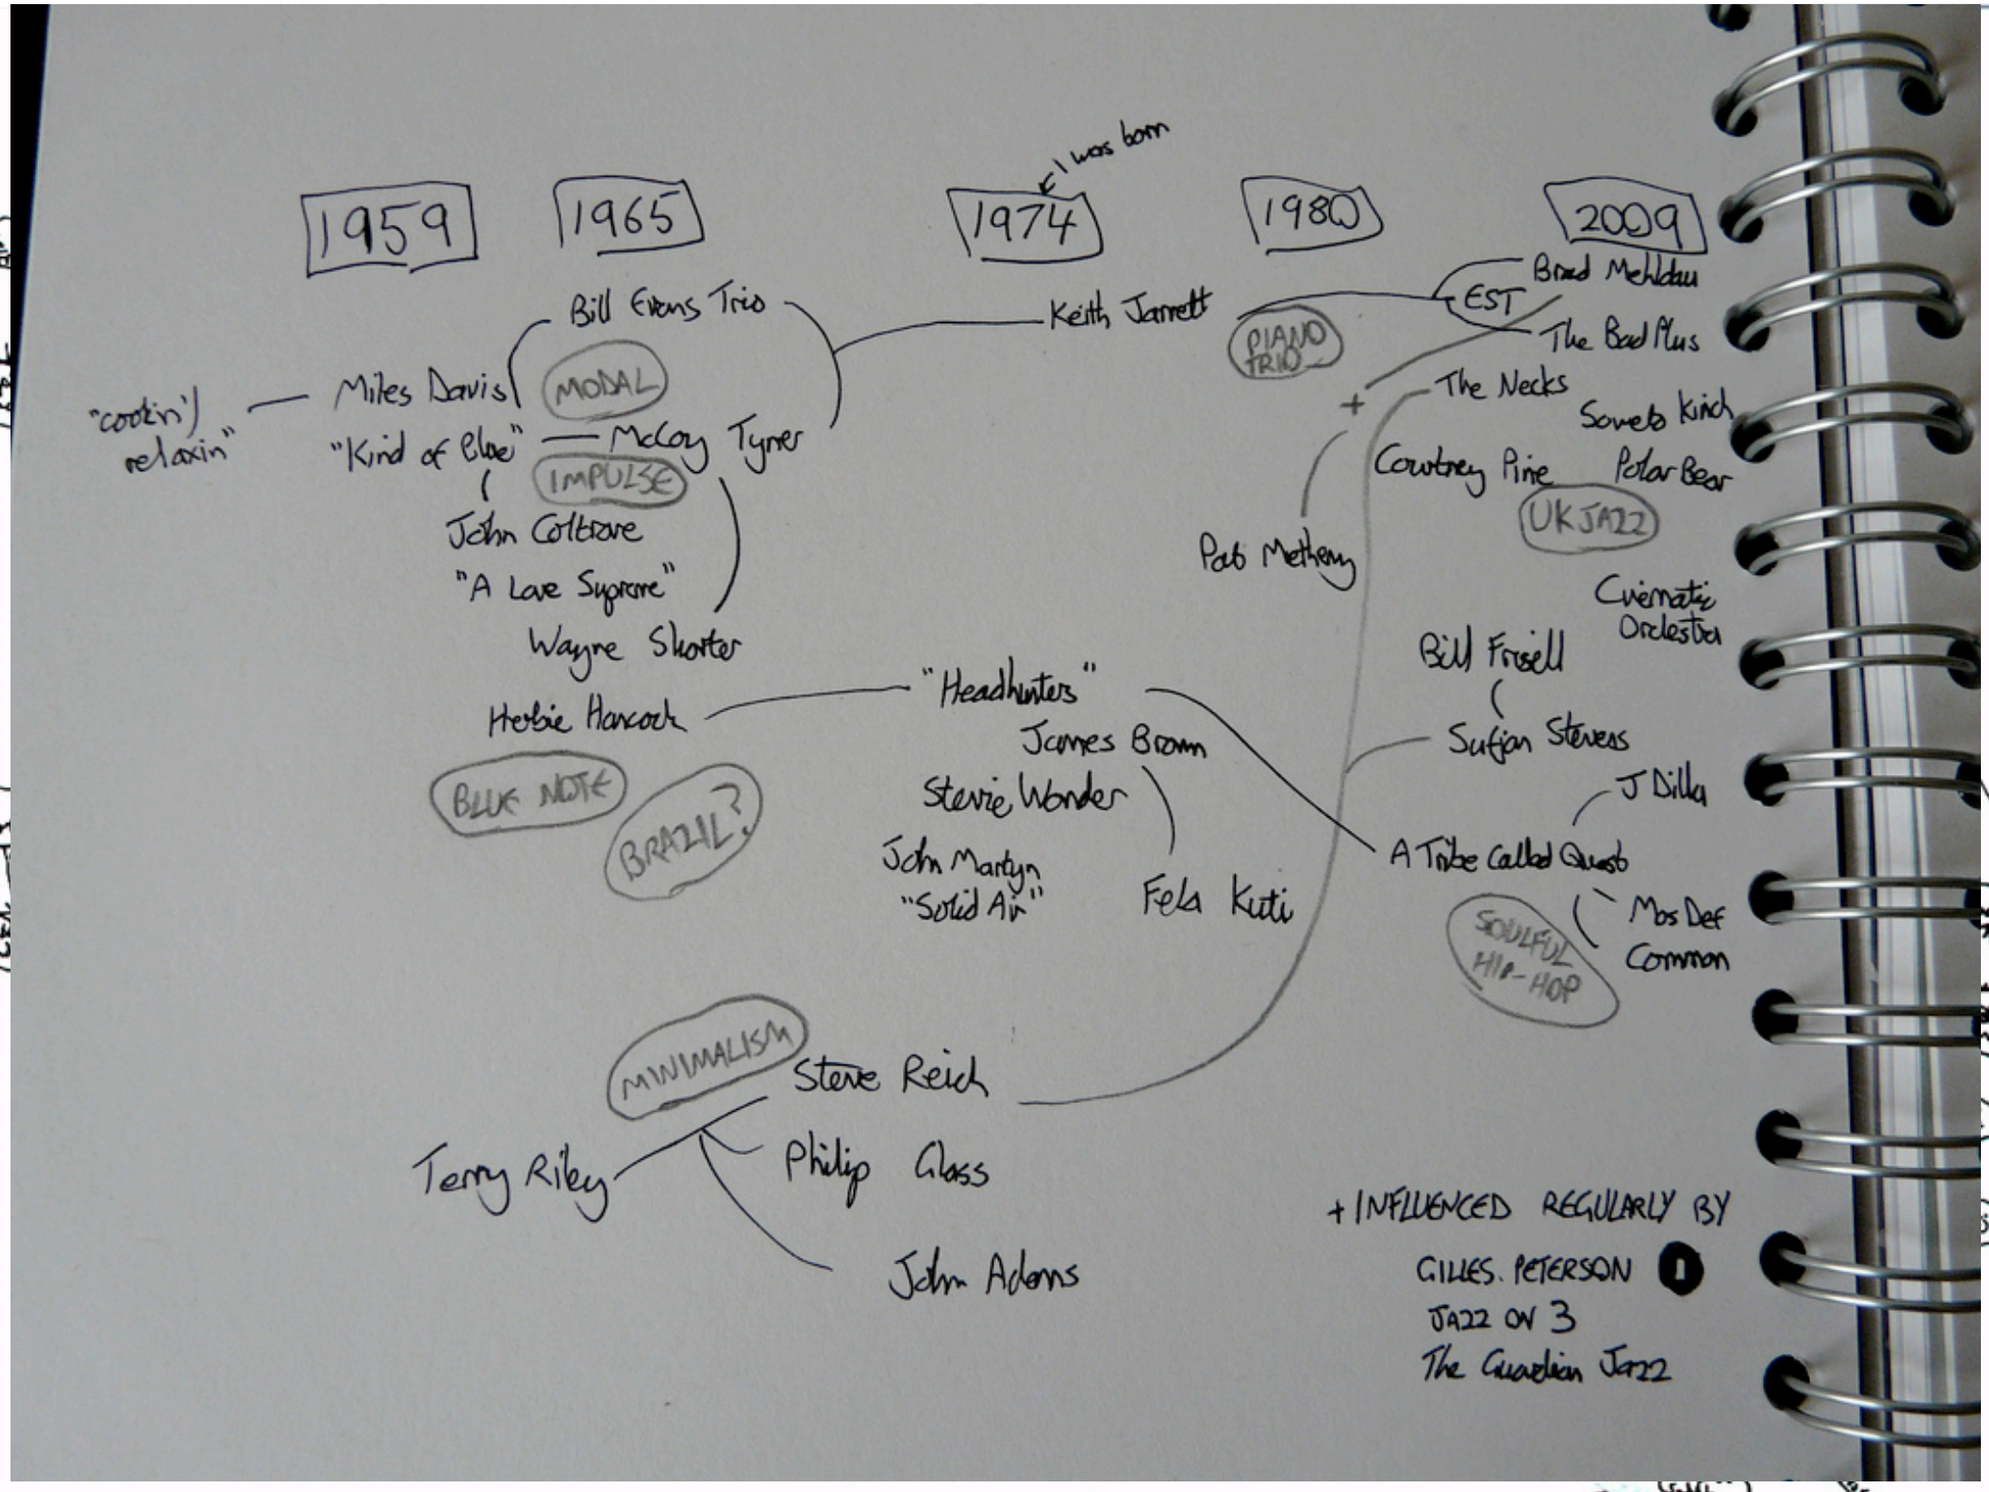

Visualizations as a graph

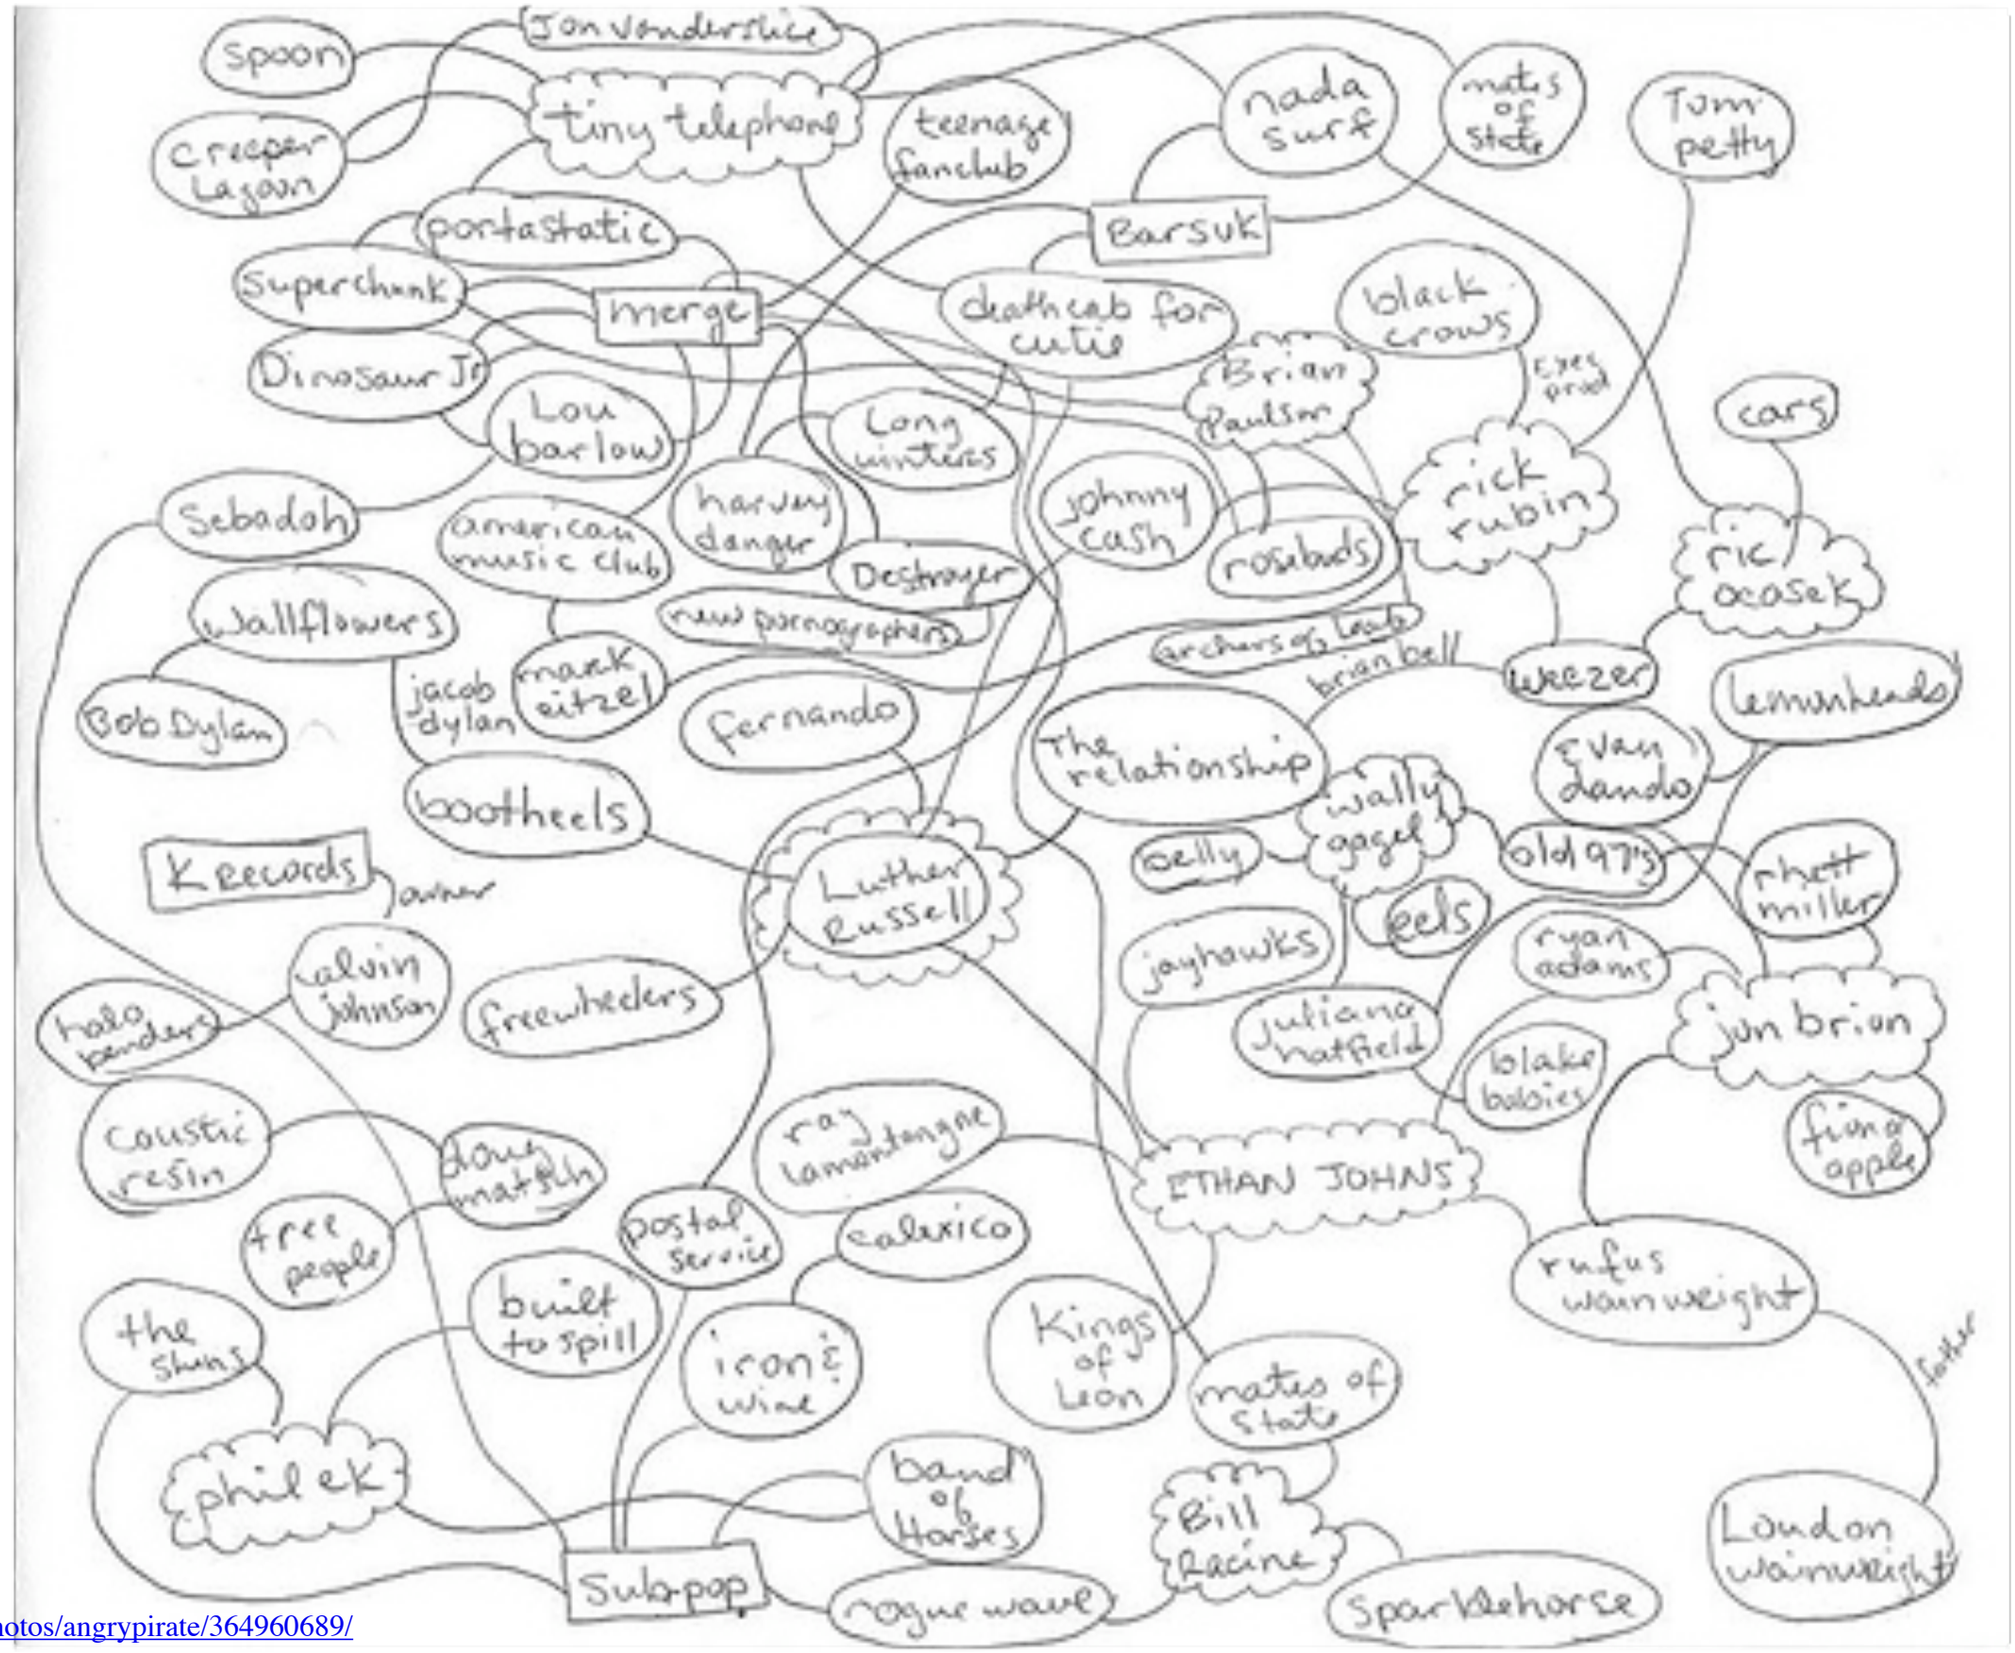

www.flickr.com/photos/angrypirate/364960689/

Visualizations as a graph

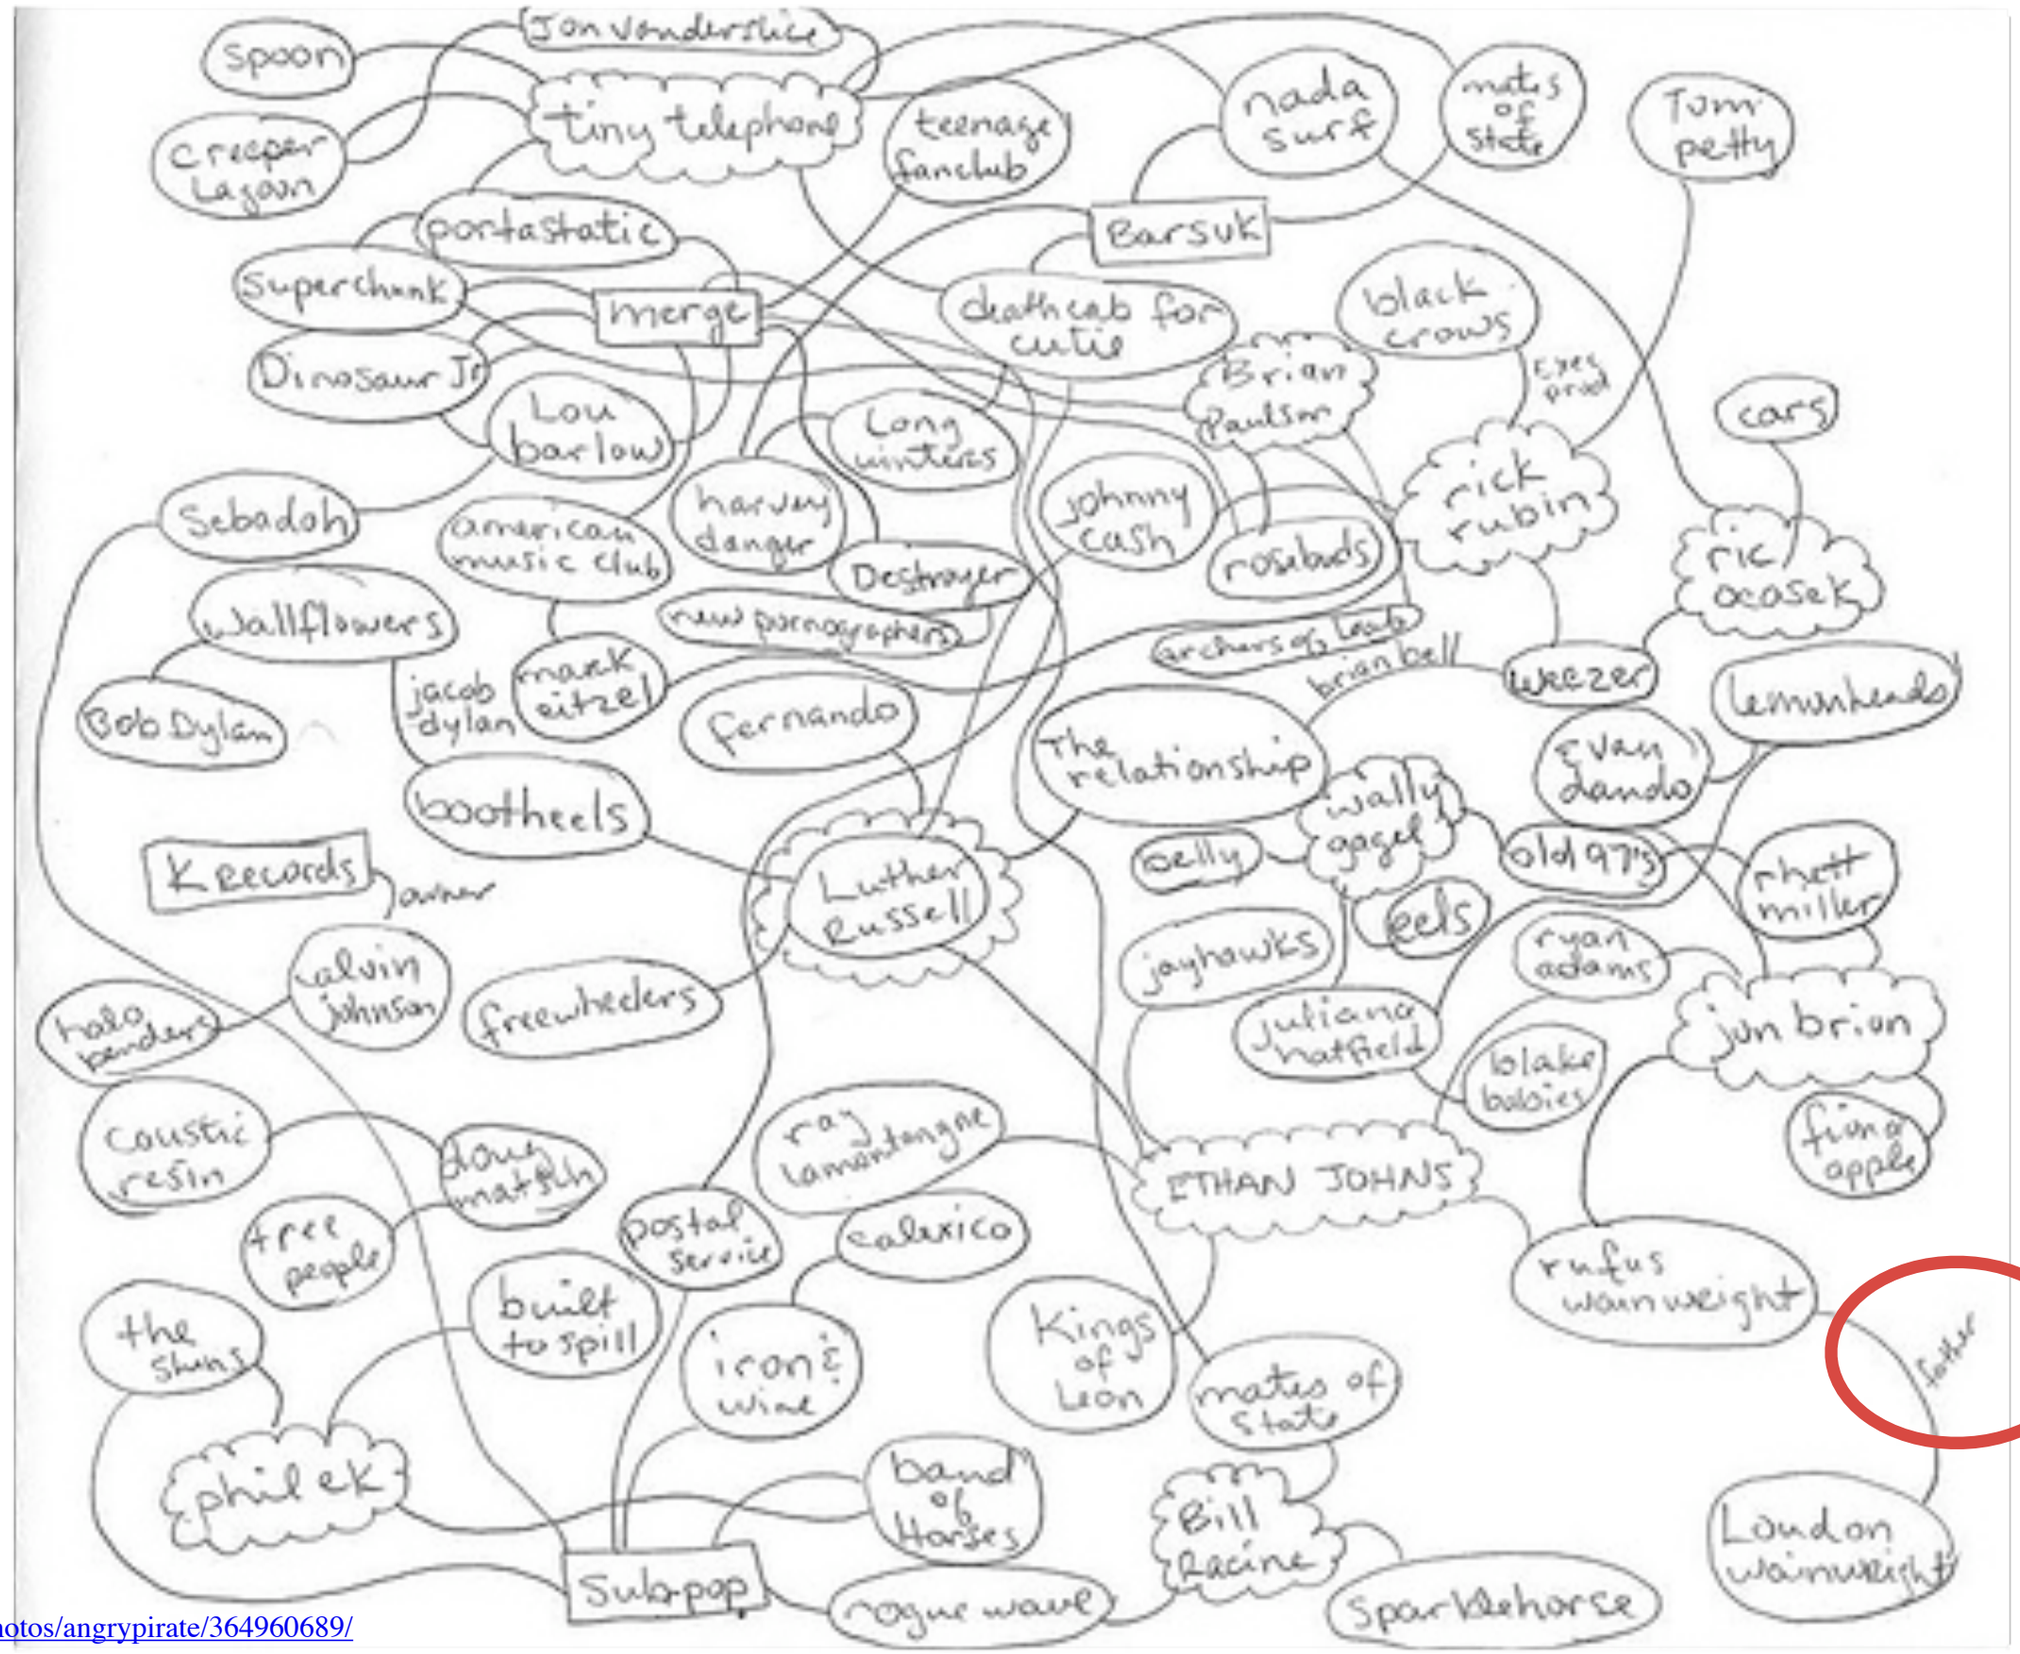

www.flickr.com/photos/angrypirate/364960689/

Visualizations as a graph

www.flickr.com/photos/angrypirate/364960689/

Visualizations as a graph

www.flickr.com/photos/angrypirate/364960689/

Visualizations as a graph

www.flickr.com/photos/angrypirate/364960689/

Visual Music Guiden

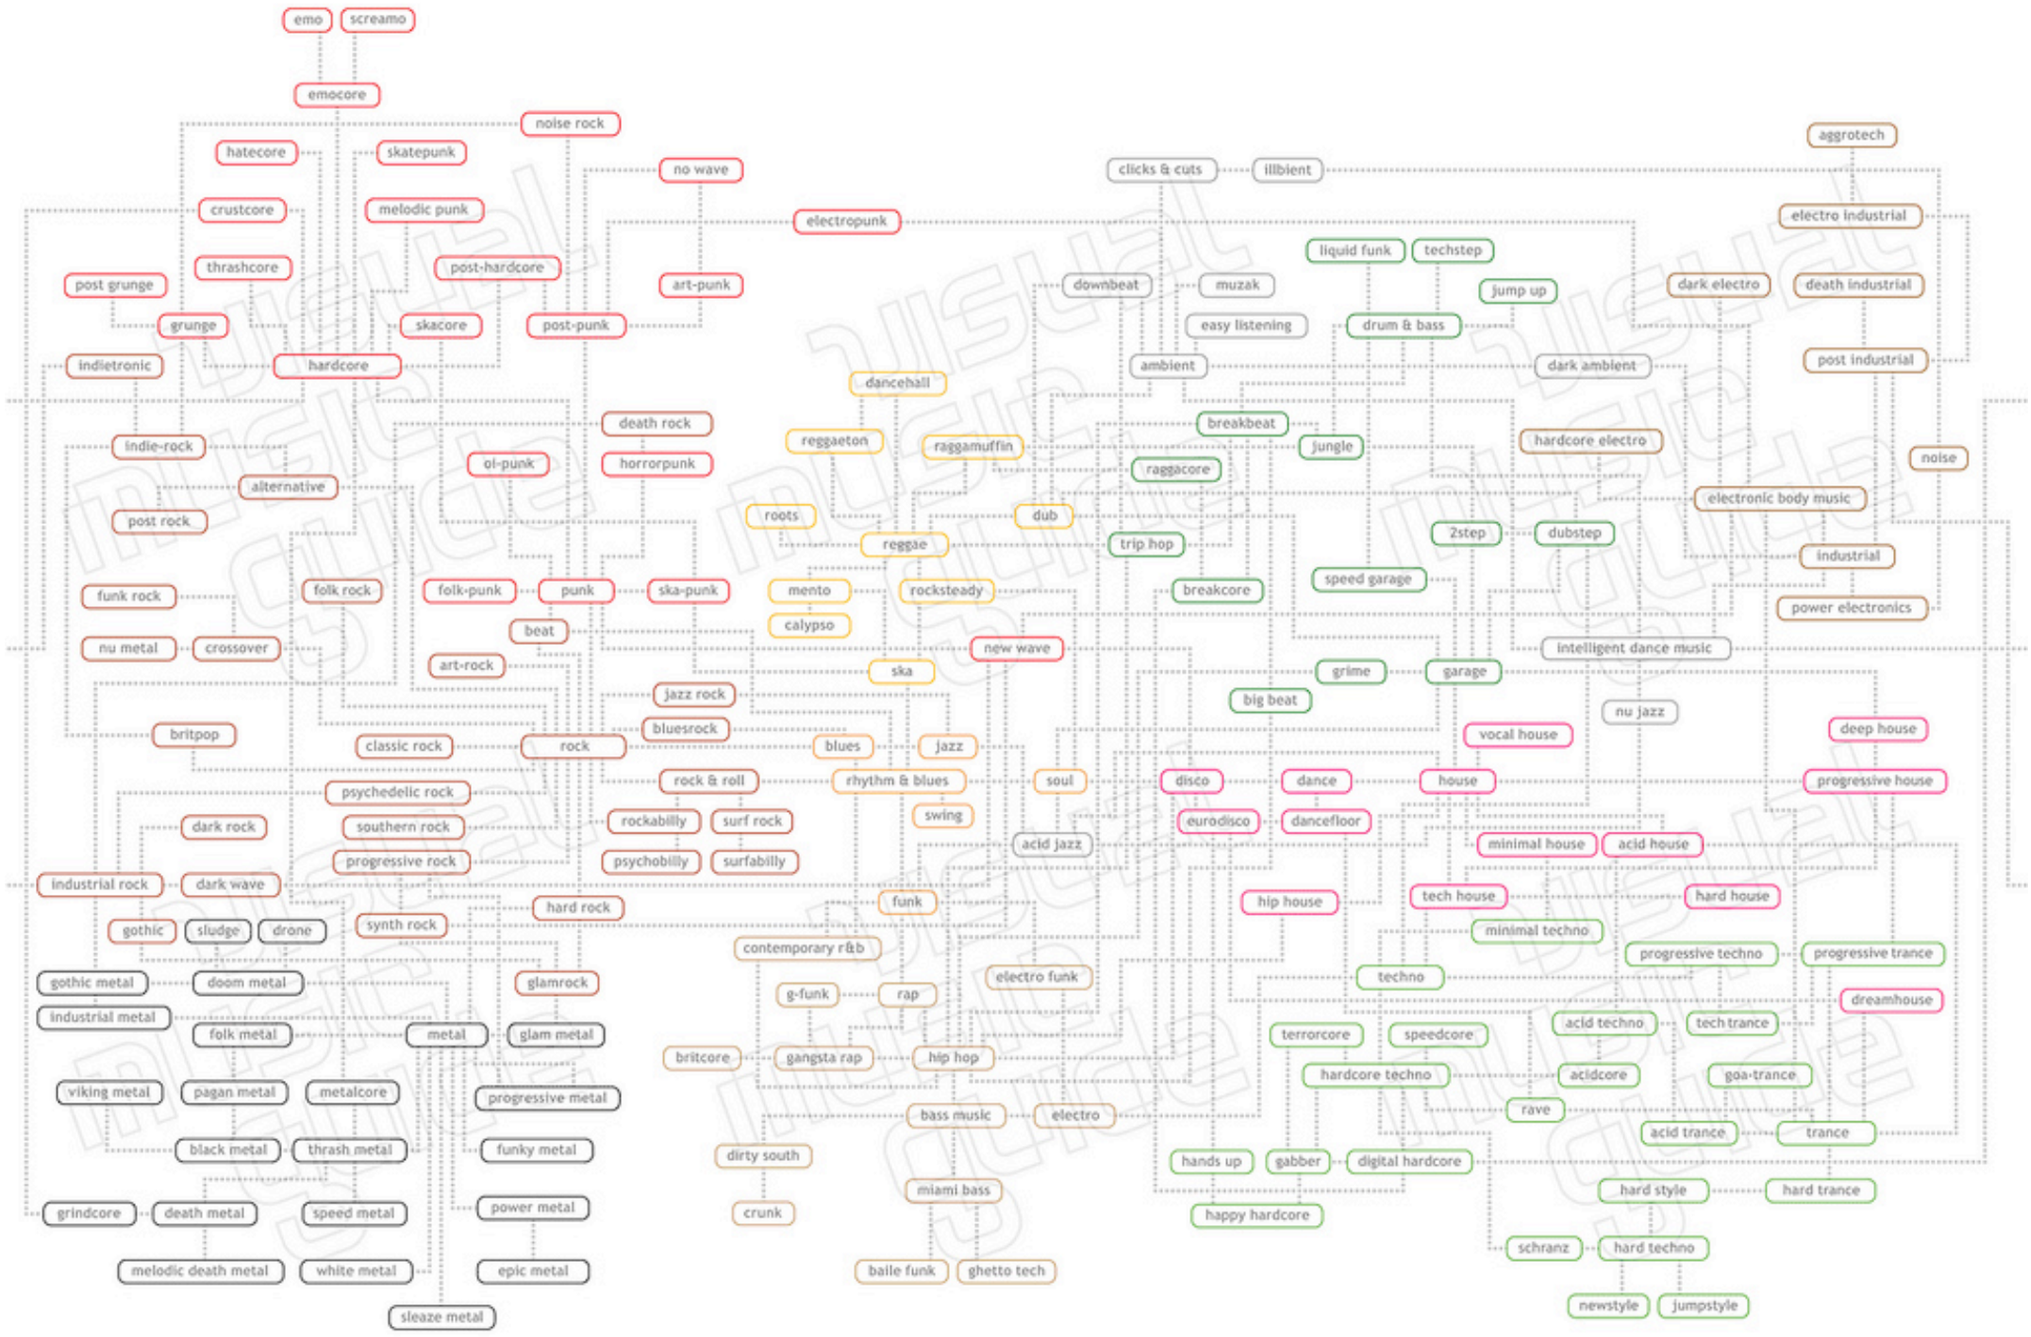

Images (cc) Fabian Schneidmadel

Fabian Schneidmadel - http://www.firutin.de/blog/

Visual Music Guiden

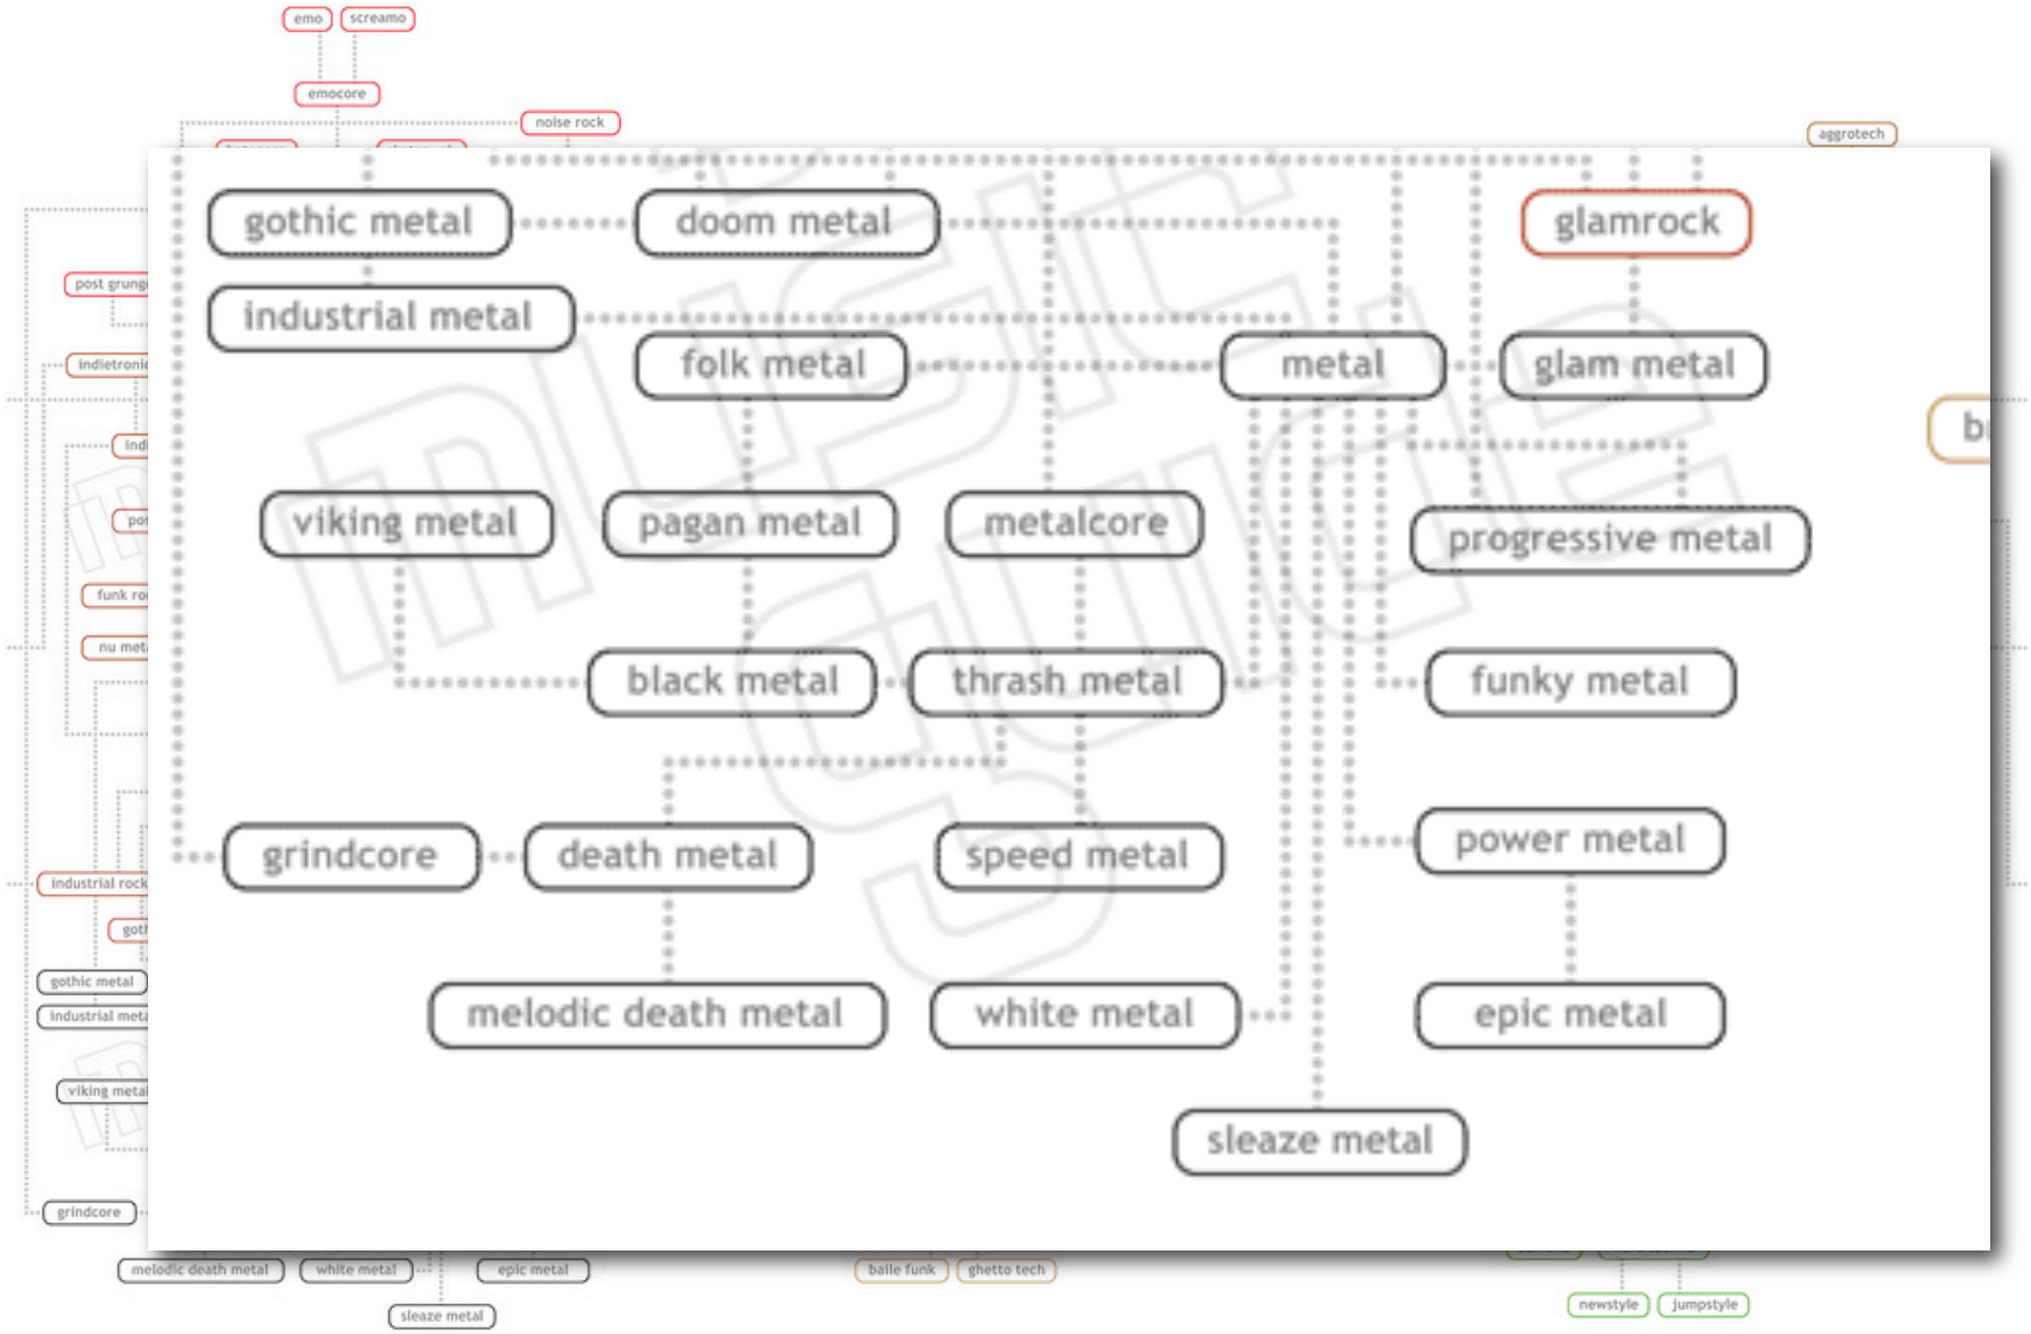

Images (cc) Fabian Schneidmadel

Fabian Schneidmadel - http://www.firutin.de/blog/

Visualizations as a tree

Visualizations as a tree

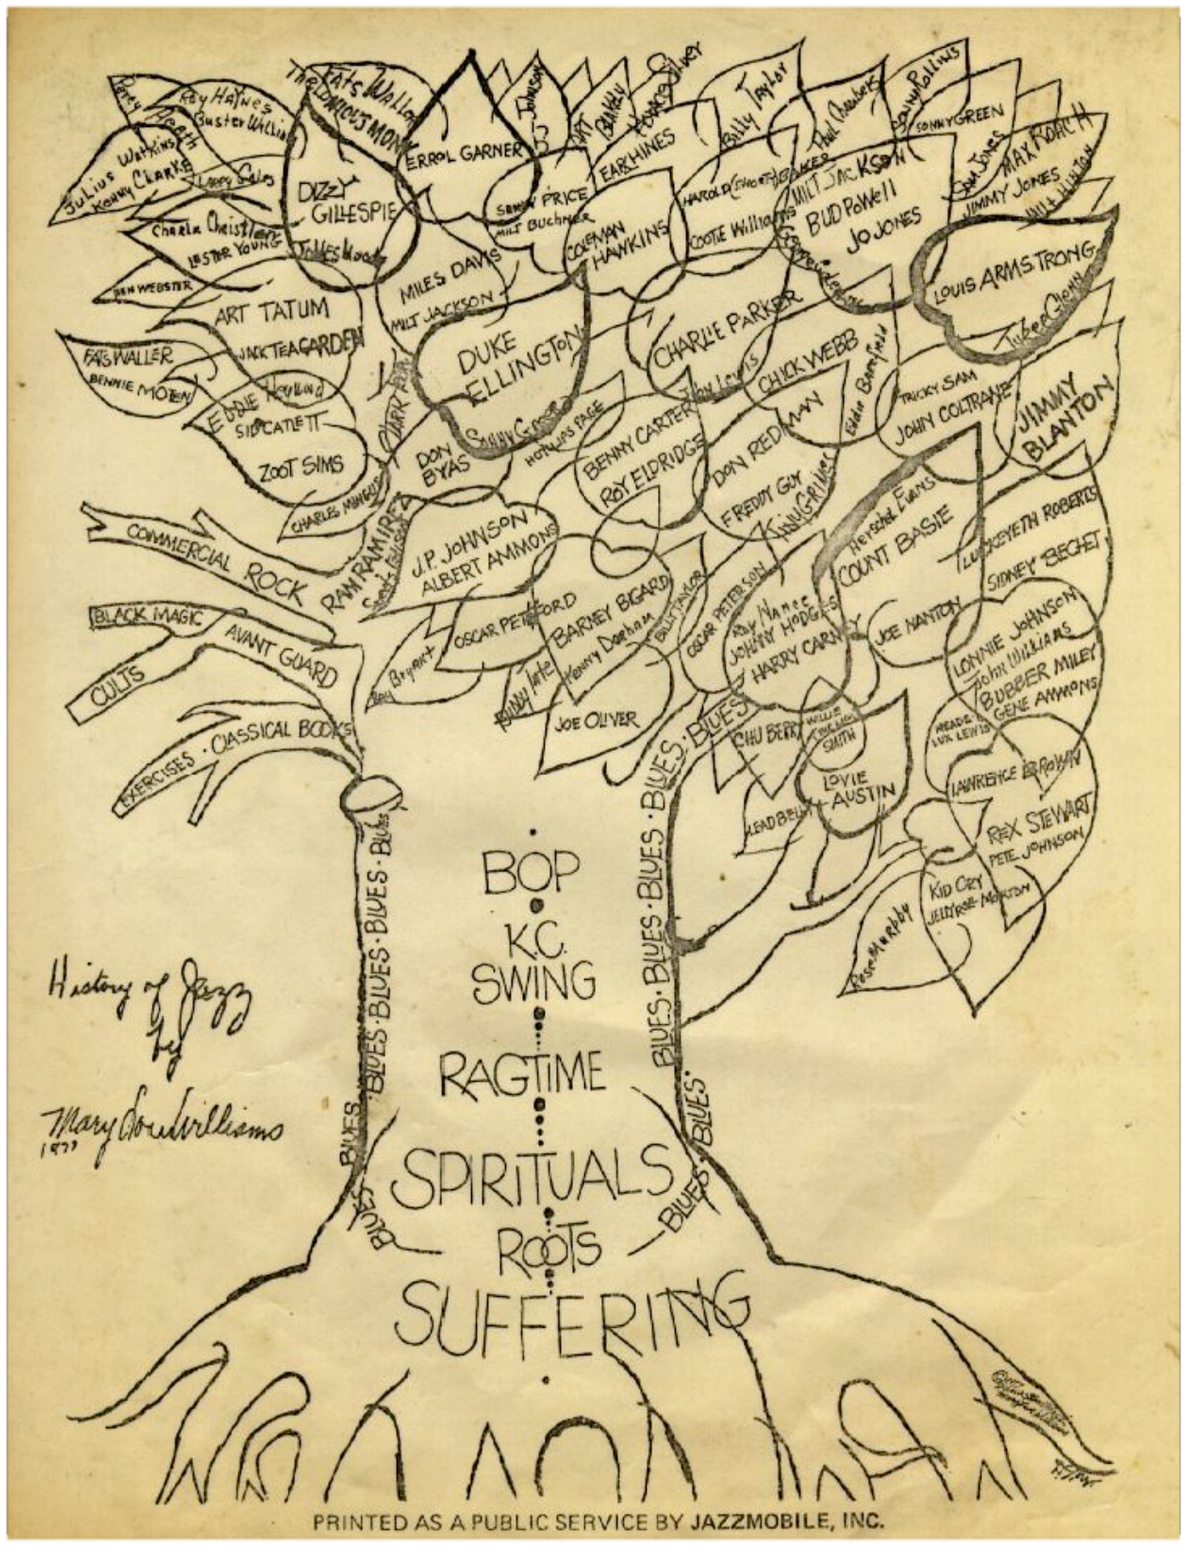

Visualizations as a tree

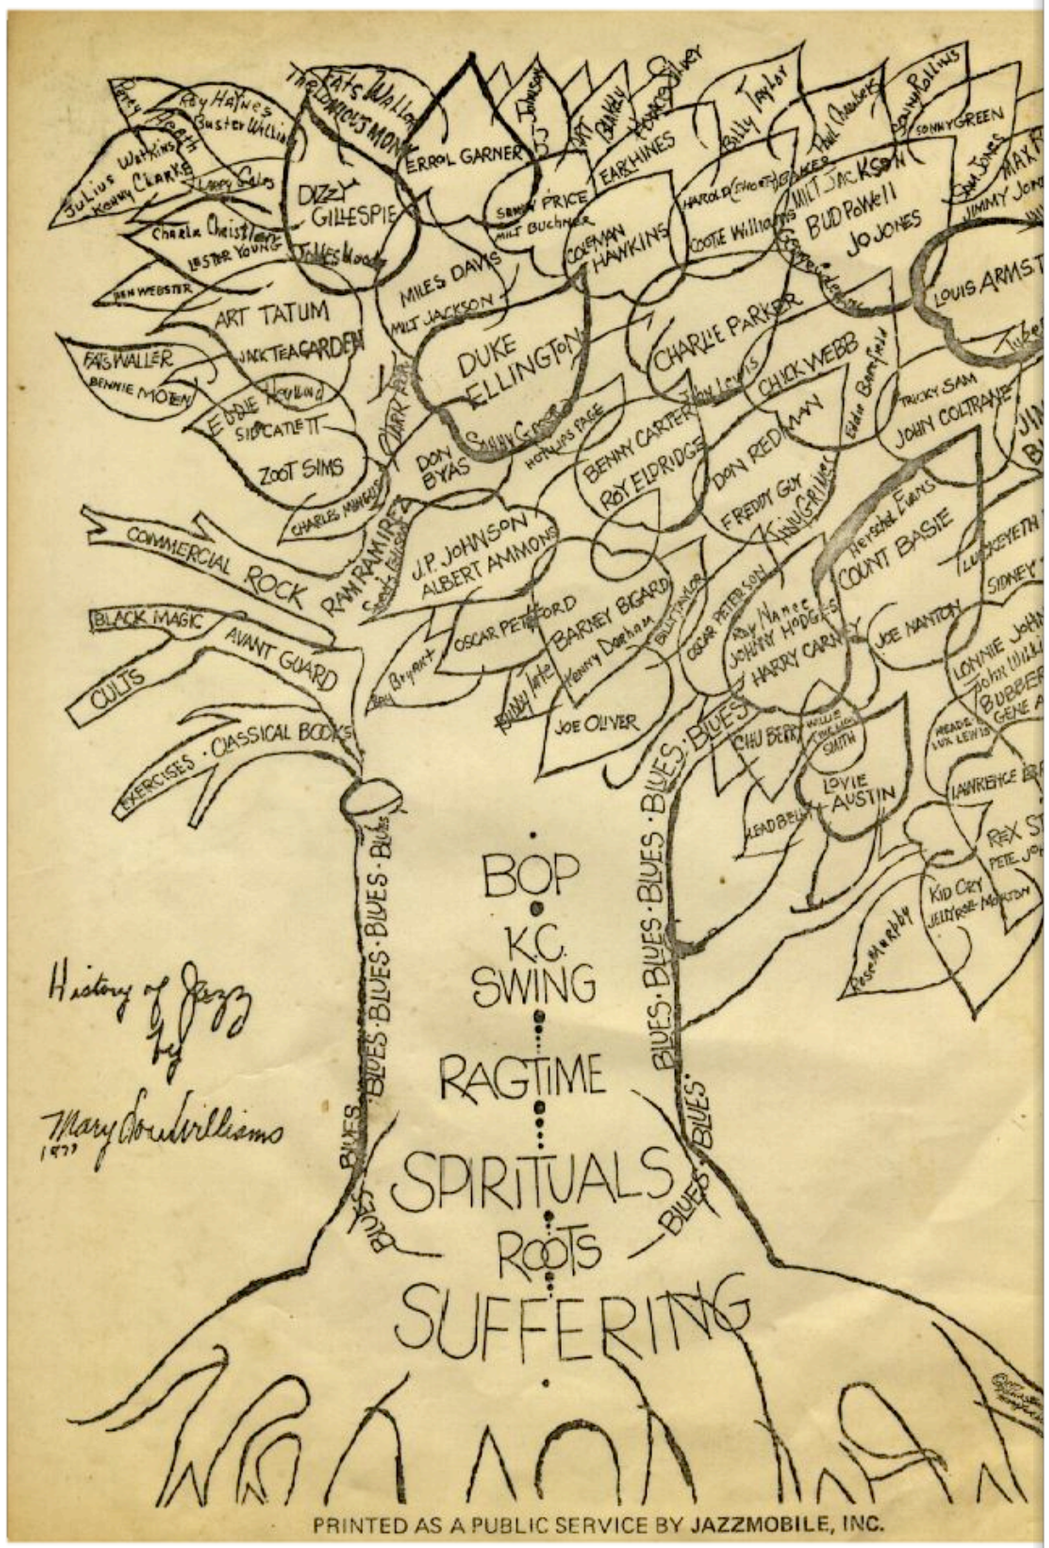

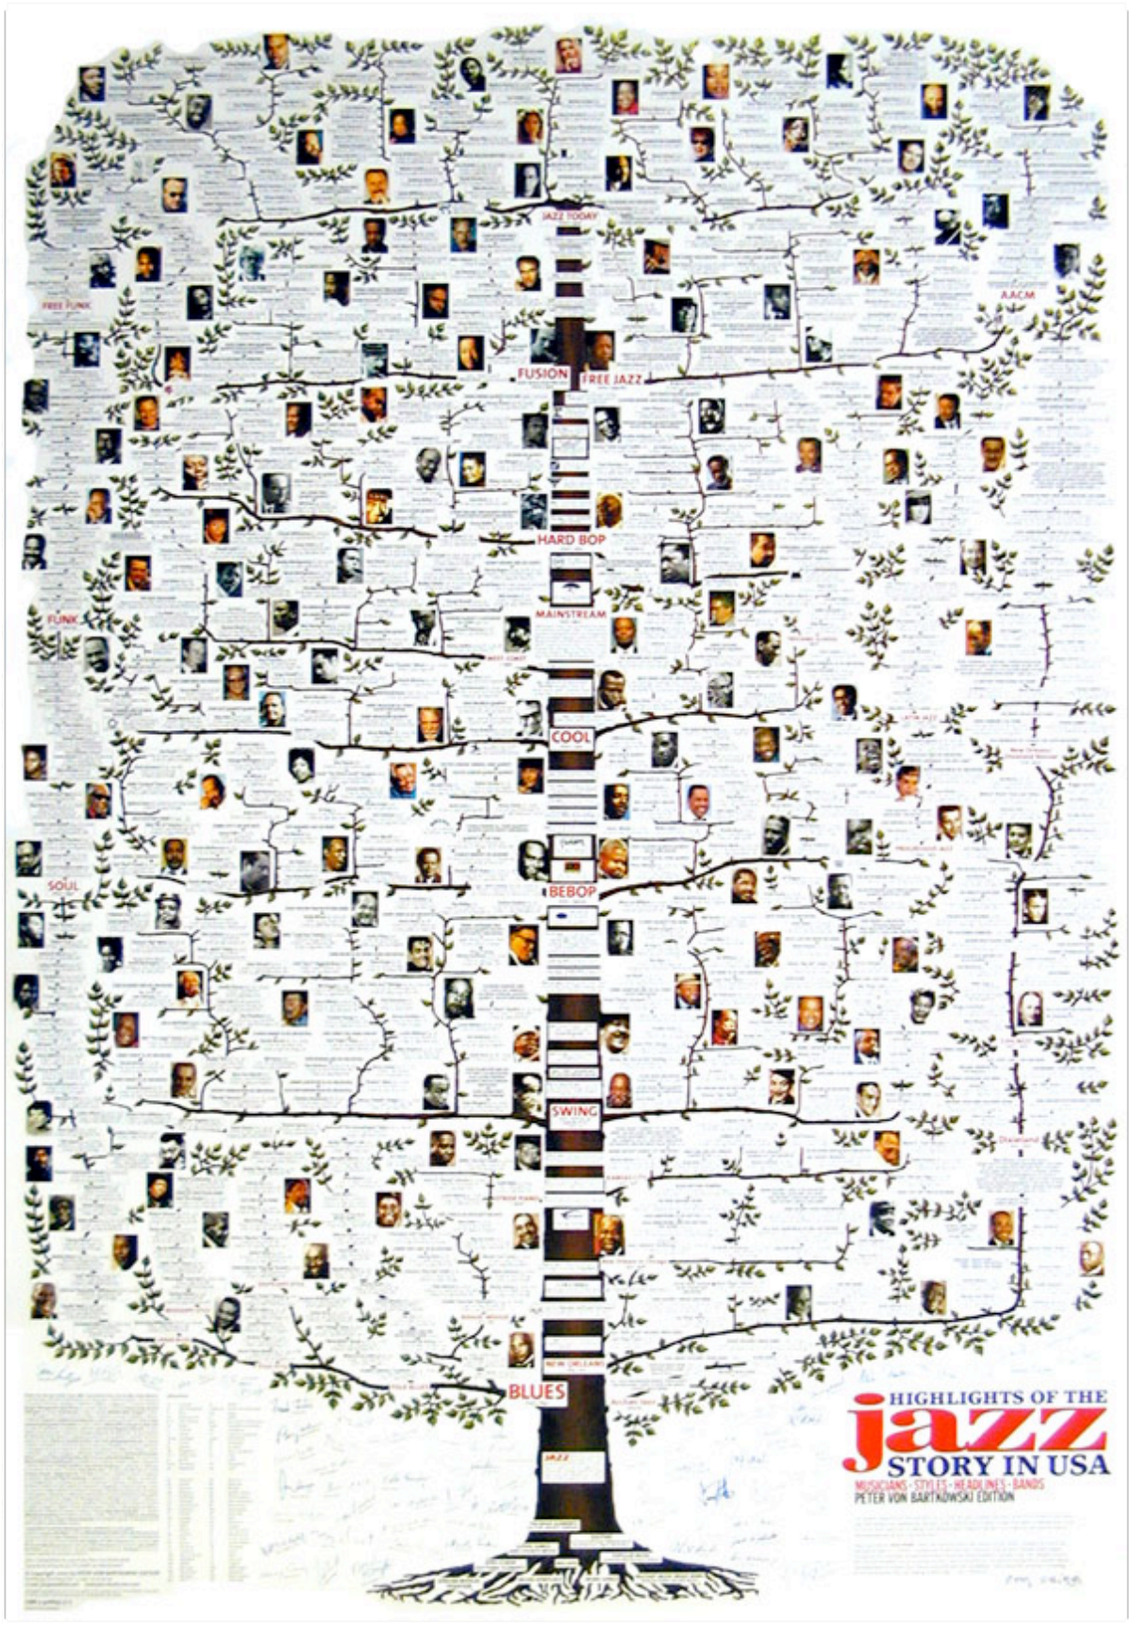

Visualizations as a tree

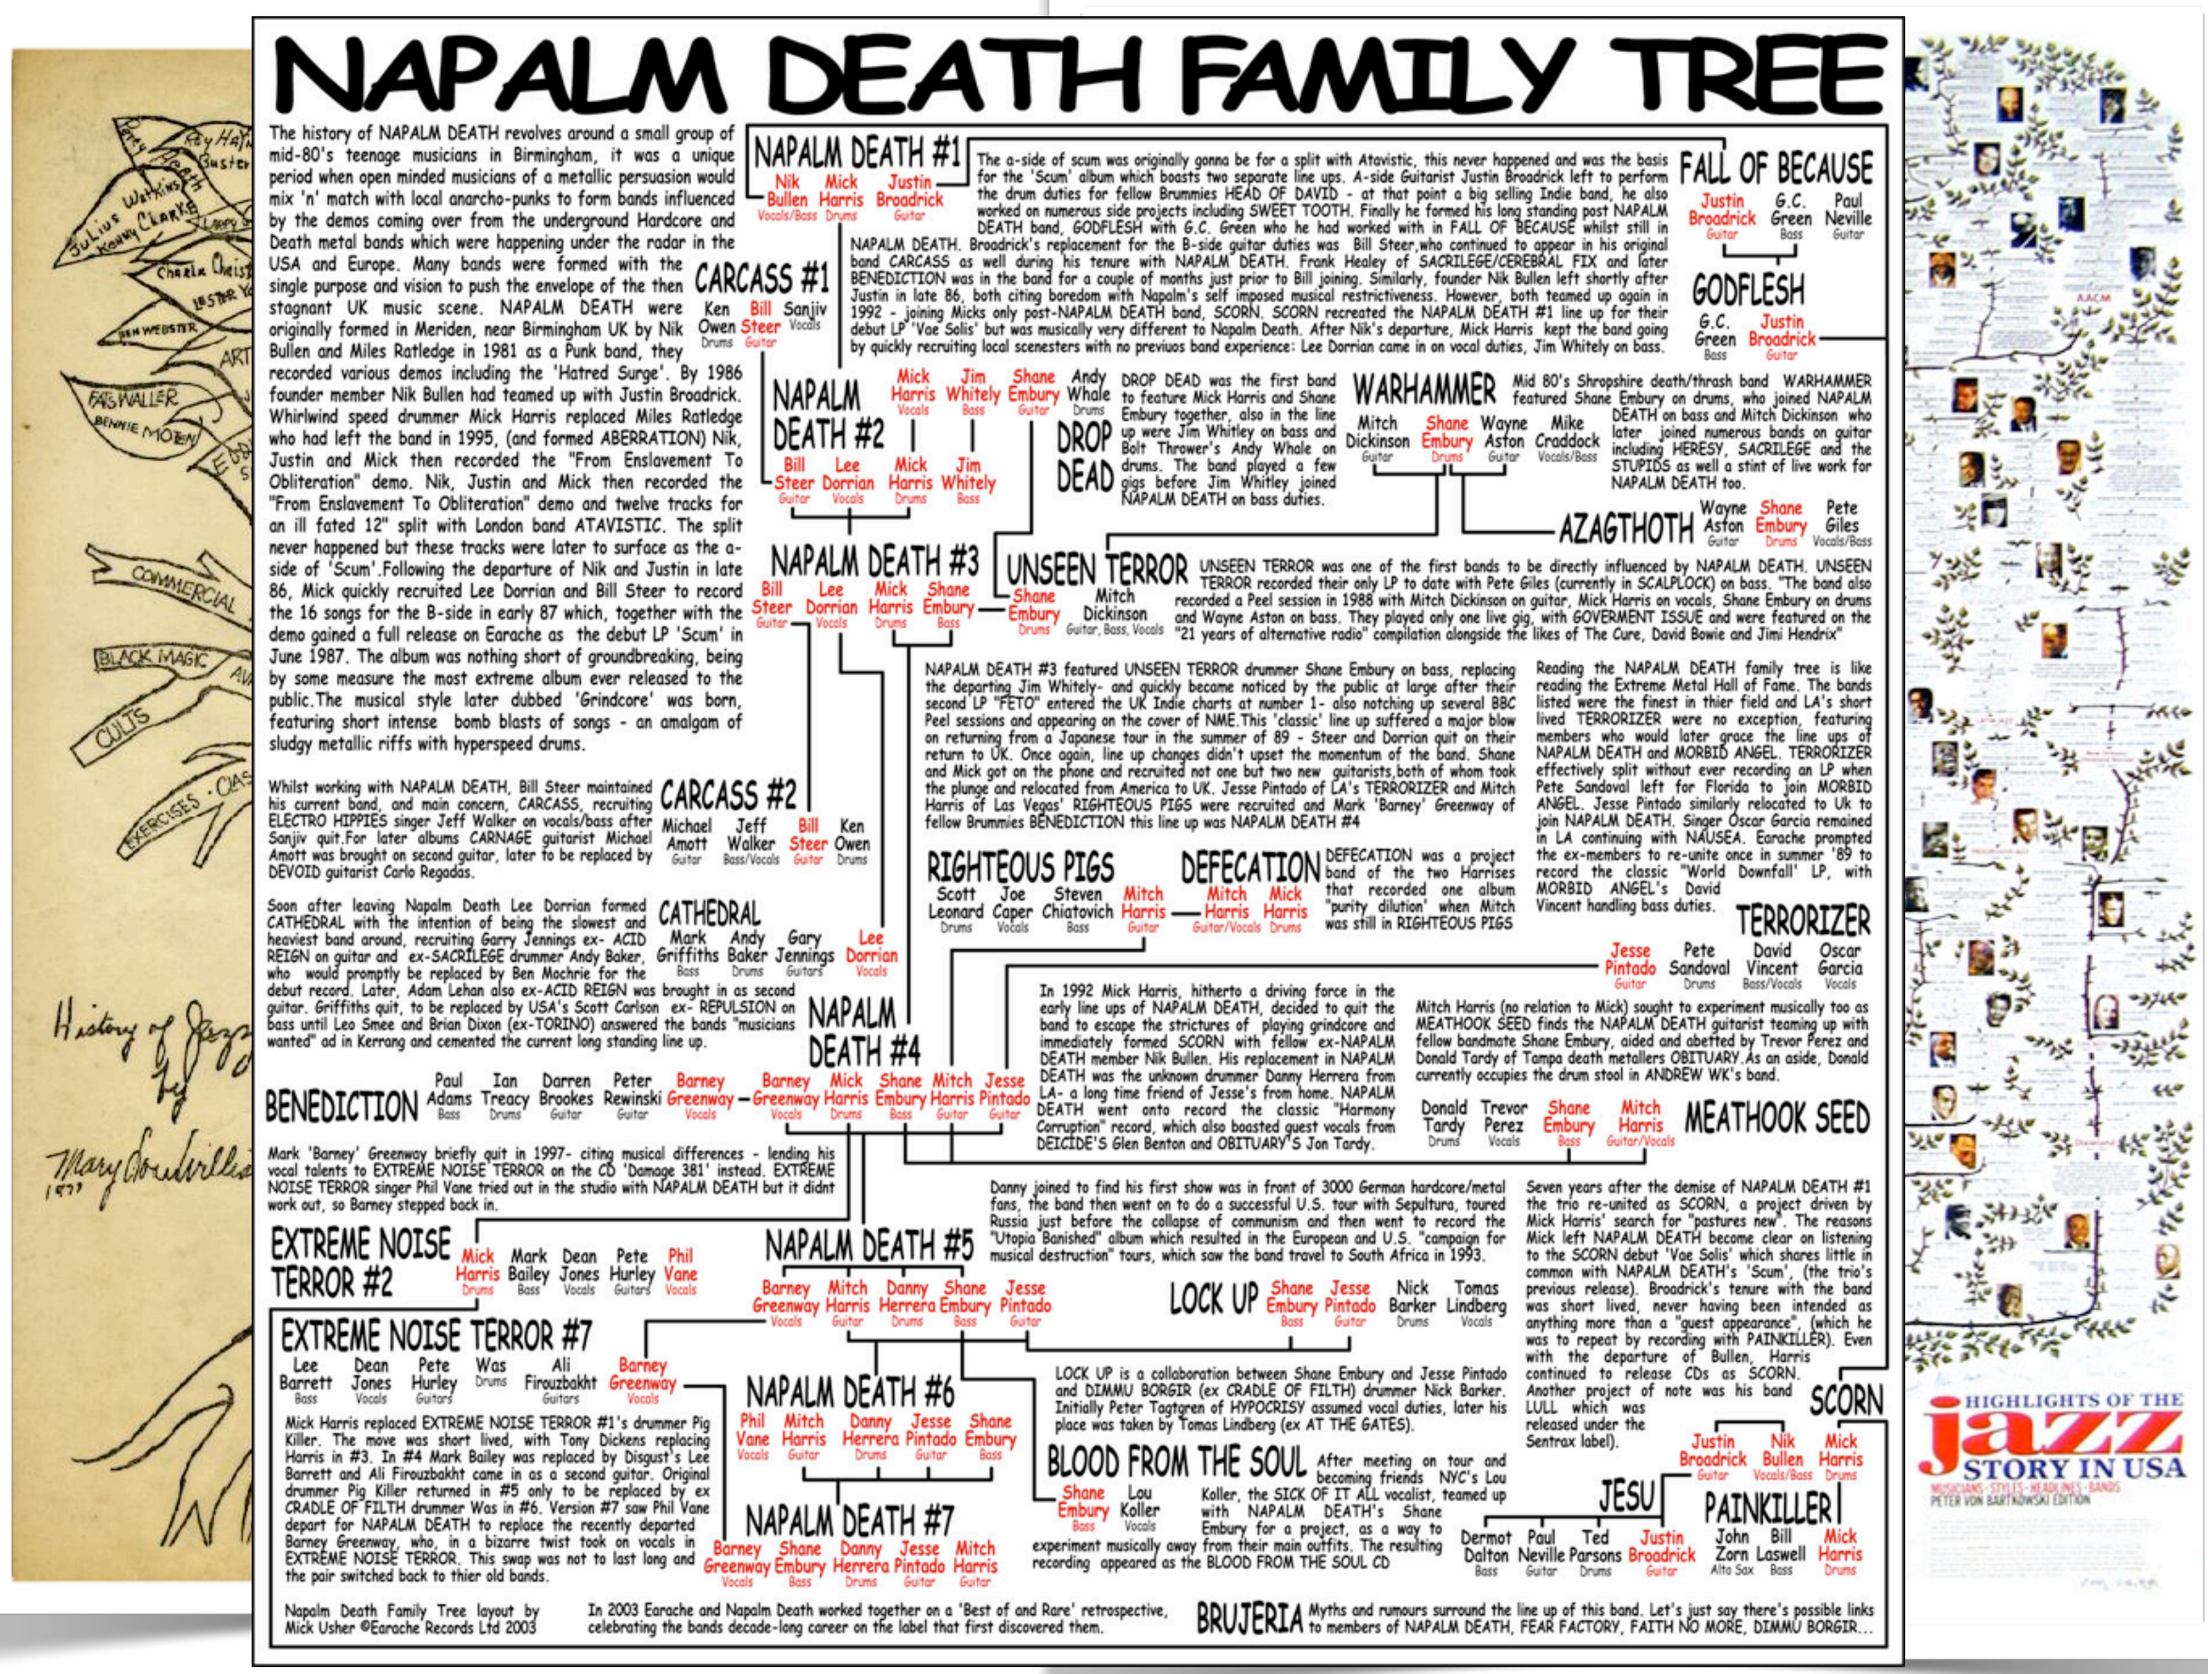

Handmade Genre Maps and Trees

Handmade Genre Maps and Trees

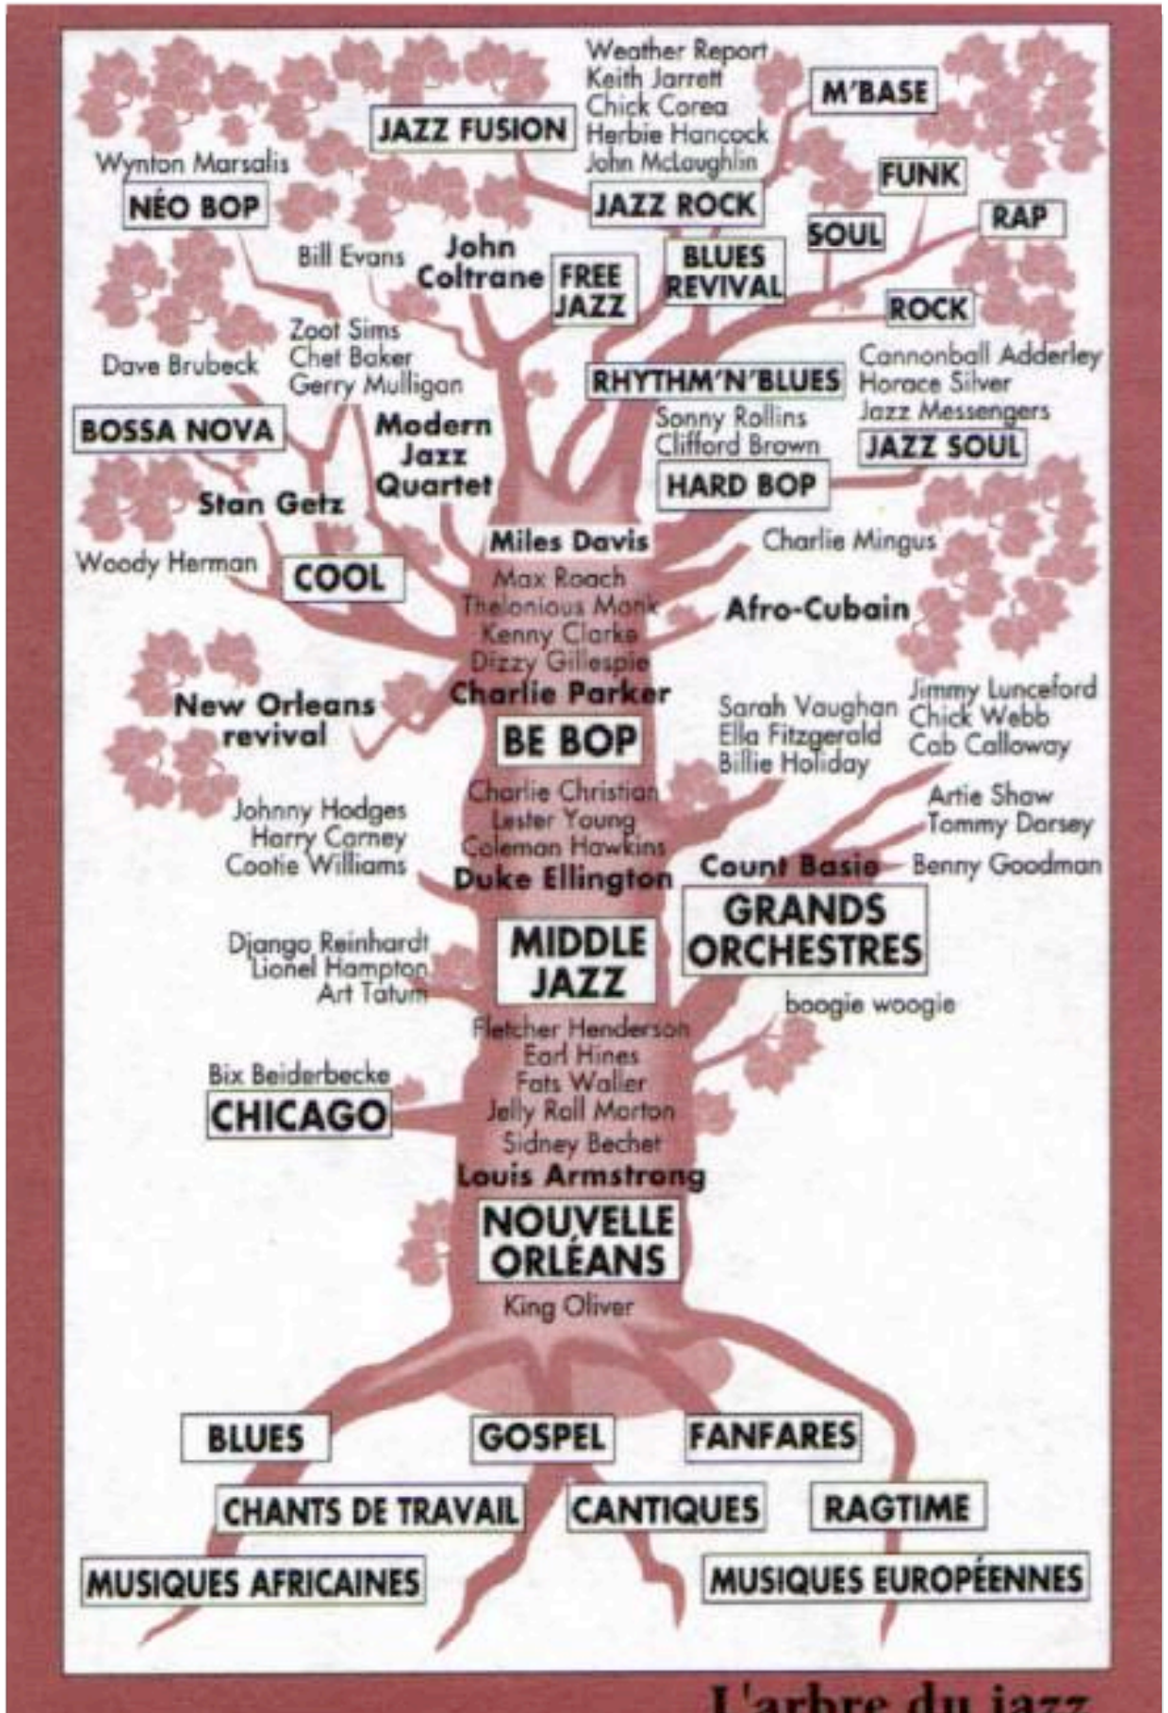

L'arbre du jazz

Handmade Genre Maps and Trees

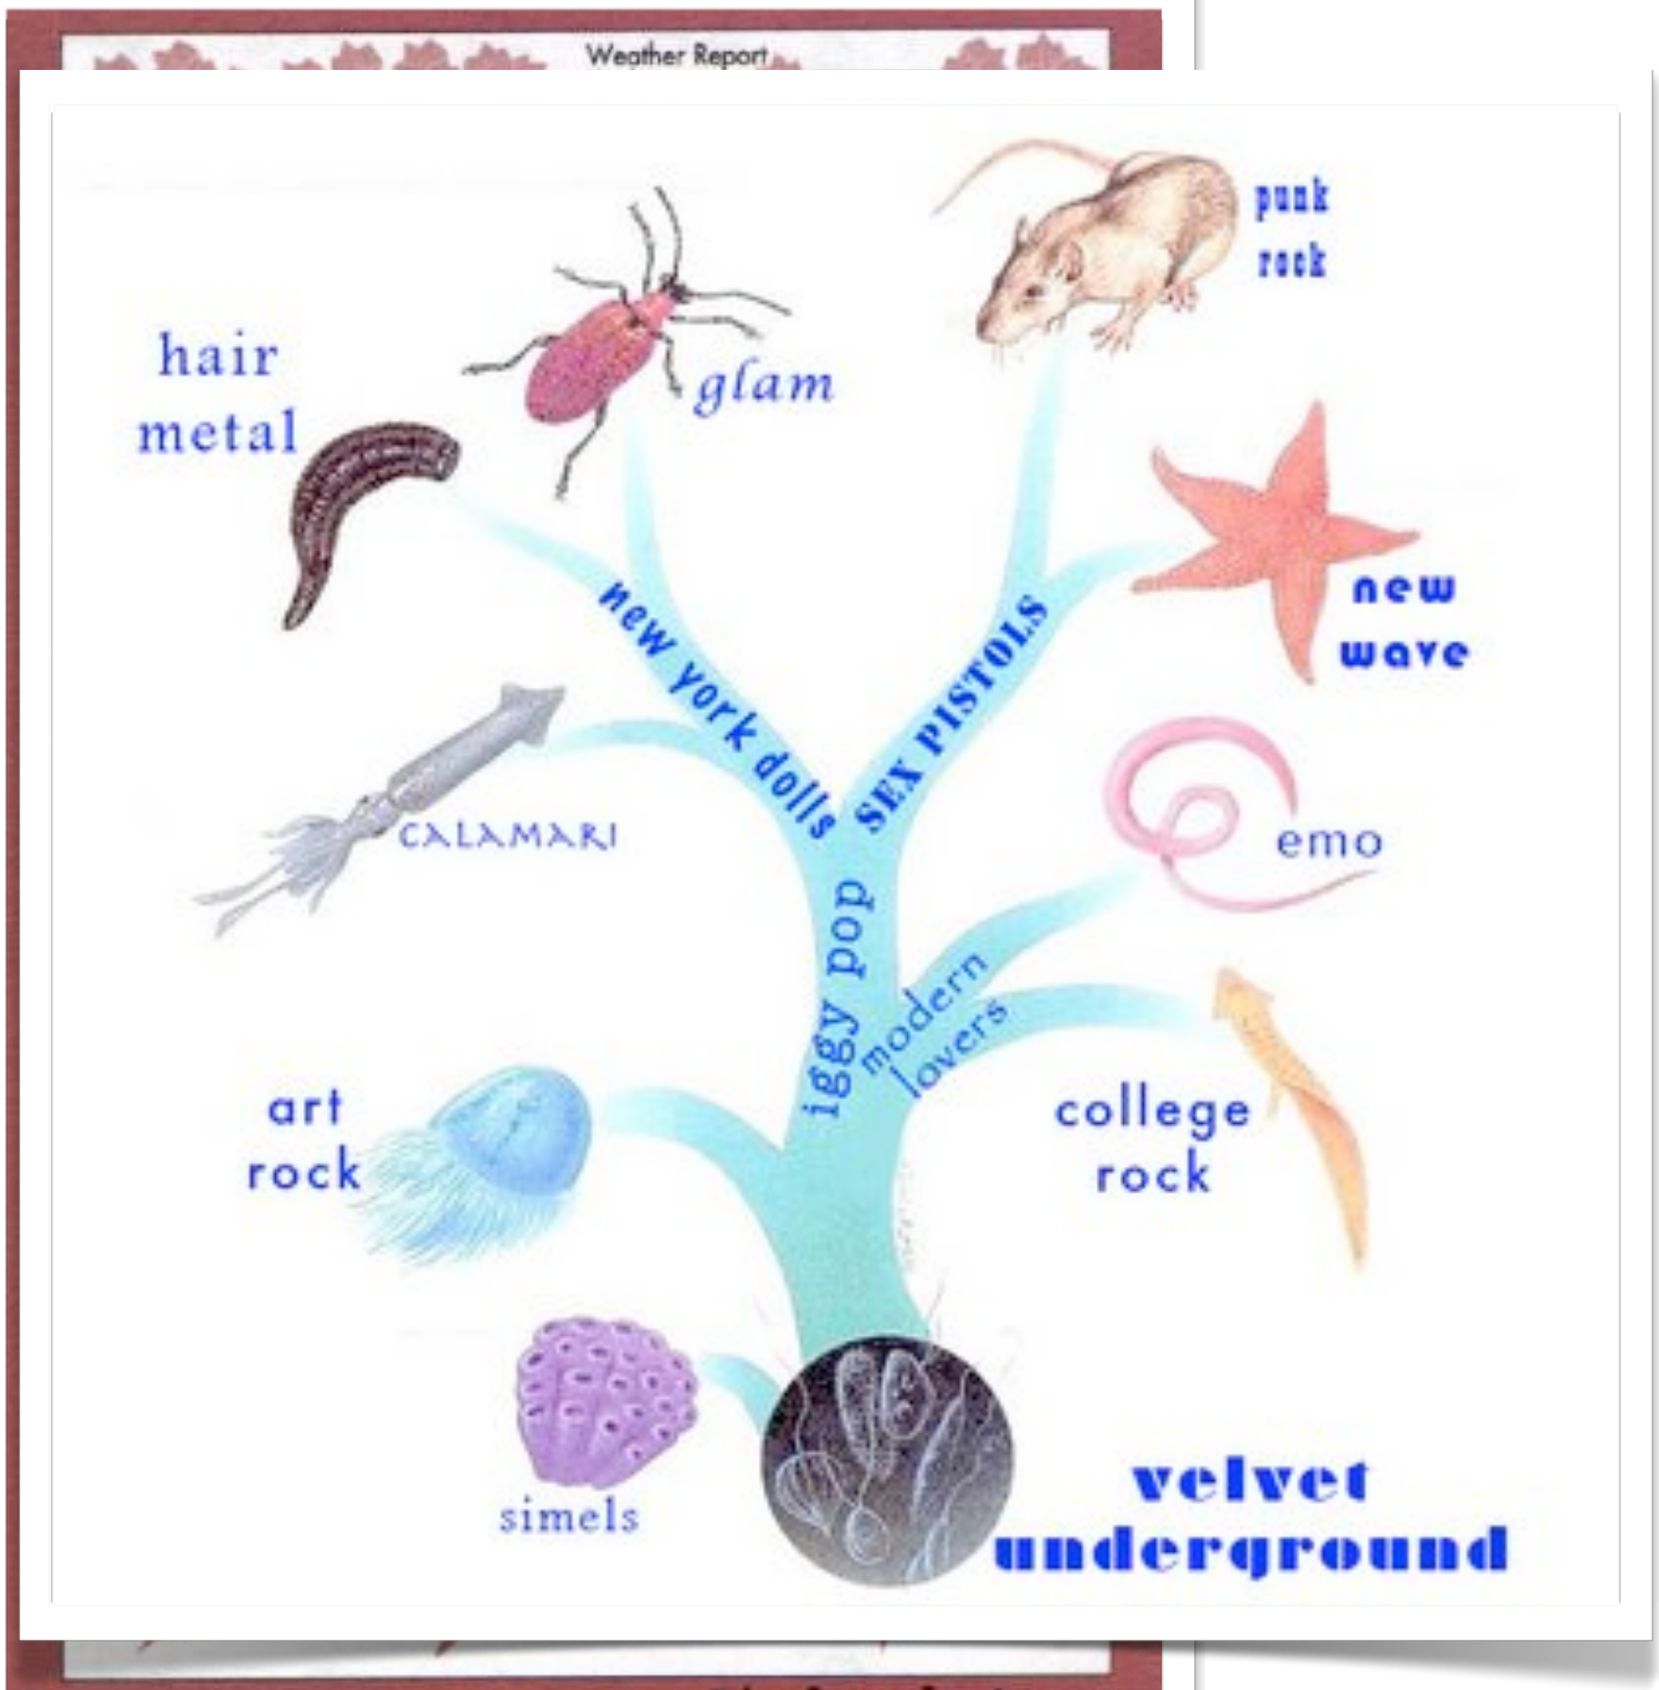

Handmade Genre Maps and Trees

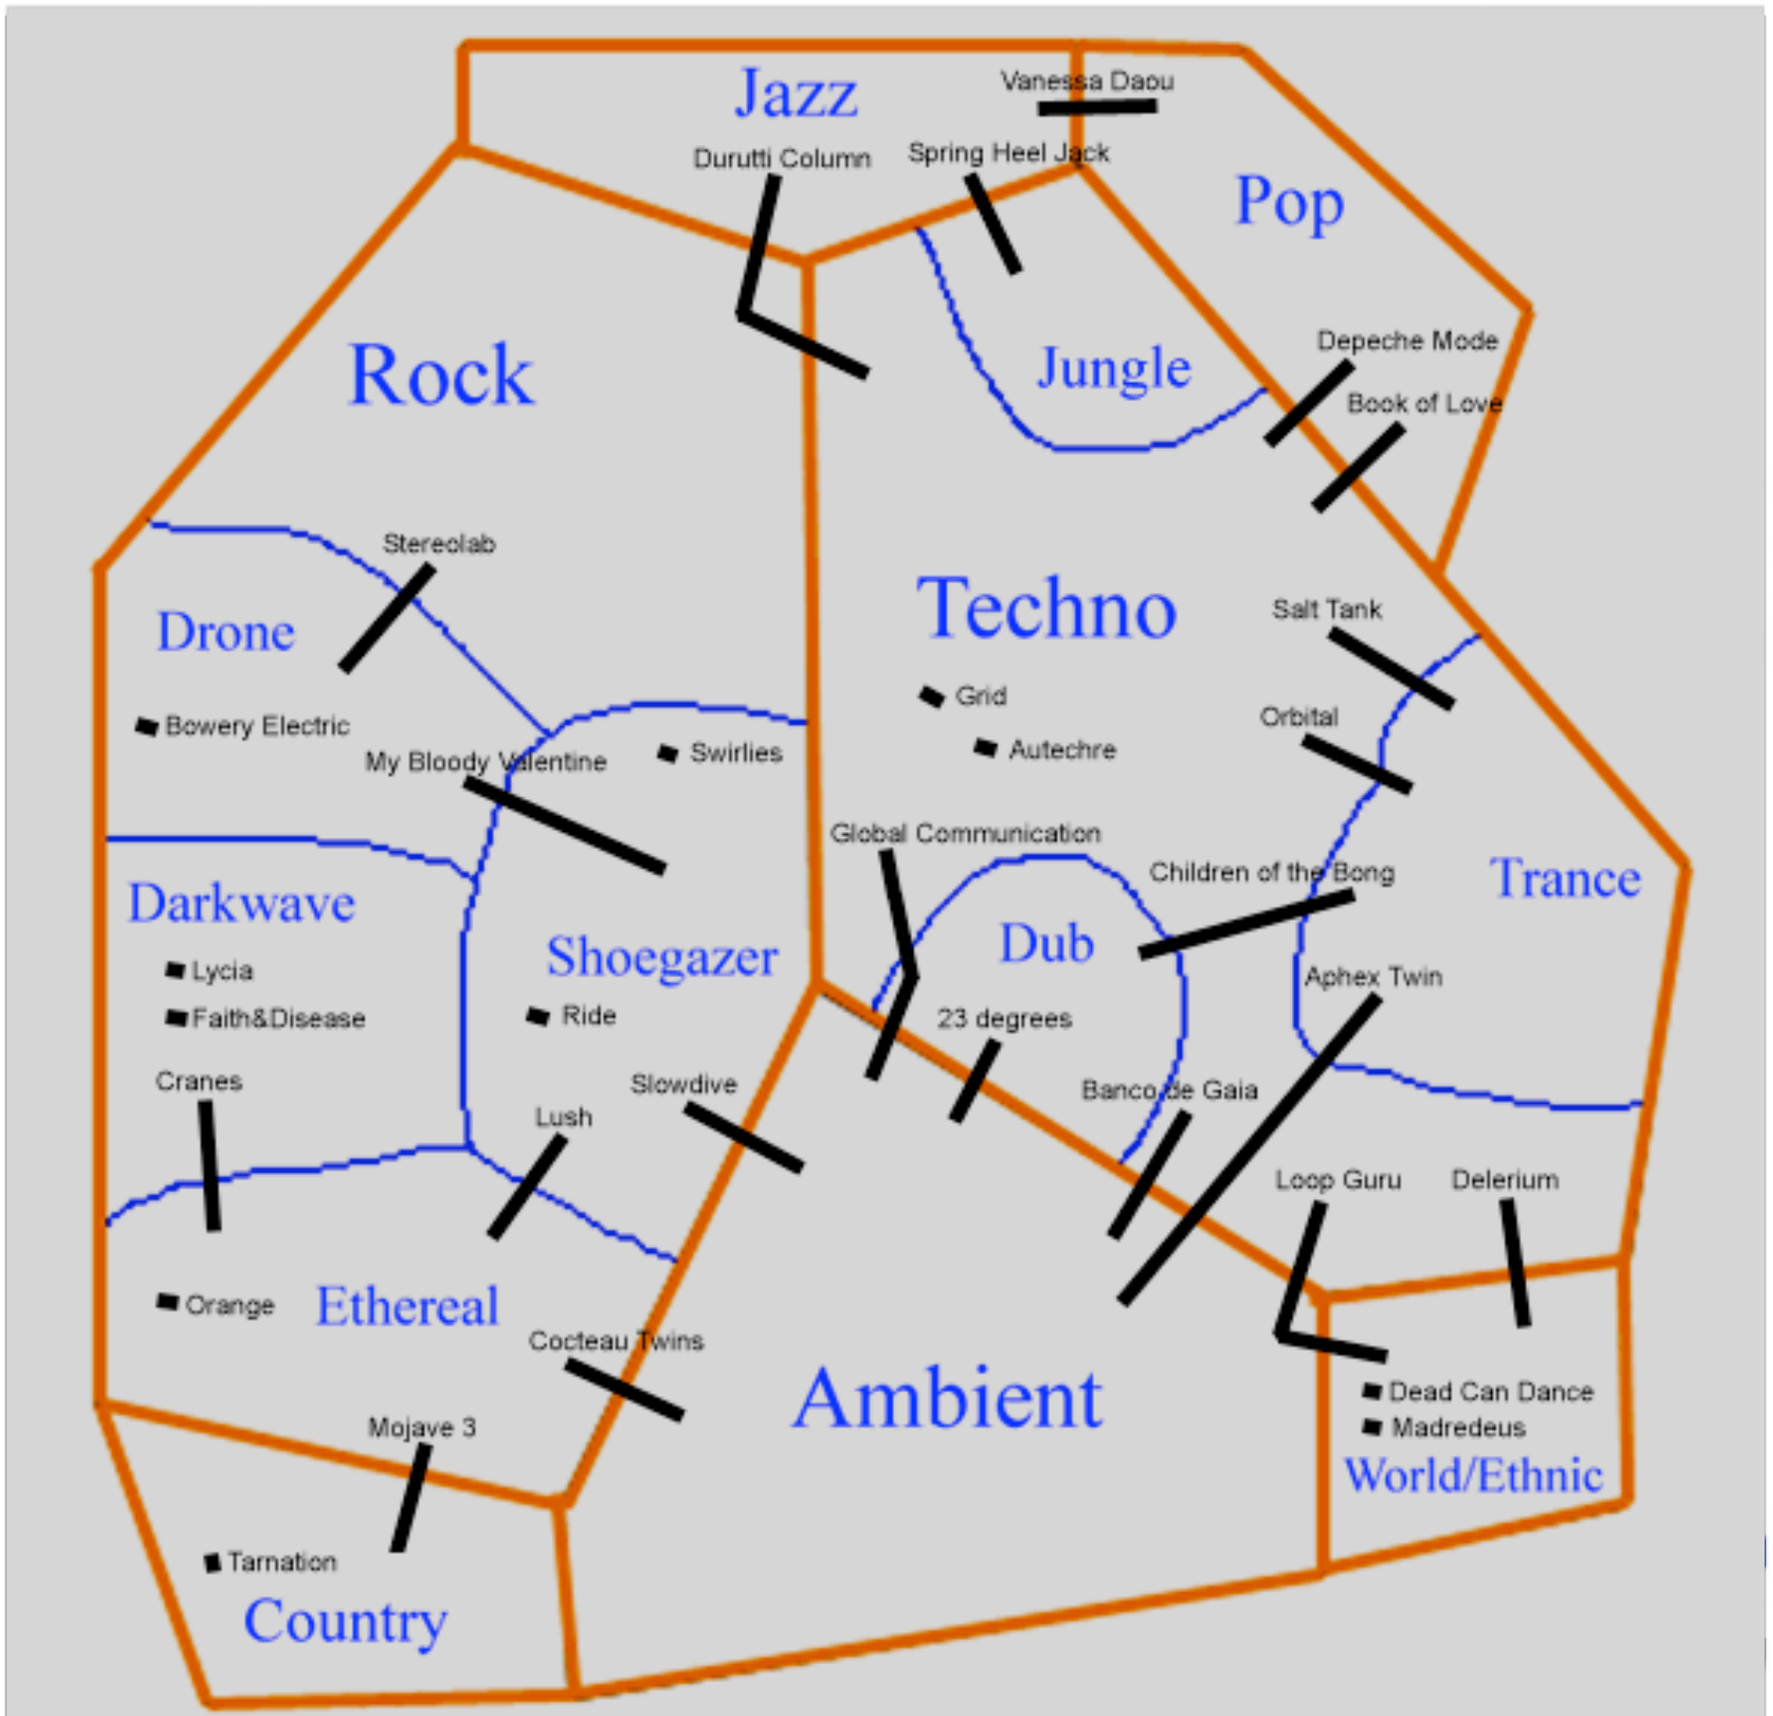

Handmade Genre Maps and Trees

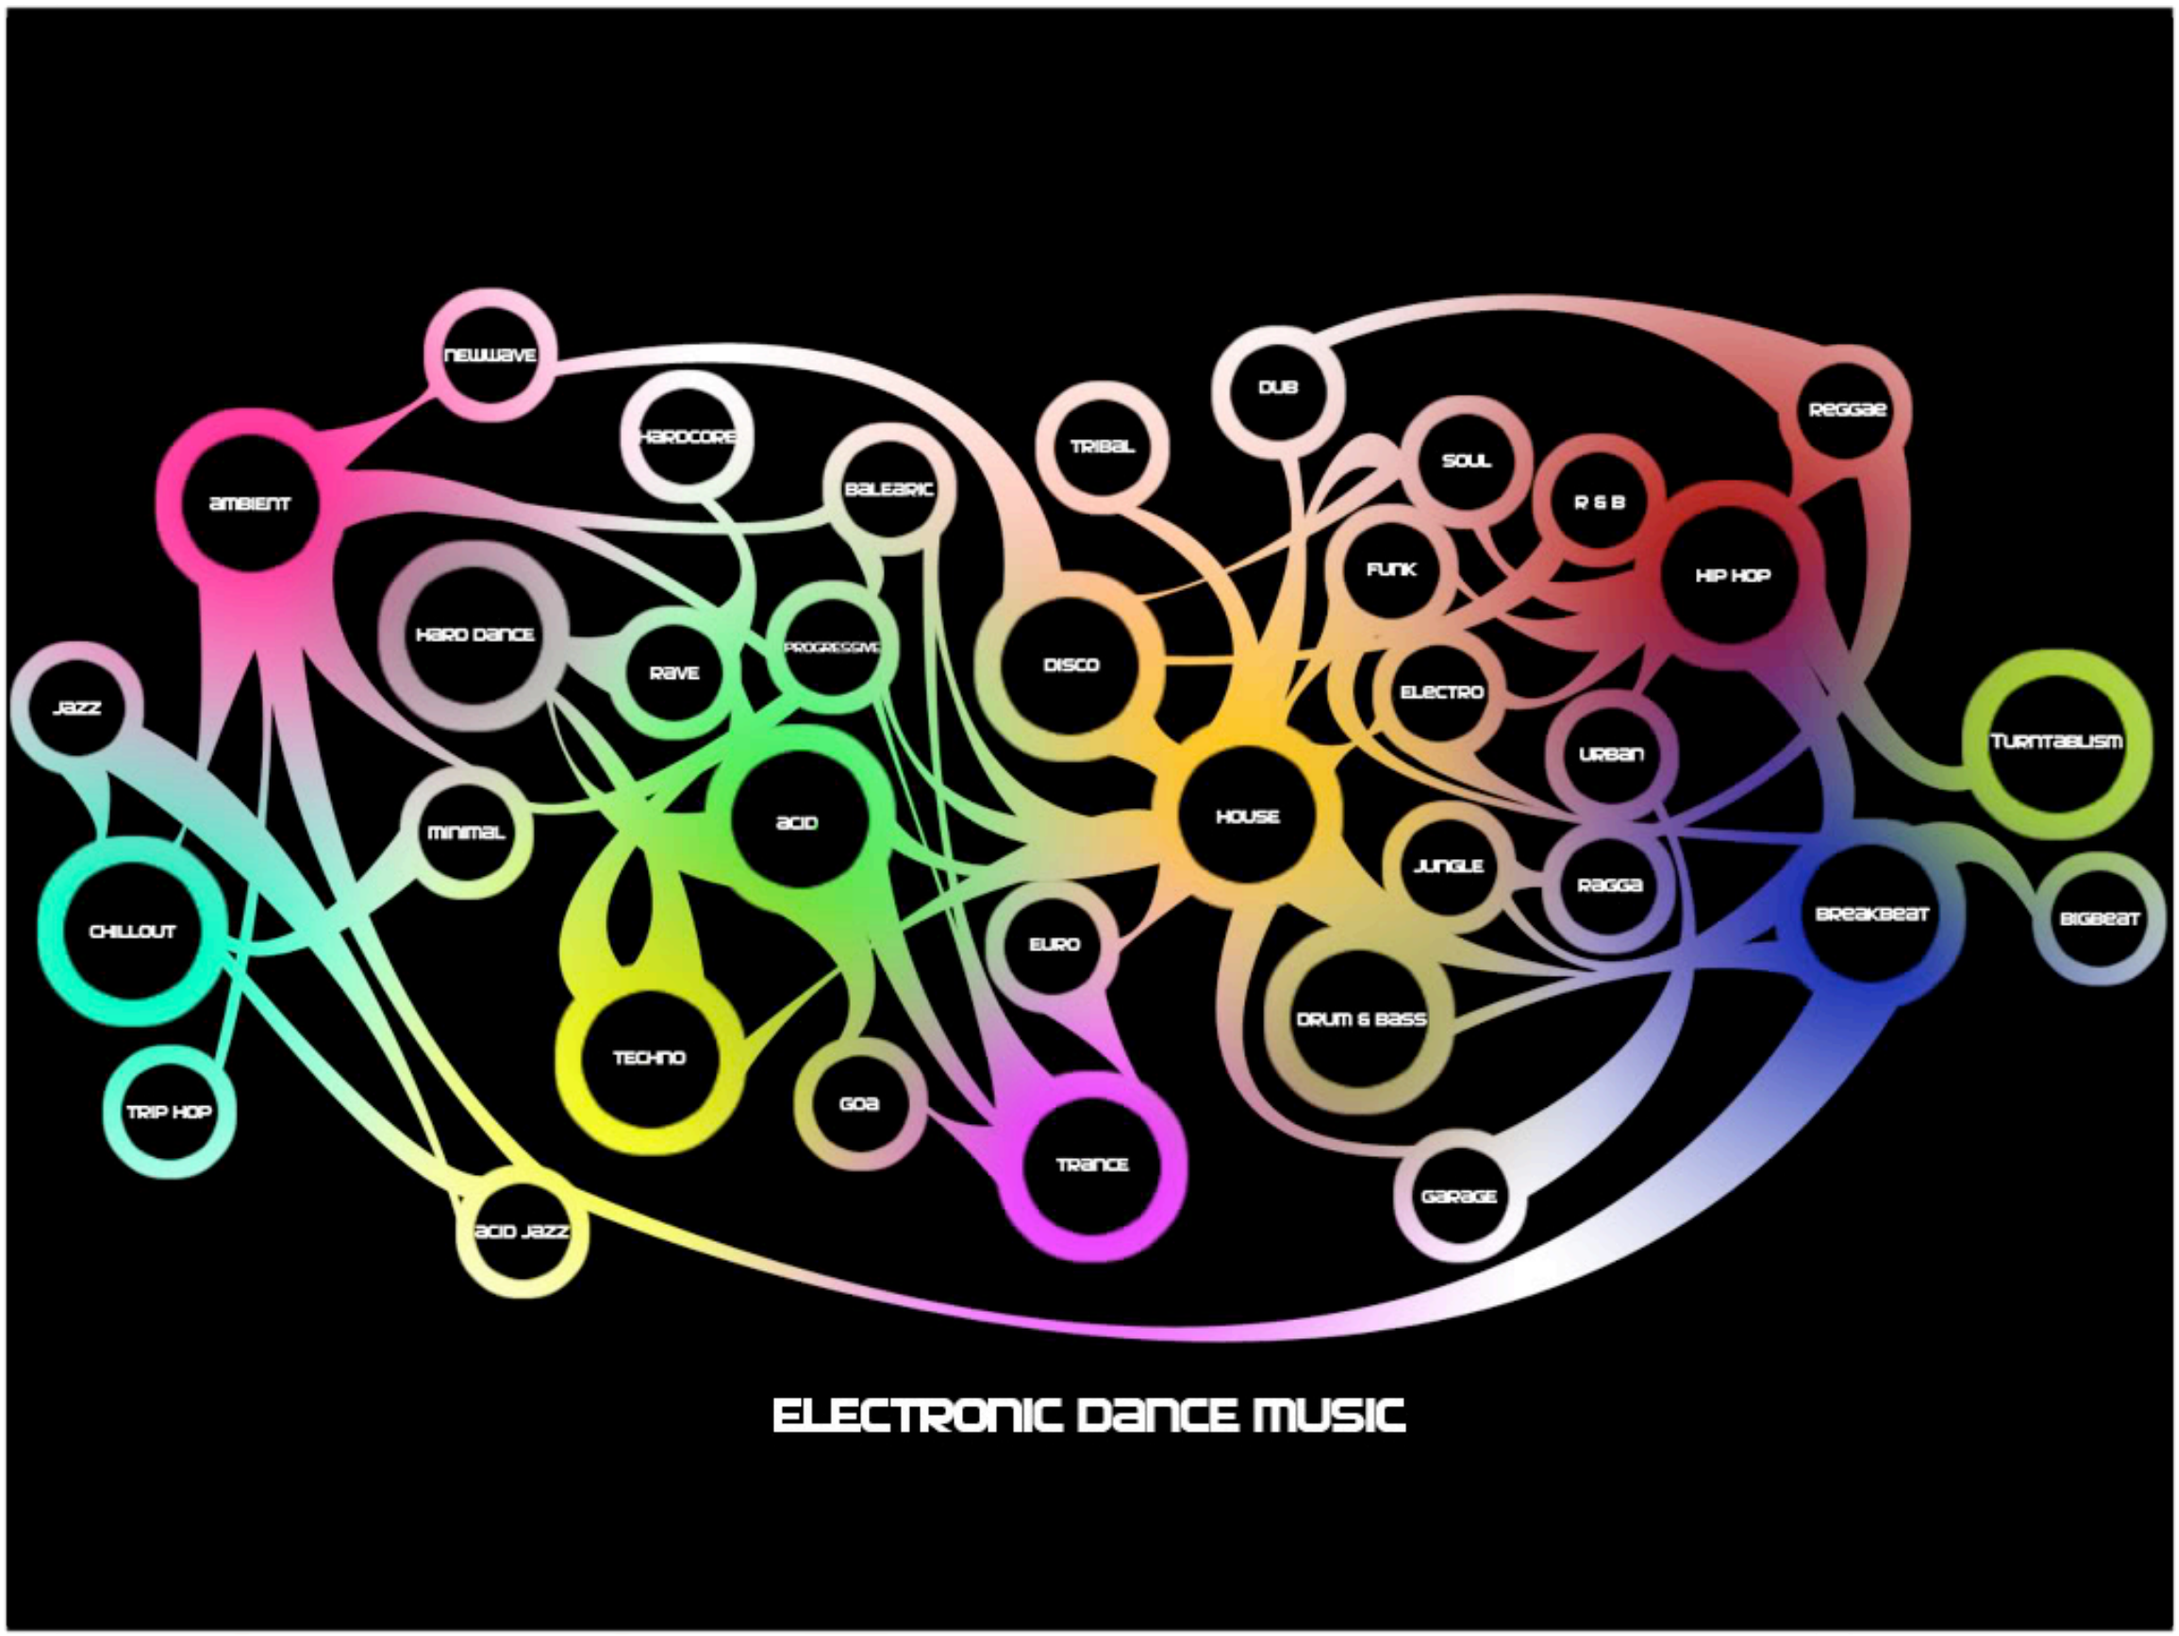

Handmade Genre Maps and Trees

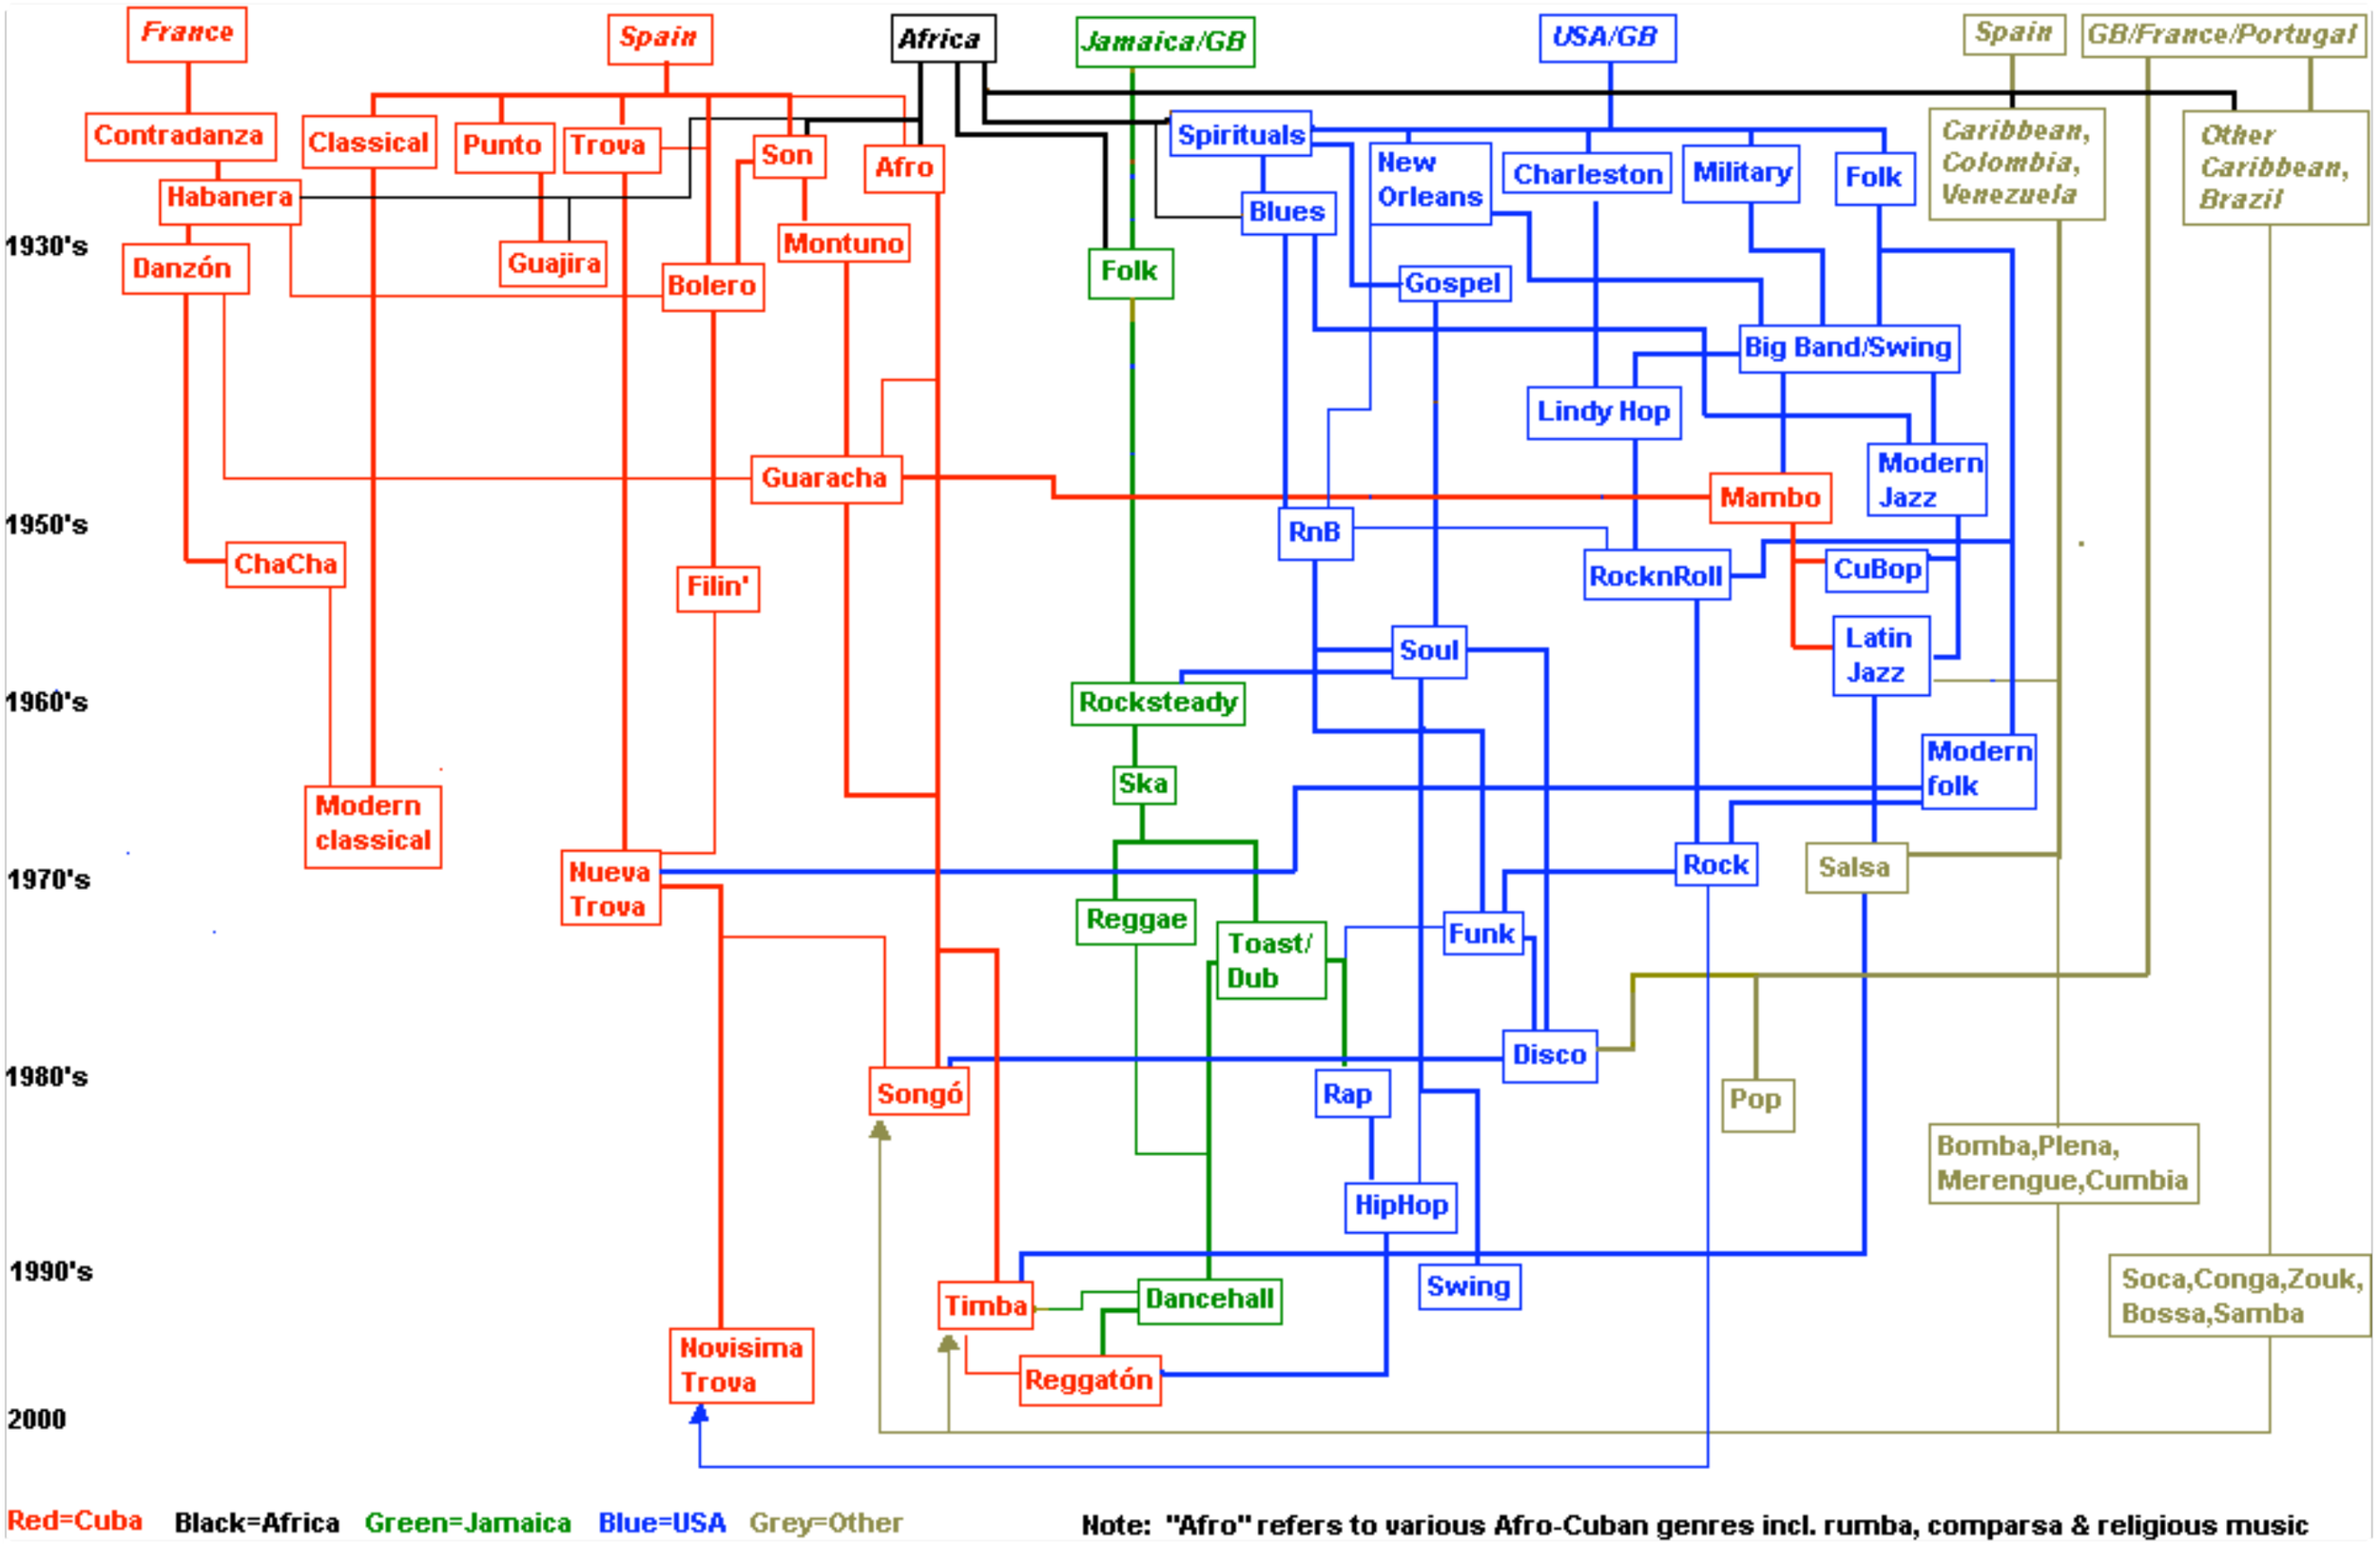

Handmade Genre Maps and Trees

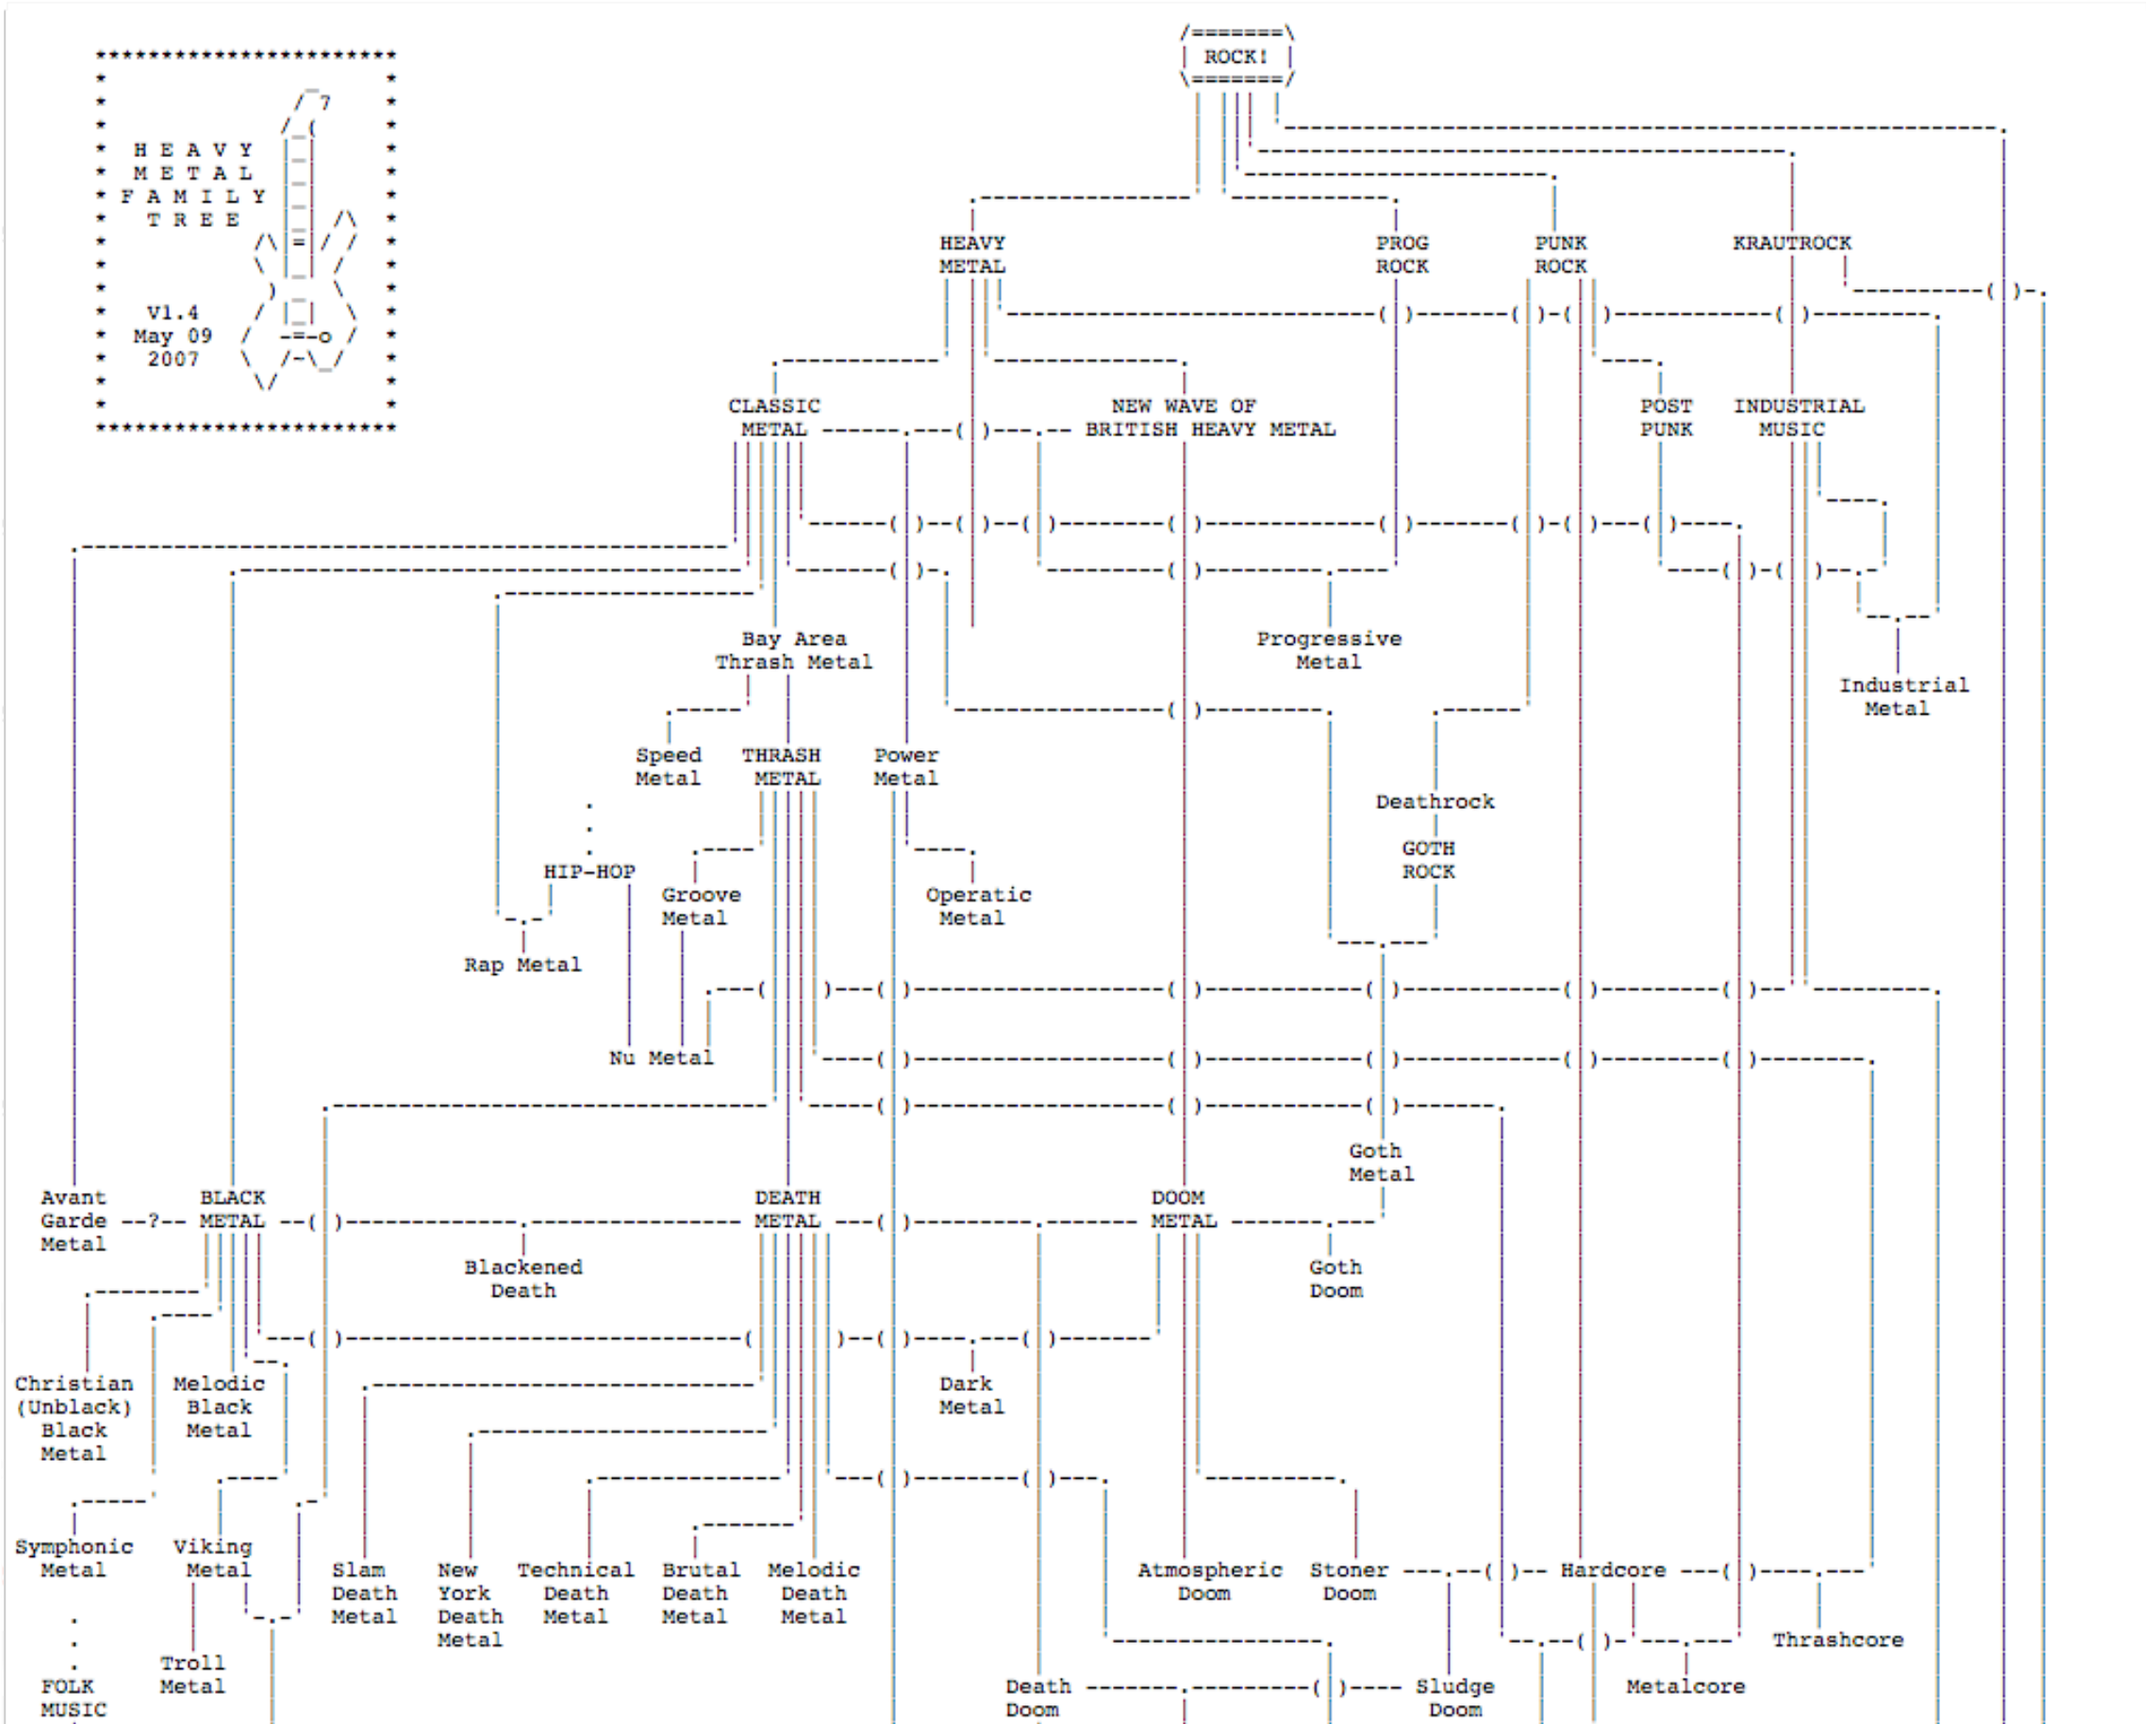

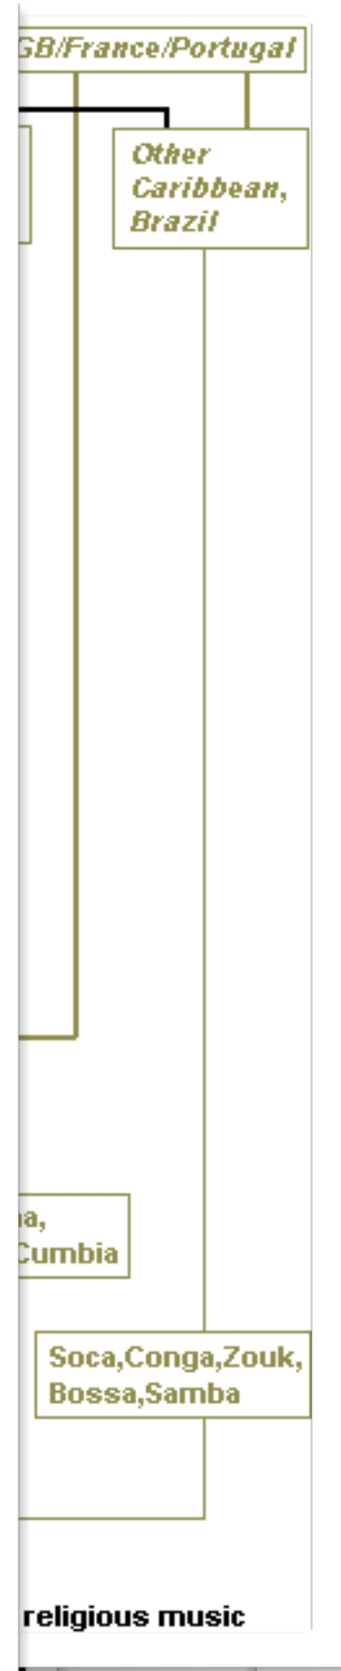

Handmade Genre Maps and Trees

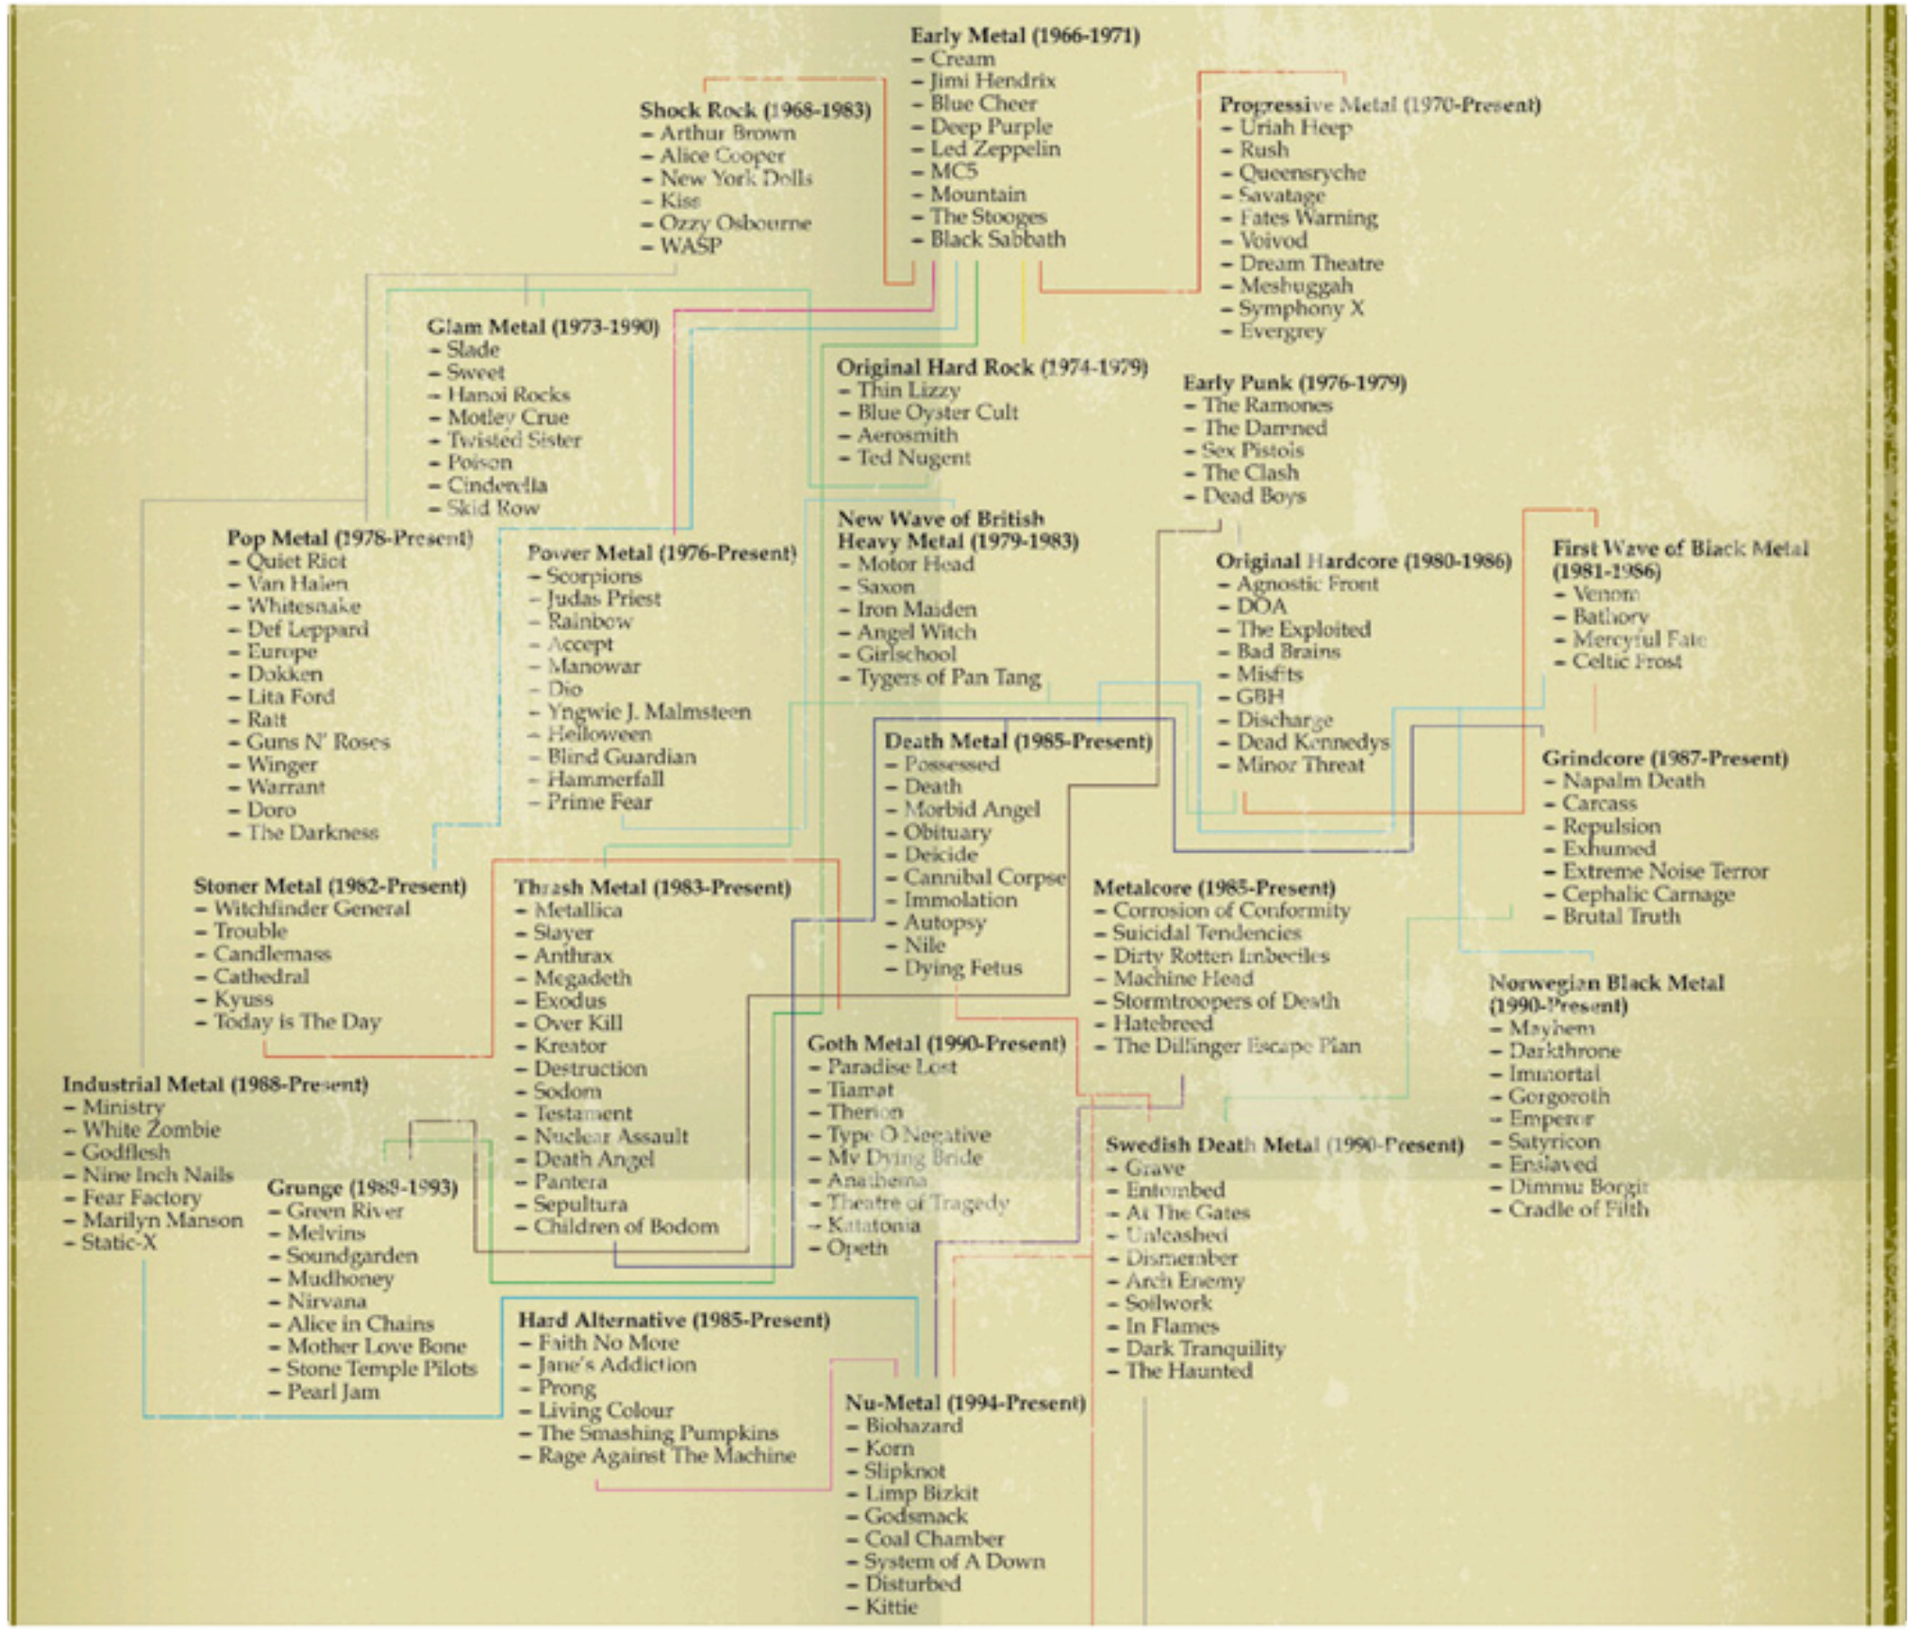

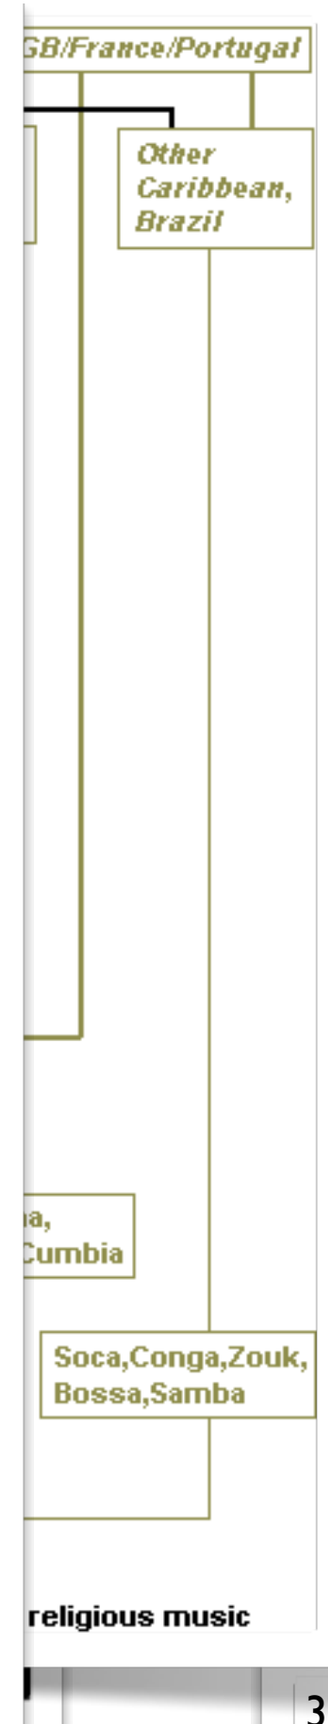

Handmade Genre Maps and Trees

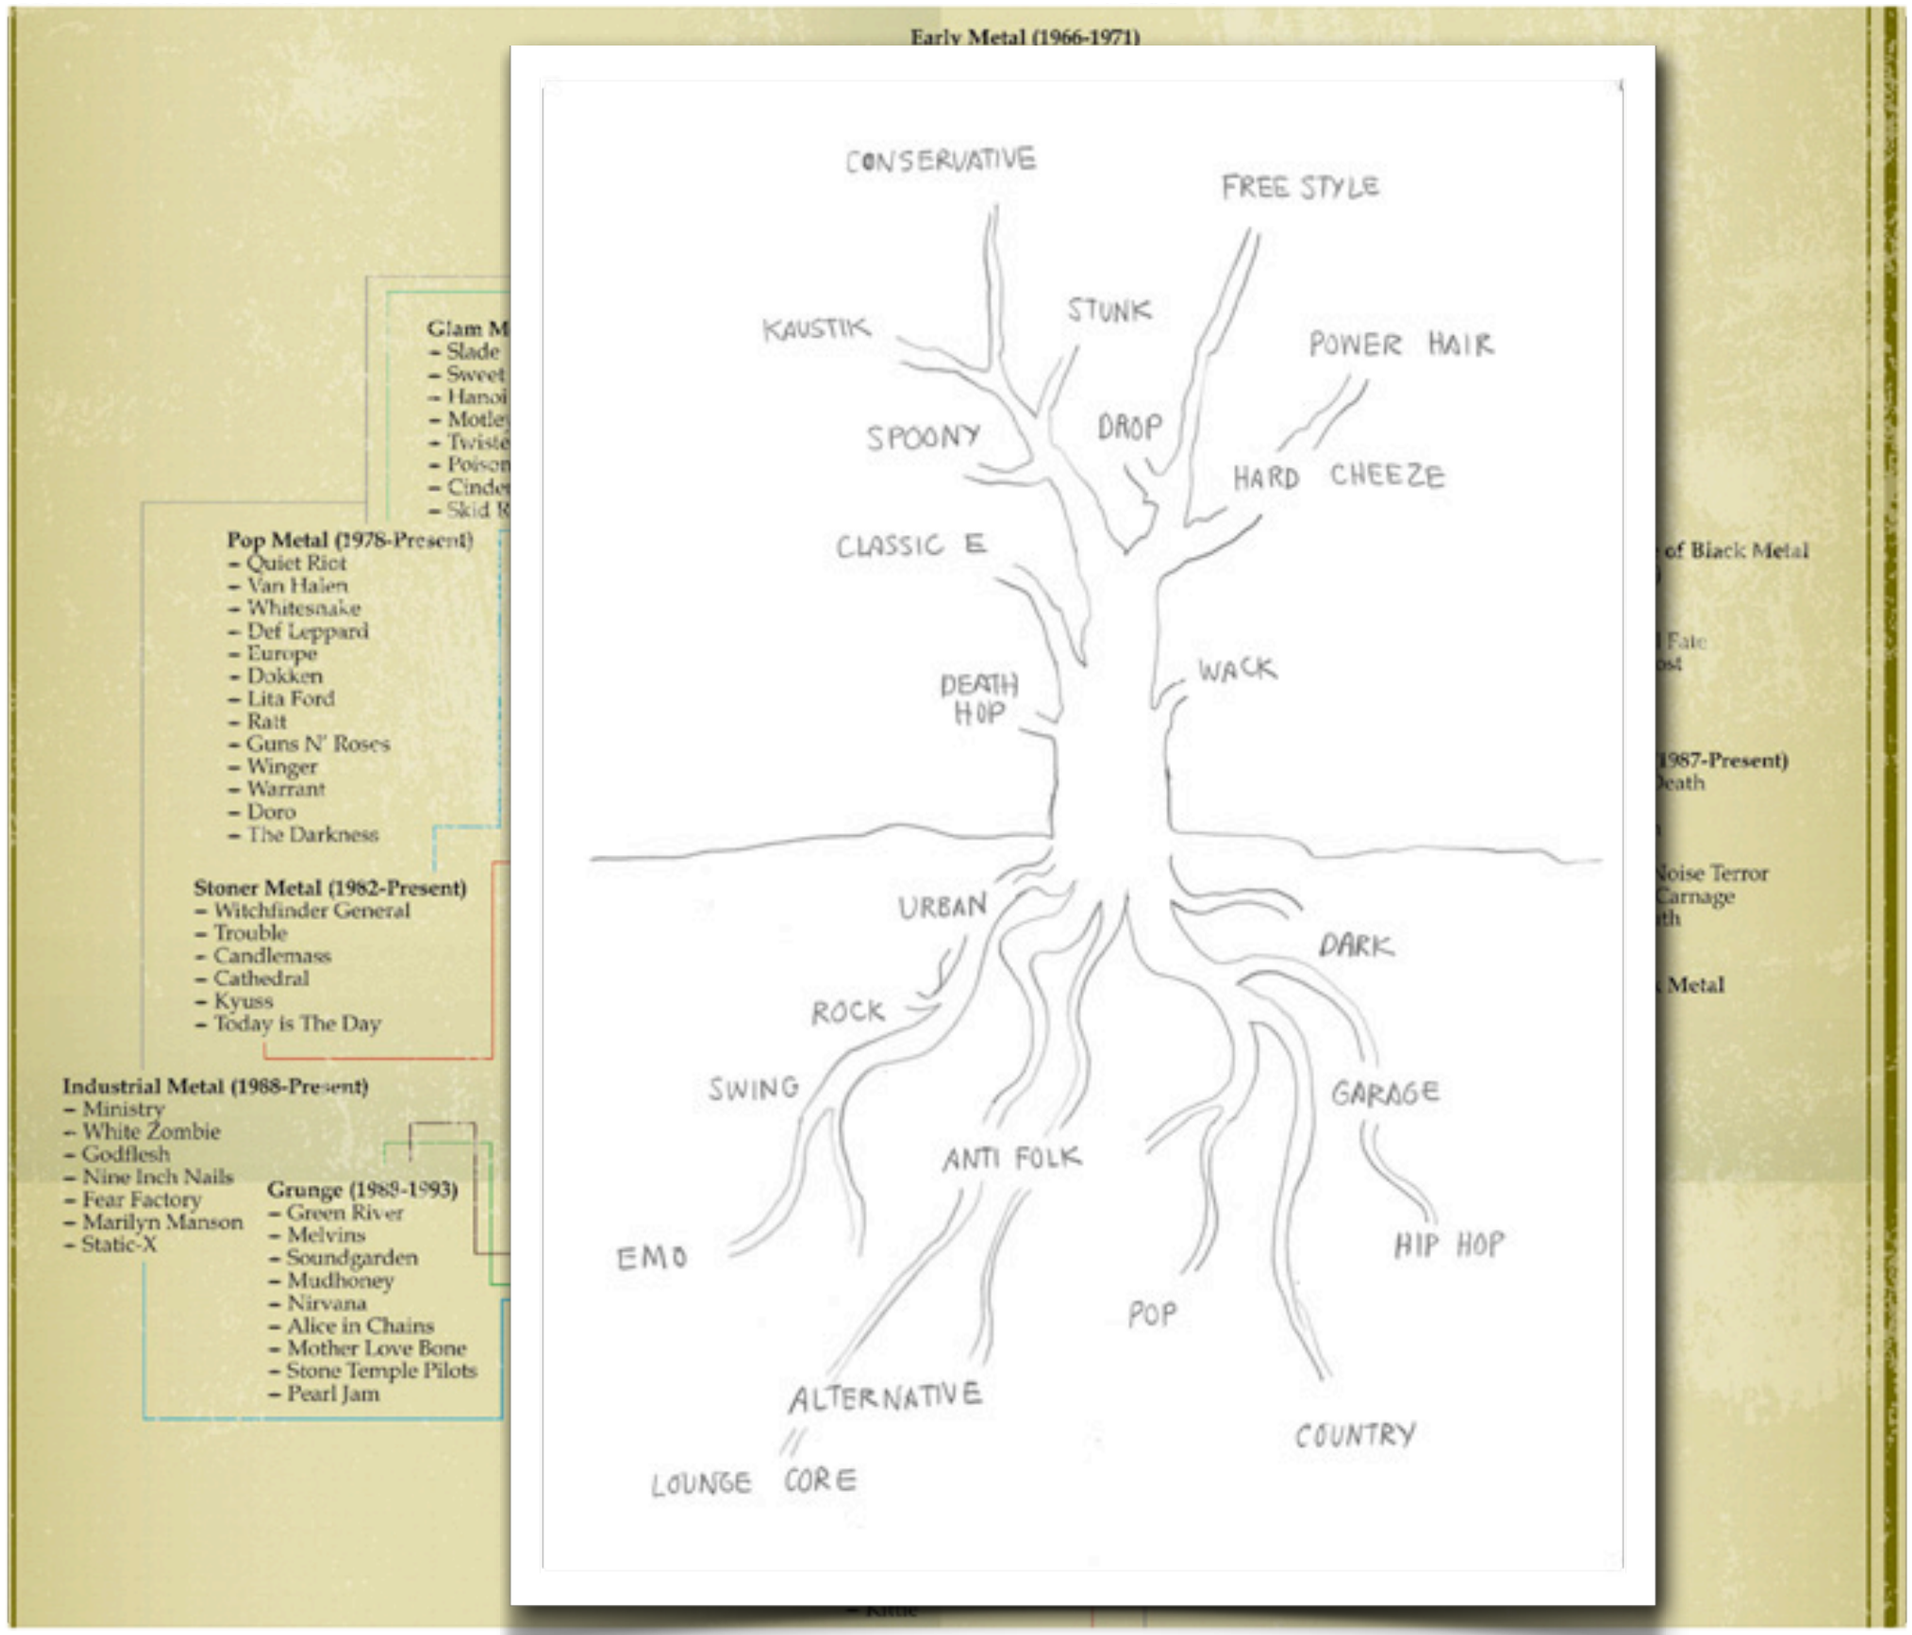

Other types of visualizations

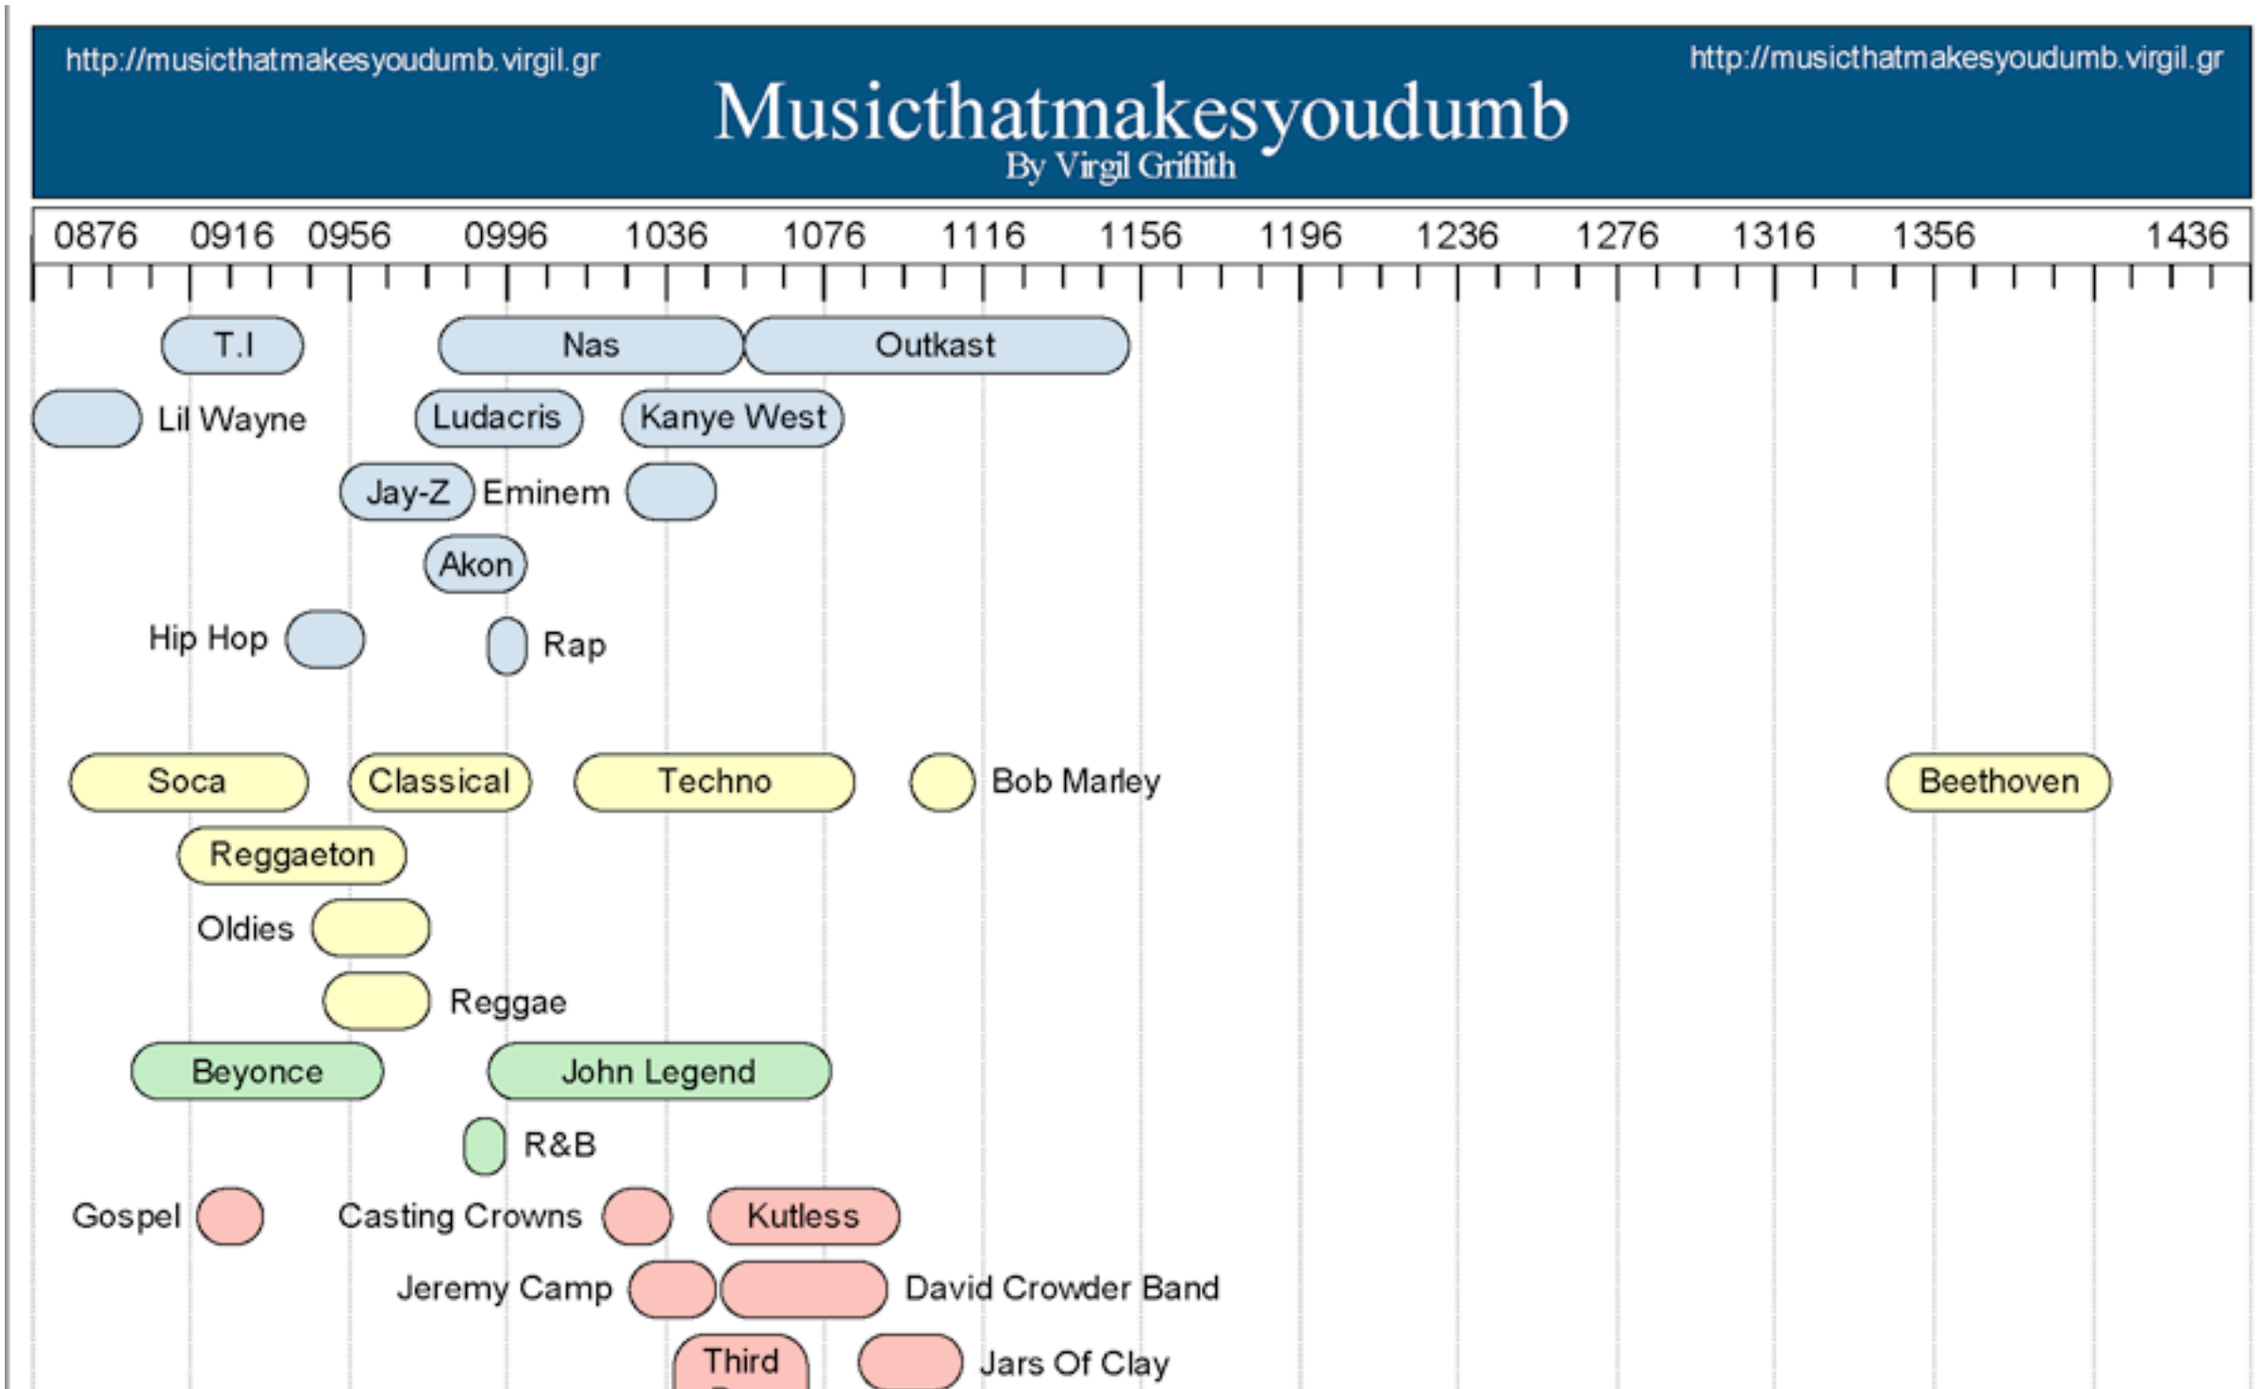

Other types of visualizations

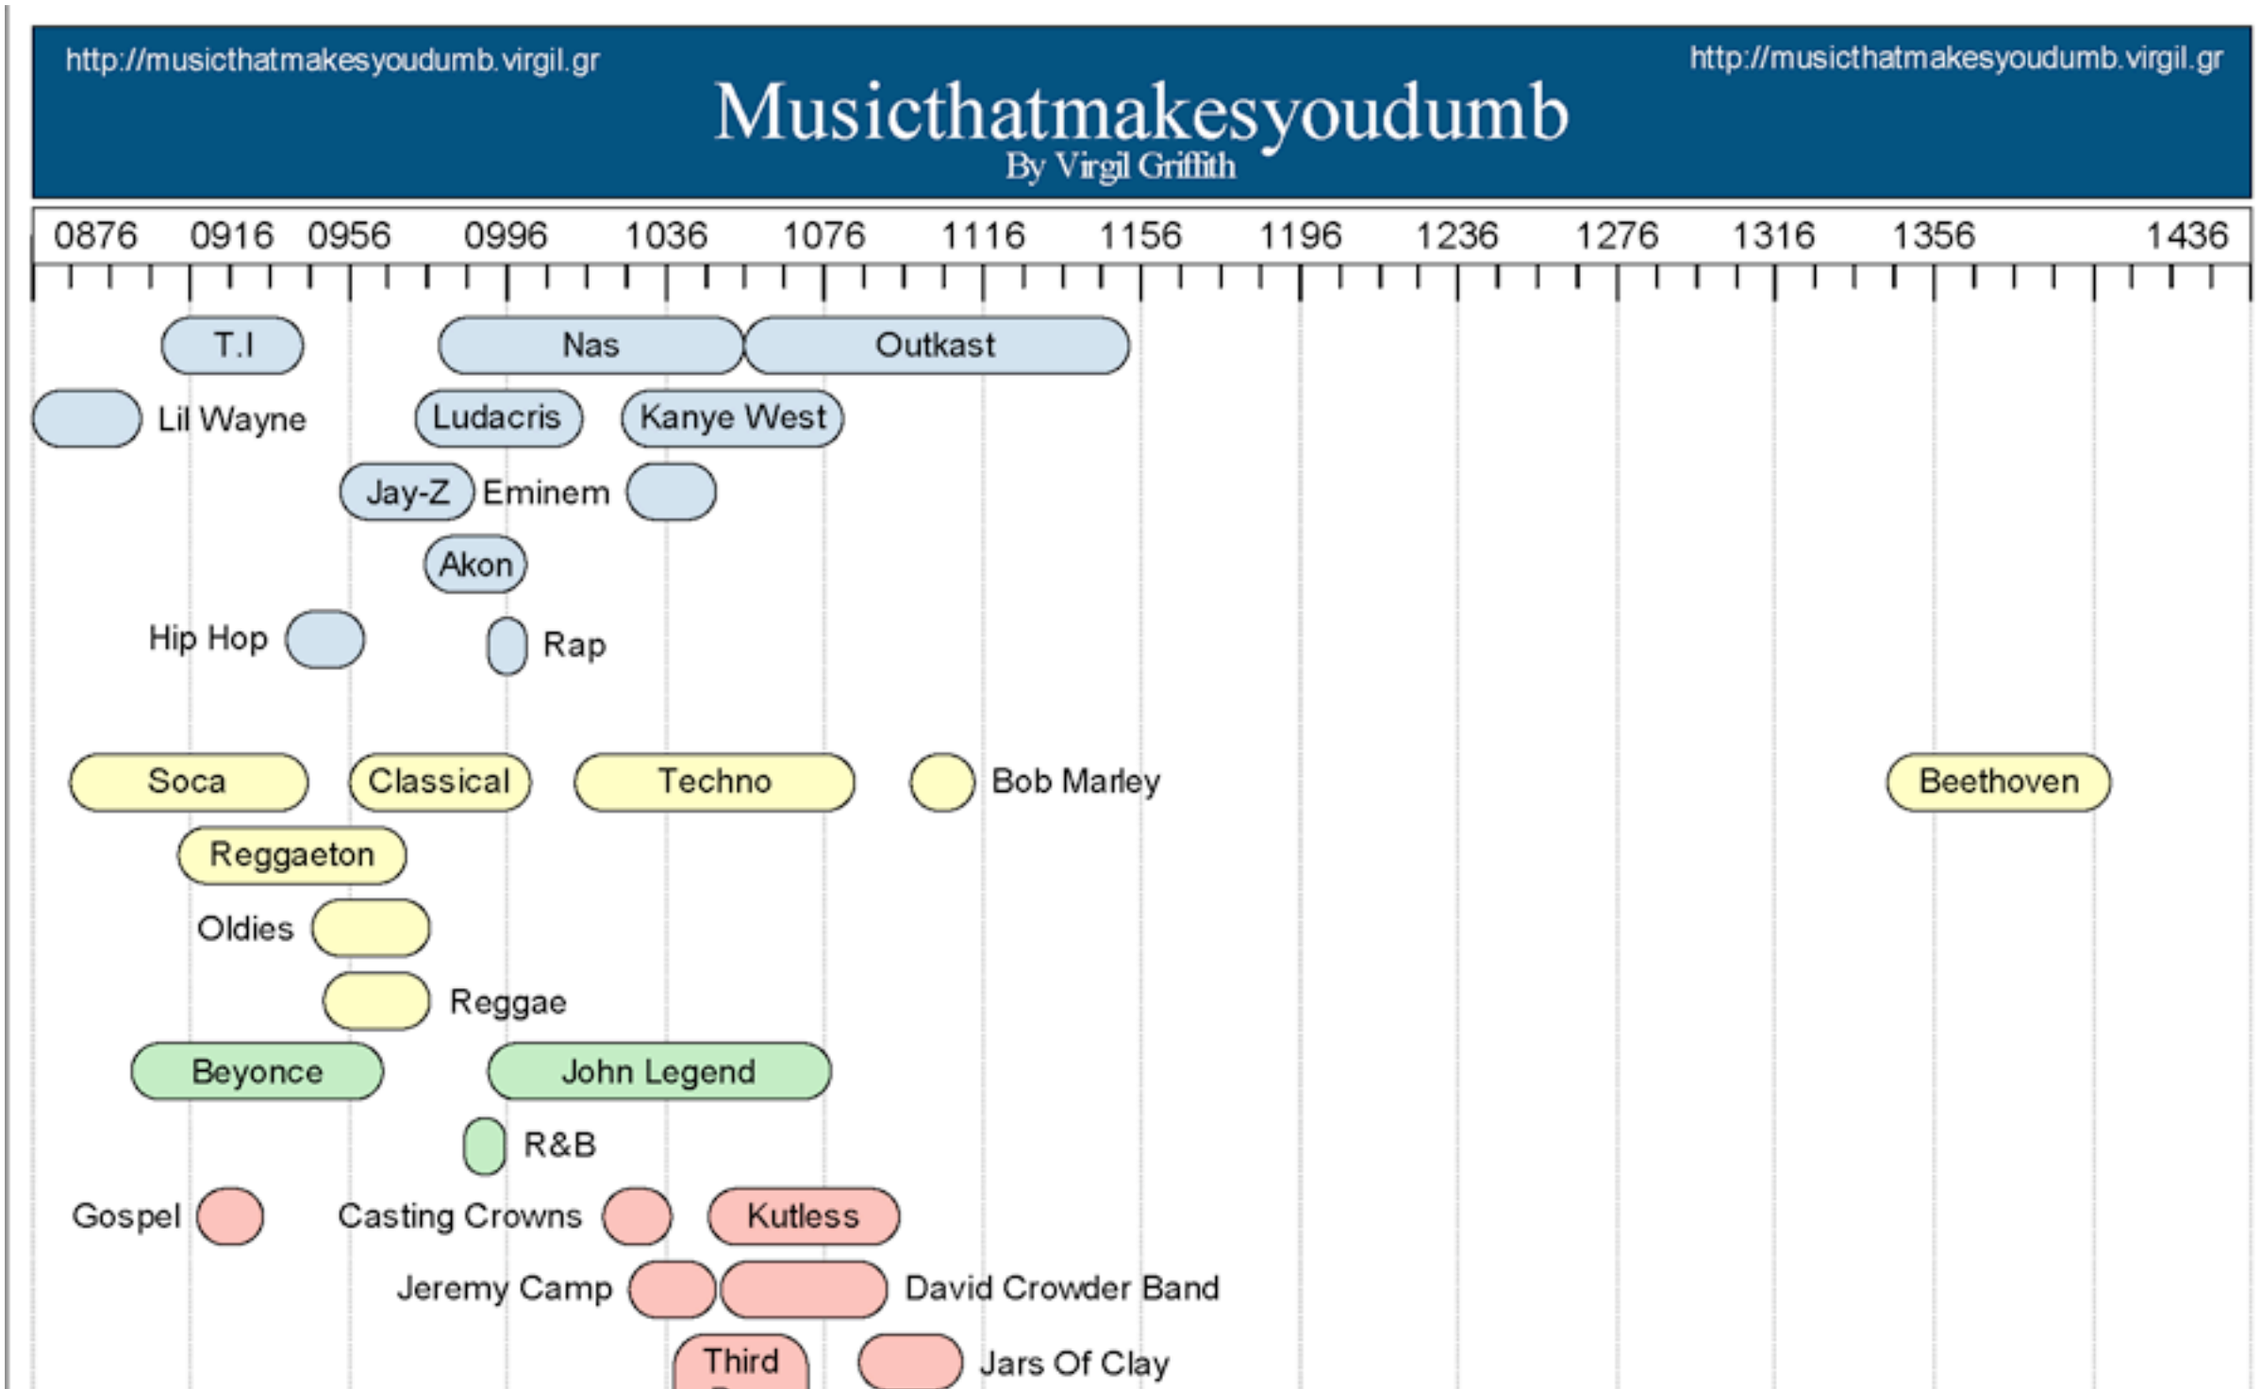

Other types of visualizations

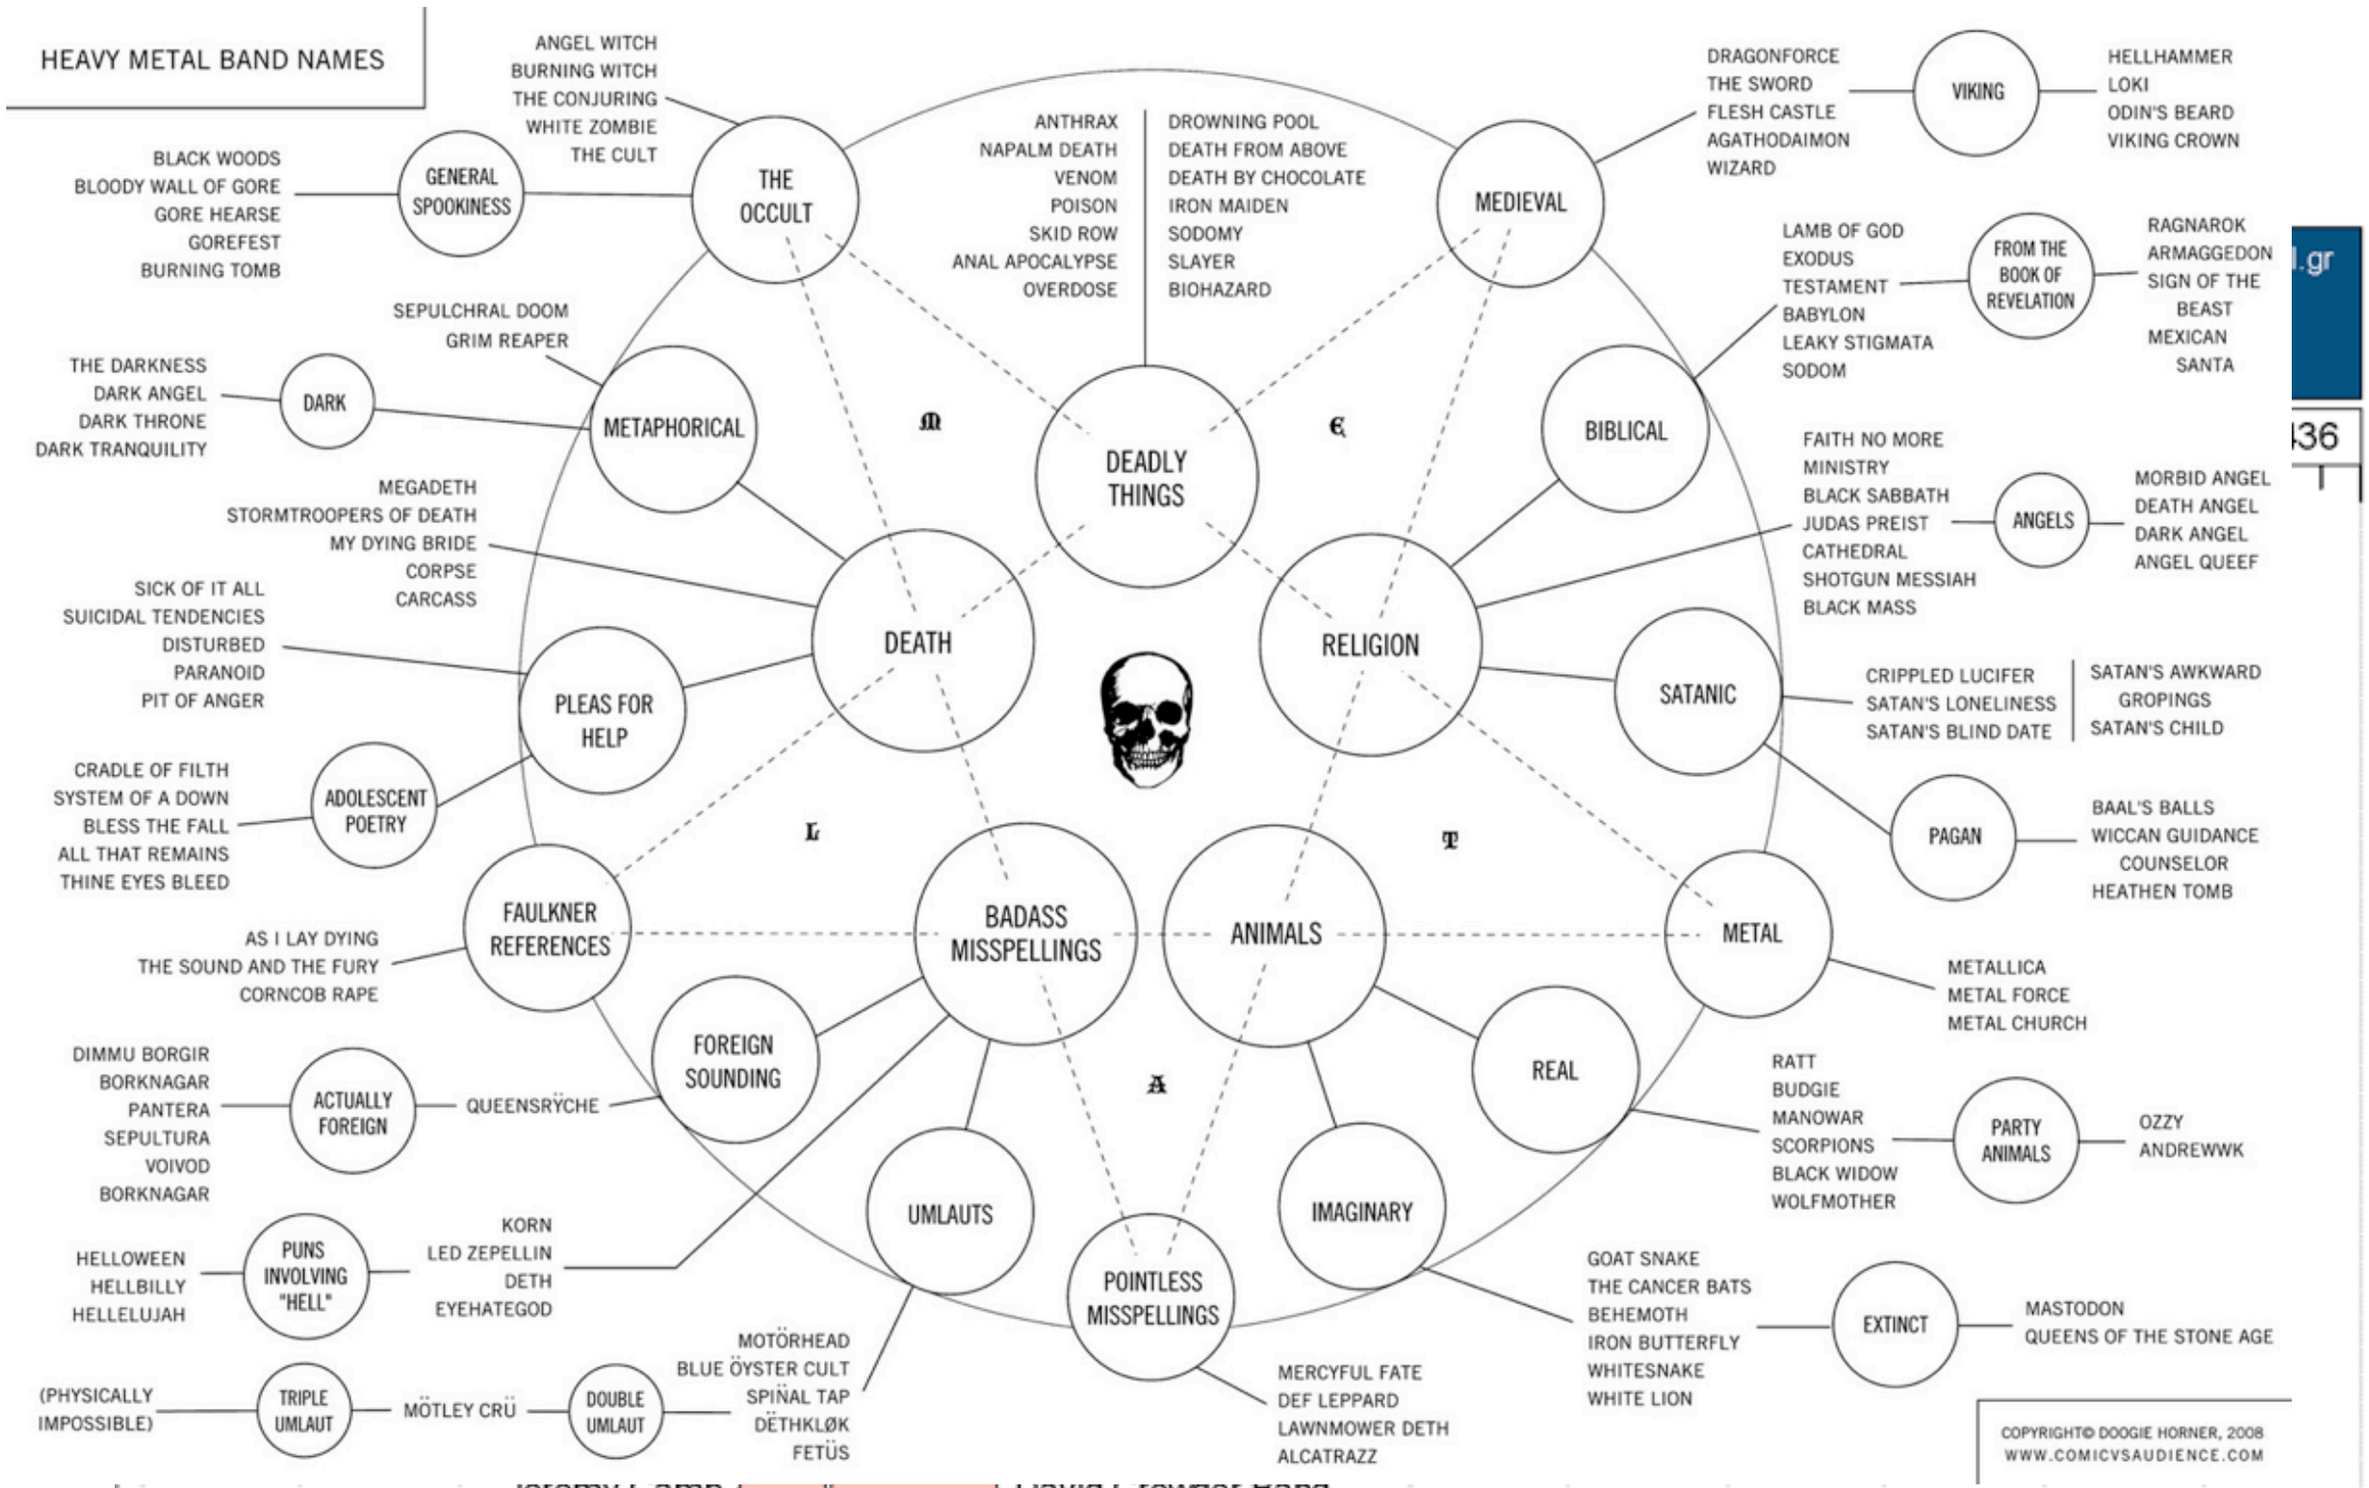

Other types of visualizations

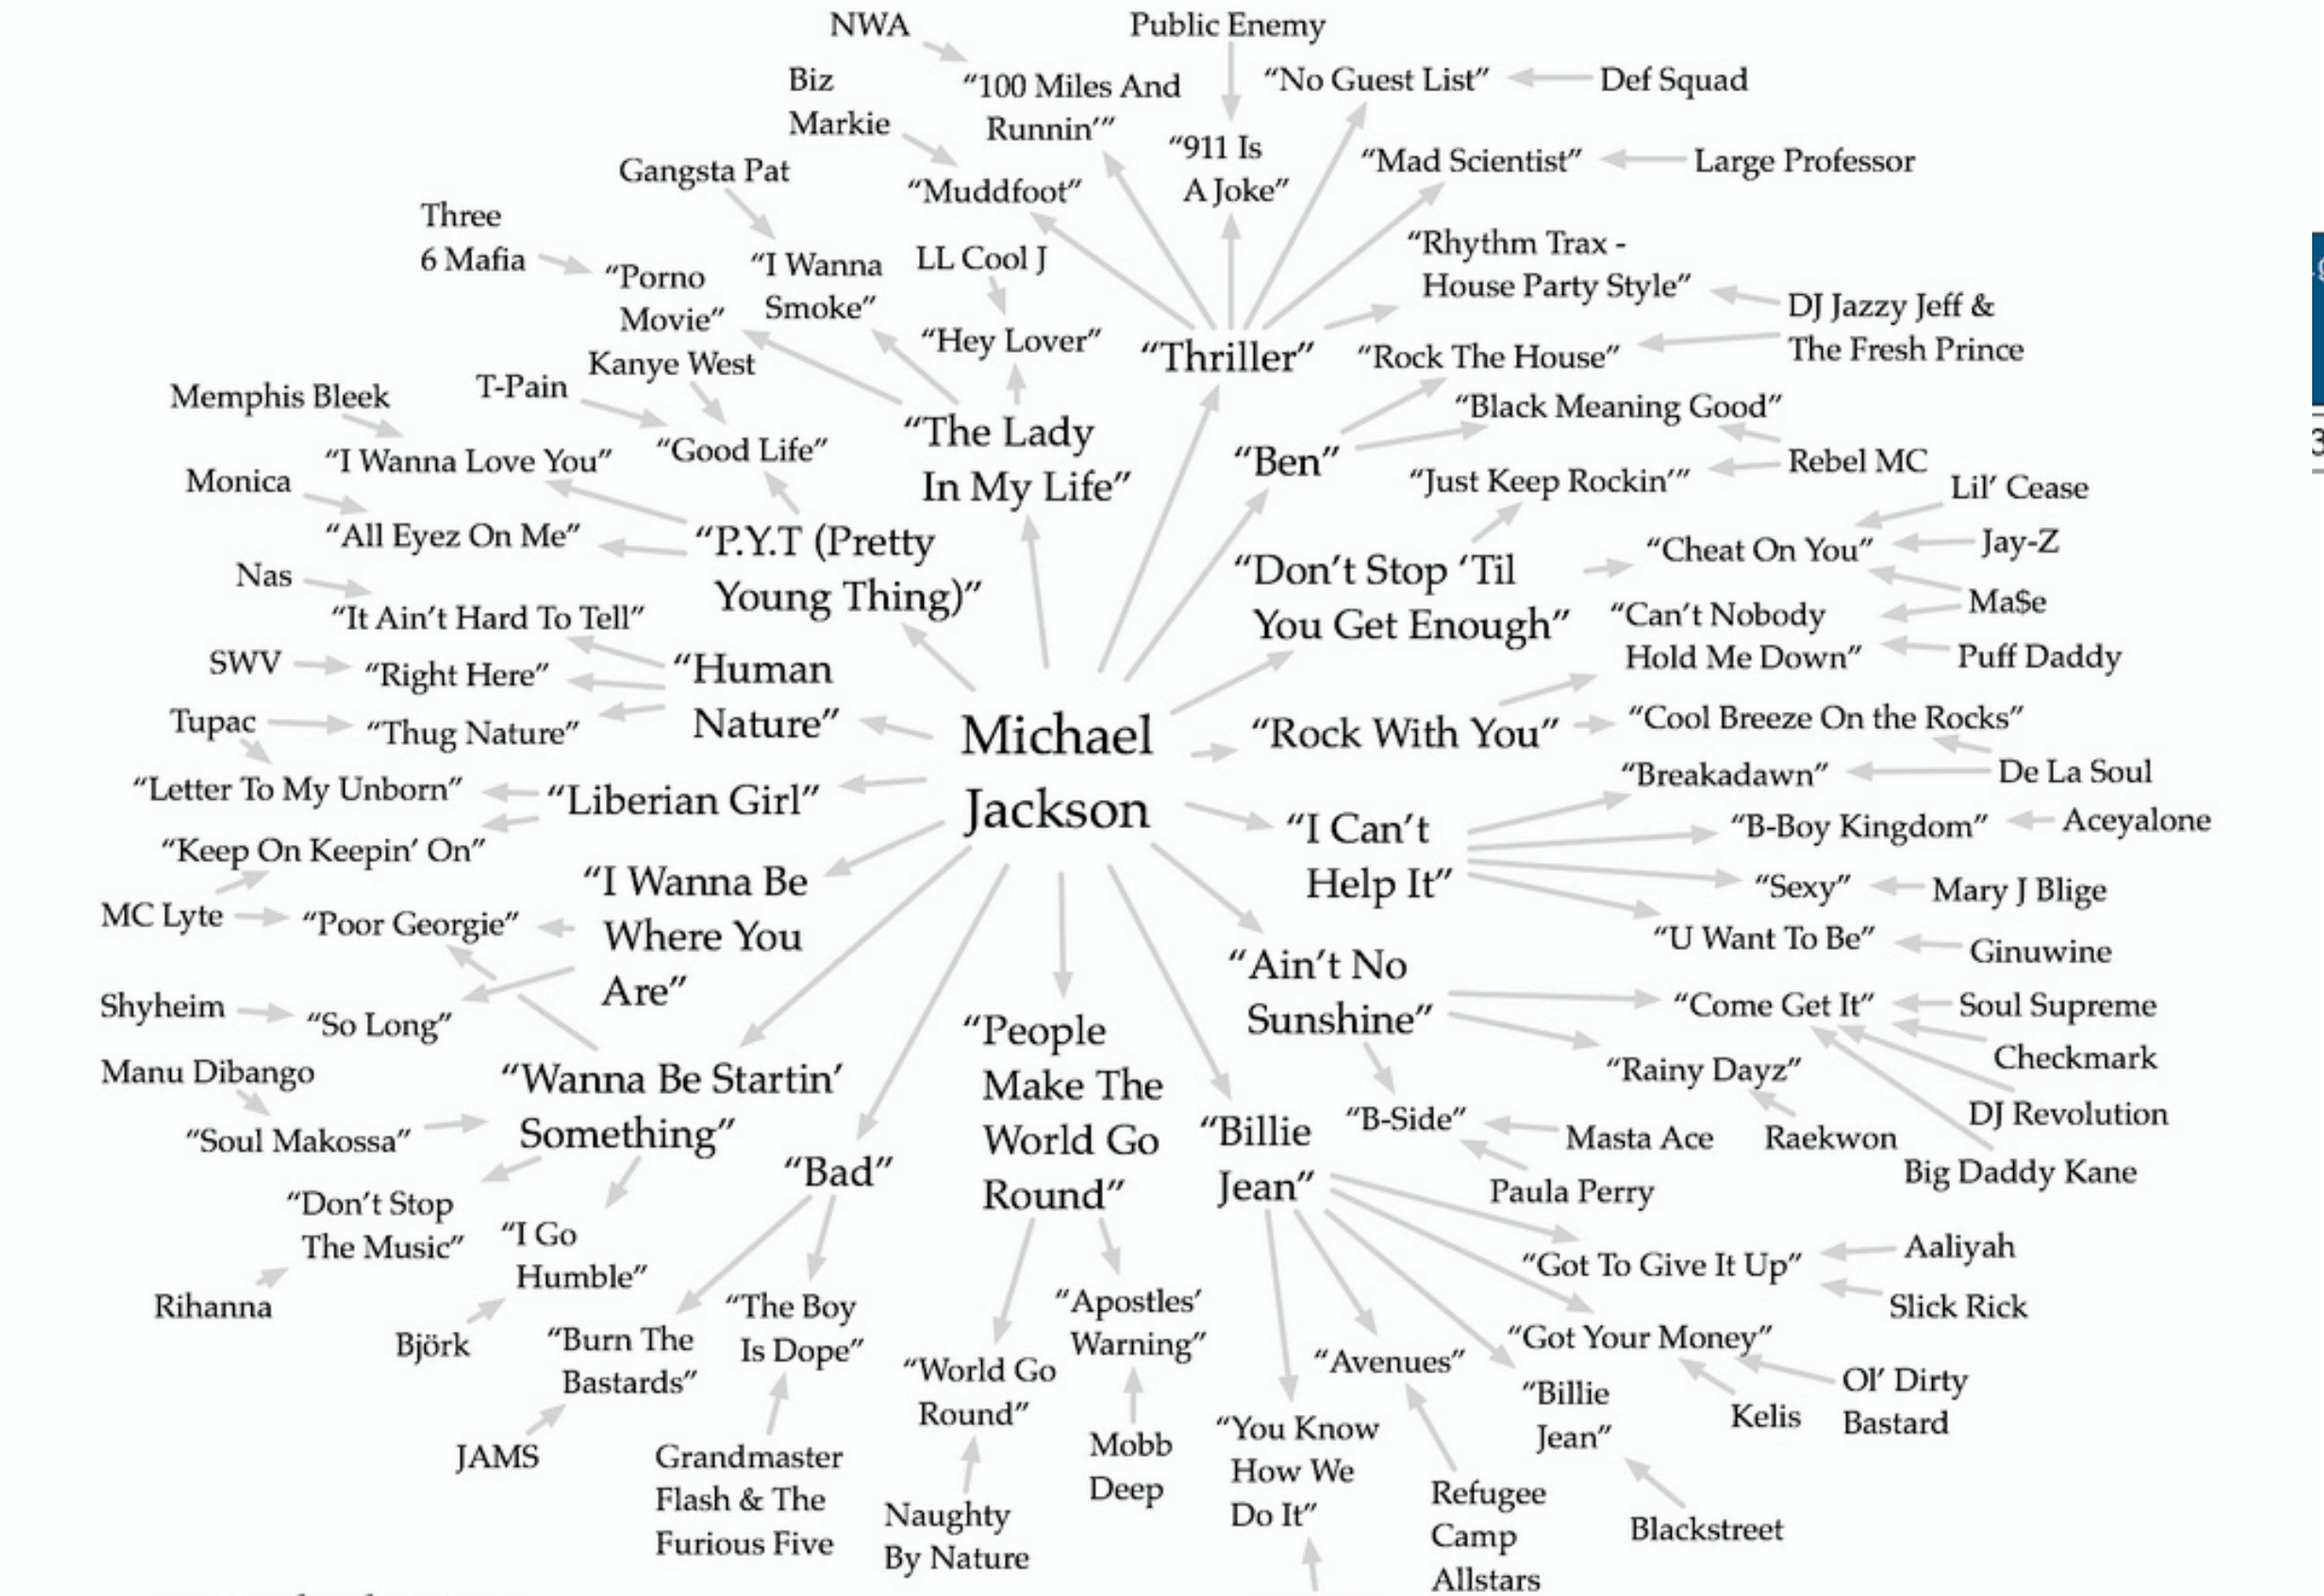

Other types of visualizations

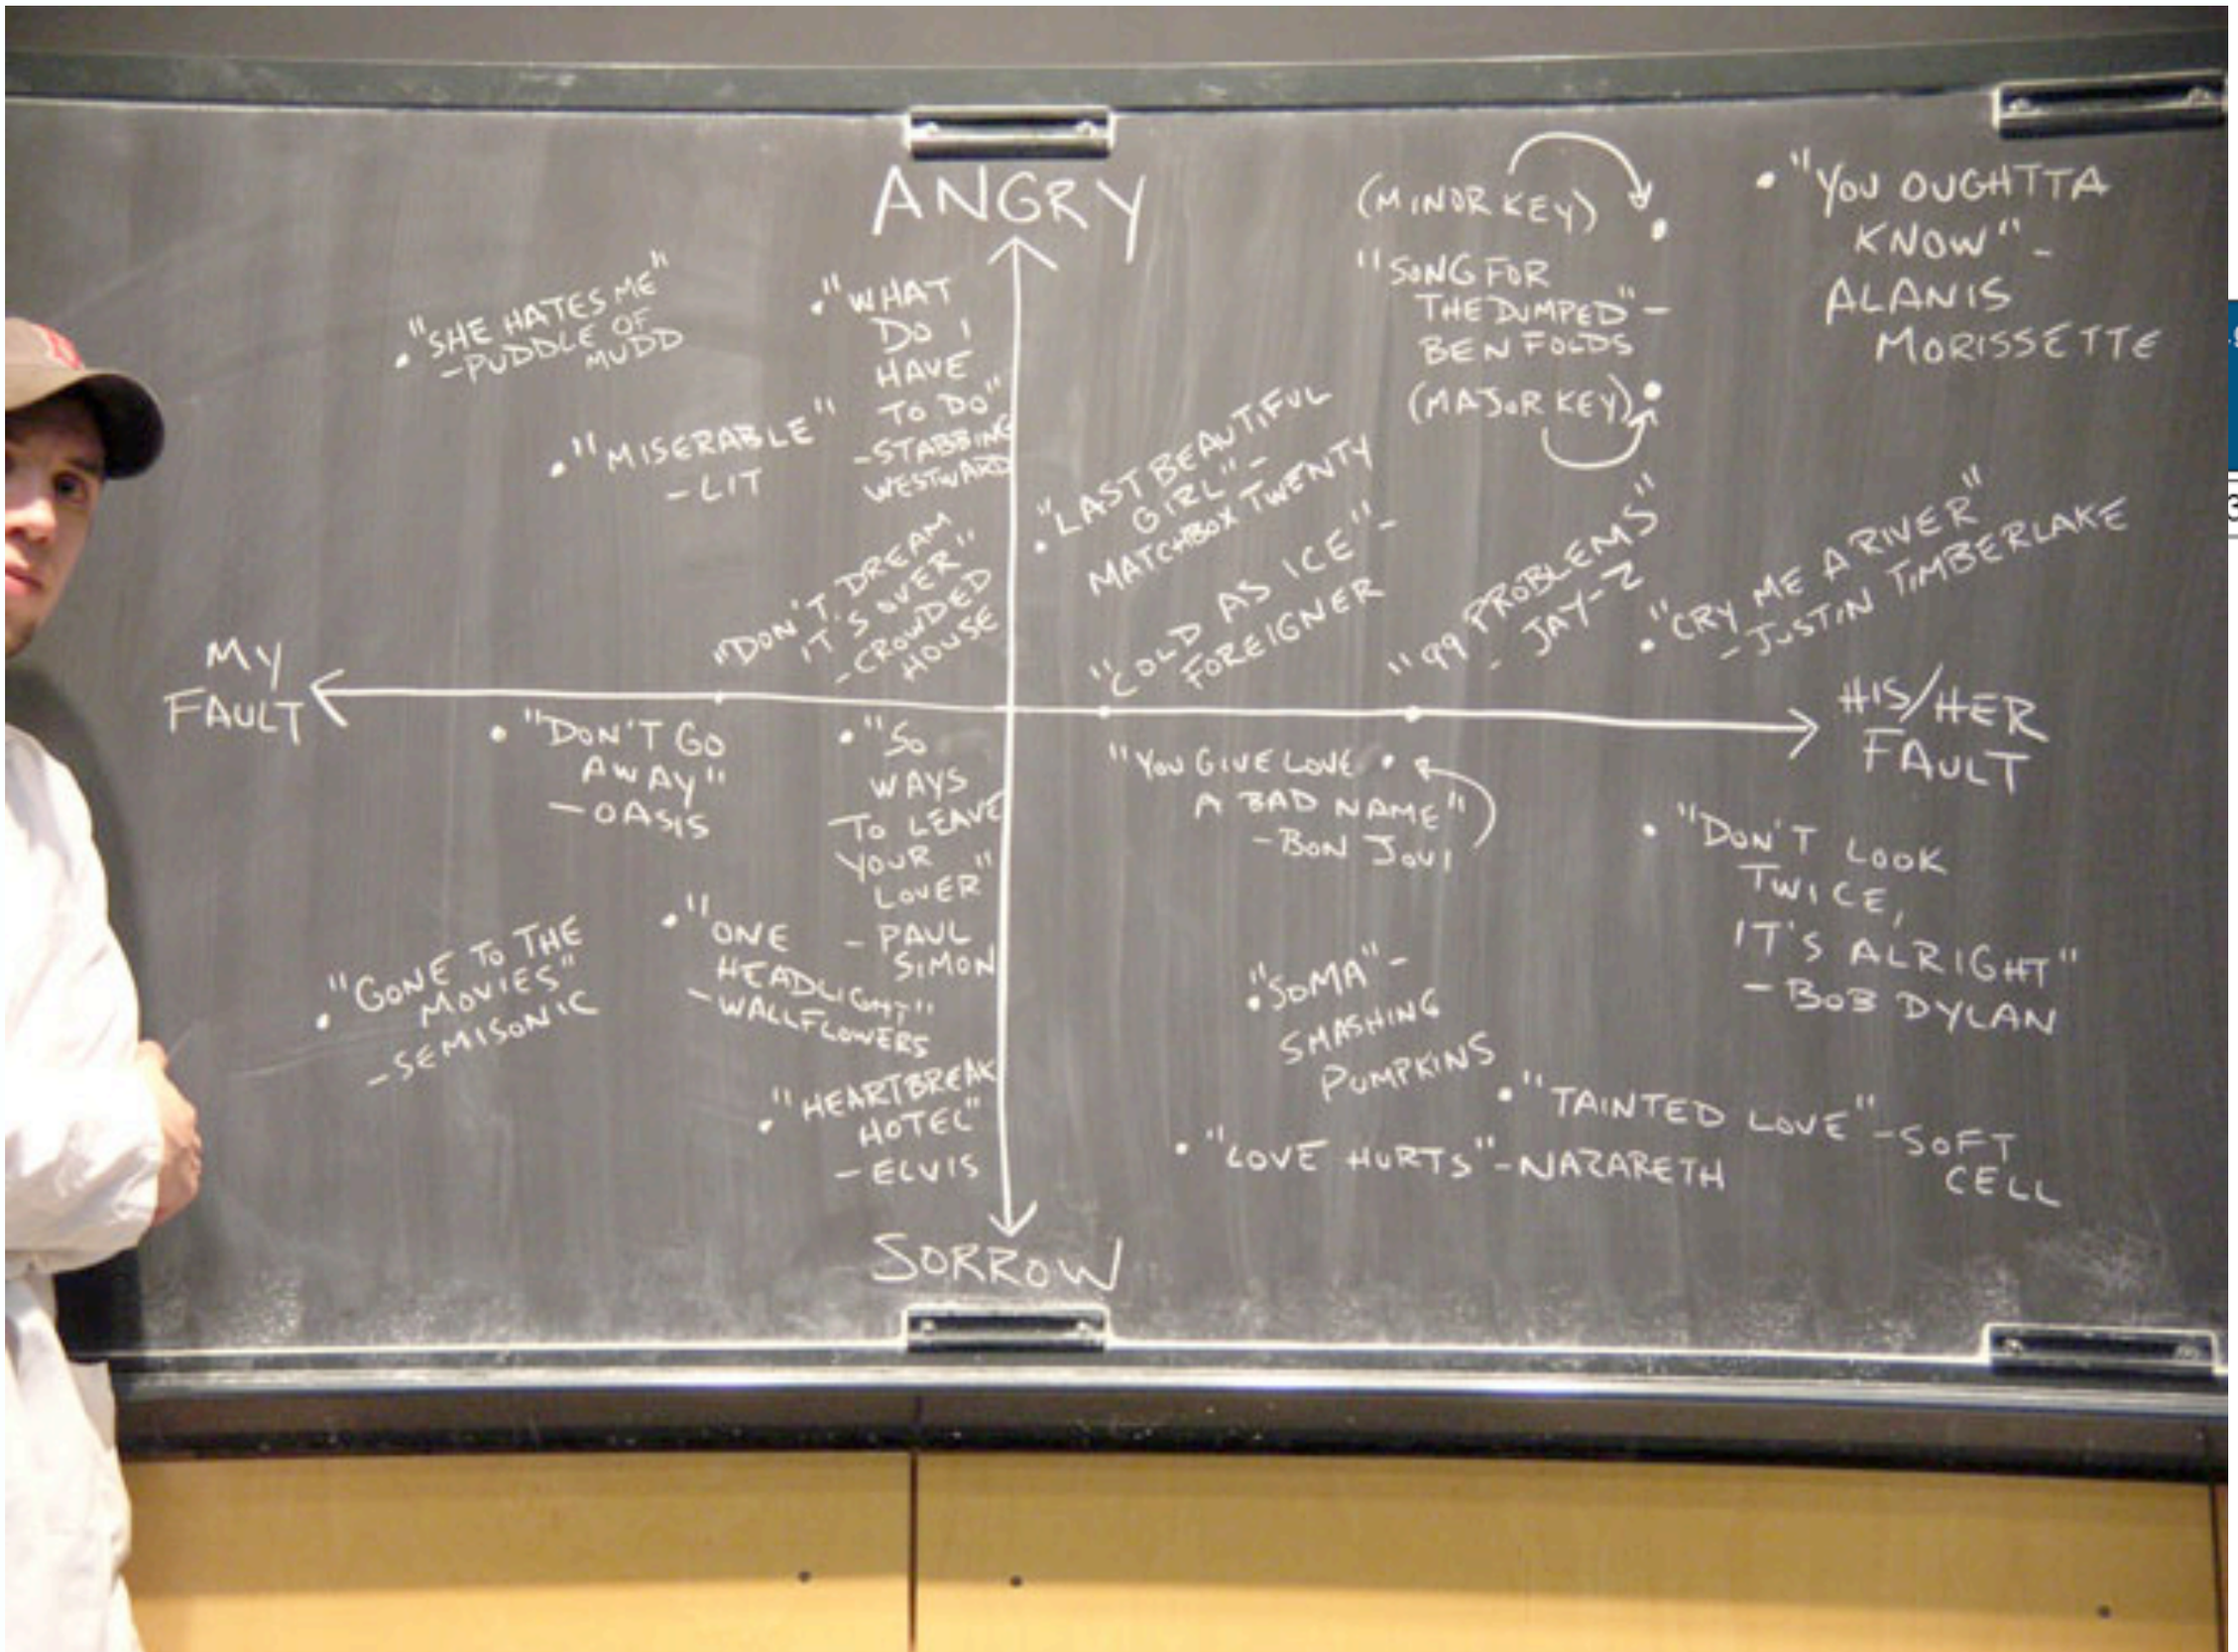

Other types of visualizations

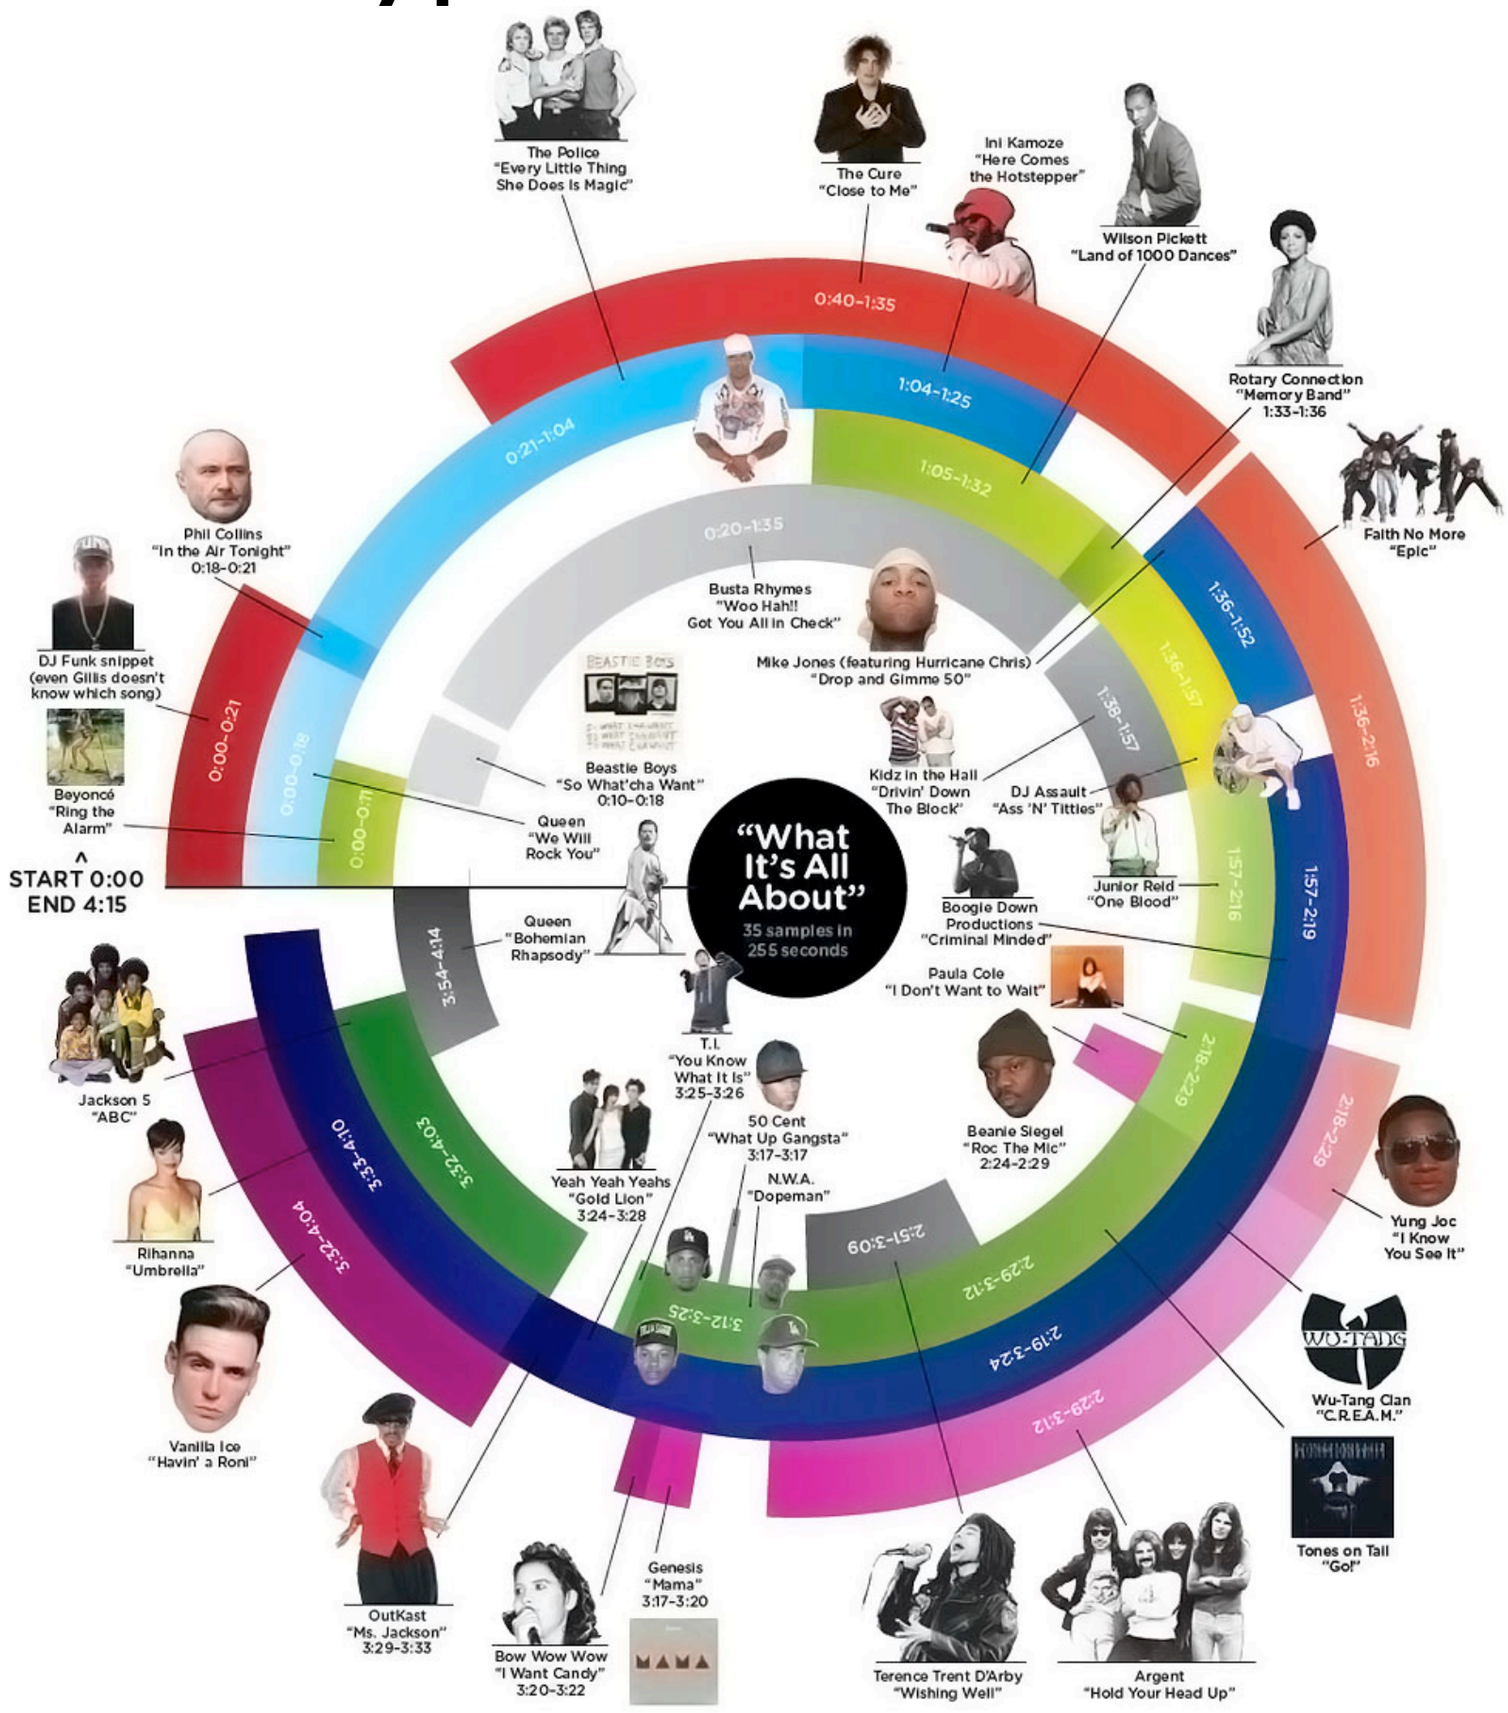

Other types of visualizations

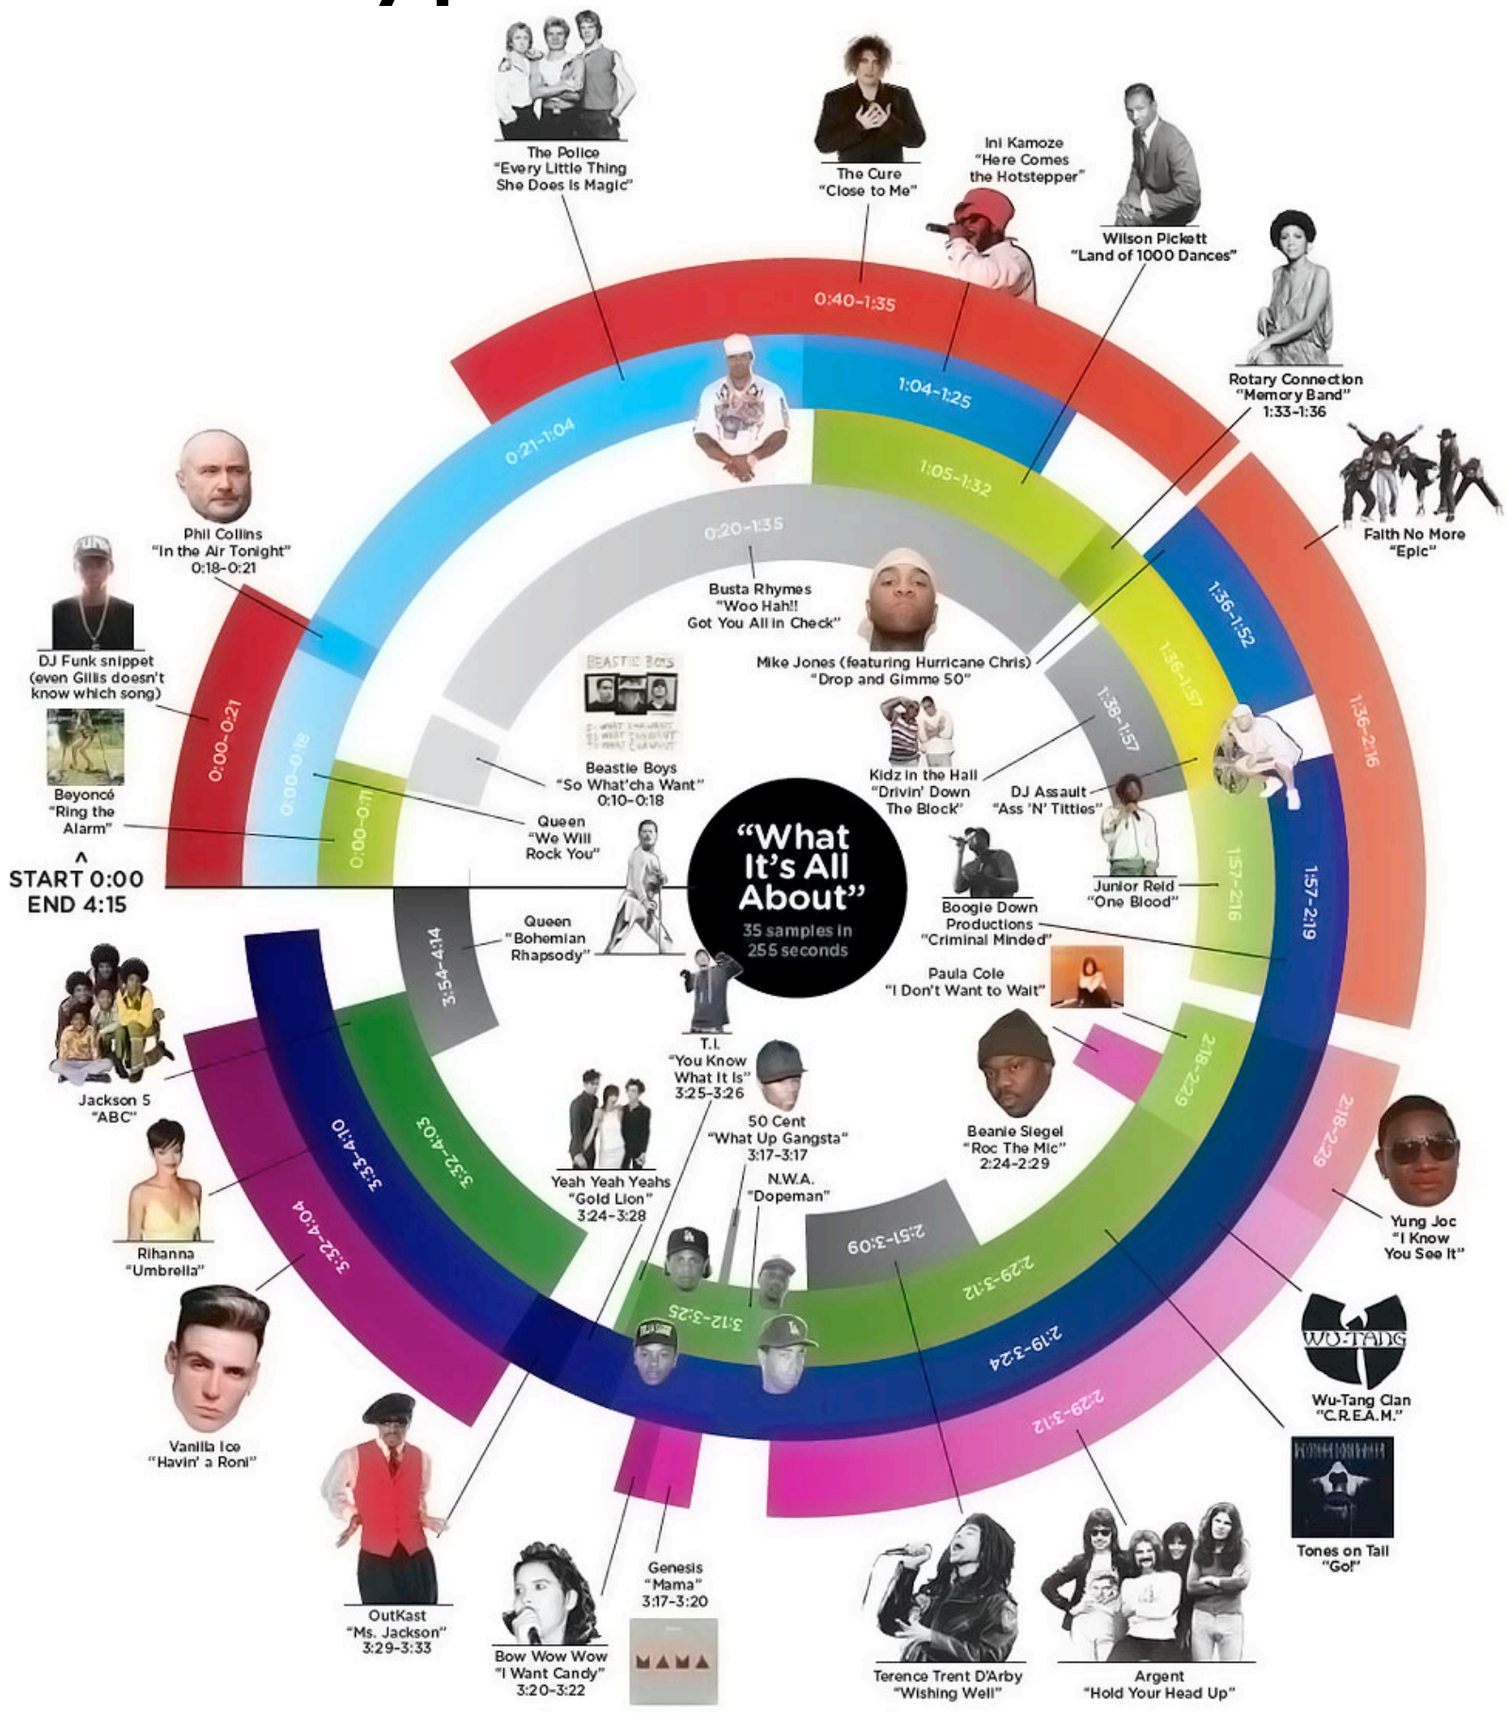

Artist as a function of temperature

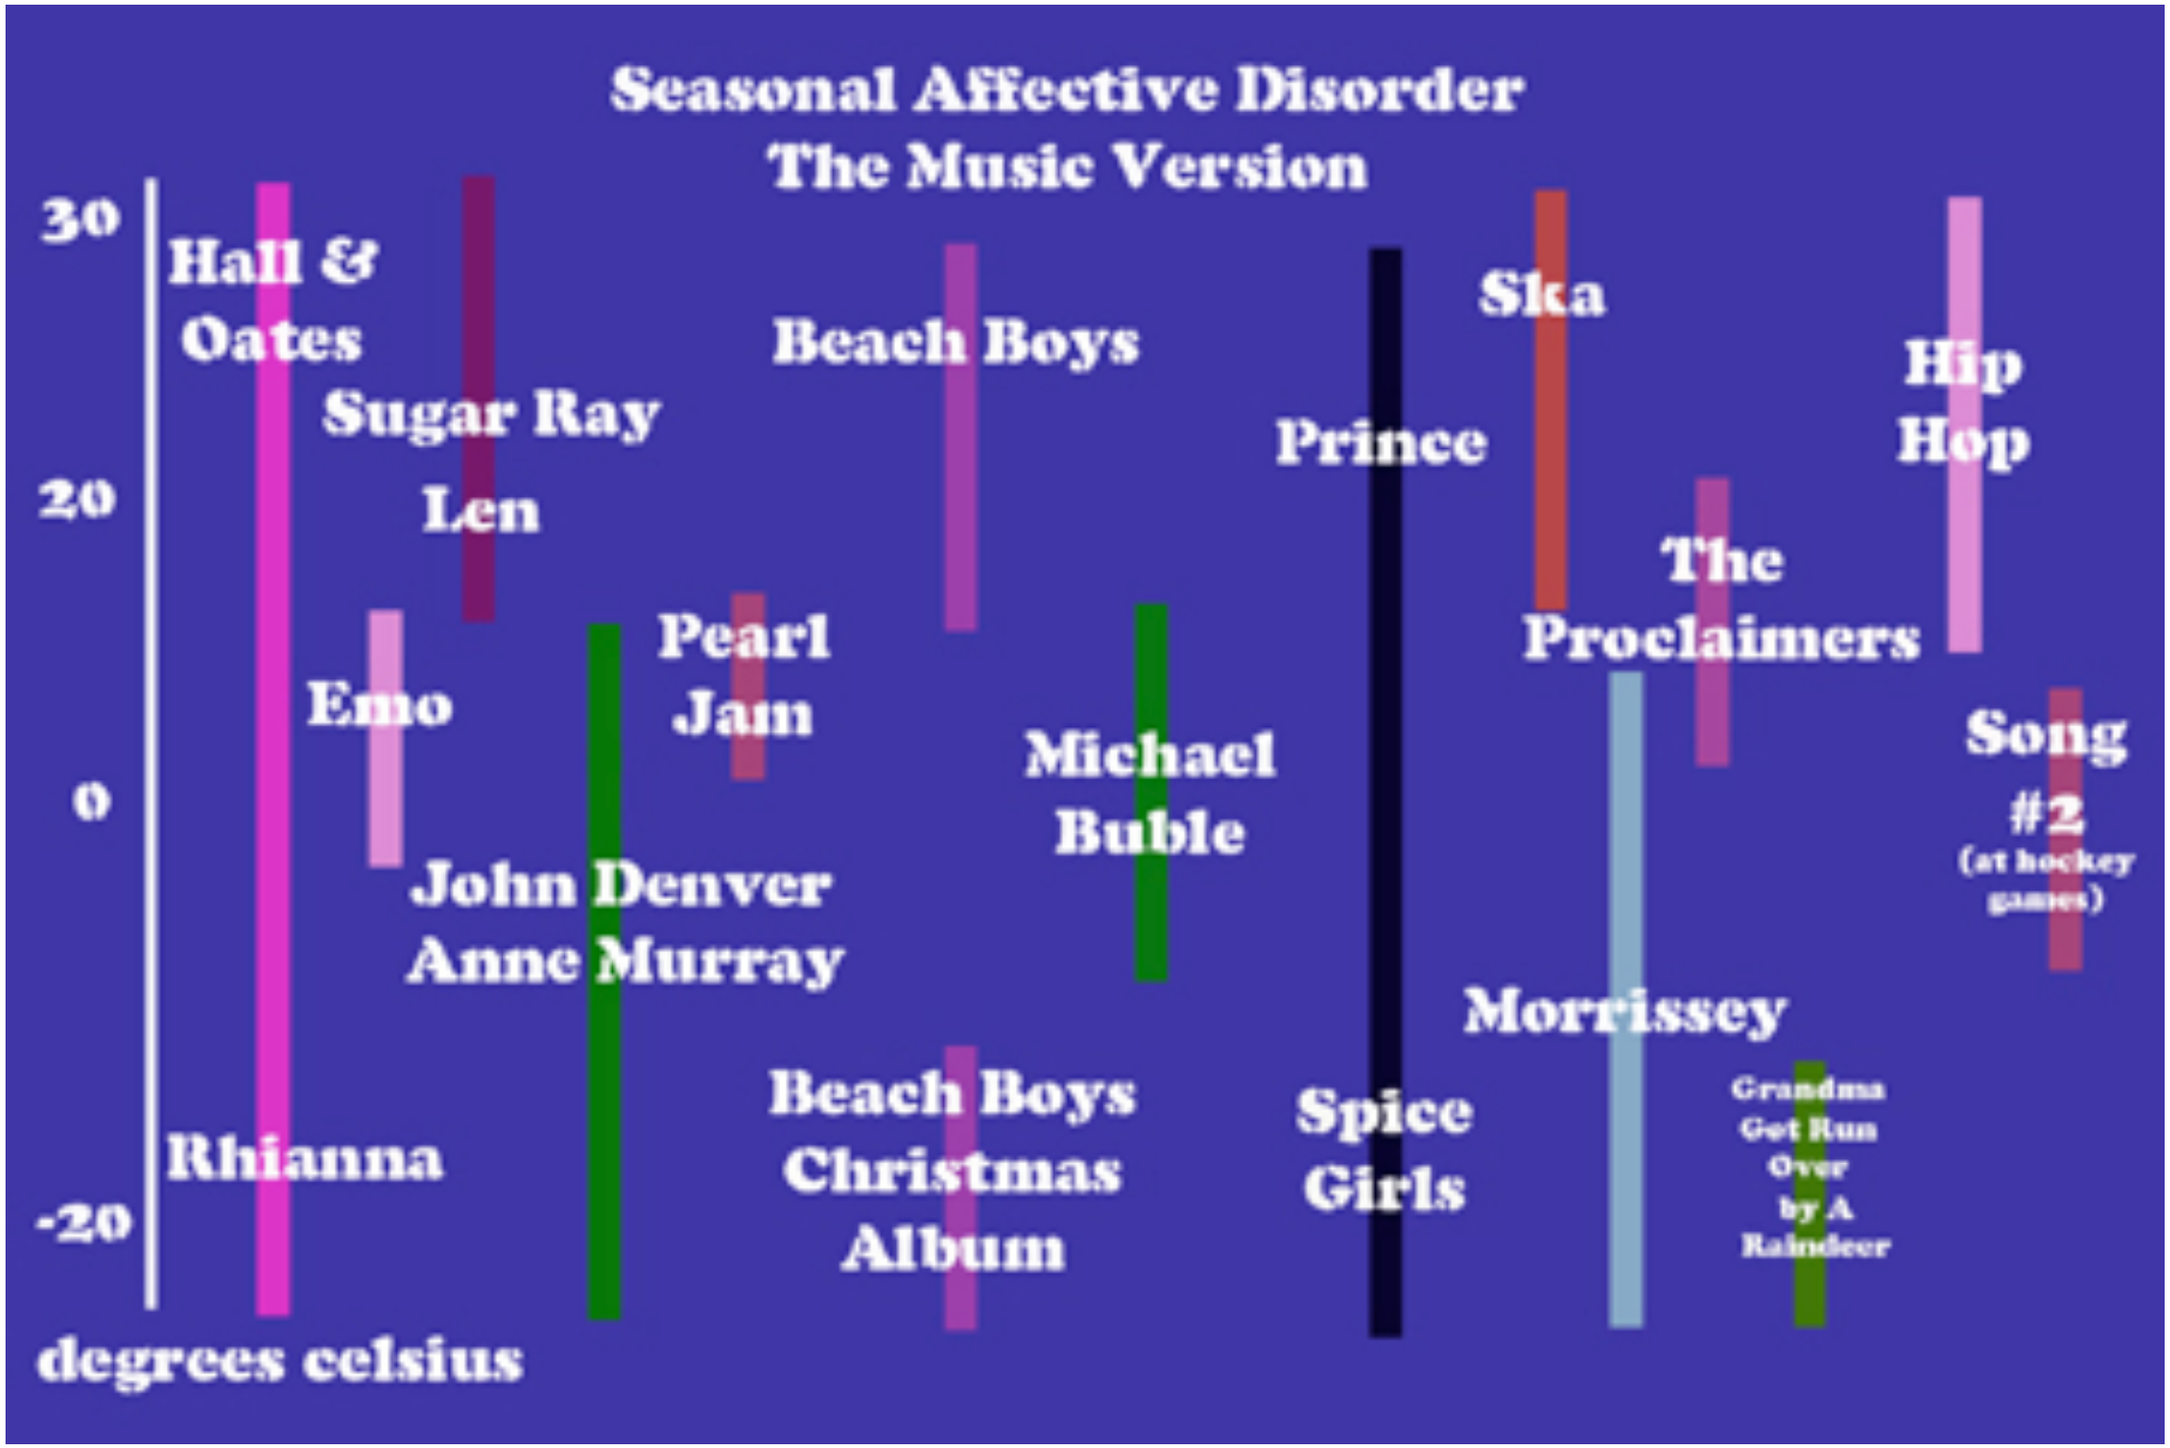

ISMIR Tutorial 2005

Visualization for Music IR

Tutorial II,part 2

ISMIR2005 London UK

Stephan.Baumann@dfki.de

(Thanks to the authors of original contributions!)

Senior Researcher >> German Research Center for AI

Creative Visionary >>> computationalculture.org

Stephan Bauman -Senior Researcher - German Research Center for AI

ISMIR Tutorial 2005

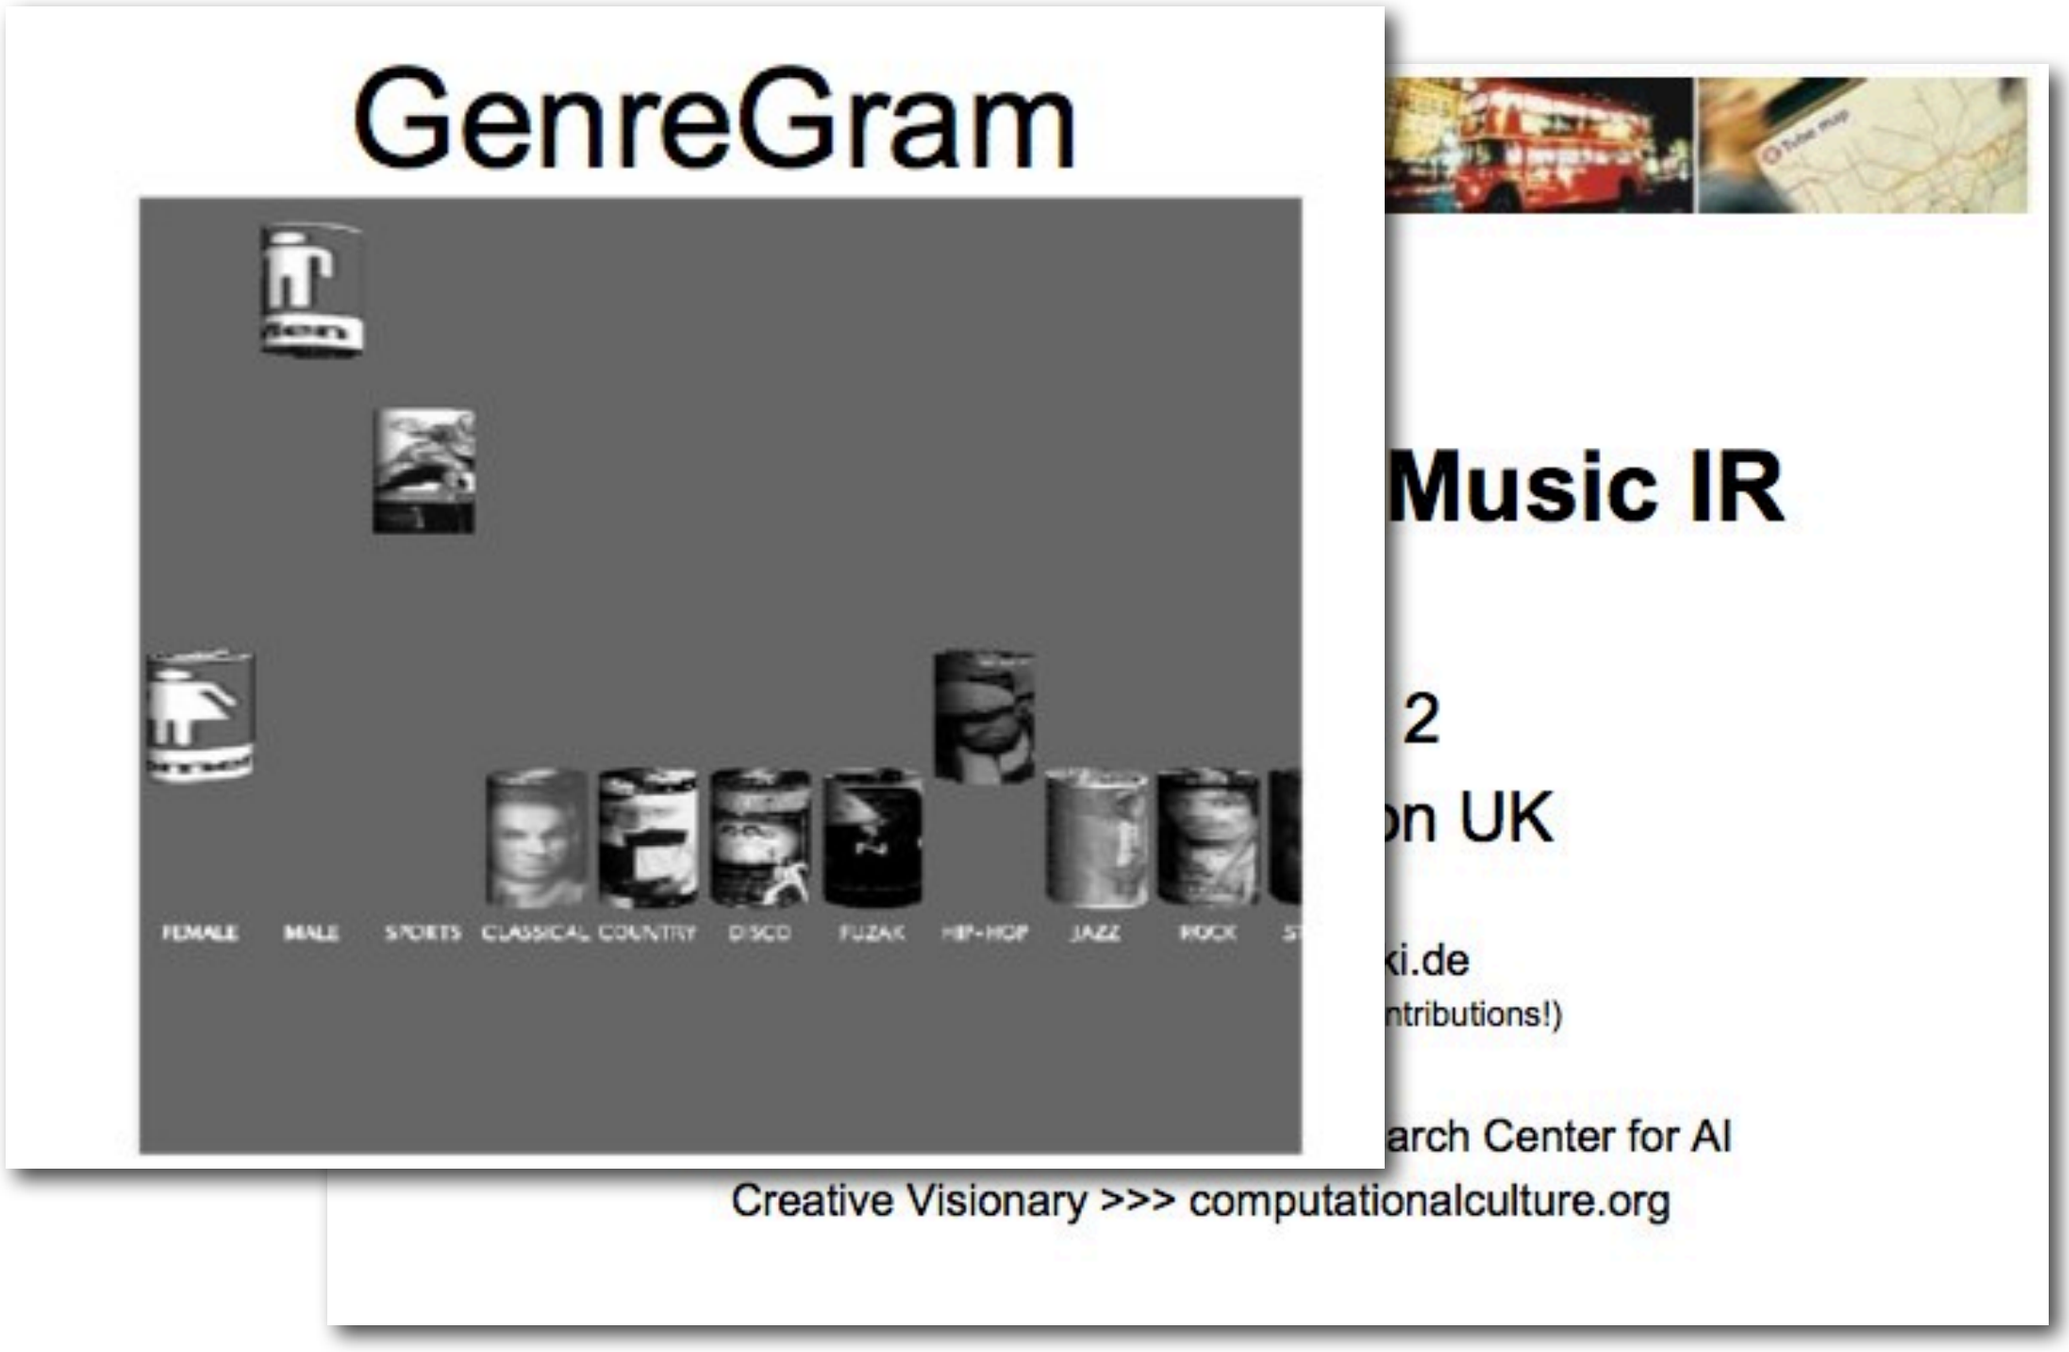

Stephan Bauman -Senior Researcher - German Research Center for AI

ISMIR Tutorial 2005

GenreGram

Powerwall of Tzanetakis

Stephan Bauman -Senior Researcher - German Research Center for AI

ISMIR Tutorial 2005

GenreGram

Powerwall of Tzanetakis

Weathercharts, Islands of Music

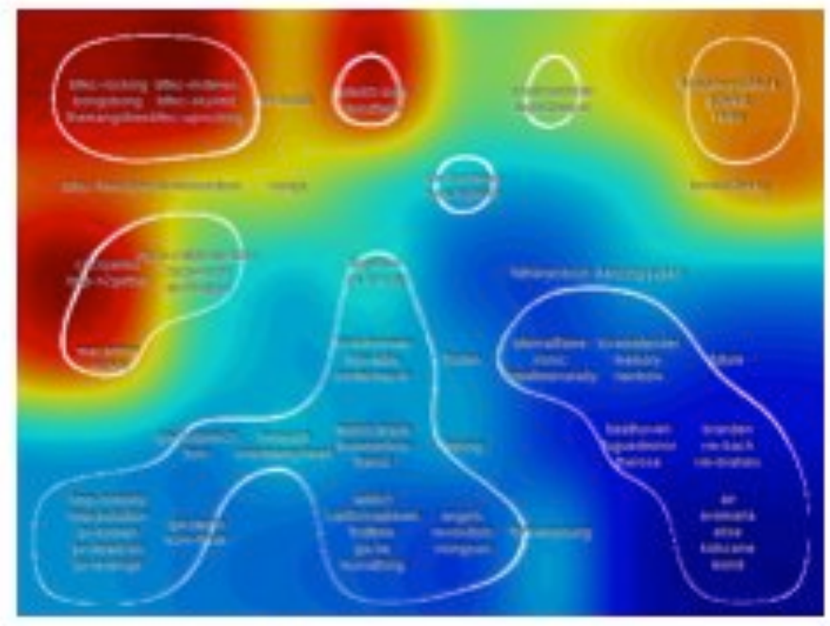

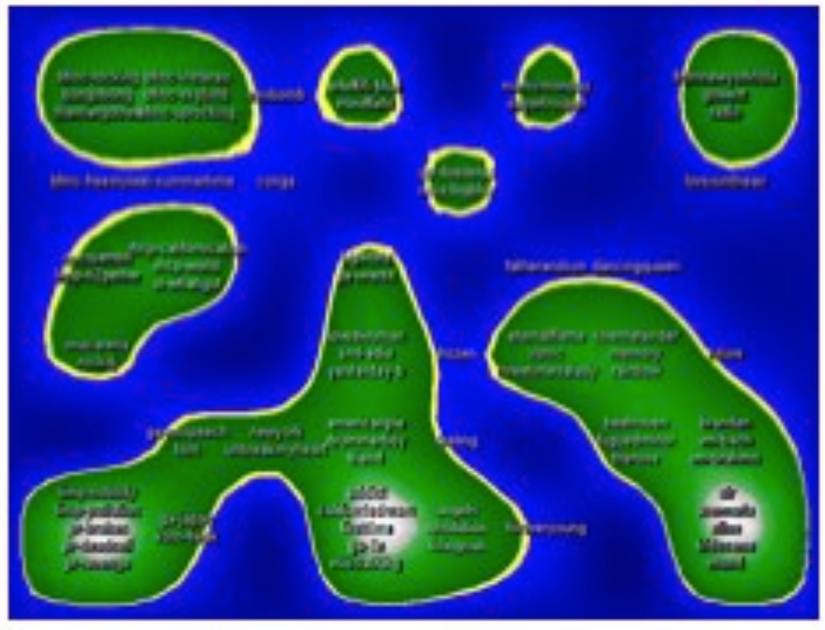

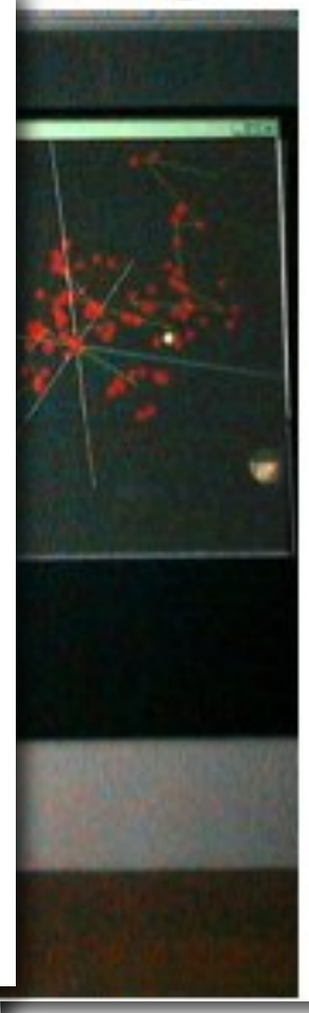

Stephan Bauman -Senior Researcher - German Research Center for AI

ISMIR Tutorial 2005

GenreGram

Weath

Visualization tool for ESOMs

MusicMiner

Kis

Stephan Bauman -Senior Researcher - German Research Center for AI

Core Concepts

How do computers represent music?

- Score - As MIDI data
- Signal - As digital wave form
- Association - Through associations with tags, people, or other songs

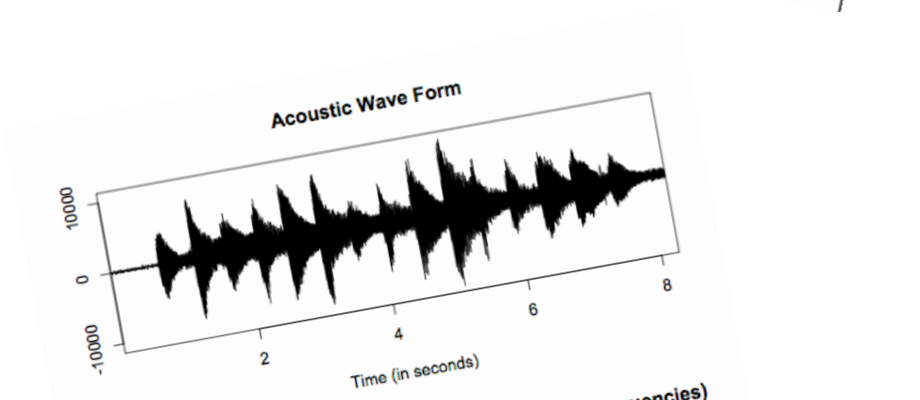

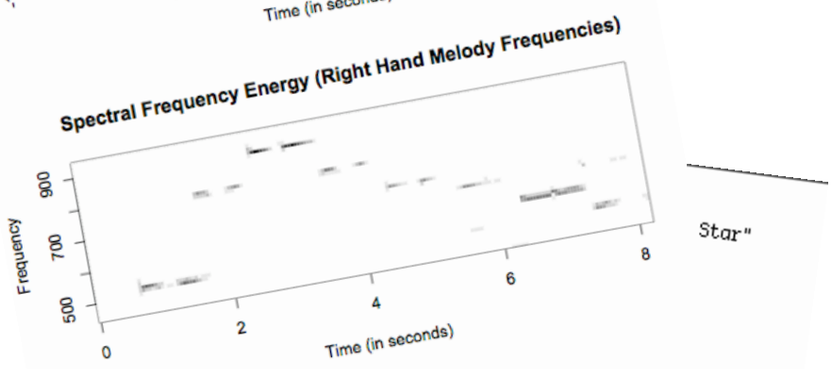

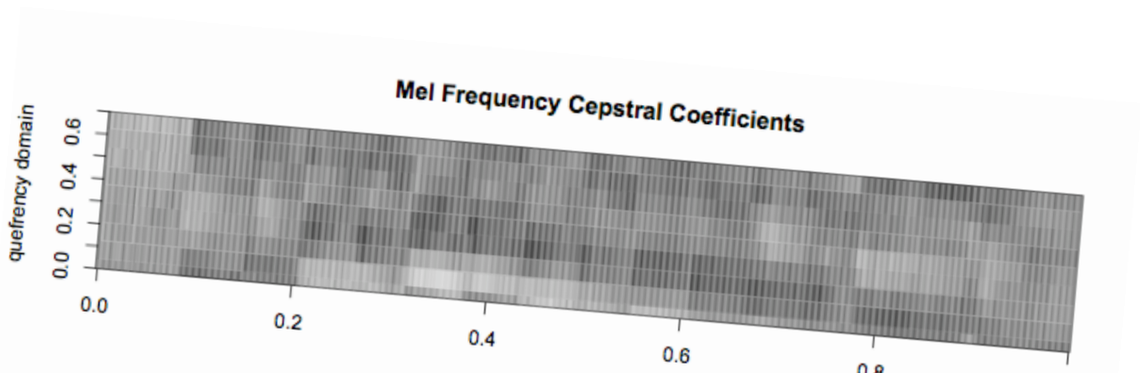

(a brief note on) Symbolic score data

- Typically provides an explicit compositional structure, and sometimes even explicit performance directions.
- “Easy” to analyze/extract meaningful musicological features (time signature, key, etc.)
- Scores are typically not present for most music, so other forms of representation are becoming popular.
- ***We won’t cover score based representations of music today.***

Acoustic Data

- Easy (and even occasionally legal) to find digital wave form representations of music
- New techniques, methods, and even services for music signal analysis.

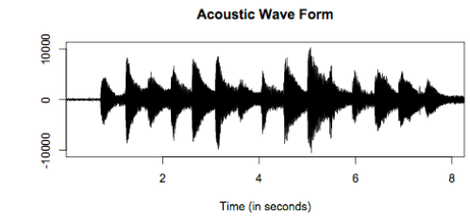

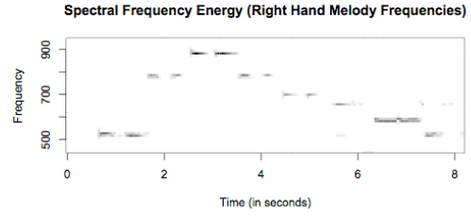

http://www.echonest.com/ http://www.magnatune.com/

Association Data

- Relatively new form of music analysis
- New methods/services as well
- What connections/ associations/relationships does music share with tags, people, playlists, other music?

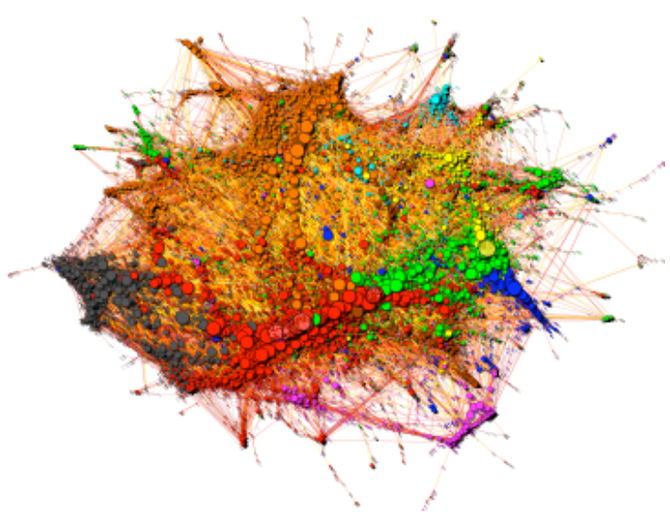

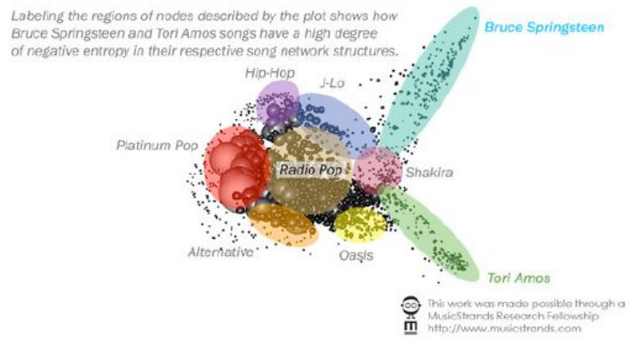

http://www.last.fm
http://sixdegrees.hu/last.fm/
http://www.flickr.com/photos/scwn/153098444/
http://manyeyes.alphaworks.ibm.com/manyeyes/visualizations/matchbox-20-song-lyrics-tag-cloud

Representing "Lots of Songs"

- Many different types of music-related information.
- Problems of sparsity, noise, data set size.
- We need a common appropriate *format of representation* and *method of comparison.*
- We would like to work with matrices.

Song	Feature1	Feature2
"Love me do"	0.4	0.9
"Hound Dog"	3.2	1.4
"Paint it Black"	3.4	1.5

Distances and Similarity

- To form a two/three dimensional visualization, we will need to express a *distance* or *dissimilarity** between the songs in vector form.

- Distance *must be* symmetric in order to represent a metric space.

- Distance must be "non-degenerate" (distance from song A to song A == 0)

Song	"Love me do"	"Hound Dog"	"Paint it Black"
"Love me do"	n/a	3.2	3.4
"Hound Dog"	3.2	n/a	1.5
"Paint it Black"	3.4	1.5	n/a

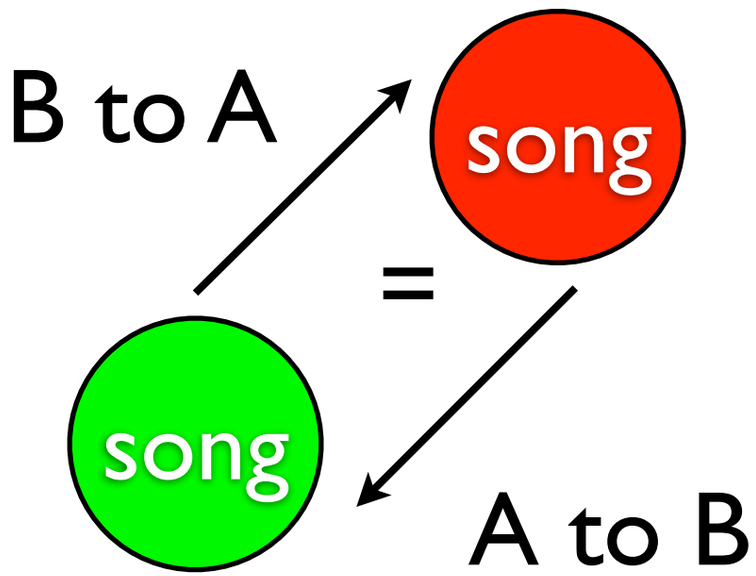

"Big Picture" Visualization

- A couple of "simple" visualizations help to assess variance, etc. across the relevant dimensions.
- Boxplots, matrix plots, image plots, etc.
- Selectively normalizing data at this point to have more control, and removing obvious outlier (points/dimensions).

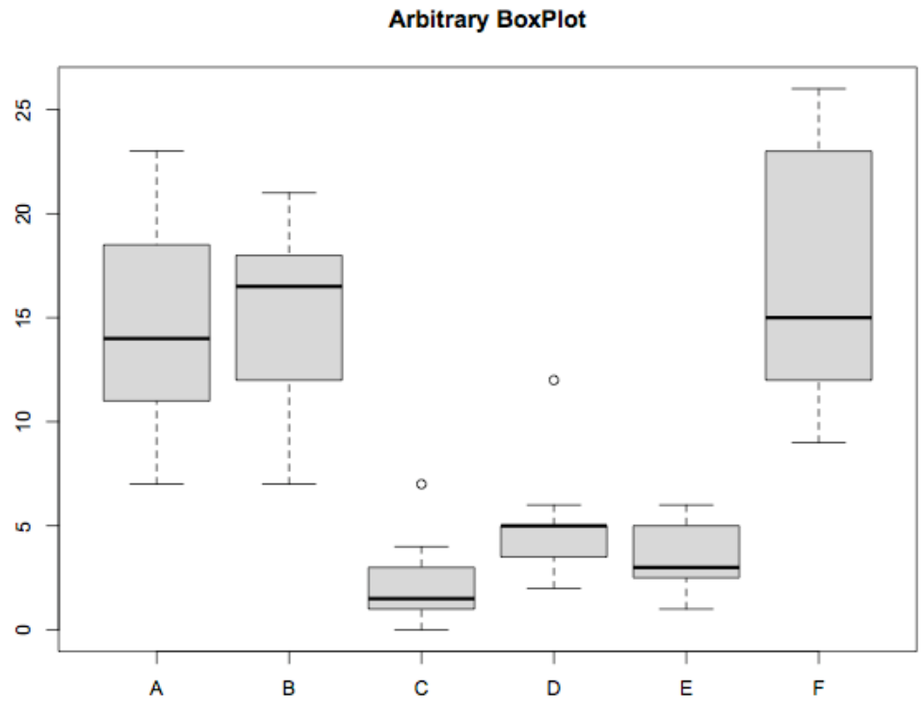

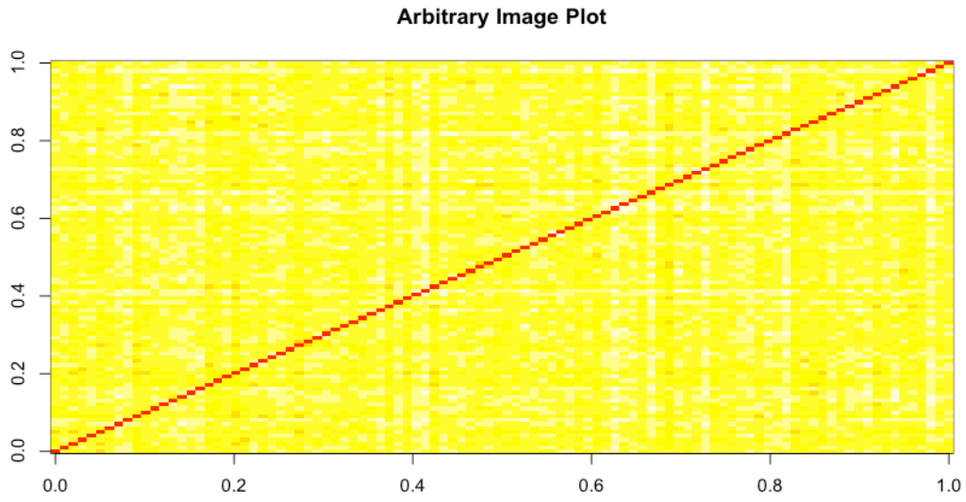

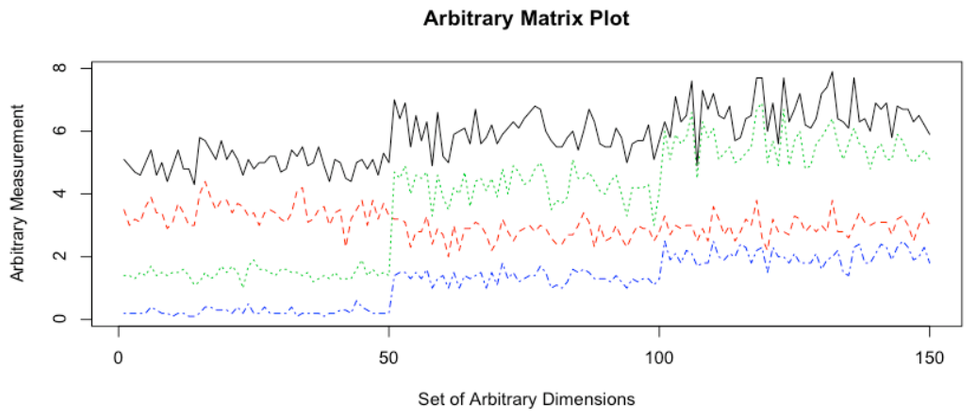

Multidimensional Scaling

- *Reducing* the number of dimensions of the music vector.

	Athens	Barcelona	Brussels	Calais	Cherbourg
Athens	0	3313	2963	3175	3339
Barcelona	3313	0	1318	1326	1294
Brussels	2963	1318	0	204	583
Calais	3175	1326	204	0	460
Cherbourg	3339	1294	583	460	0
Cologne	2762	1498	206	409	785

- Similar to "rotating" the data to find large dimensions of variation/ dissimilarity in the data.

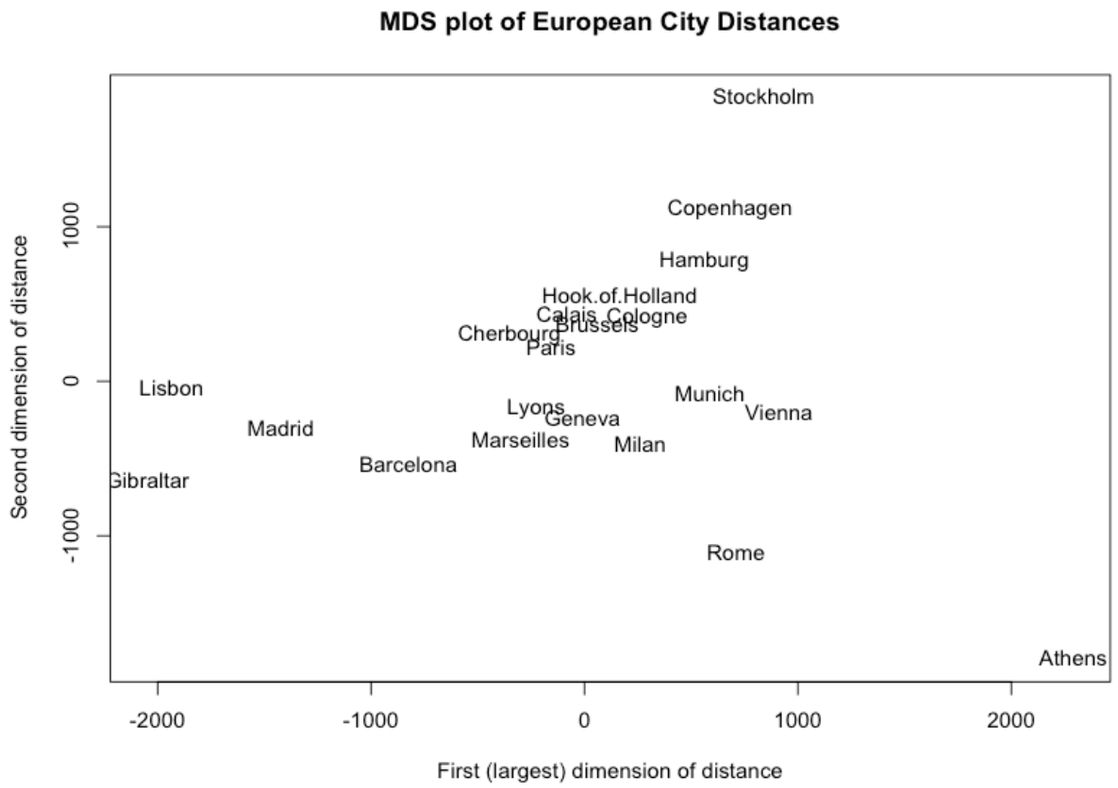

http://www.mathpsyc.uni-bonn.de/doc/delbeke/delbeke.htm

DEMO

Demo: Acoustic Feature Vectors

- Working with *Magnatagatune dataset*
- Combination of:
  - Magnatune creative commons licensed songs (clips)
  - Tagatune project data
  - **Echo Nest feature data**
- Freely available

The dataset consists of the following files:

clip_info_final.csv.bz2	audio clips information, such as title, artist, album, url, start and end time, download URL for the mp3 file (entire song), and path to the mp3 clip.
annotations_final.csv.bz2	tags associated with each audio clip, and path to the clip mp3 file.
comparisons_final.csv.bz2	similarity judgments (number of people who voted that a particular clip is the most different) associated with a tuple of audio clips, and paths to the mp3 clips.
clips.tar	TAR archive with all the audio clips as 16kHz, 32kbps, mono mp3.
mp3_echonest_xml.zip	The Echo Nest analysis of each of the clips in XML format.

http://tagatune.org/Magnatagatune.html

Demo: Feature Vector

Echonest data includes:

- loudness
- tempo
- segments with:
  - pitch information
  - timbre information


```
– <analysis decoder="Quicktime" version="0x7608000">
  – <track duration="29.12331" endOfFadeIn="0.00000" startOfFadeOut="29.12331"
    mode="1" modeConfidence="1.000">
    + <meter></meter>
    + <sections></sections>
    – <segments>
      – <segment start="0.00000" duration="1.64082">
        – <loudness>
            <dB time="0">-58.190</dB>
            <dB time="0.12190" type="max">-22.718</dB>
          </loudness>
        – <pitches>
            <pitch class="0">0.164</pitch>
            <pitch class="1">0.146</pitch>
            <pitch class="2">0.061</pitch>
            <pitch class="3">0.252</pitch>
            <pitch class="4">0.984</pitch>
            <pitch class="5">0.249</pitch>
            <pitch class="6">1.000</pitch>
            <pitch class="7">0.330</pitch>
            <pitch class="8">0.430</pitch>
            <pitch class="9">0.069</pitch>
            <pitch class="10">0.075</pitch>
            <pitch class="11">0.412</pitch>
          </pitches>
        – <timbre>
            <coeff dim="0">22.041</coeff>
            <coeff dim="1">-61.527</coeff>
            <coeff dim="2">51.959</coeff>
            <coeff dim="3">-42.123</coeff>
            <coeff dim="4">17.854</coeff>
            <coeff dim="5">185.841</coeff>
            <coeff dim="6">13.933</coeff>
            <coeff dim="7">0.999</coeff>
            <coeff dim="8">-76.186</coeff>
            <coeff dim="9">99.808</coeff>
            <coeff dim="10">-42.628</coeff>
            <coeff dim="11">16.706</coeff>
          </timbre>
```


http://tagatune.org/Magnatagatune.html

Demo: Feature Vector

Demo motivation

- Genre classification has become a prominent trend in the IR community. Explore the data pertaining to such systems.
- Current approaches can identify "textural" properties of acoustic sound, but can lack cultural or structural properties of music. West & Lamere note that:
  "...a common example might be the confusion of a classical lute timbre, with that of an acoustic guitar string that might be found in folk, pop, or rock music..."

K. West and P. Lamere, "A model-based approach to constructing music similarity functions," *EURASIP Journal on Advances in Signal Processing* 2007 (2007).

Demo: Feature Vector

Demo motivation

- Magnatagatune happens to have a large amount of Classical music, and Classical Lute music to boot.
- We'll try to visualize this part of the magnatagatune data set, along with folk music.

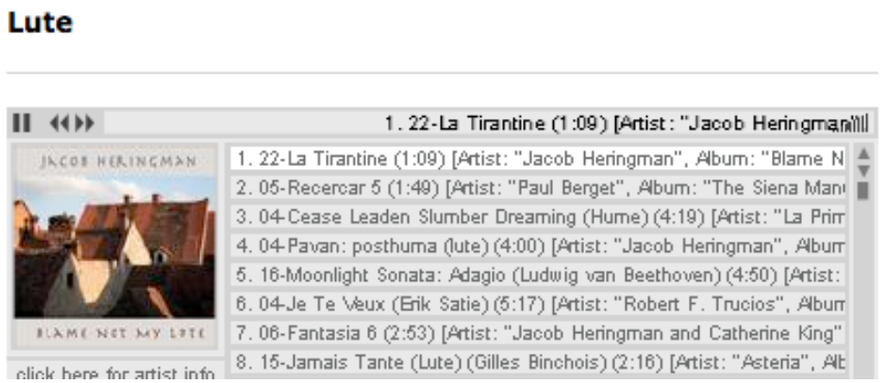

Tools

Tools: Feature Vector

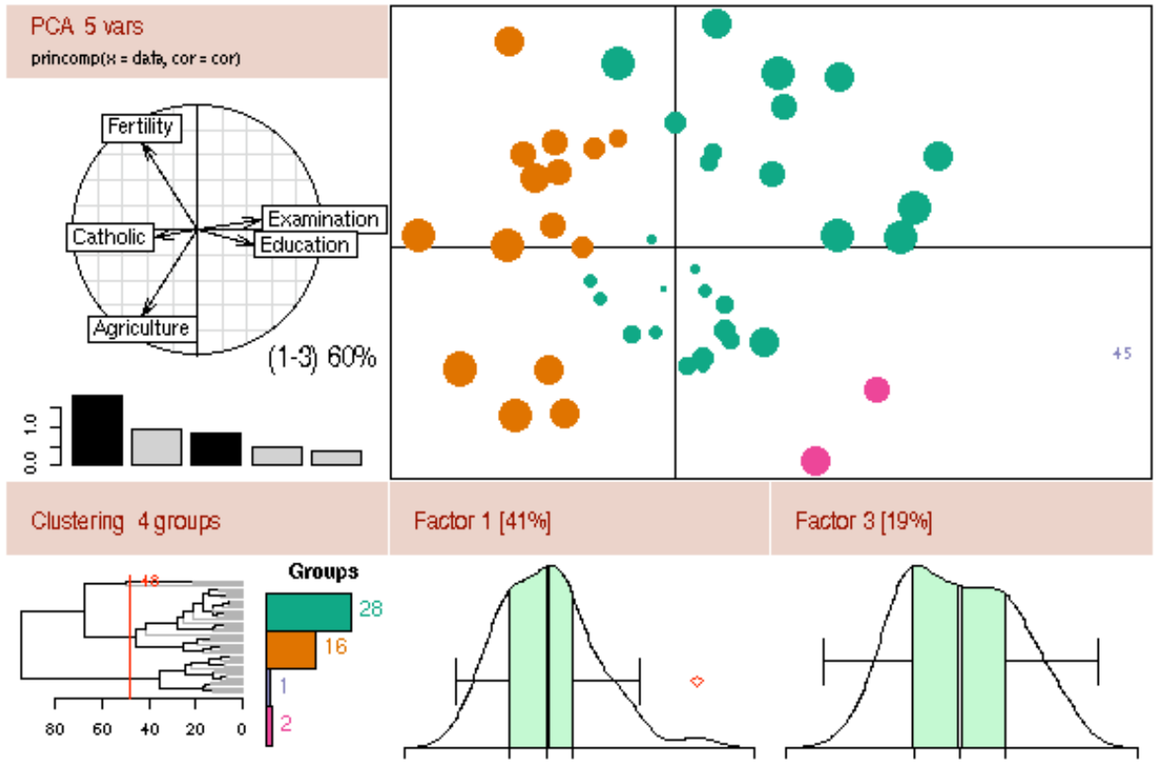

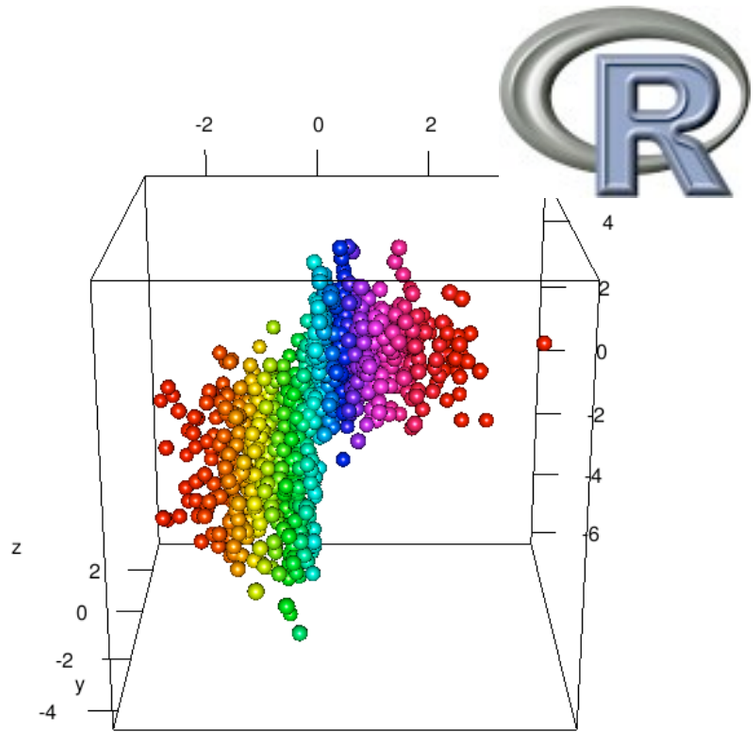

- Download Data
- Use **R** to process, analyze, and visualize data
- Free! (Libre!)

http://cran.r-project.org/

http://www.cyclismo.org/tutorial/R/

http://cran.r-project.org/doc/manuals/R-intro.html

Demo: Feature Vectors

Demo: Feature Vector

Basic Echo Nest Features

- (overall) loudness
- tempo
- pitch means (12) *
- pitch stdev (12) *
- timbre means (12) *
- timber stdev (12) *


```
– <analysis decoder="Quicktime" version="0x7608000">
  – <track duration="29.12331" endOfFadeIn="0.00000" startOfFadeOut="29.12331"
    mode="1" modeConfidence="1.000">
    + <meter></meter>
    + <sections></sections>
    – <segments>
      – <segment start="0.00000" duration="1.64082">
        – <loudness>
            <dB time="0">-58.190</dB>
            <dB time="0.12190" type="max">-22.718</dB>
          </loudness>
        – <pitches>
            <pitch class="0">0.164</pitch>
            <pitch class="1">0.146</pitch>
            <pitch class="2">0.061</pitch>
            <pitch class="3">0.252</pitch>
            <pitch class="4">0.984</pitch>
            <pitch class="5">0.249</pitch>
```


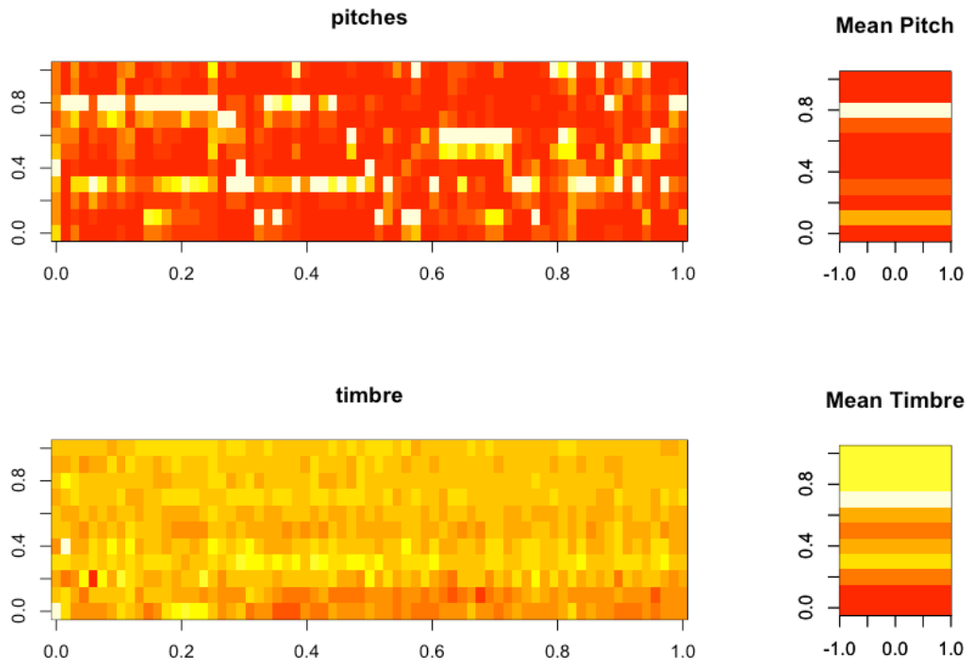

* weighted by segment *onset* loudness

Boxplots

Boxplots show the *five-number summaries* of a distribution:

- sample maximum
- upper quartile
- median
- lower quartile
- sample minimum
- (outliers)

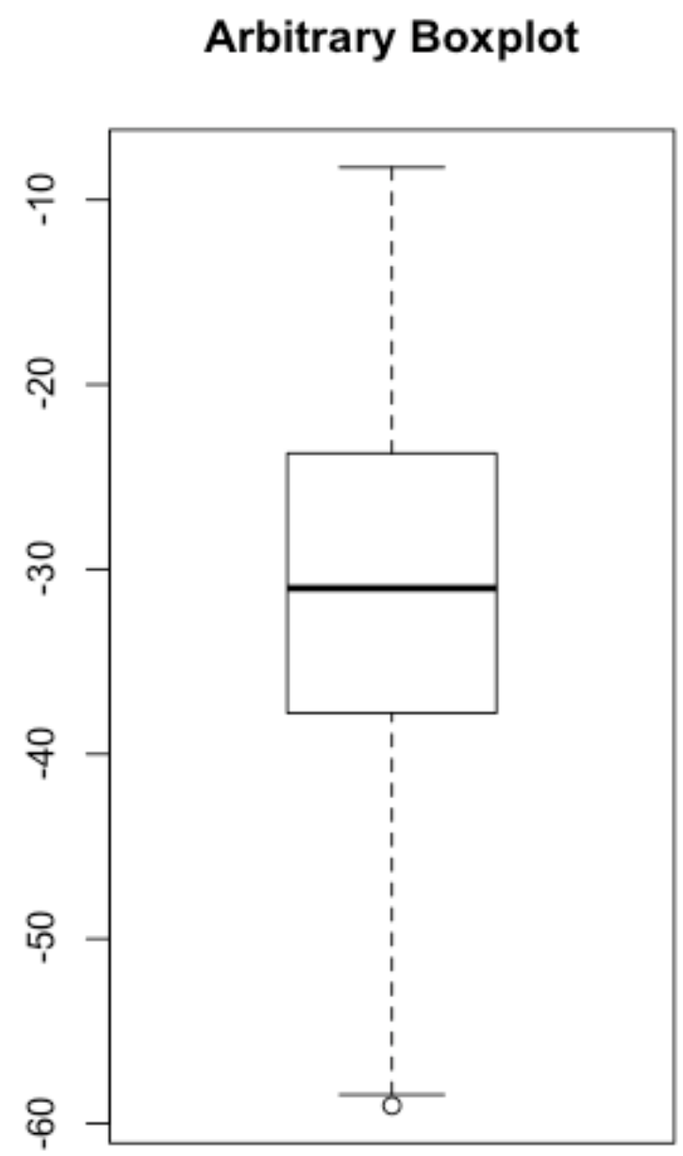

Boxplots

Boxplots show the *five-number summaries* of a distribution:

- sample maximum
- upper quartile
- median
- lower quartile
- sample minimum
- (outliers)

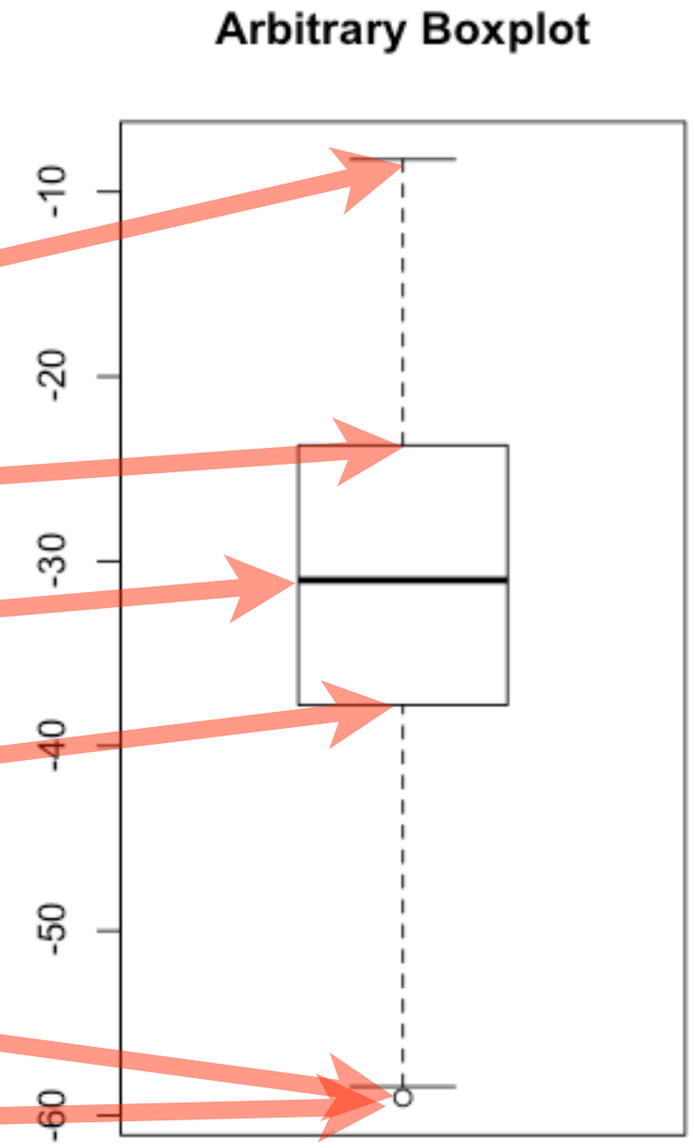

Demo: Feature Vector

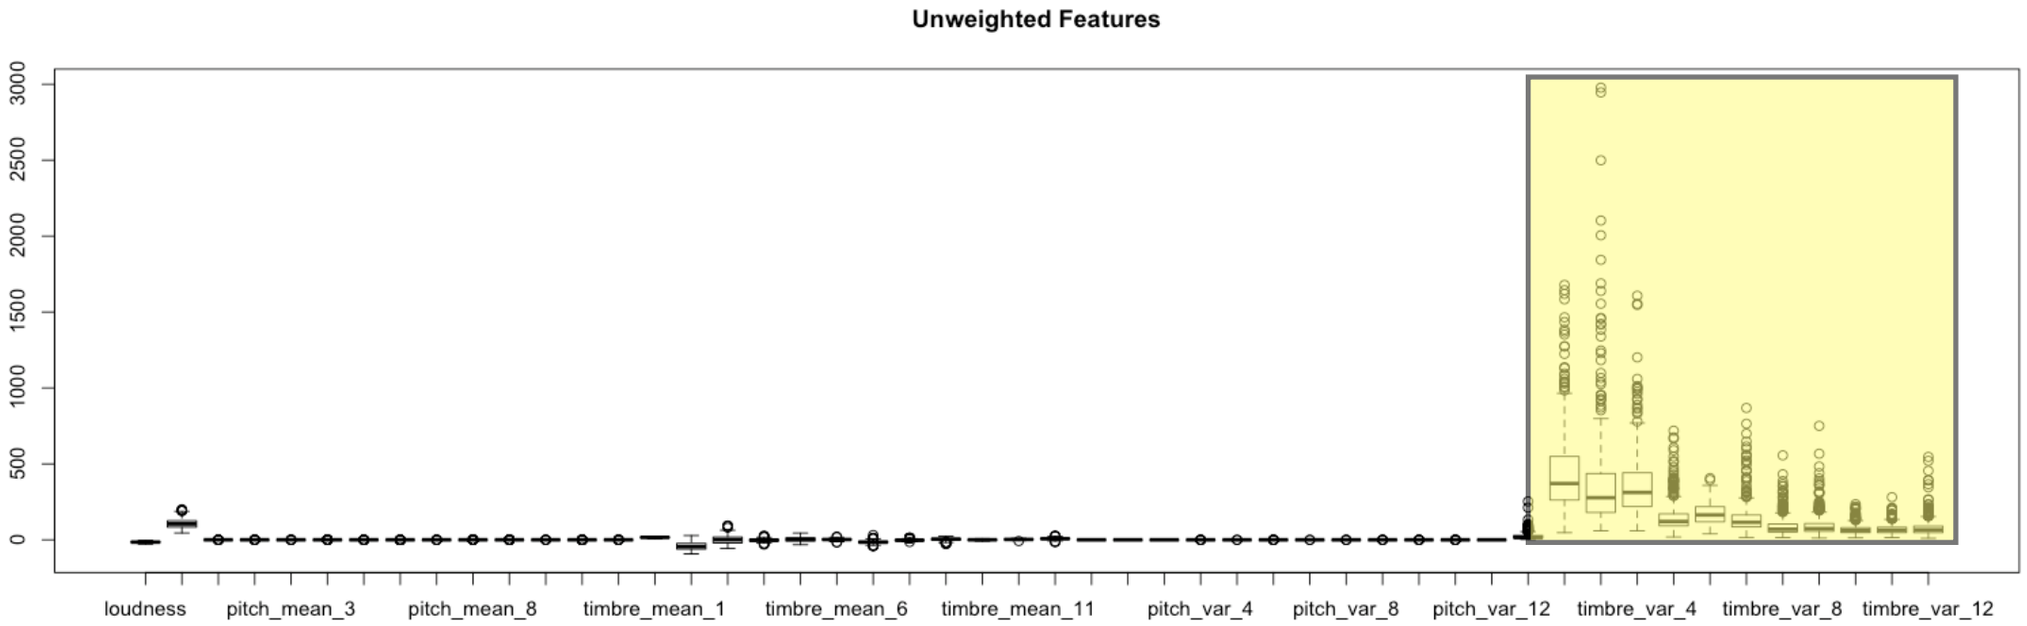

Raw Feature Boxplots...

- Are these all normal distributions?
- Are all of these weighted evenly?
- Scaling becomes important... why?

Demo: Feature Vector

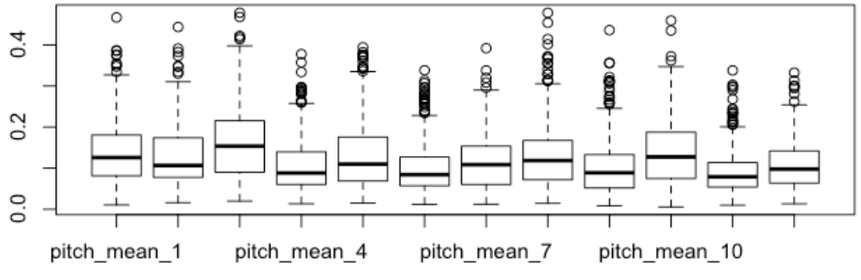

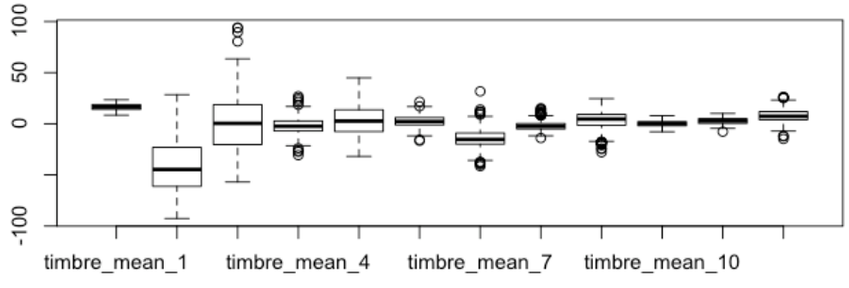

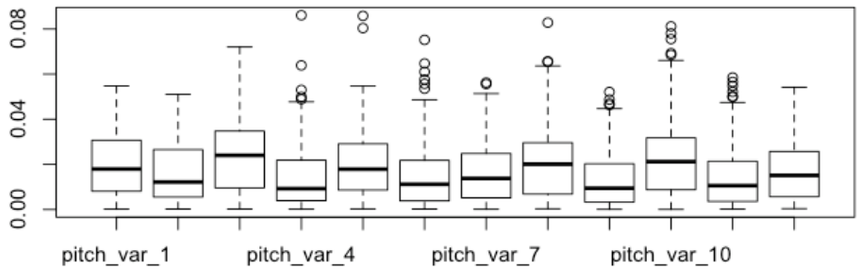

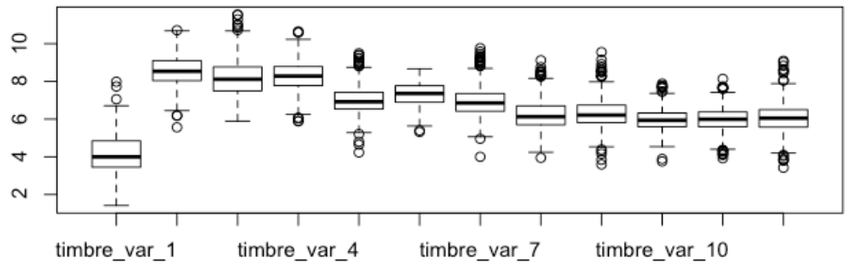

Looking at scaled summary statistics

These are the 'expected' values for any song, so they're 'normal' for a song.

We want to enhance relationships between songs that *differ from the norm*.

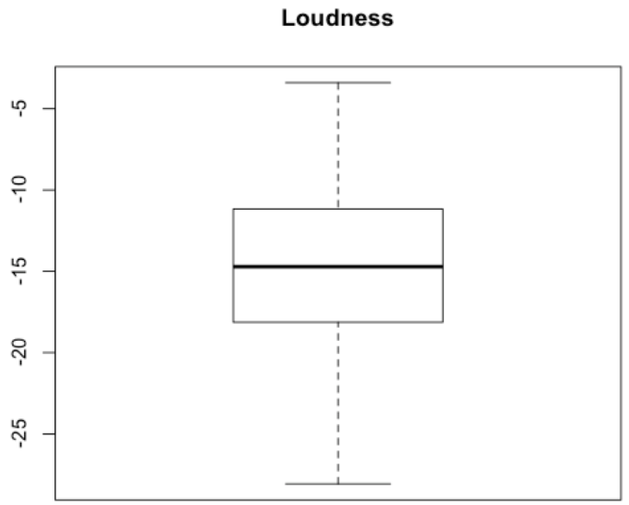

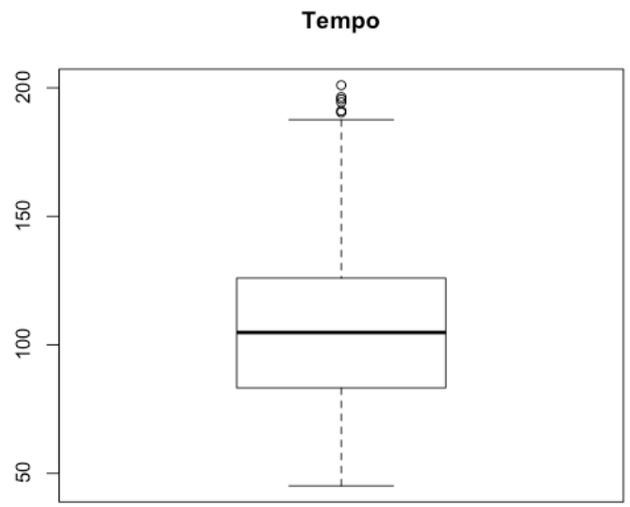

Demo: Feature Vector

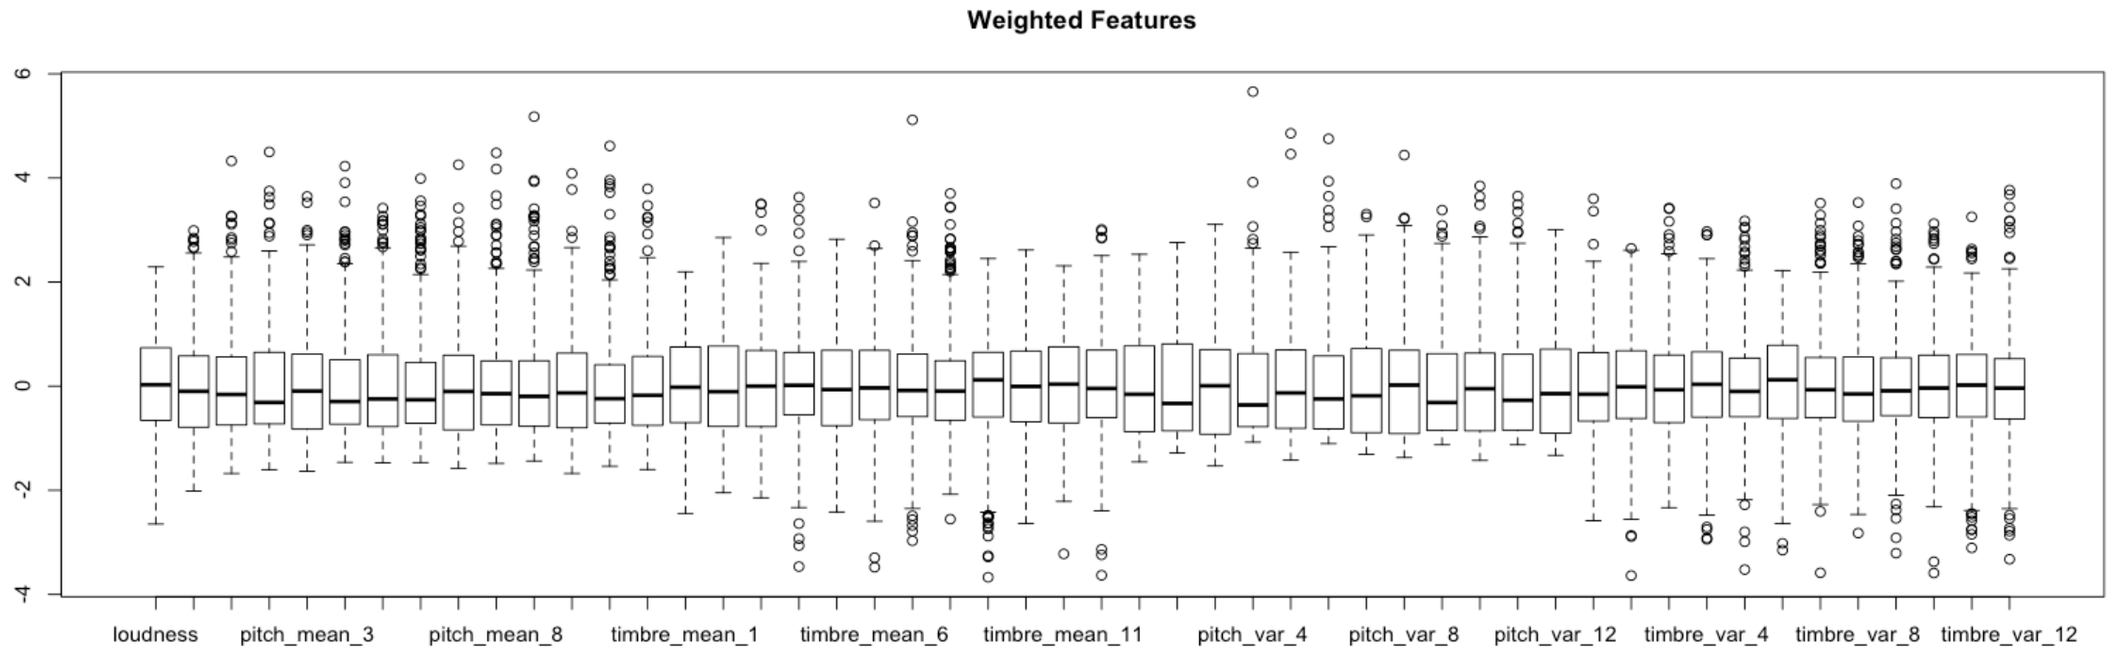

Weighted Feature Boxplots...

- Standard normalized data. All means are 0, all standard deviations are 1.
- We can find 'interesting' relationships between songs better this way. Distances can be understood in terms of *standard deviations from the mean.*
- However... 24 pitch and 24 timbre features, vs. just 1 tempo...

Demo: Feature Vector

Principle Component Analysis (PCA) of timbre and pitch information reveal a high degree of covariance.

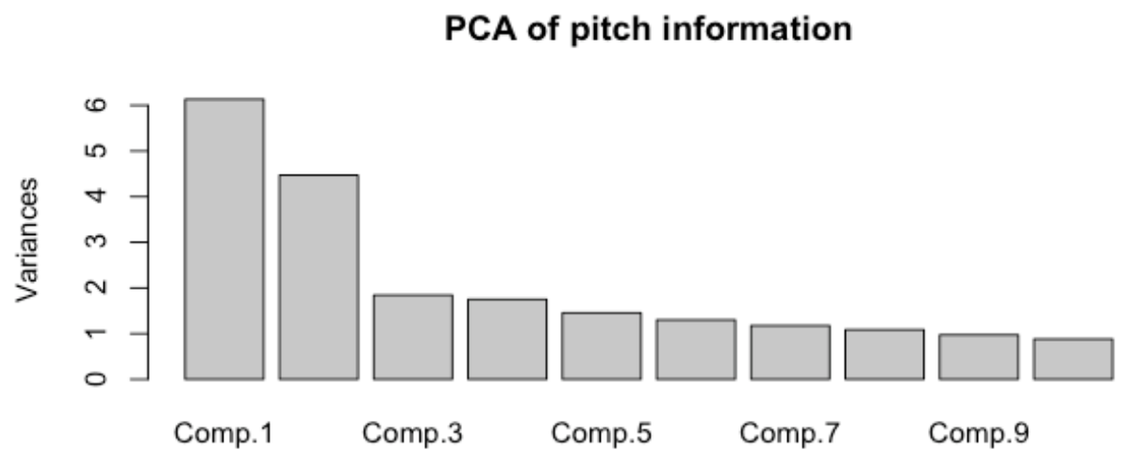

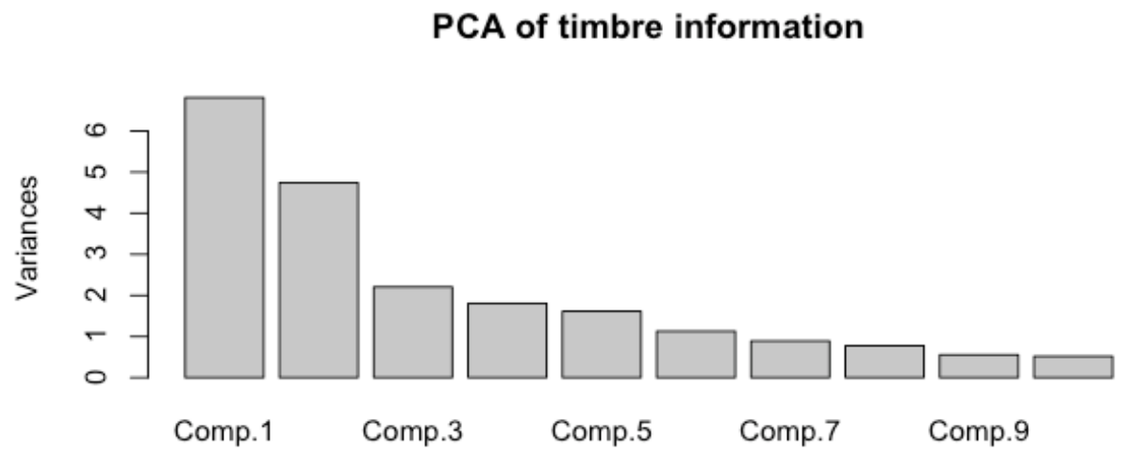

Demo: Feature Vector

Plotting the PCA loadings show the 'structure' of how the pitch profiles co-vary across the music.

Blue = "Classical Guitar"
(almost entirely lute music)

Red = "Folk"

Green= "Both"

(flips and rotations)

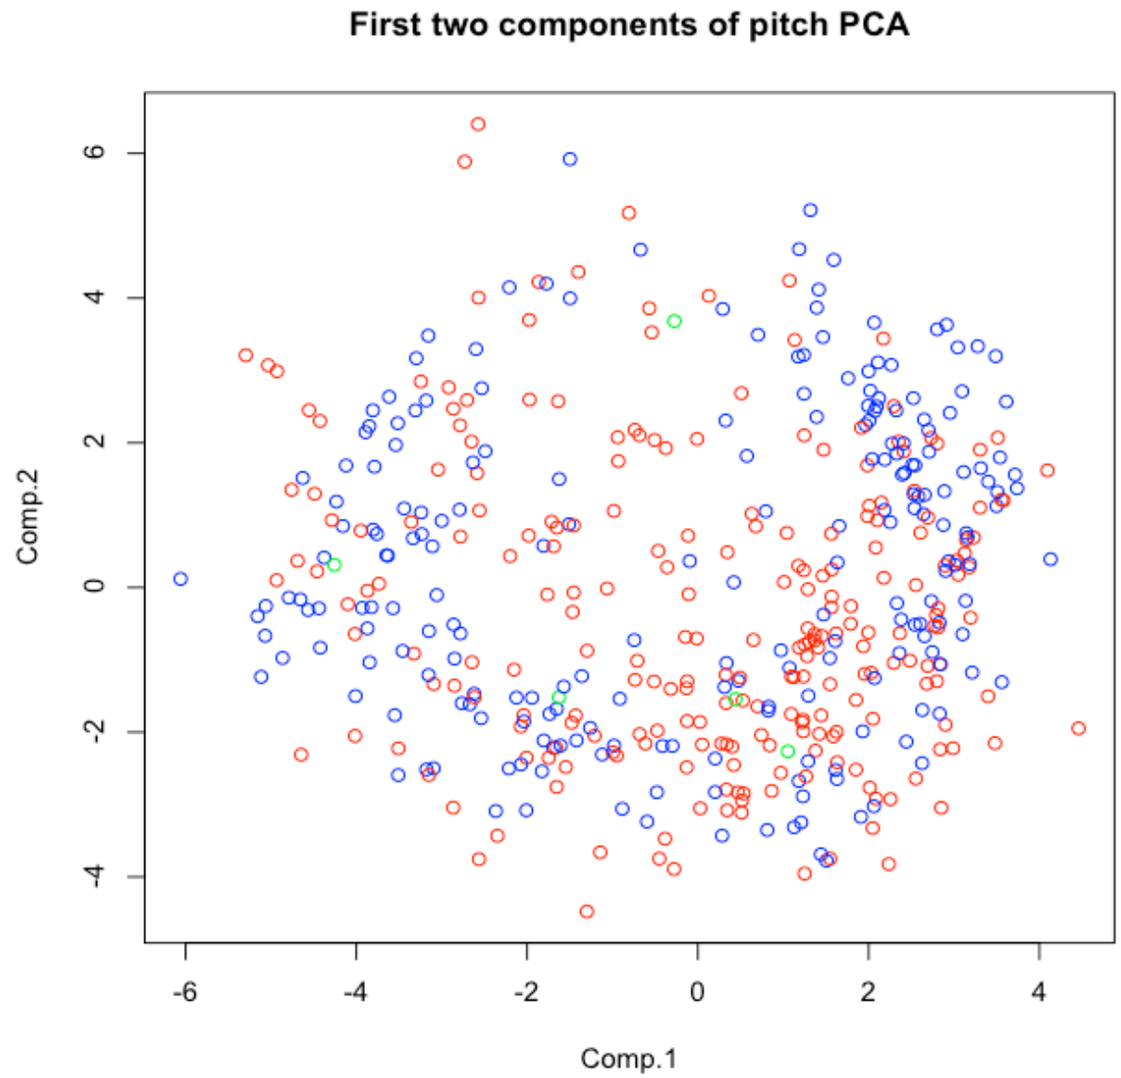

Demo: Feature Vector

Performing a Factor analysis shows evidence of "musicological" factors.

(alternating positive/negative loadings for both means and variances, indicating certain "keys" are followed for songs)

Should pitch profile distances be figured another way?

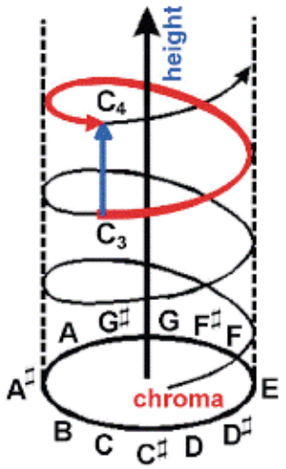

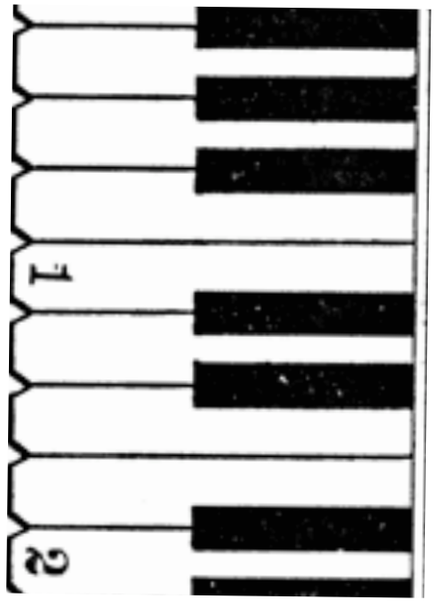

```
Loadings:
                Factor1 Factor2
pitch_mean_1    -0.117   0.625
pitch_mean_2     0.576  -0.270
pitch_mean_3    -0.585
pitch_mean_4     0.677   0.126
pitch_mean_5    -0.103  -0.487
pitch_mean_6             0.473
pitch_mean_7     0.332  -0.504
pitch_mean_8    -0.512   0.217
pitch_mean_9     0.527  -0.124
pitch_mean_10   -0.520  -0.274
pitch_mean_11    0.475   0.452
pitch_mean_12    0.216  -0.359
pitch_var_1     -0.318   0.689
pitch_var_2      0.638  -0.377
pitch_var_3     -0.661   0.100
pitch_var_4      0.695   0.289
pitch_var_5     -0.169  -0.536
pitch_var_6              0.658
pitch_var_7      0.363  -0.550
pitch_var_8     -0.660   0.315
pitch_var_9      0.669  -0.104
pitch_var_10    -0.620  -0.213
pitch_var_11     0.385   0.593
pitch_var_12     0.124  -0.399
```

Demo: Feature Vector

Timbre features clearly split the different tags we were concerned with.

Blue = “Classical Guitar”

Red = “Folk”

Green= “Both”

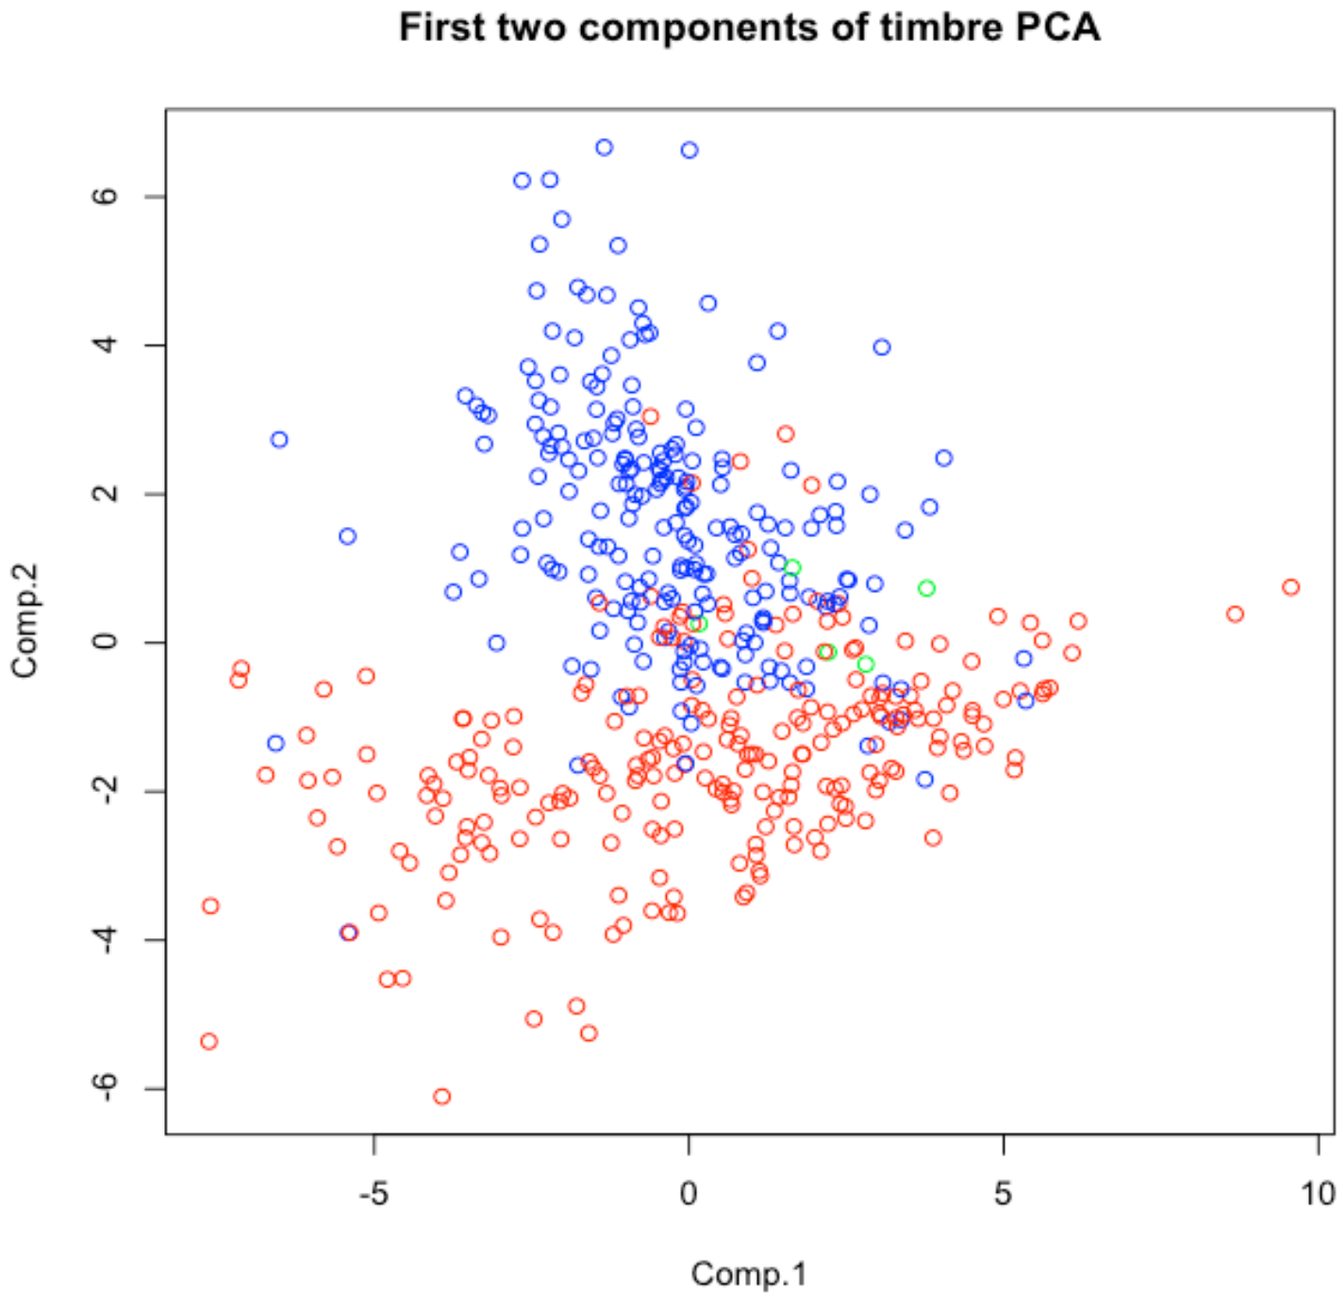

Demo: Feature Vector

Timbral factors are a lot more straightforward. Certain songs have a lot of timbral variance. Other songs simply have strong overall timbral means.

```
Loadings:
                Factor1 Factor2
timbre_mean_1             0.987
timbre_mean_2             0.489
timbre_mean_3   0.151     0.728
timbre_mean_4  -0.139     0.349
timbre_mean_5            -0.570
timbre_mean_6
timbre_mean_7  -0.126     0.248
timbre_mean_8   0.175    -0.556
timbre_mean_9   0.122     0.589
timbre_mean_10  0.257    -0.461
timbre_mean_11 -0.363    -0.238
timbre_mean_12
timbre_var_1    0.582     0.614
timbre_var_2    0.681    -0.286
timbre_var_3    0.822    -0.114
timbre_var_4    0.680    -0.387
timbre_var_5    0.779
timbre_var_6    0.241    -0.690
timbre_var_7    0.810
timbre_var_8    0.712
timbre_var_9    0.792
timbre_var_10   0.591    -0.358
timbre_var_11   0.667    -0.140
timbre_var_12   0.736
```

Demo: Association Data

Demo: Association Data

	clip_id	no.voice	singer	duet	plucking	hard.rock	world	bongos	harpsichord	female.singing	classical
1	2	0	0	0	0	0	0	0	0	0	0
2	6	0	0	0	0	0	0	0	1	0	0
3	10	0	0	0	1	0	0	0	0	0	0
4	11	0	0	0	0	0	0	0	0	0	0
5	12	0	1	0	0	0	0	1	0	0	1
6	14	0	0	0	0	0	0	0	0	0	0

- Magnatagatune has ‘tag’ data for songs
- Tags only exist if *two* people used the same tag for the same song
- 188 total tags

Demo: Association Data

guitar	classical	slow	techno	strings	drums	electronic	rock
4852	4272	3547	2954	2729	2598	2519	2371
fast	piano	ambient	beat	violin	vocal	synth	female
2306	2056	1956	1906	1826	1729	1717	1474
indian	opera	male	singing	vocals	no.vocals	harpsichord	loud
1395	1296	1279	1211	1184	1158	1093	1086

- Tags are distributed by a *power-law*
- It's not as meaningful to share a "guitar" tag. A "soprano" tag is far more informative.
- We need to form a comparison of tags based on the *weighted relevance* of the terms.

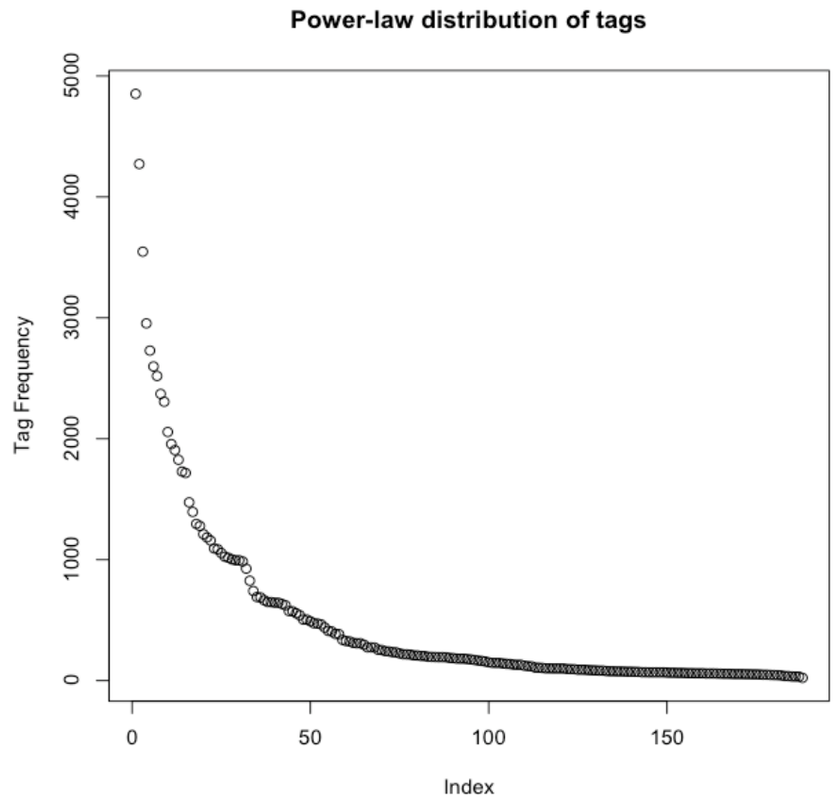

Demo: Association Data

- TF/IDF (term frequency/document frequency)
- n(i,j) = number of times a term i appears in a ‘document’ j.
- n(k,j) = total number of terms in ‘document’ j.
- |D| = total number of ‘documents’
- {d:...} = total documents containing i
- documents = tags for song

$$\mathrm{tf_{i,j}} = \frac{n_{i,j}}{\sum_k n_{k,j}}$$

$$\mathrm{idf_i} = \log \frac{|D|}{|\{d : t_i \in d\}|}$$

http://en.wikipedia.org/wiki/Tf%E2%80%93idf

Demo: Association Data

- Forming a similarity/dissimilarity between weighted tags commonly involves finding the *angle* between two vectors.
- Cosine dissimilarity is one (of many) ways of expressing a distance between vectors.
- A, B = vectors of attributes
- ||A||, ||B|| = magnitudes of vectors

$$\text{similarity} = \cos(\theta) = \frac{A \cdot B}{\|A\|\|B\|}.$$

$$\|\mathbf{x}\| := \sqrt{x_1^2 + \cdots + x_n^2}.$$

http://en.wikipedia.org/wiki/Cosine_similarity

Demo: Association Data

Cosine similarity of tags also (not surprisingly) clearly separates the two tags

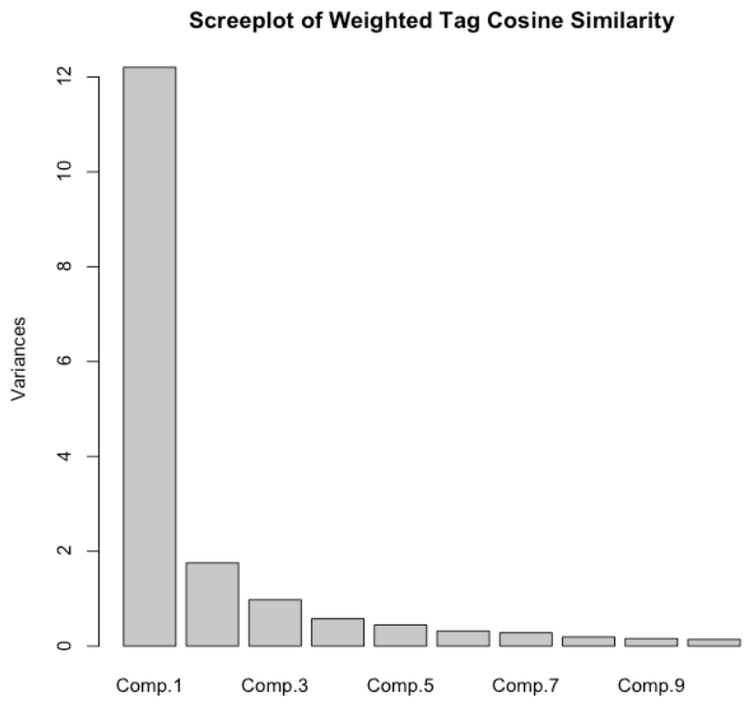

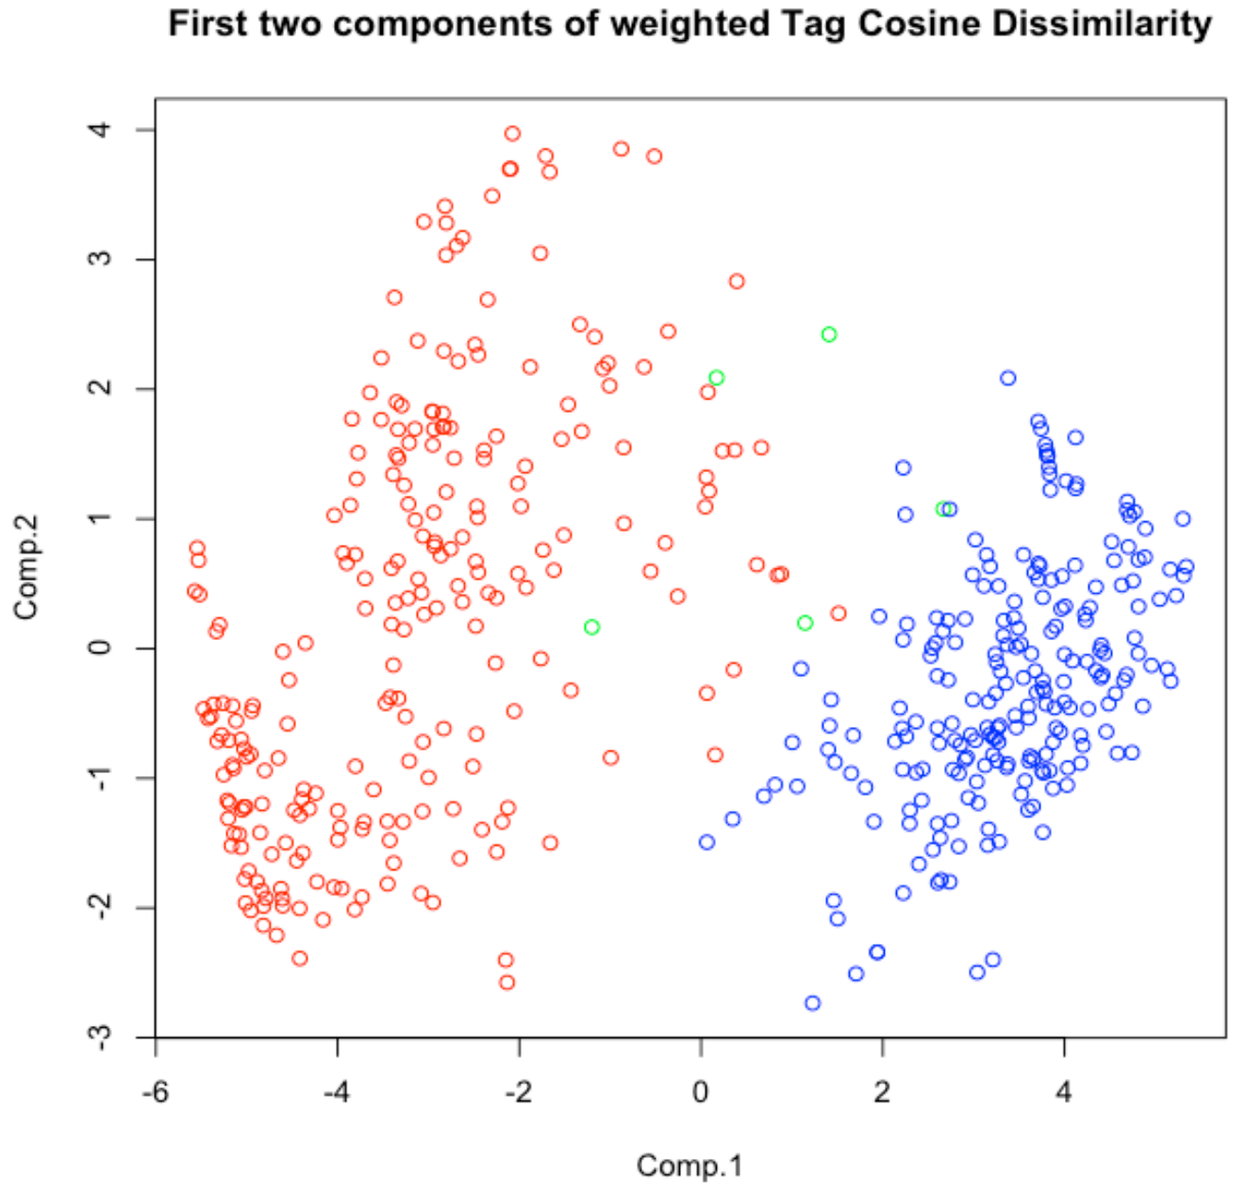

Demo: Hybrid

Demo: Feature Vector

Timbre features clearly split the different tags we were concerned with.

Blue = “Classical Guitar”

Red = “Folk”

Green= “Both”

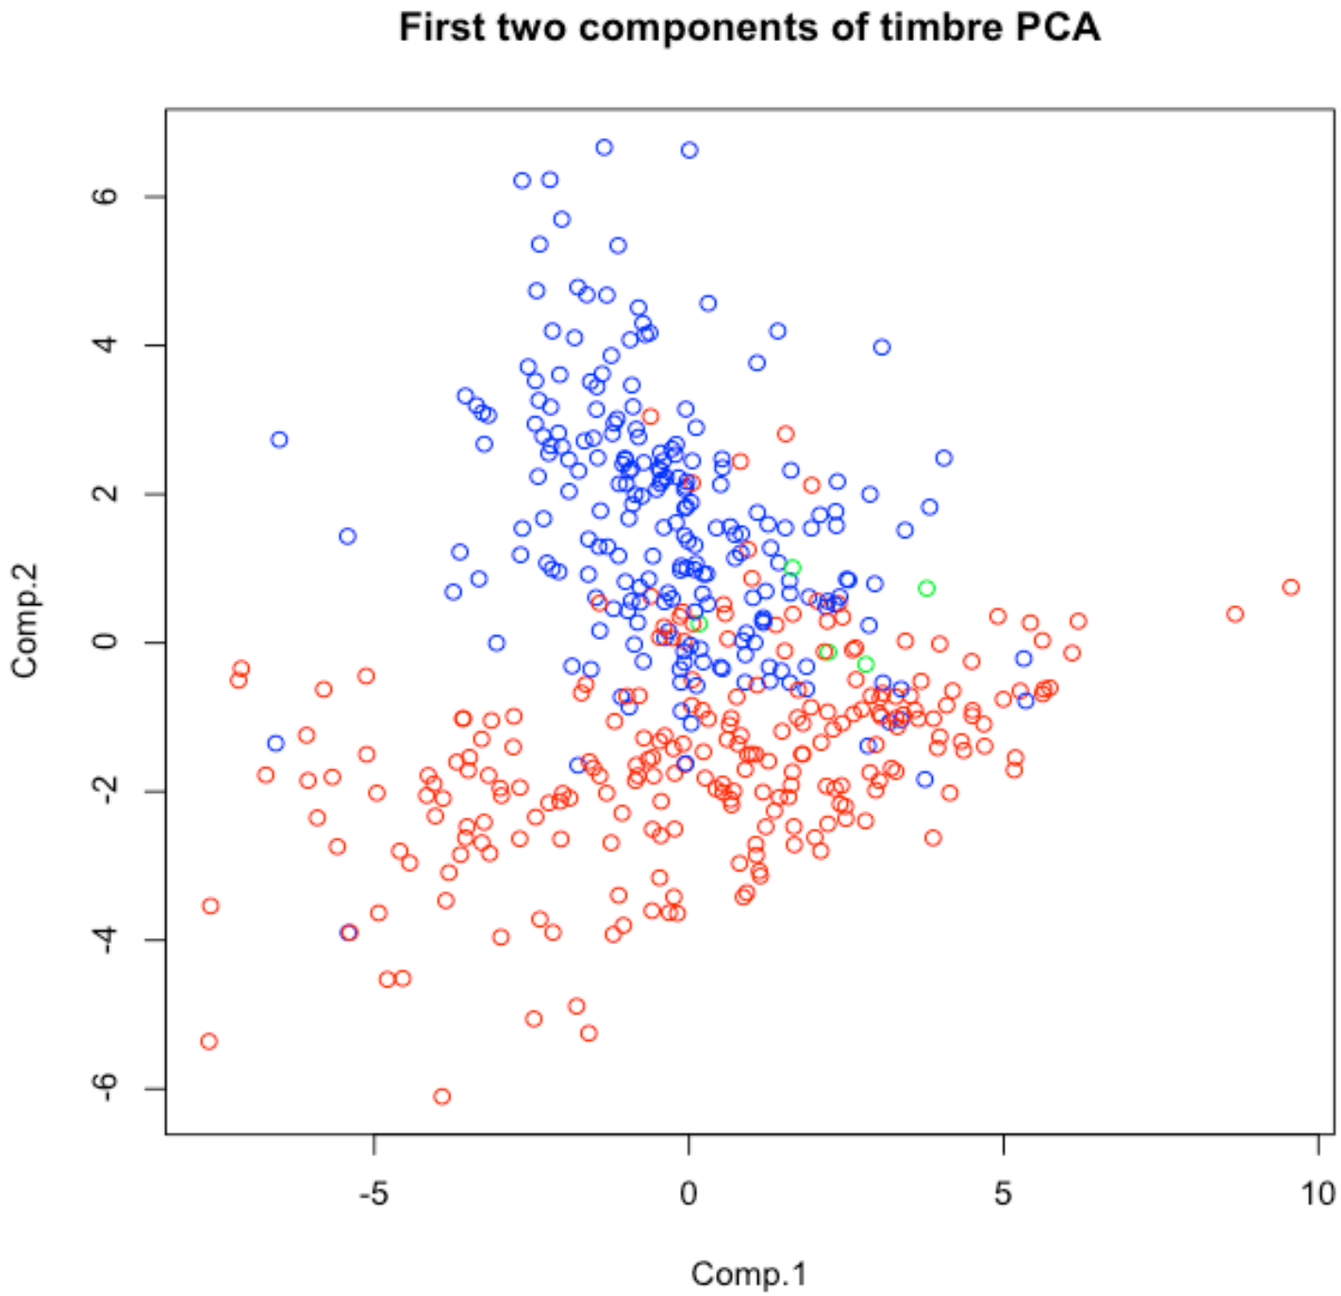

Conclusions (demo)

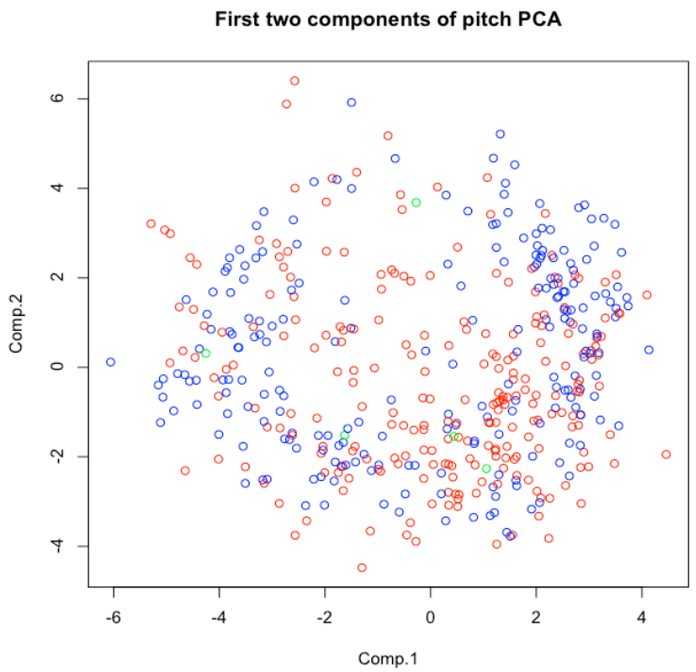

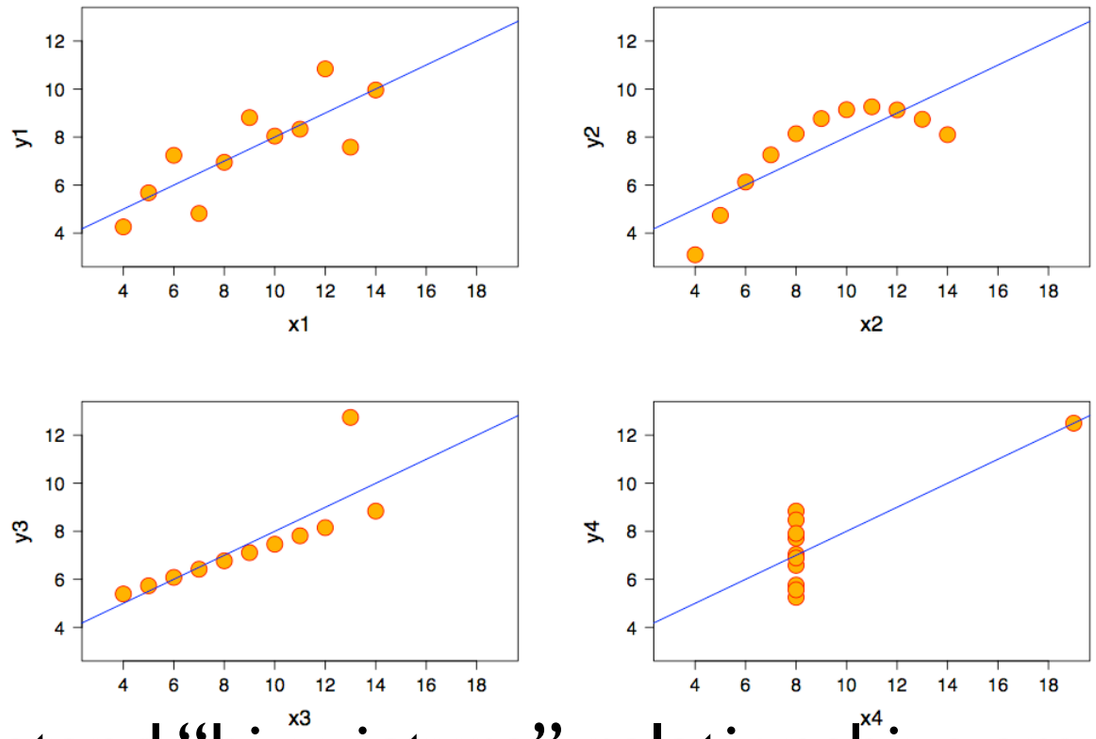

- Visualizations can help understand "big picture" relationships among songs, and support that the relationships are meaningful. (left plot)
- Visualizations can help understand "mismatches" in modeling/clustering/ classification approaches. (x2)
- Visualizations can help detect noise (x1) and outliers. (x3, x4)
- Visualizations can help find a way to *improve the underlying model*.

Opportunity: Hybrid Maps

Mixing information from both timbre & term spaces has a great opportunity for forming a 'better' representation of genre

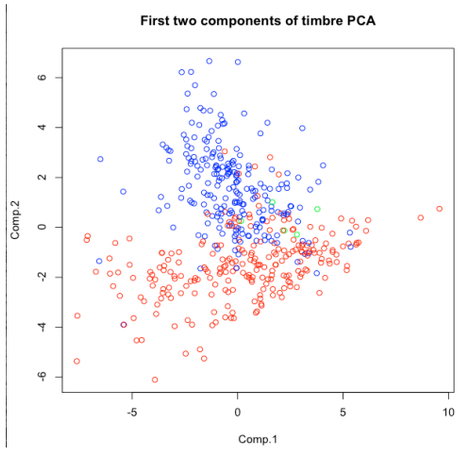

+

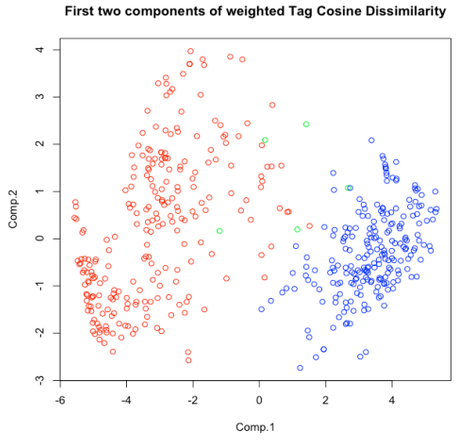

=

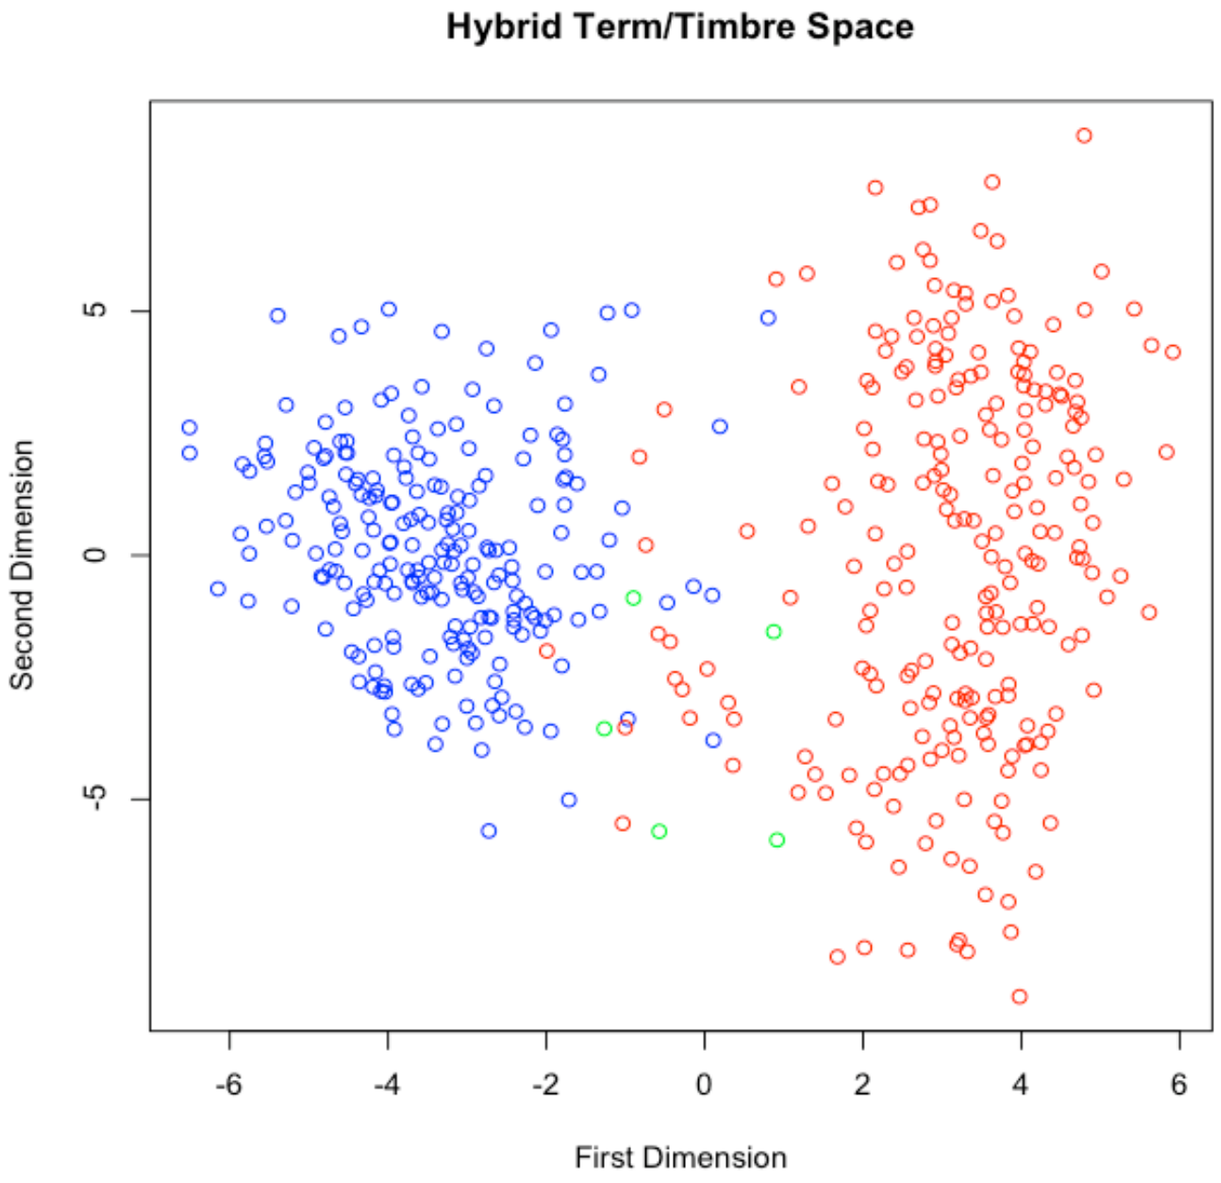

"Reading" Conventional Dimensionality Reductions

- *Large* dissimilarity is preserved over *small dissimilarity.*

- Try to understand the plot in terms of its two biggest *orthogonal* distances. If those don't make sense, then the plot is suspect.

- Next, look for local structure, clustering, etc.

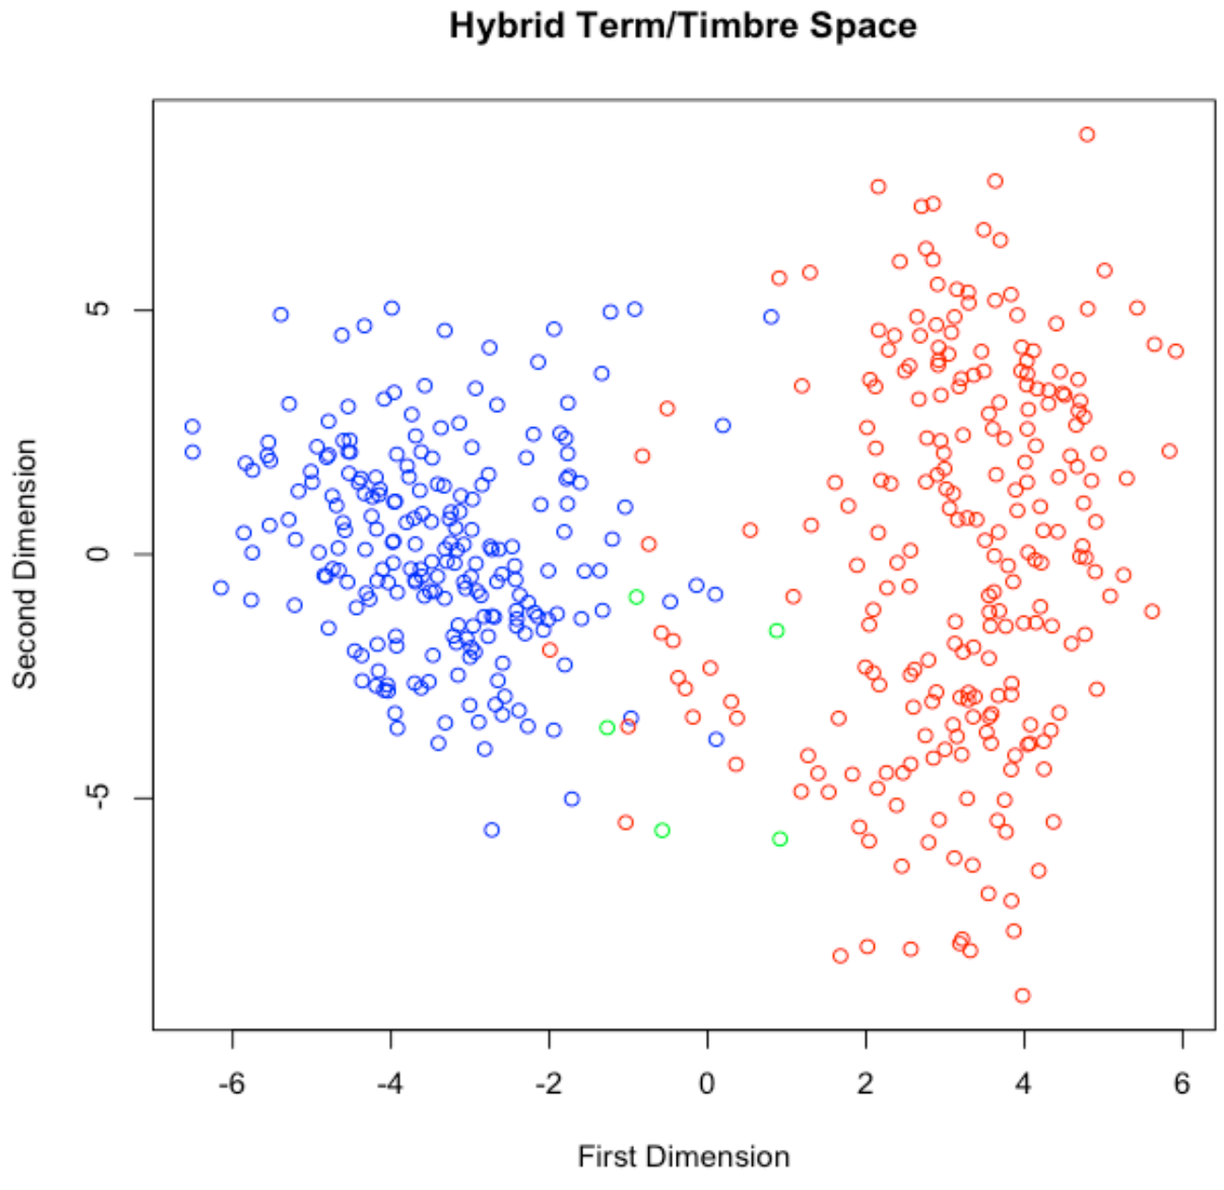

Other Techniques

Other Techniques

- Locally Linear Embedding : Only consider ‘close’ points in distance matrix. Good at recovering from “kinks”, “folds”, and “spirals” in data (chroma)

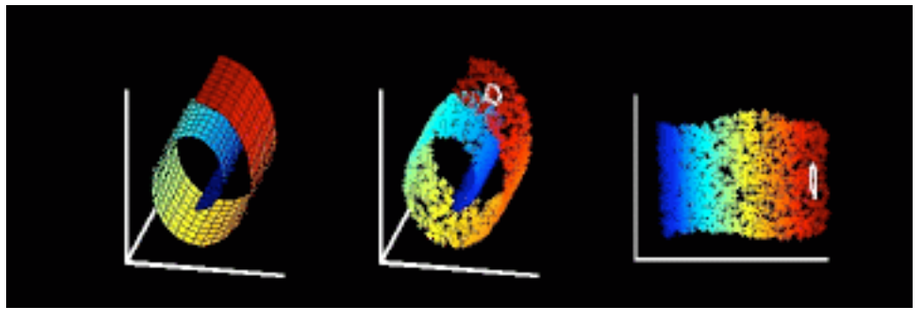

- Isomap : Establish geodesic “neighborhood” distances using Dijkstra’s algorithm, and solve with mds.

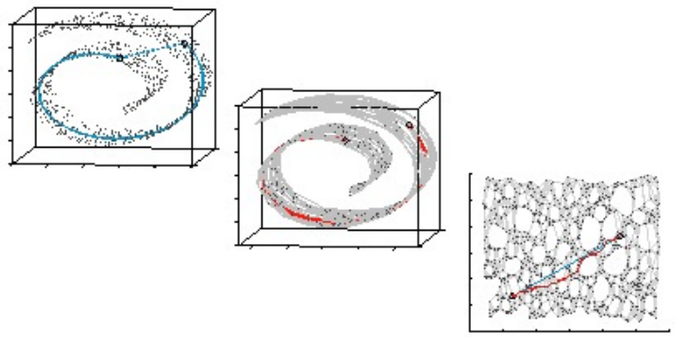

- T-SNE : Convert neighbor distances to conditional probabilities. Perform gradient descent on Kullback-Leibler Divergence. Good at recovering clusters at different *scales*.

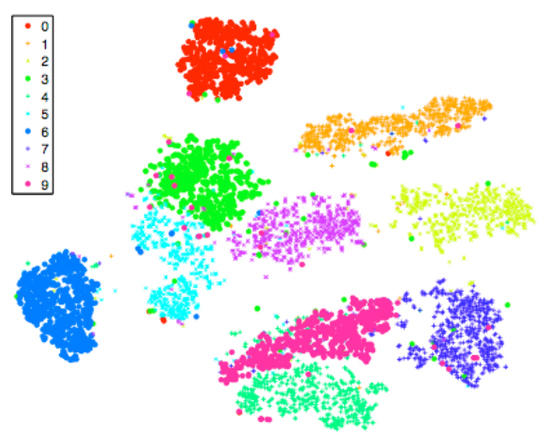

(a) Visualization by t-SNE.

Other Techniques

(observations)

- Many advanced techniques sacrifice "large scale" dissimilarities in order to preserve/enhance "local" features.

- Sometimes large scale information is useful (to determine if scales are wrong, or if there are significant outliers)

- Interpretation of advanced techniques is much more complicated. (What are the loadings, factors, of features etc.?)

- It can be harder to use advanced dimensionality reduction techniques to inform classifier design.

Building an artist map

The Goal - build an artist graph suitable for music exploration

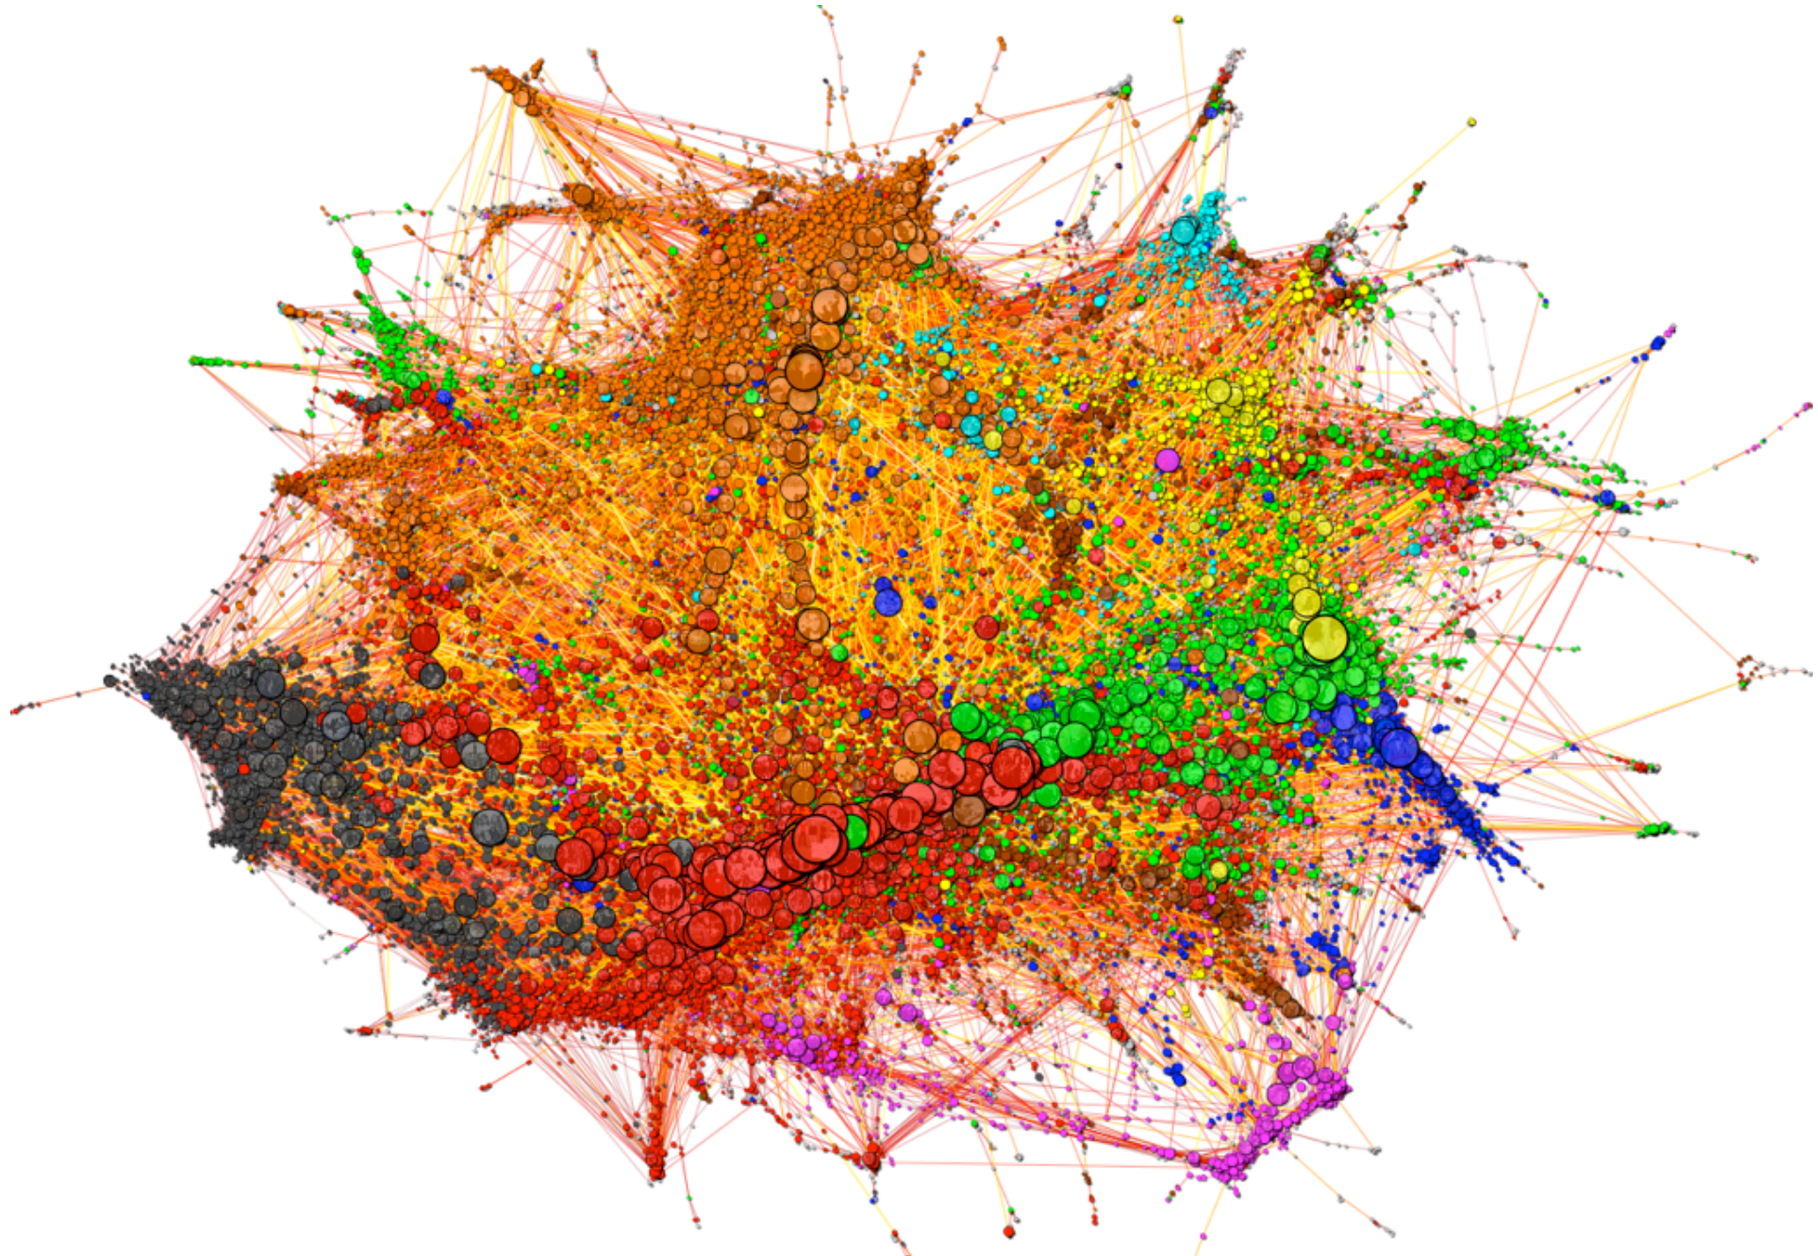

Problem: The artist space is very complex

The Data

- Crawled The Echo Nest API for:
  - Artist names
  - Artist hotness
  - Artist familiarity
  - 15 most similar artists
- 70,000 artists
- 250,000 artist-artist links
- Available at musicviz.googlepages.com

Version 3
Overview
Data Types
Version History

API Methods

Artist:
get_audio
get_biographies
get_blogs
get_familiarity
get_hotttnesss
get_image
get_news
get_profile
get_reviews
get_similar
get_top_terms
get_urls
get_video

Making a good graph

- Minimize edge crossings, bends and curves

Good

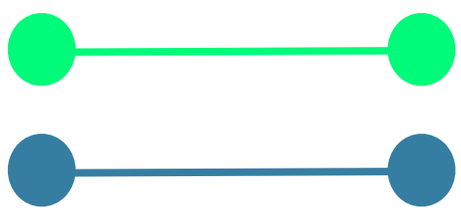

Bad

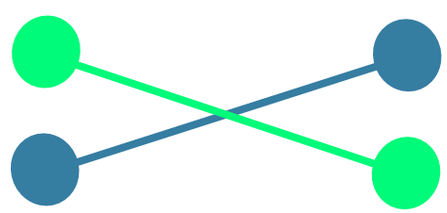

- Maximize angular resolution, symmetry

Good

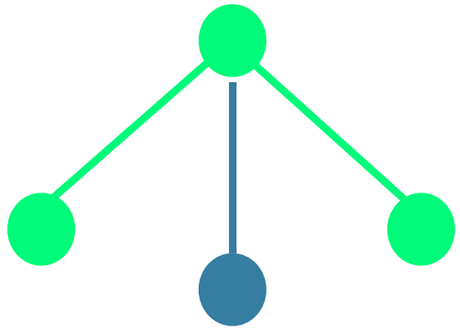

Bad

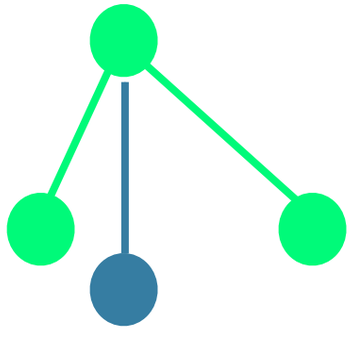

Cognitive measurements of graph aesthetics, Ware, Purchase, Colpoys, McGill

Graphing the Beatles Neighborhood

```
digraph {
    "The Beatles" ->  "Electric Light Orchestra";
    "The Beatles" ->  "The Rutles";
    "The Beatles" ->  "The Dave Clark Five";
    "The Beatles" ->  "The Tremeloes";
    "The Beatles" ->  "Manfred Mann";
    "The Beatles" ->  "The Lovin' Spoonful";
    "The Beatles" ->  "The Rolling Stones";
    "The Beatles" ->  "Creedence Clearwater Revival";
    "The Beatles" ->  "Bee Gees";
    "The Beatles" ->  "the Hollies";
    "The Beatles" ->  "Badfinger";
    "The Beatles" ->  "Paul McCartney";
    "The Beatles" ->  "Julian Lennon";
    "The Beatles" ->  "George Harrison";
    "The Beatles" ->  "John Lennon";
}
```

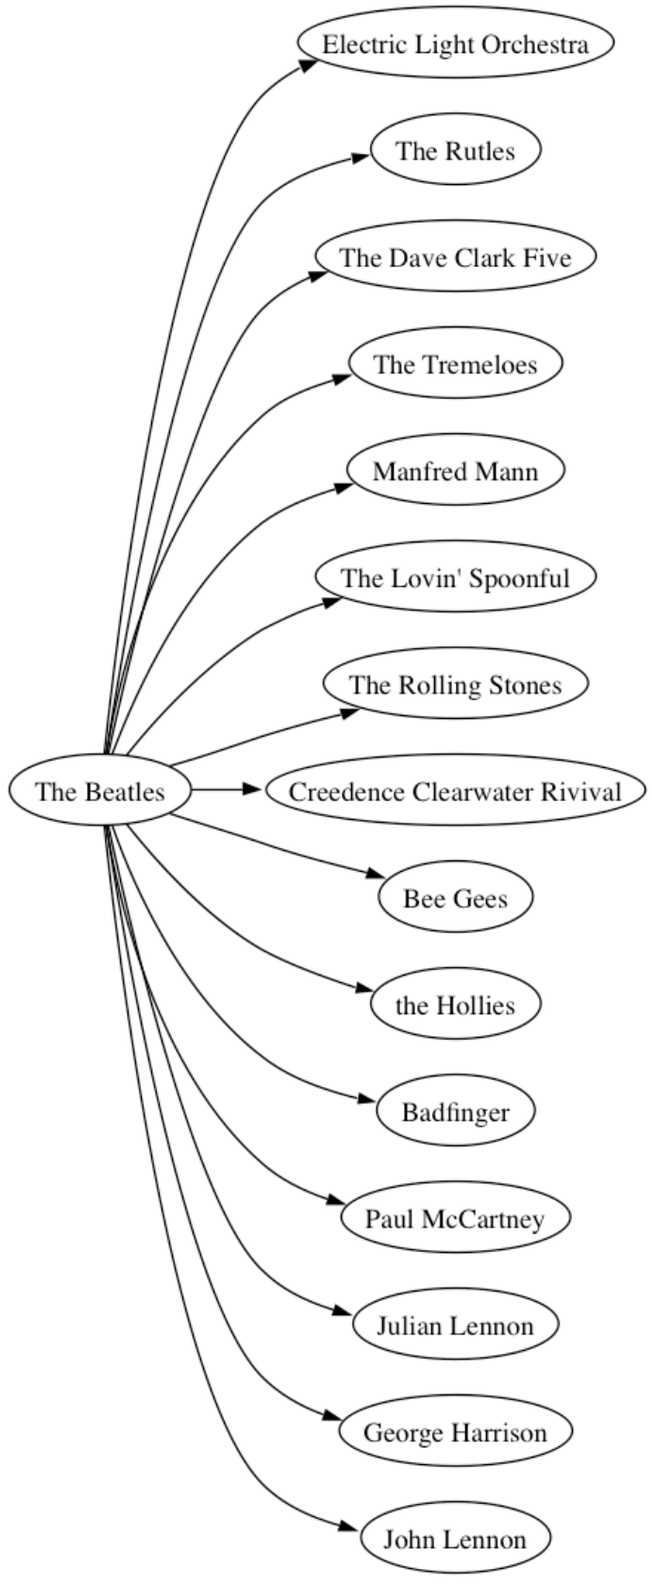

http://www.graphviz.org/

Artist map - level 1 out

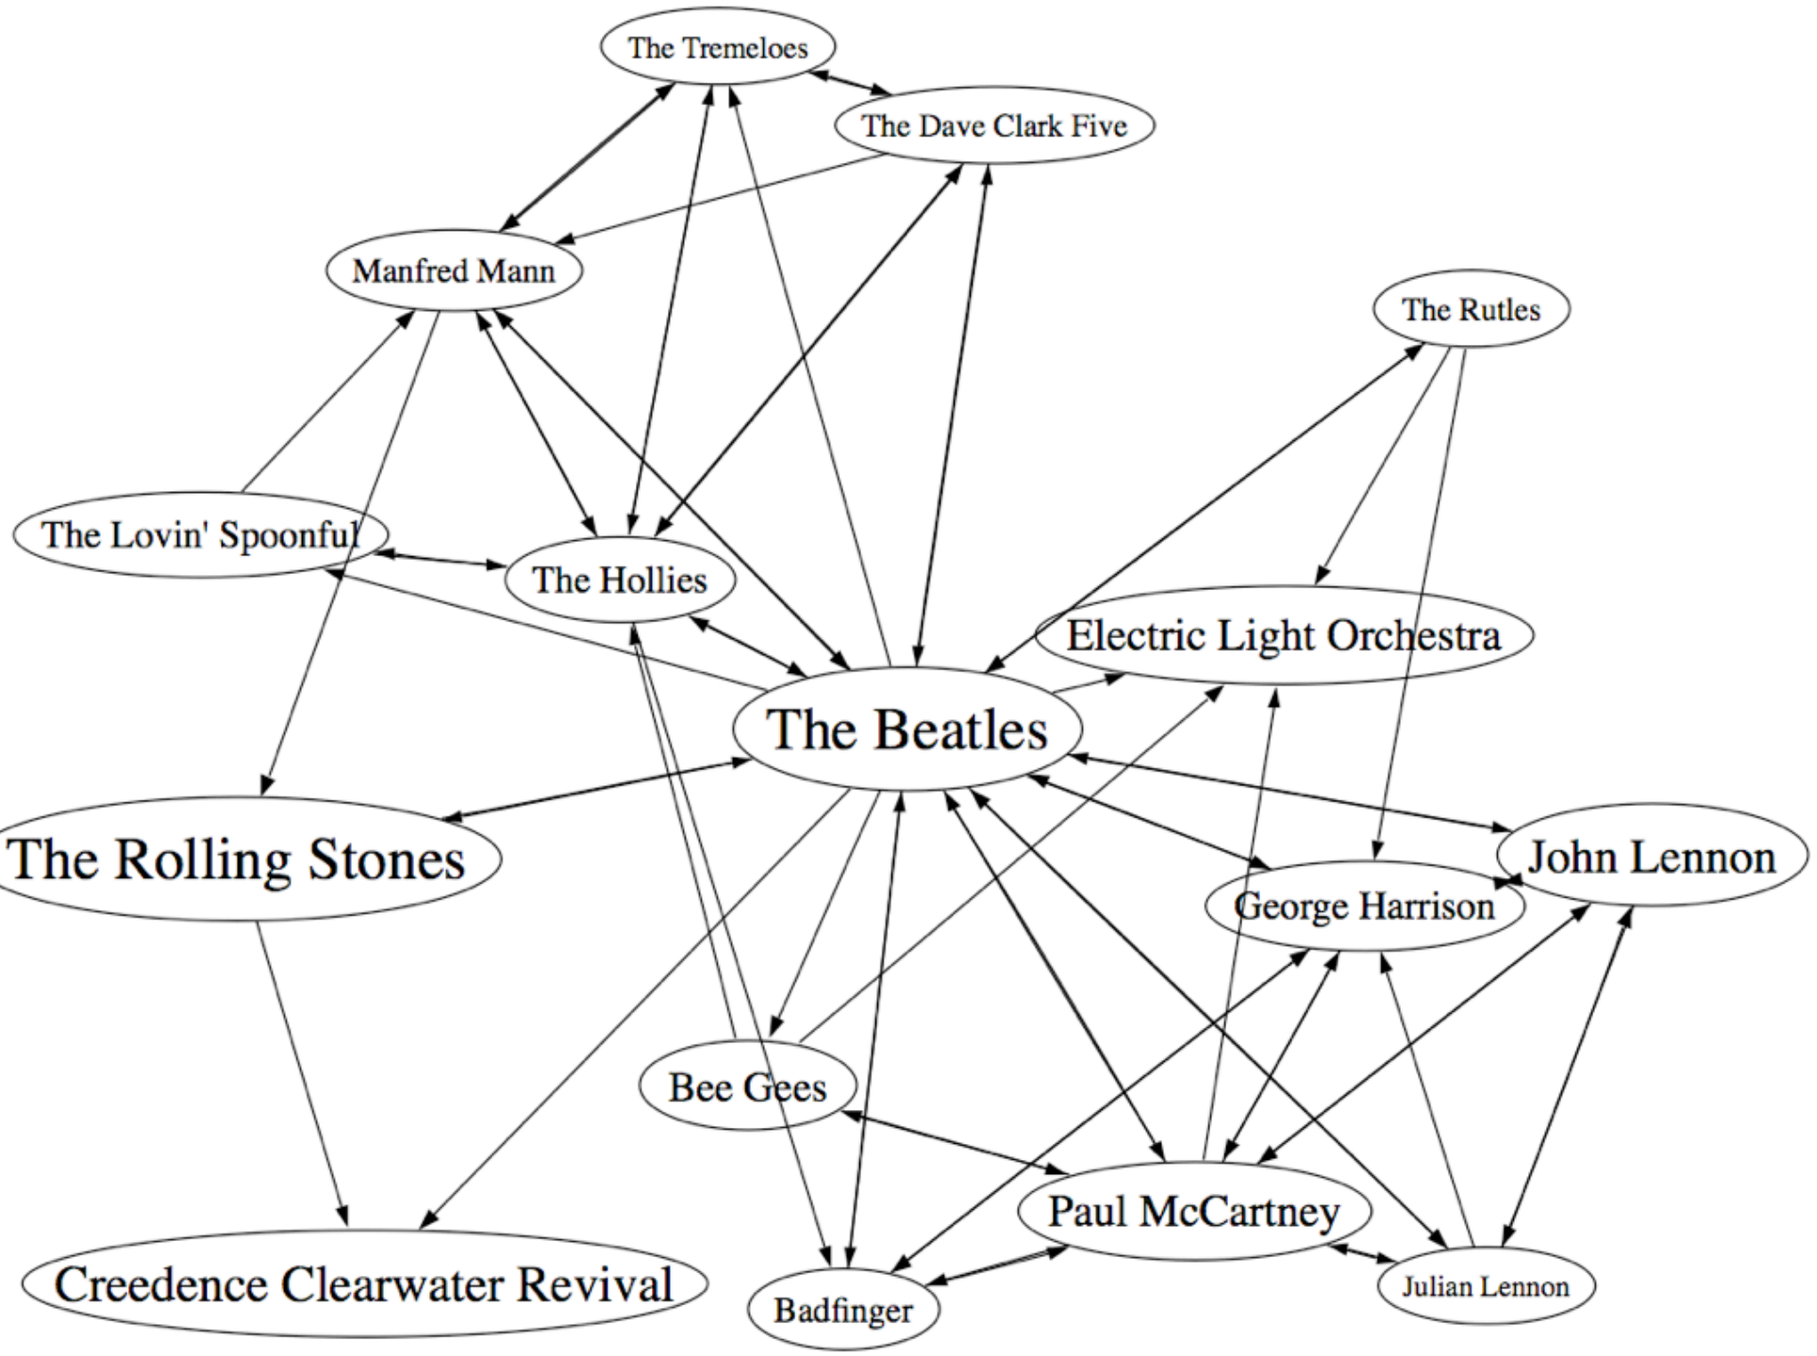

Already interconnections start to overwhelm

Artist map - level 1 in/out

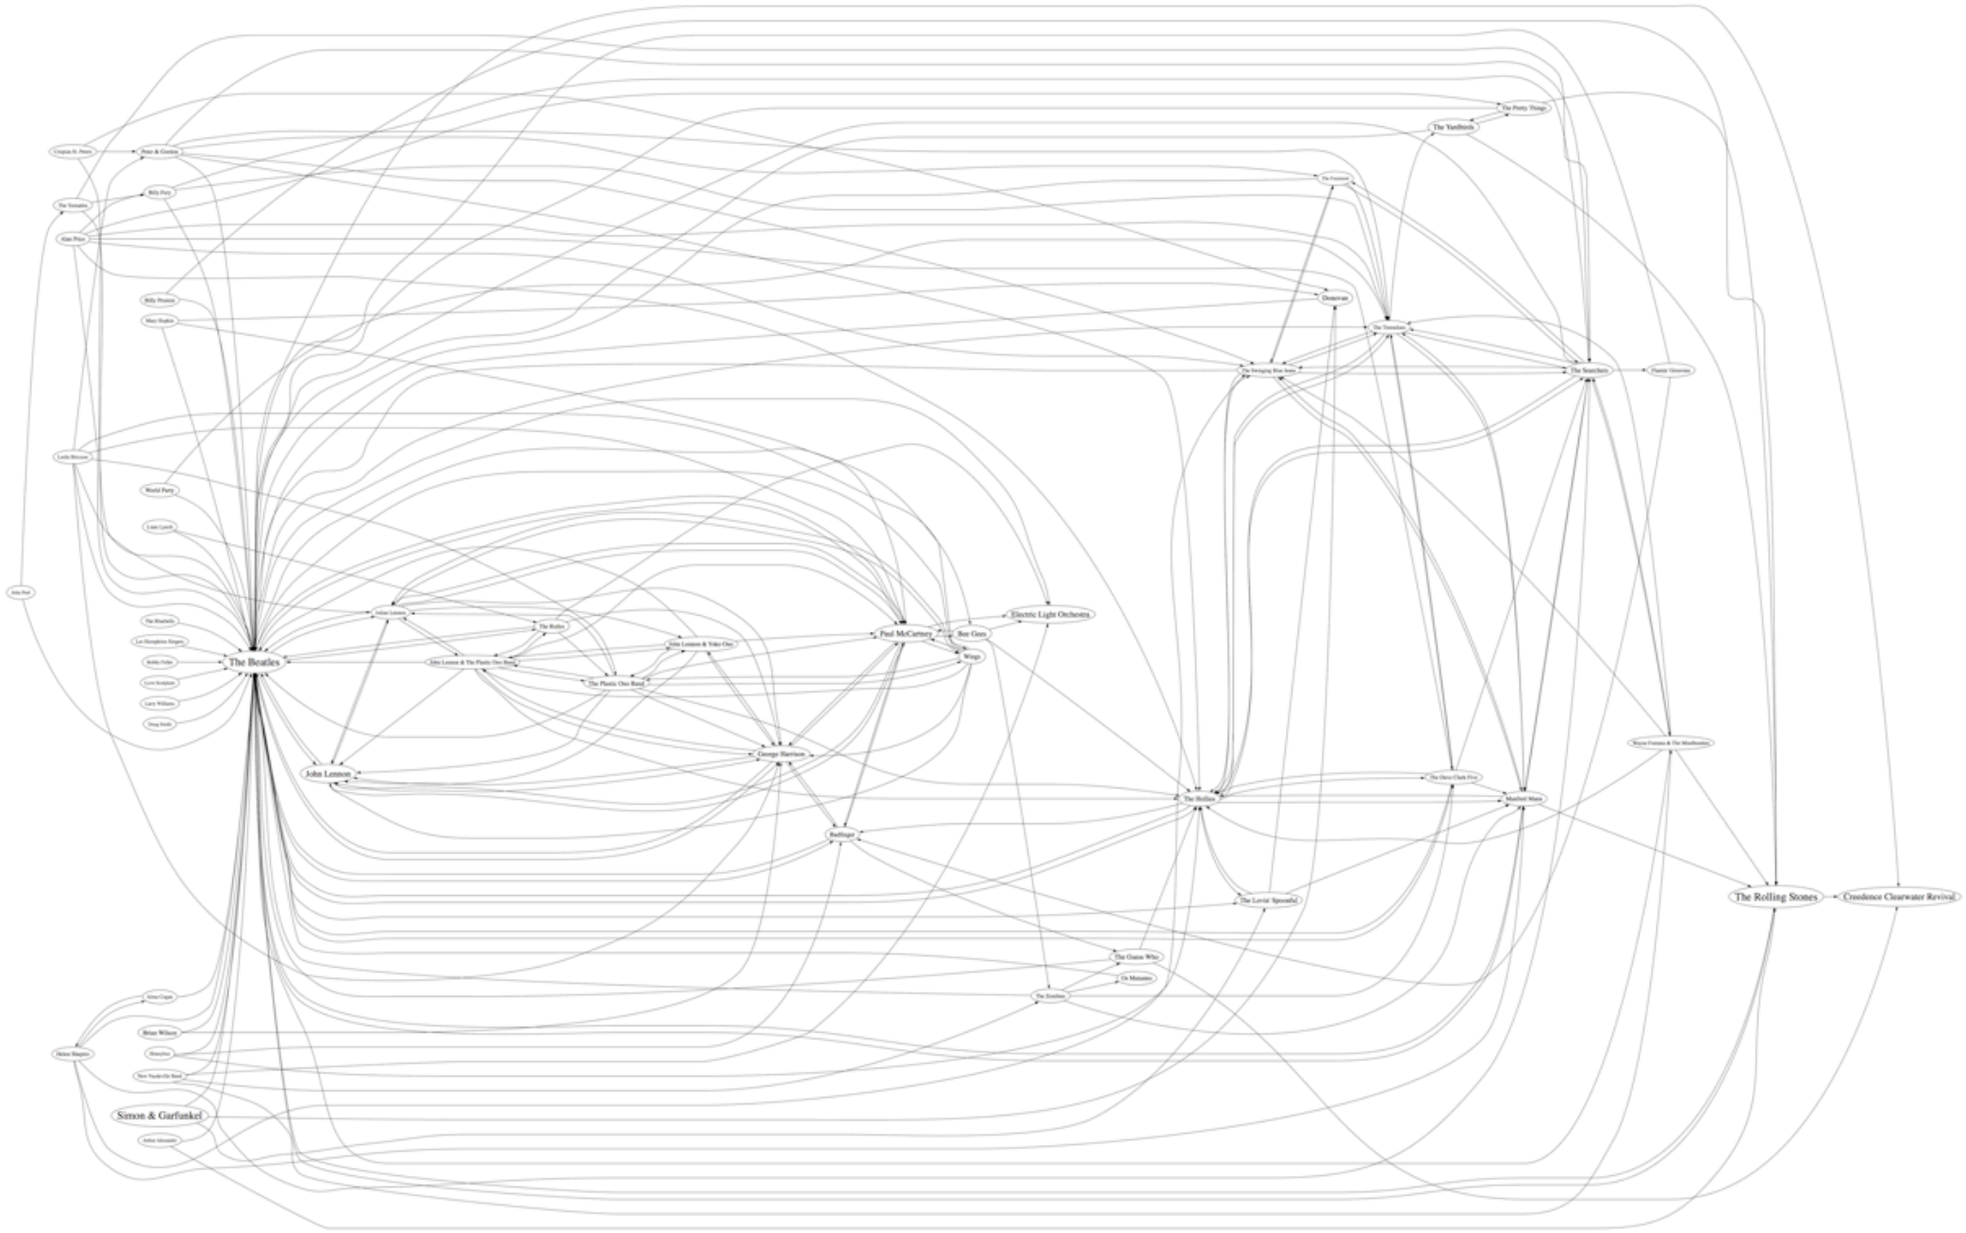

Artist map - level 1 in/out

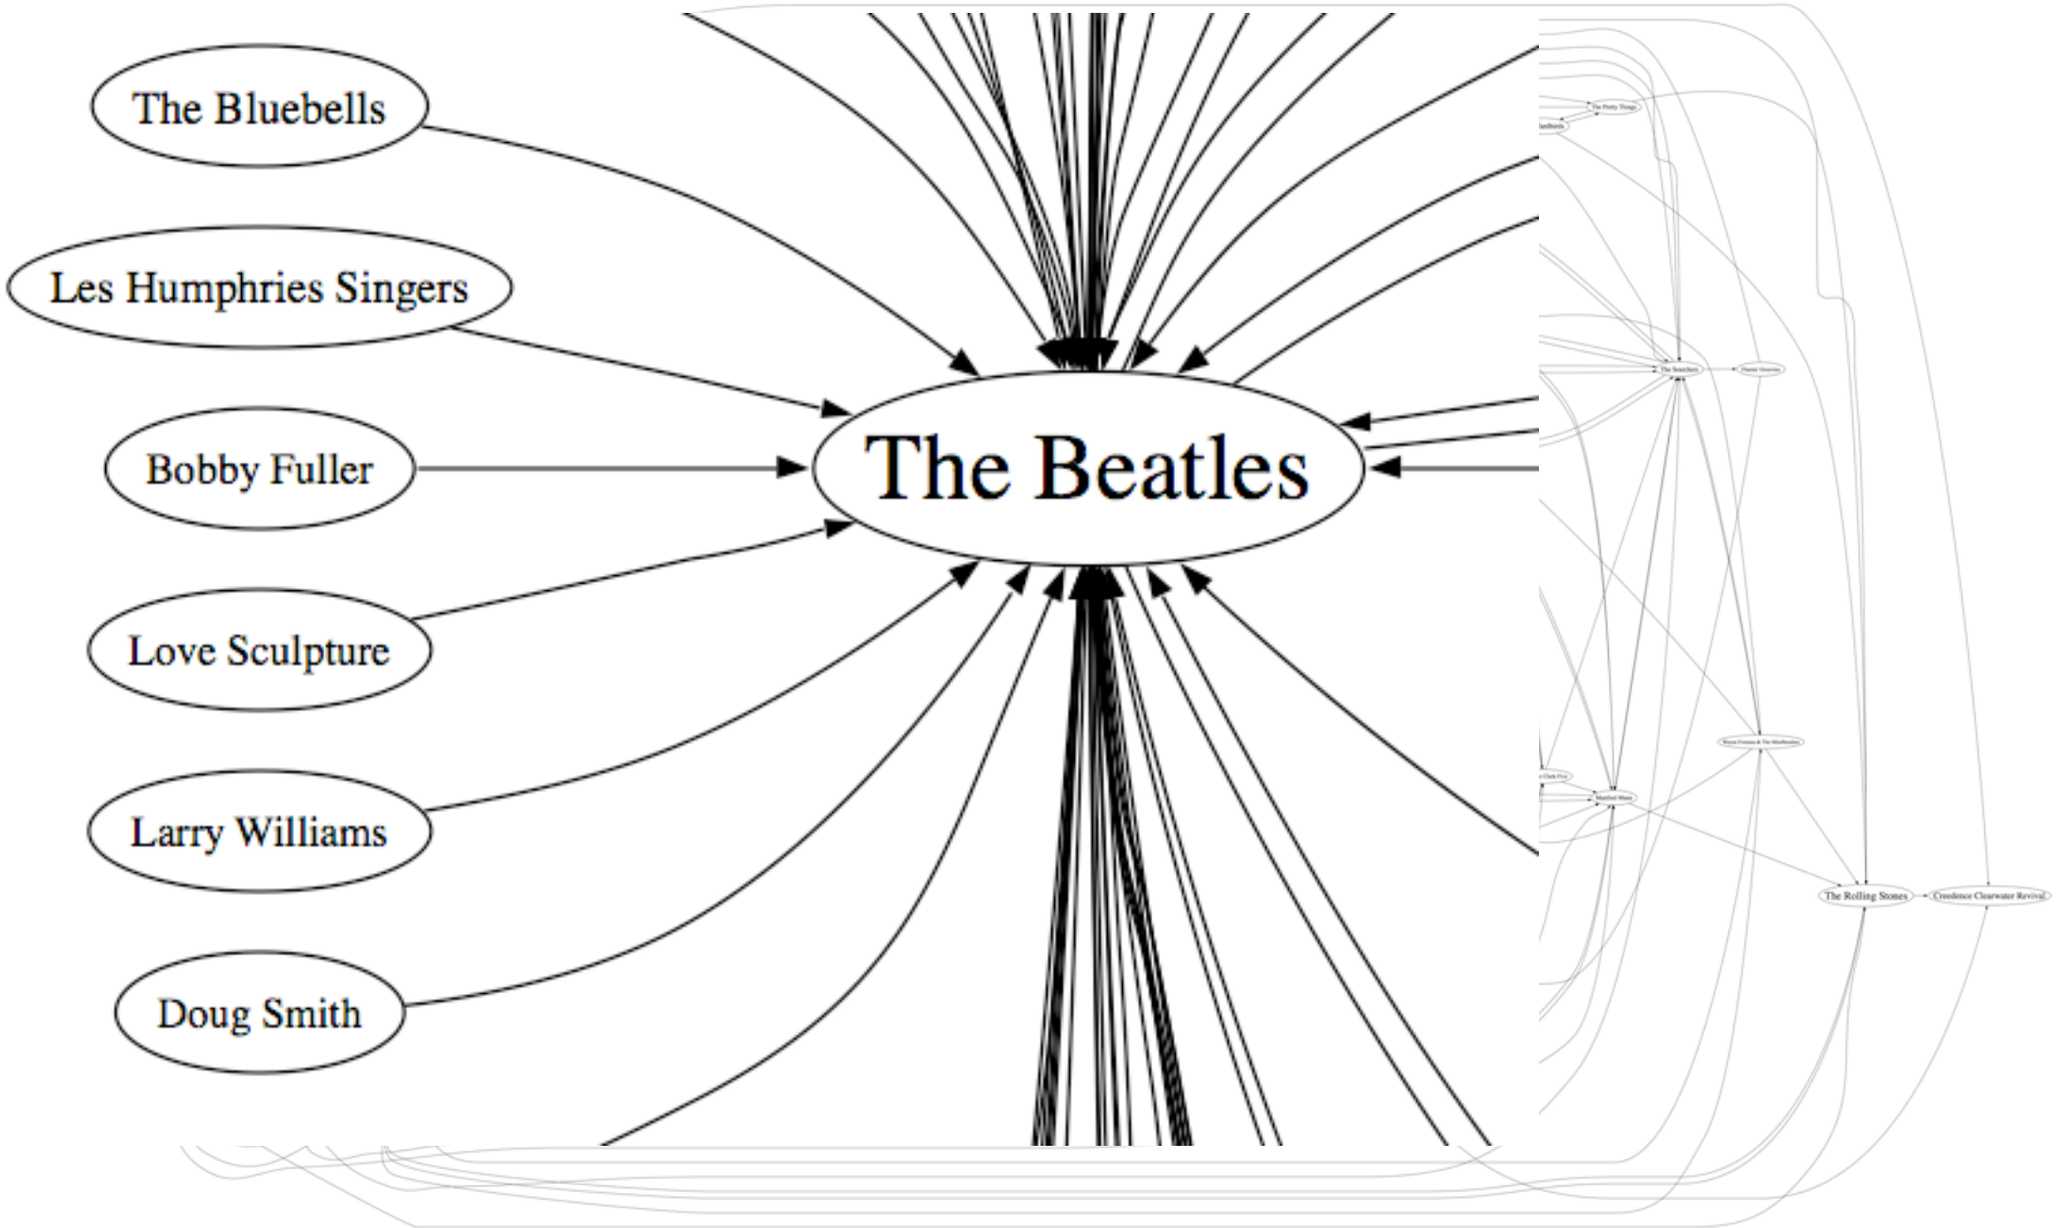

Artist map - level 1 in/out

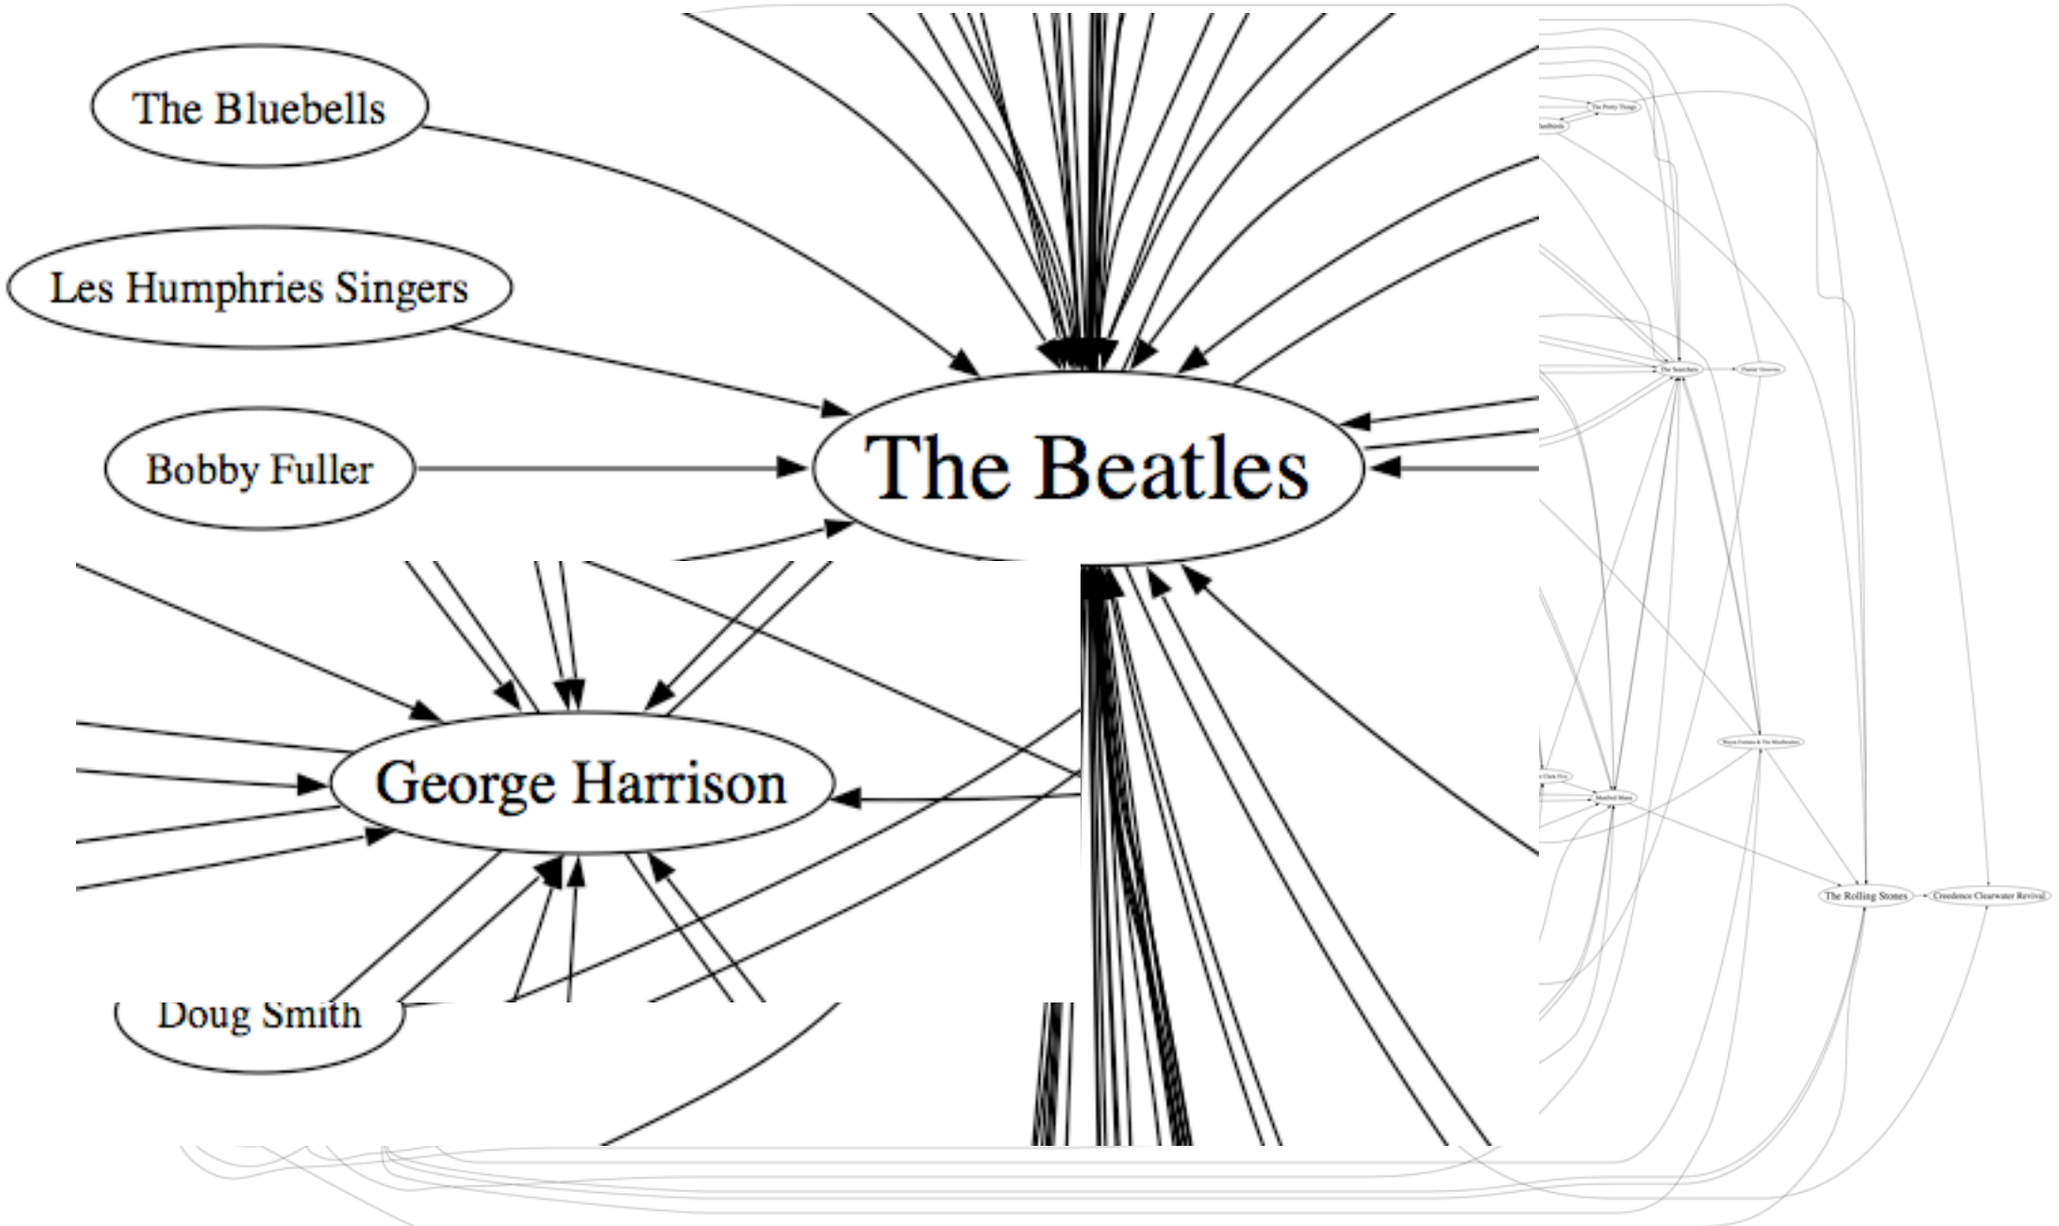

Artist map - level 1 in/out

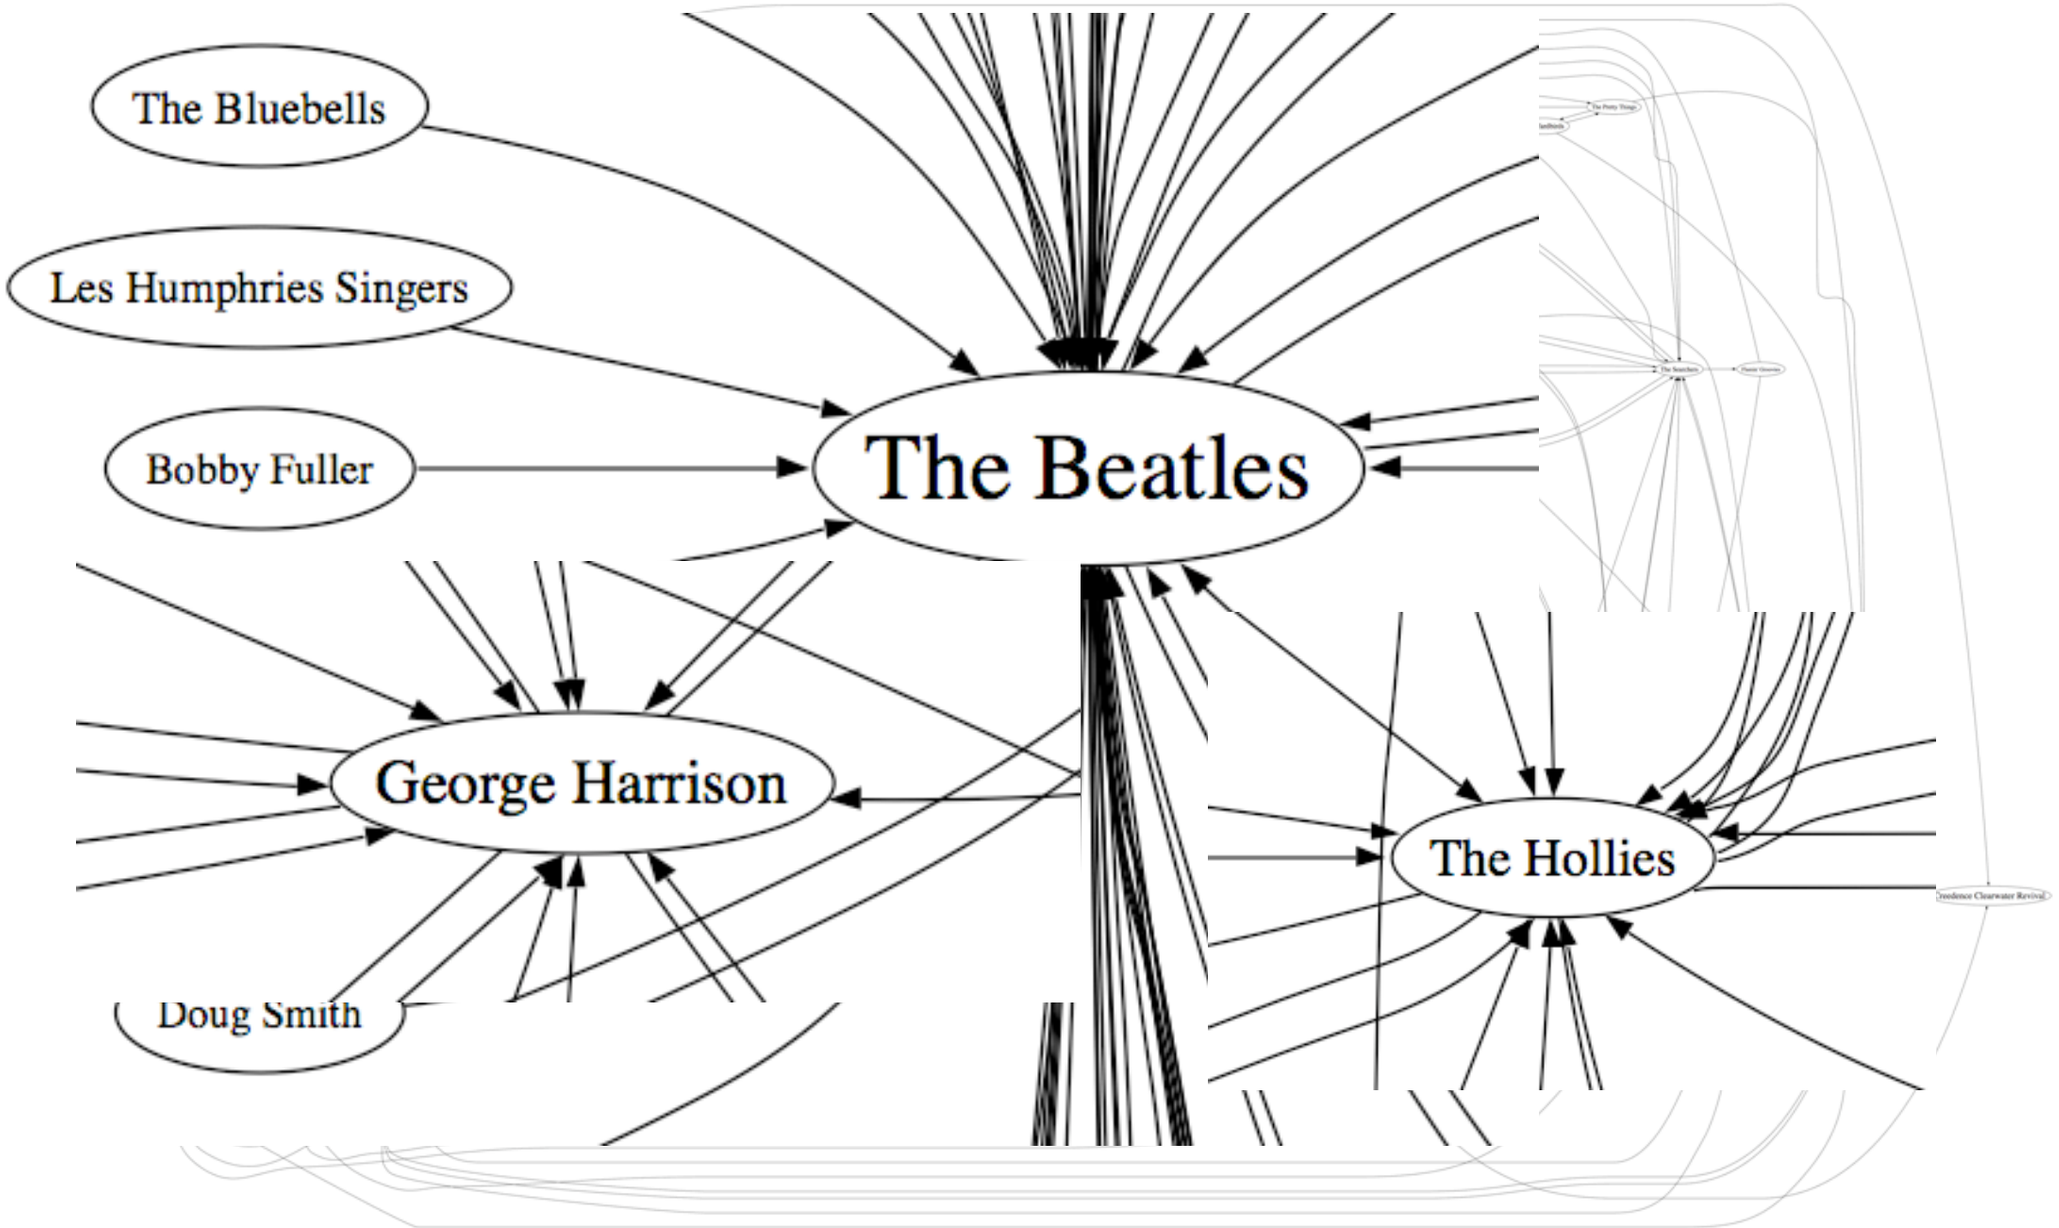

Artist map - level 1 in/out

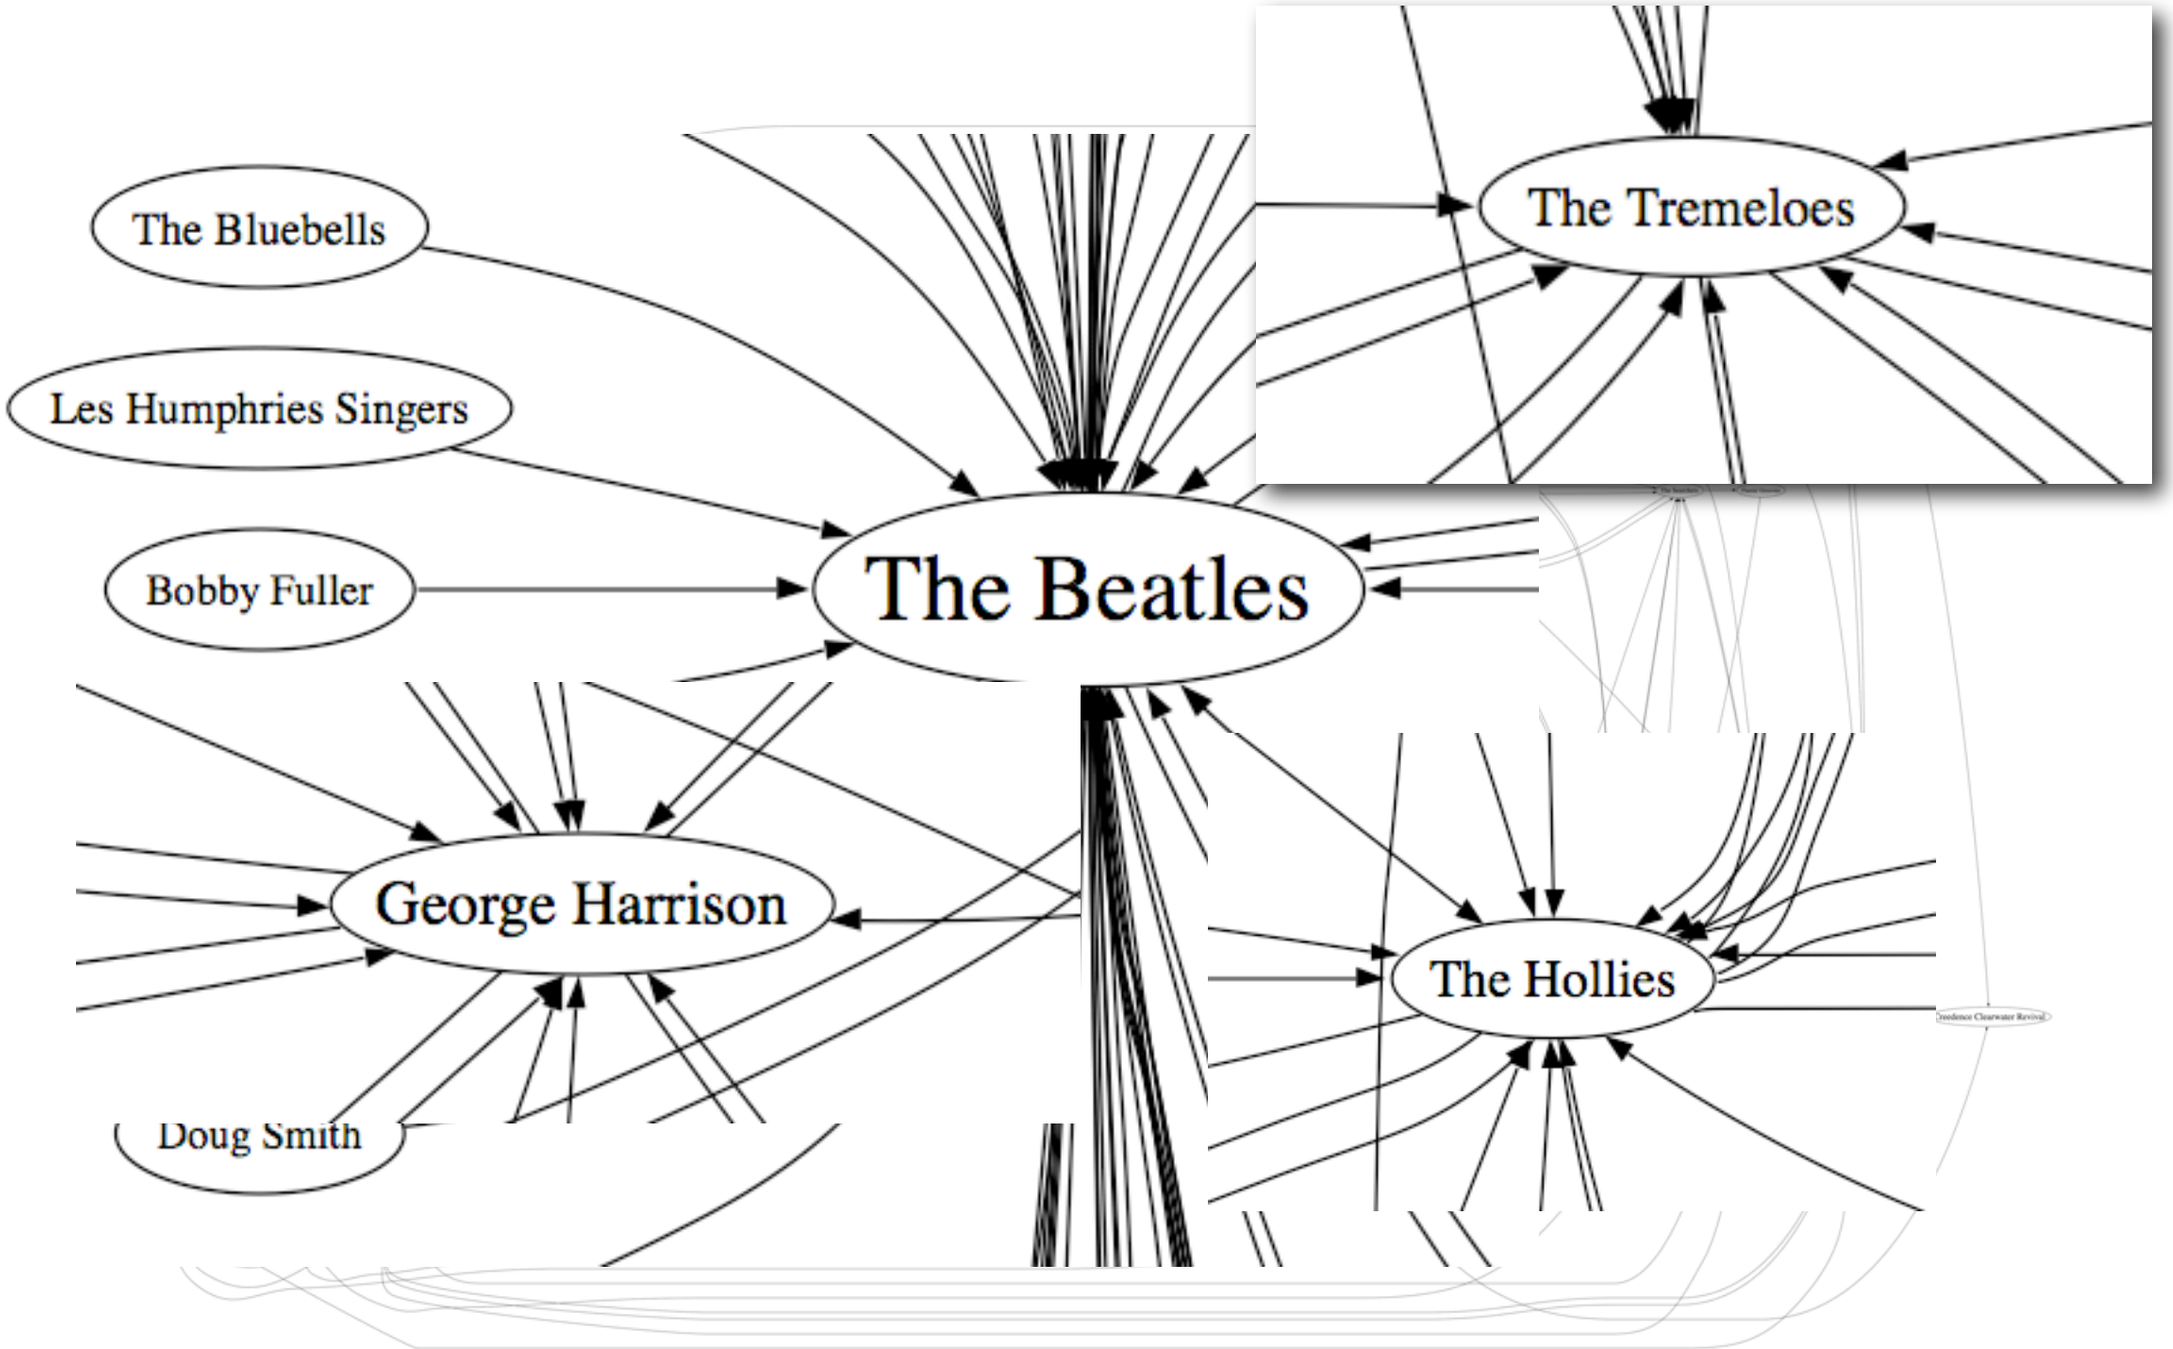

Artist map - level 1 in/out

Artist map - level 1 in/out

Artist map - level 2

116 artists and 665 edges. We are totally overwhelmed.

Artist map - level 2

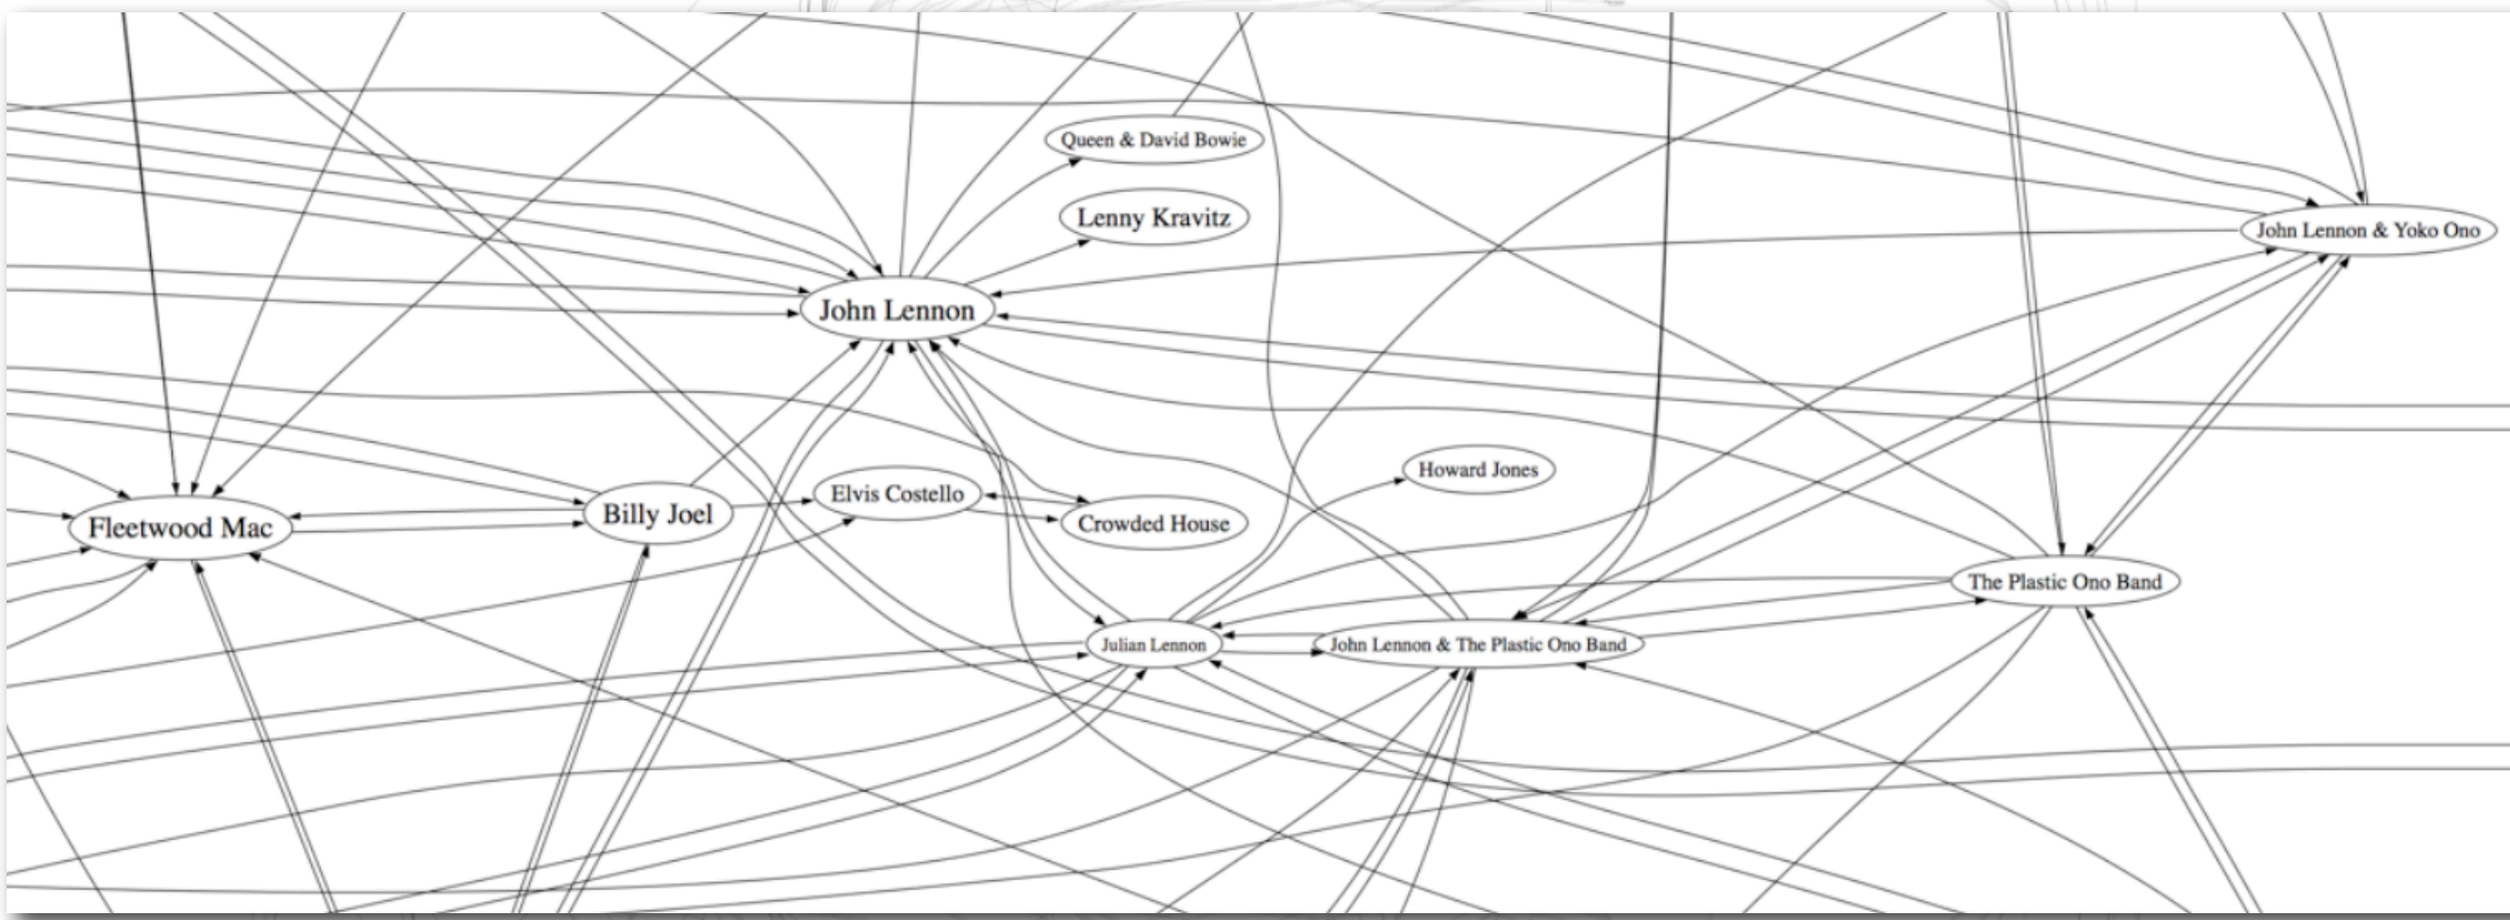

116 artists and 665 edges. We are totally overwhelmed.

Simplifying the graph

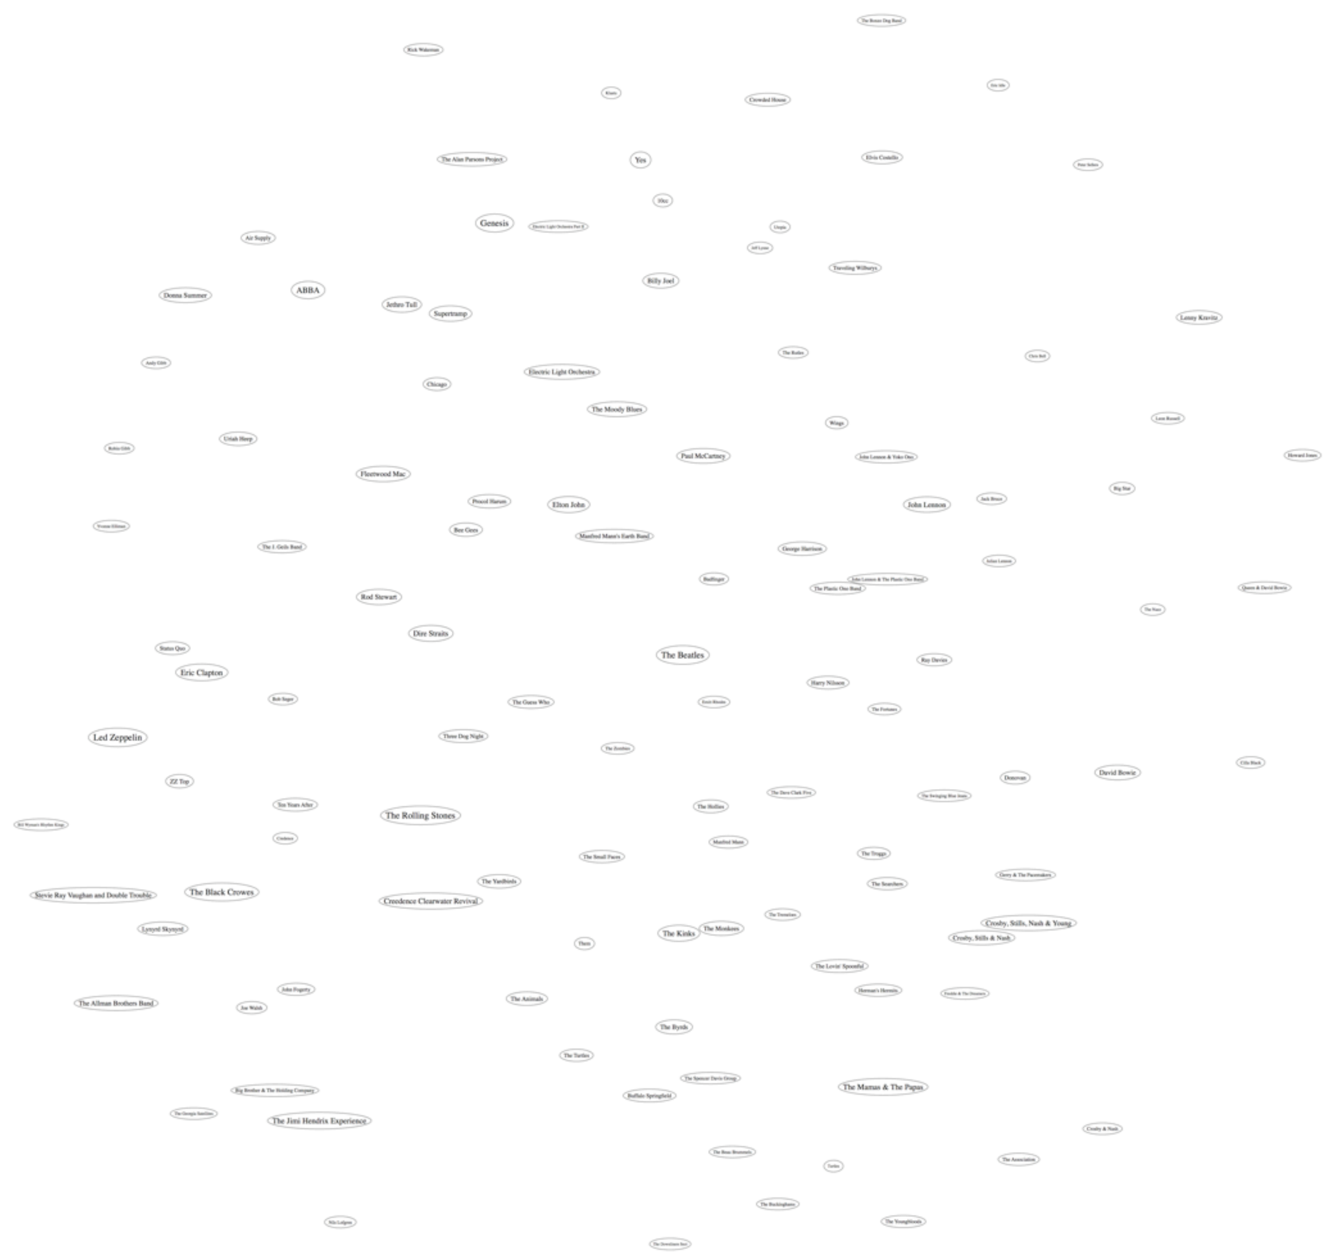

Eliminate the edges and use spring force embedding to position the artists

Simplifying the graph

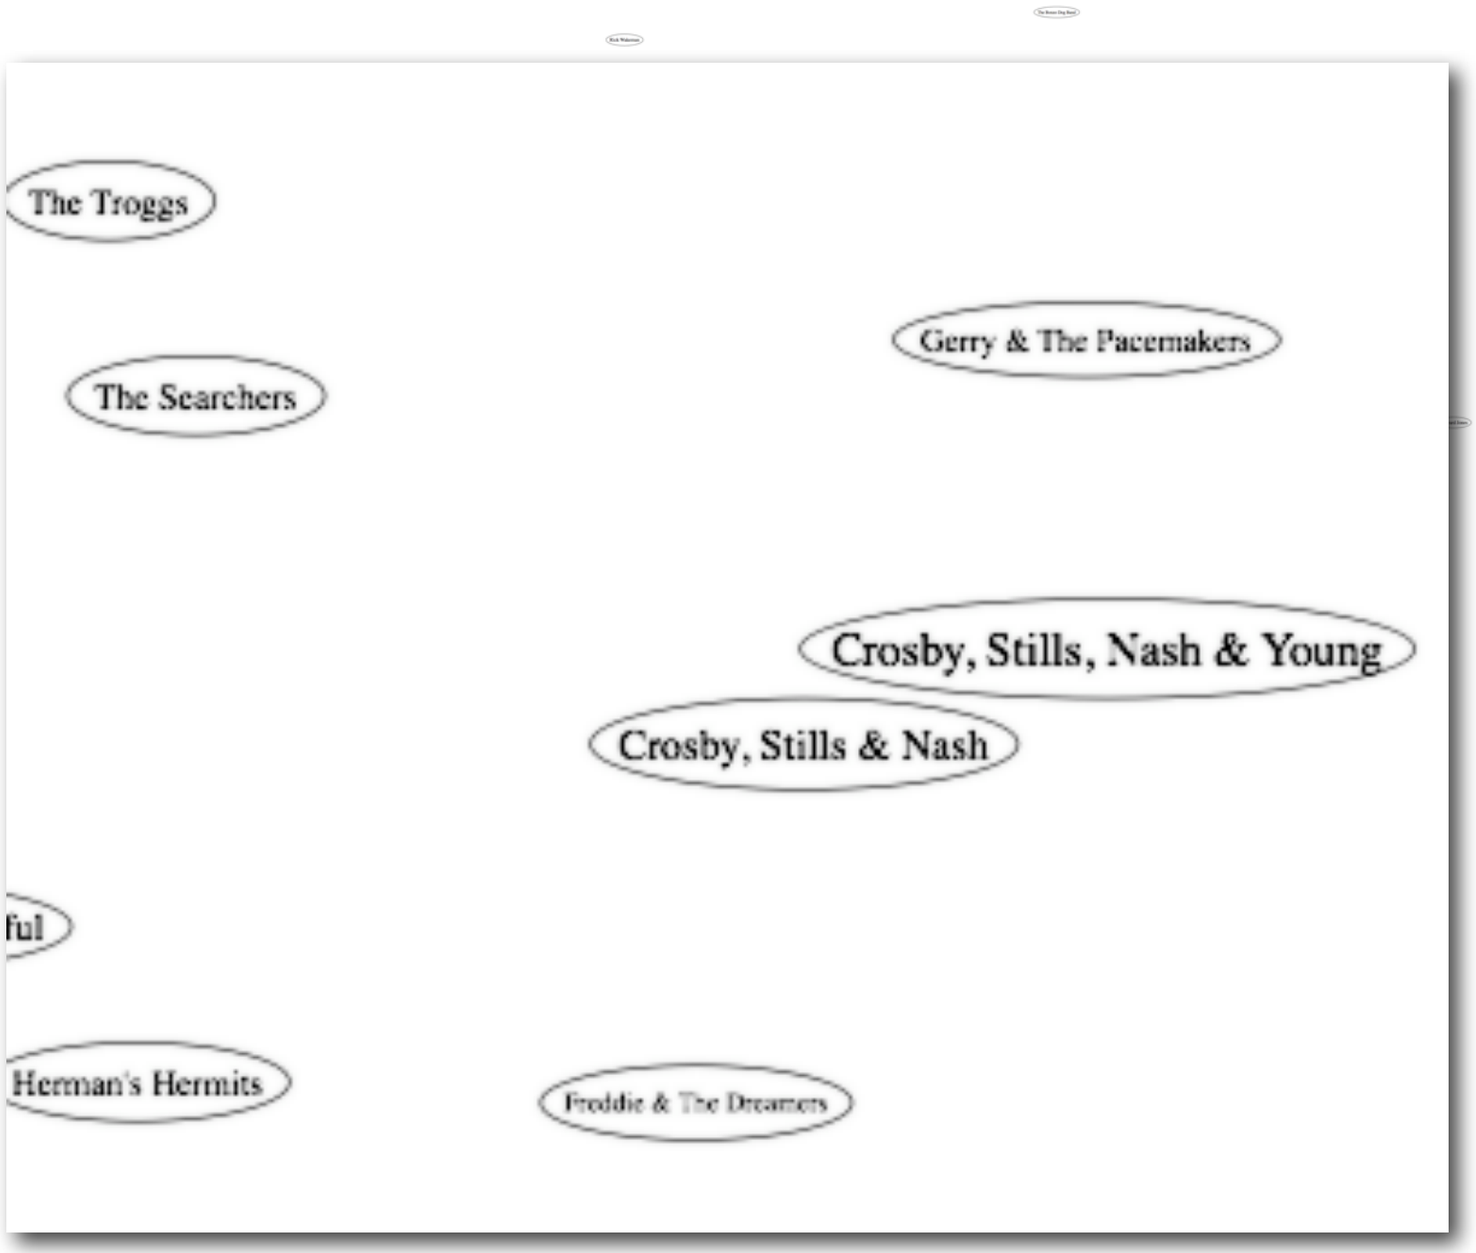

Eliminate the edges and use spring force embedding to position the artists

Reducing the edges

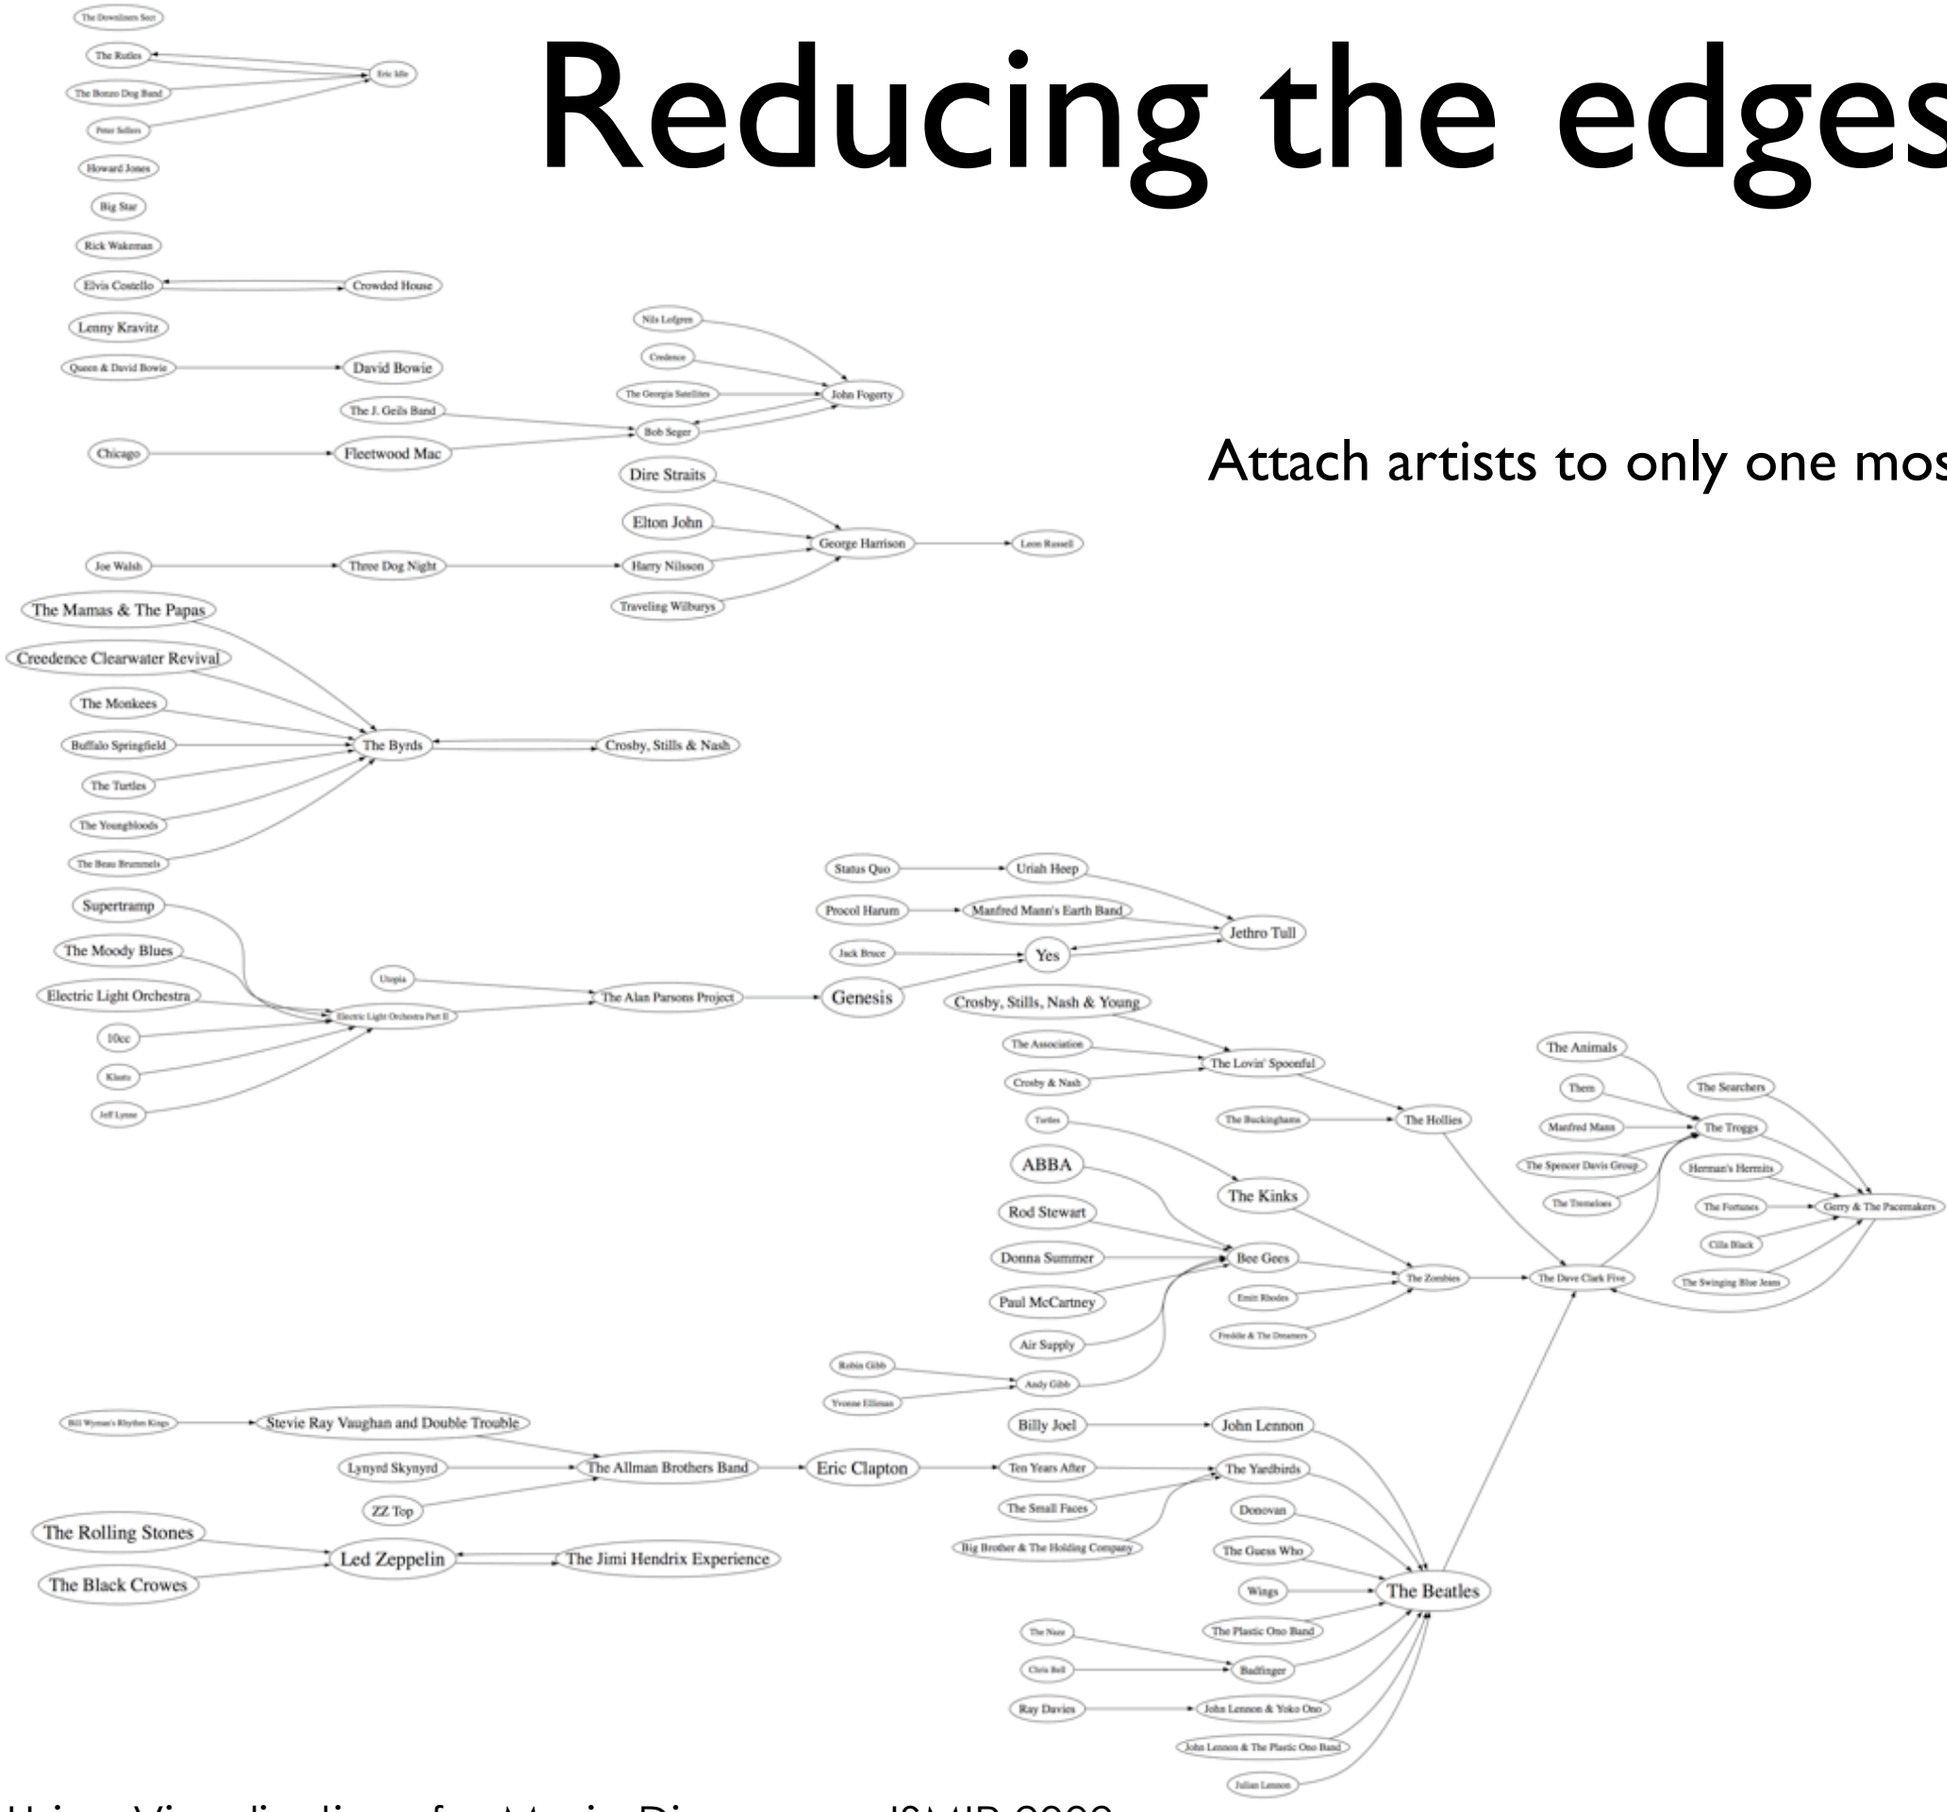

Attach artists to only one most similar artist

Reducing the edges

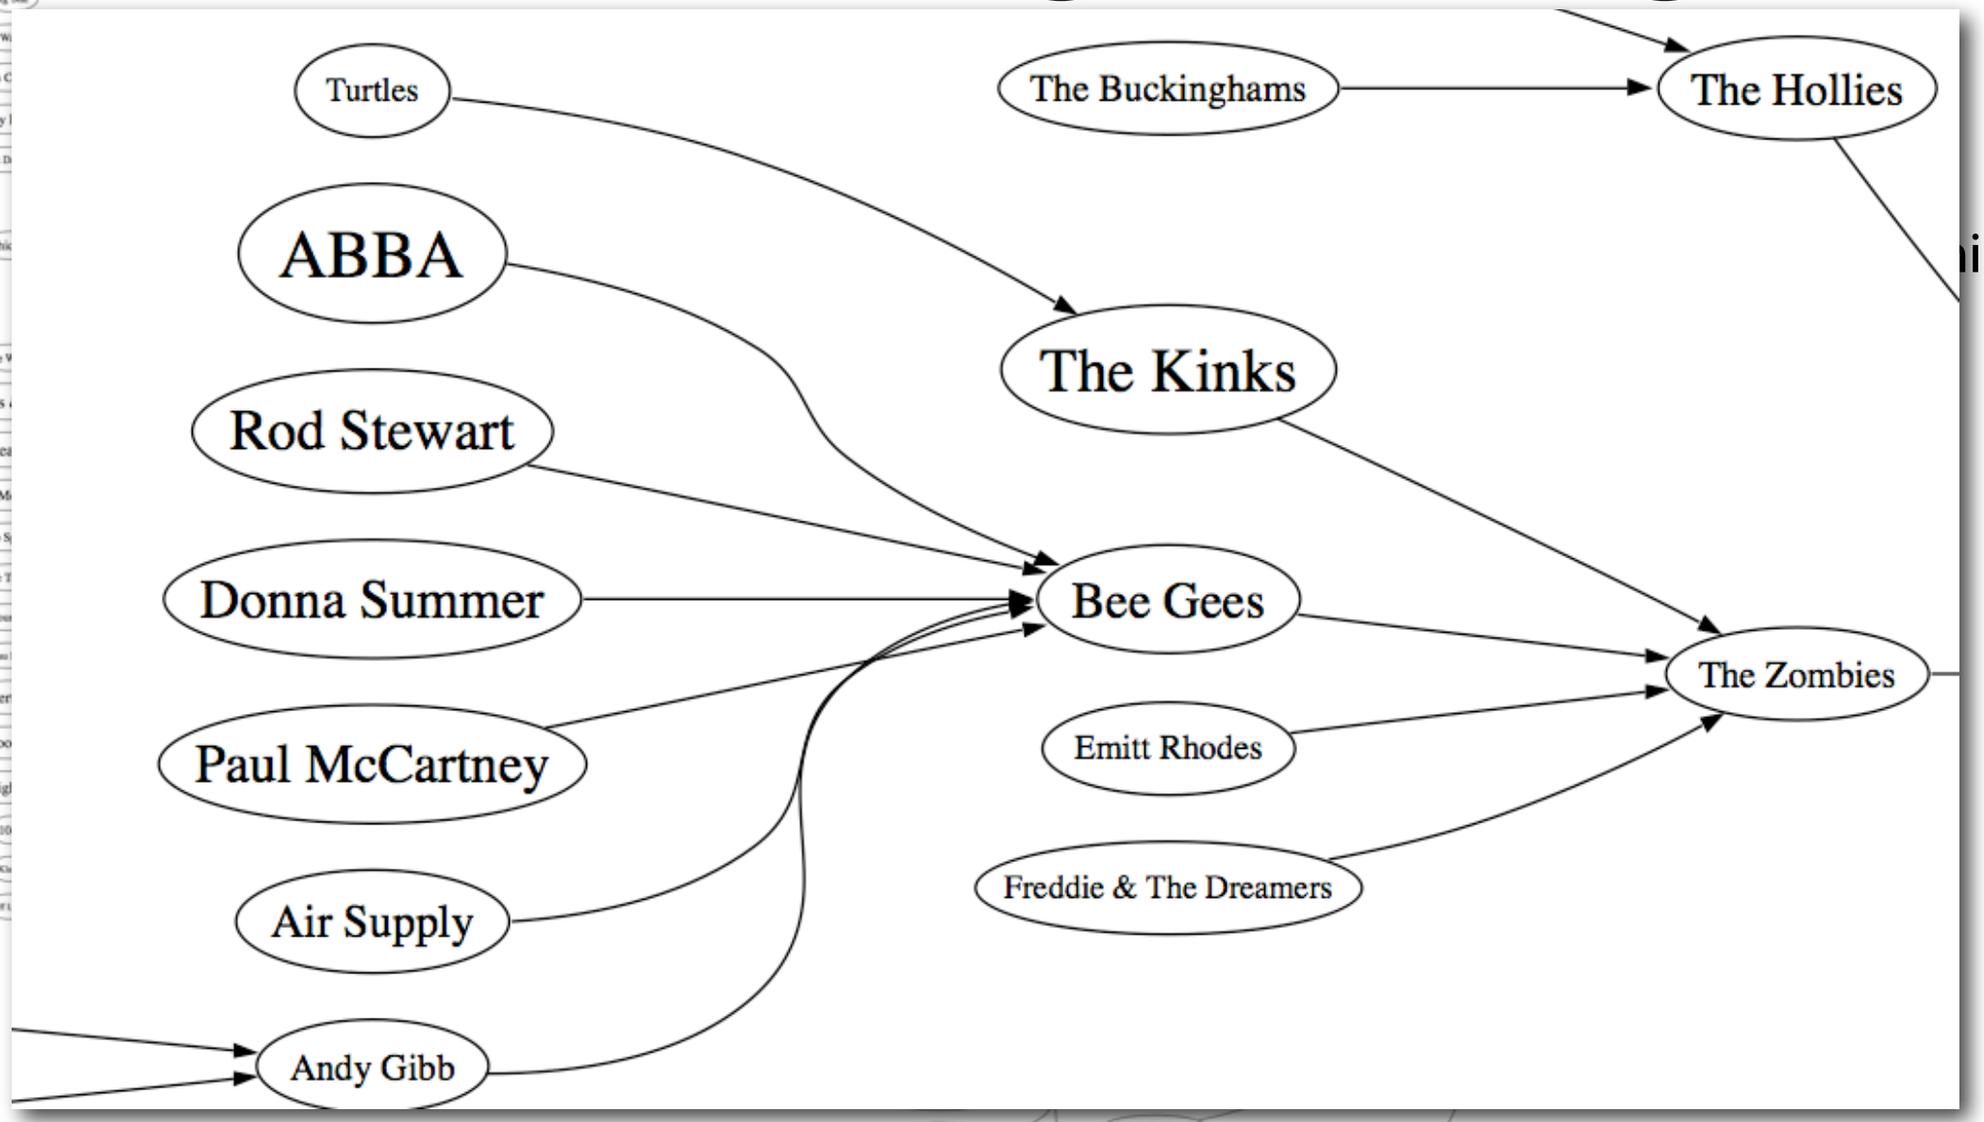

ilar artist

Reducing the edges

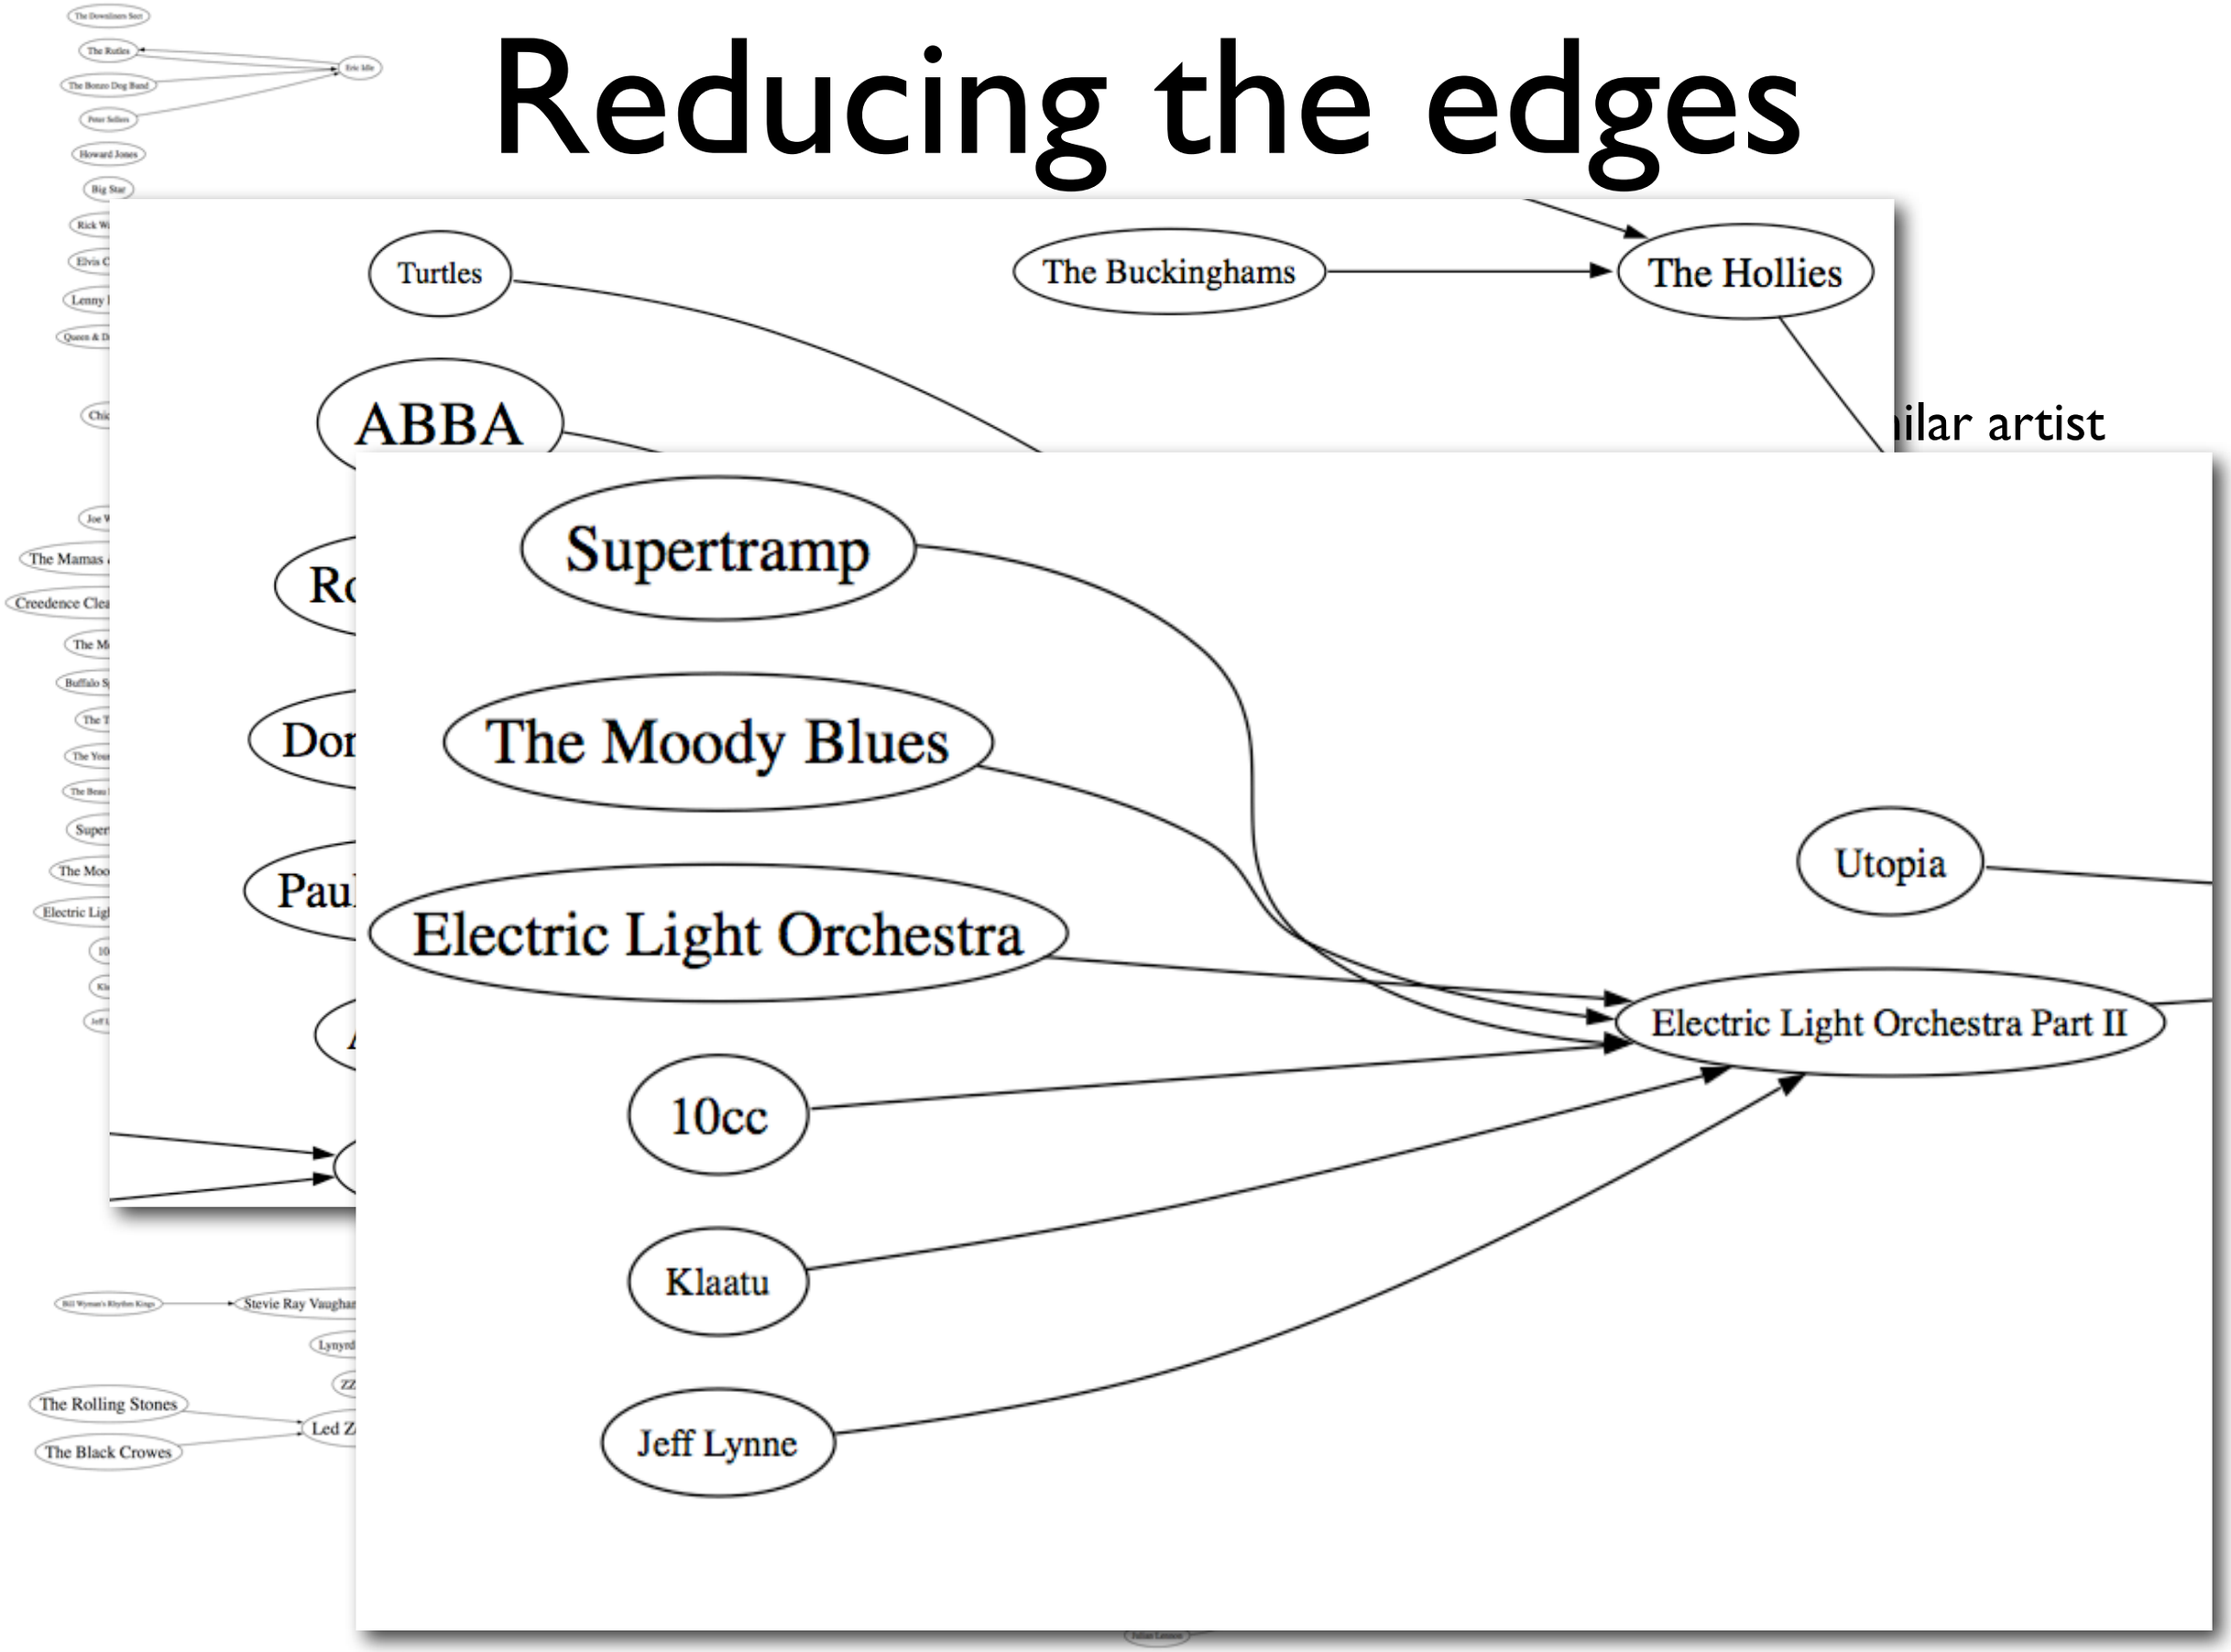

Reducing the edges

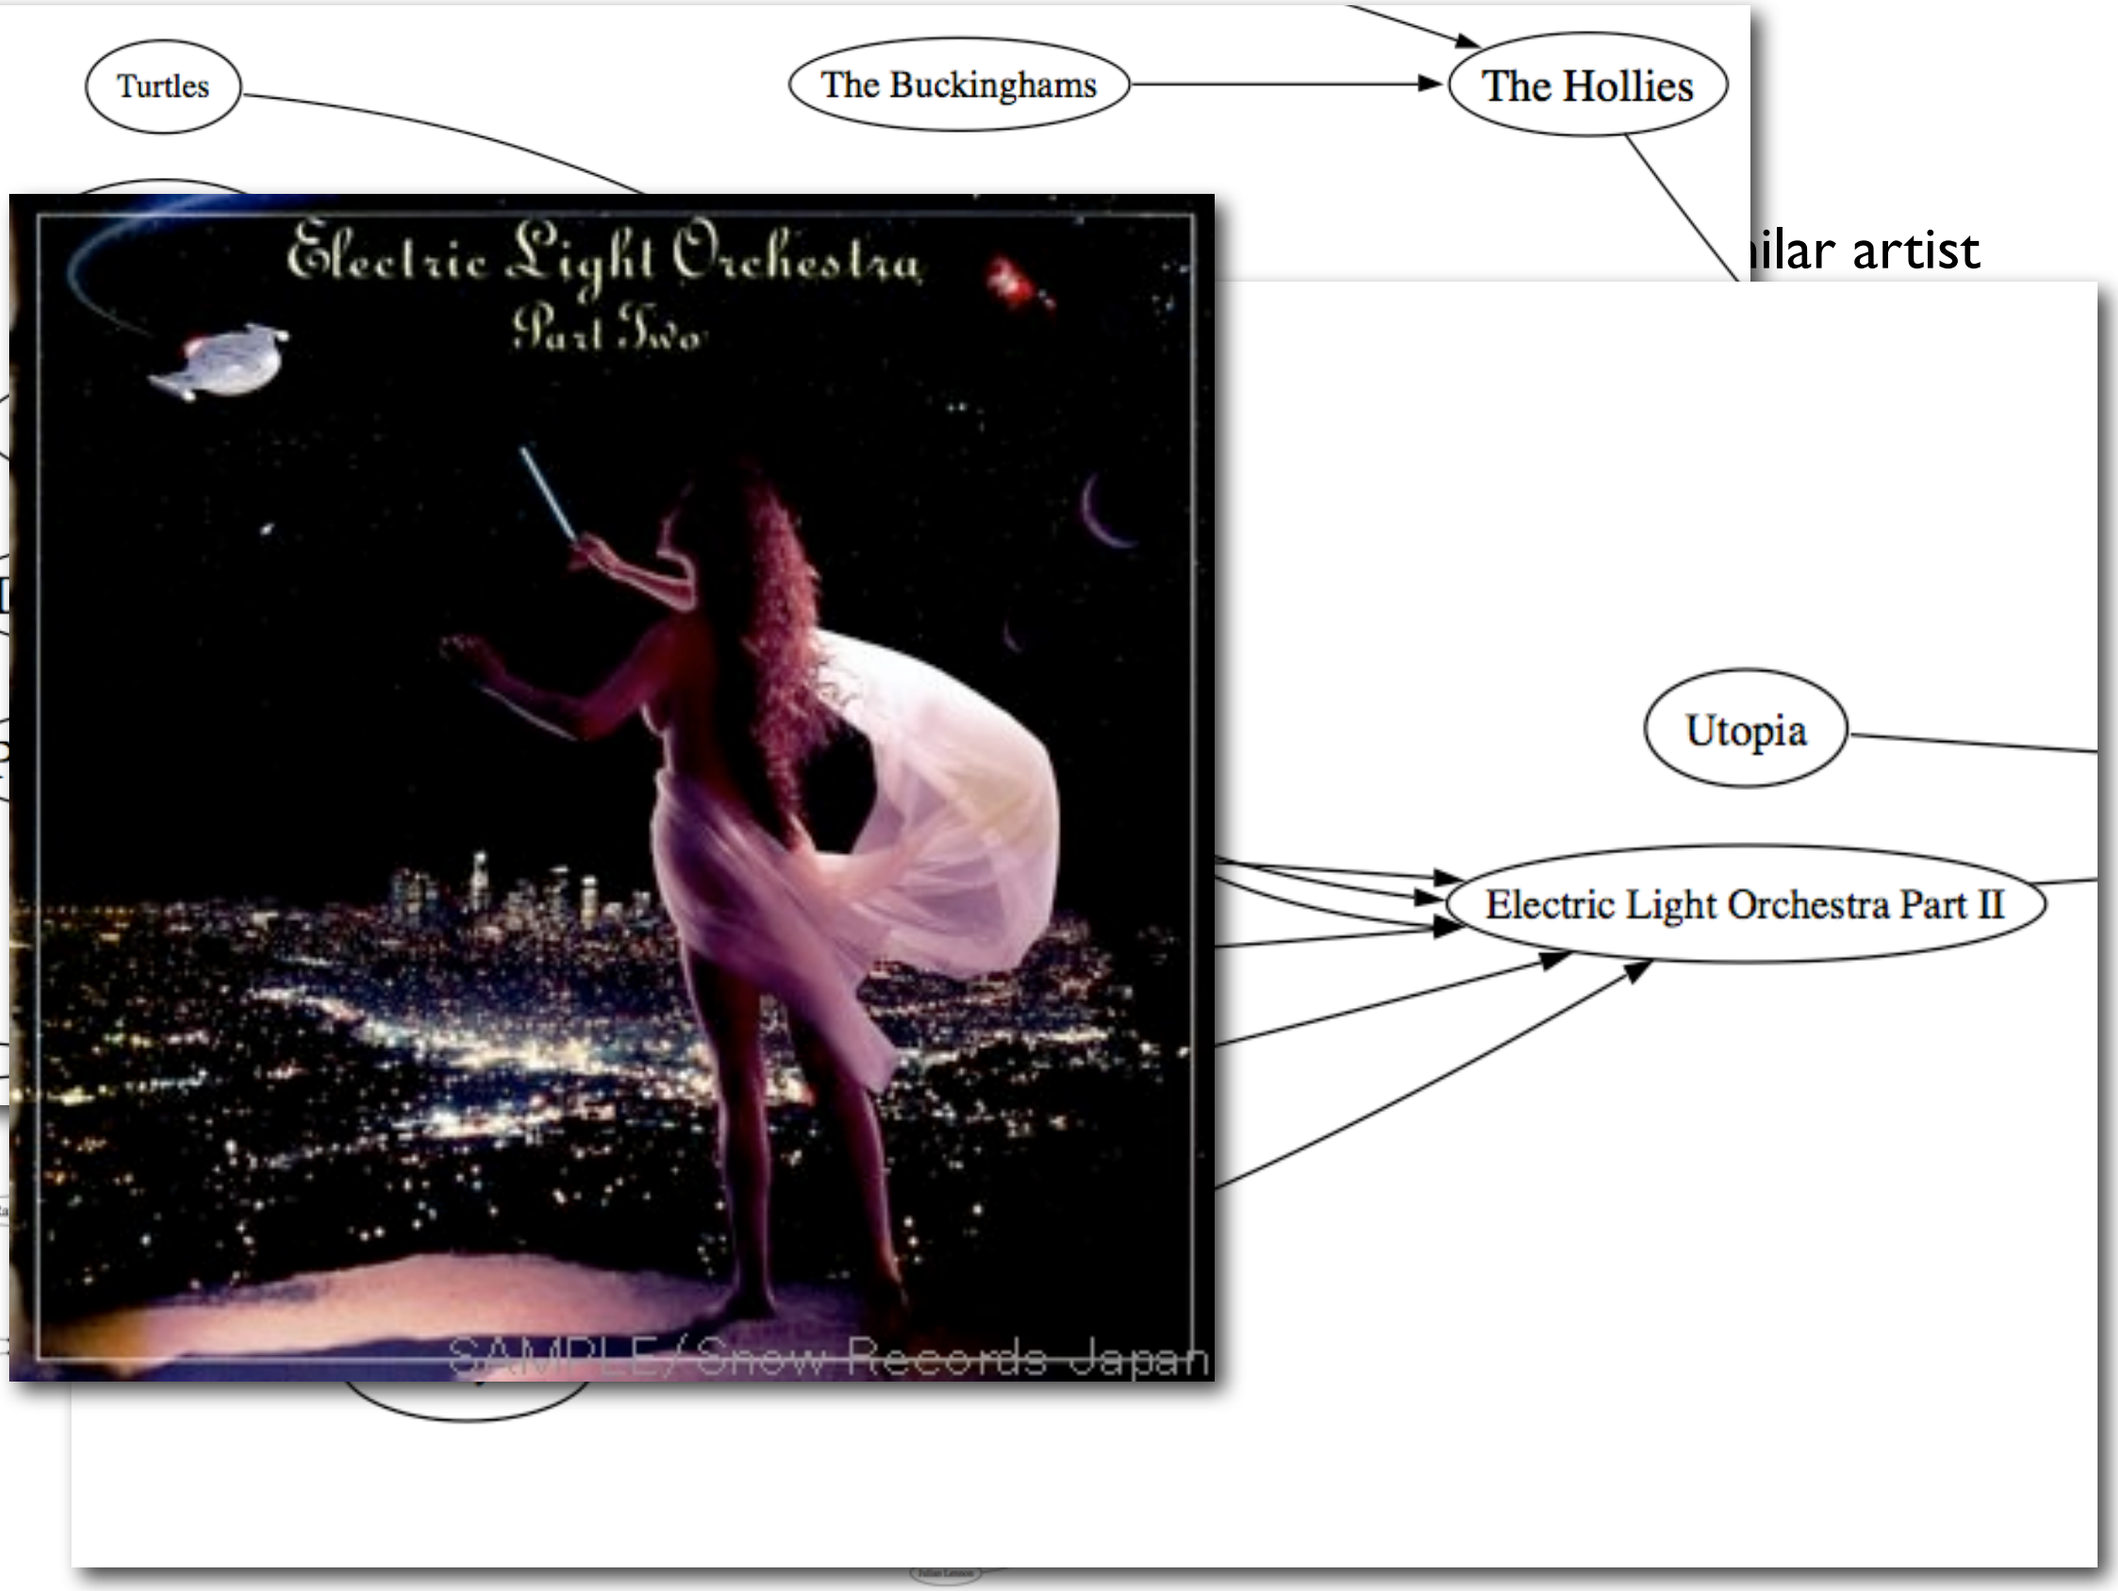

Putting ELO II in its place

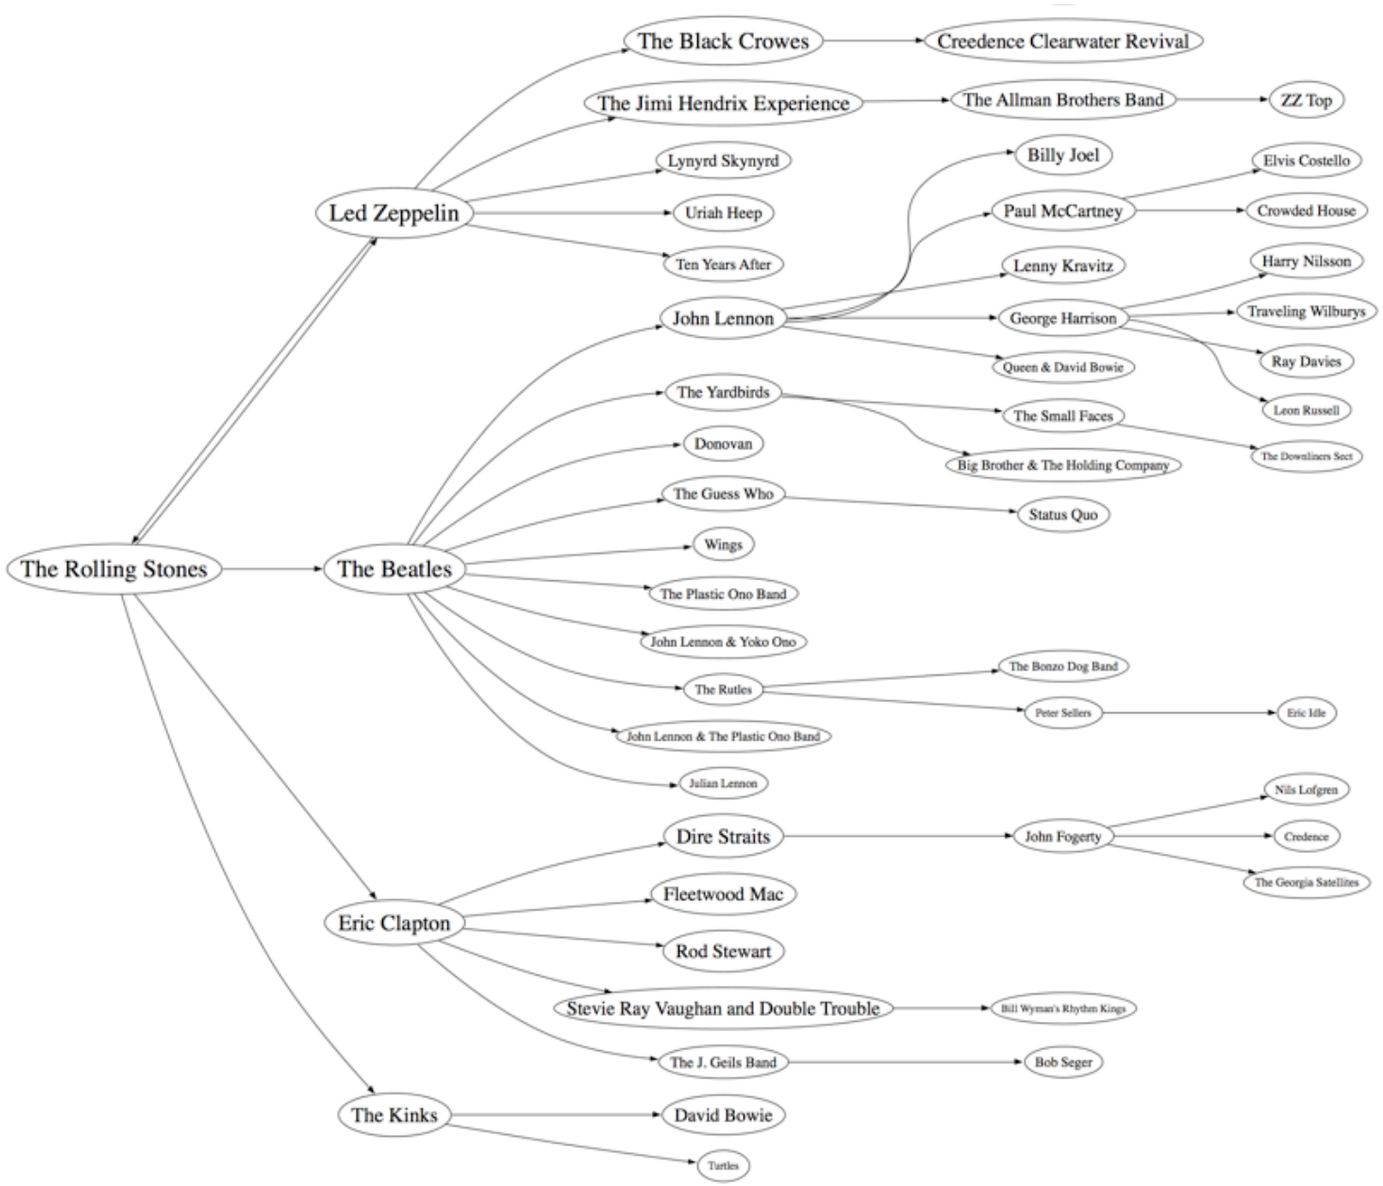

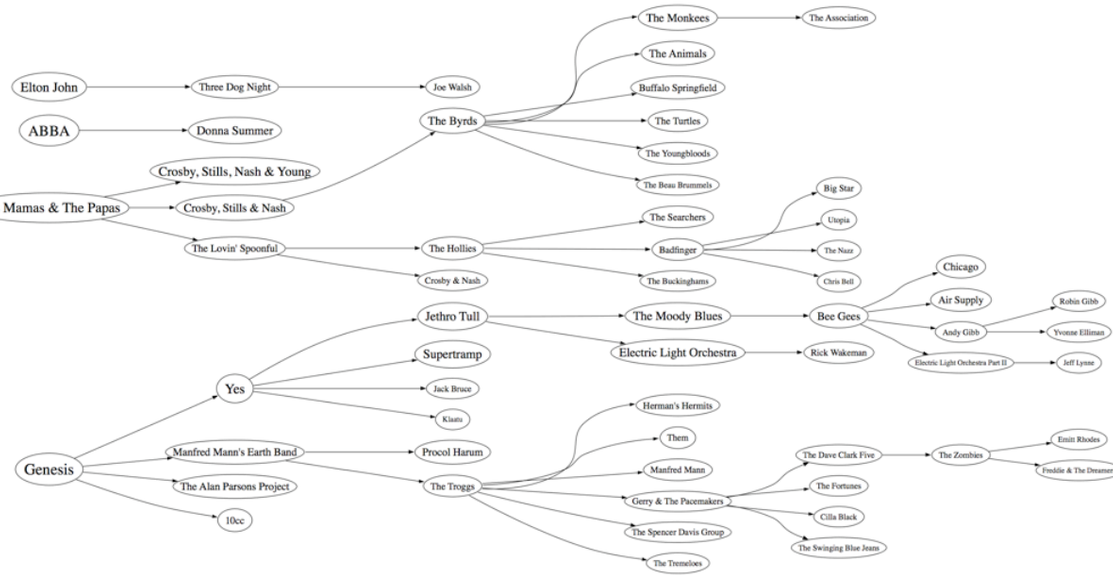

Attach artists to the most similar artist that has greater familiarity

Putting ELO II in its place

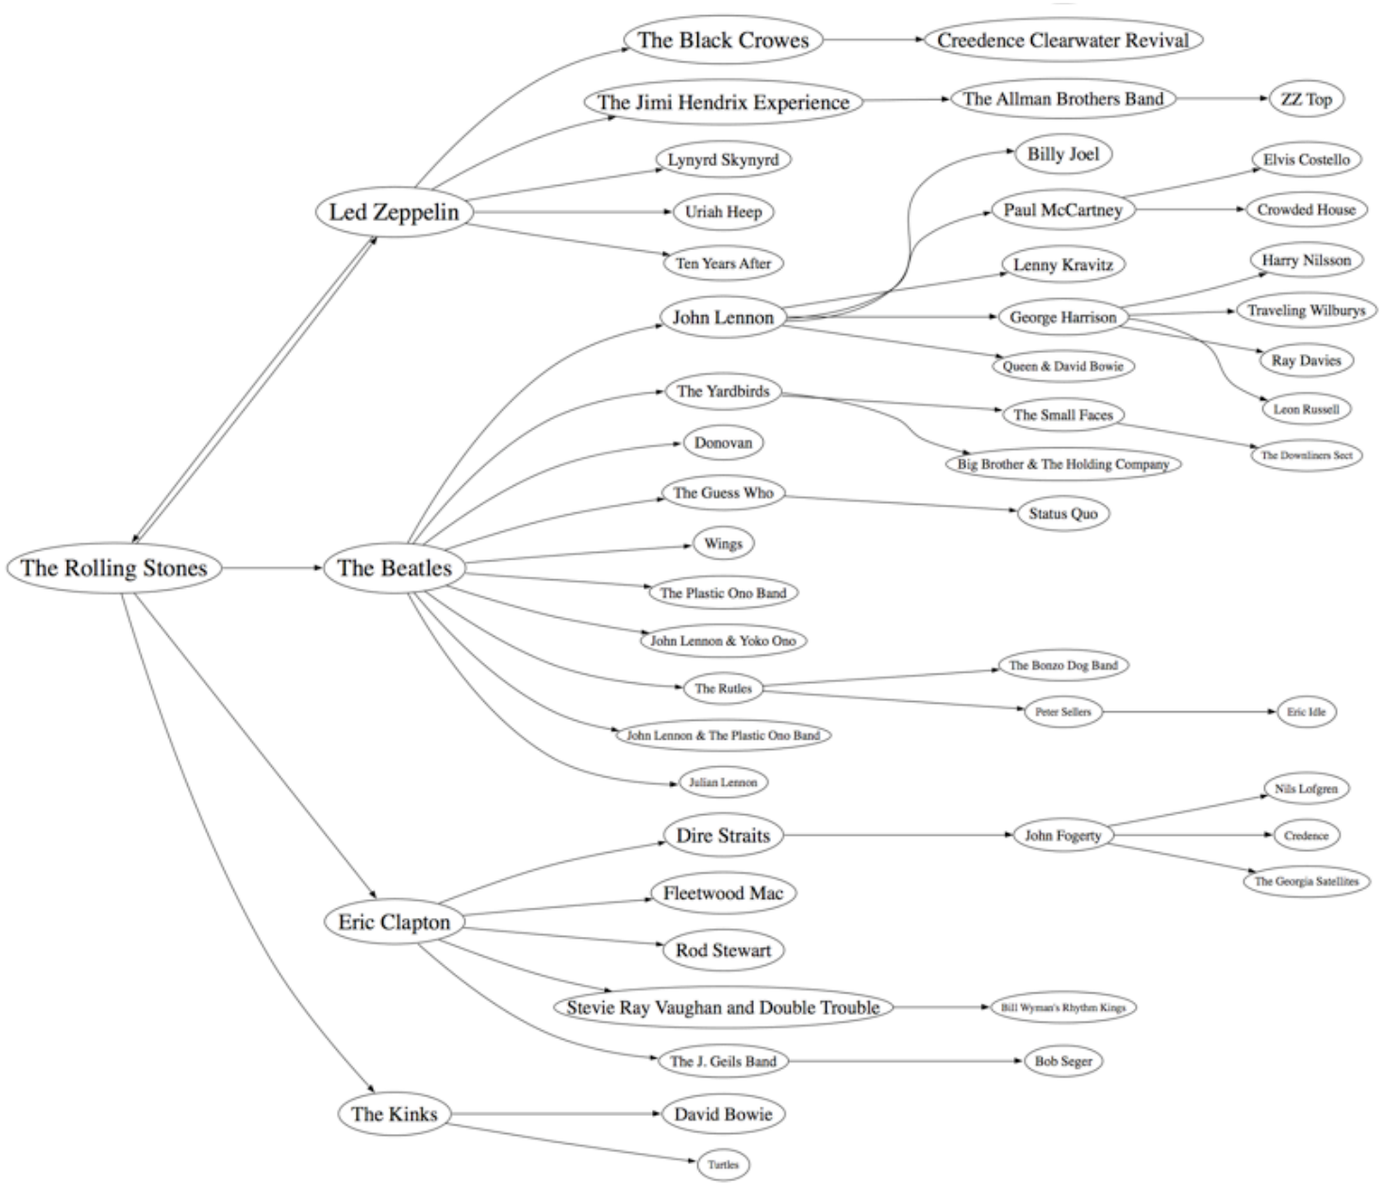

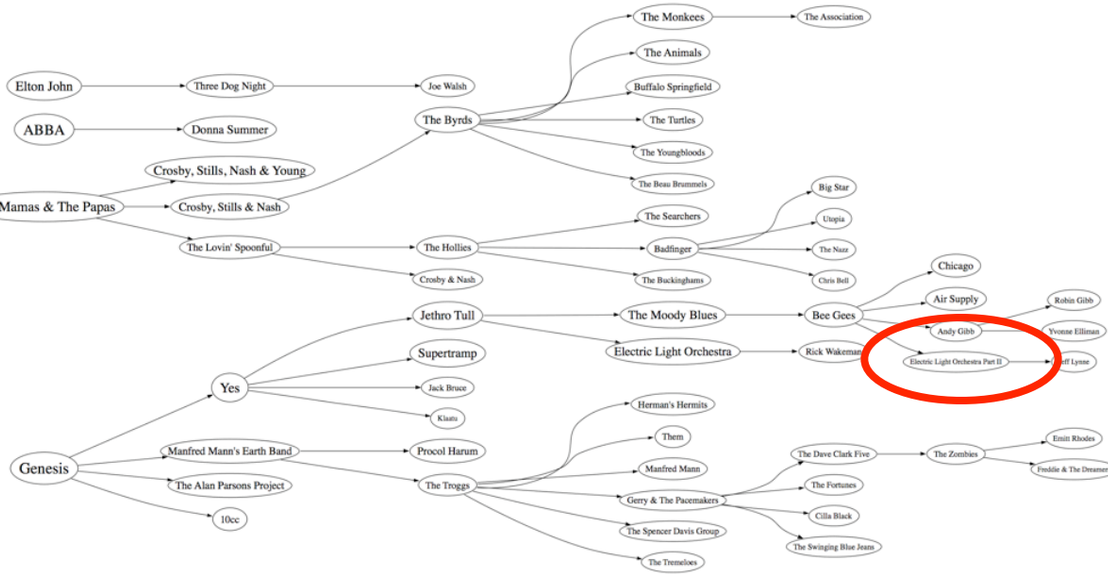

Attach artists to the most similar artist that has greater familiarity

Putting ELO II in its place

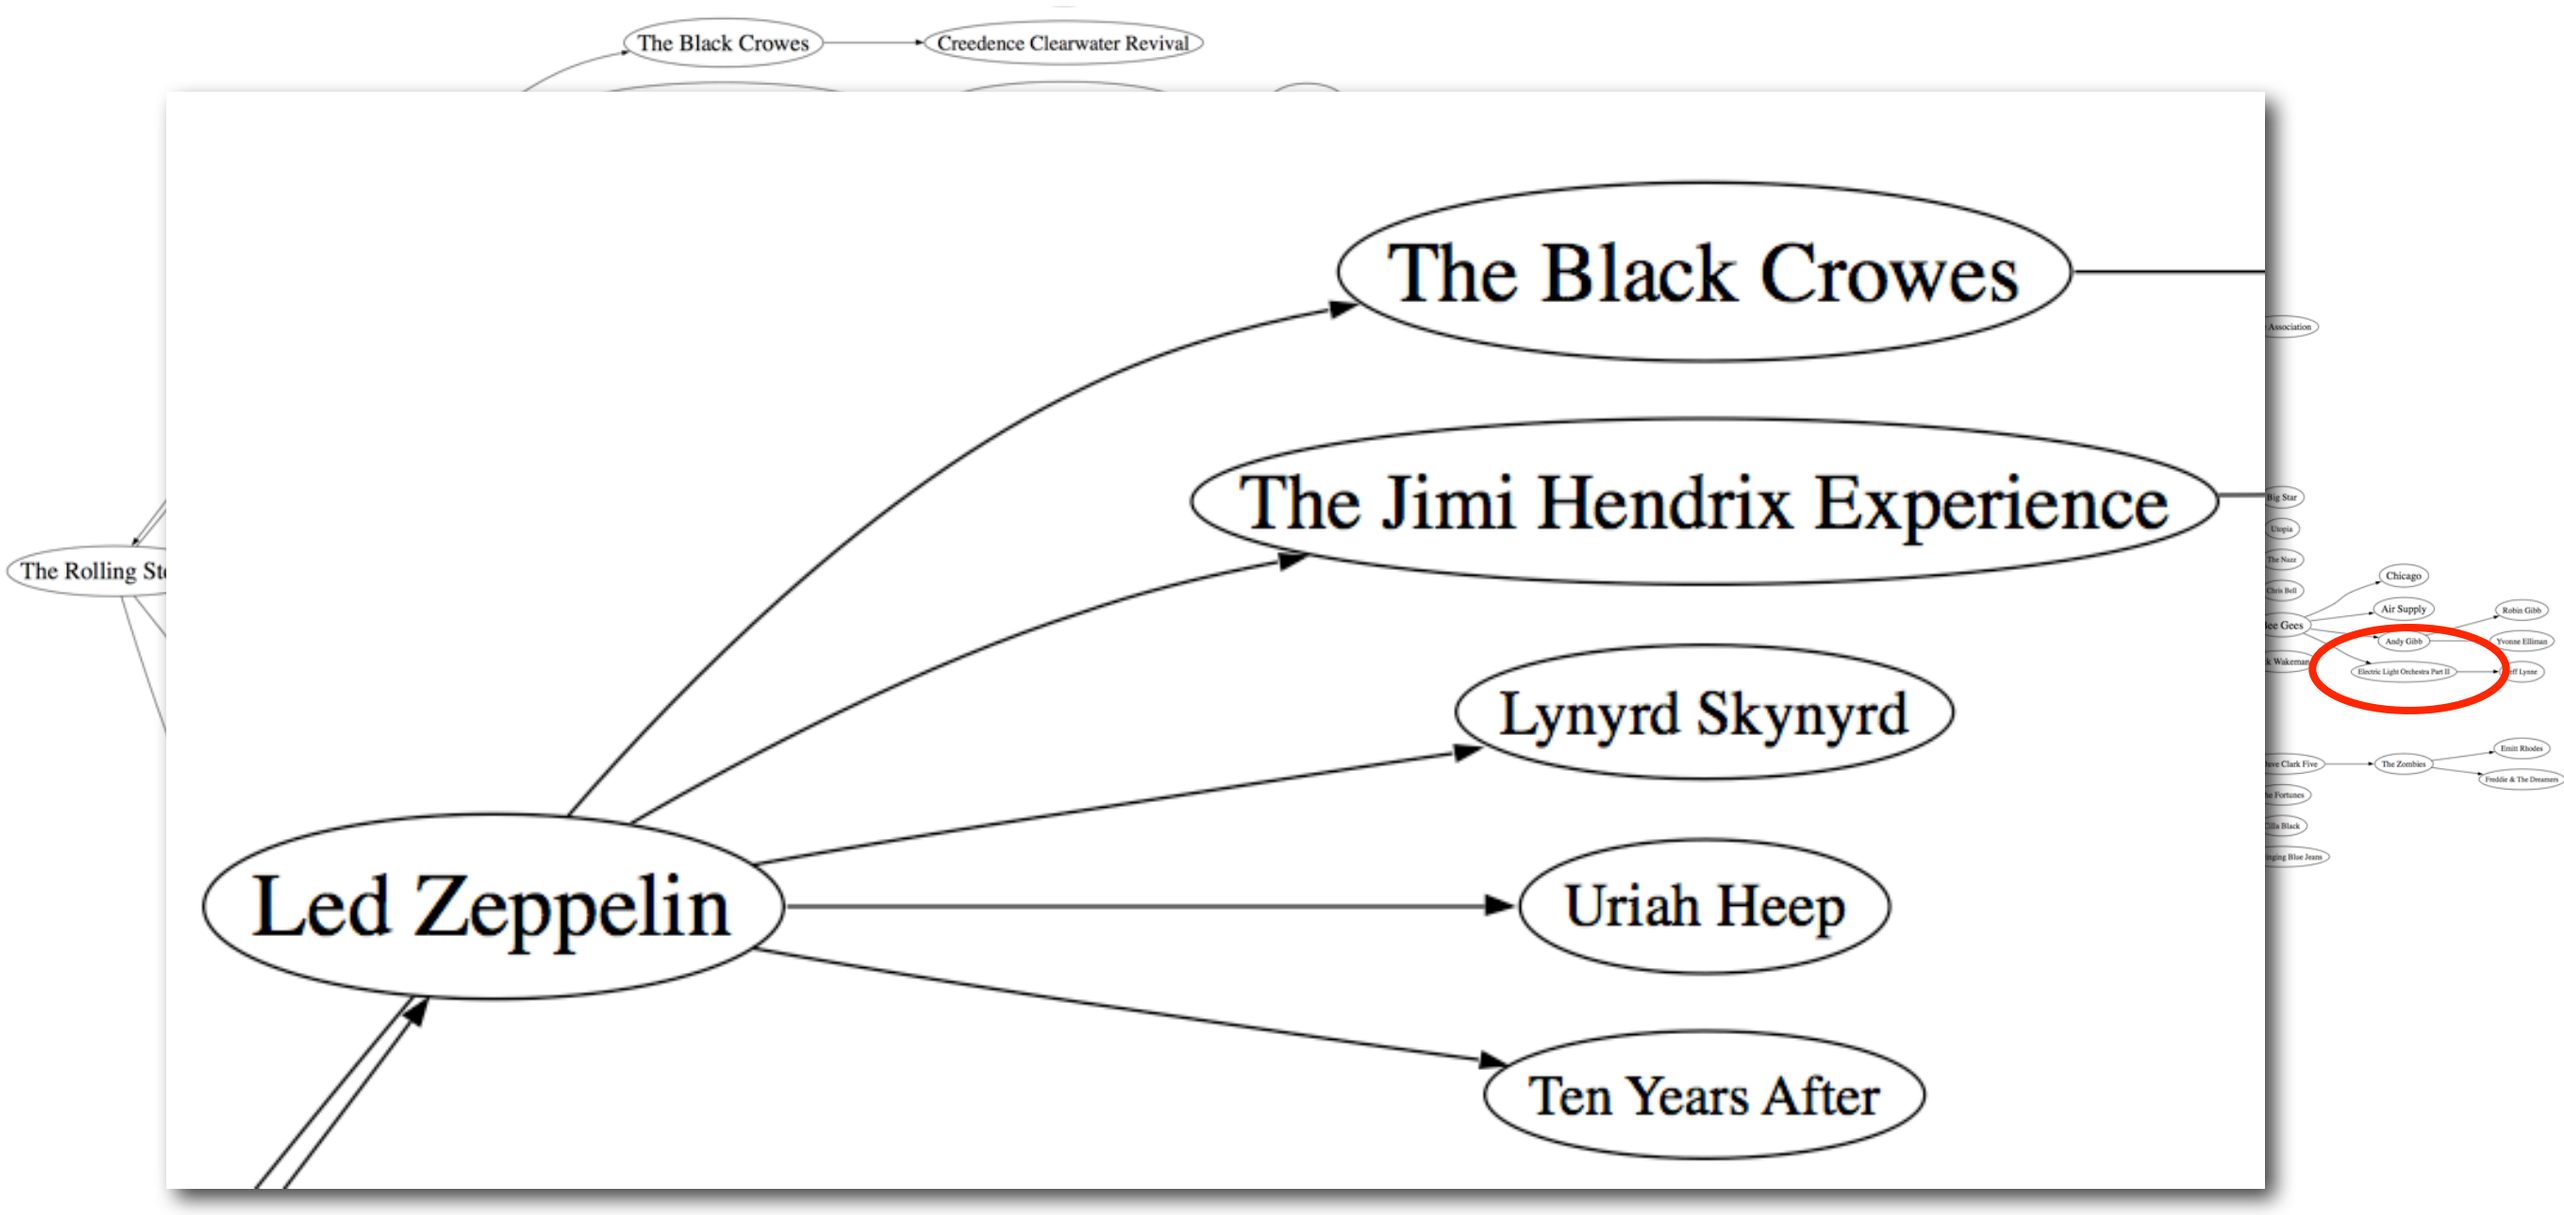

Attach artists to the most similar artist that has greater familiarity

Making the graph connected

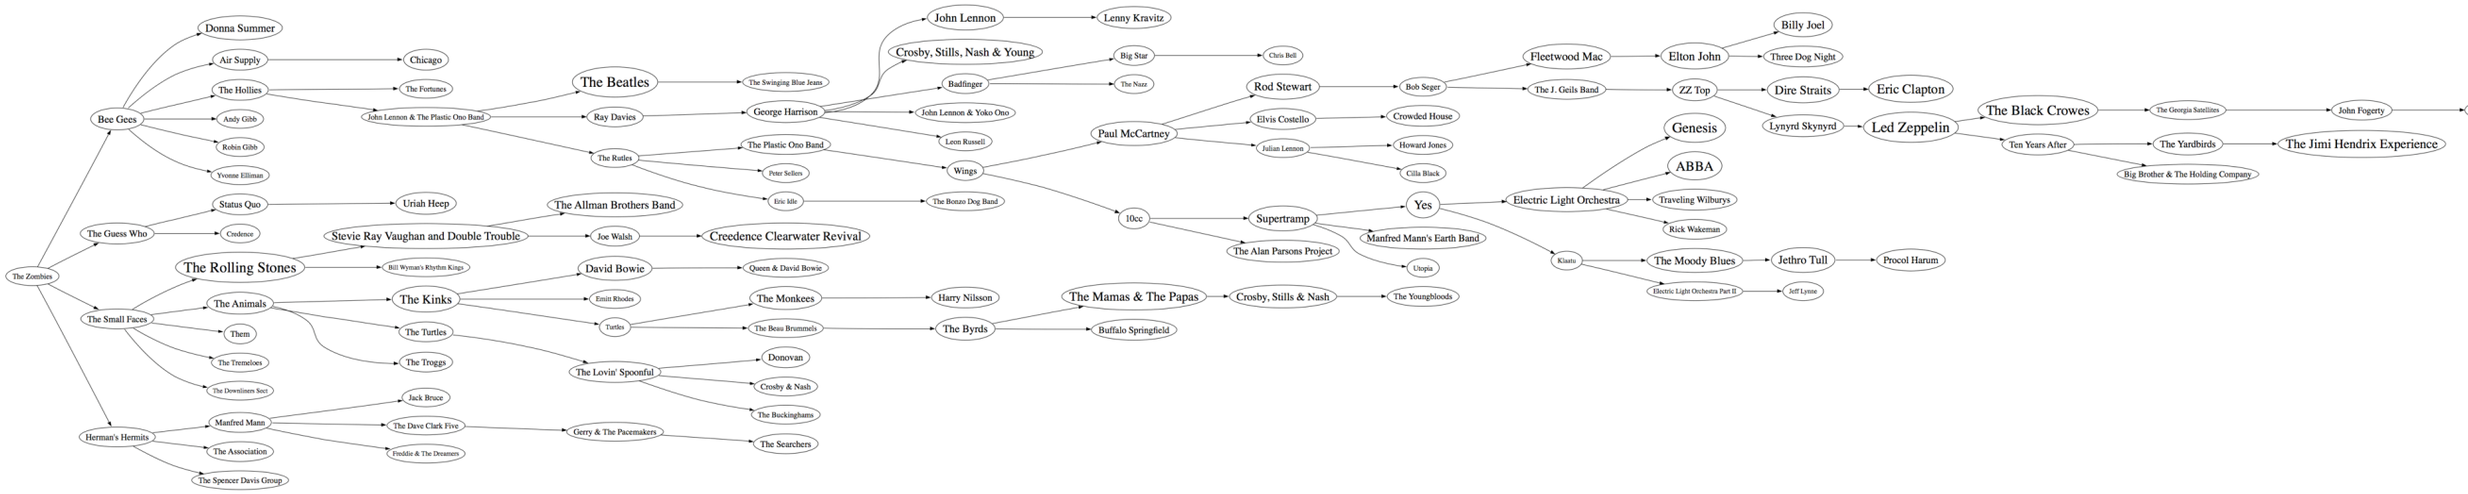

Build a minimum spanning tree from the graph

Making the graph connected

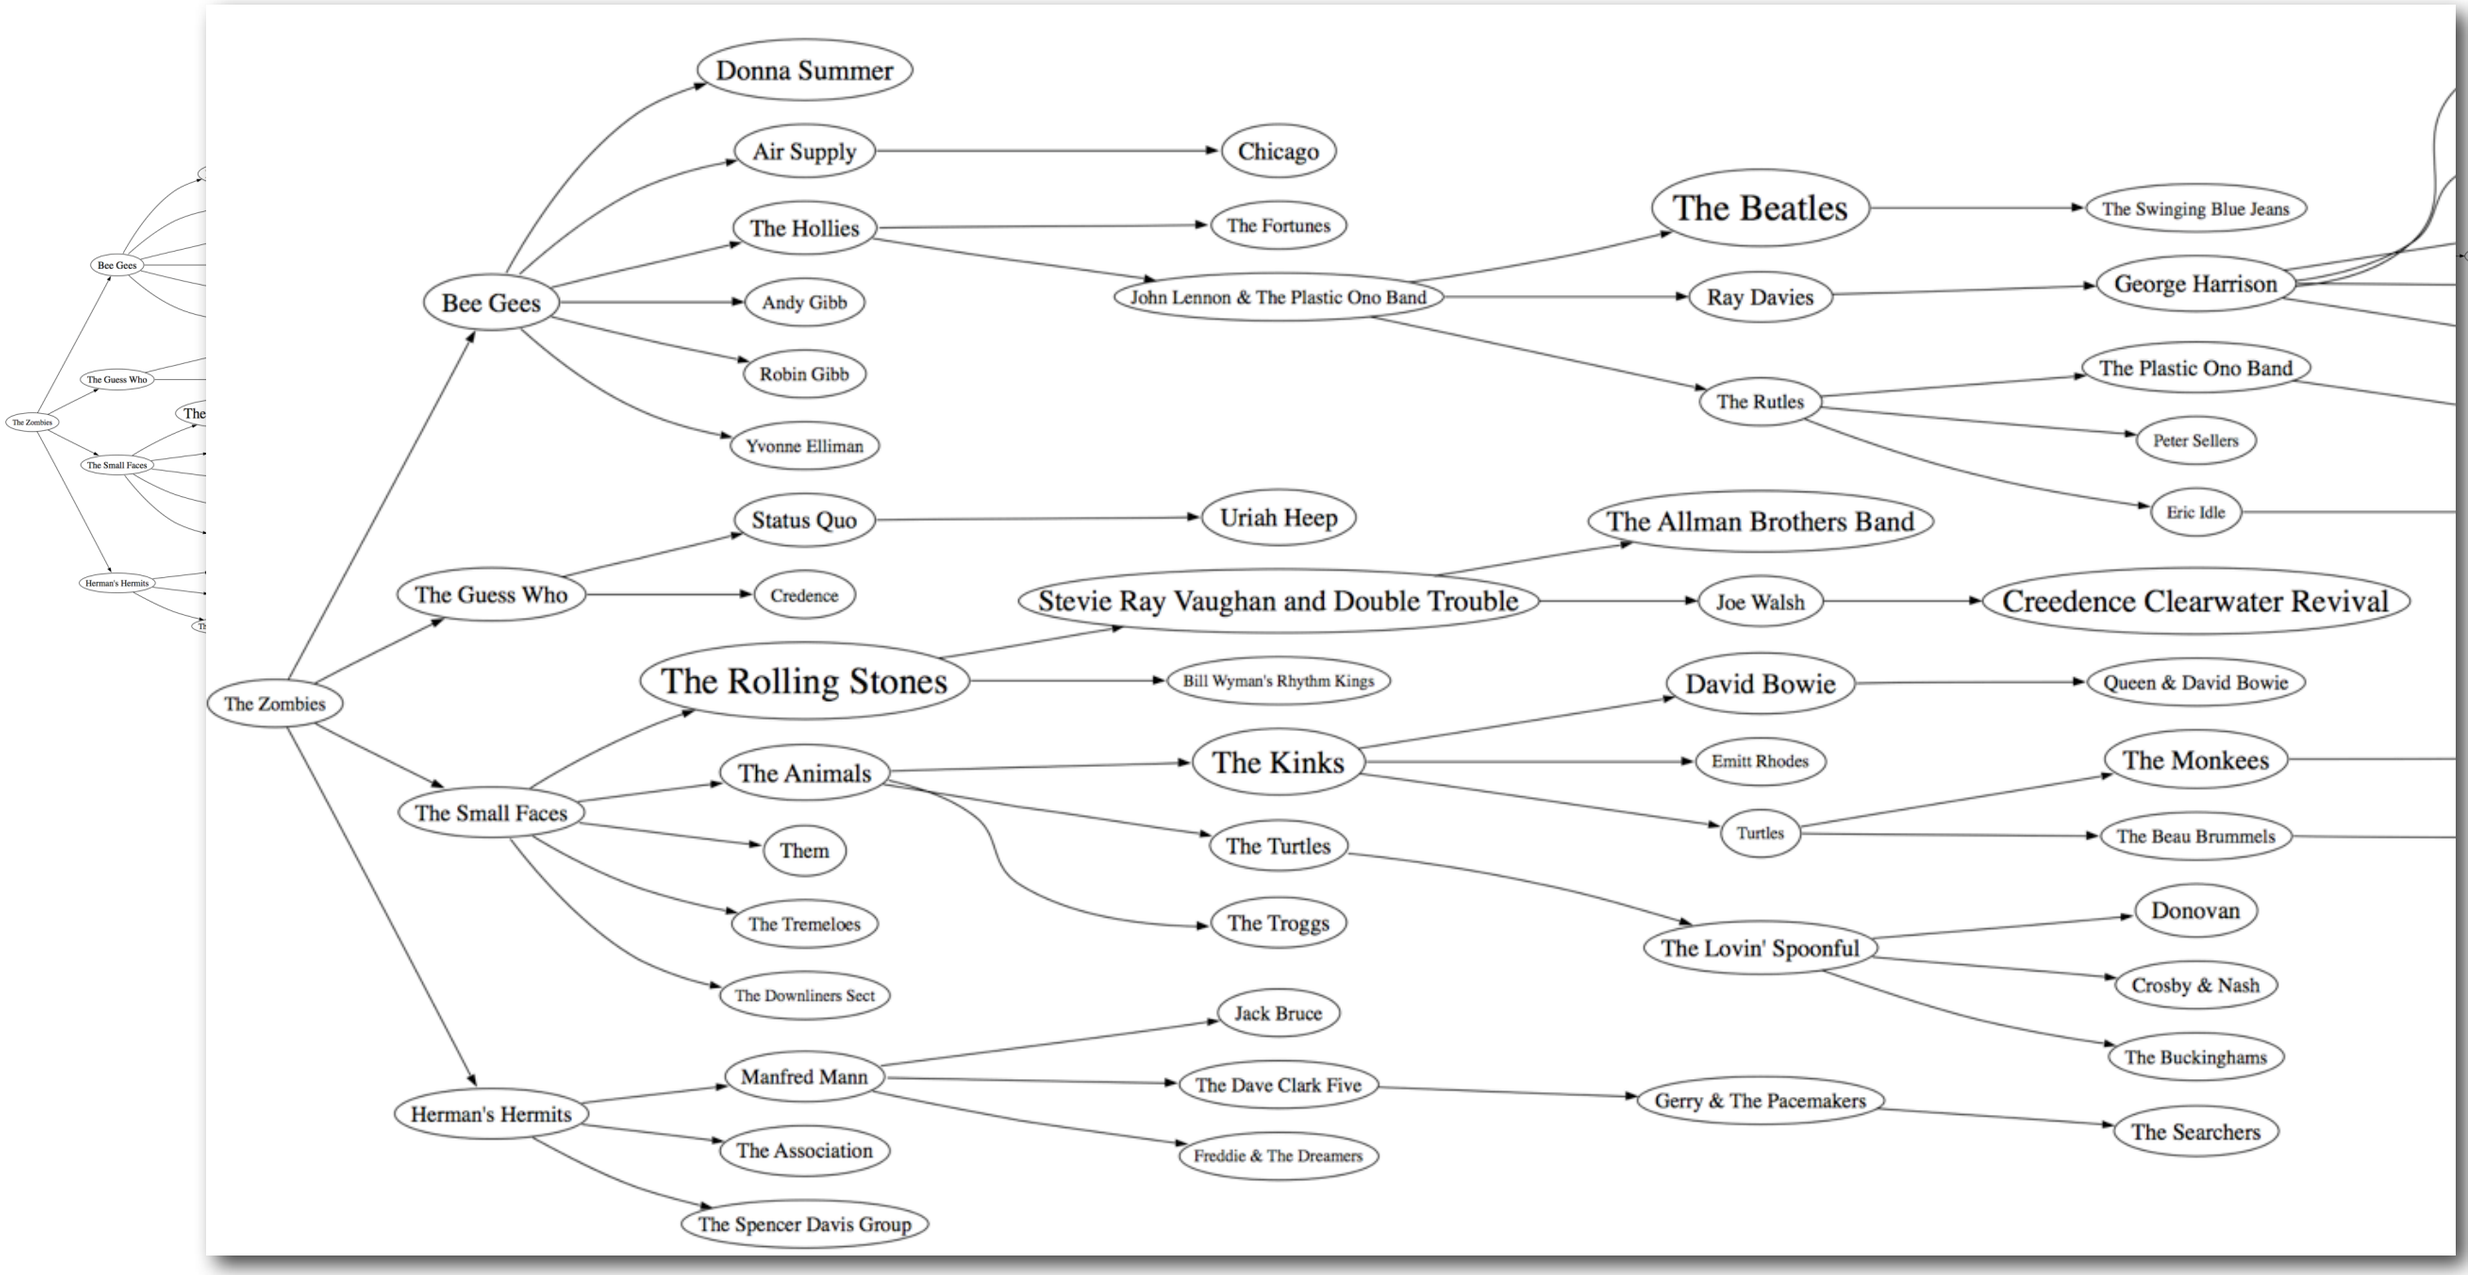

Build a minimum spanning tree from the graph

Making the graph connected

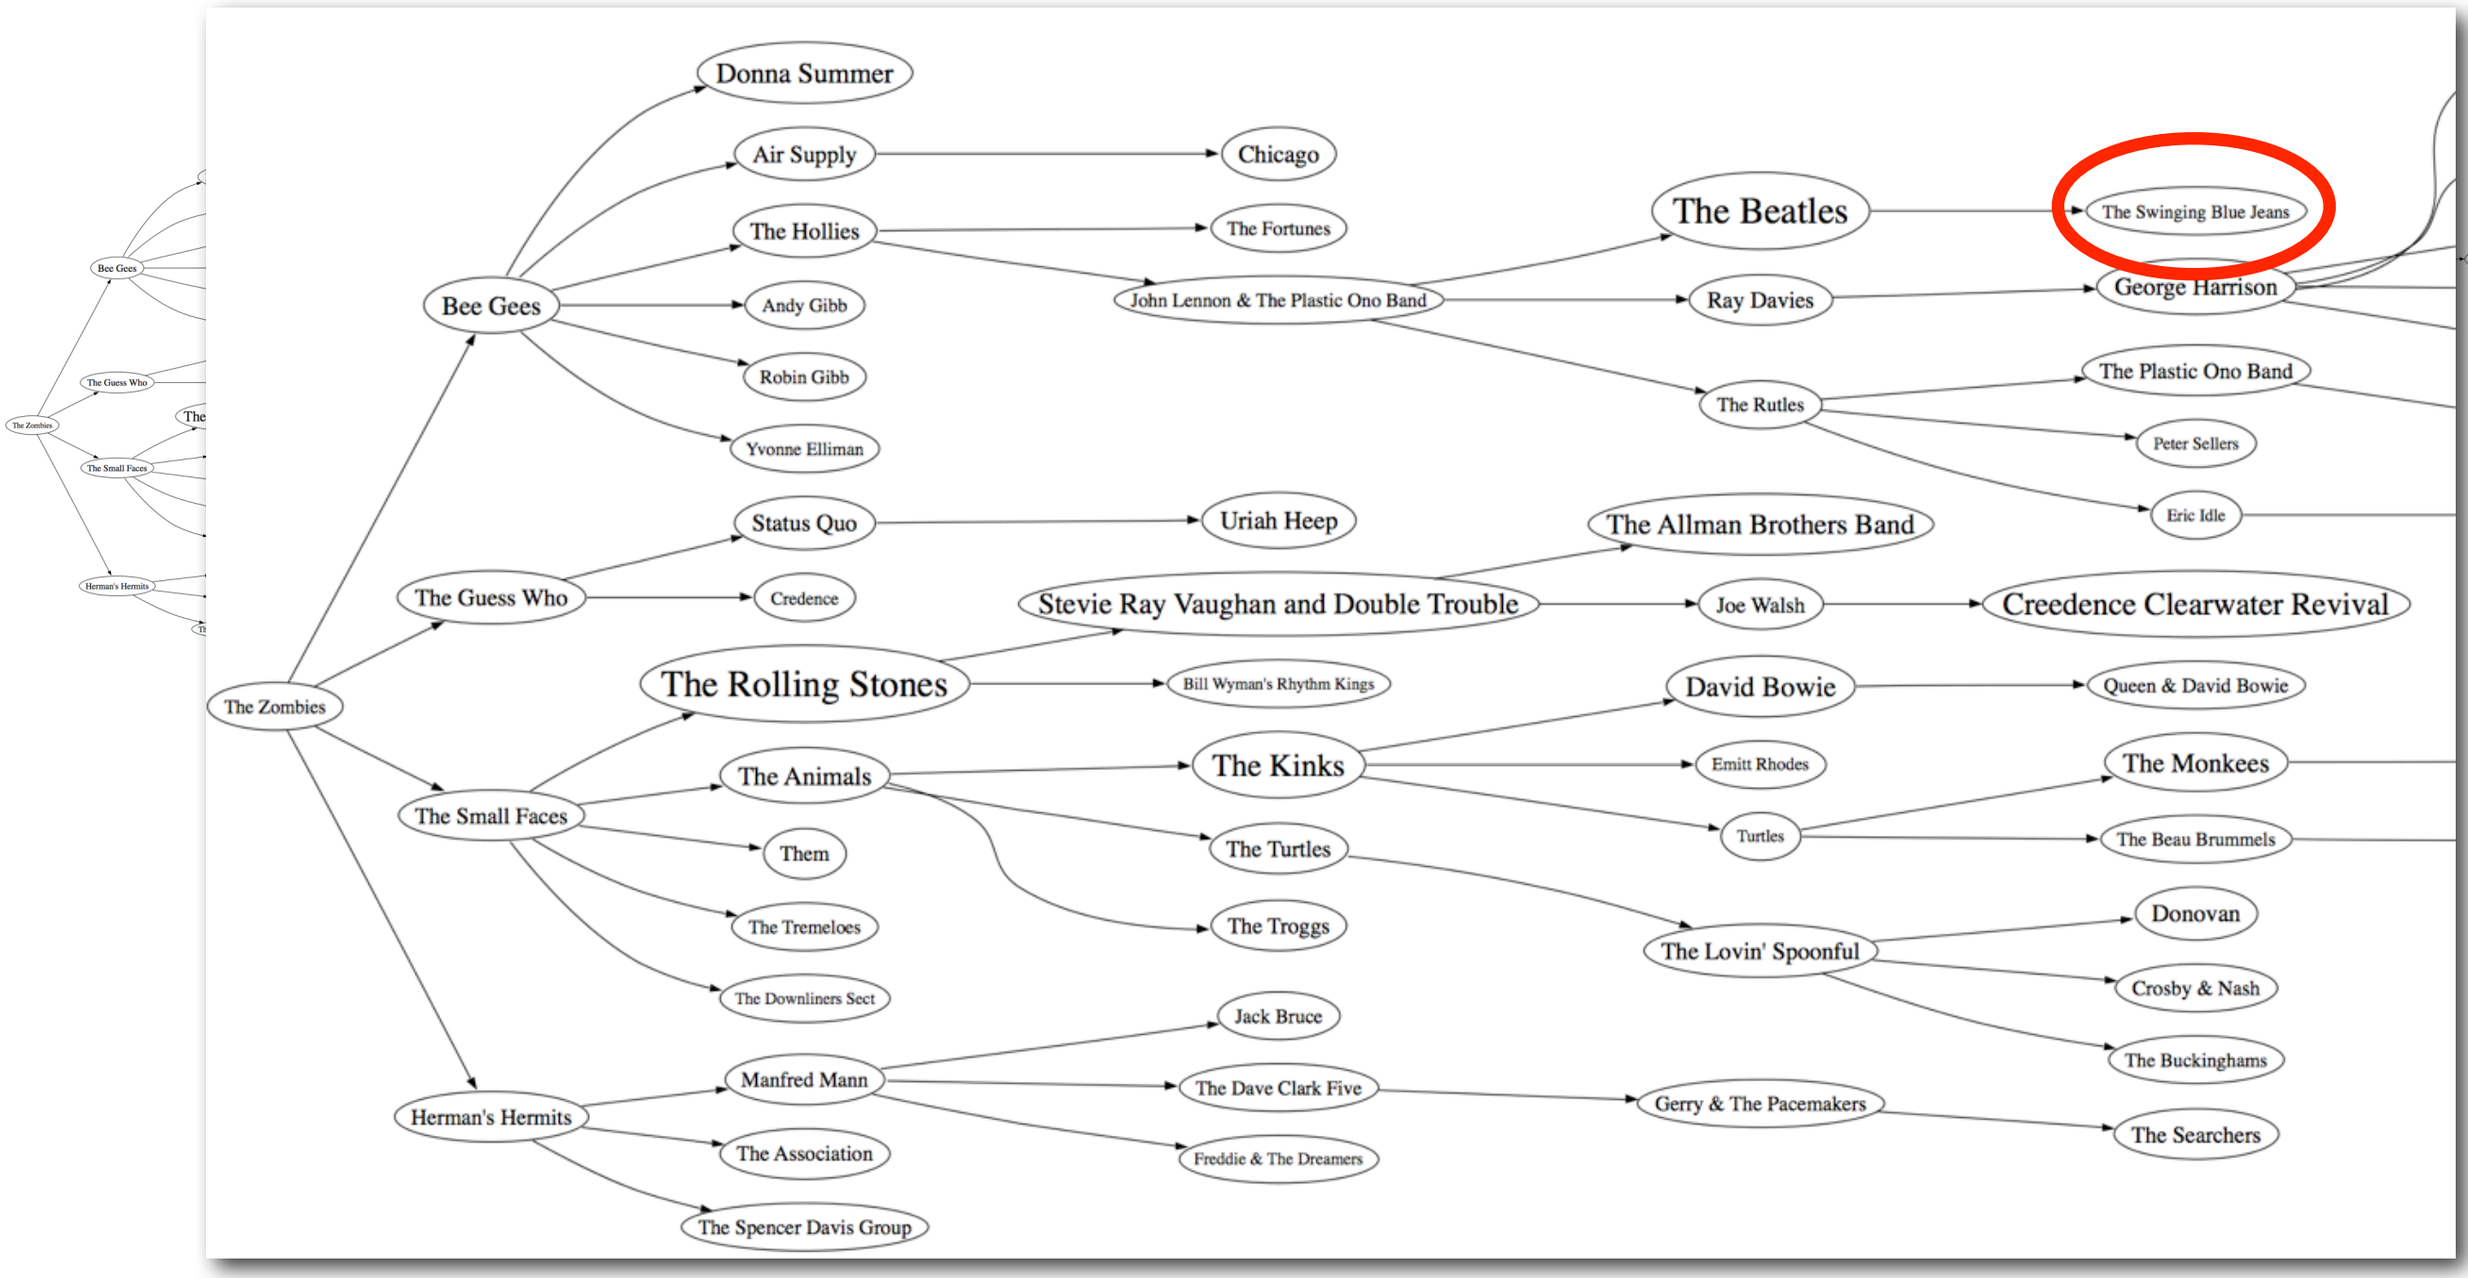

Build a minimum spanning tree from the graph

Fixing the swinging blue jeans problem

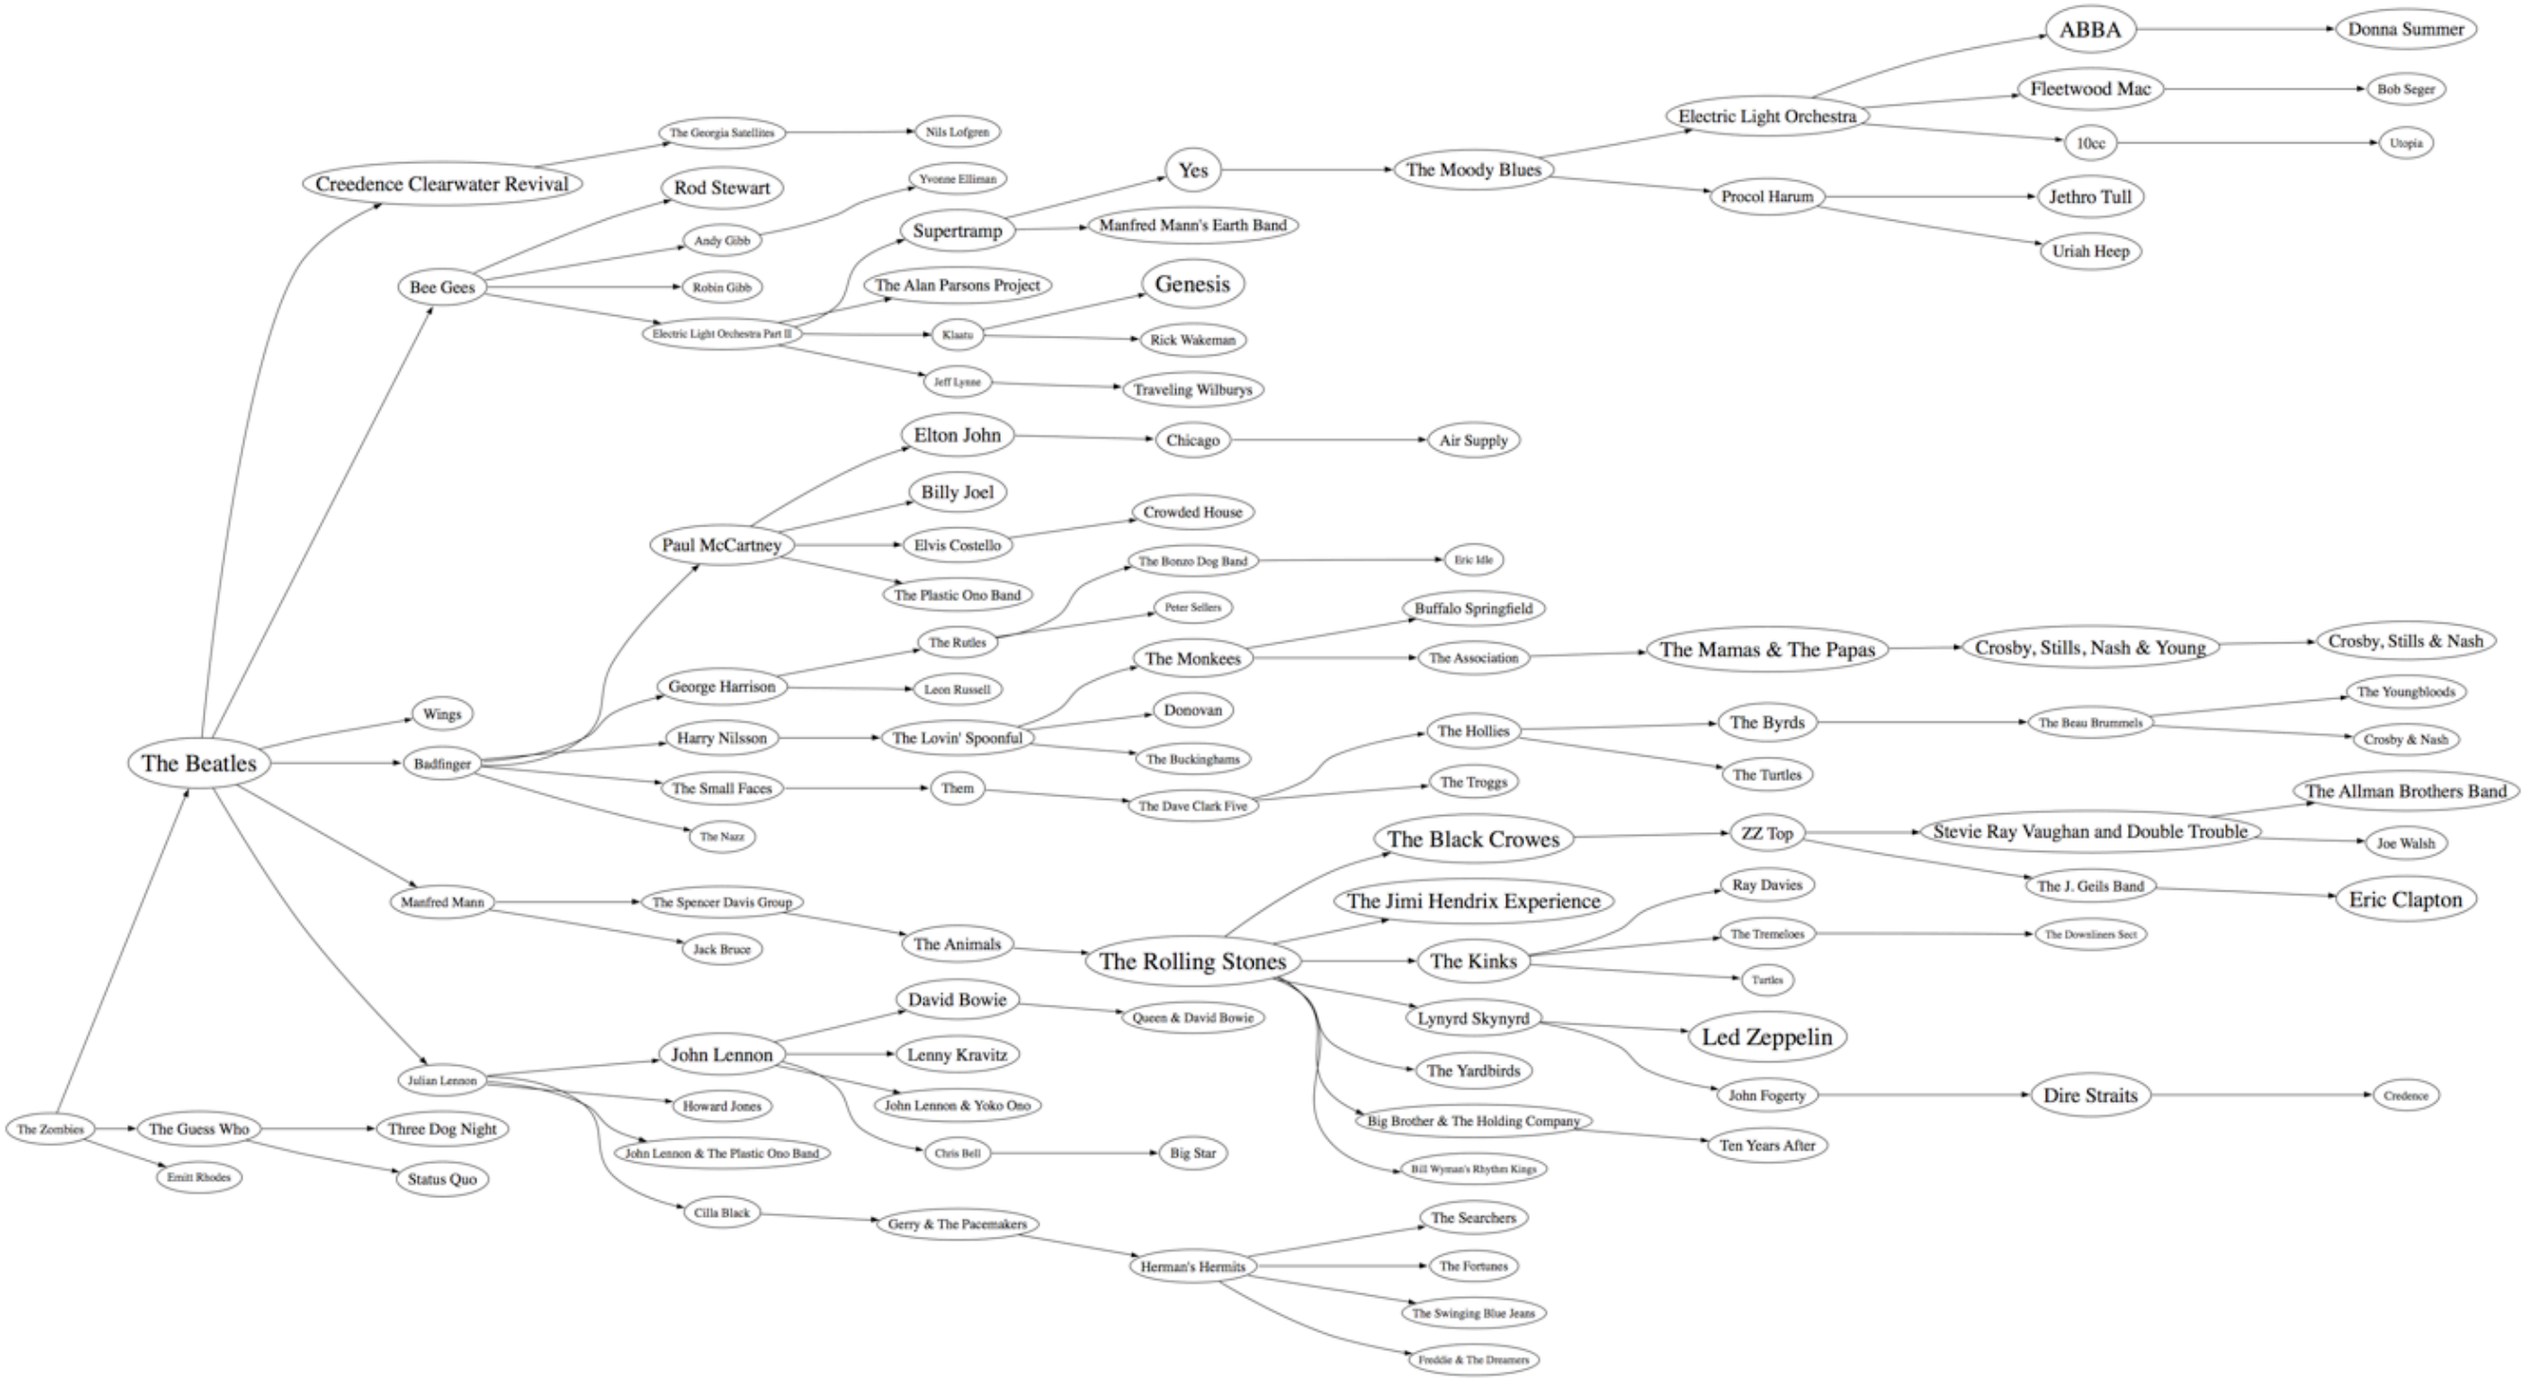

Prefer connections to bands of similar familiarity

Other near-neighbor graphs

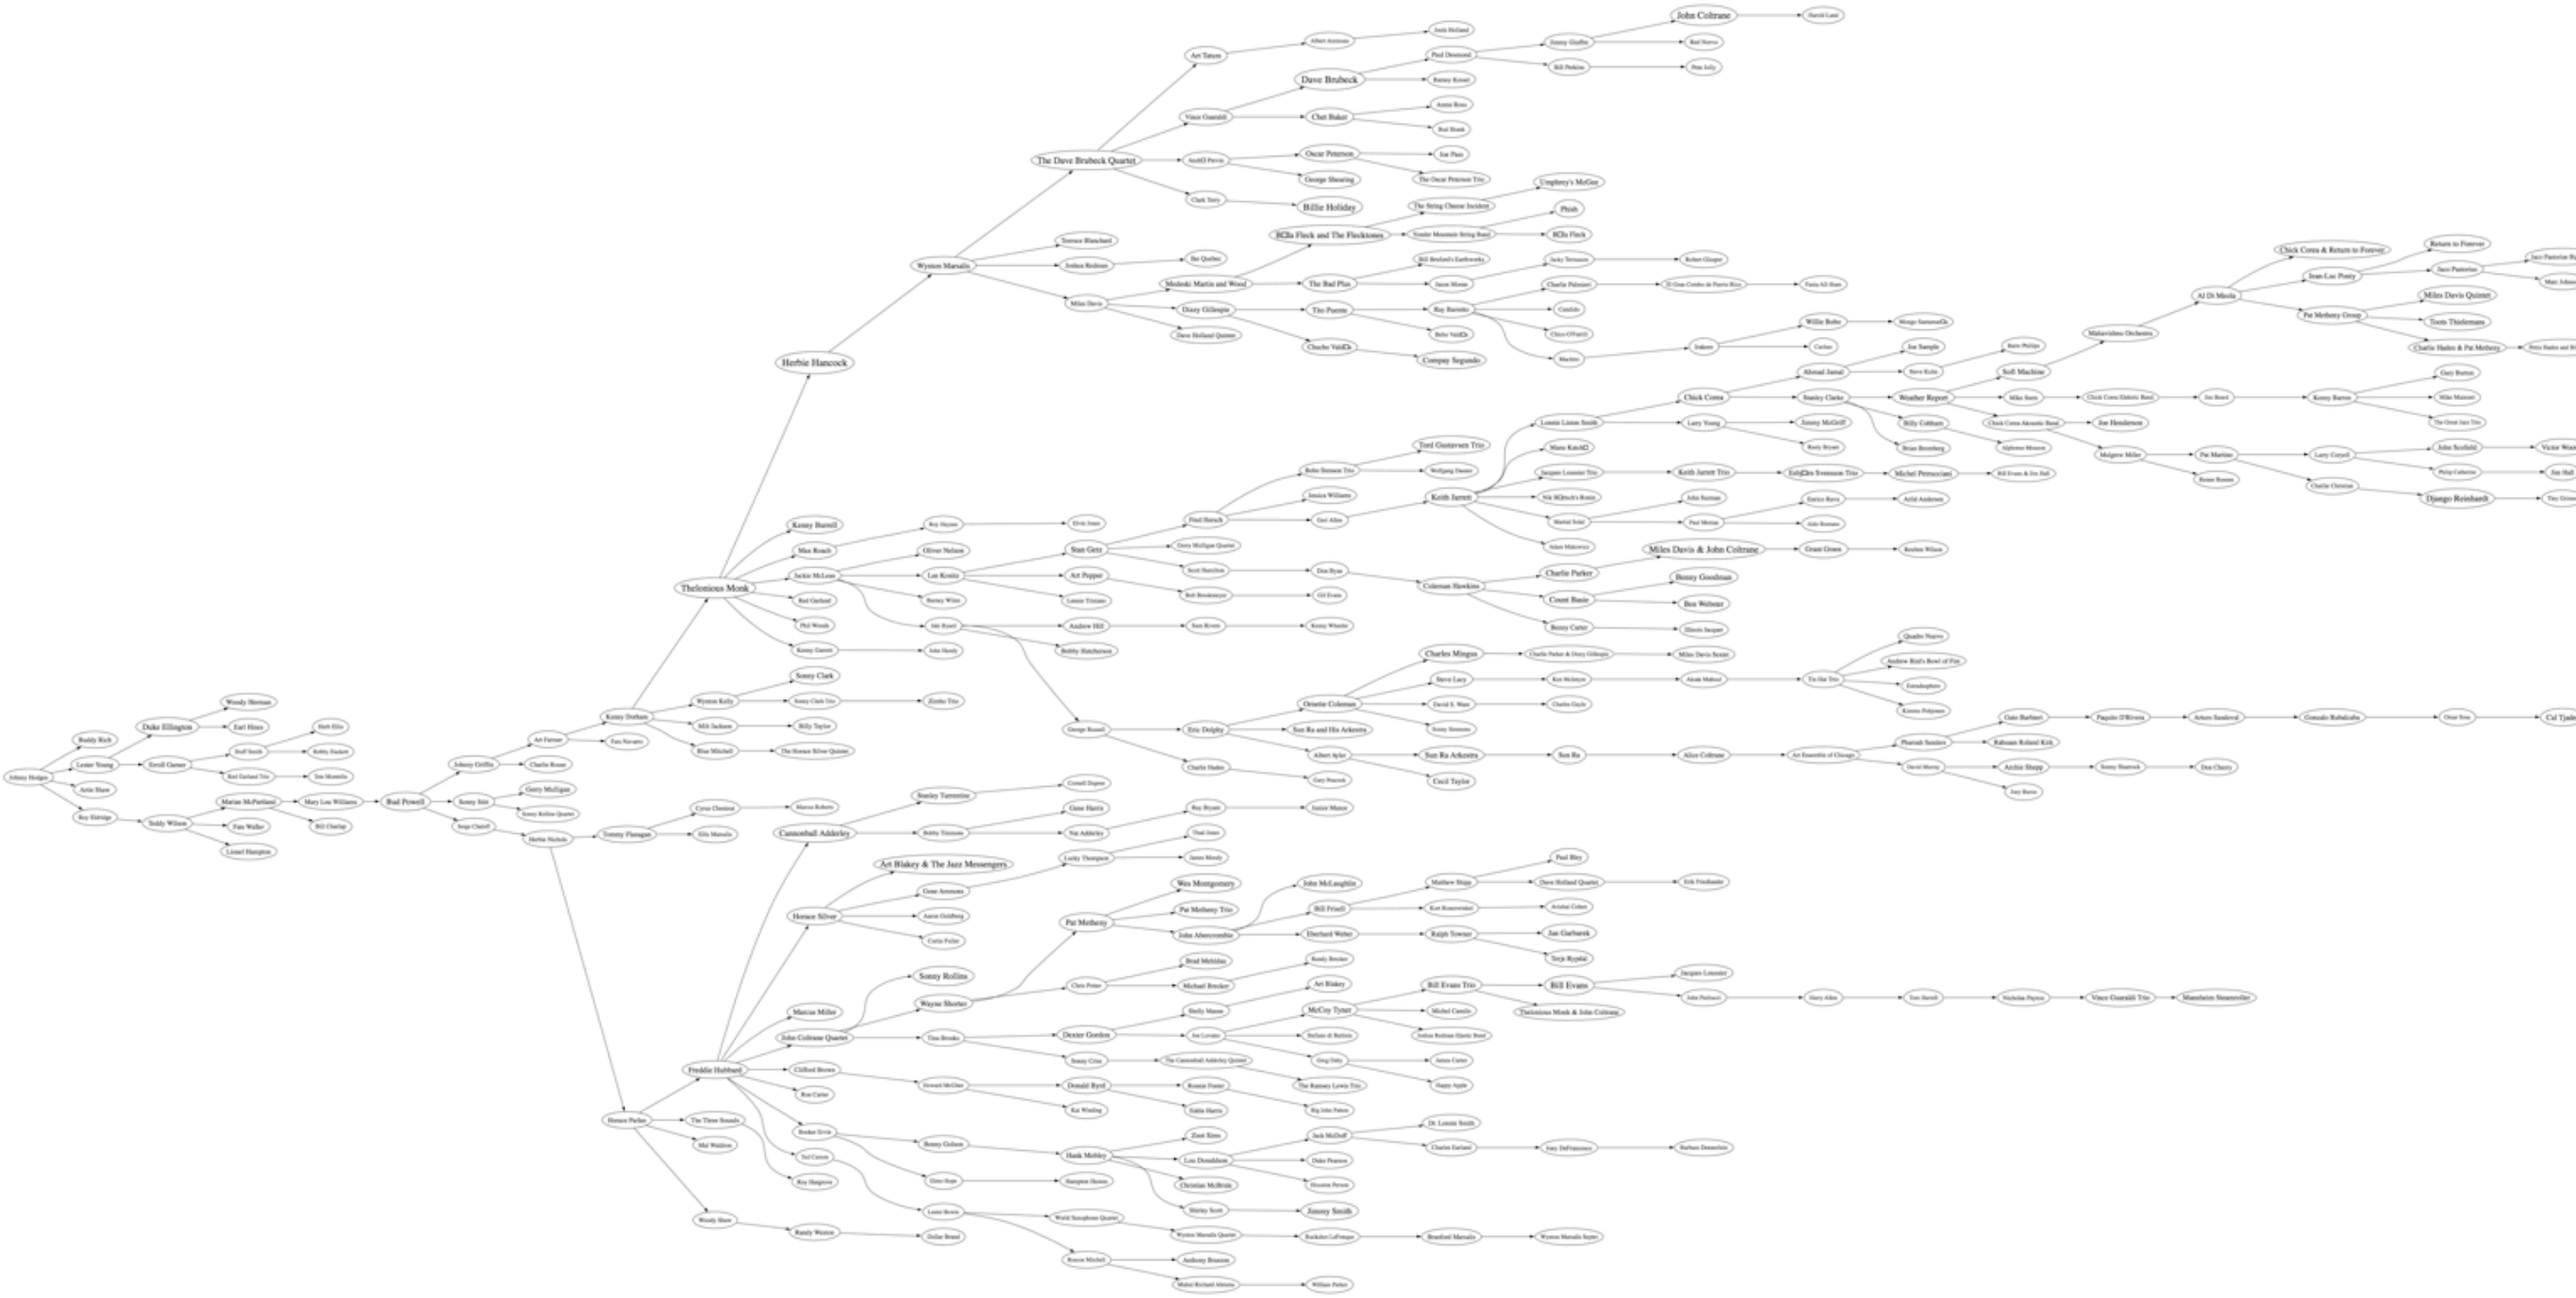

Miles Davis

Other near-neighbor graphs

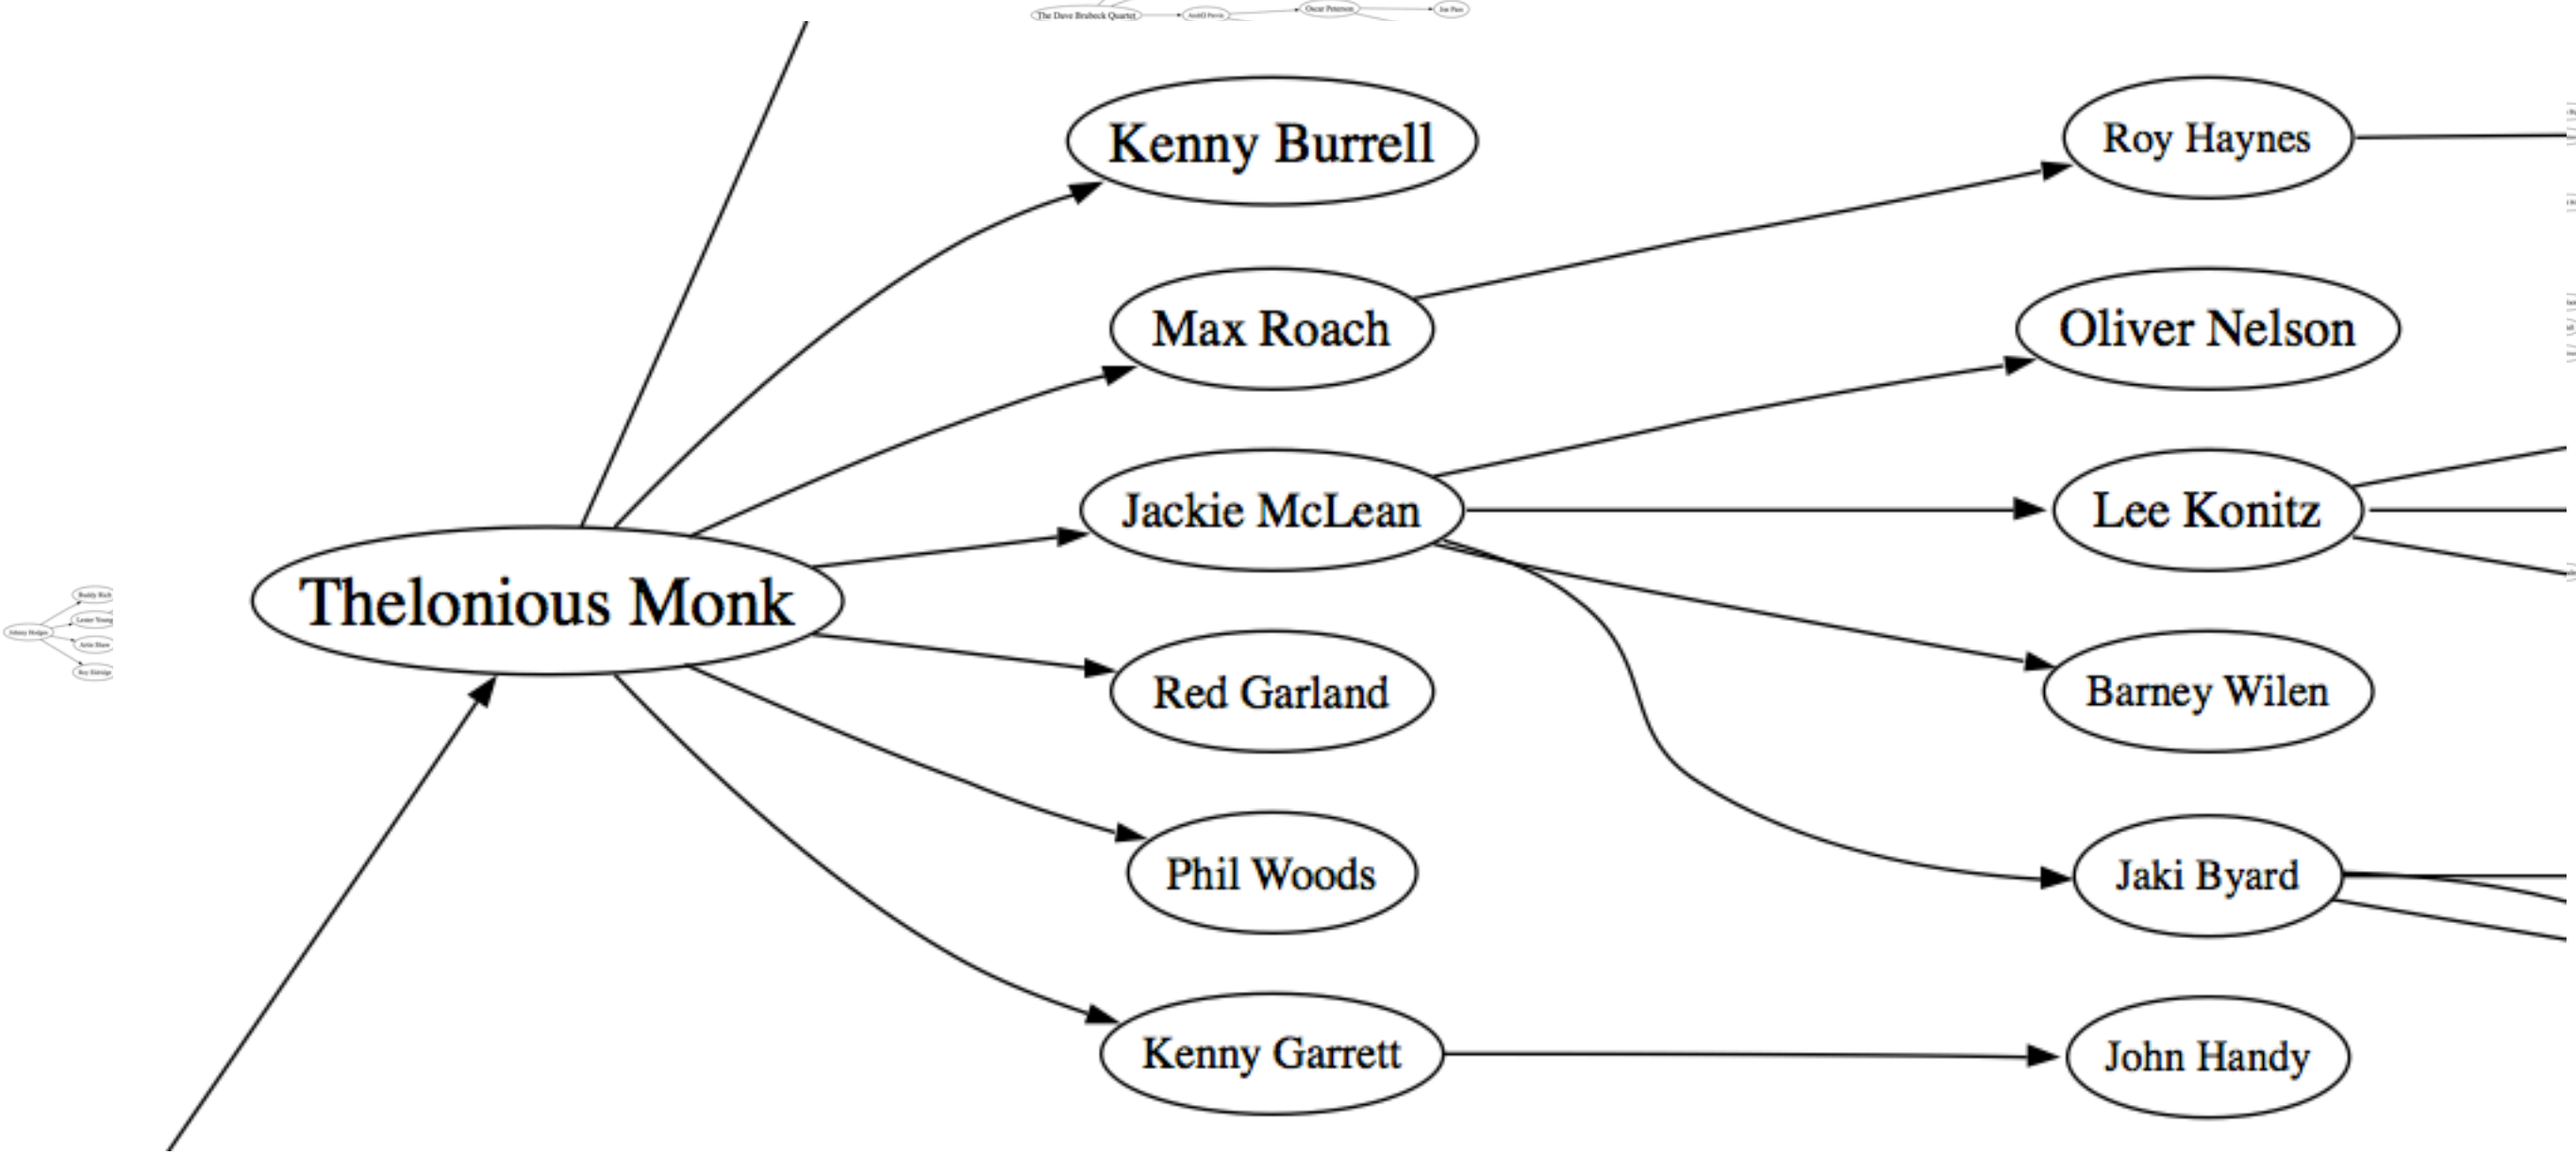

Miles Davis

2K artists

A browsing network for music discovery

2K artists

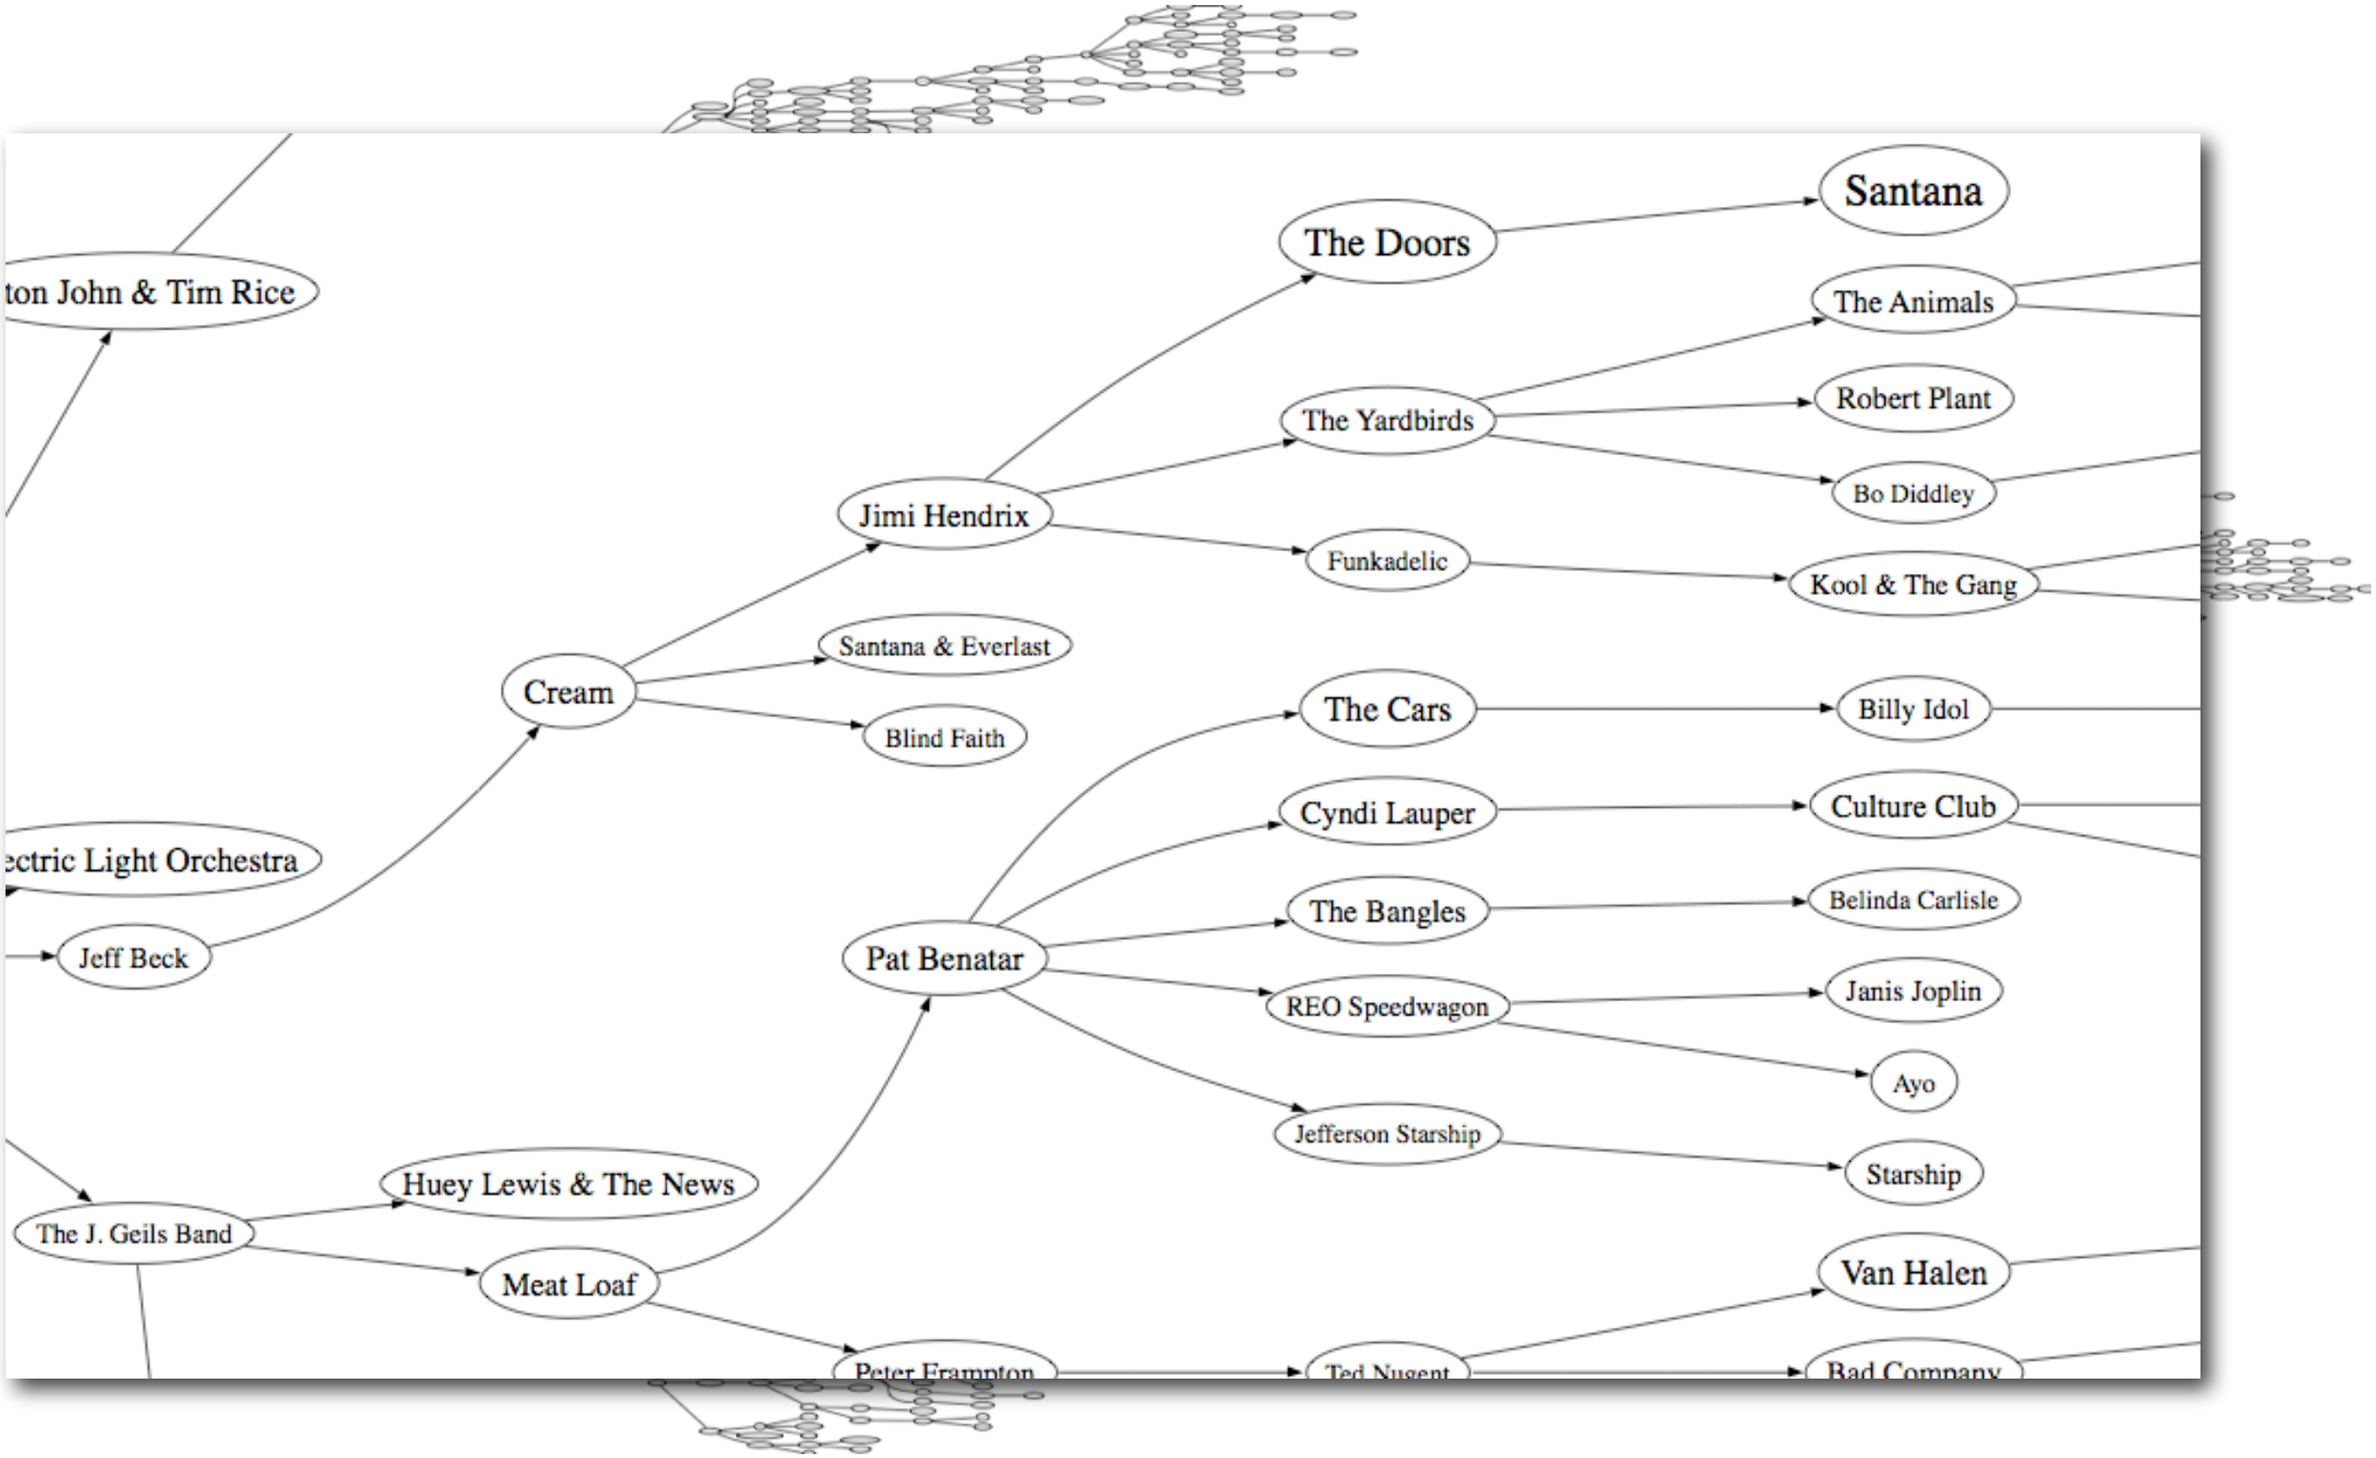

A browsing network for music discovery

Making the graph interactive

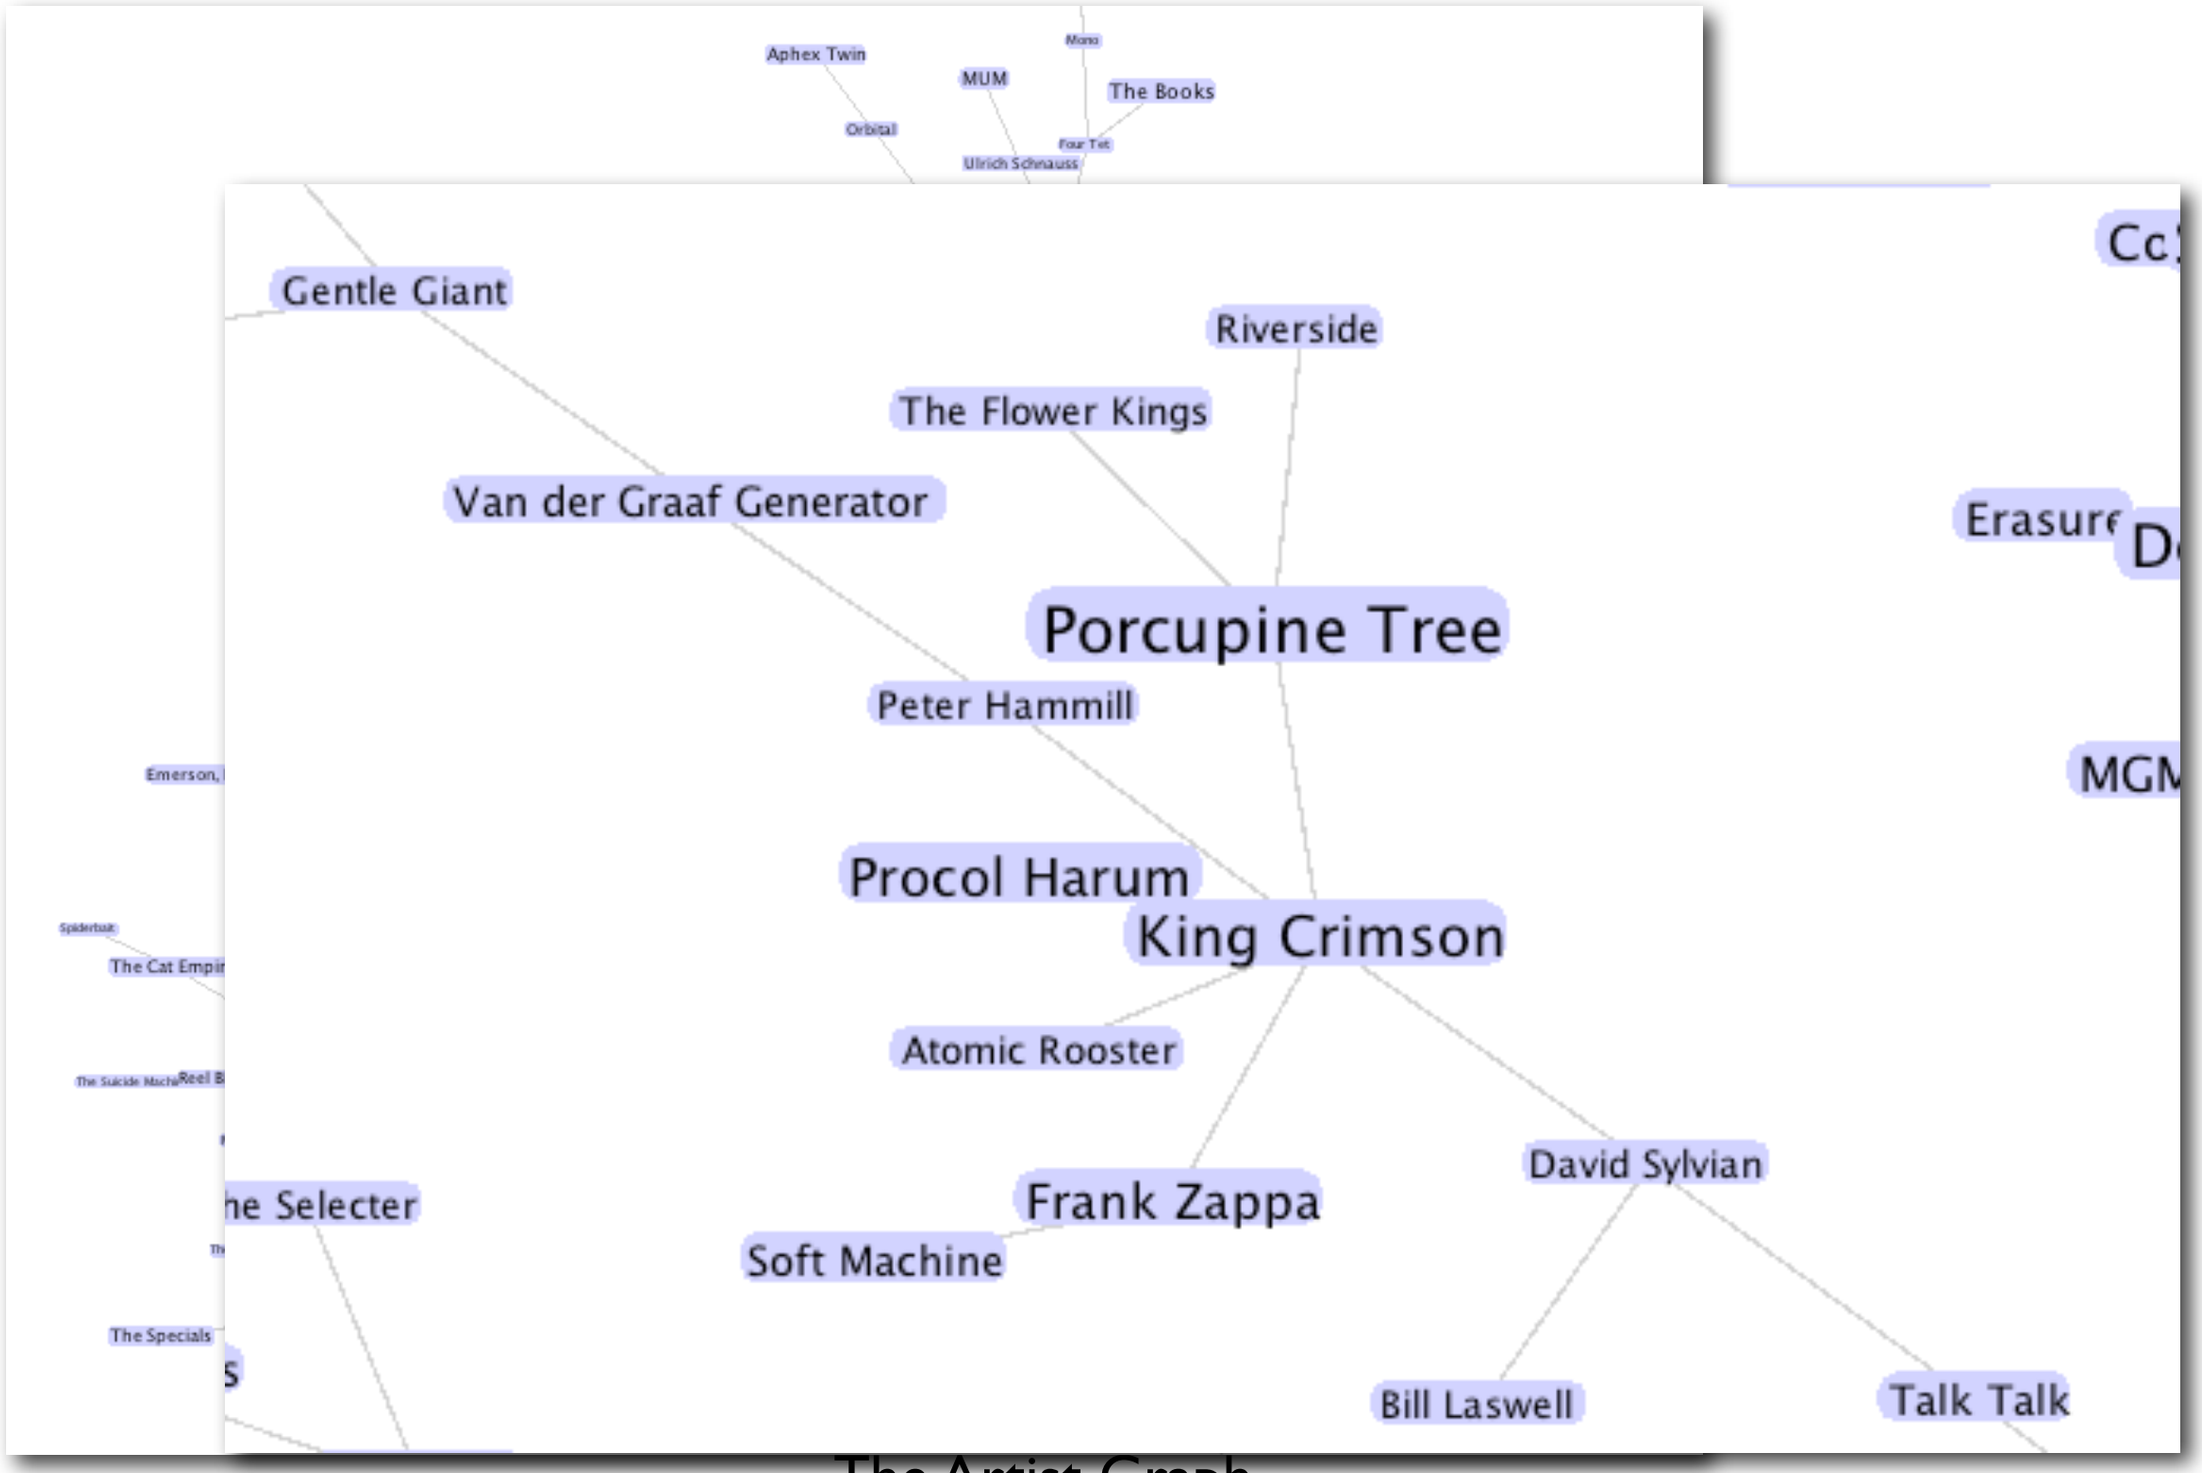

The Artist Graph

Force-Directed Edge Bundling for Graph Visualization

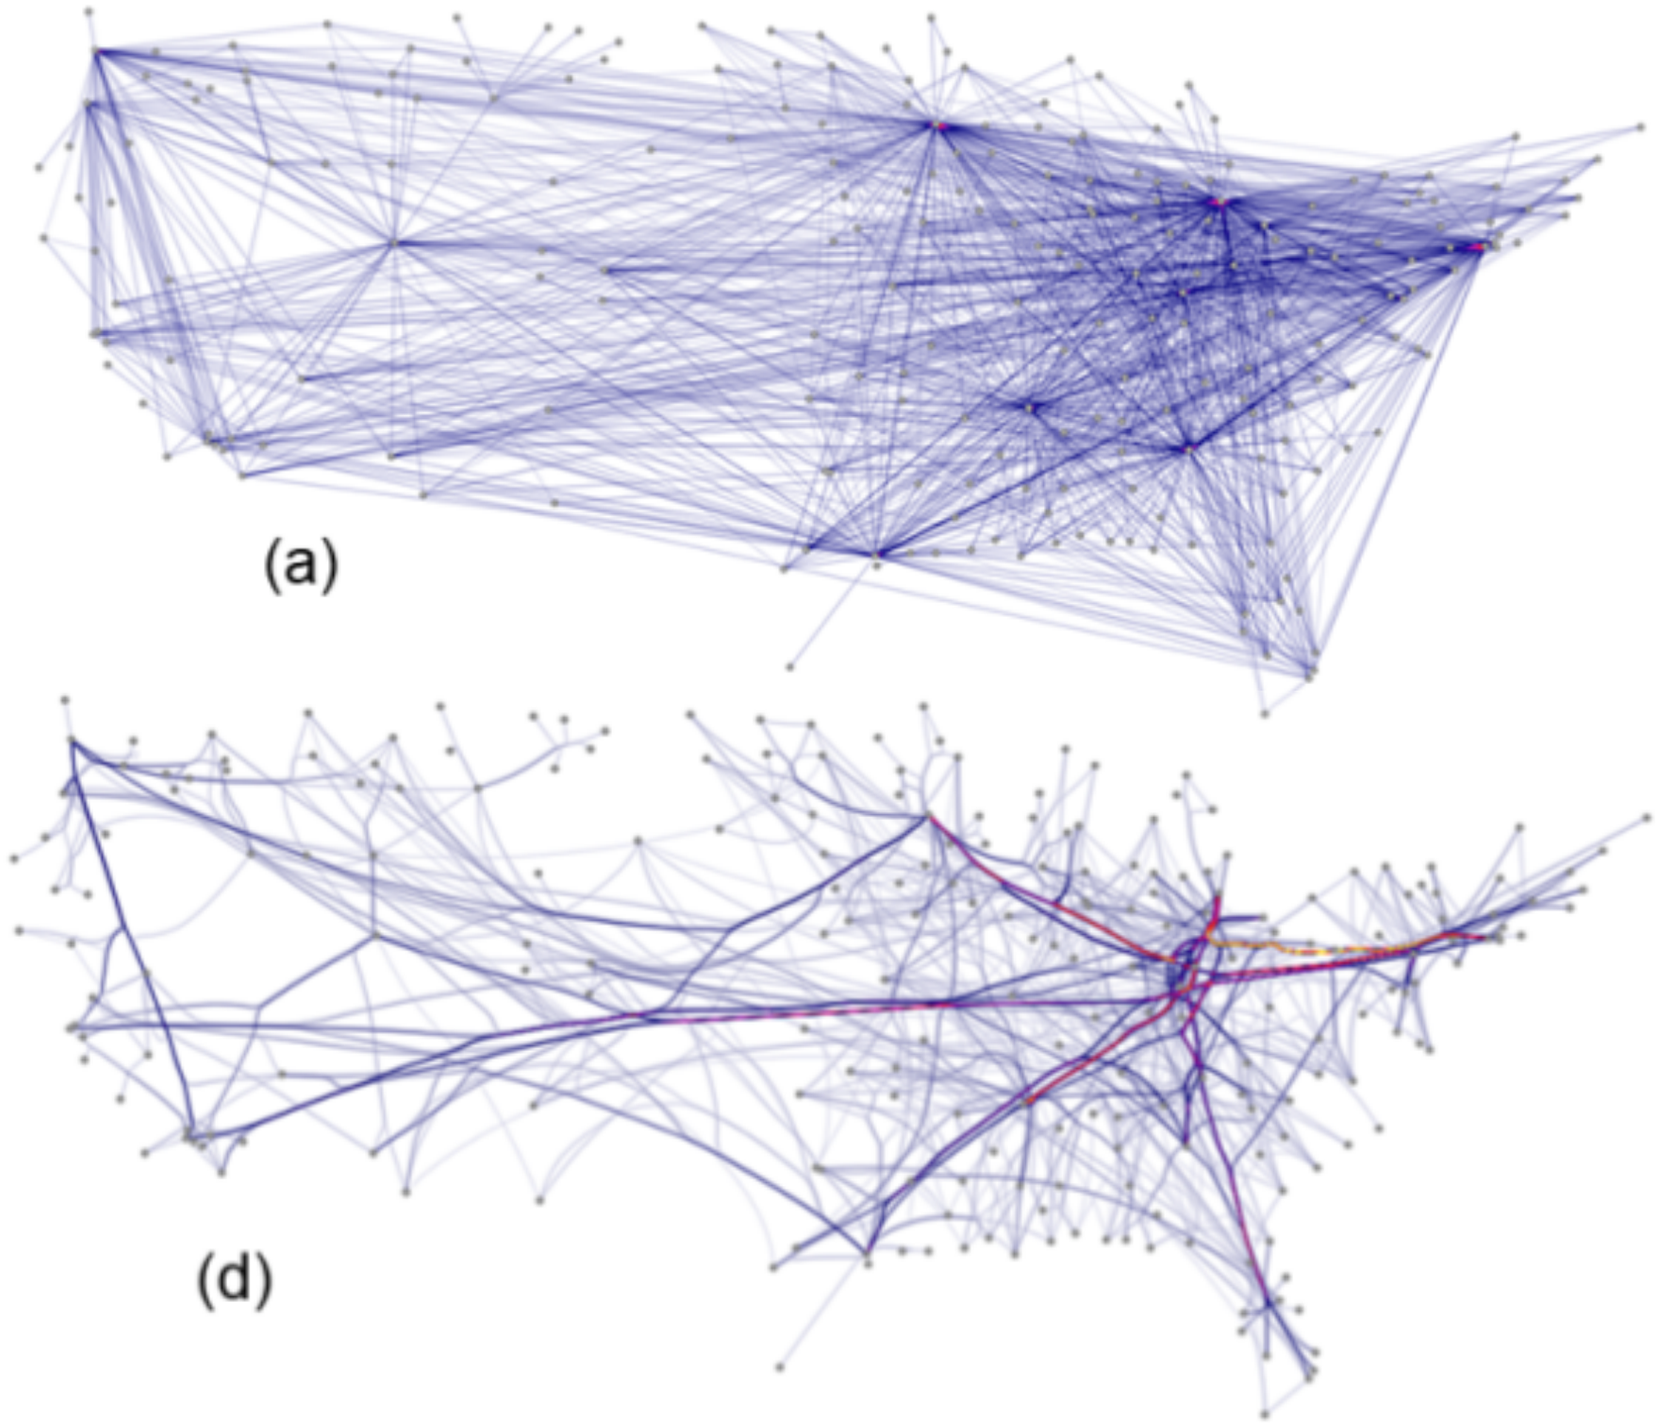

Danny Holten, Jarke J. van Wijk

Survey of Music Visualizations

Graphs

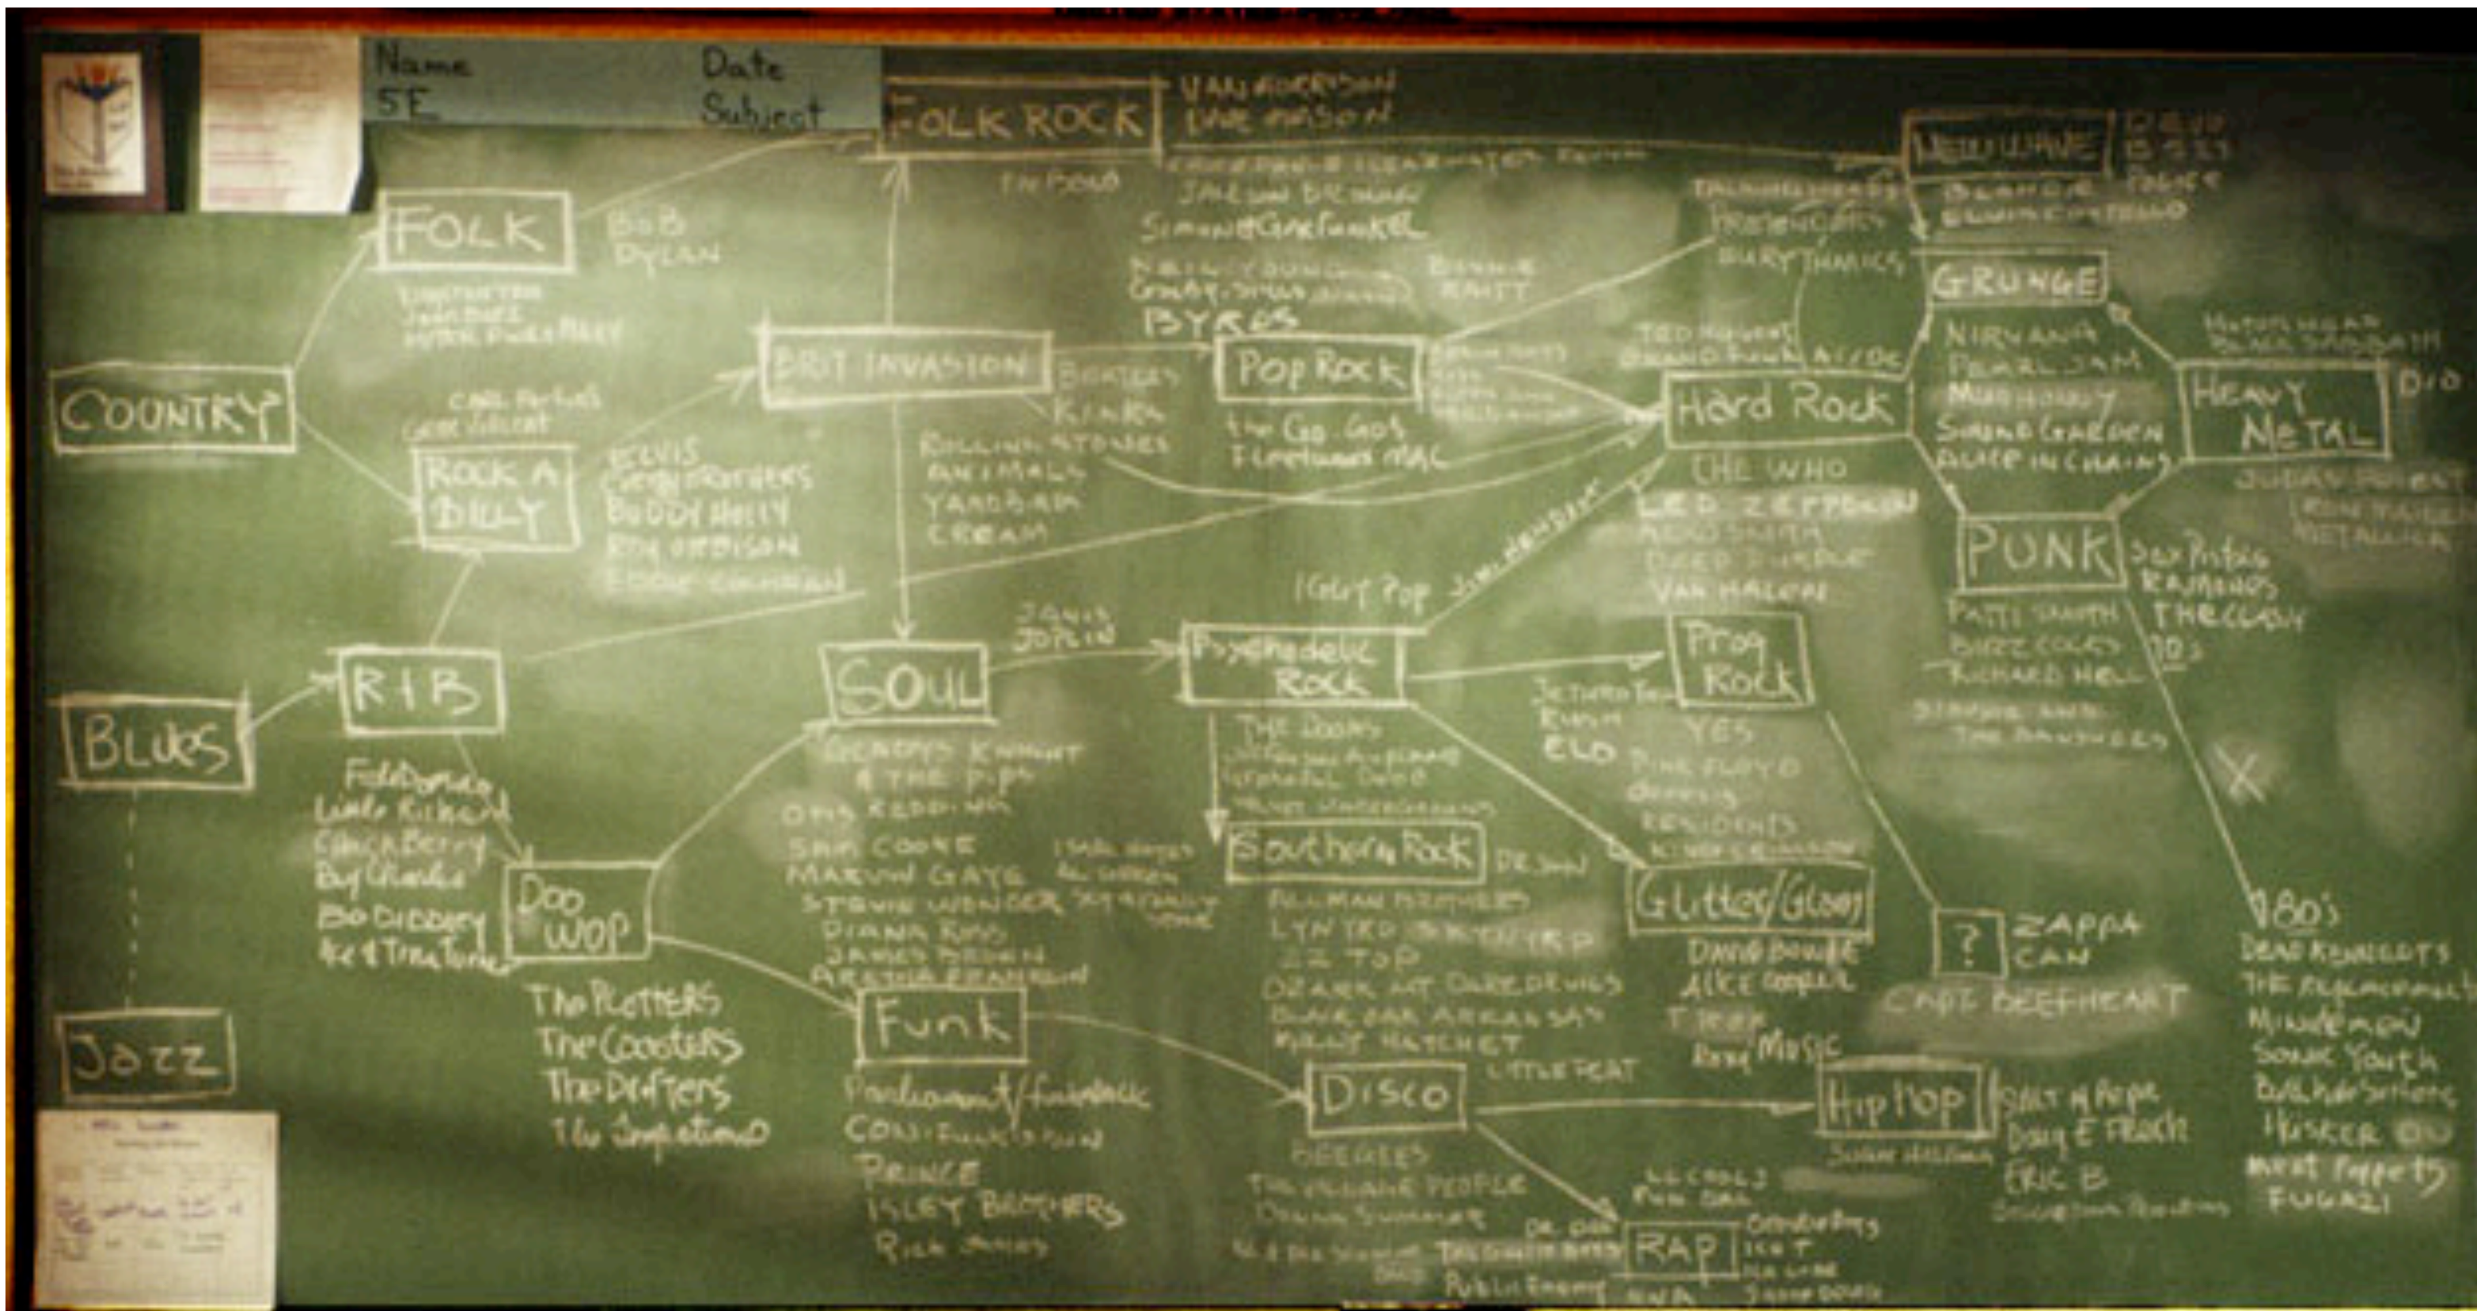

Name
Date
Subject
FOLK ROCK
FOLK
BOB DYLAN
COUNTRY
ROCK A BILLY
ELVIS
BRIT INVASION
BYRDS
Pop Rock
Hard Rock
NEW WAVE
GRUNGE
NIRVANA
PEARL JAM
HEAVY METAL
PUNK
THE WHO
LED ZEPPELIN
VAN HALEN
R+B
BLUES
SOUL
Psychedelic Rock
Prog Rock
YES
PINK FLOYD
Southern Rock
Glitter/Glam
DOO WOP
Funk
DISCO
Hip Hop
RAP
Jazz
ZAPPA
CAN
CAPT BEEFHEART
80's
PRINCE

musicmap

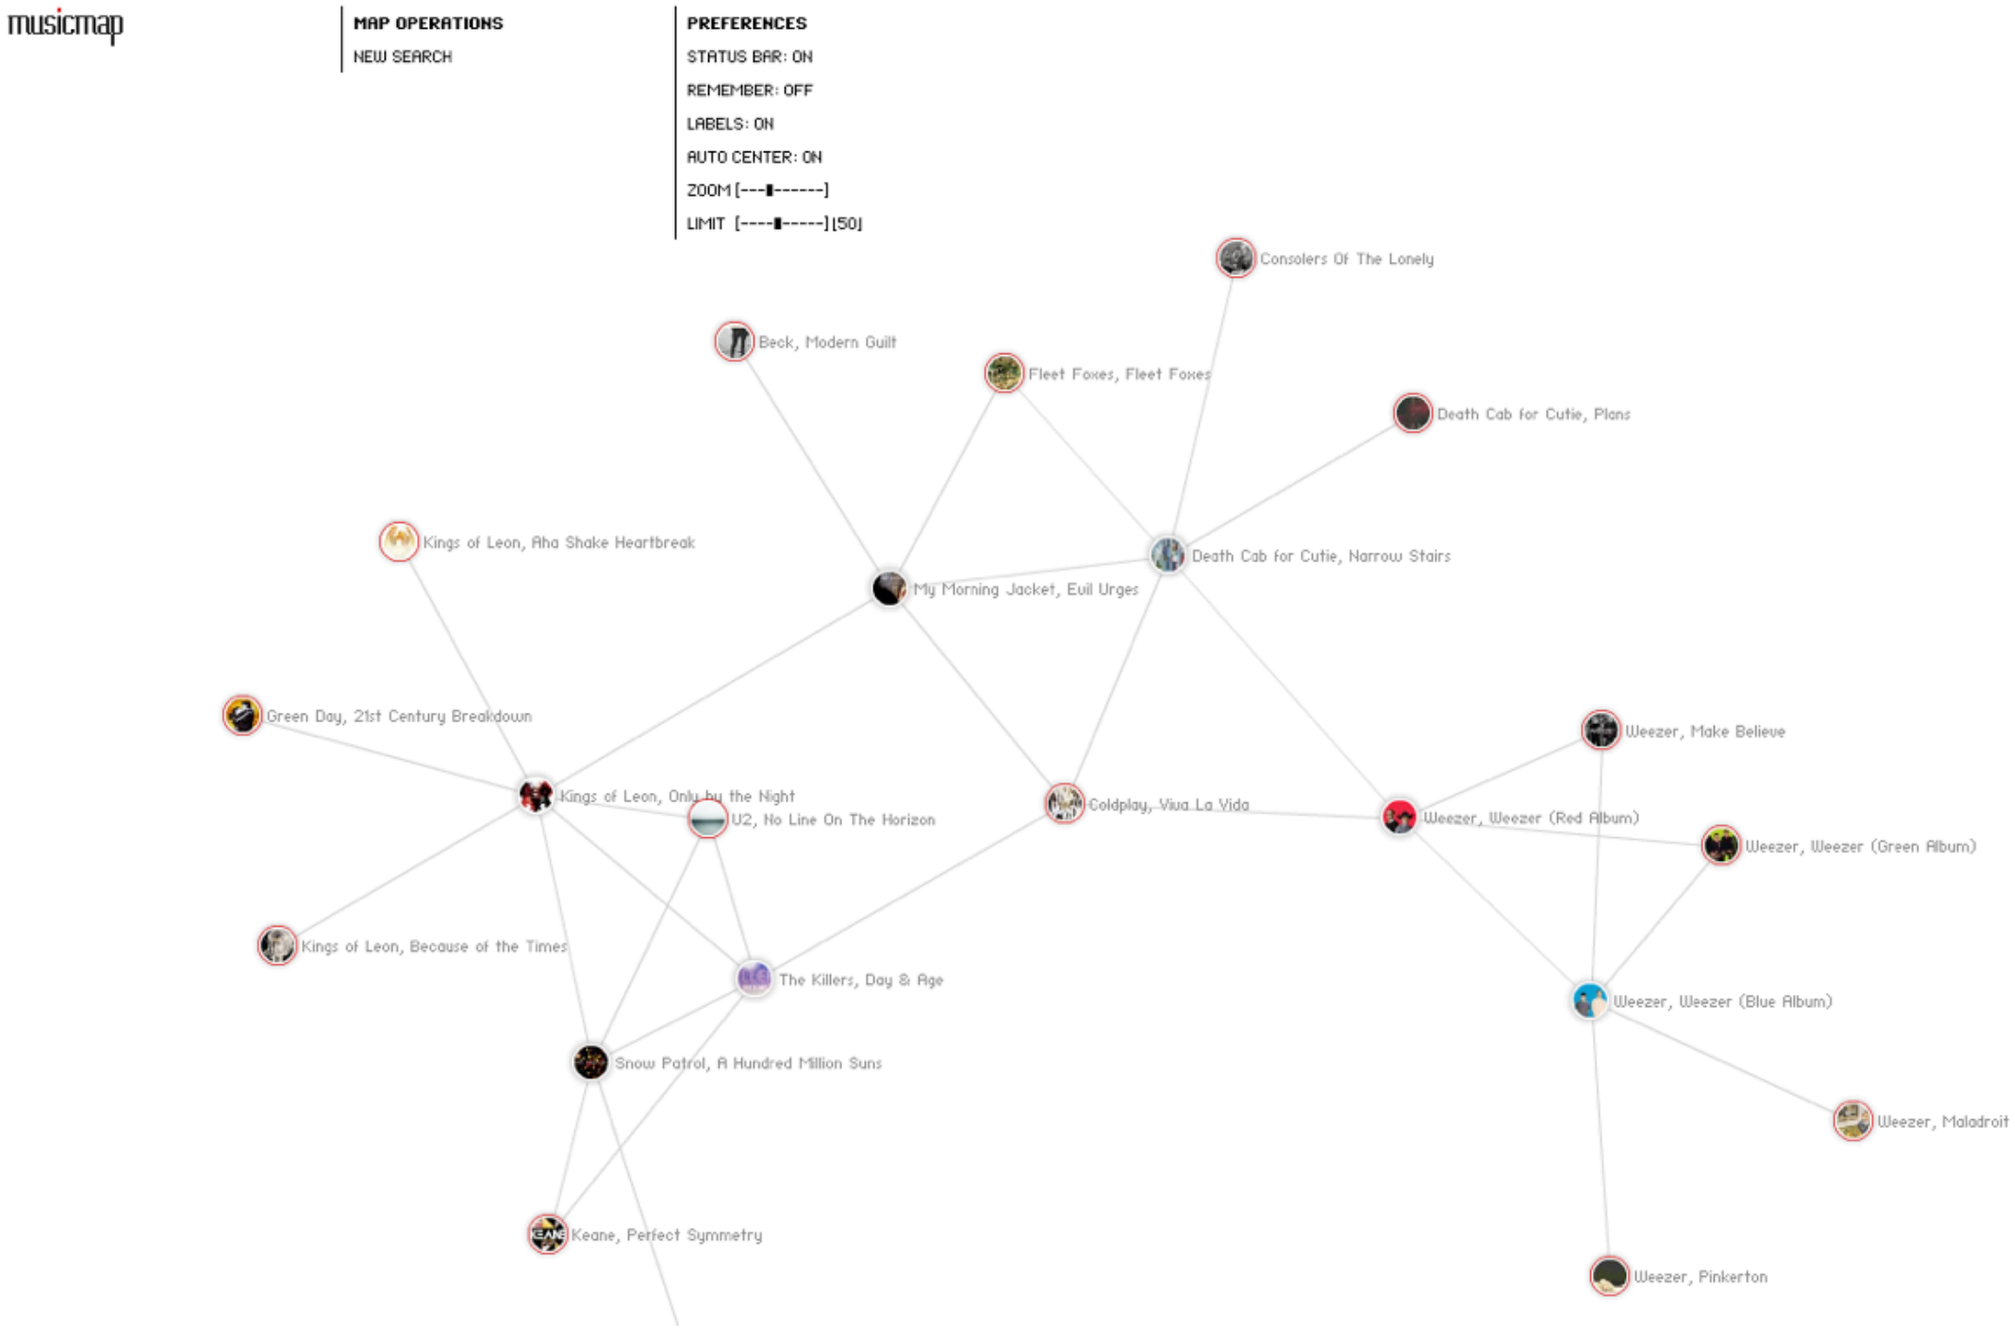

http://www.dimvision.com/musicmap/

Music Plasma

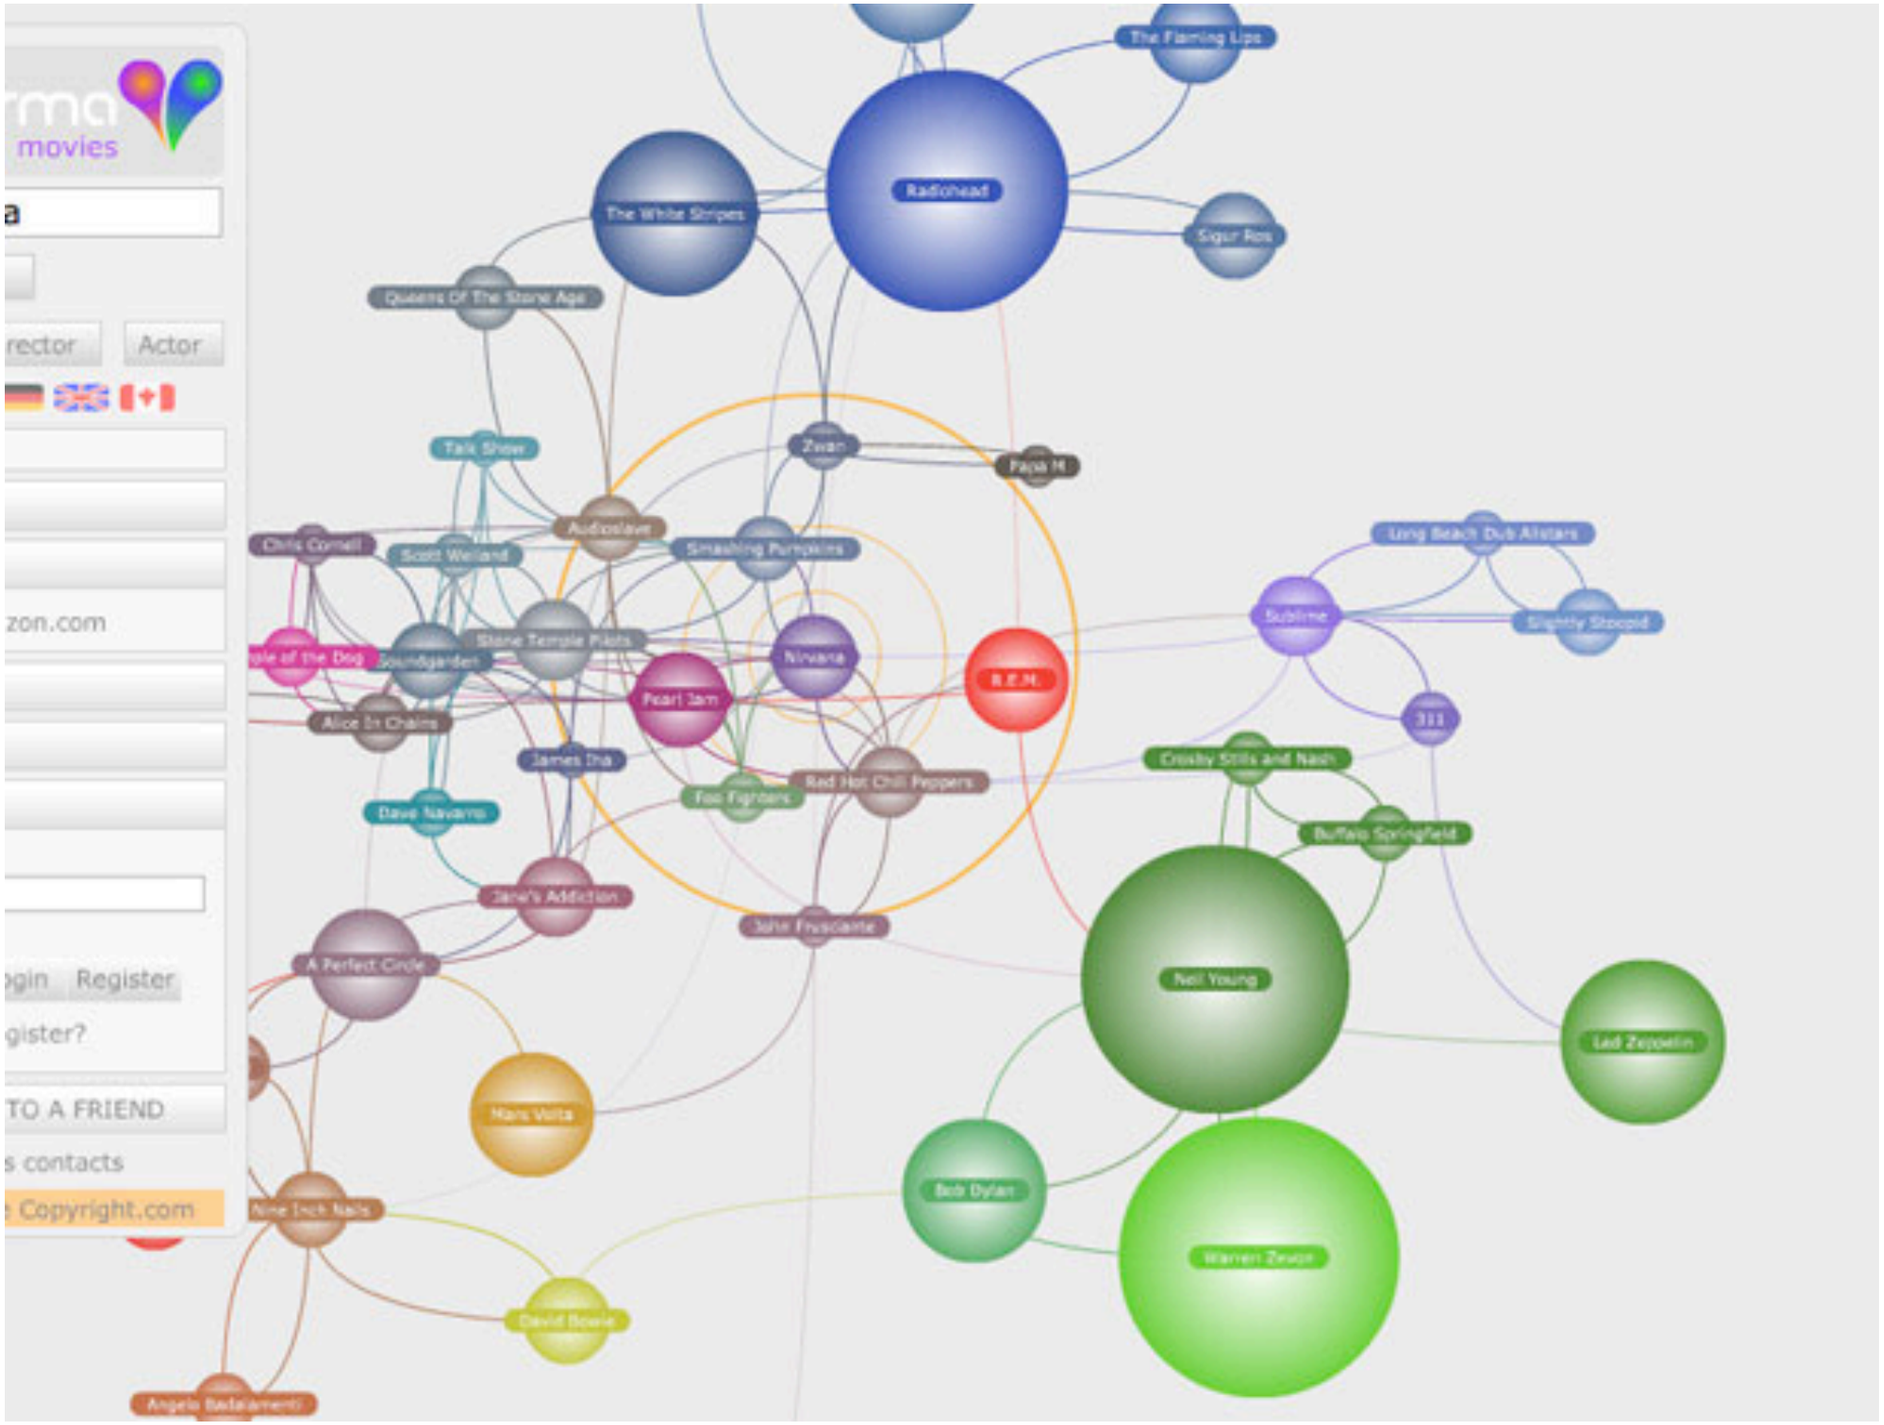

MusicPlasma.com

tuneglue

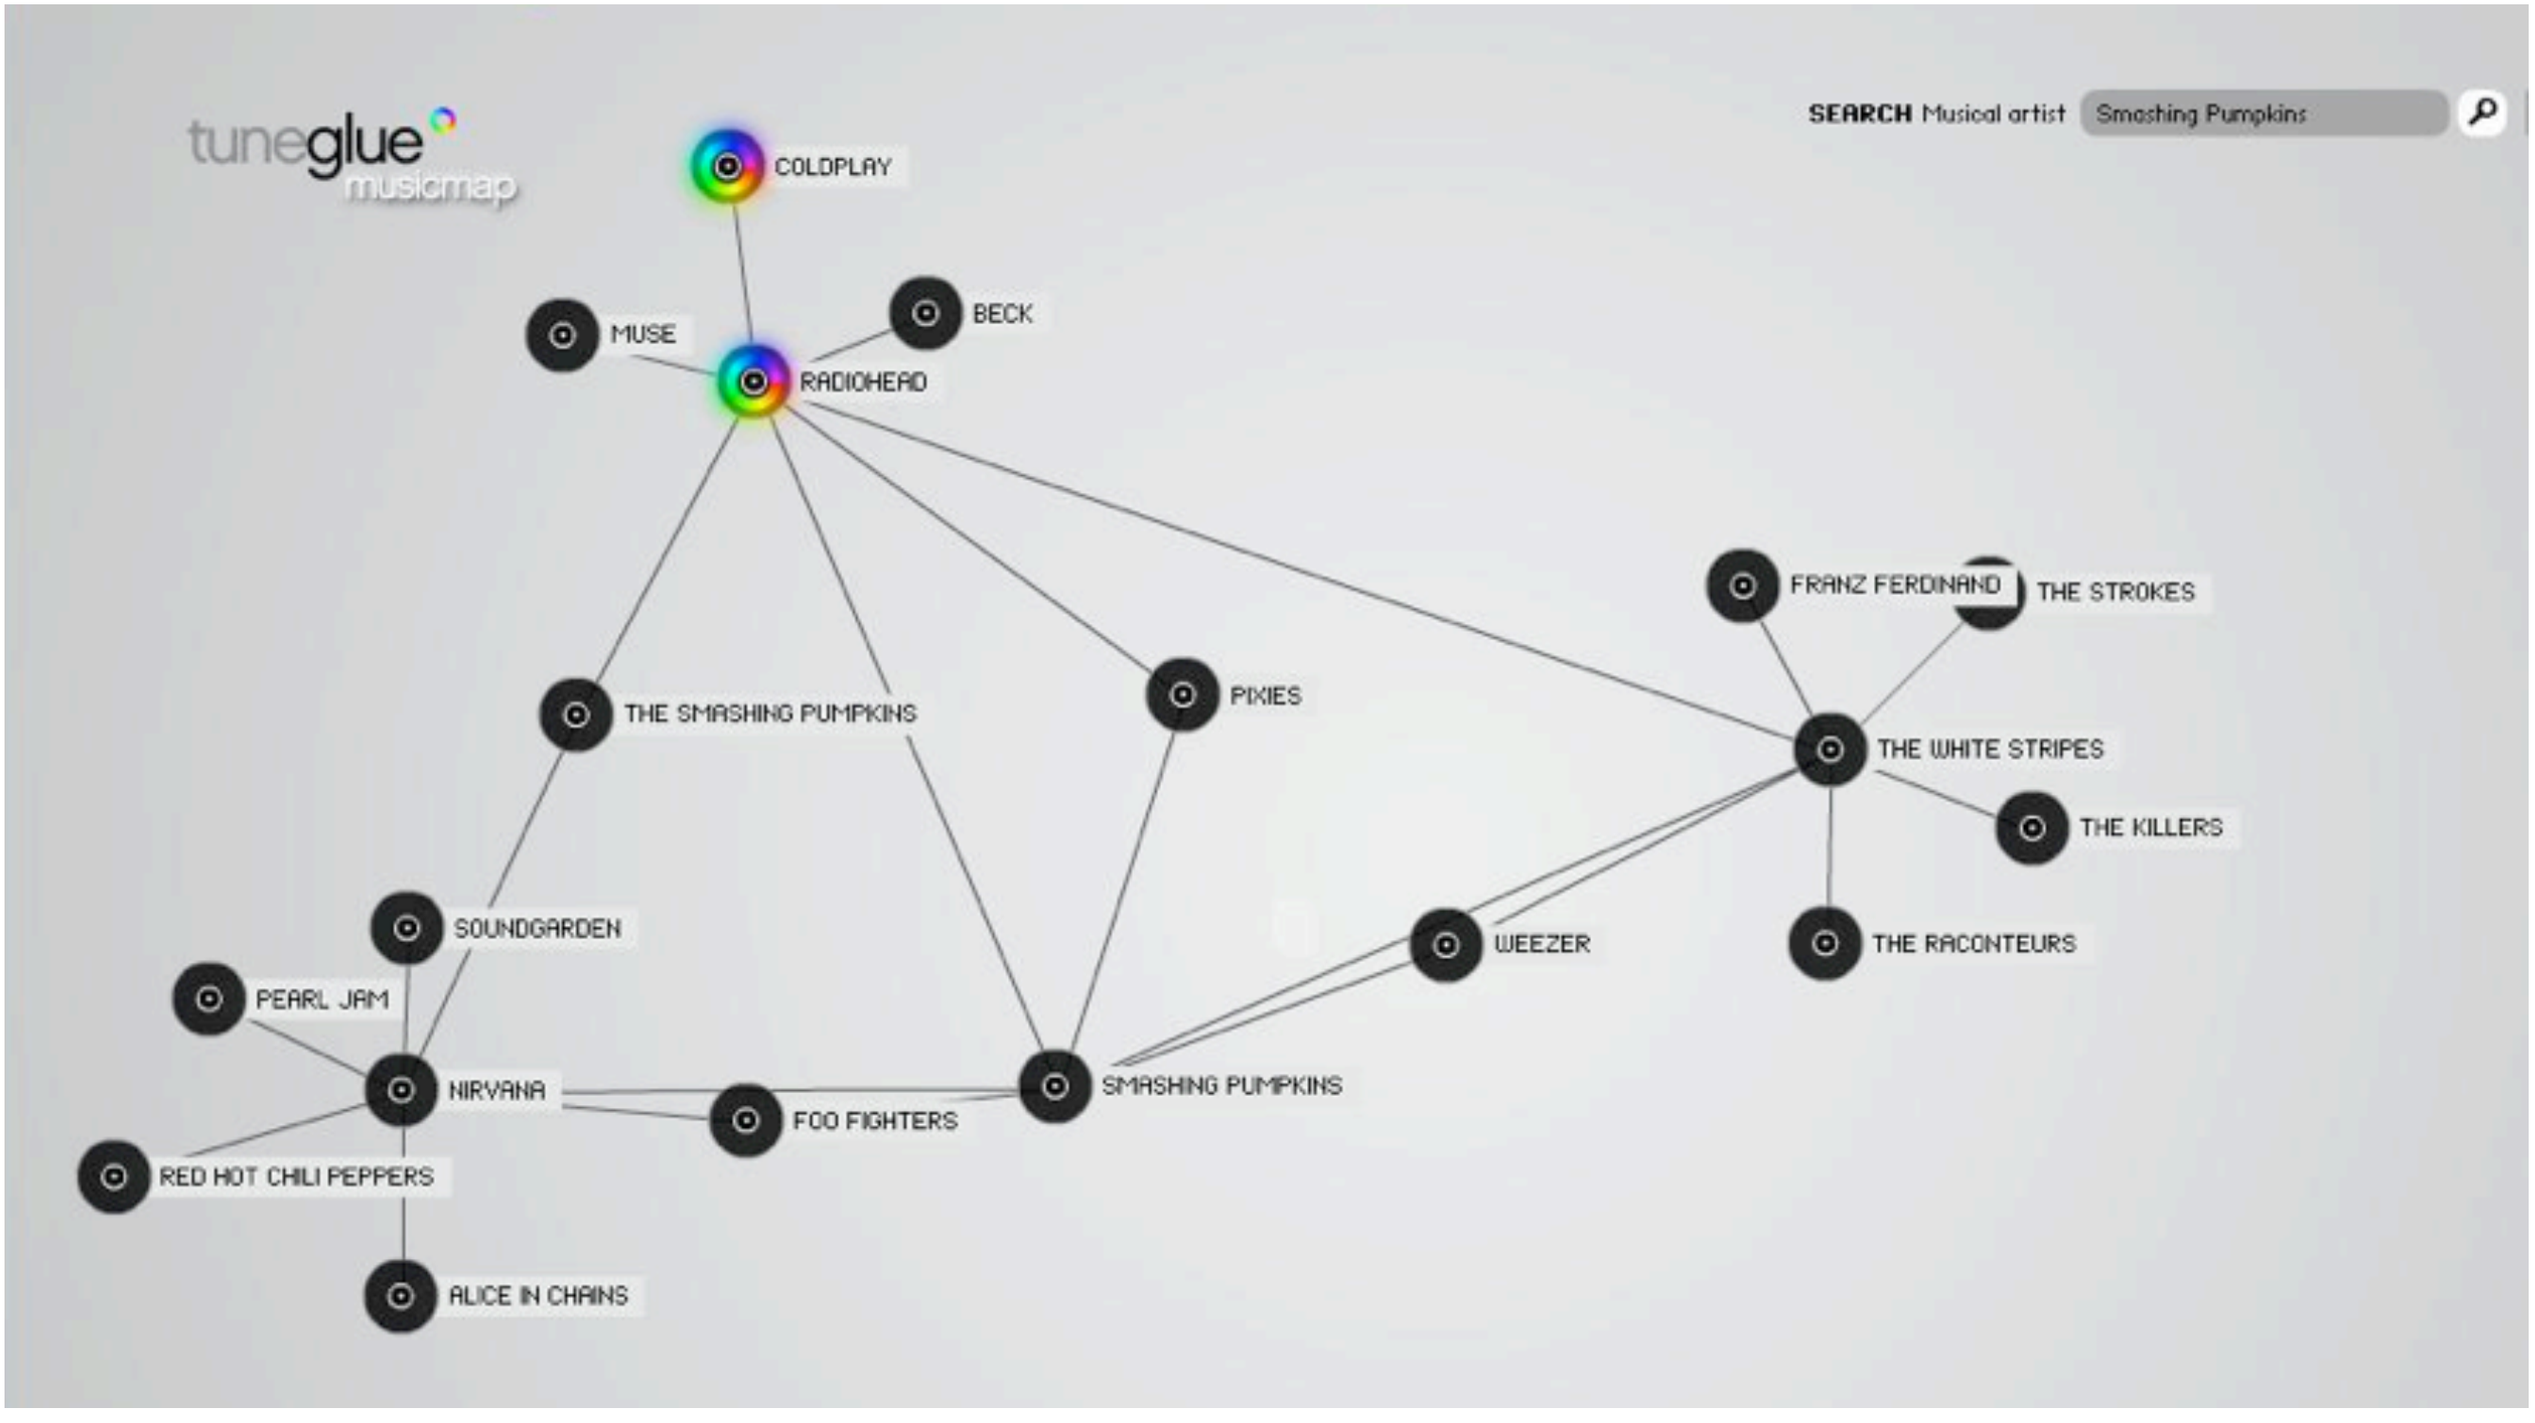

http://audiomap.tuneglue.net/

Last.fm artist map

http://lastfm.dontdrinkandroot.net/

Last.fm artist map

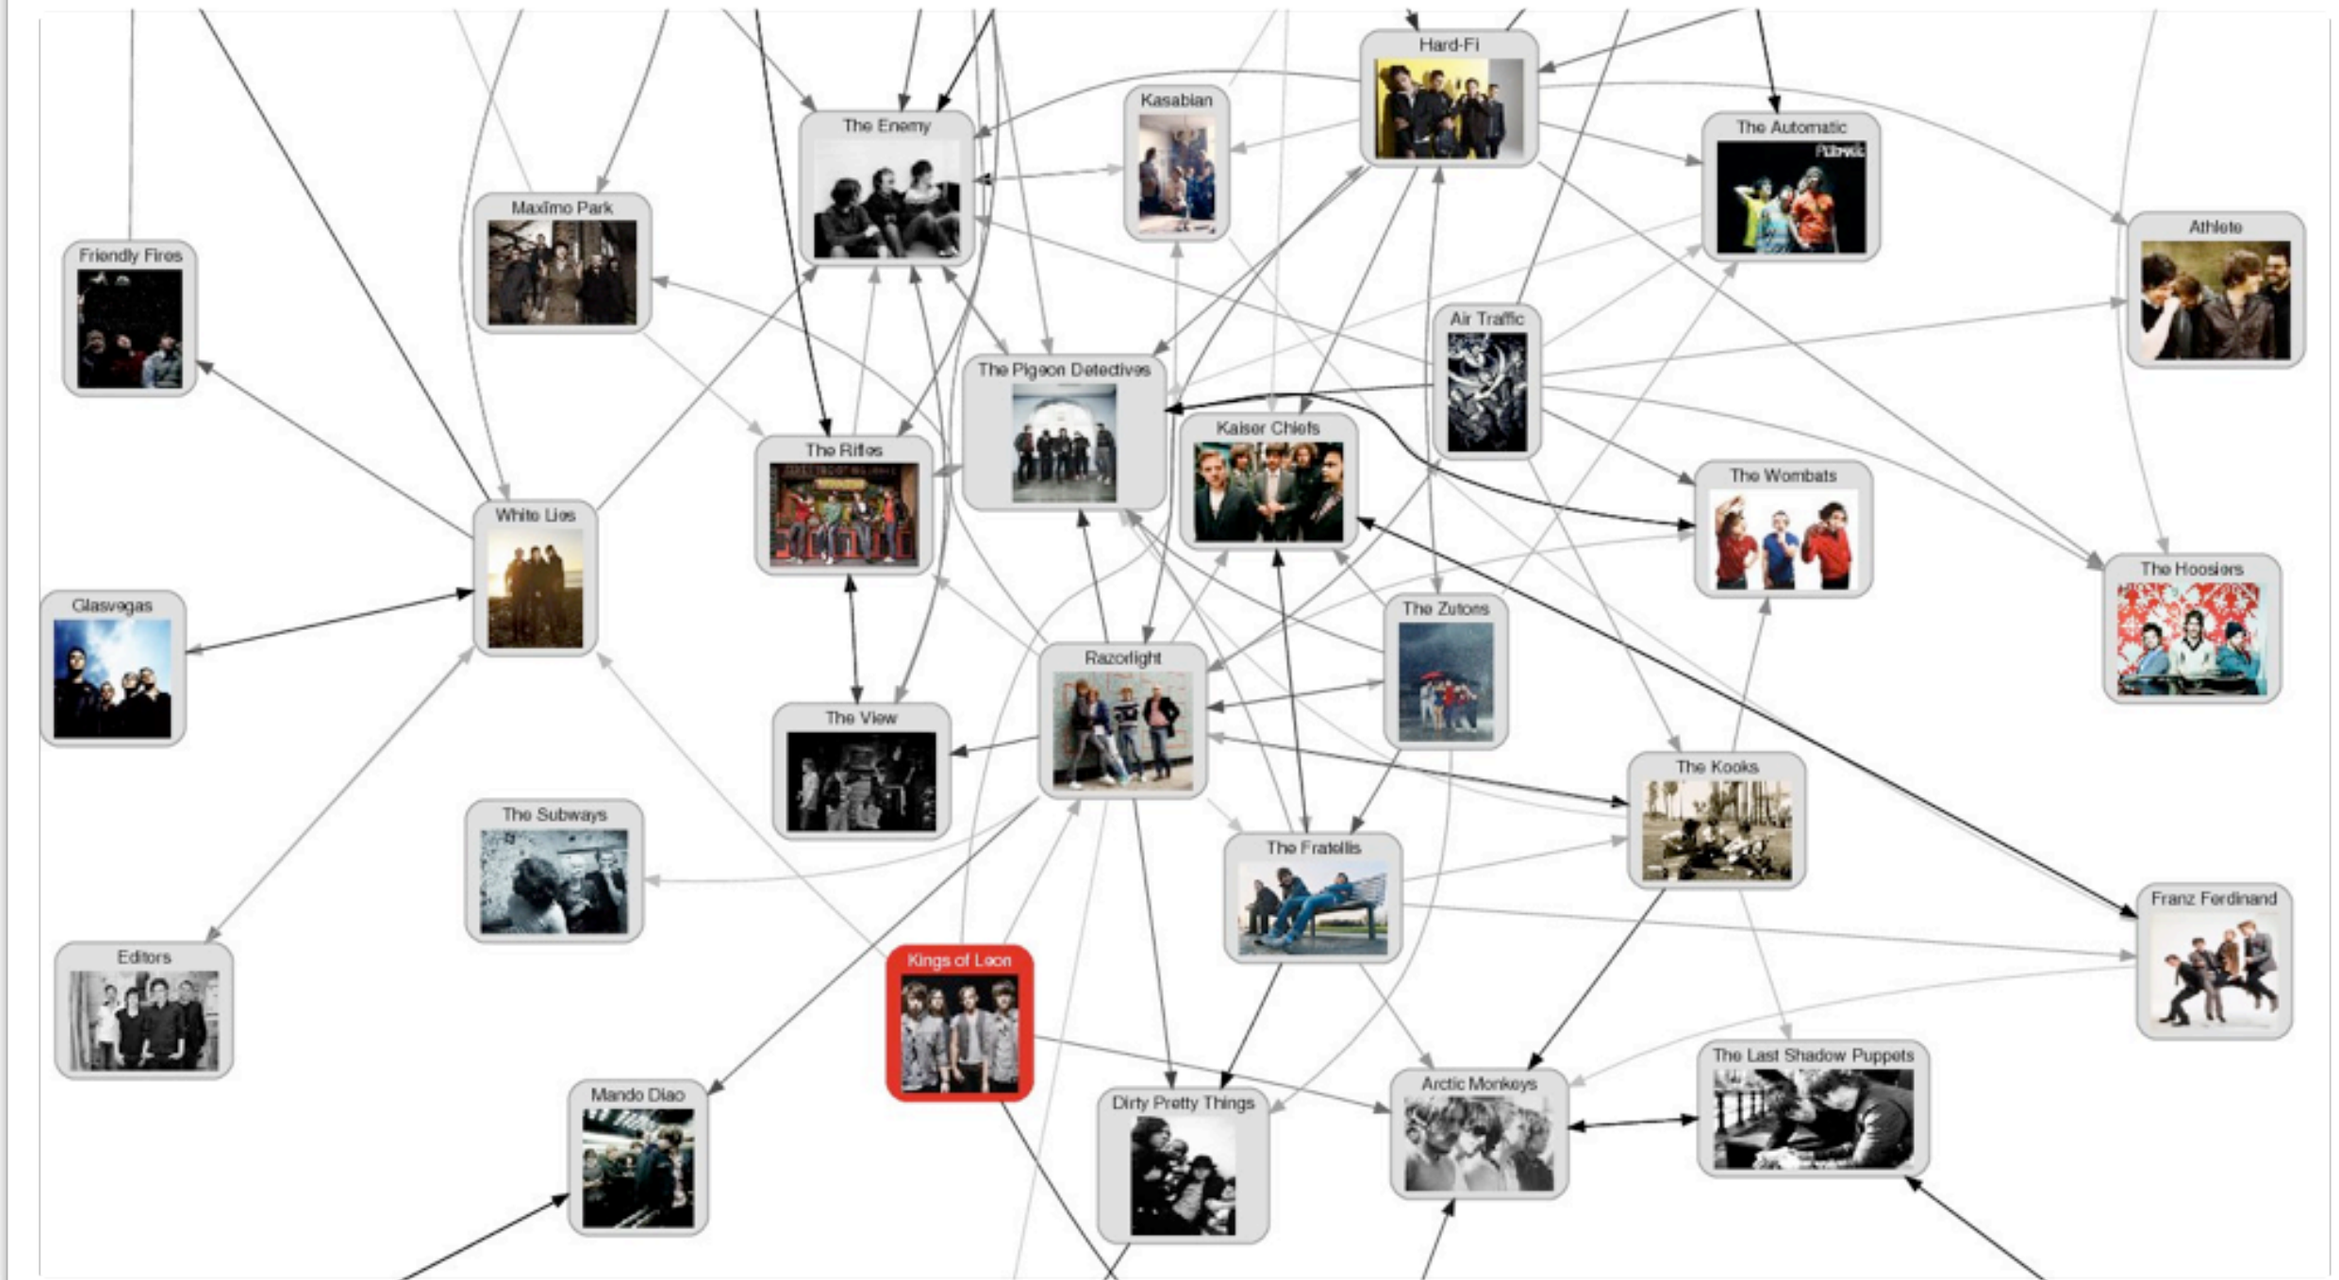

http://lastfm.dontdrinkandroot.net/

RAMA - Relational Artist Maps

Late-Breaking Demo Session

http://rama.inescporto.pt/

gnod music-map

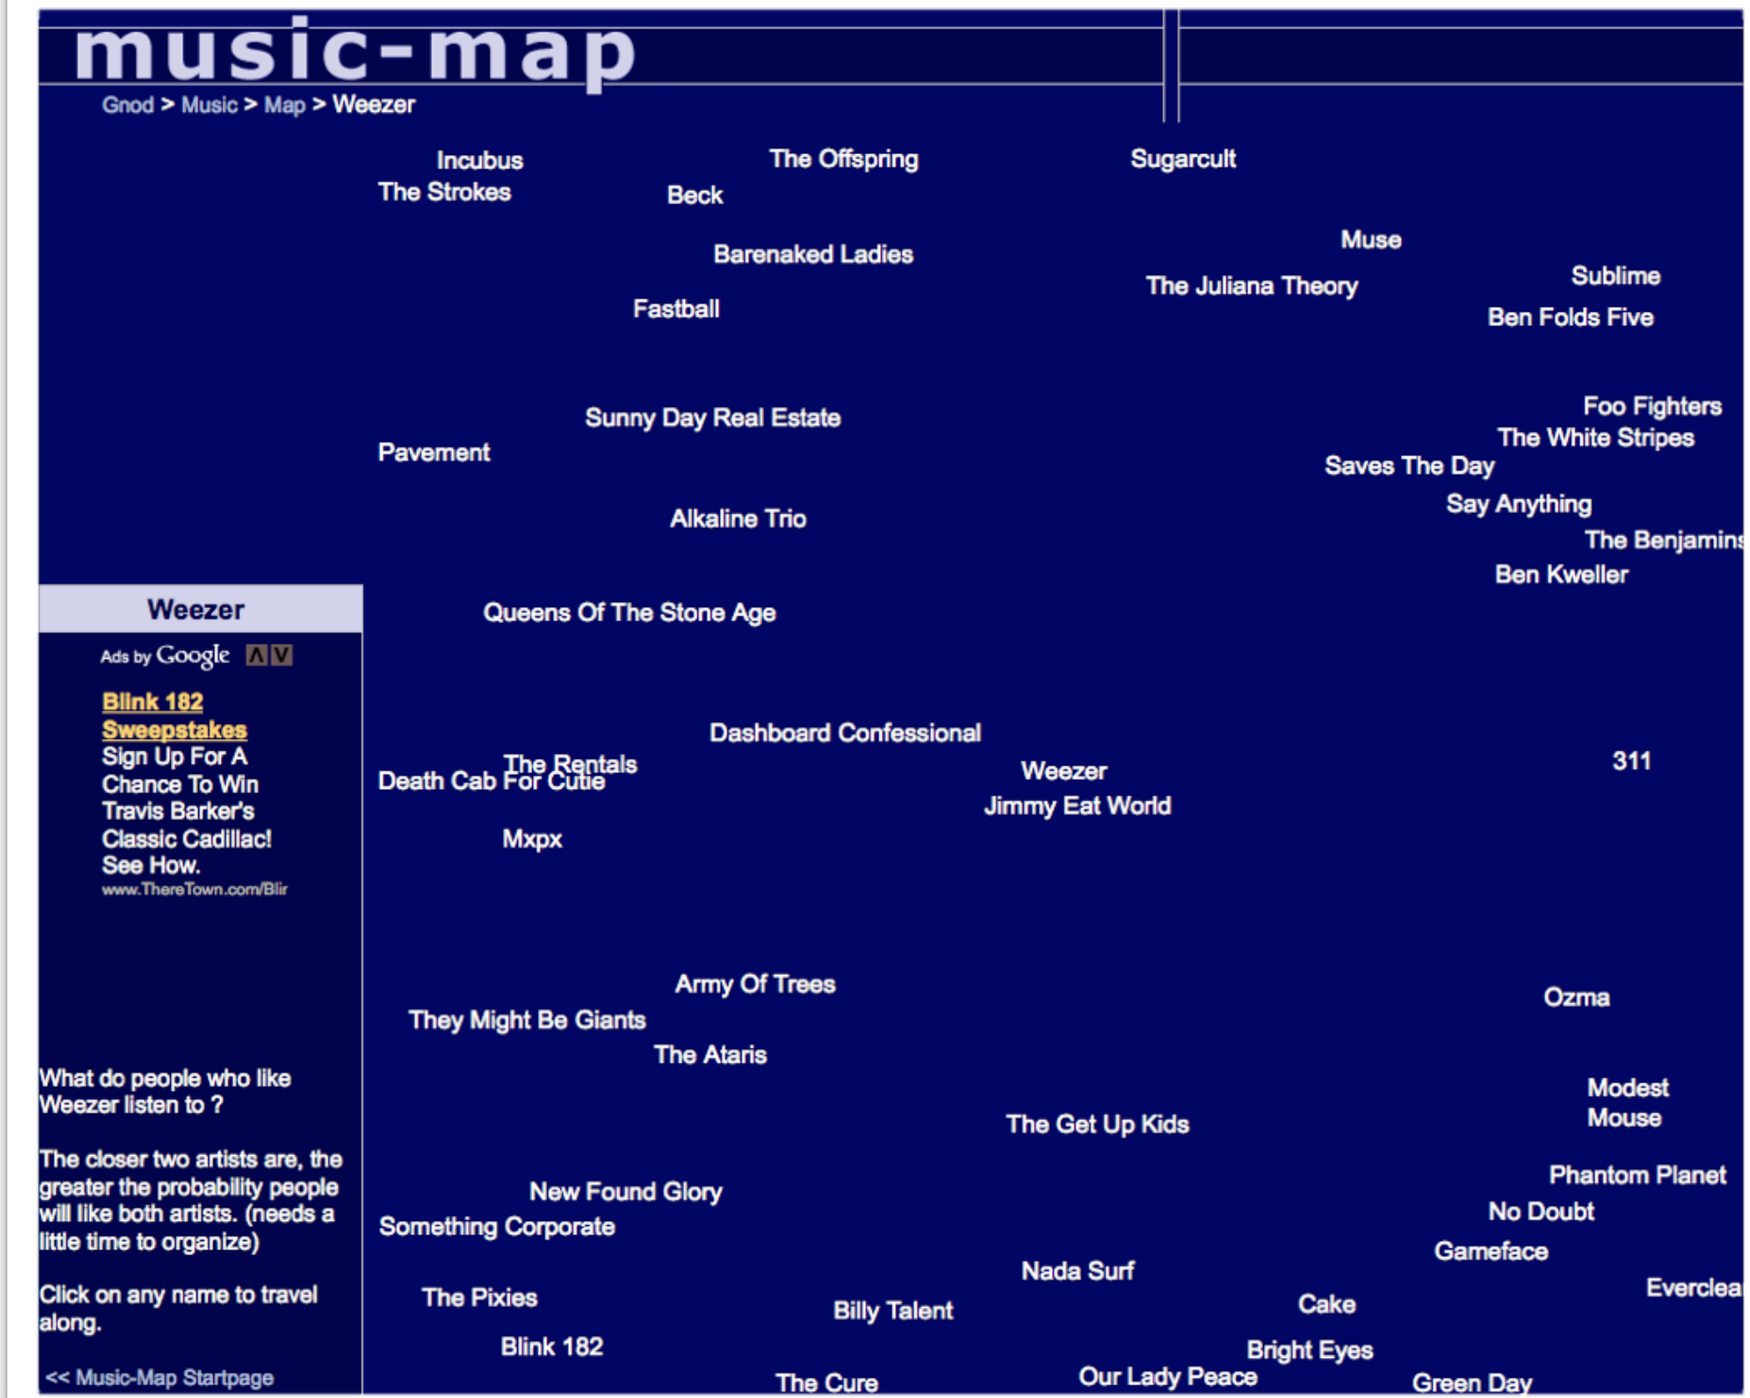

http://www.music-map.com/

musicmesh

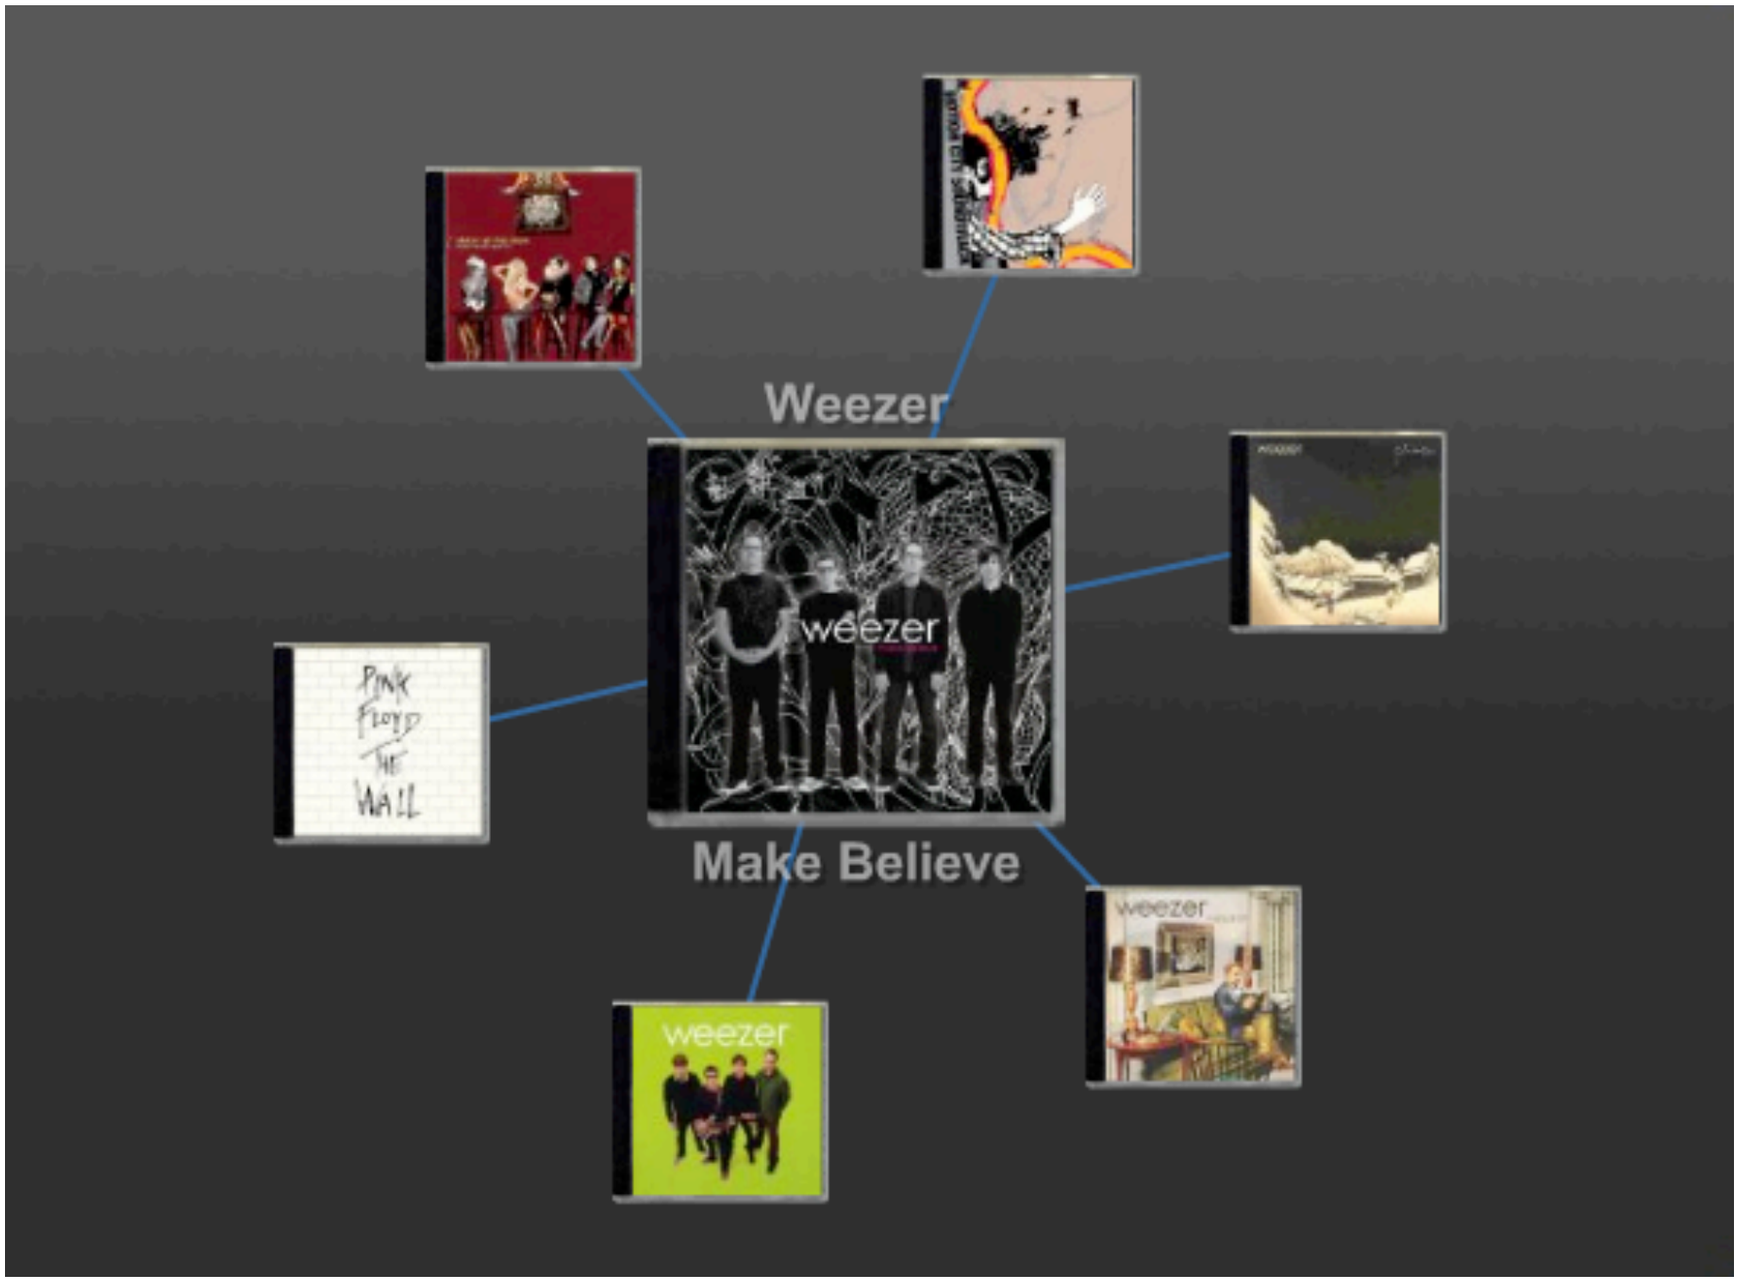

http://www.musicmesh.net/

musicovery

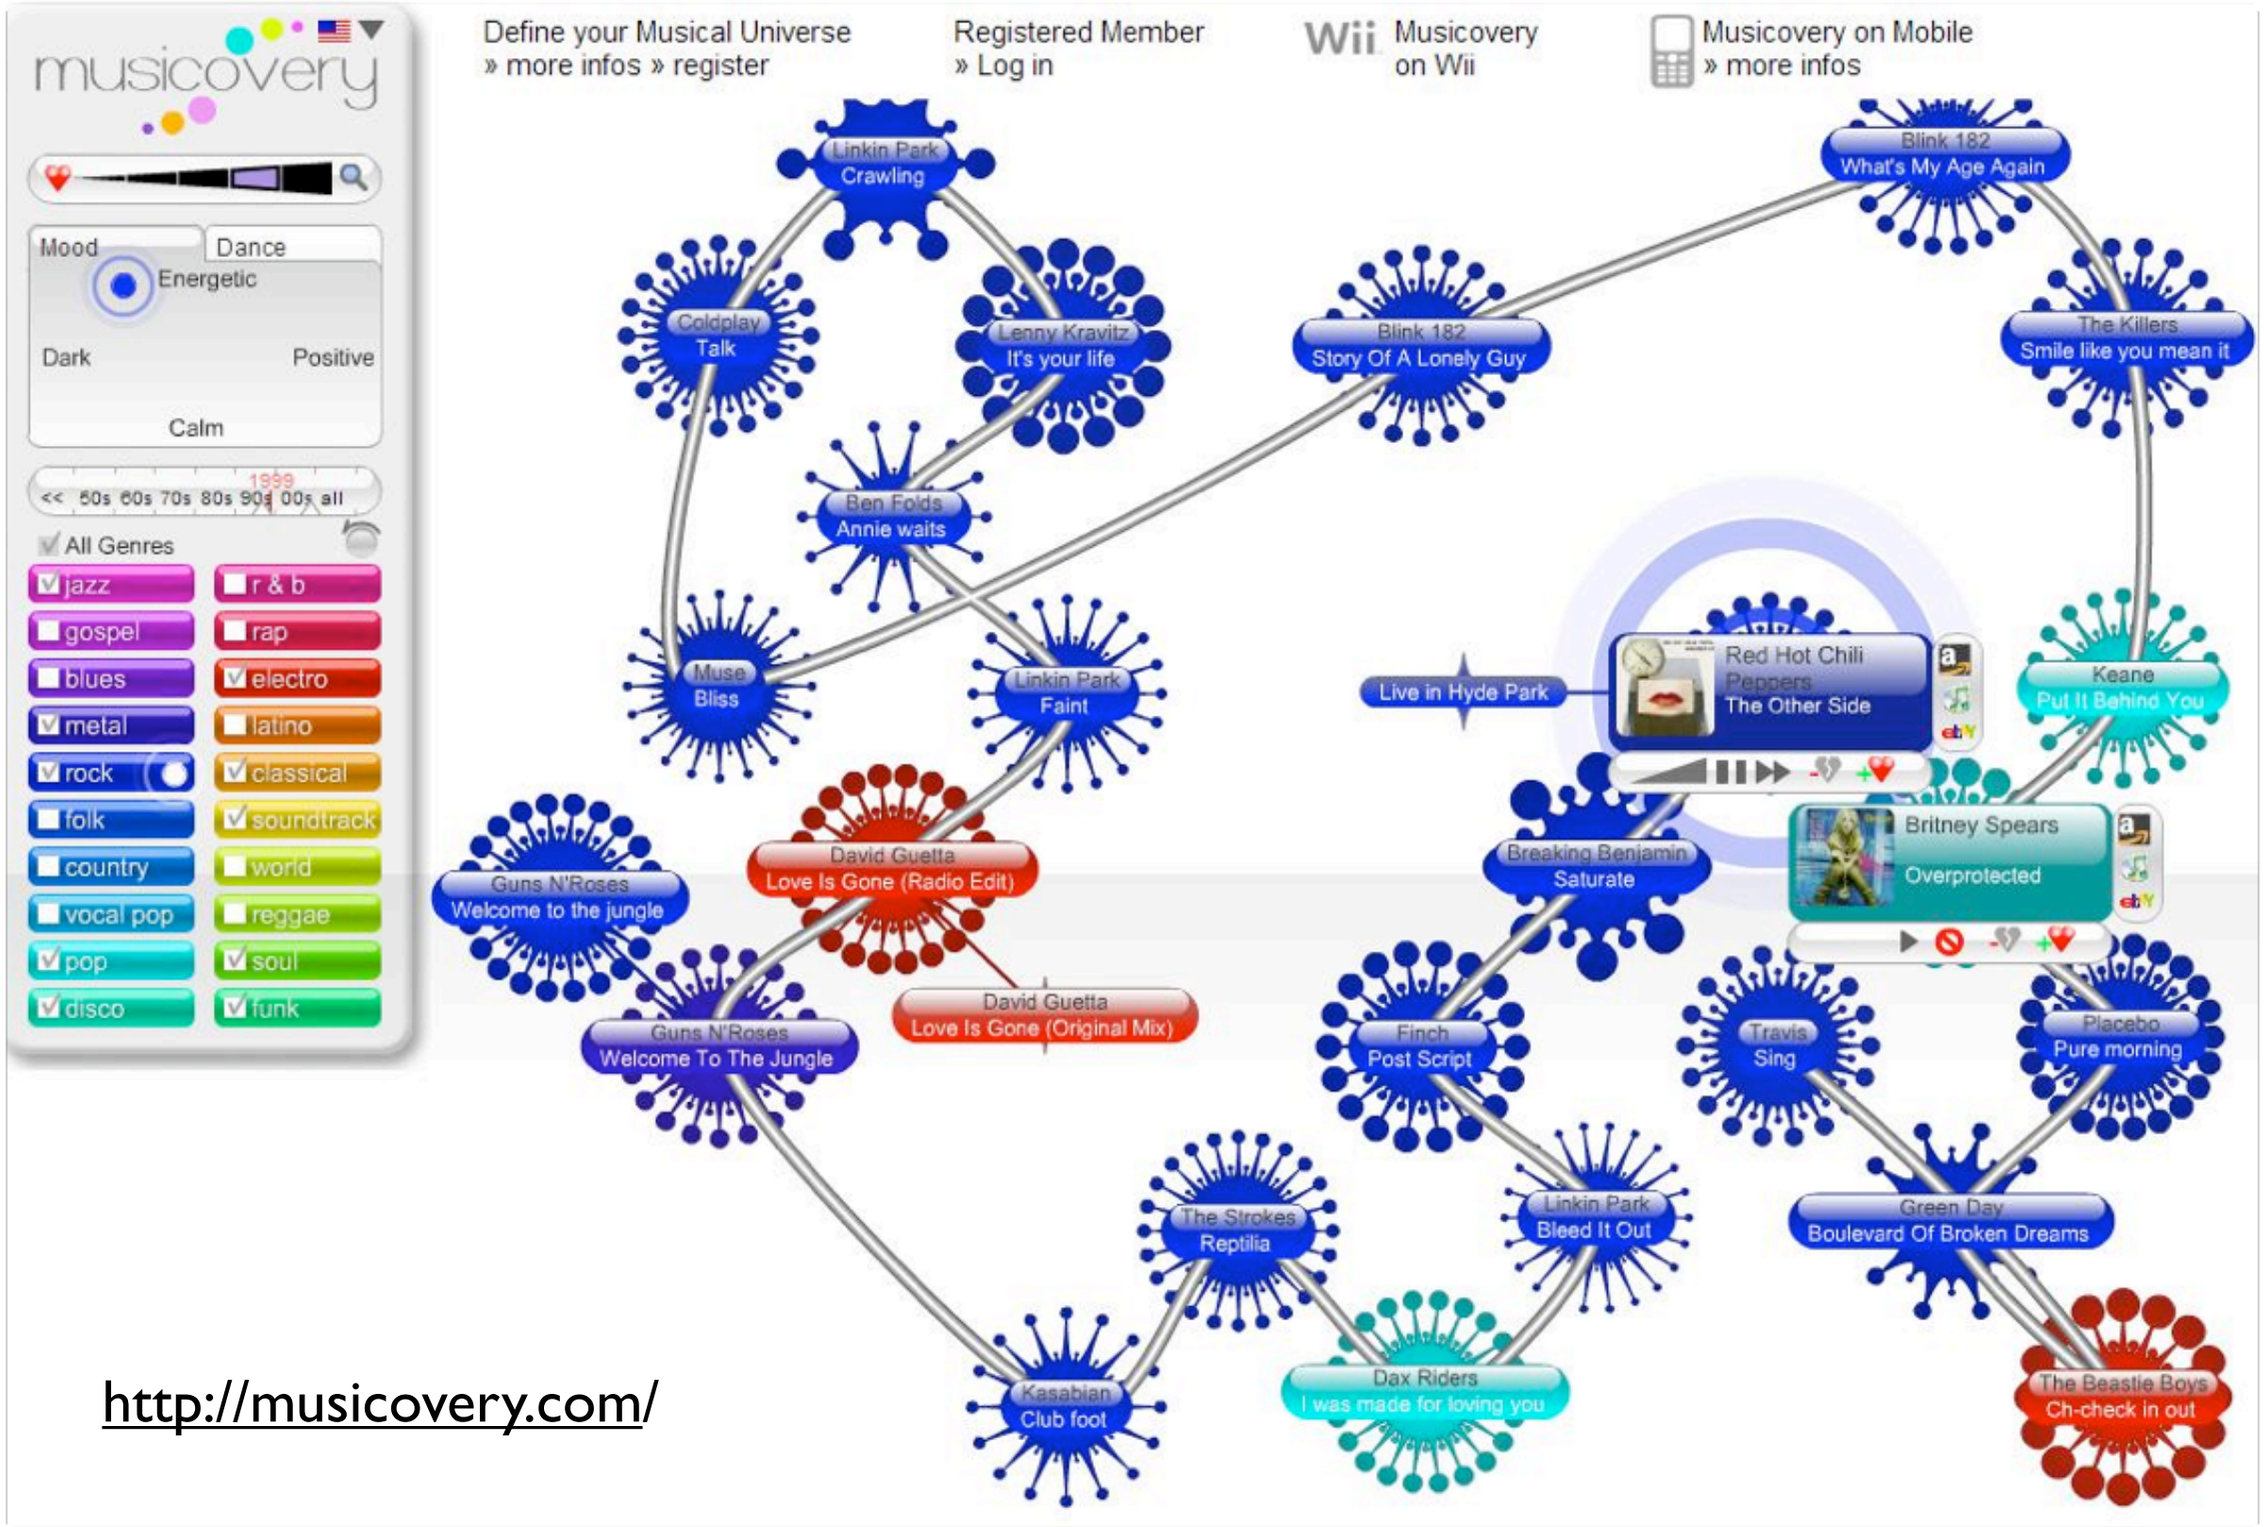

http://musicovery.com/

SoundBite for SongBird

http://www.repeatingbeats.com/thesis.php - Steve Lloyd - QMUL

SoundBite for SongBird

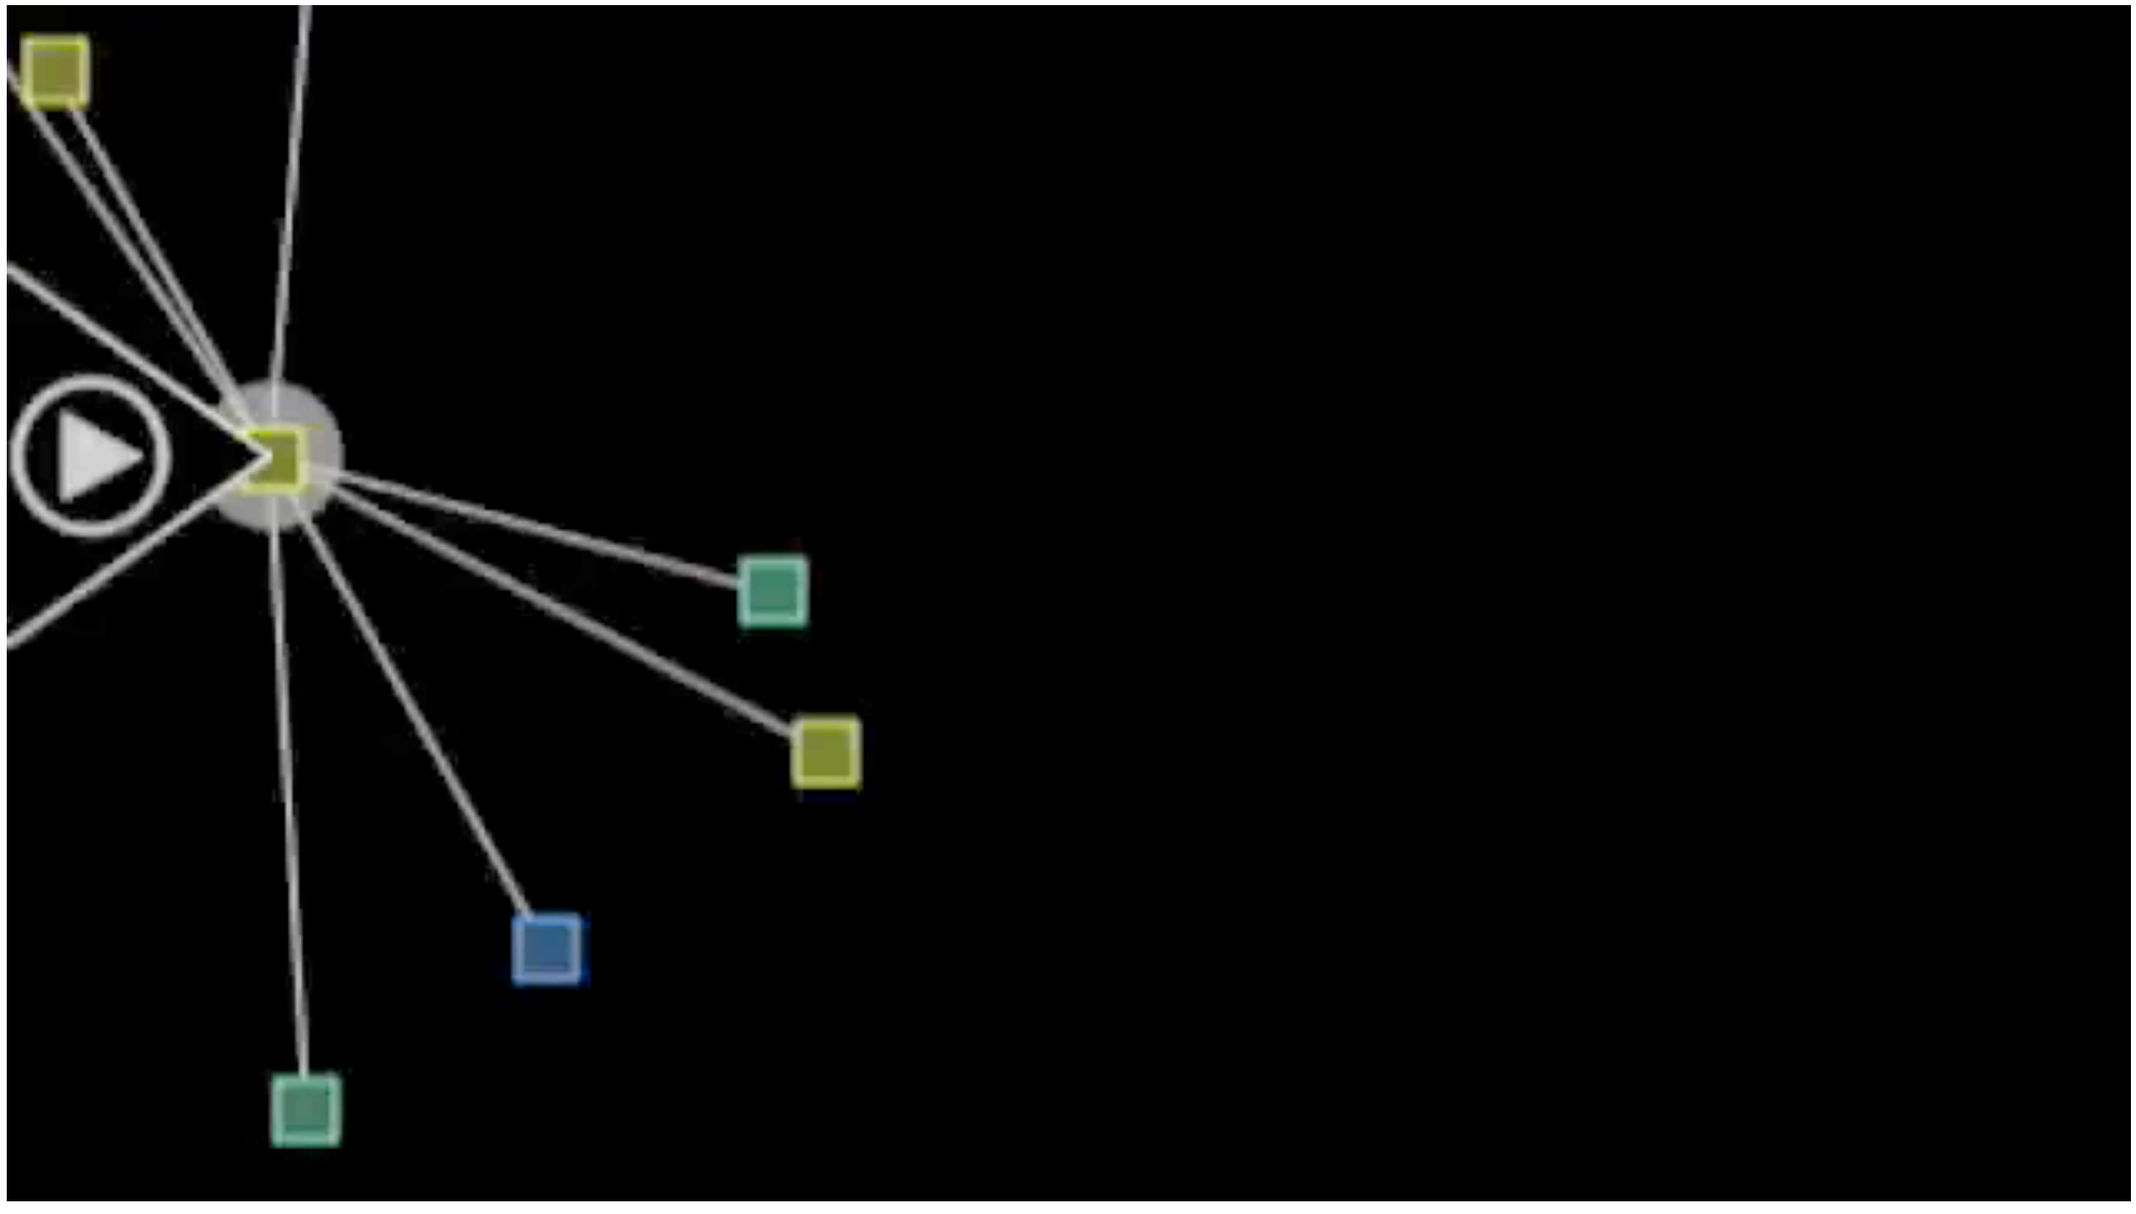

http://www.repeatingbeats.com/thesis.php - Steve Lloyd - QMUL

Amaznode

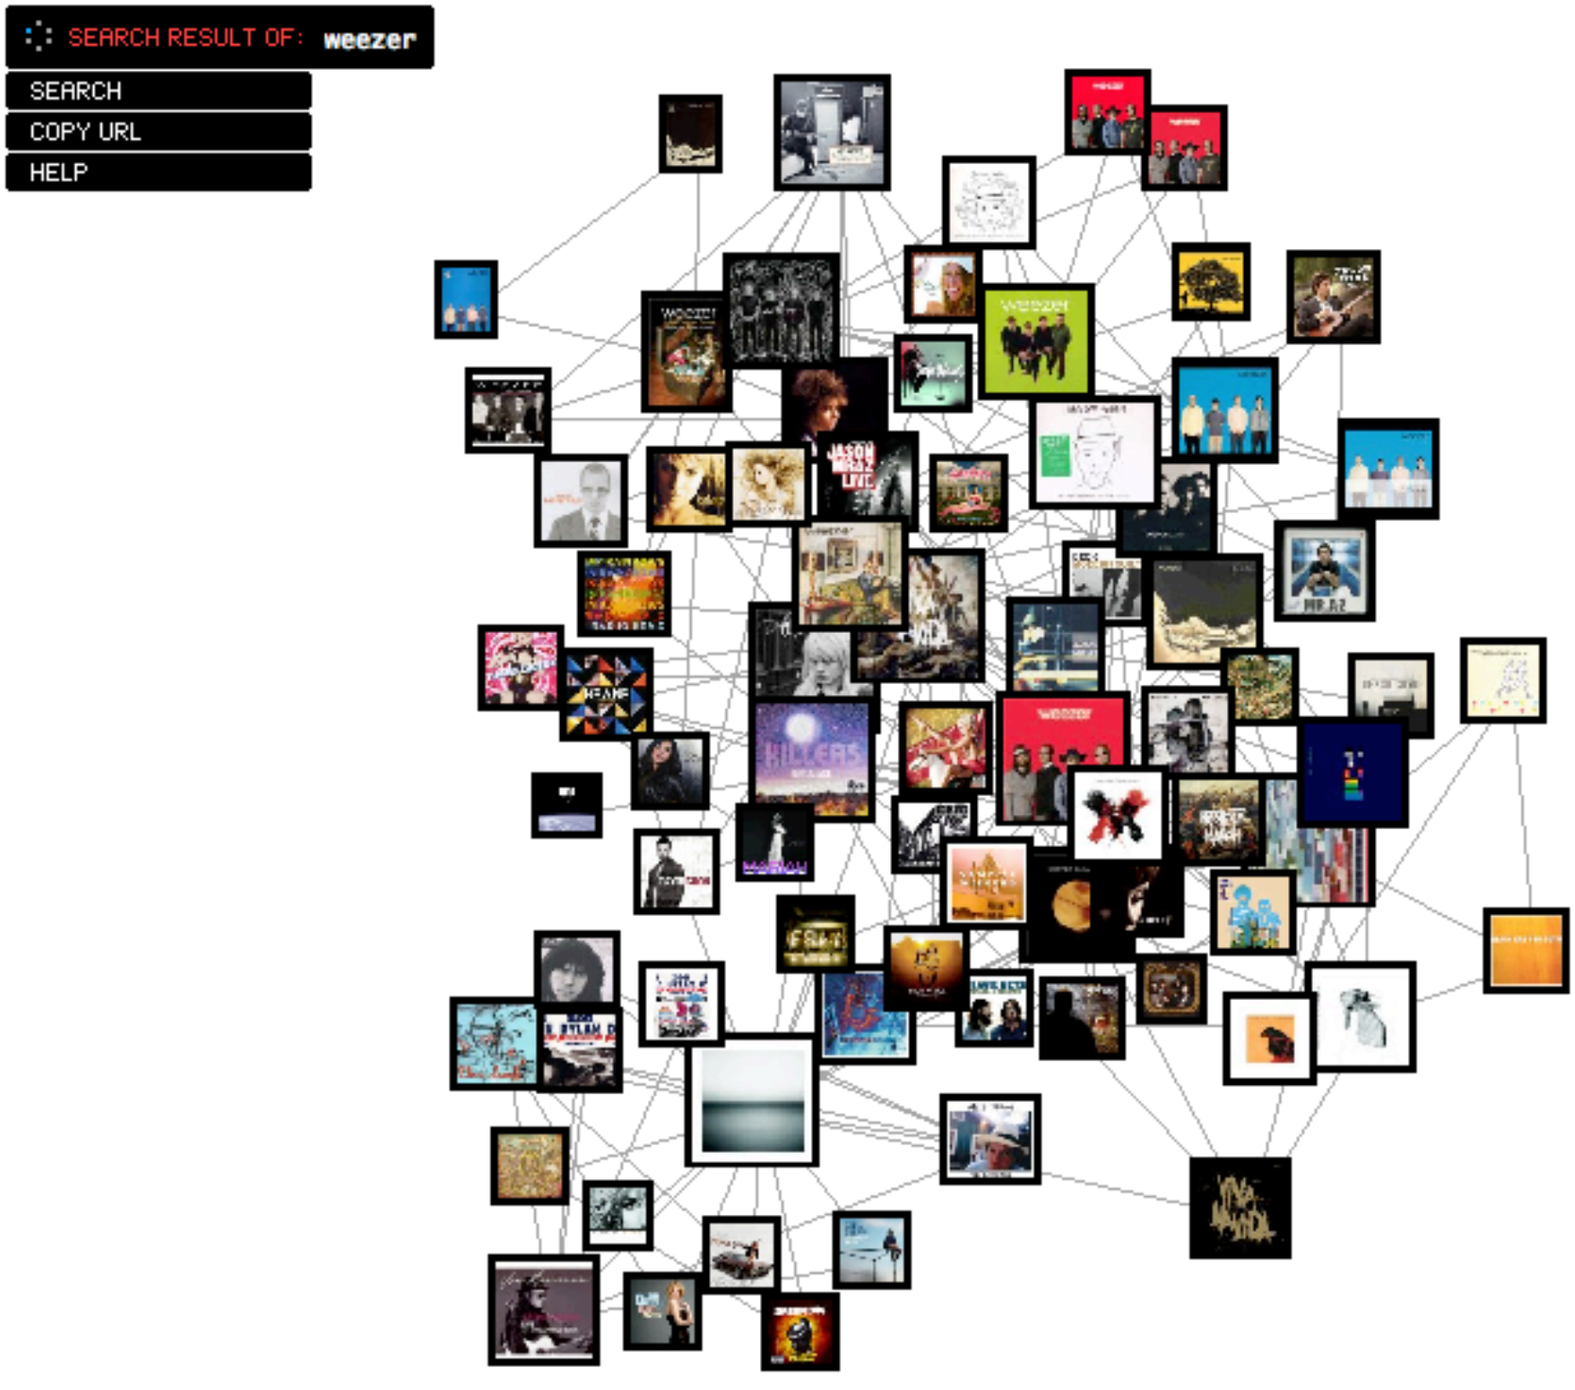

http://amaznode.fladdict.net/

Visualizing and Exploring Personal Libraries

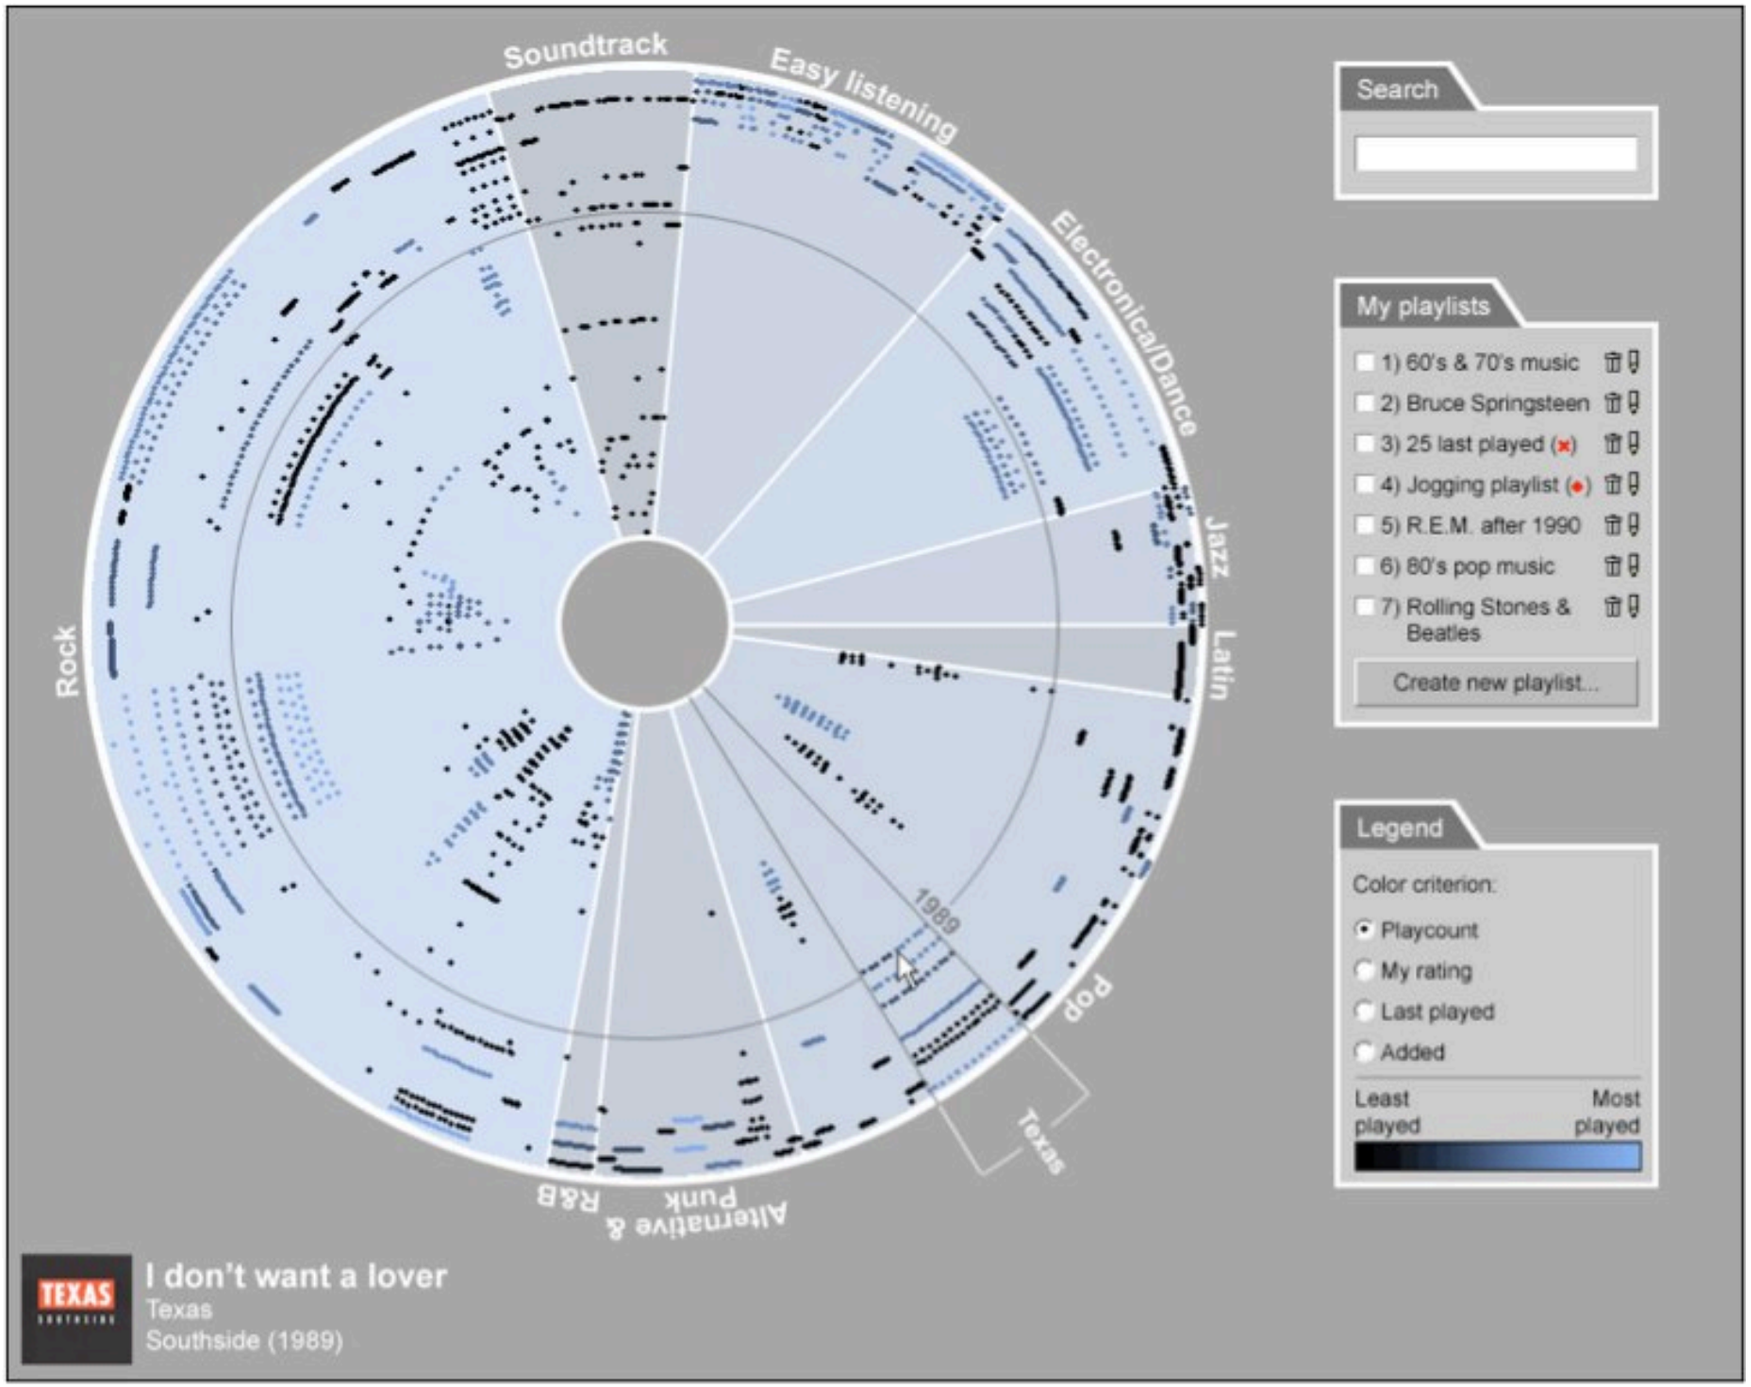

http://torrens.files.wordpress.com/2009/06/ismir.pdf
Torrens, Patrick Hertzog, Josep-Lluís Arcos

Visualizing and Exploring Personal Libraries

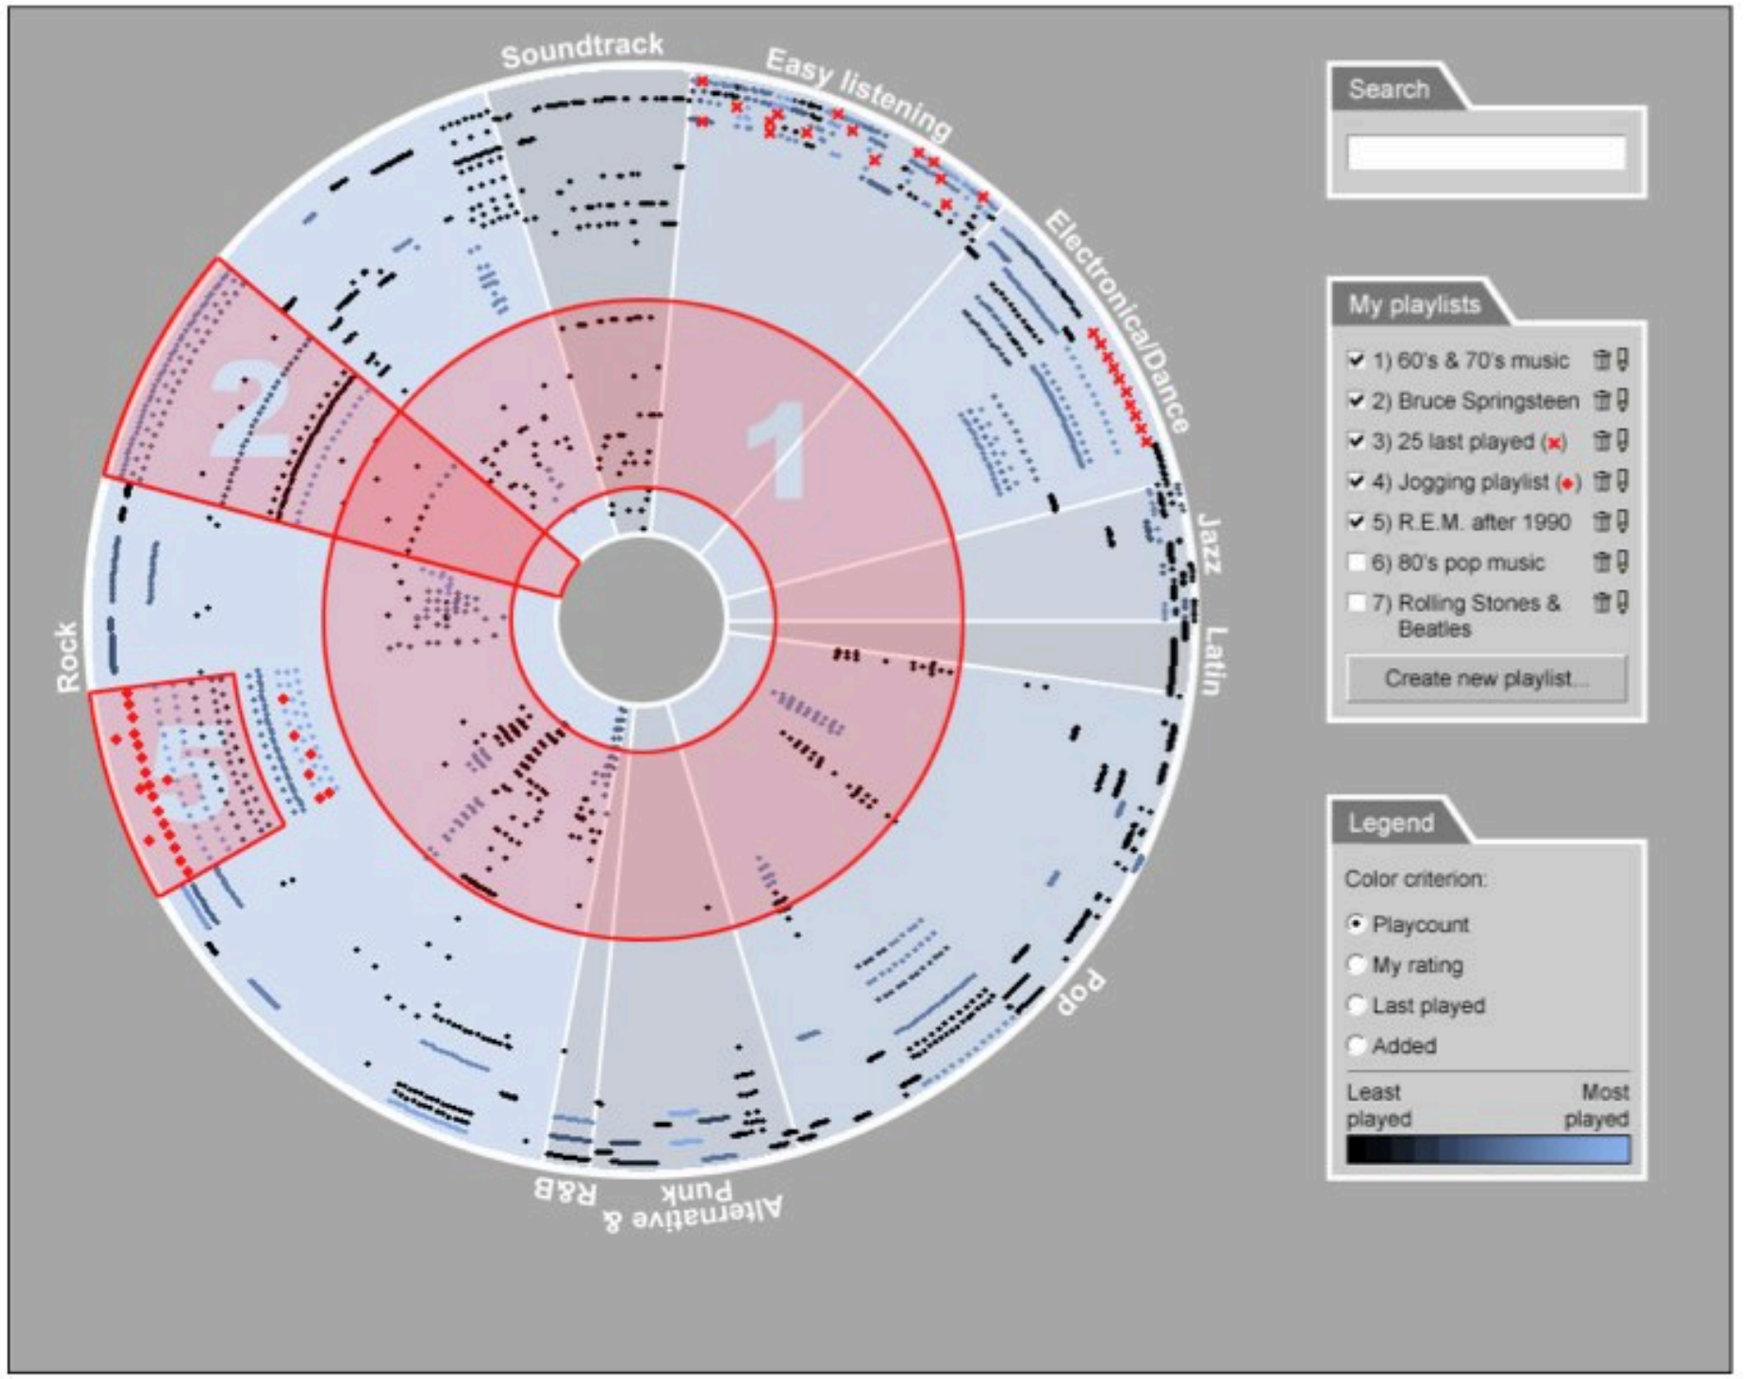

http://torrens.files.wordpress.com/2009/06/ismir.pdf
Torrens, Patrick Hertzog, Josep-Lluís Arcos

Visualizing and Exploring Personal Libraries

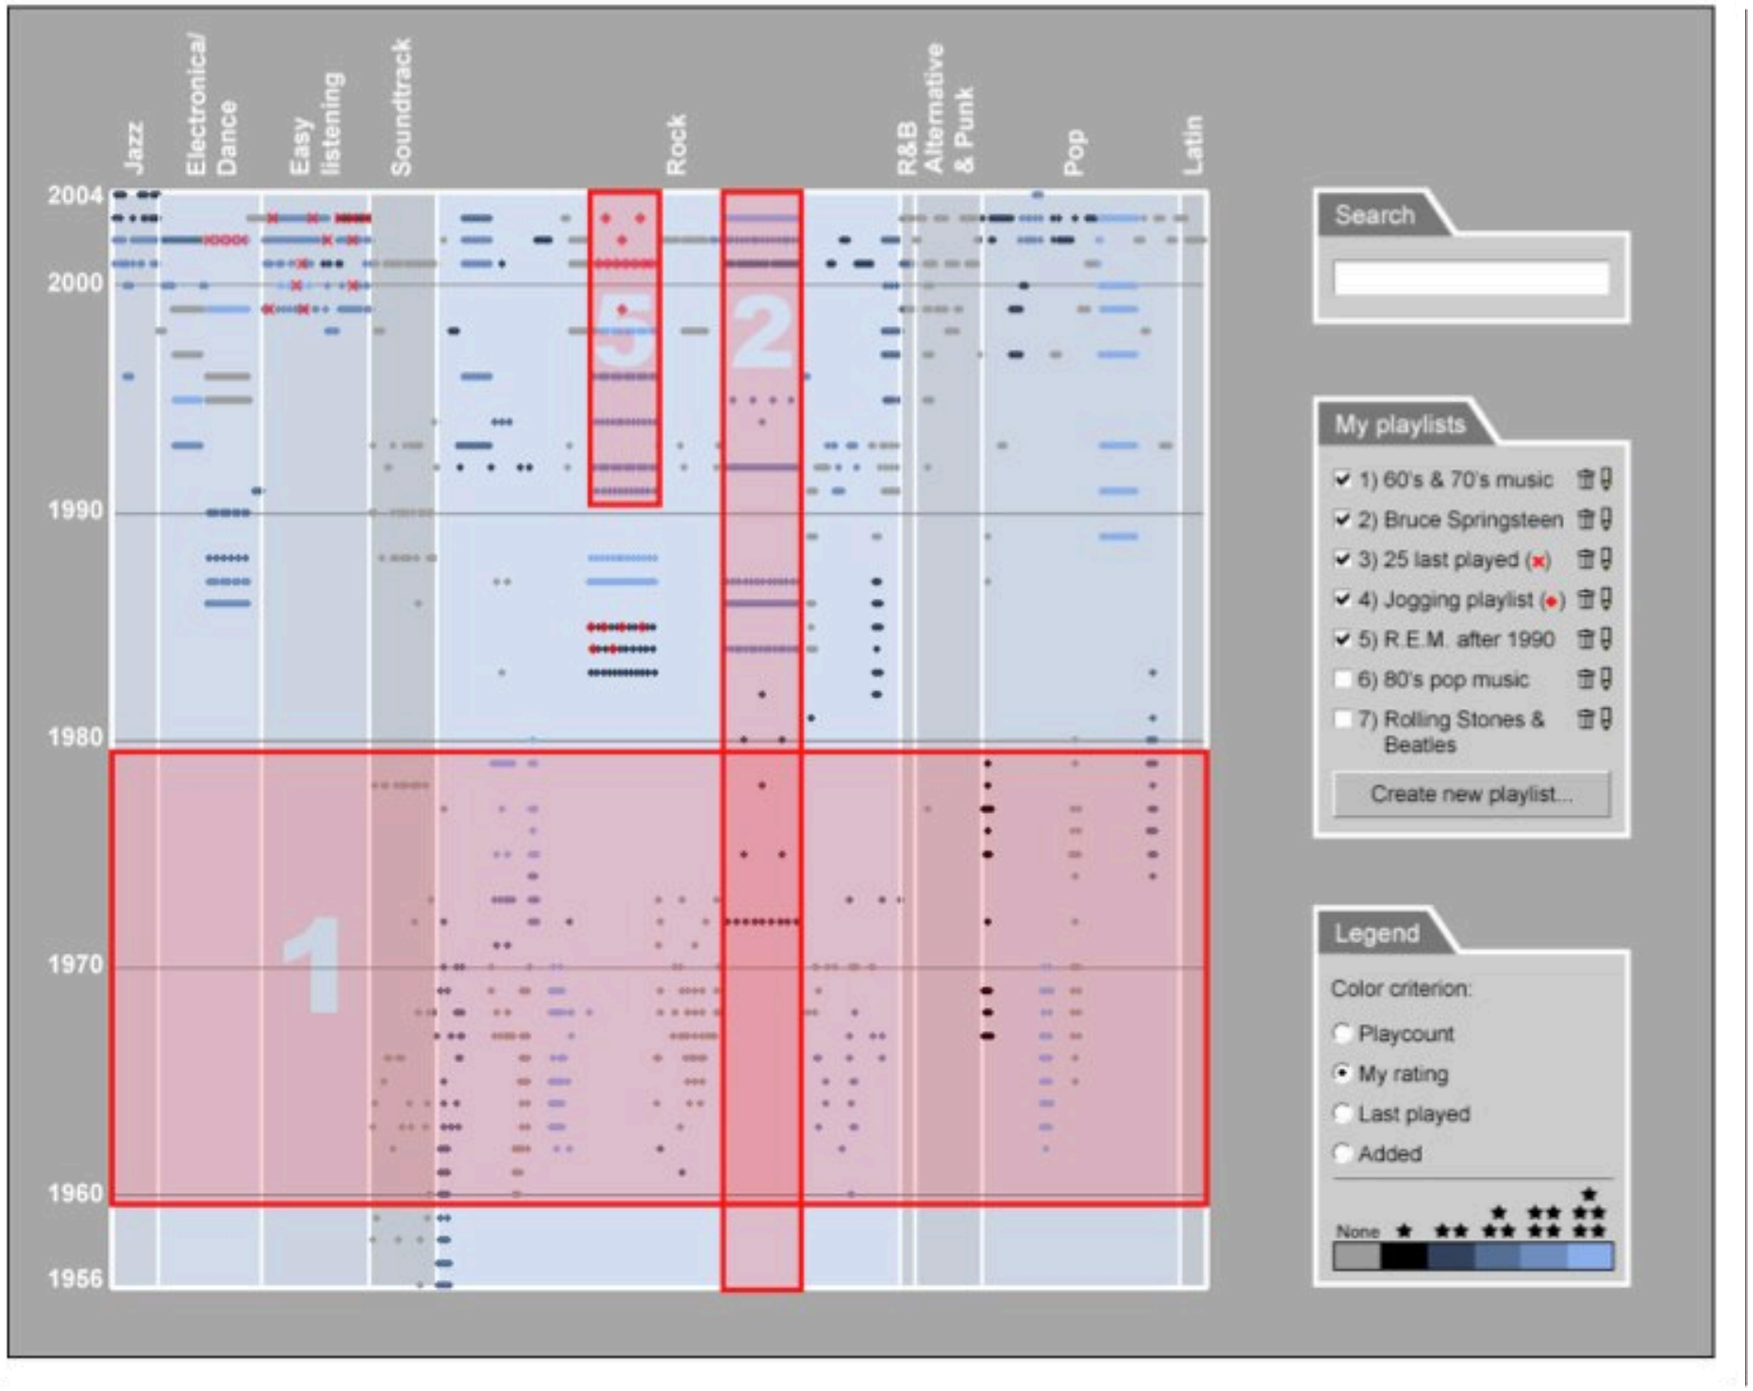

http://torrens.files.wordpress.com/2009/06/ismir.pdf
Torrens, Patrick Hertzog, Josep-Lluís Arcos

Music Explorer FX

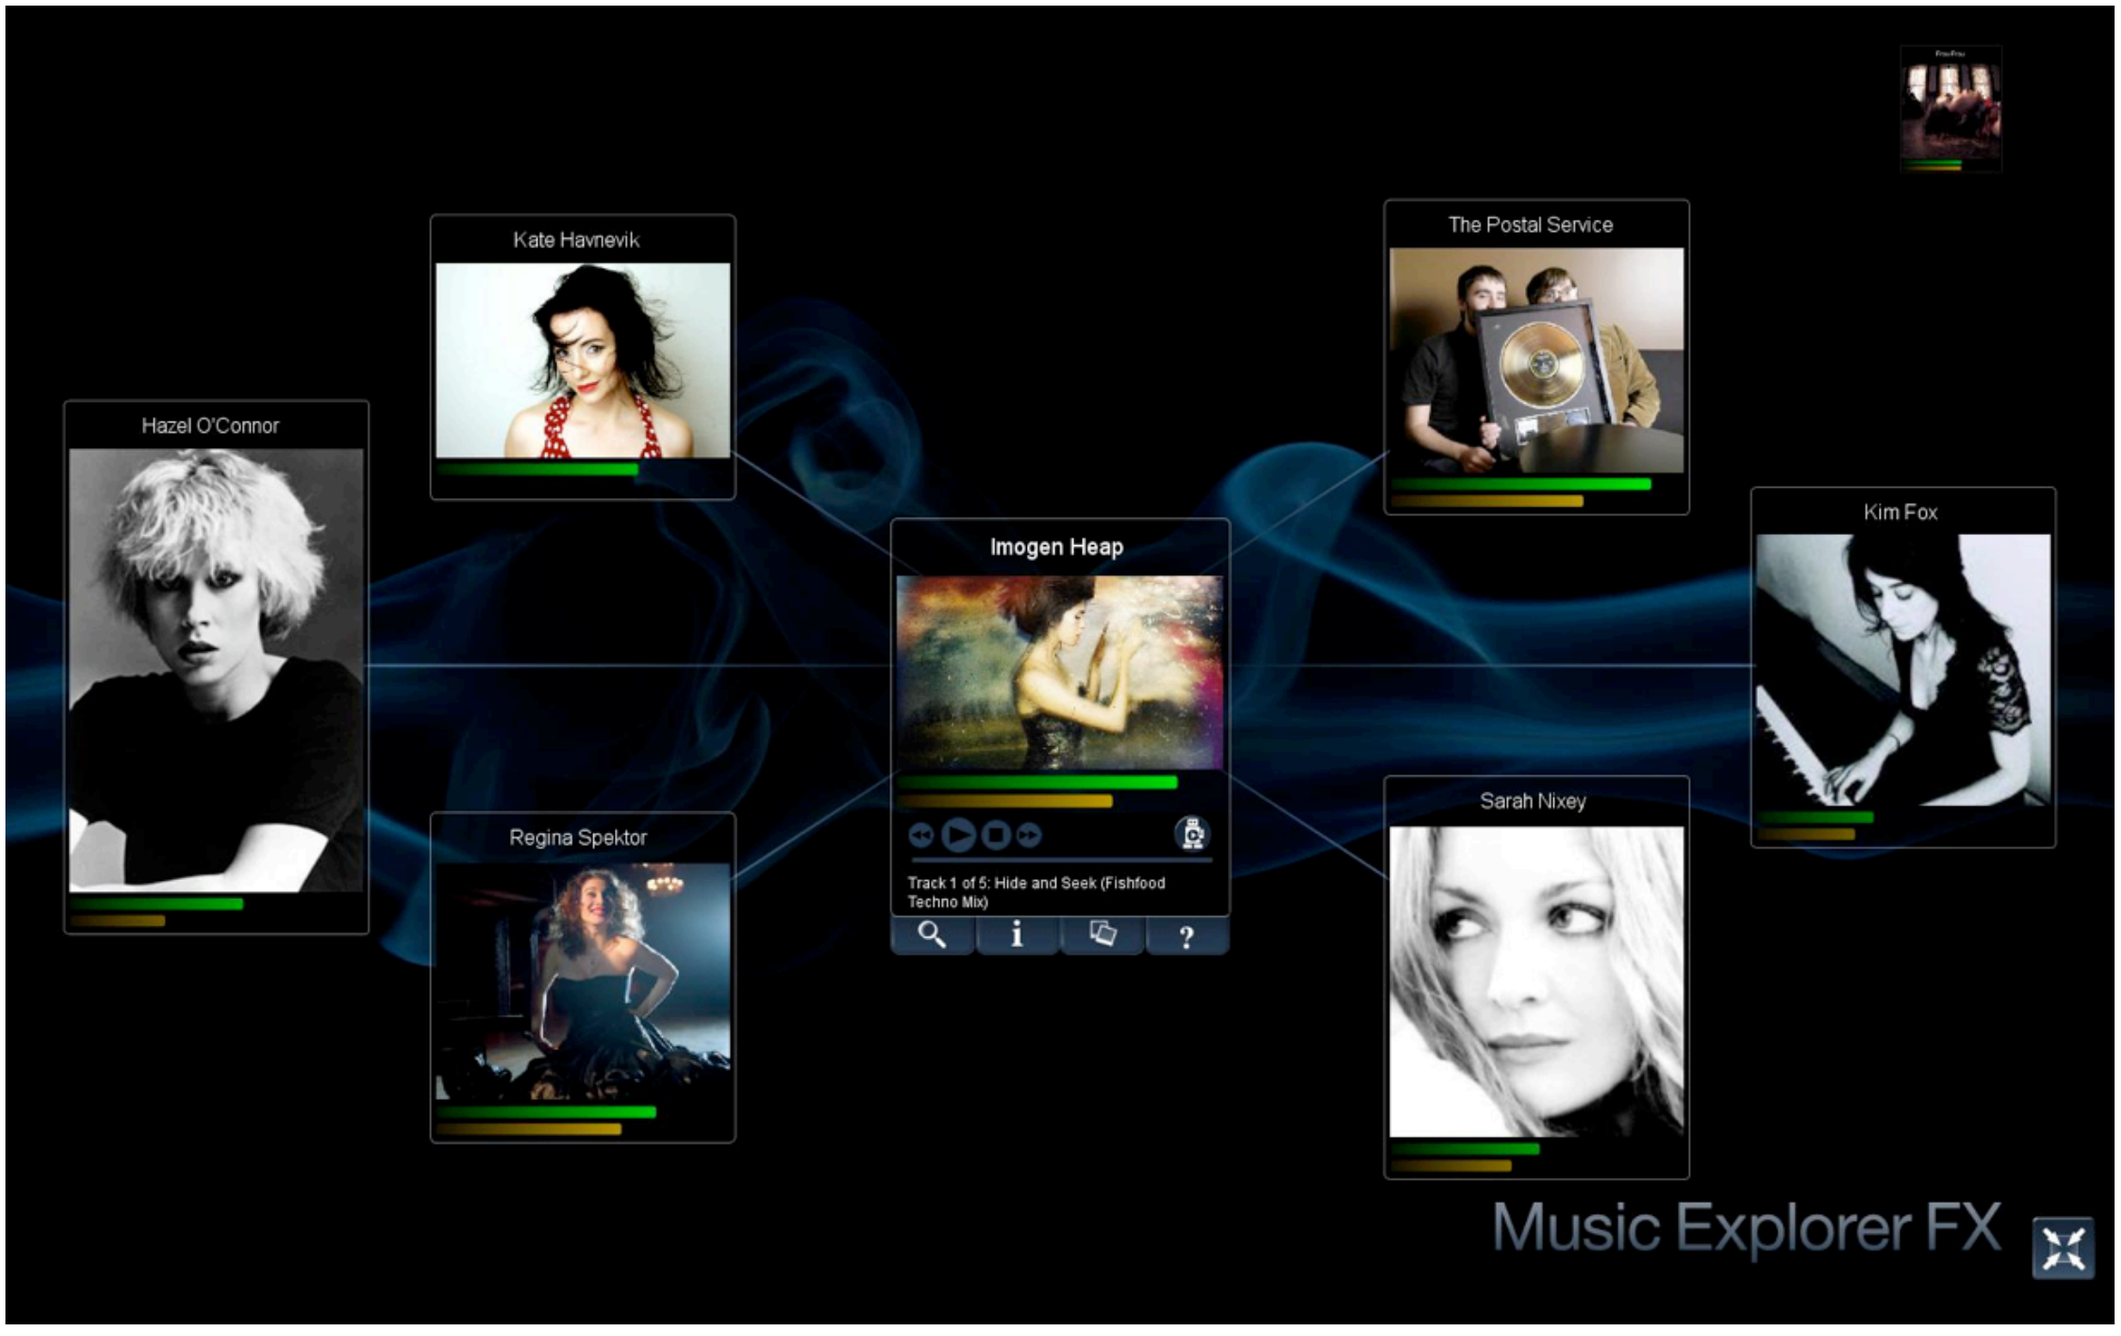

http://musicexplorerfx.citytechinc.com/

Social Music Discovery

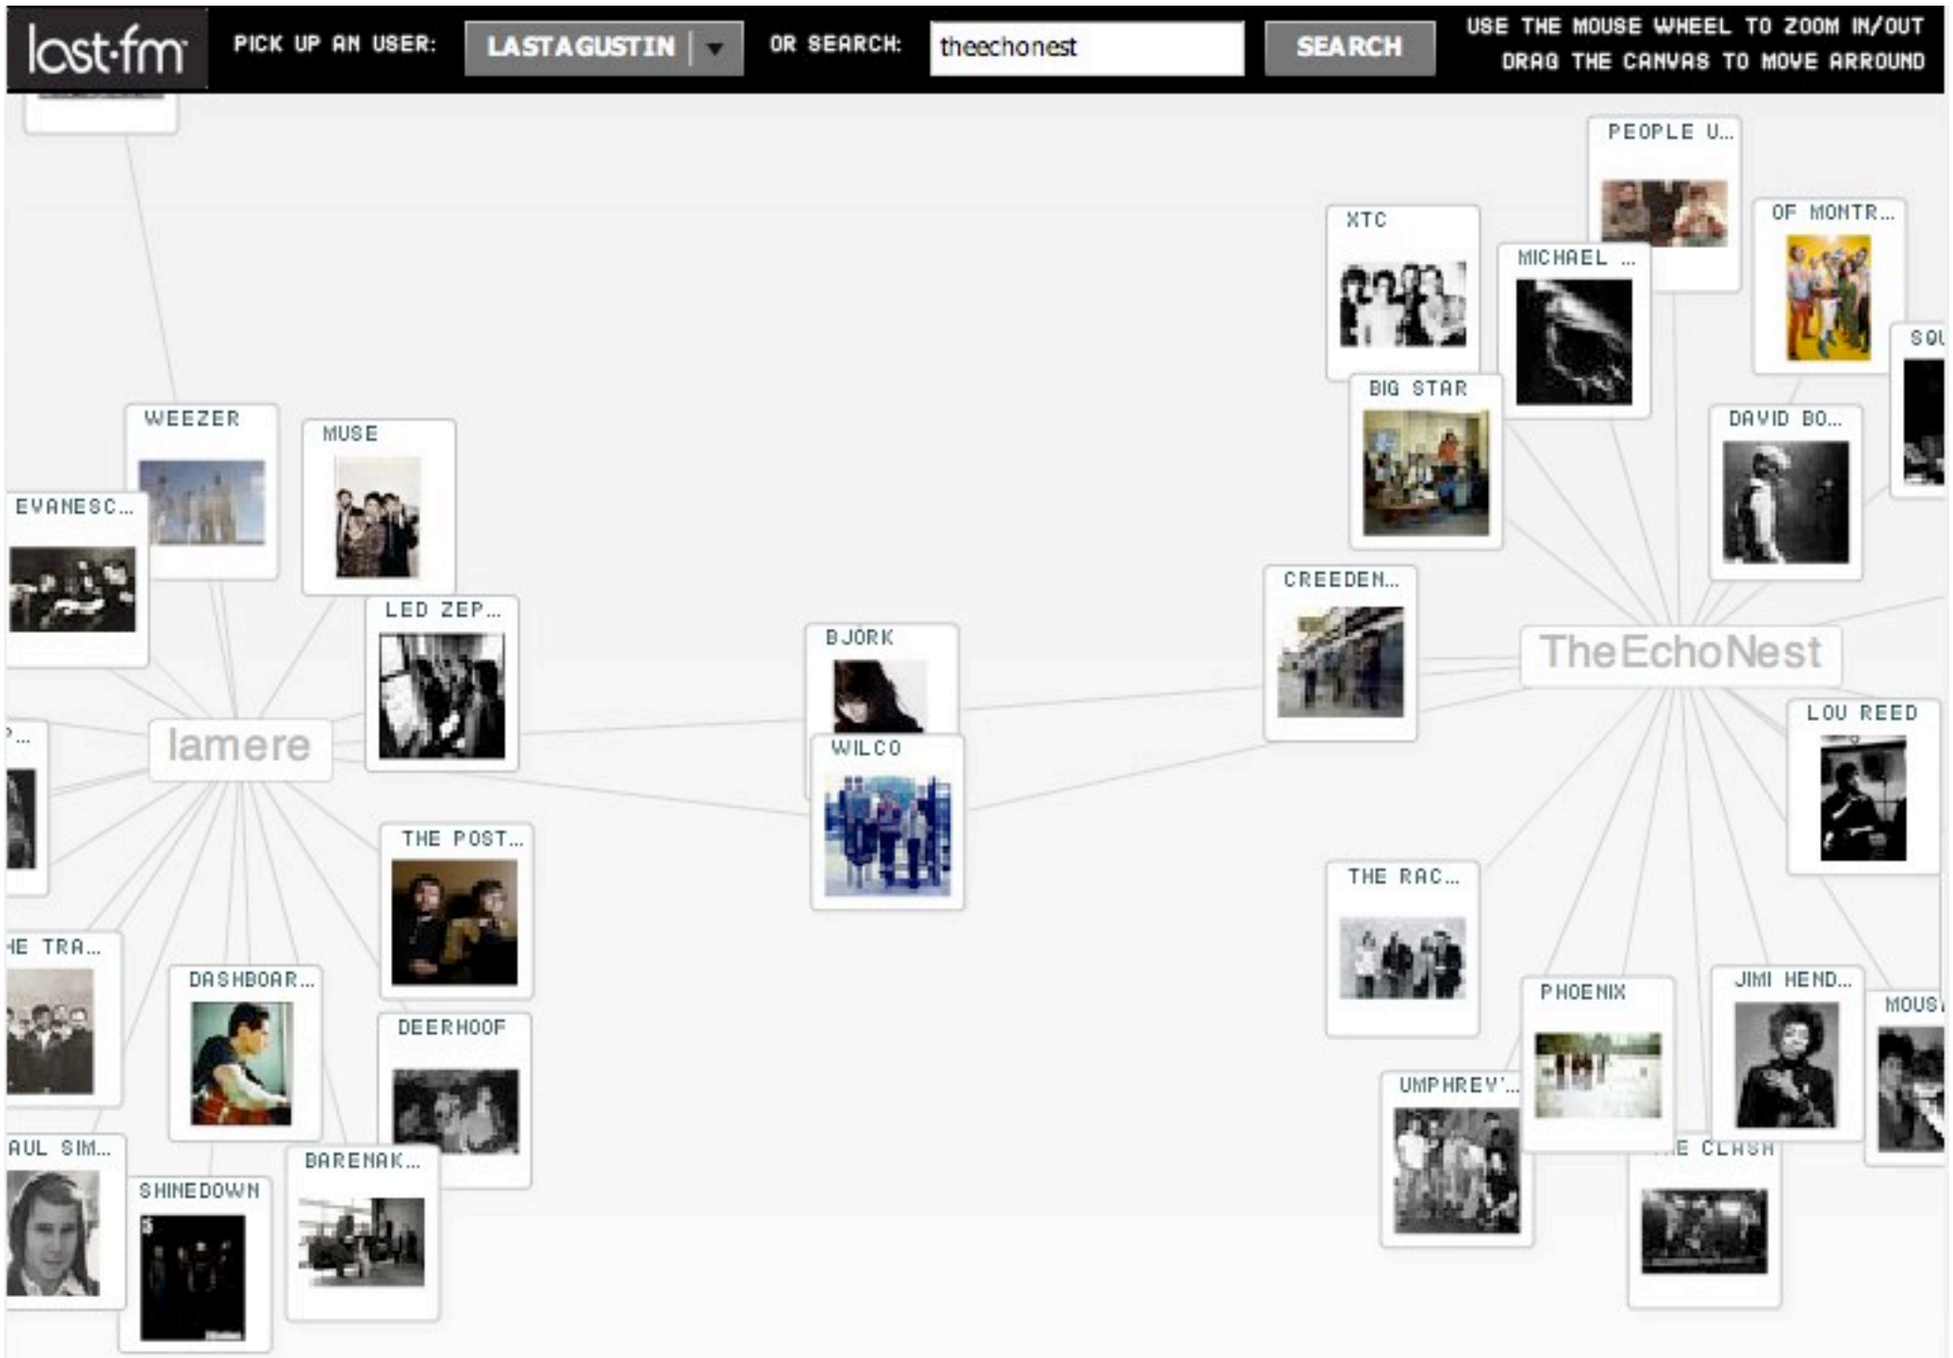

Last FM tasteOgraph

http://geniol.publishpath.com/last-fm-tasteograph

Floko-On

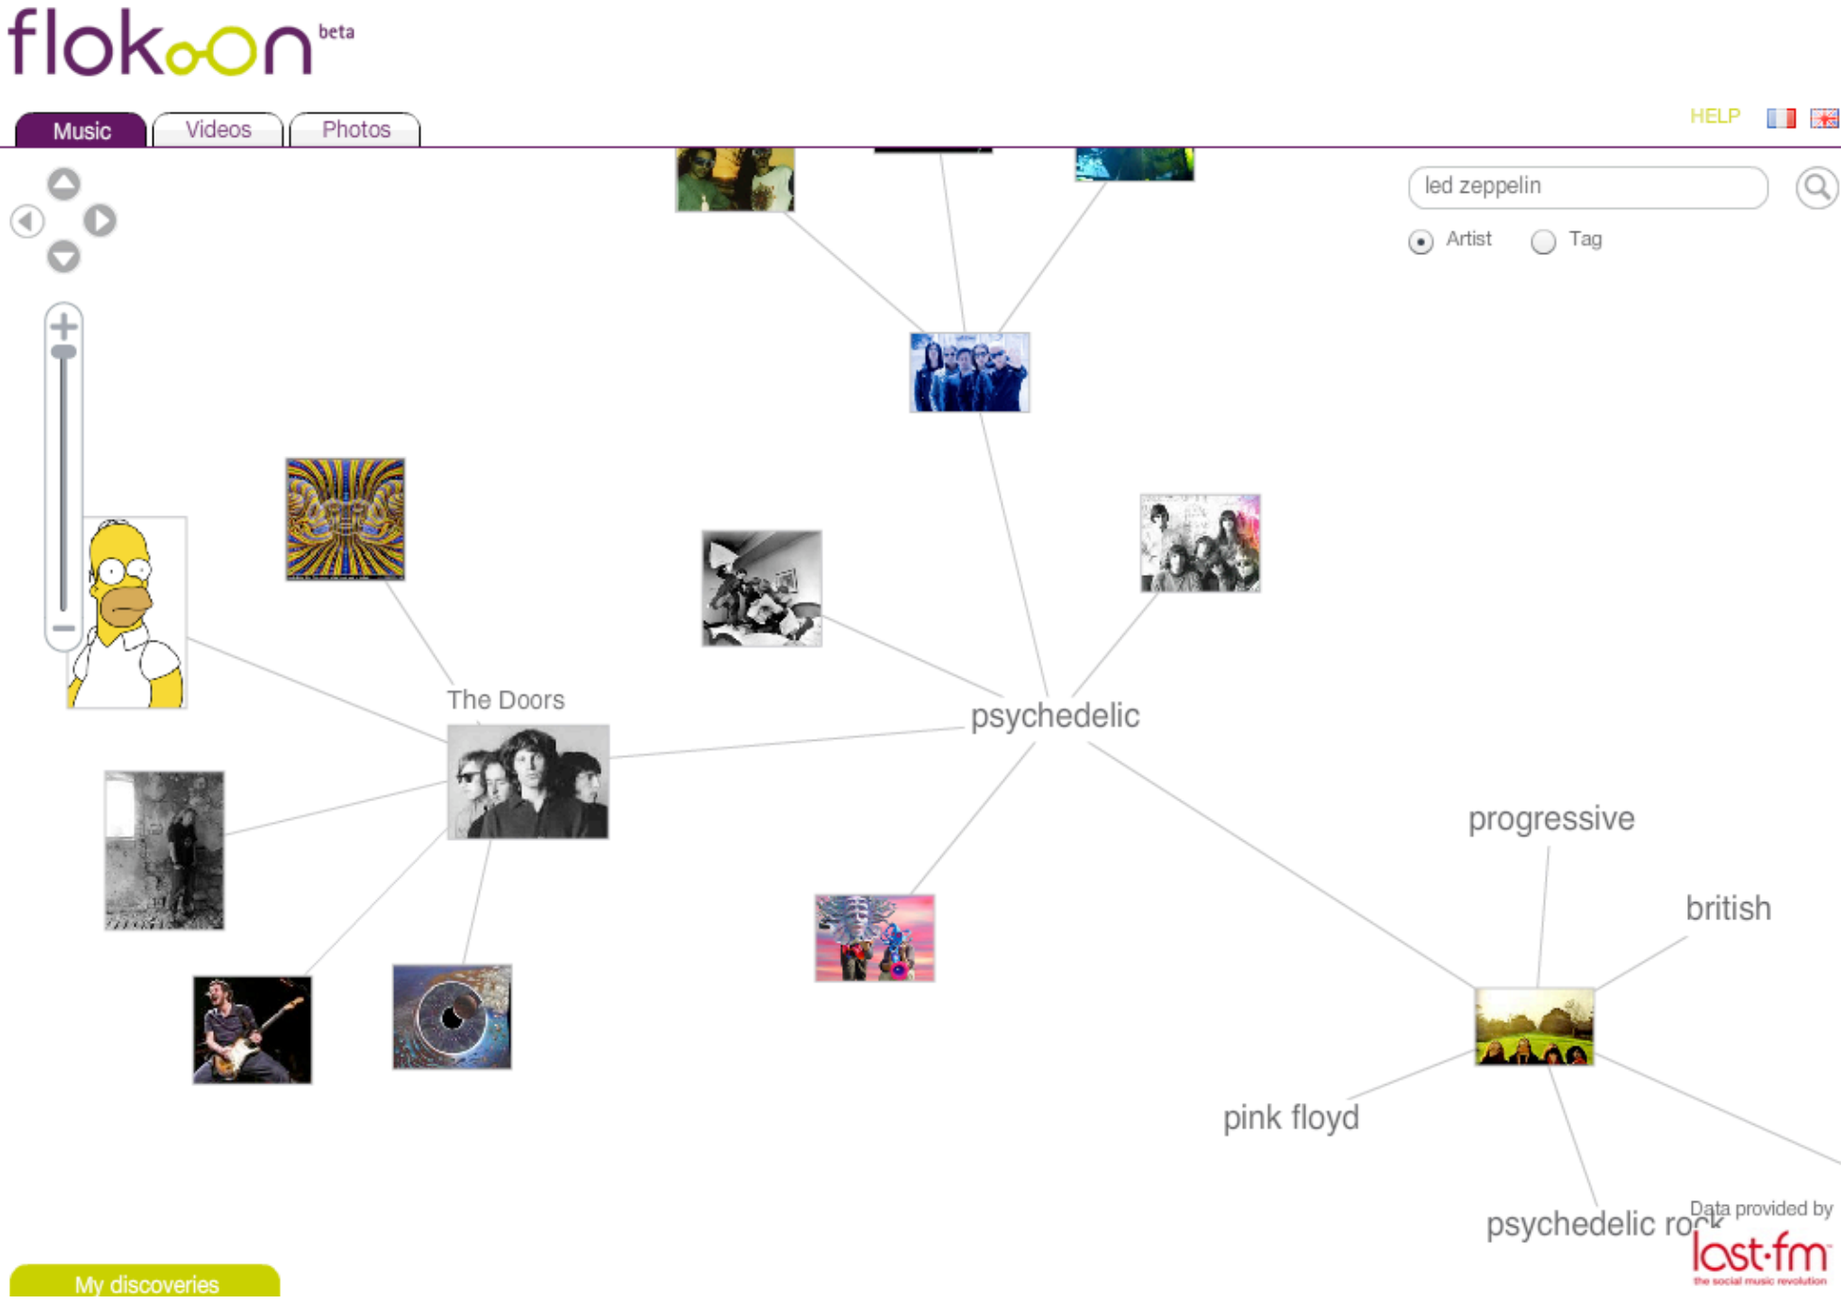

http://www.flokoon.com/#lastfm

Project Orpheus

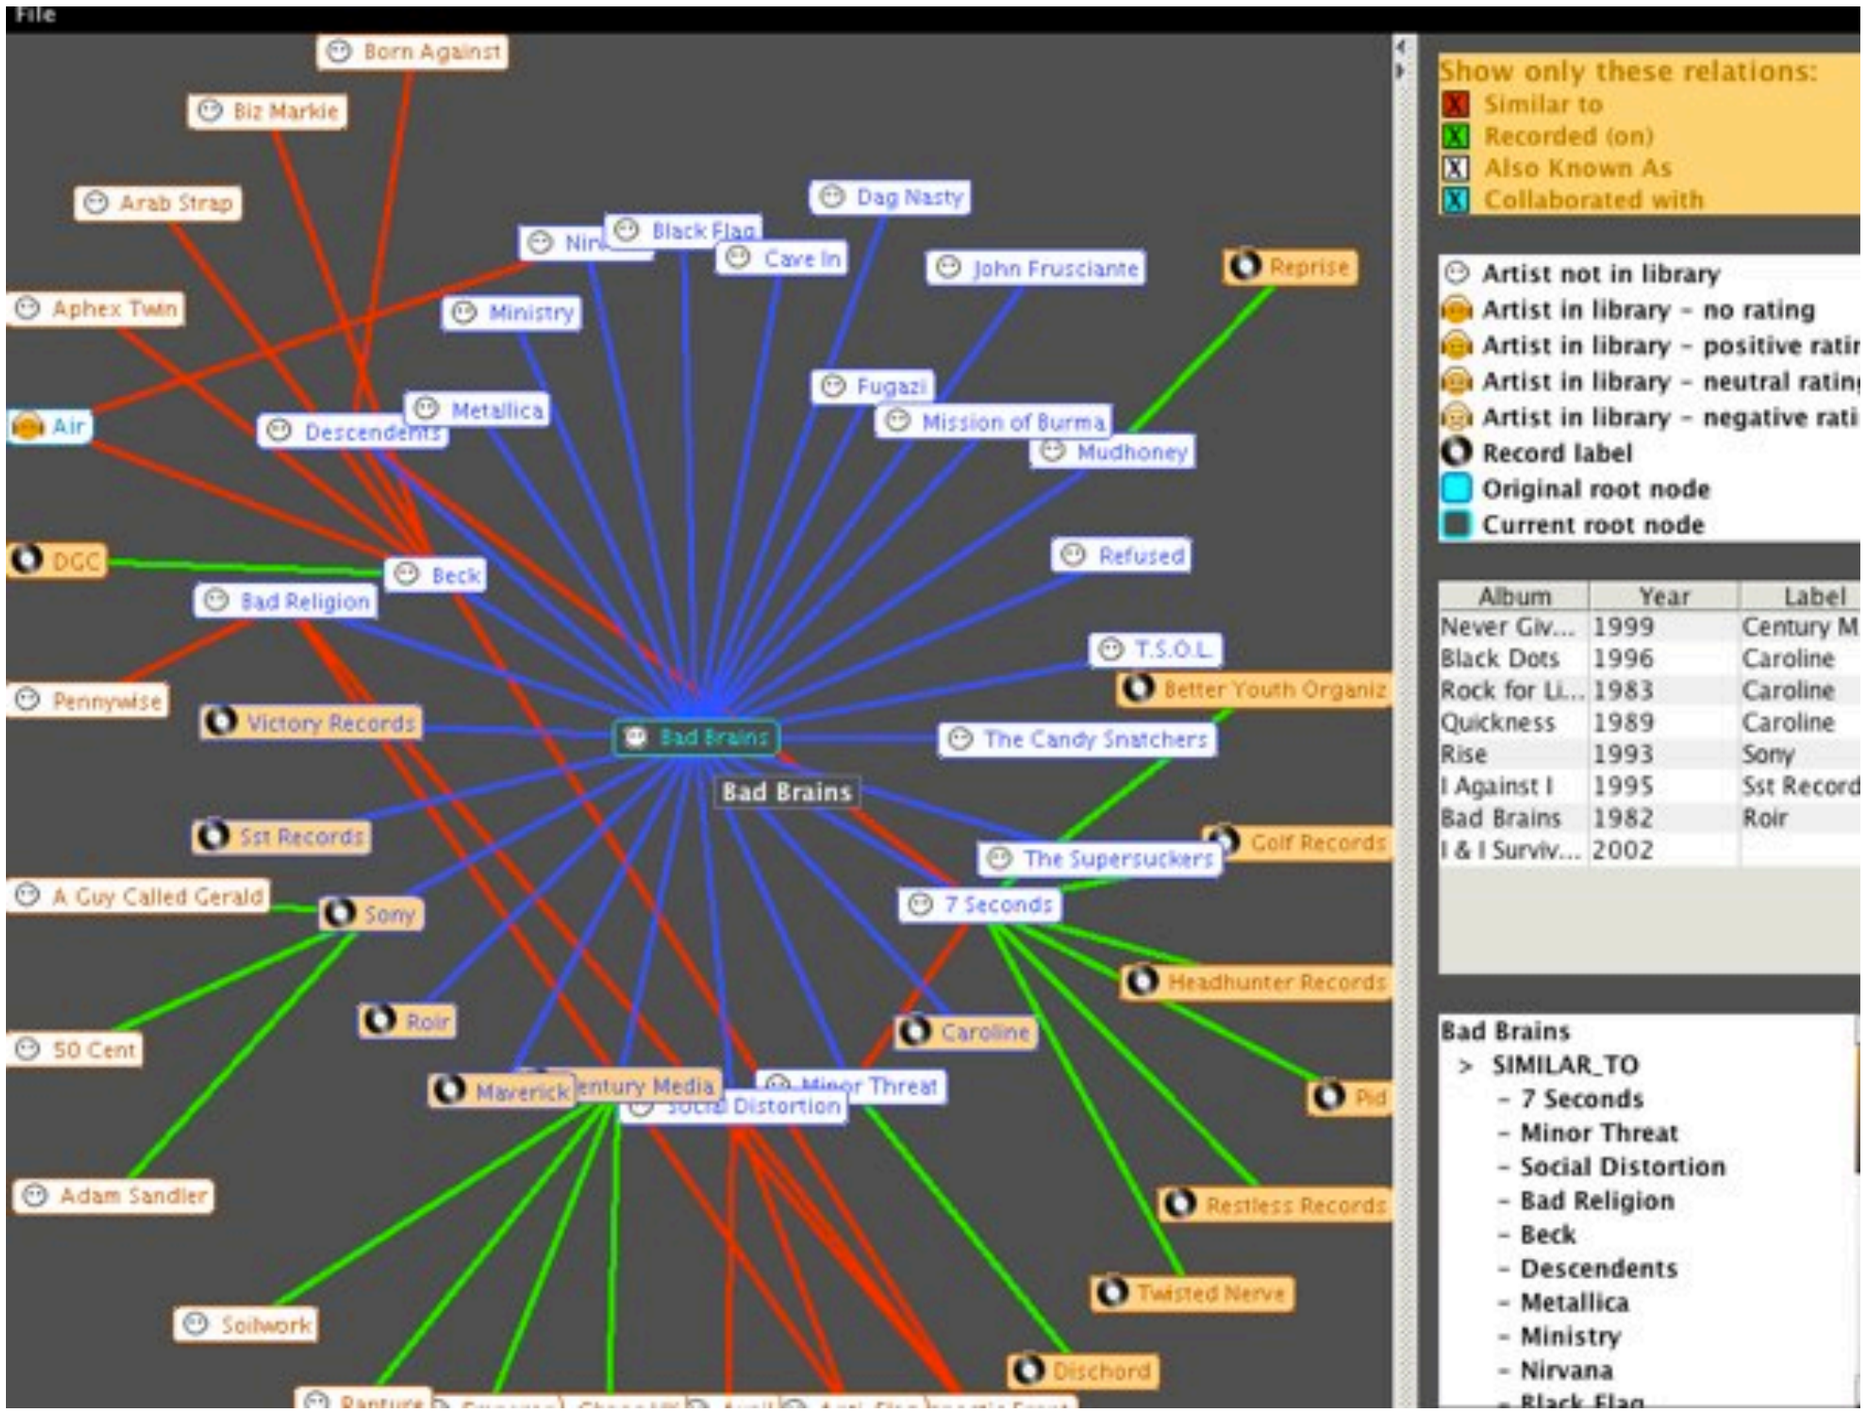

http://arcus-associates.com/orpheus/

Musician Map: visualizing music collaborations over time

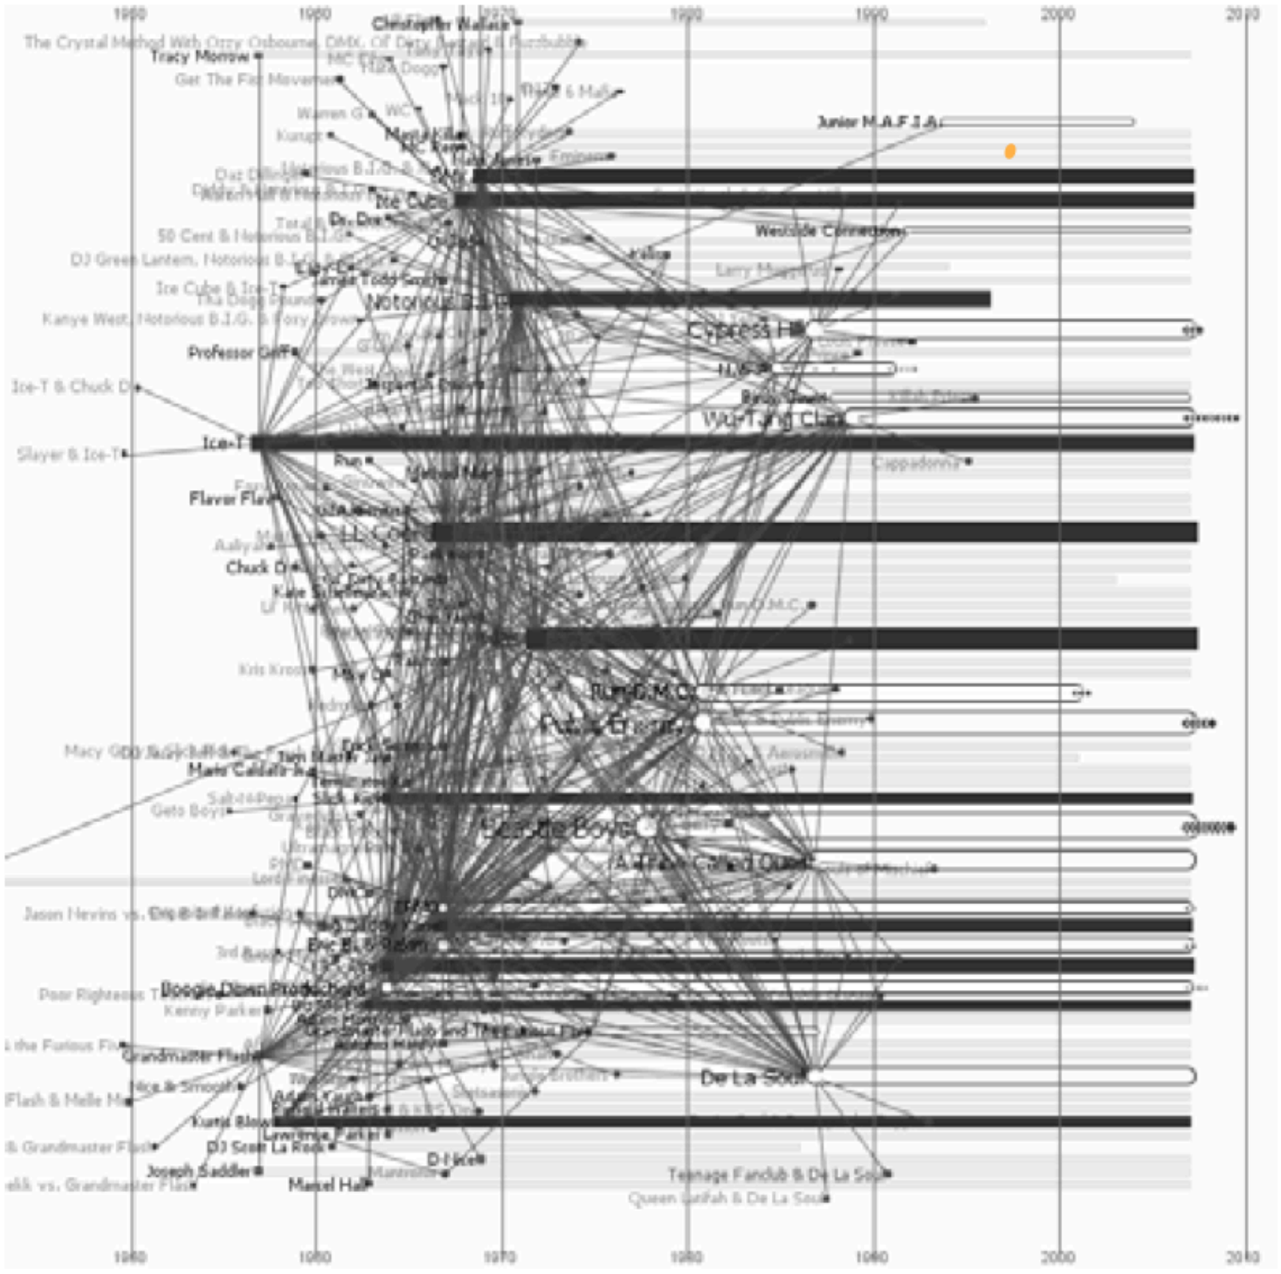

http://www.sfu.ca/~jdyim/musicianMap/

Ji-Dong Yim Chris Shaw Lyn Bartram

lyric connections

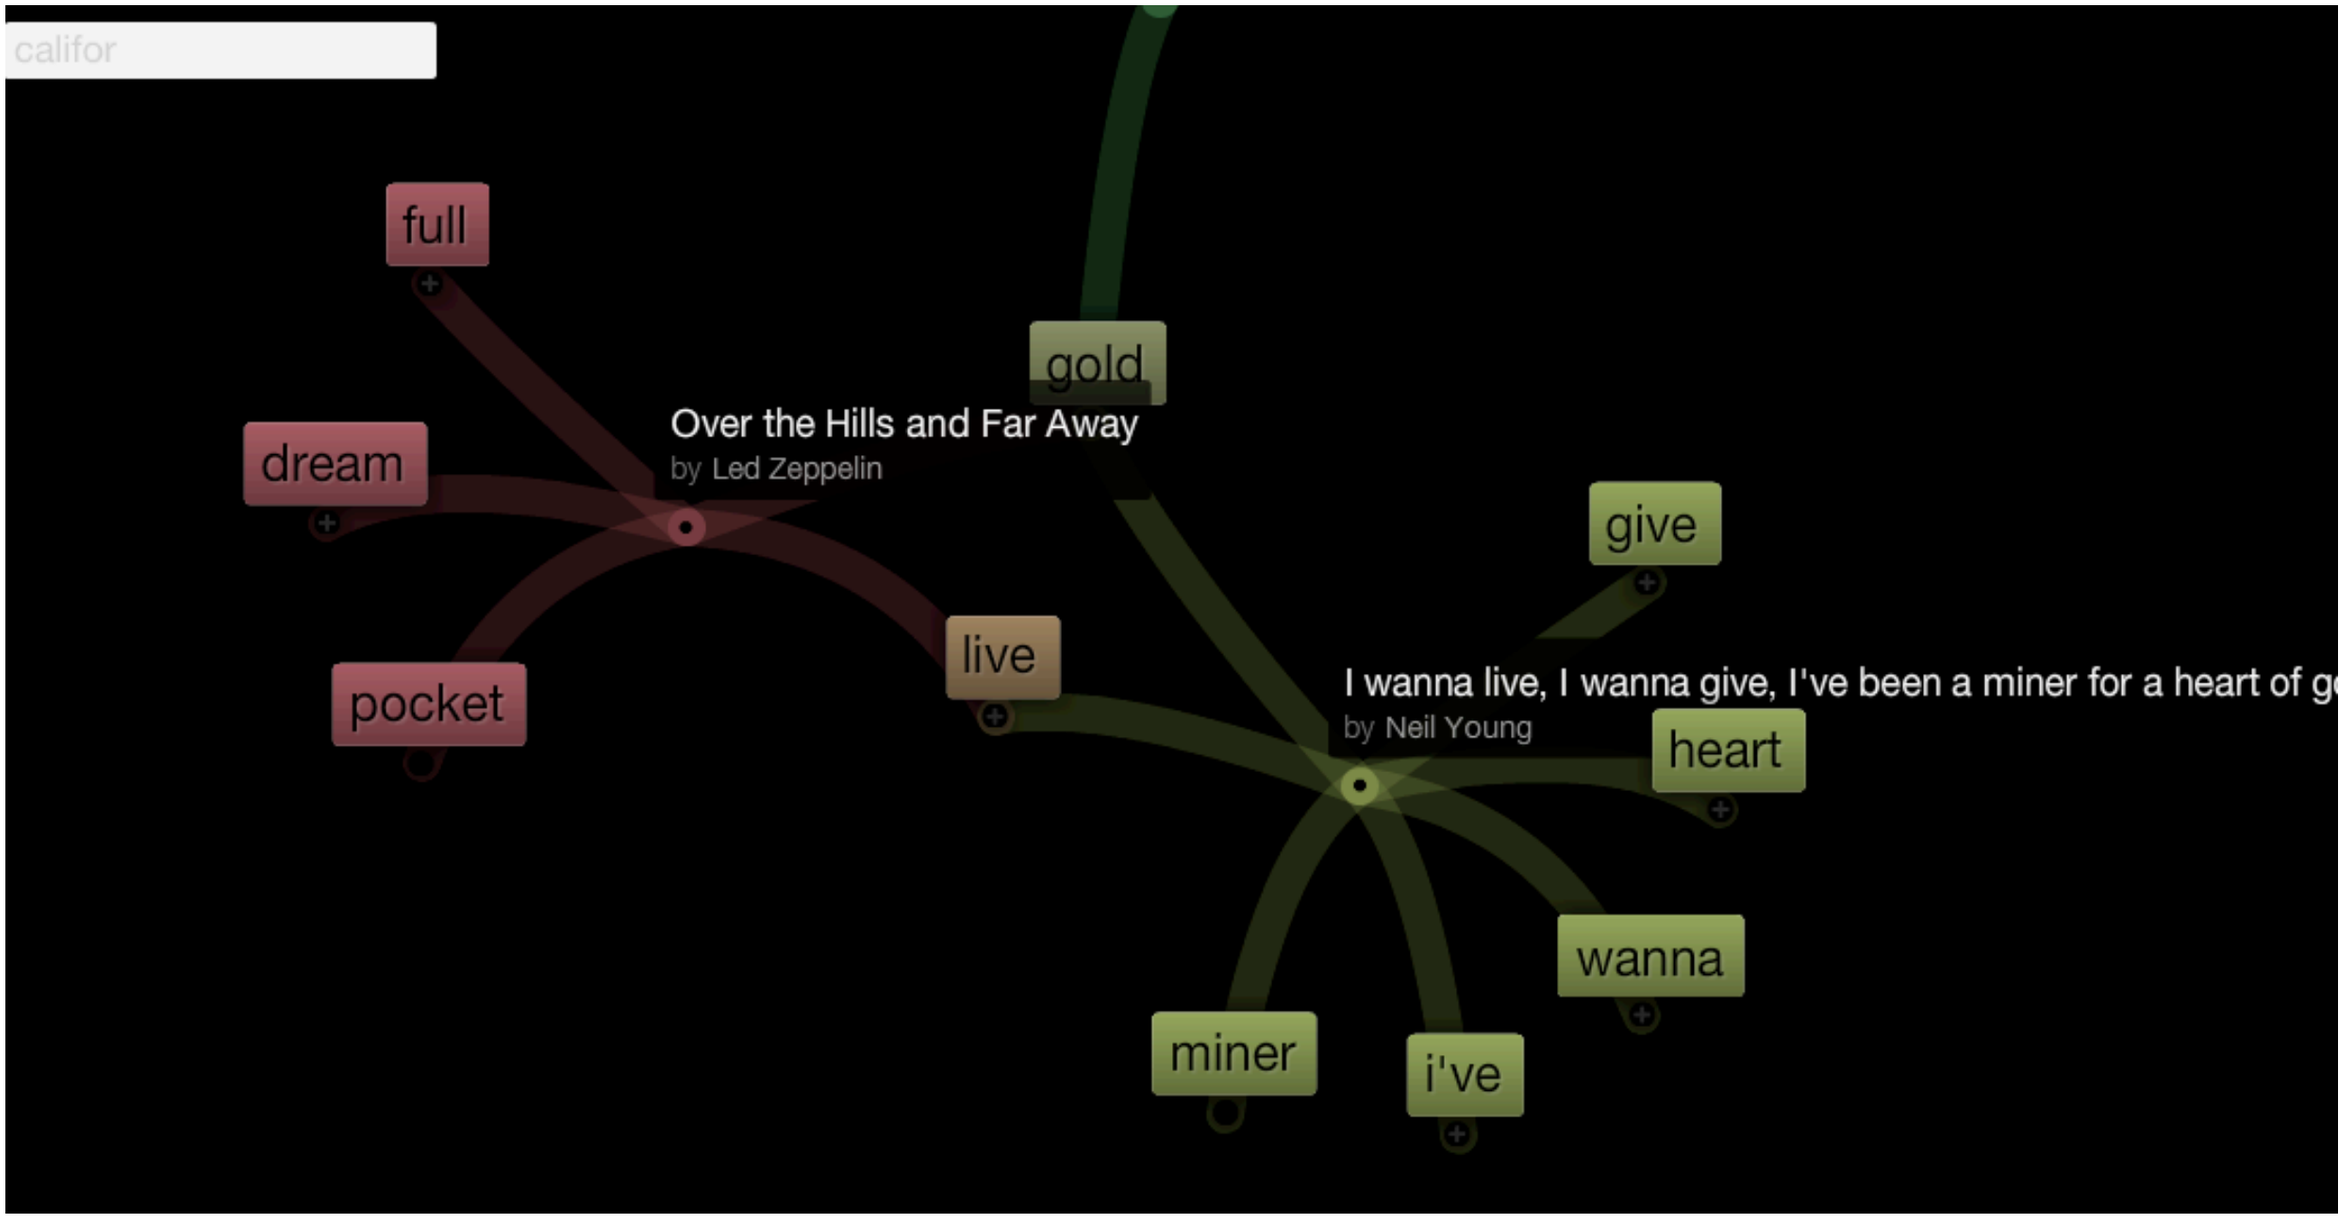

http://www.identitee.com/visualizer.php

music tees and tee shirts on i/denti/tee

Interacting with linked data about music

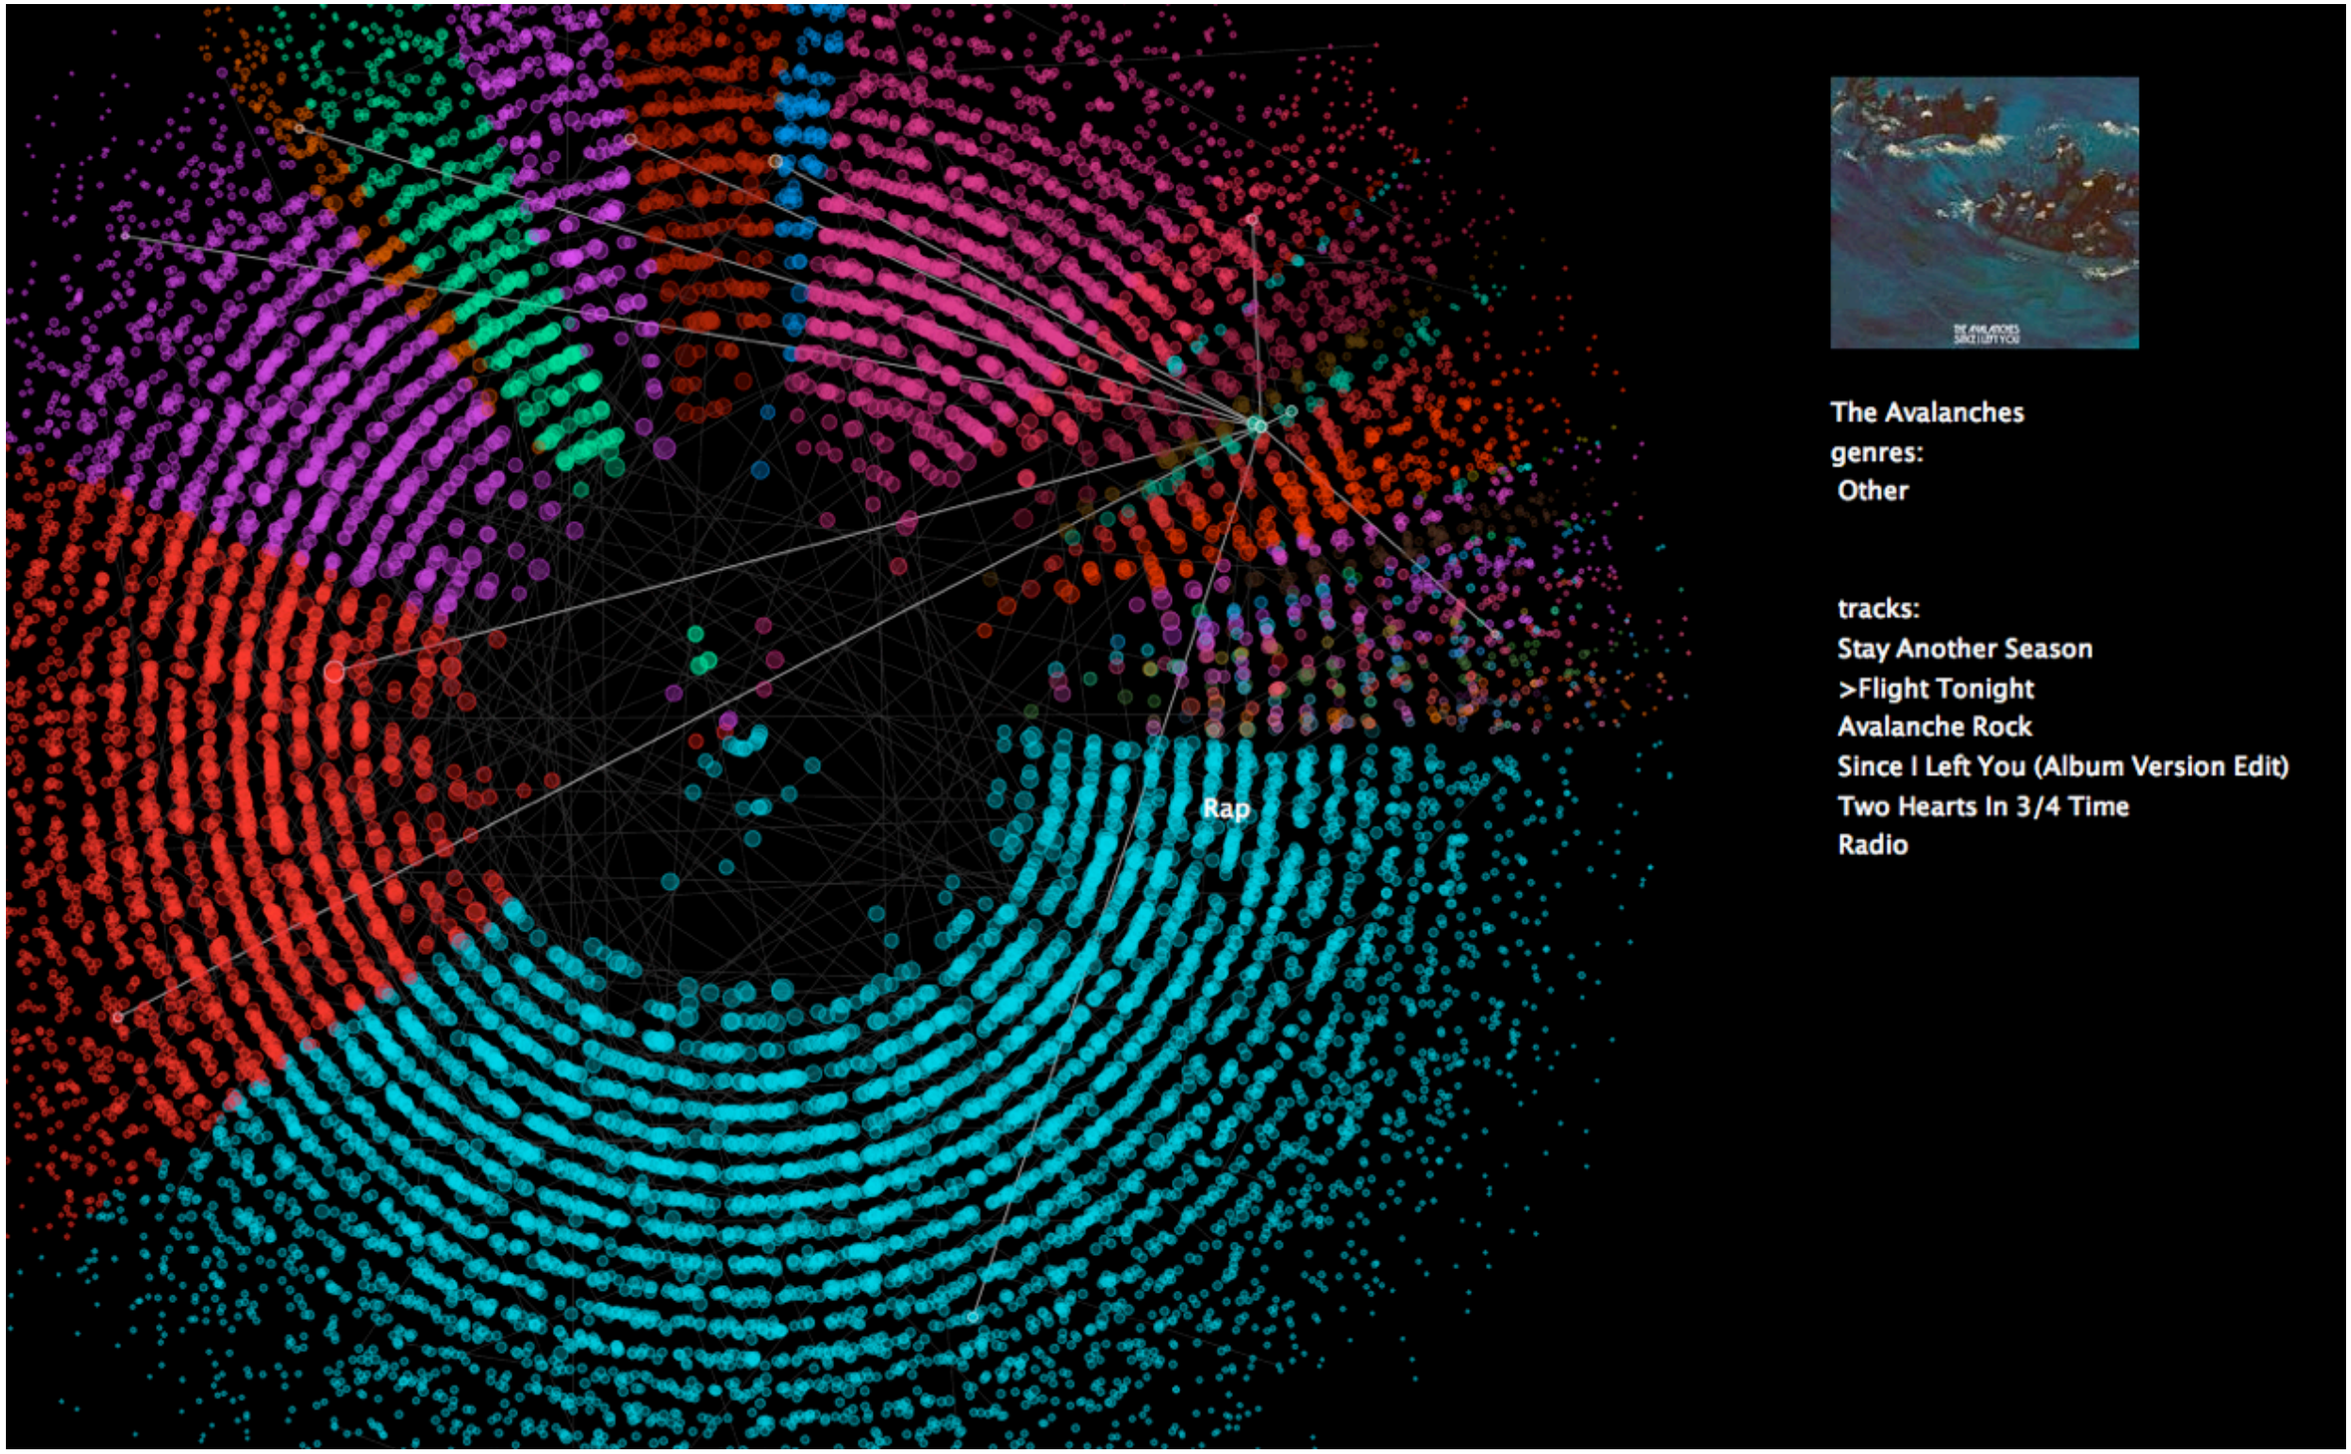

http://journal.webscience.org/110/

Interacting with linked data about music

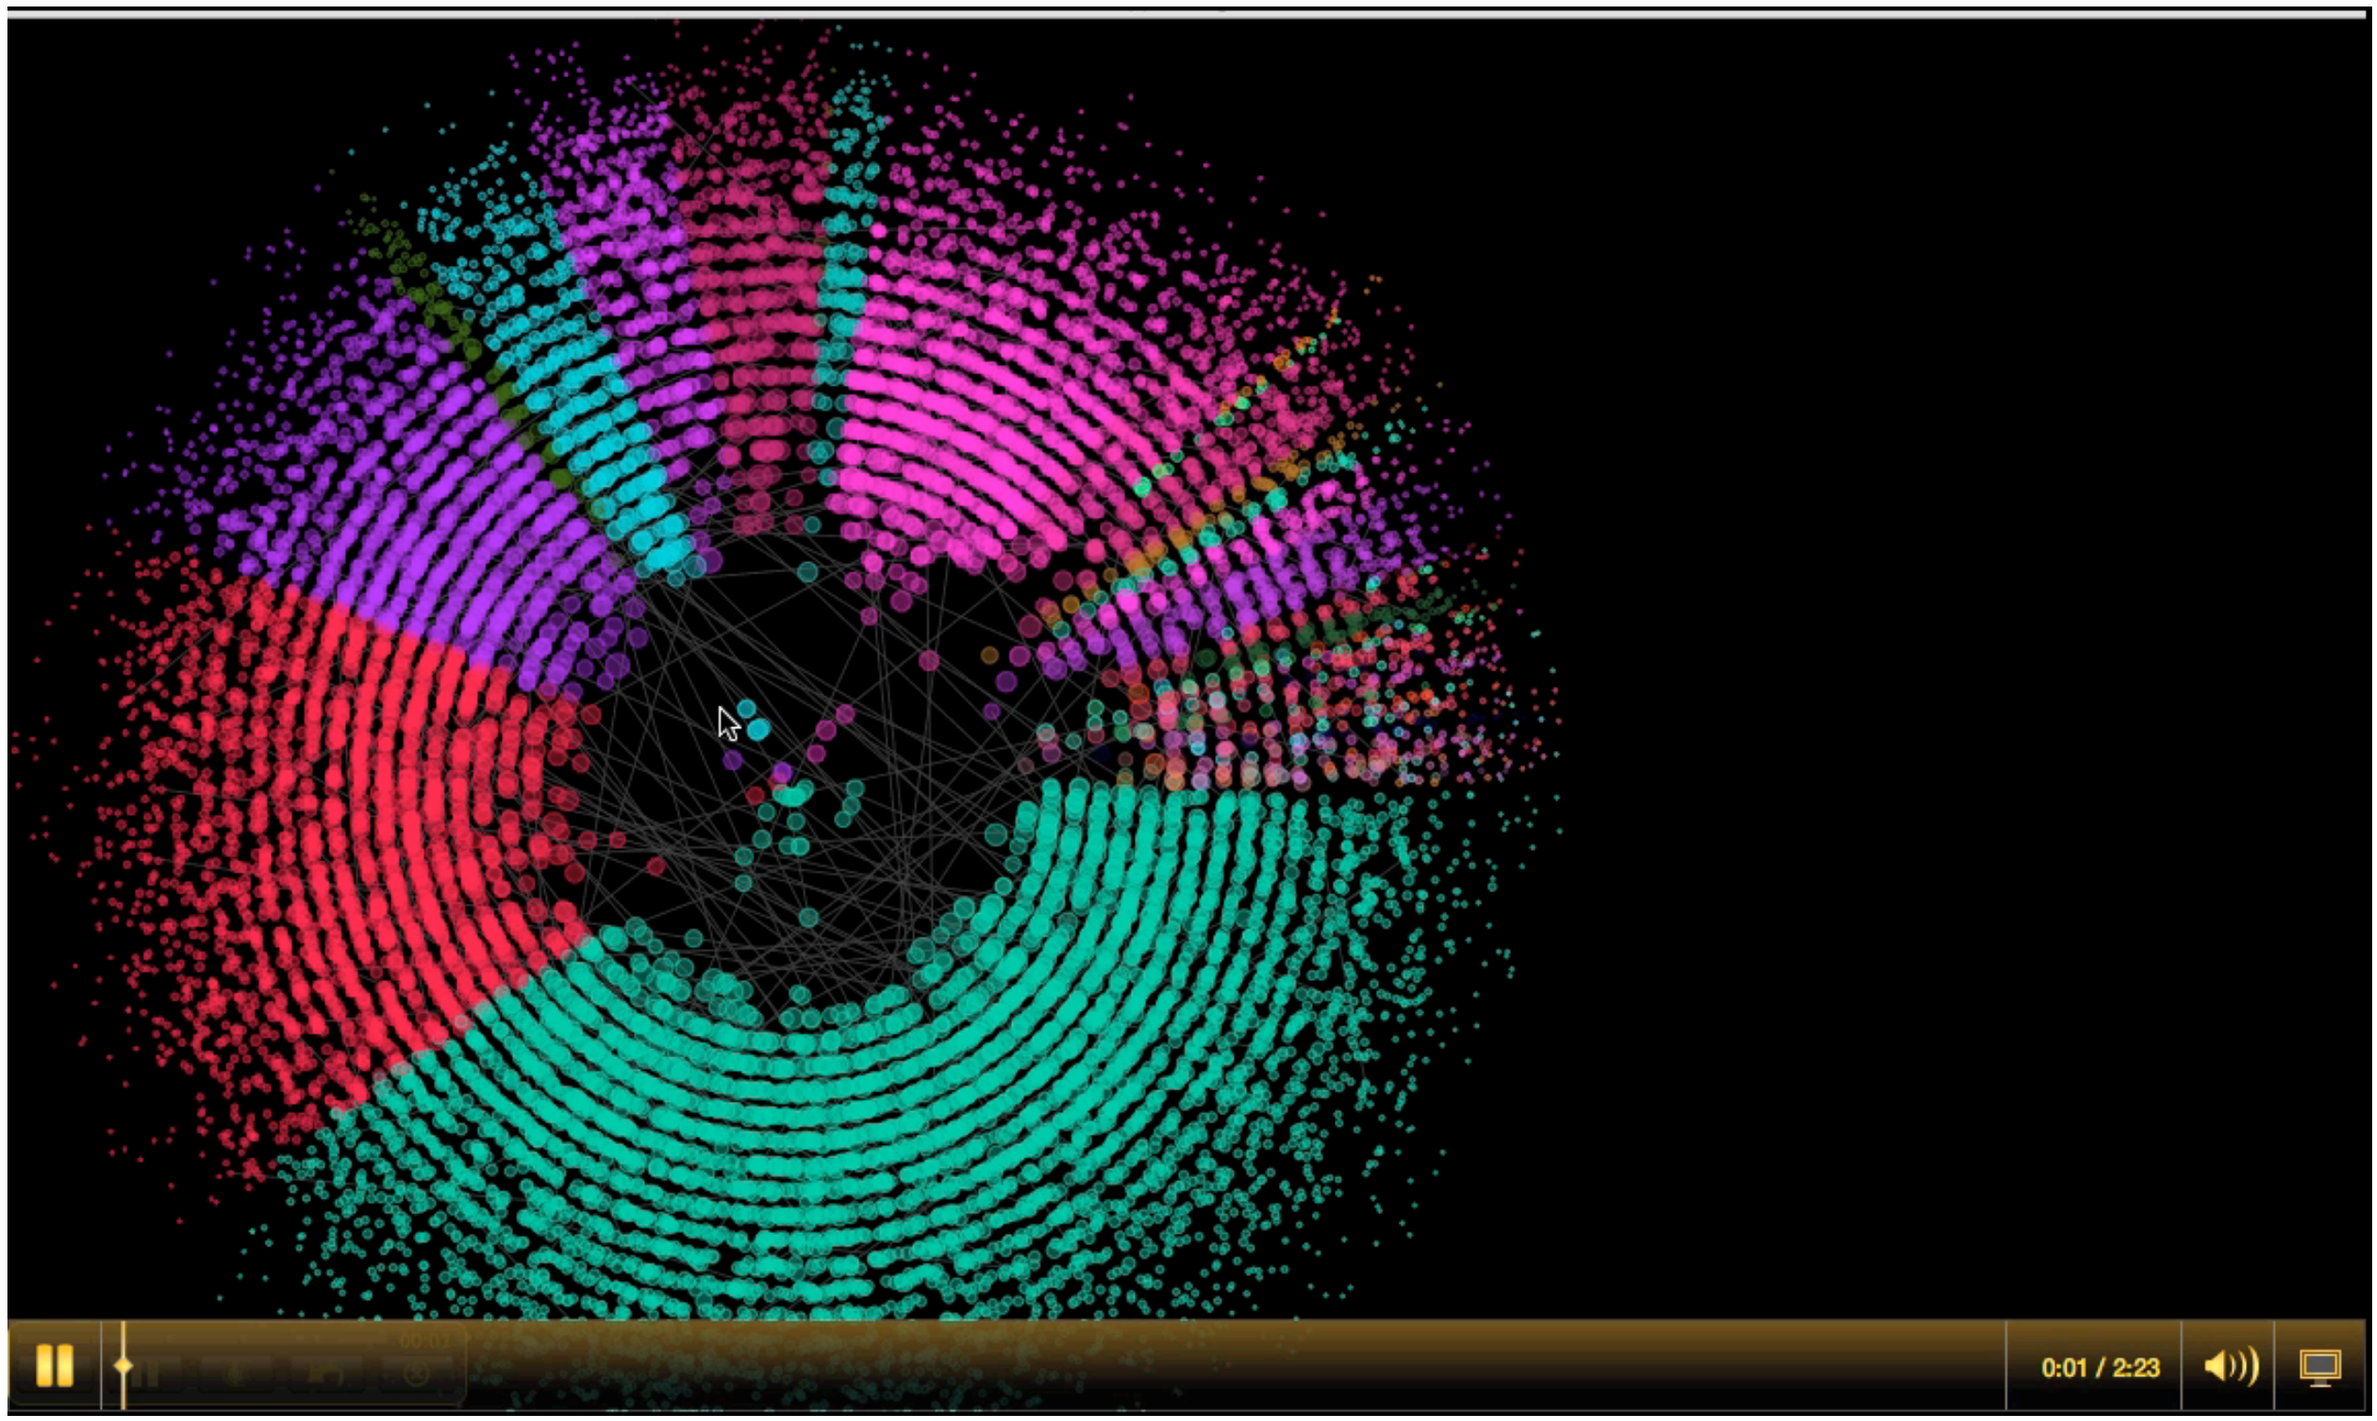

http://journal.webscience.org/110/

The World Of Music

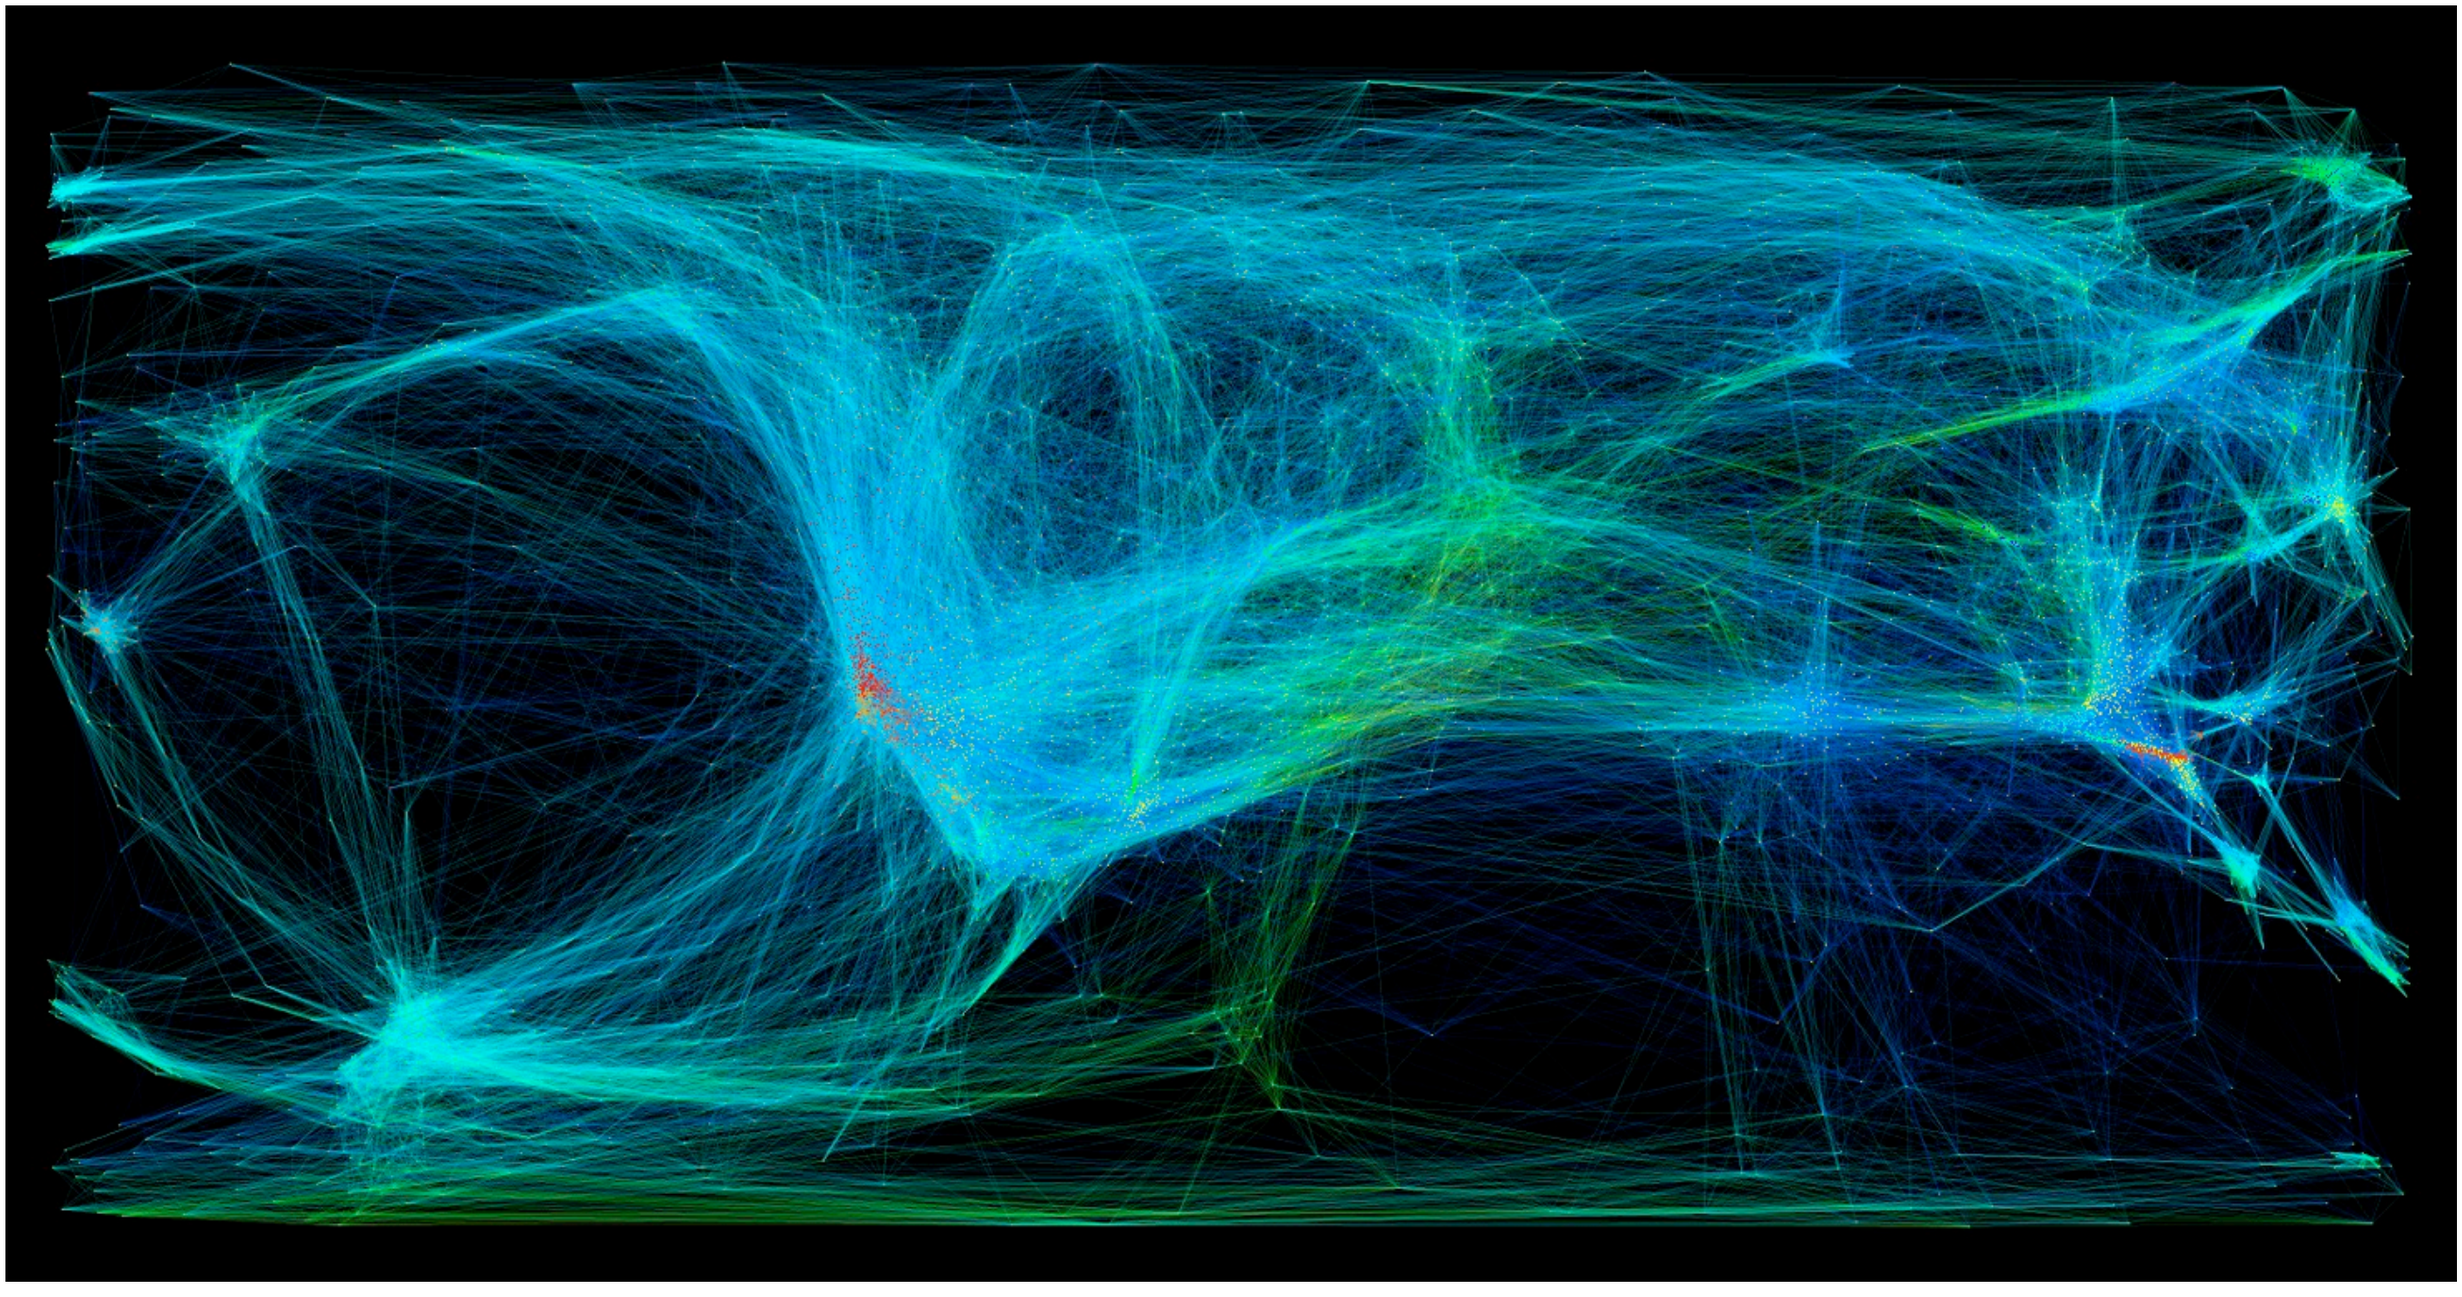

http://www.stanford.edu/~dgleich/demos/worldofmusic/WorldOfMusic.html
David Gleich, Matt Rasmussen, Leonid Zhukov, and Kevin Lang

The World Of Music

http://www.stanford.edu/~dgleich/demos/worldofmusic/WorldOfMusic.html
David Gleich, Matt Rasmussen, Leonid Zhukov, and Kevin Lang

Last.fm Artist Map

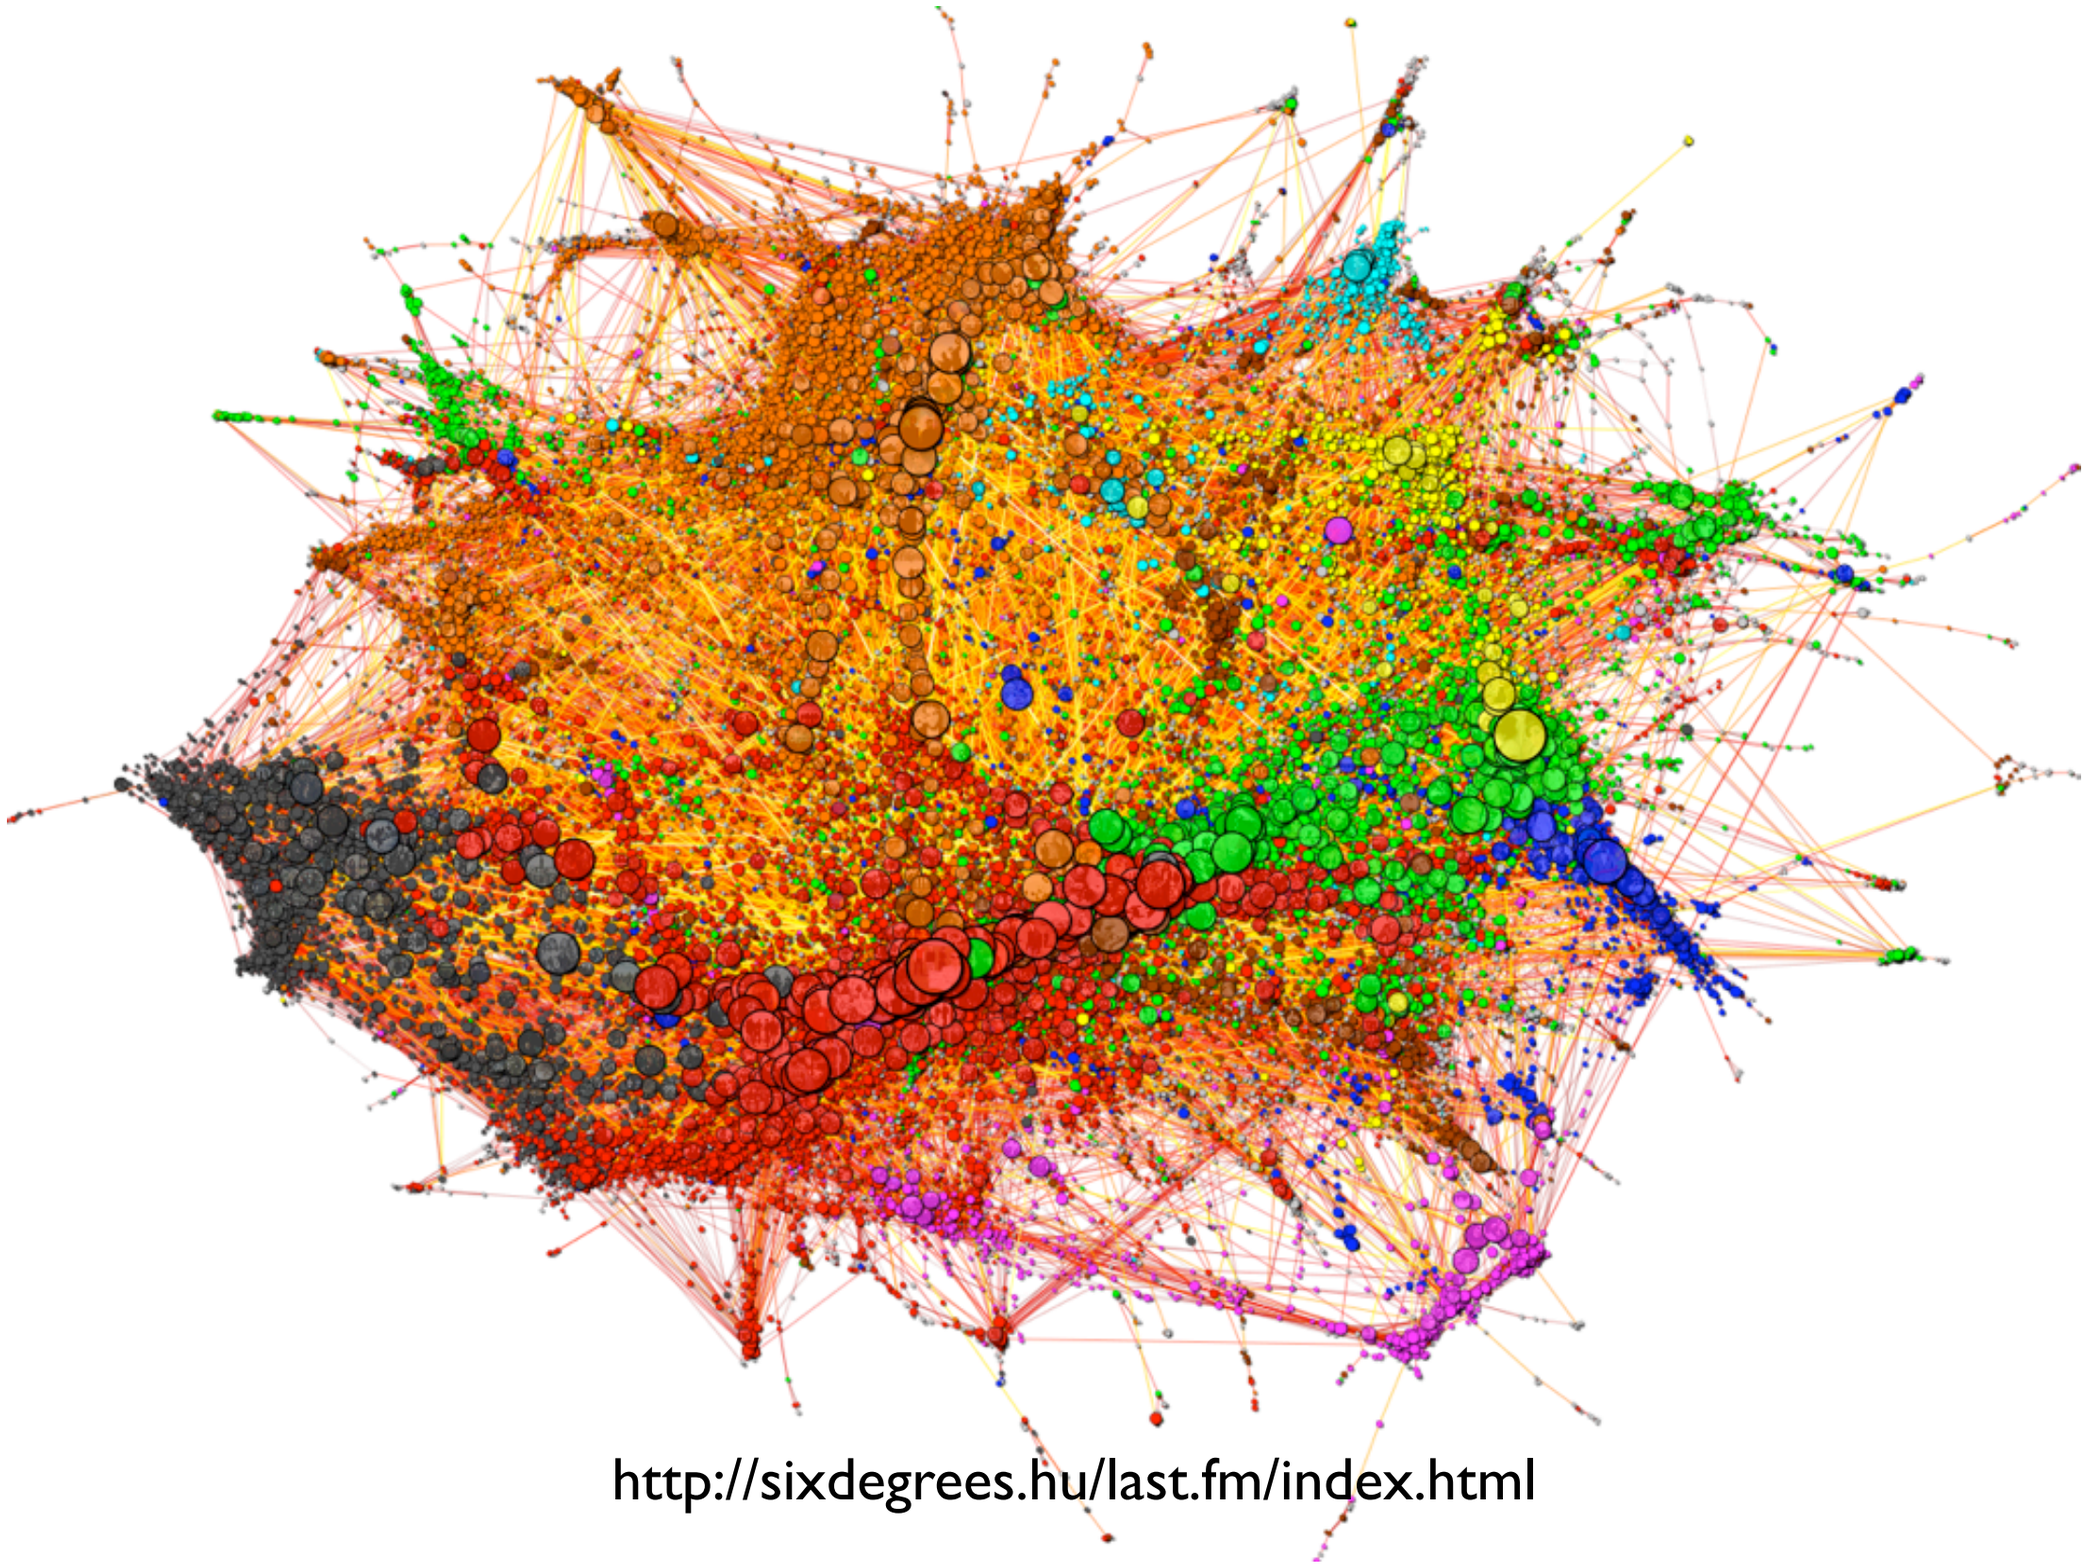

http://sixdegrees.hu/last.fm/index.html

Last.fm Artist Map

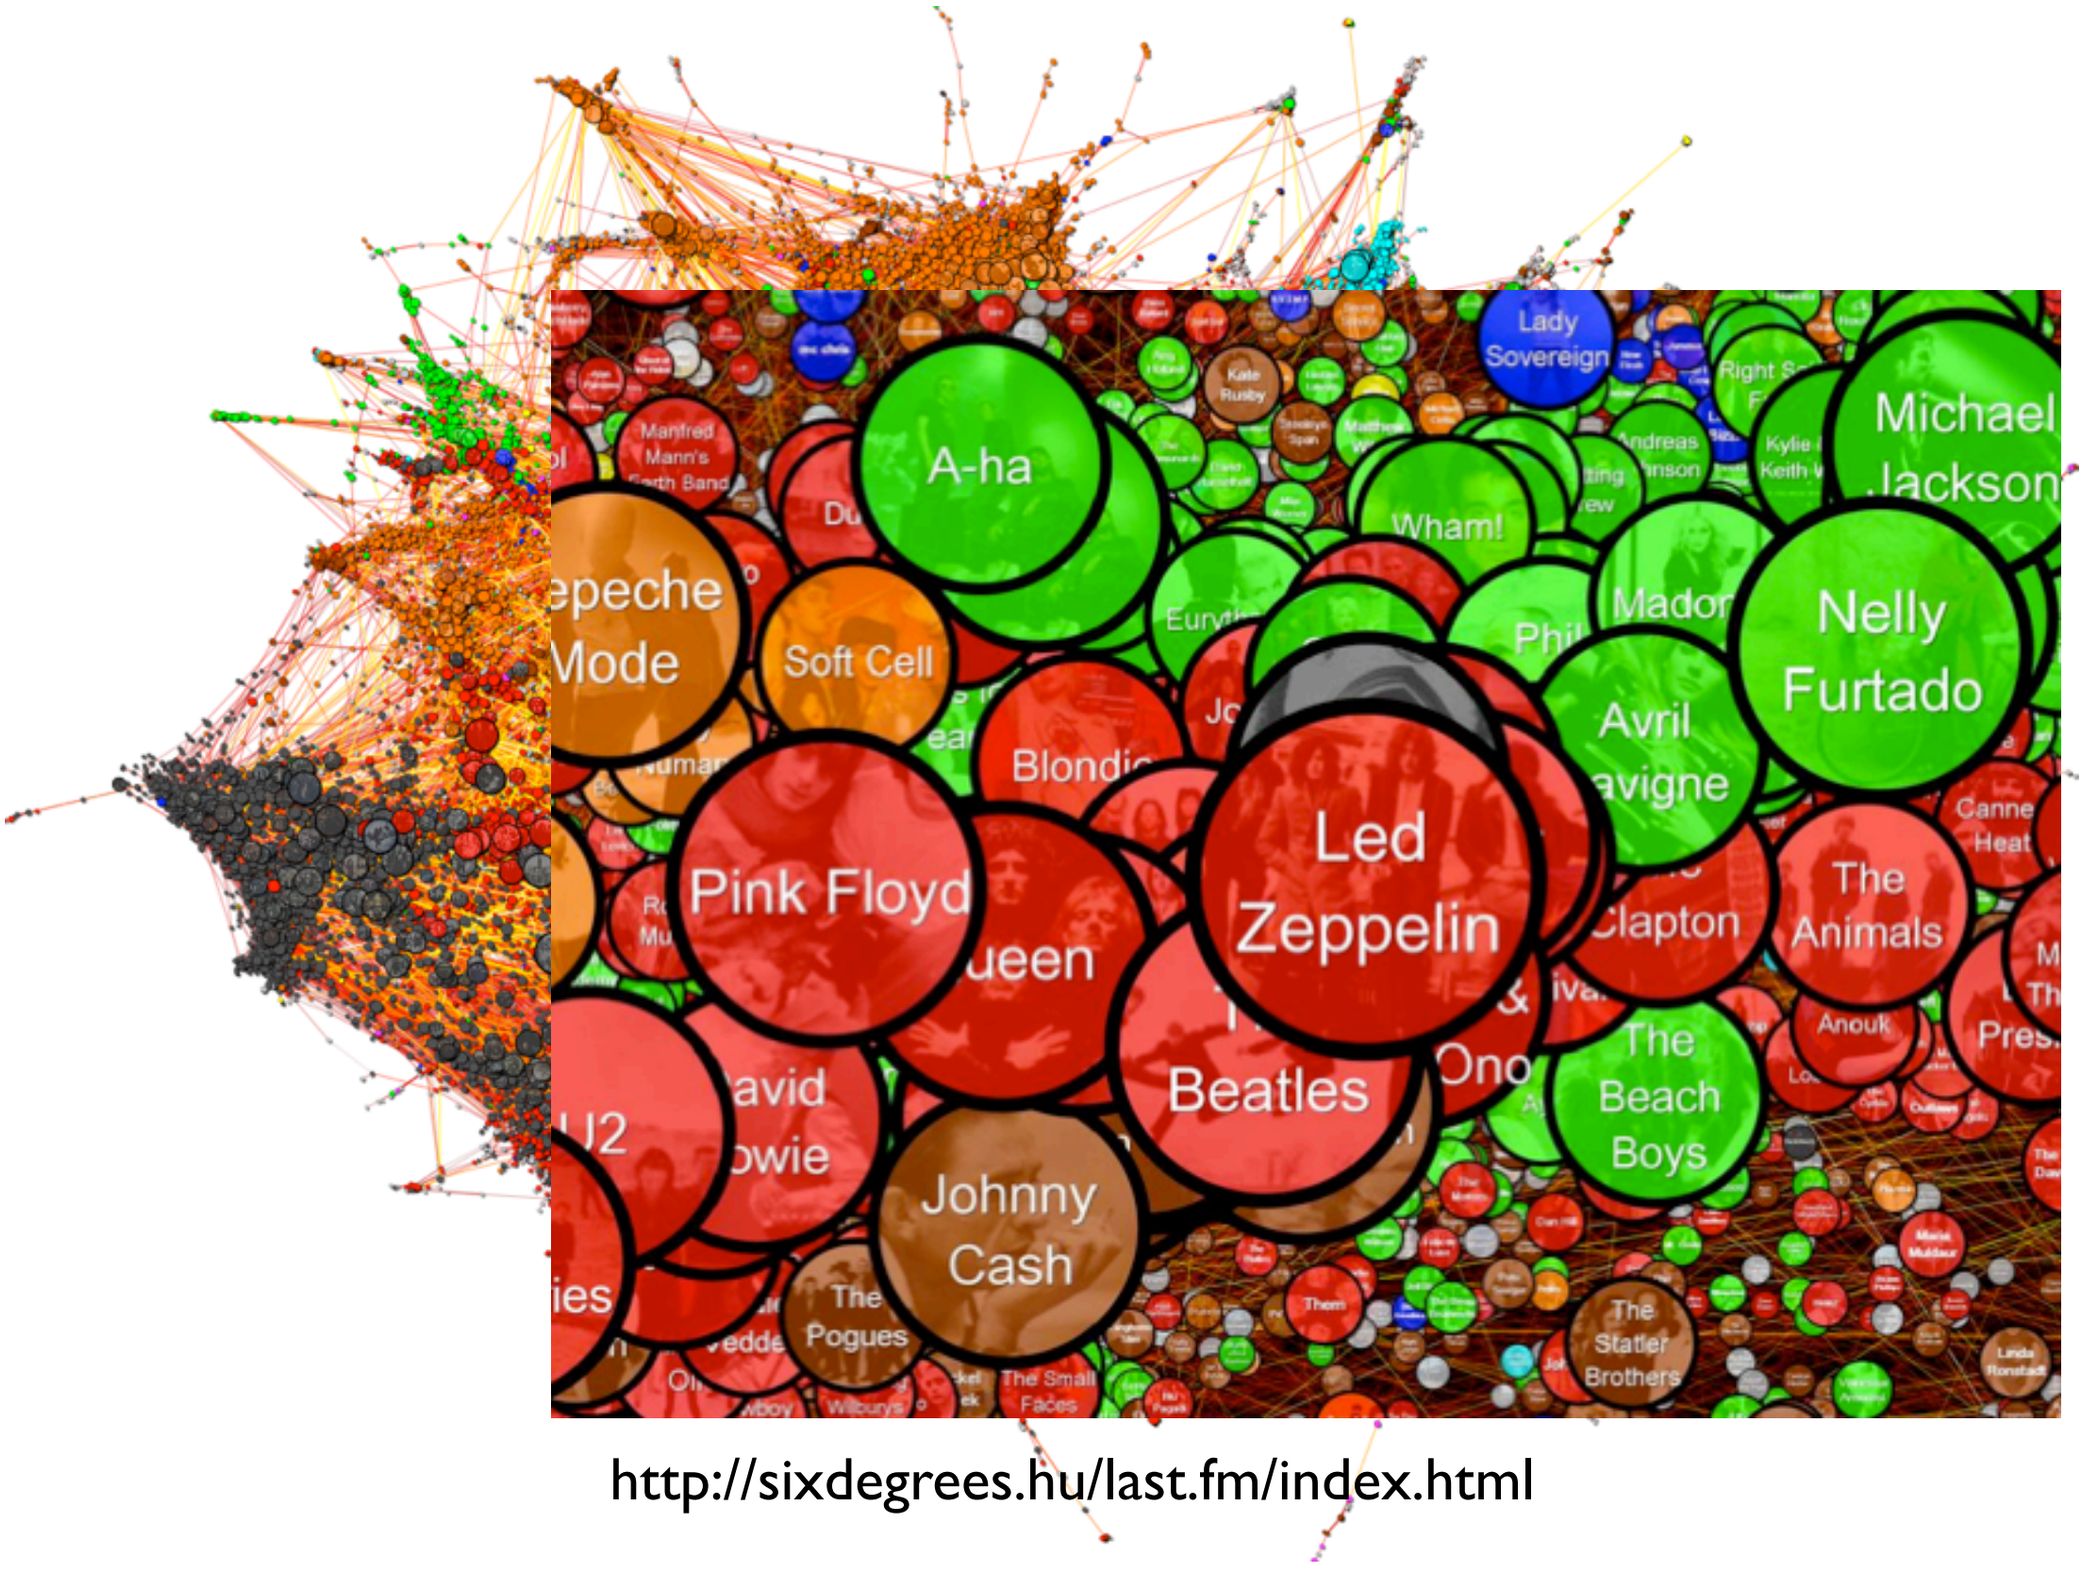

http://sixdegrees.hu/last.fm/index.html

Visualizations of the connections amongst geographically located iTunes libraries

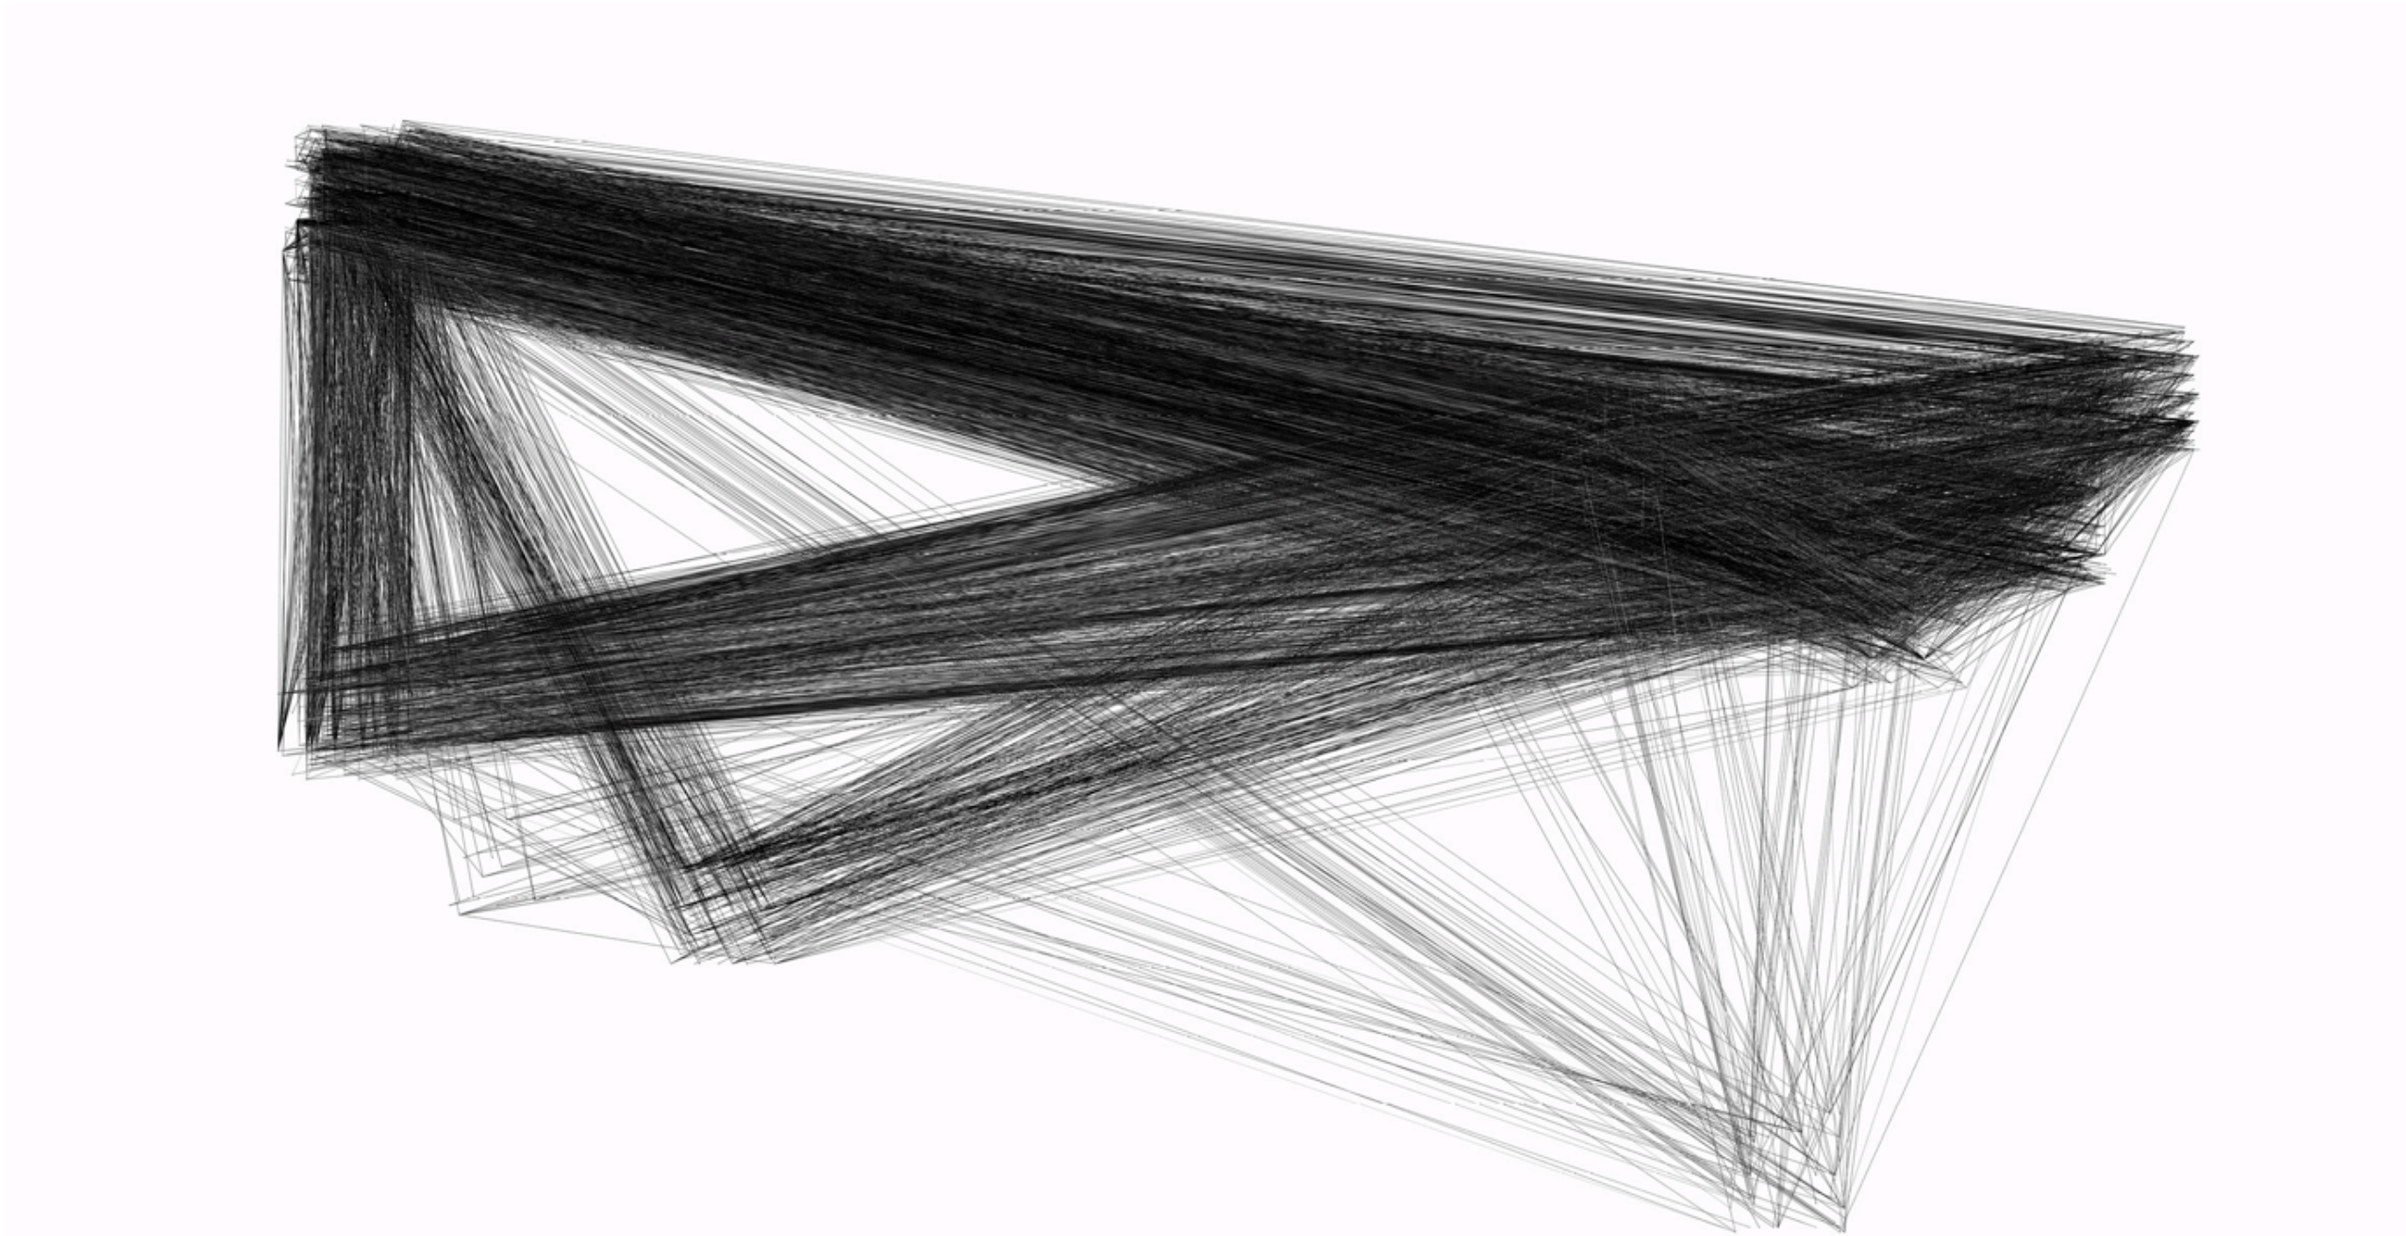

http://caleblarsen.com/projects/visualization-of-the-inherent-connections-/#0

Visualizations of the connections amongst geographically located iTunes libraries

http://caleblarsen.com/projects/visualization-of-the-inherent-connections-/#0

Maps


Quaye
Damien
'Junior Gong'
Marley
The
Fugees
Ziggy
Marley
RZA
Everly Brothers
The Kinks
Portishead
David
Bowie
Kanye
West
Blur
The Cure
OutKast

Maps

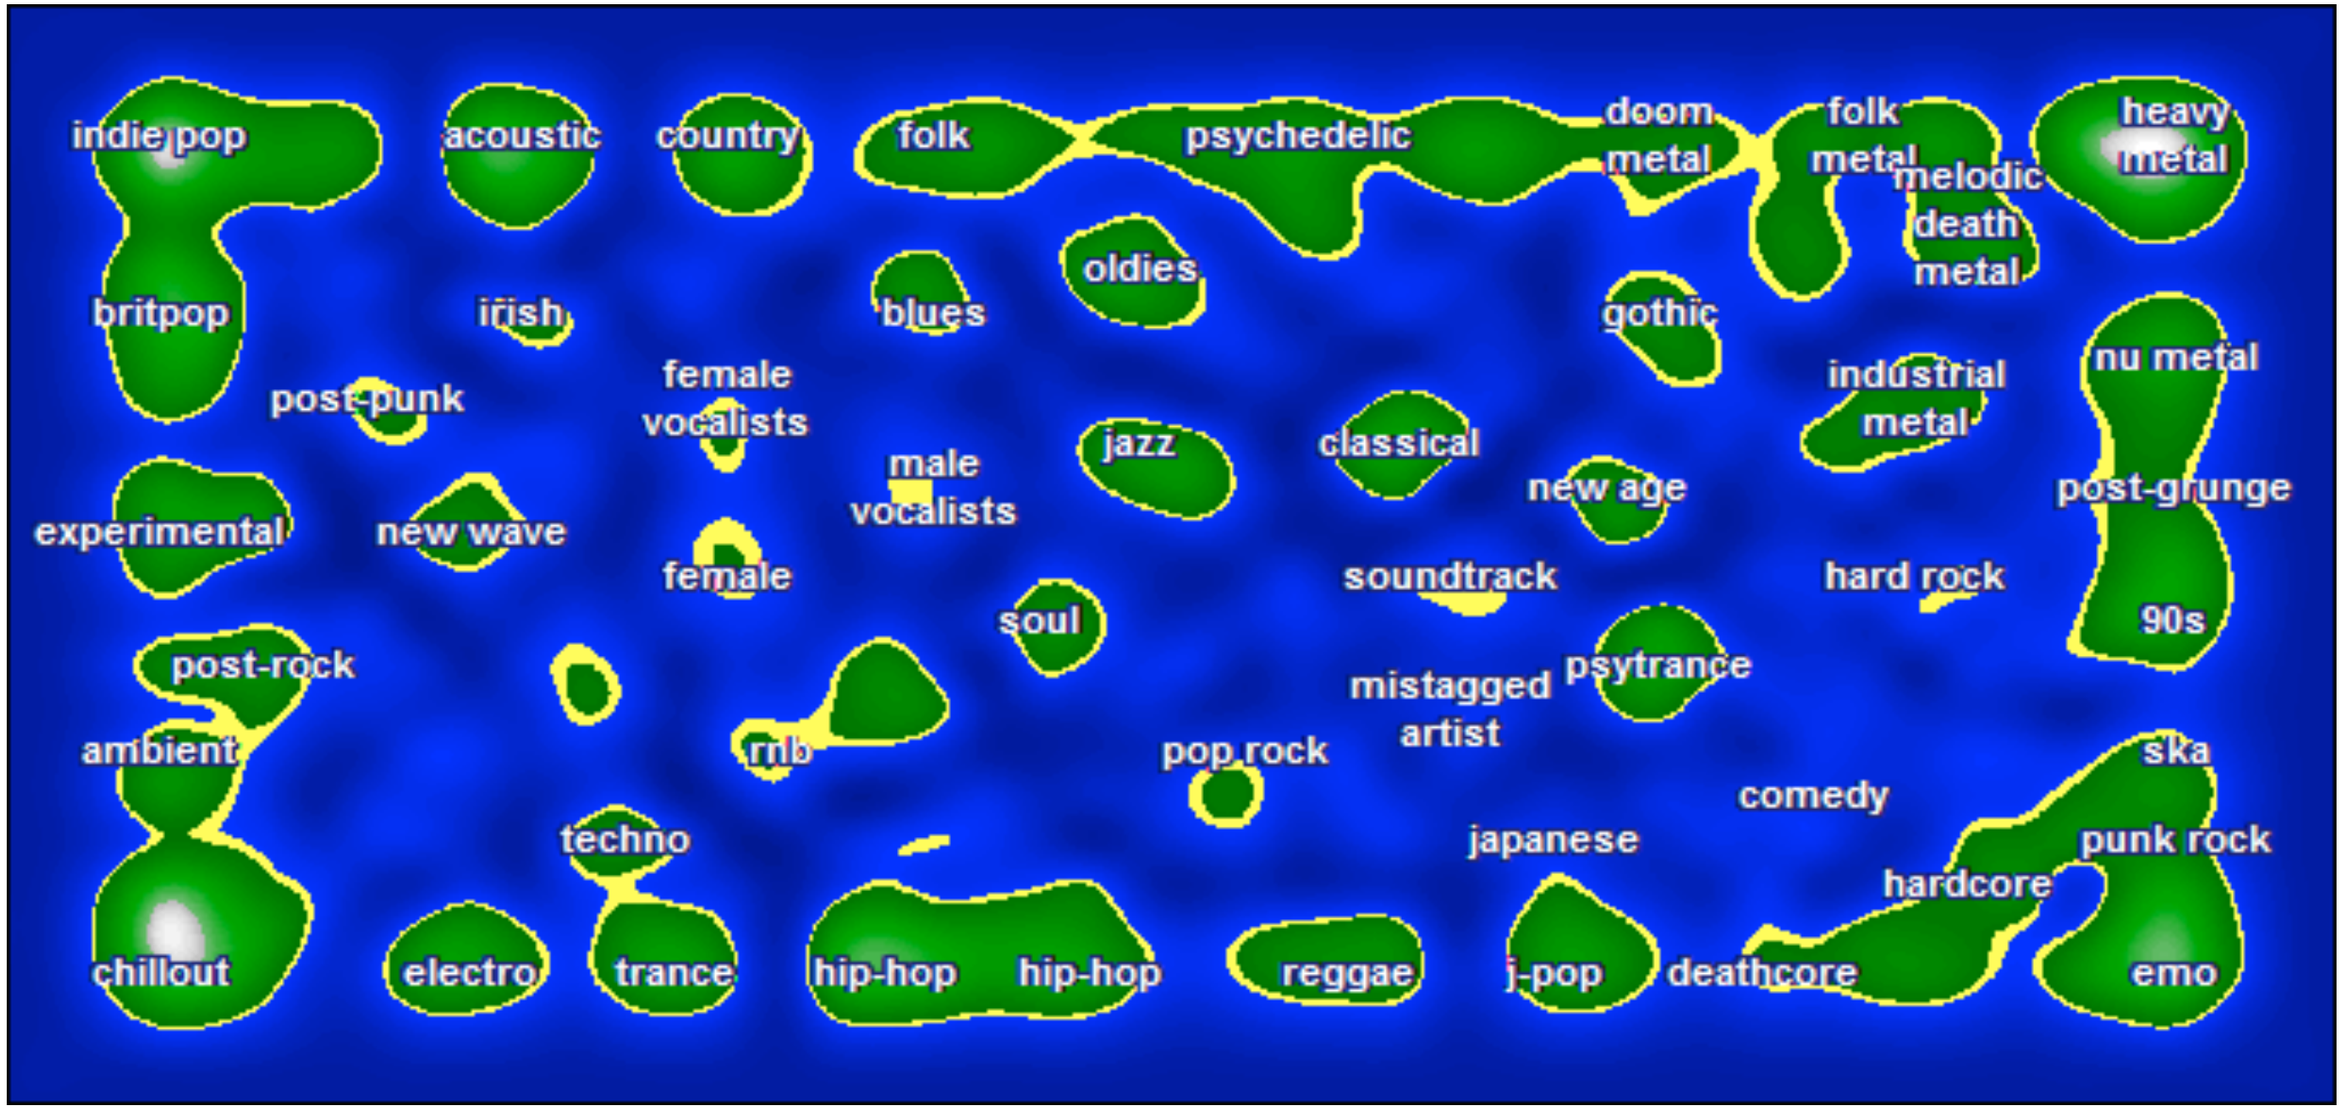

http://playground.last.fm/iom - Islands of Music - Elias Pampalk

marGrid

http://margrid.sness.net/

PlaySOM - Intuitive access to Music Archives

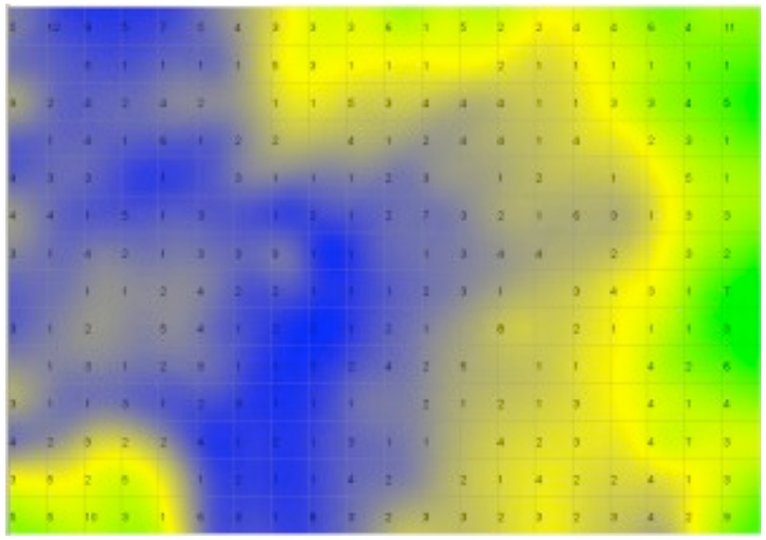

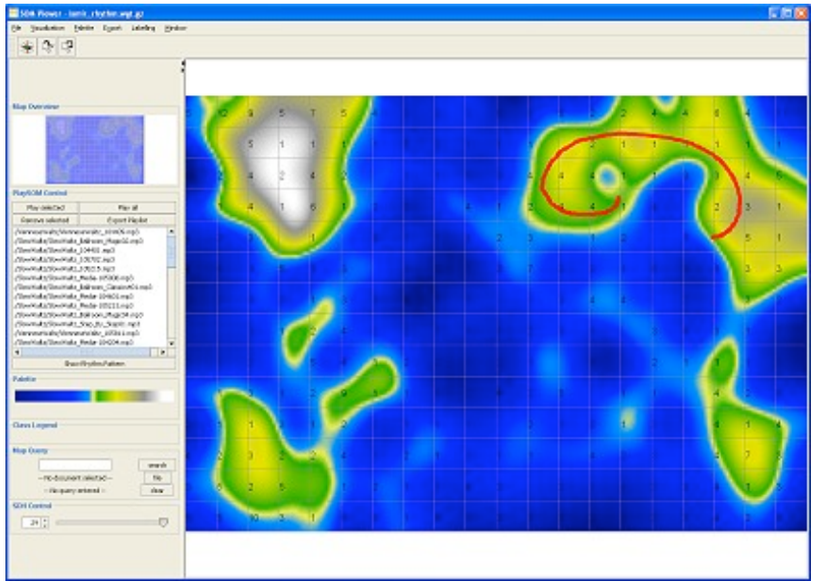

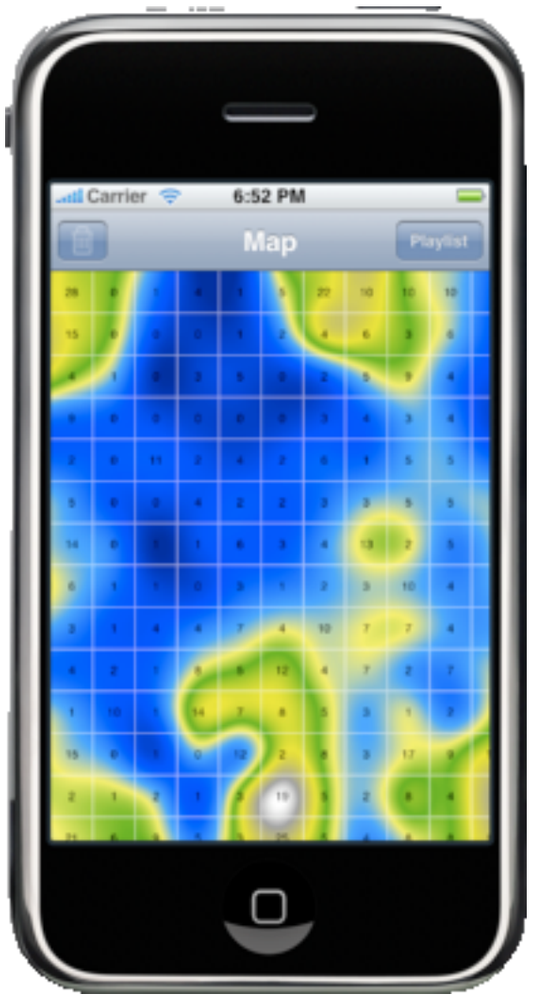

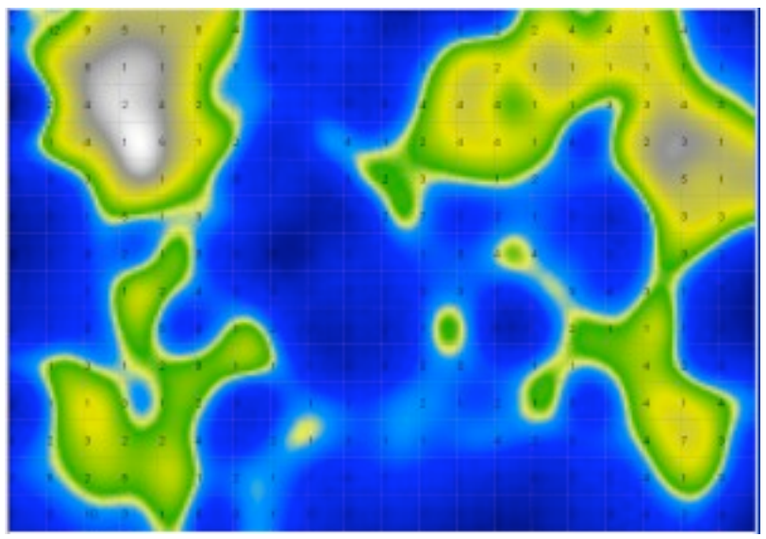

http://www.ifs.tuwien.ac.at/mir/playsom.html

Vienna University of Technology

FM4 Soundpark

http://fm4.orf.at/soundpark Martin Gasser, Arthur Flexer

nepTune

http://www.cp.jku.at/projects/nepTune/

nepTune

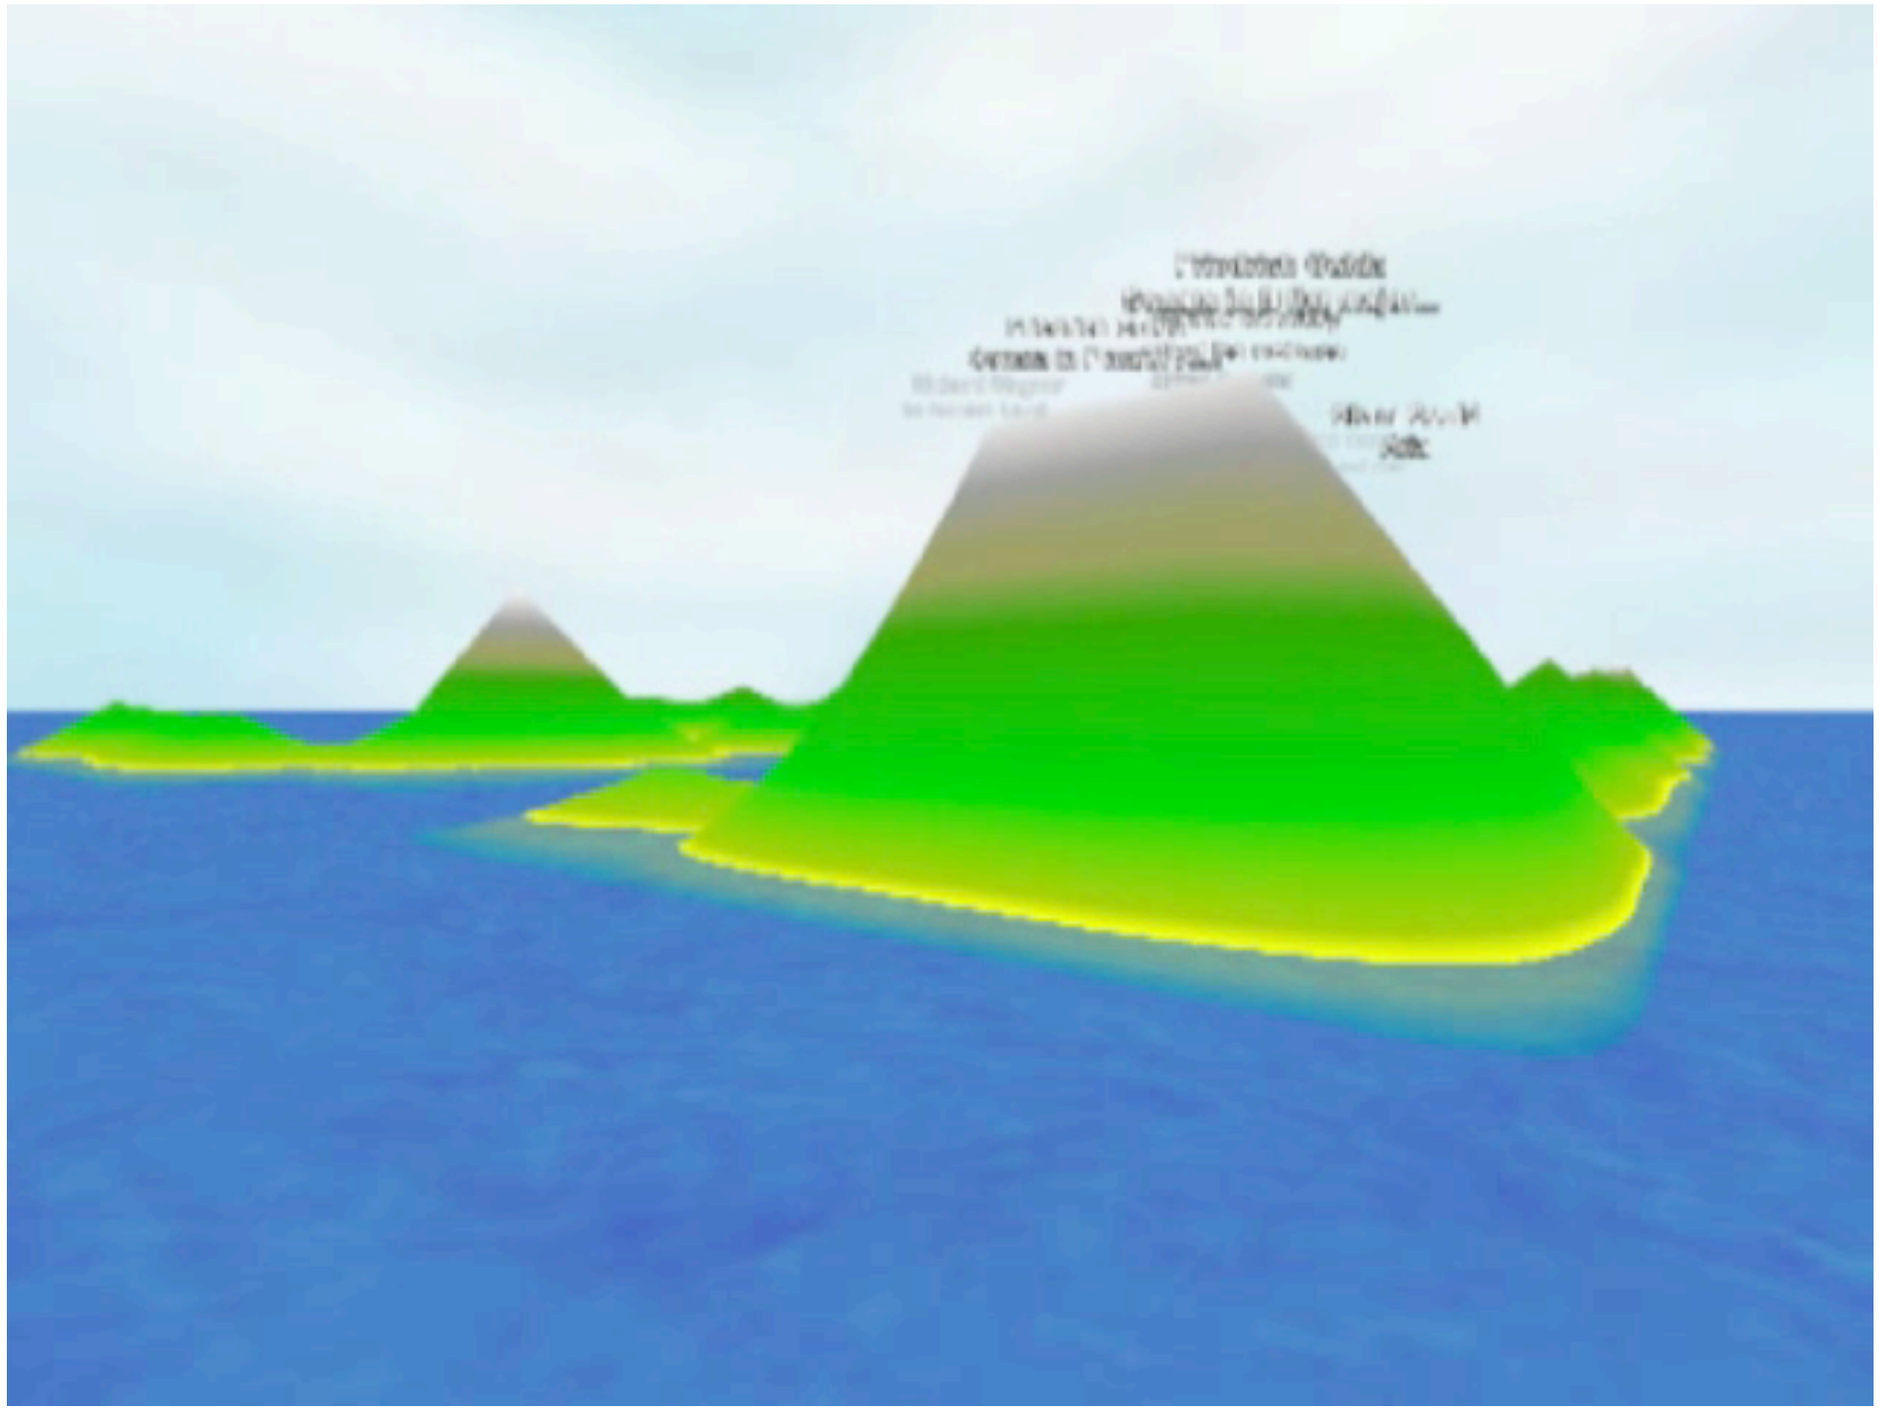

http://www.cp.jku.at/projects/nepTune/

Databionic Music Miner

http://musicminer.sourceforge.net/

Databionic Music Miner

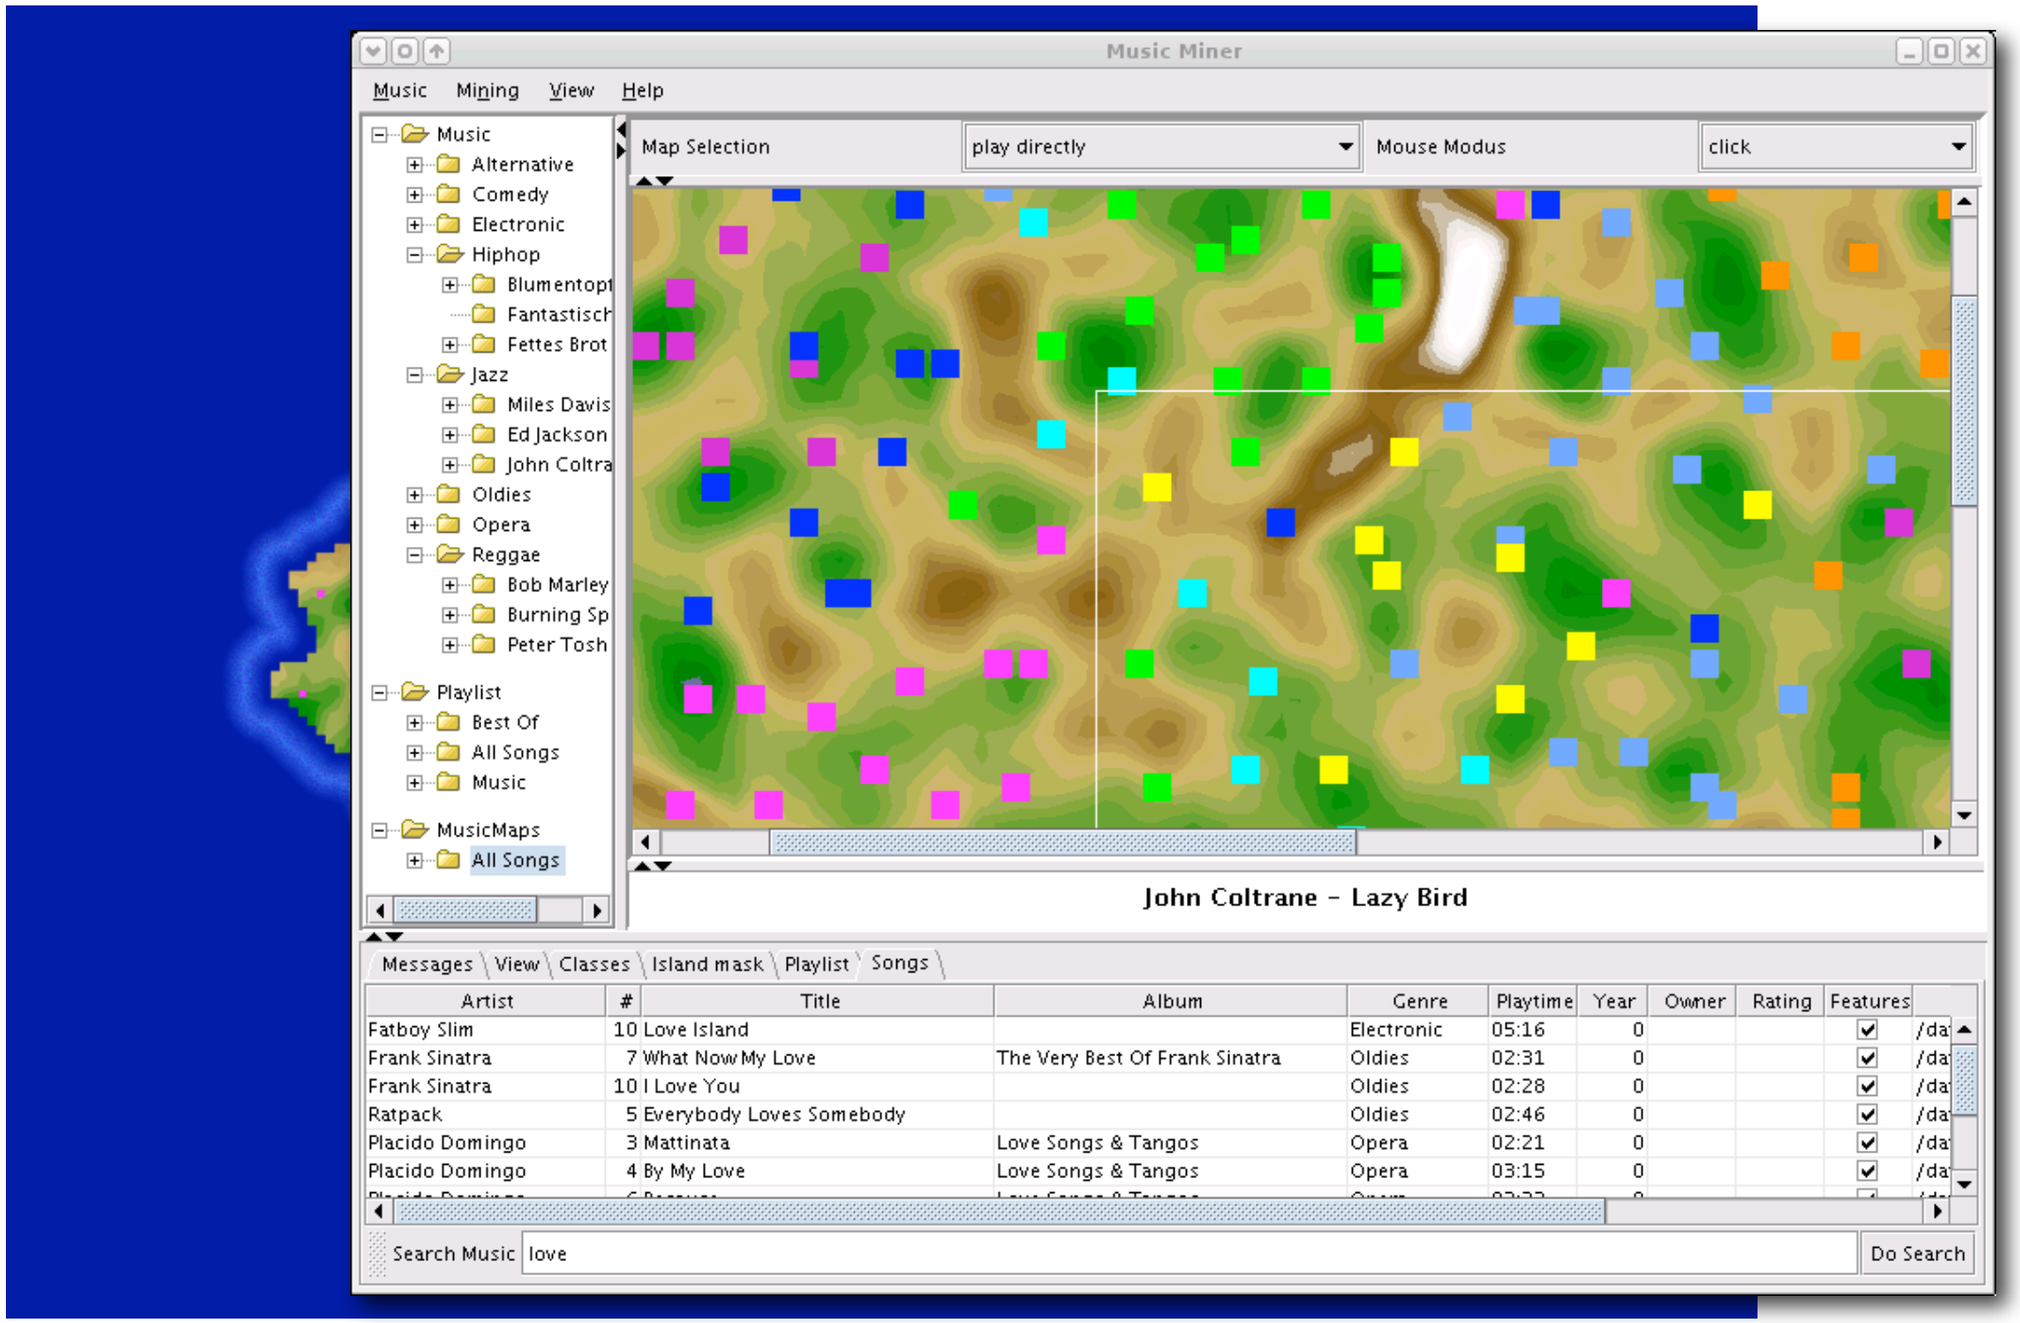

http://musicminer.sourceforge.net/

soniXplorer

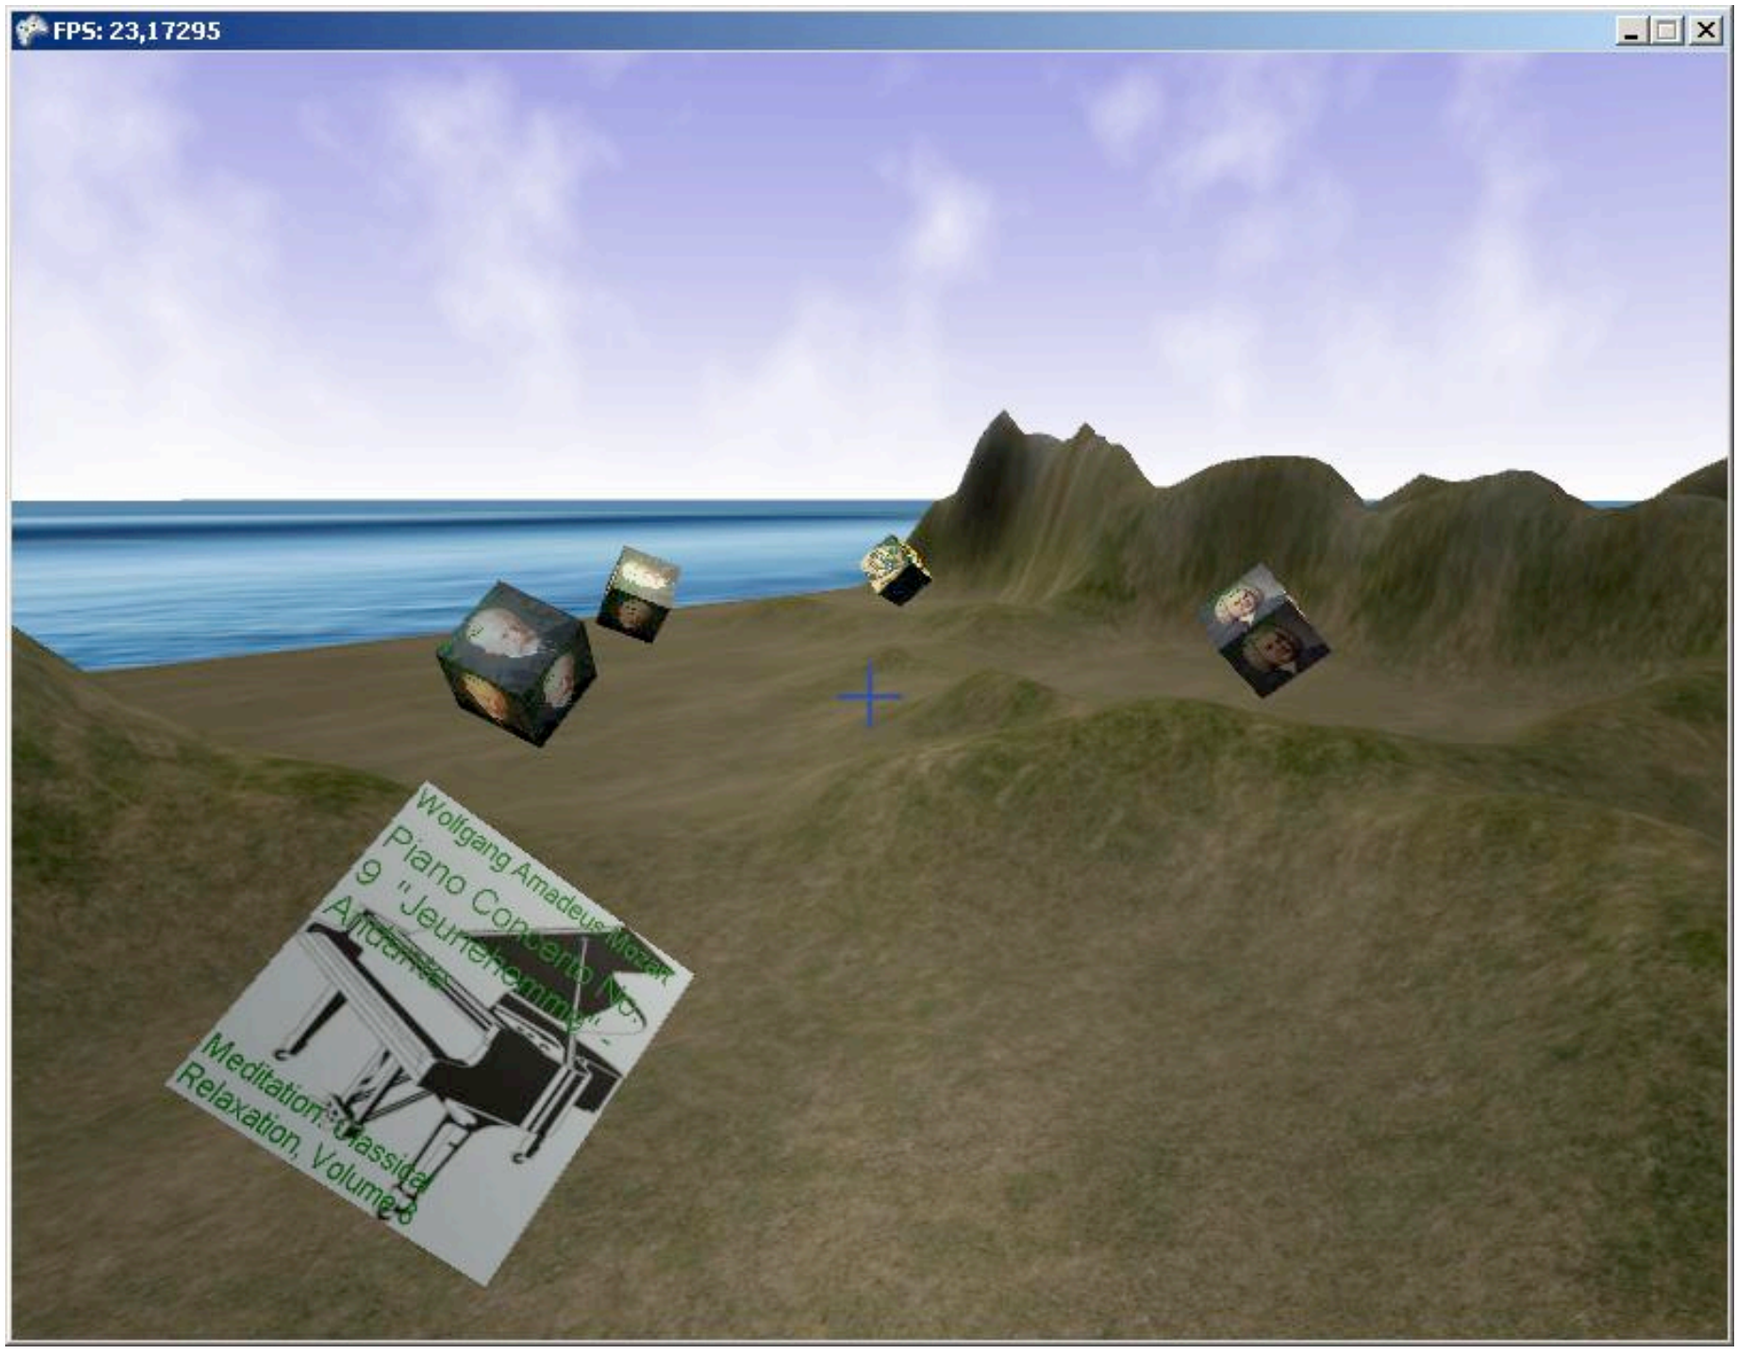

Presented at ISMIR 2009 by Dominik Lübbers

soniXplorer

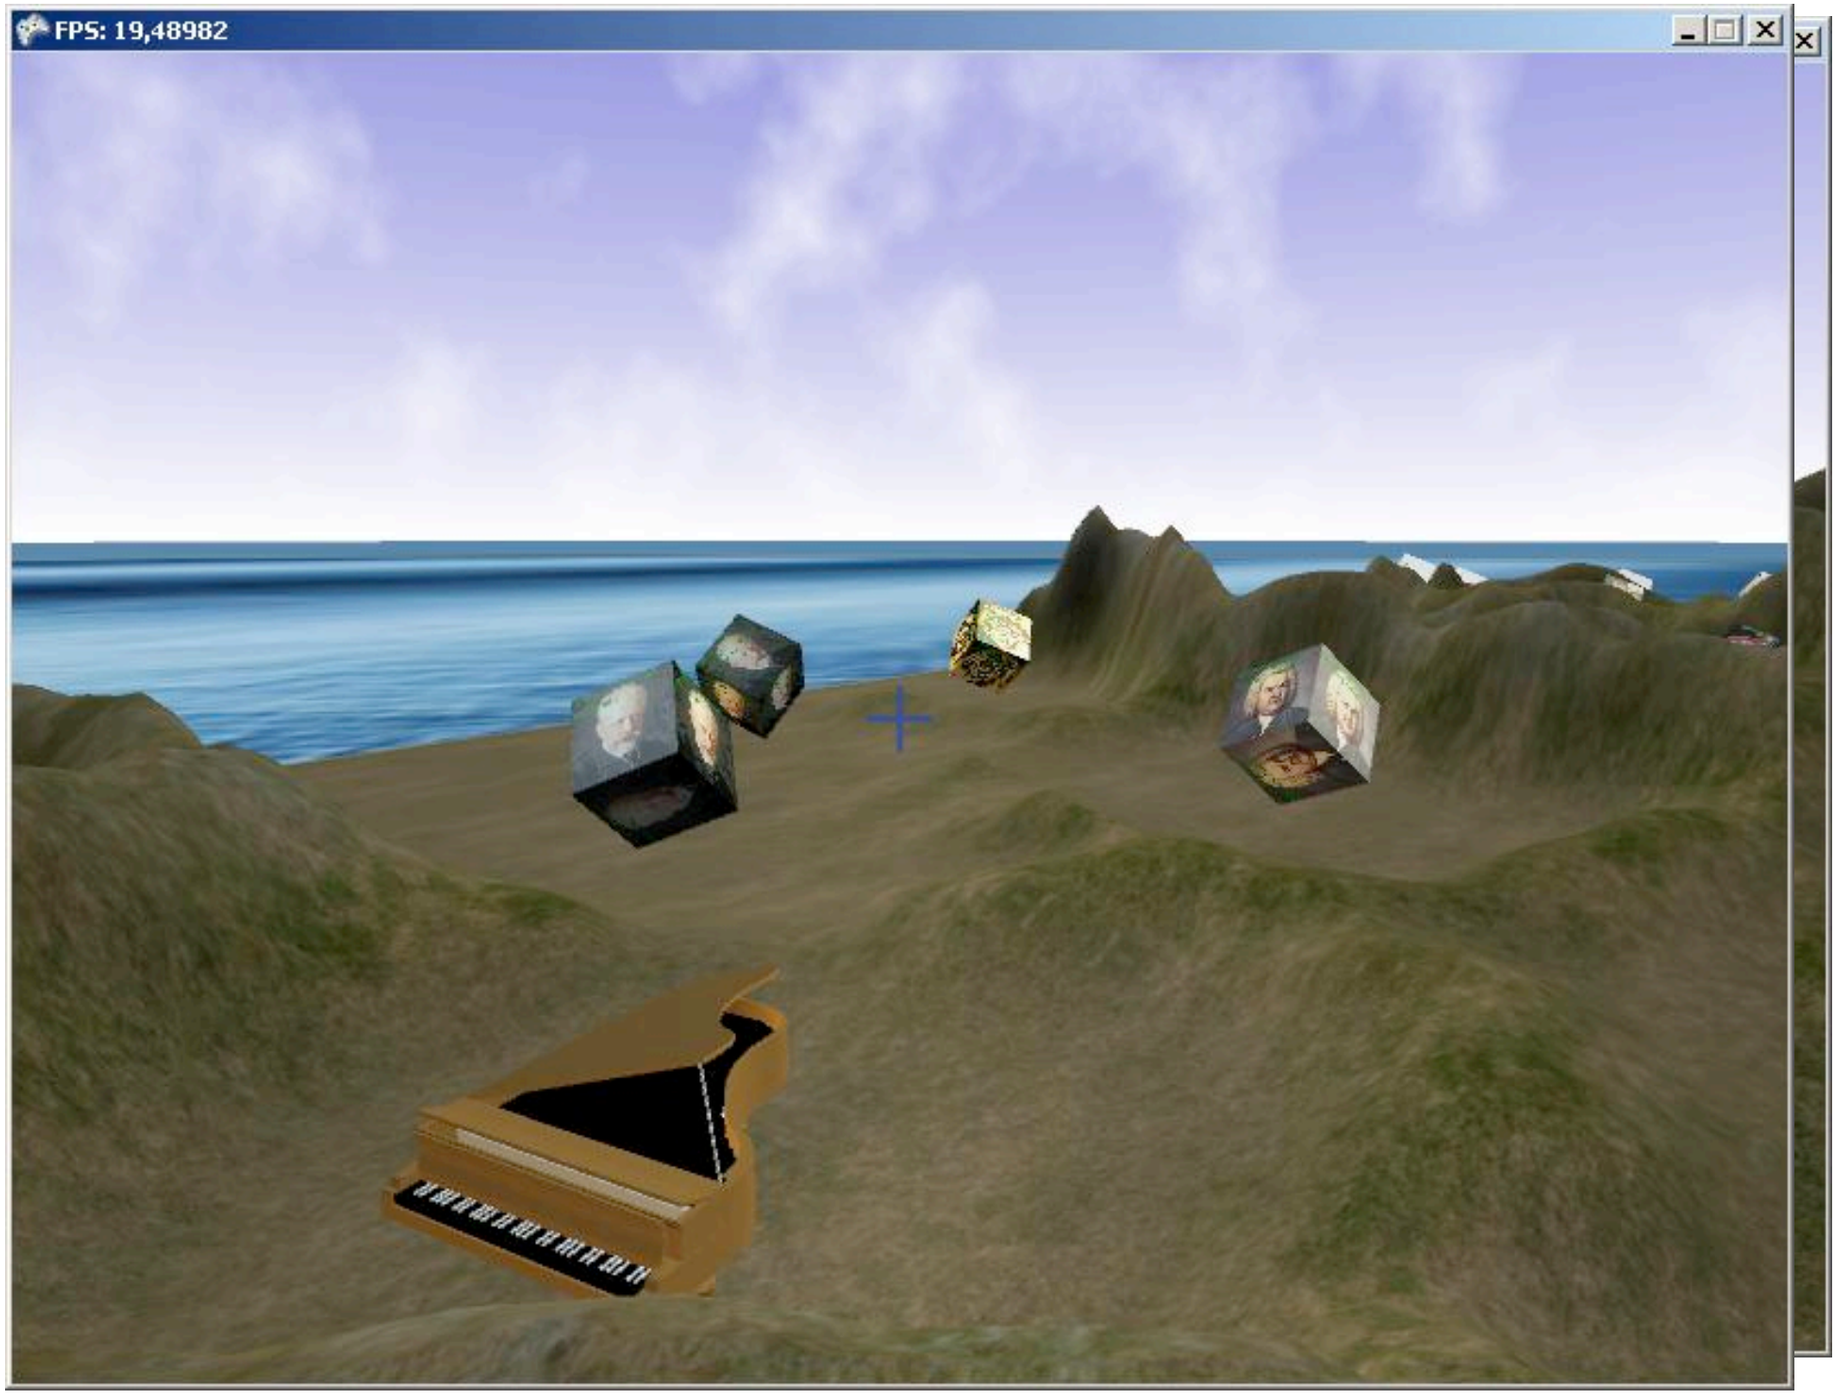

Presented at ISMIR 2009 by Dominik Lübbers

soniXplorer

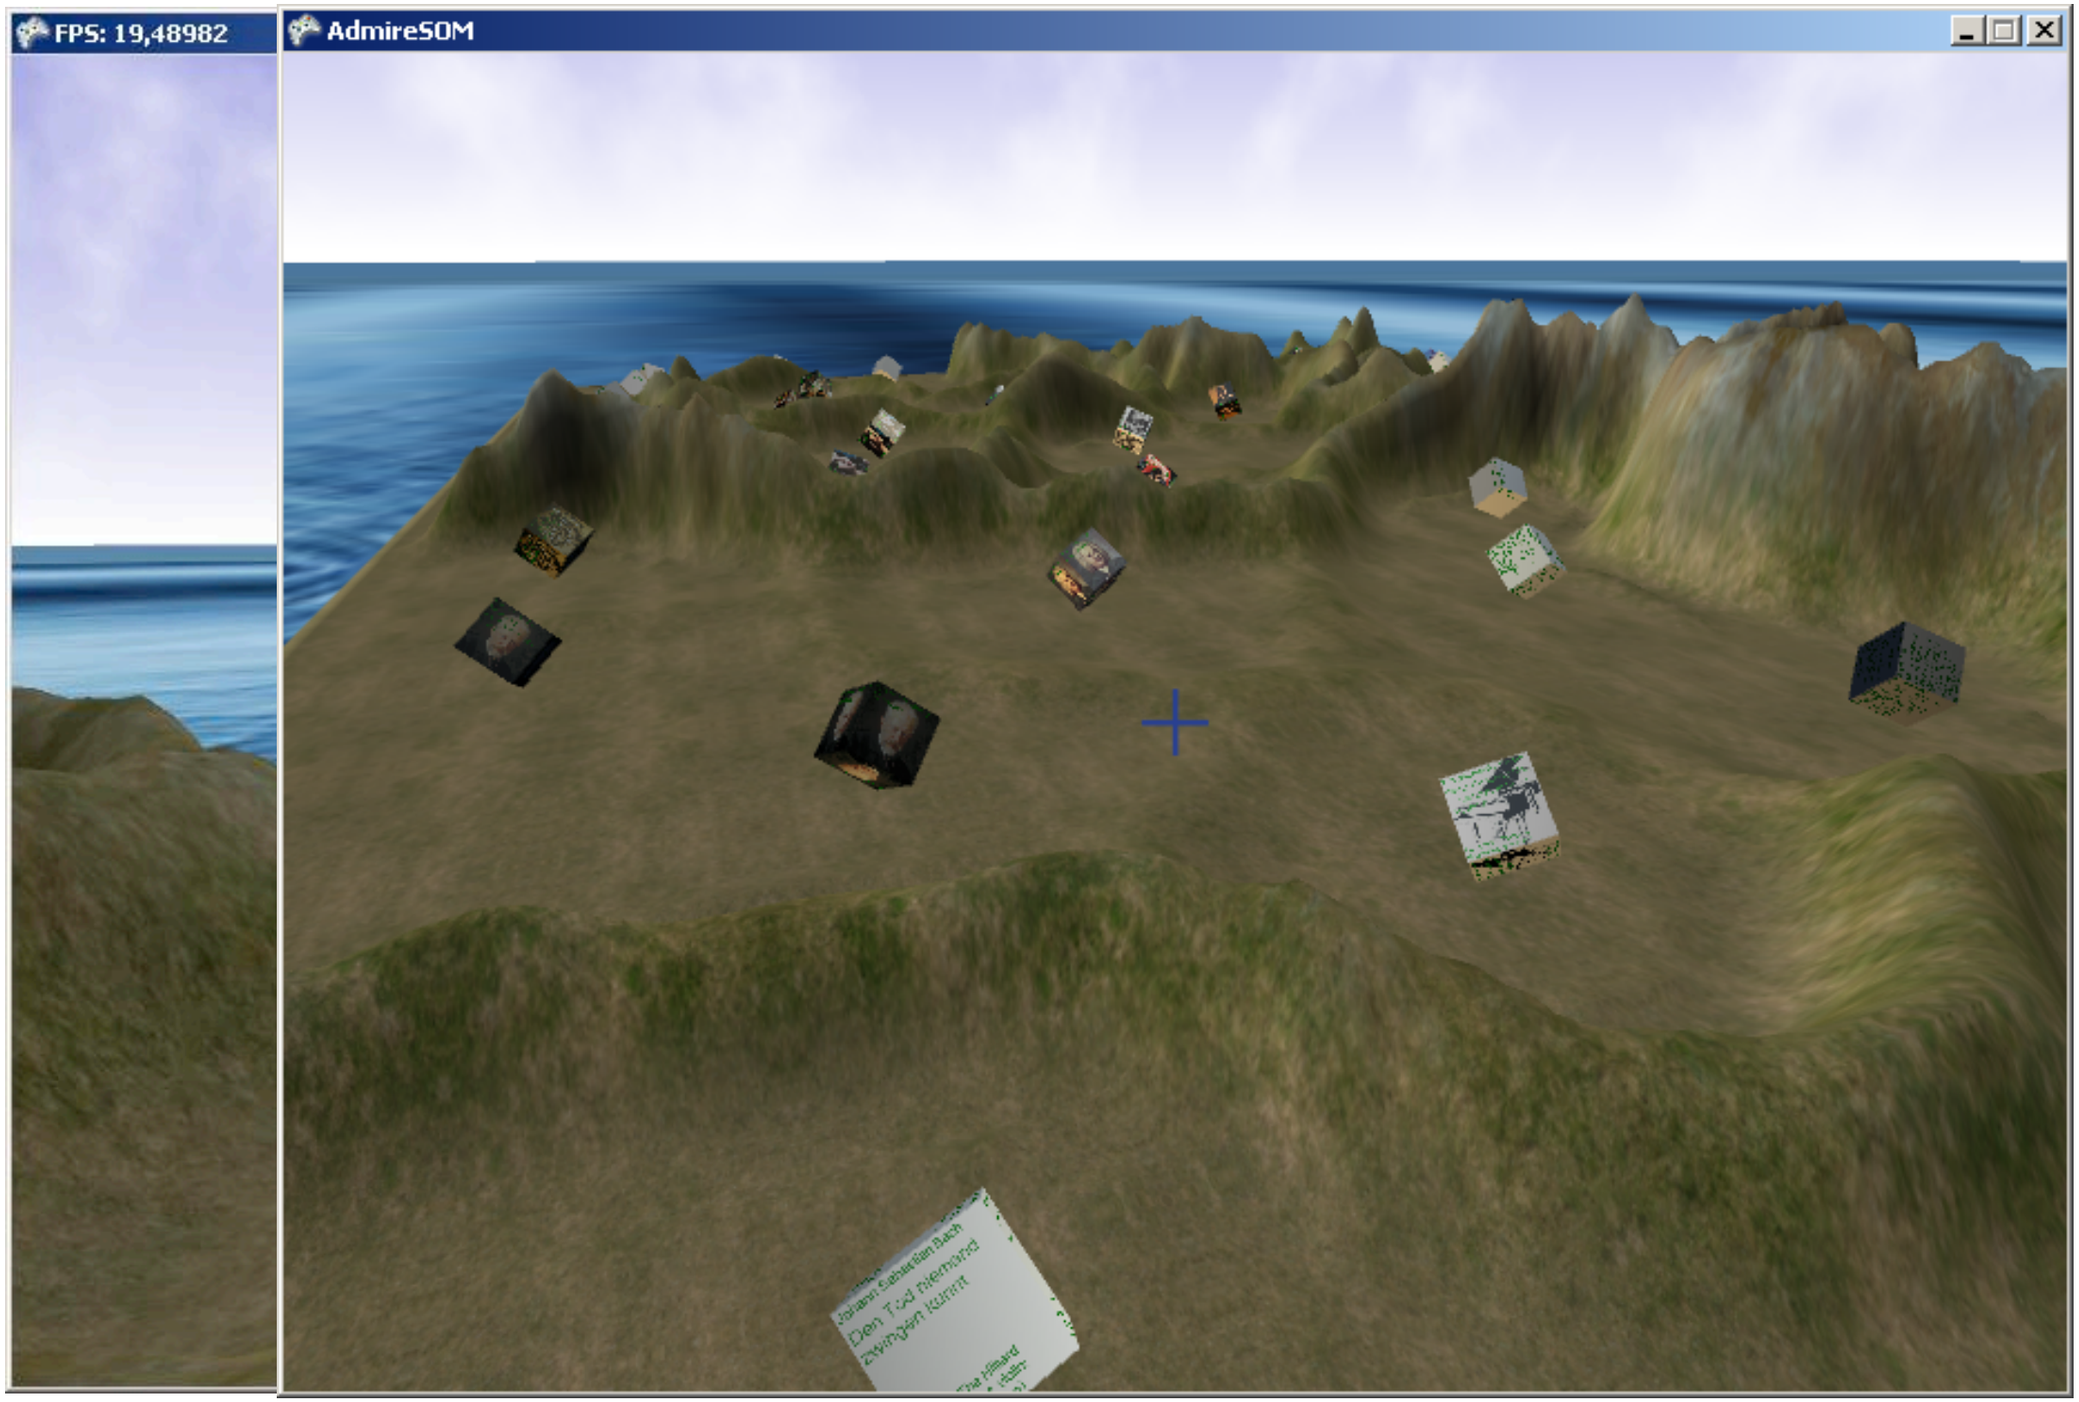

Presented at ISMIR 2009 by Dominik Lübbers

soniXplorer

Presented at ISMIR 2009 by Dominik Lübbers

mHashup

http://www.mhashup.com/

mHashup

http://www.mhashup.com/

GeoMuzik

ISMIR 2008 - Oscar Celma

Last on AM/FM

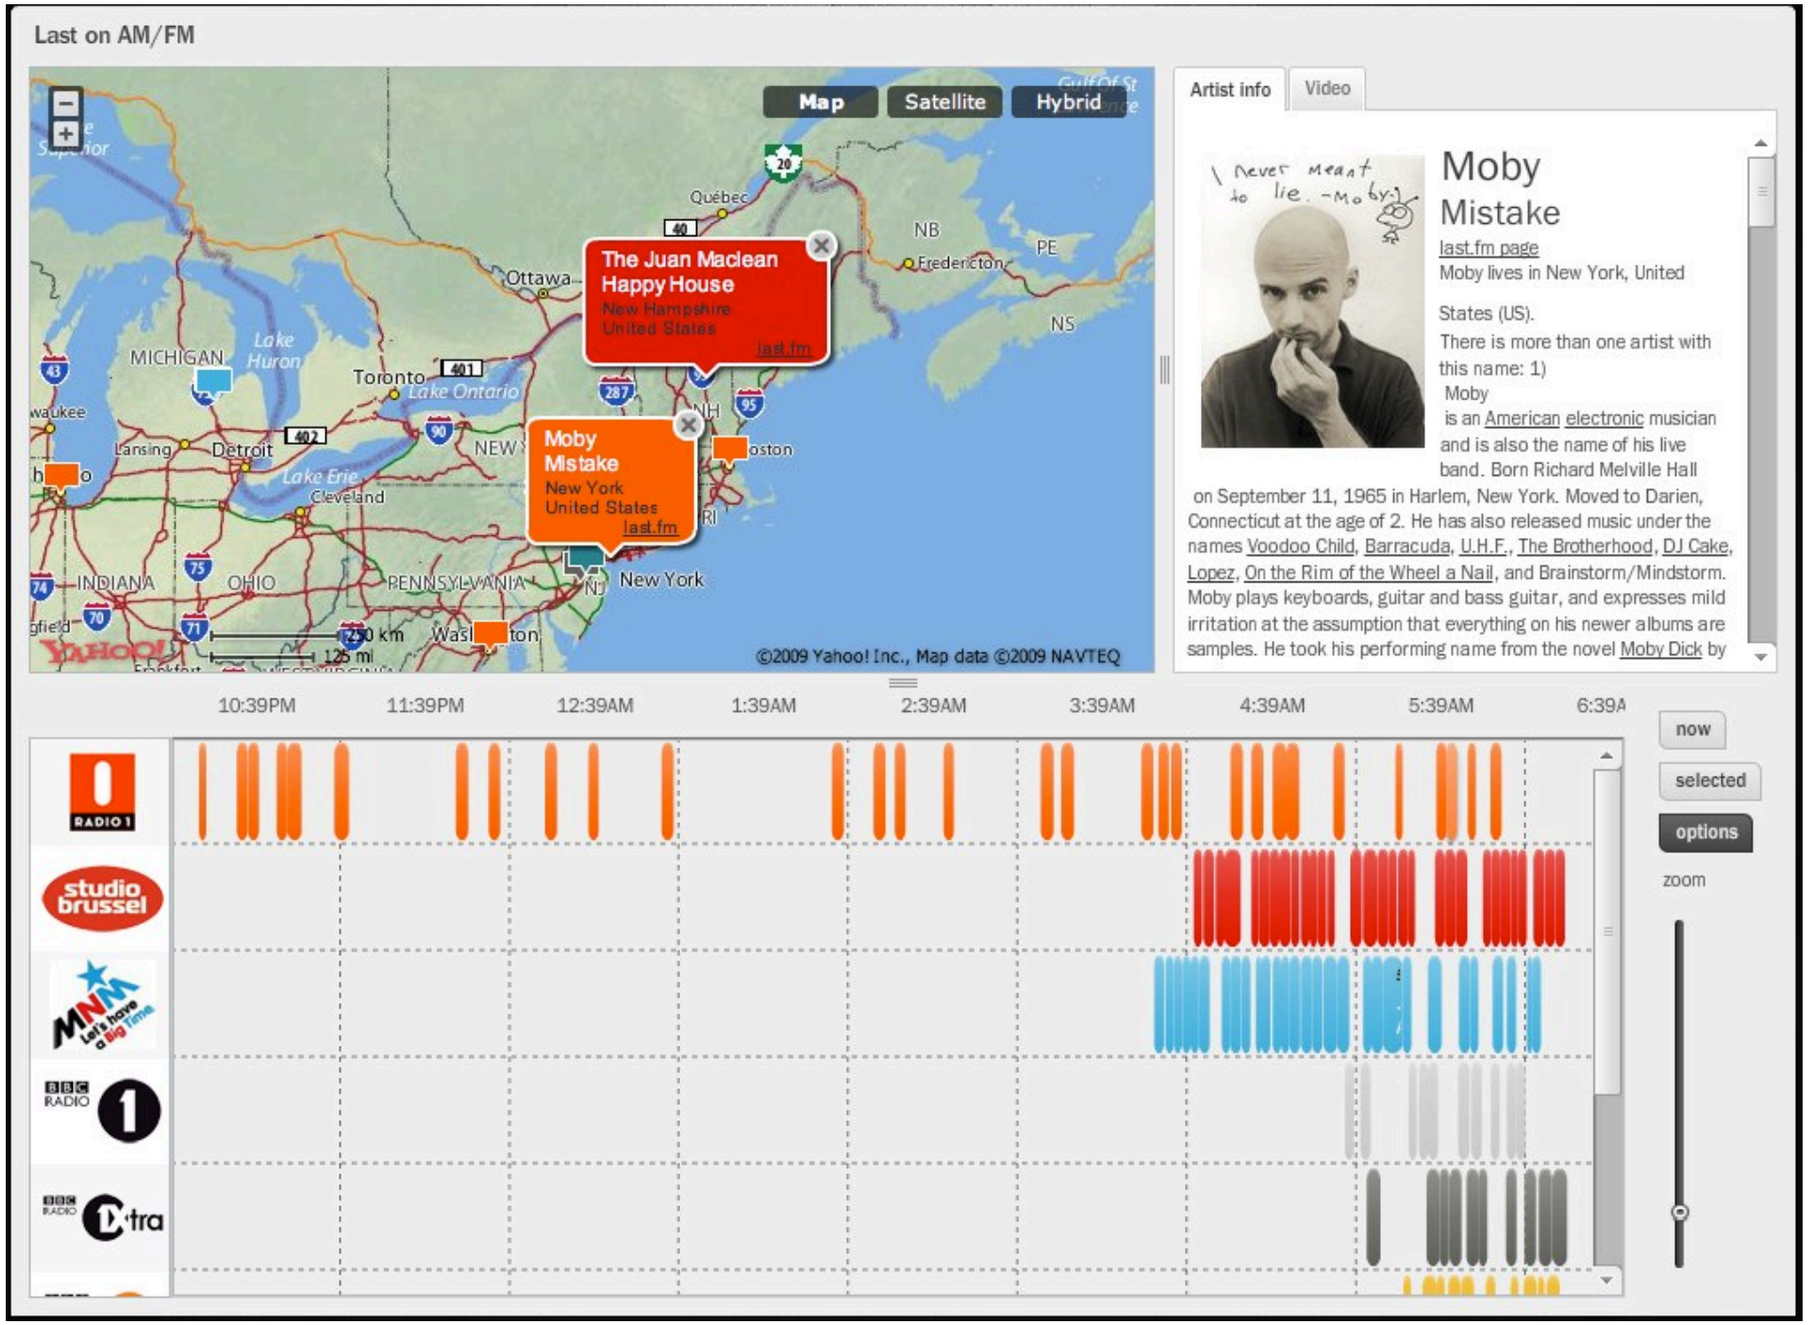

http://www.cs.kuleuven.be/~sten/lastonamfm

Last.fm tube tags

http://playground.last.fm/demo/tagstube

Last.fm tube tags

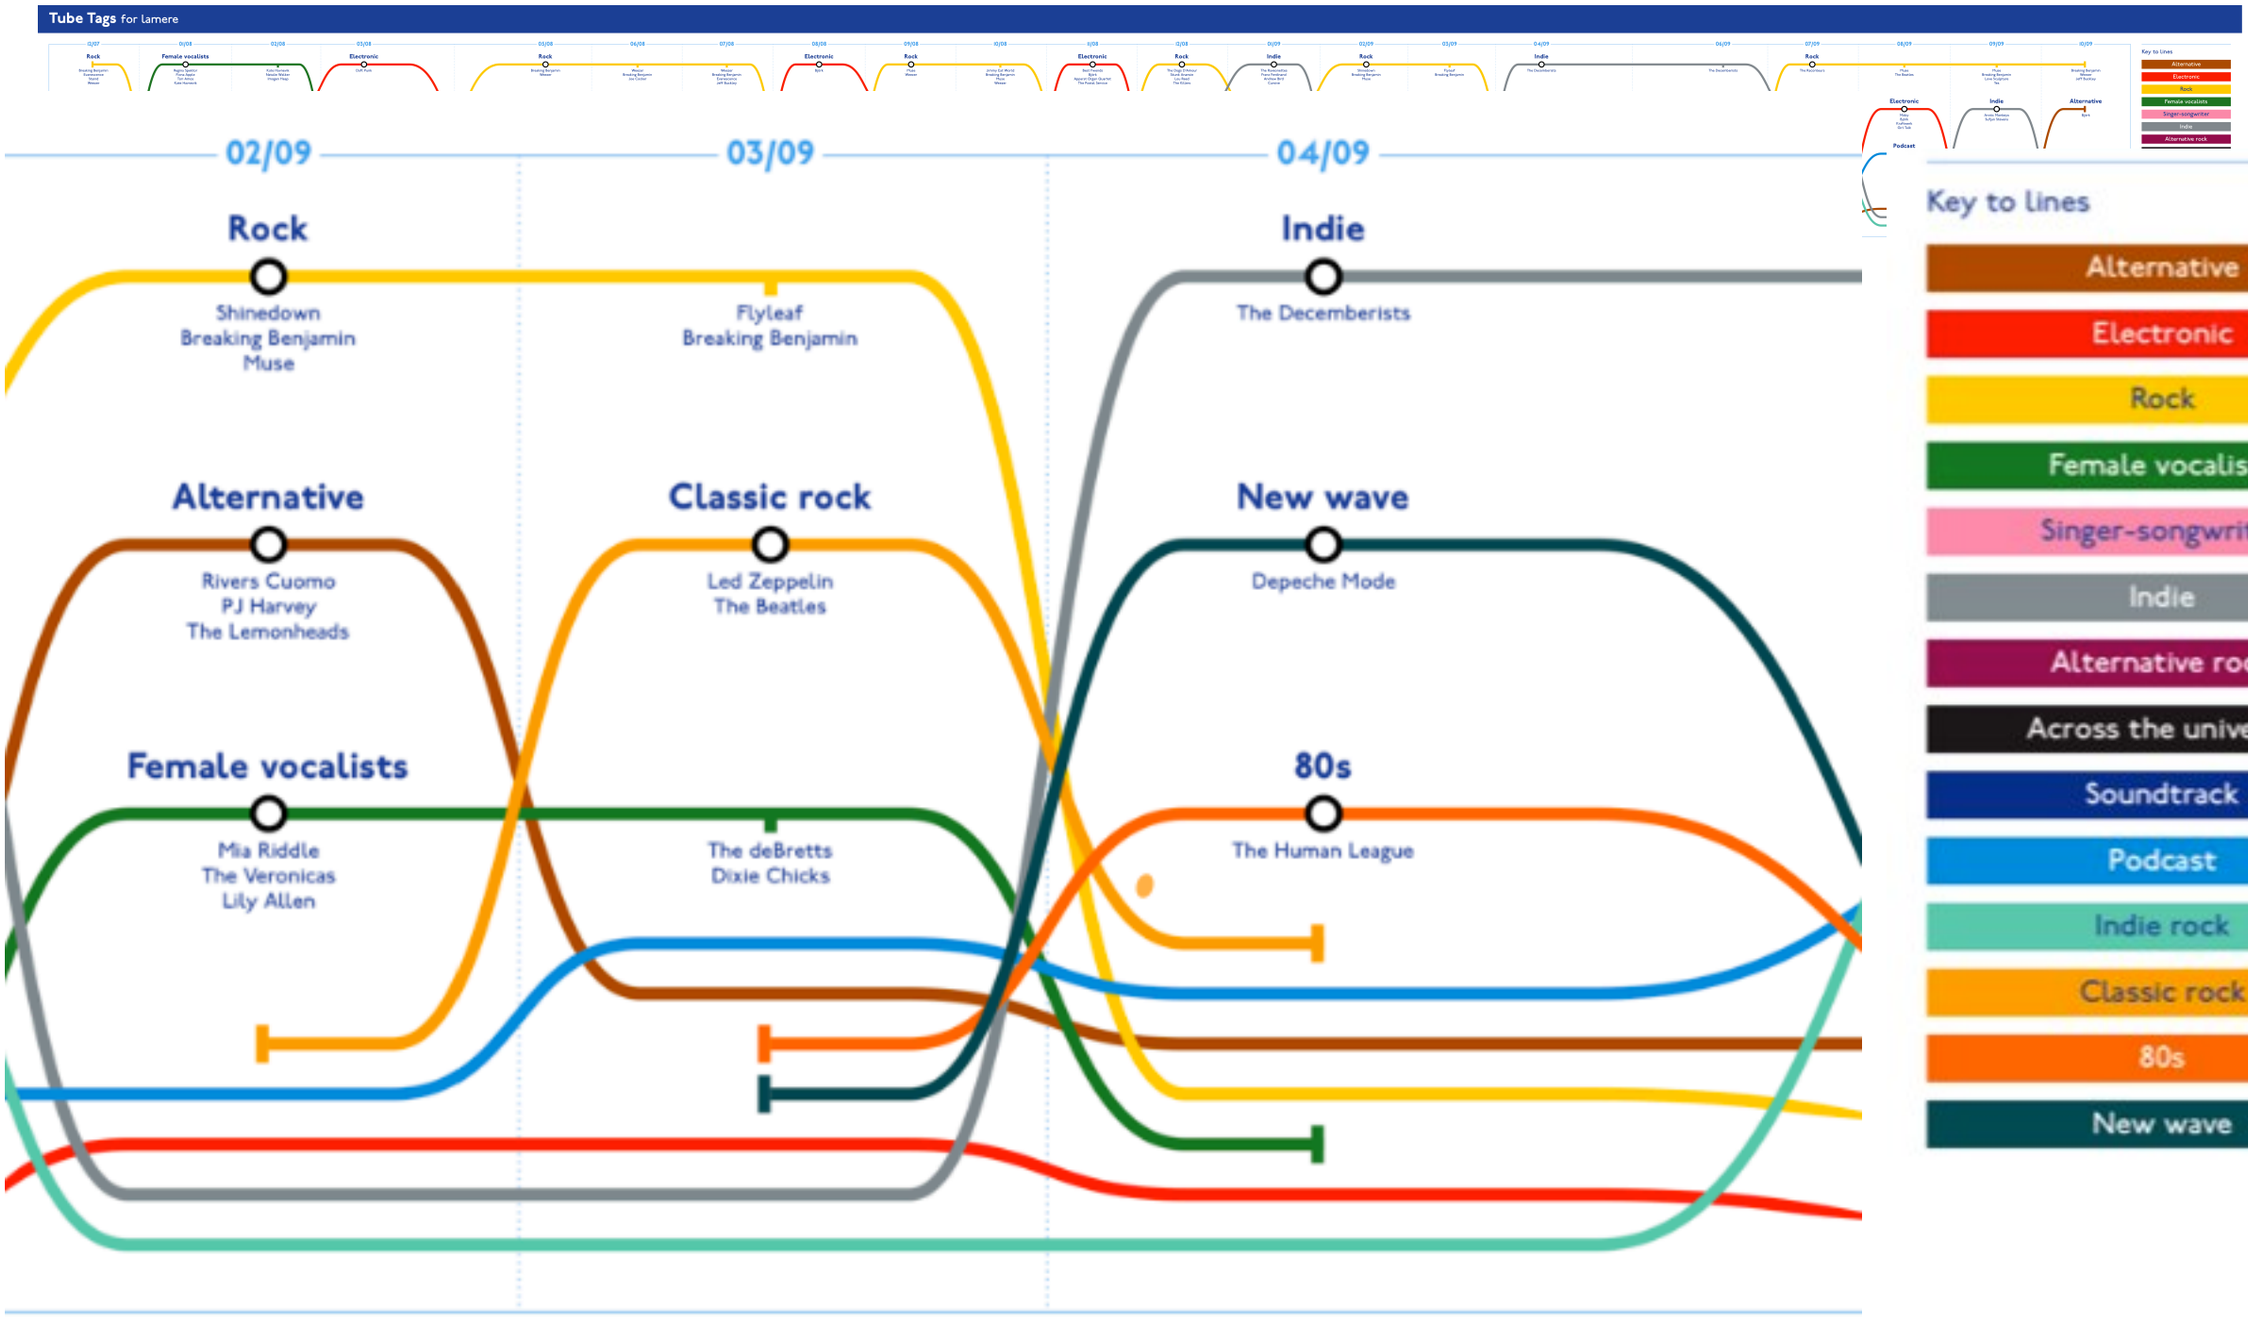

http://playground.last.fm/demo/tagstube

Scatter plots

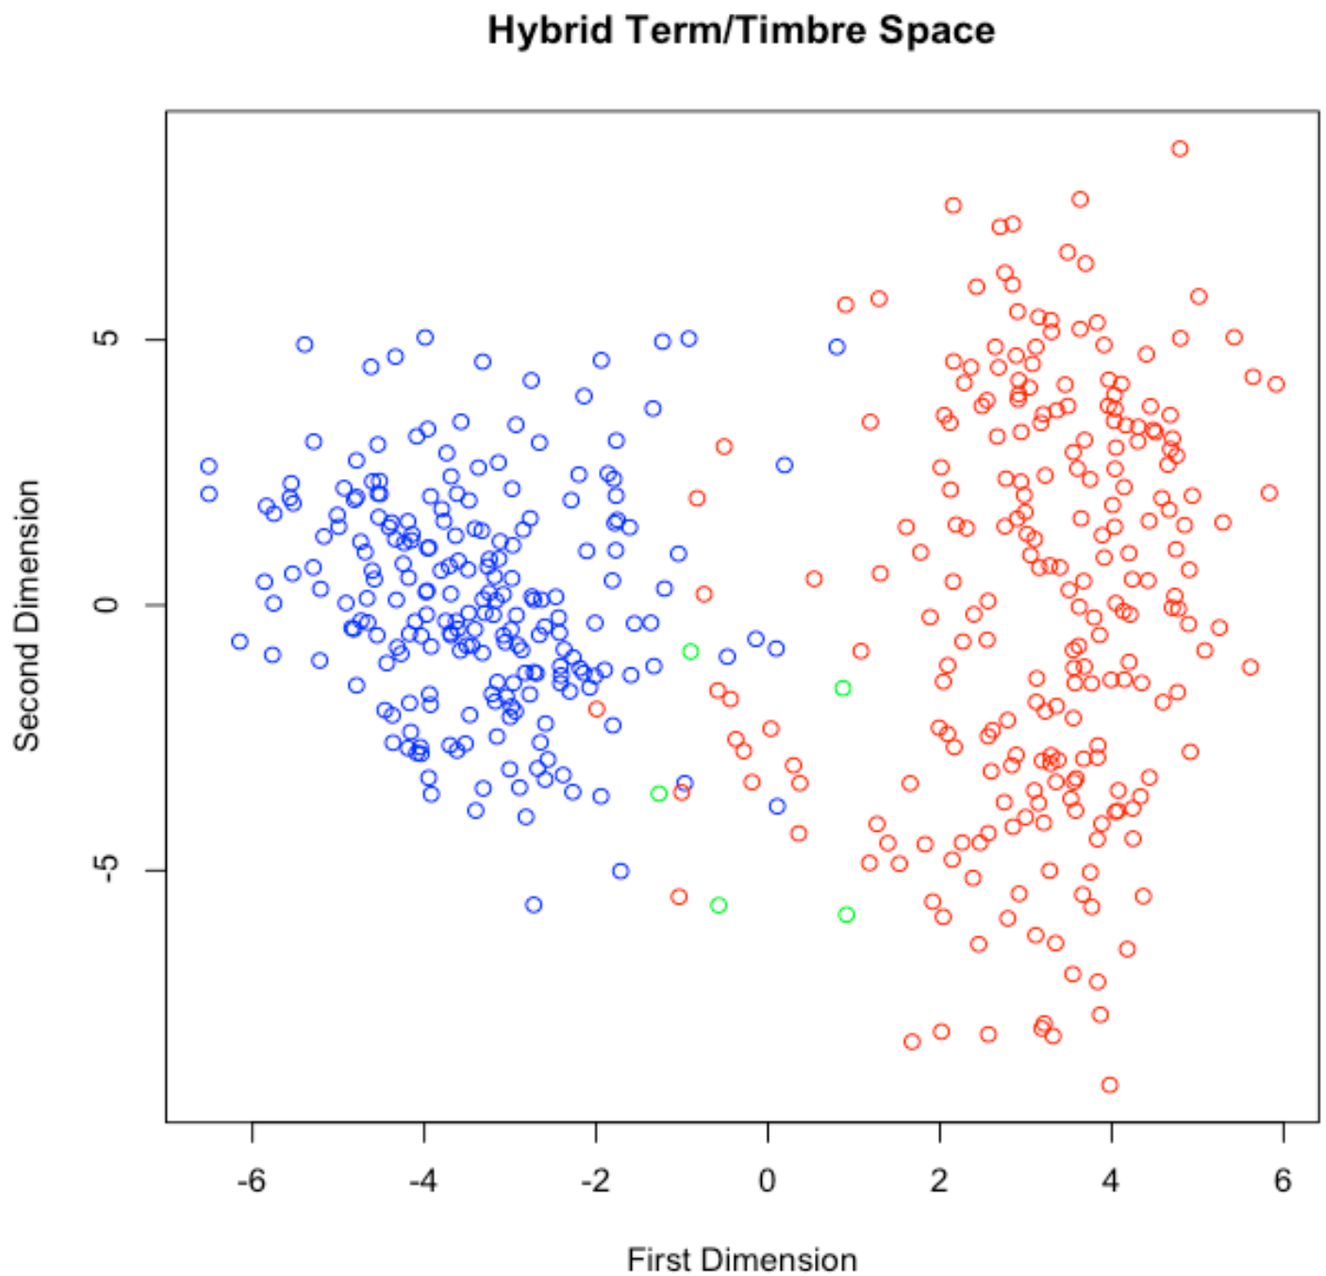

ZMDS Visualization

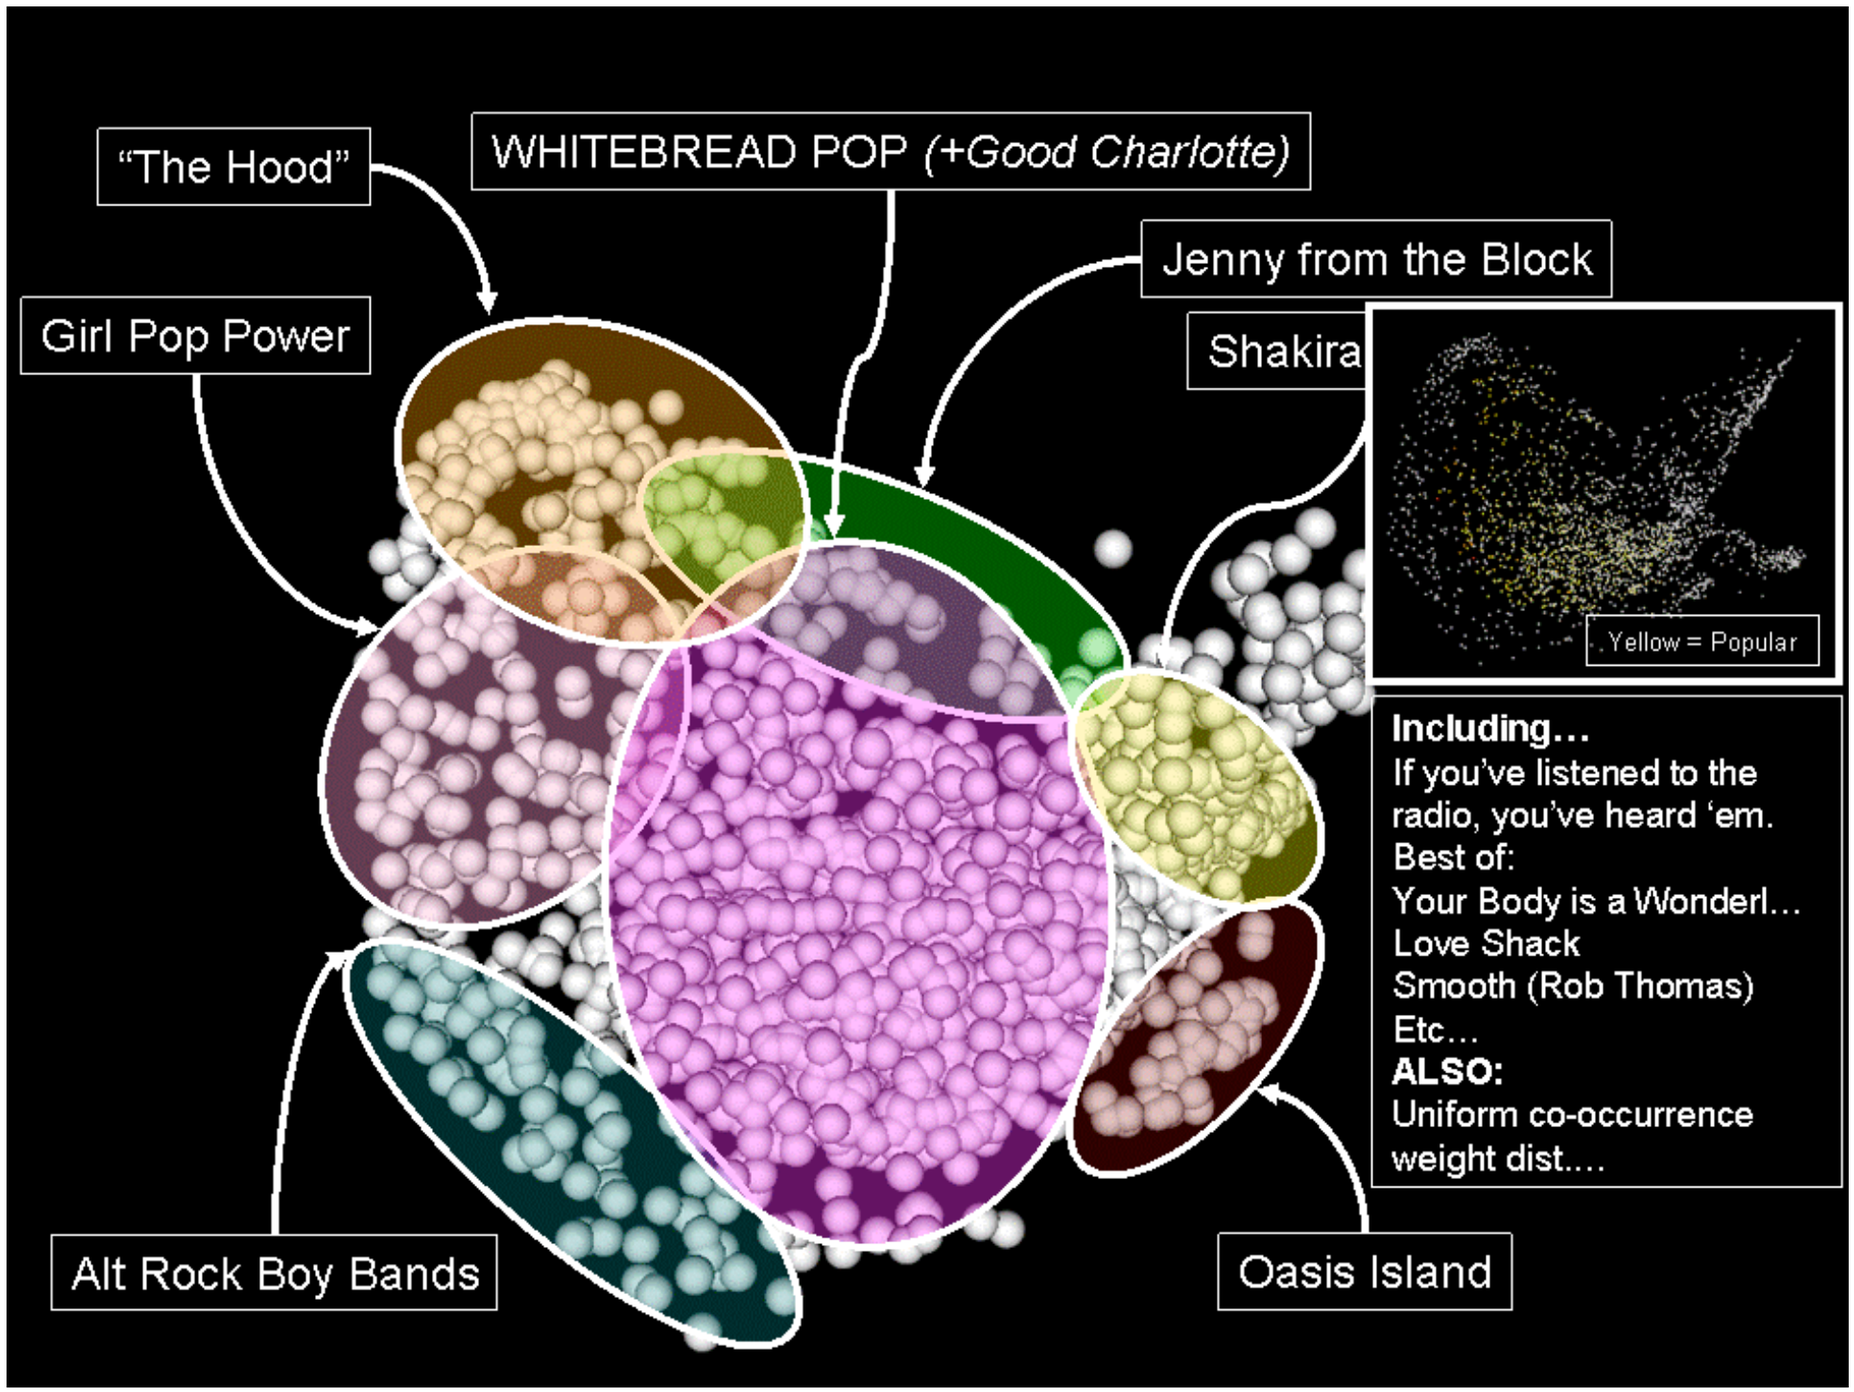

Justin Donaldson

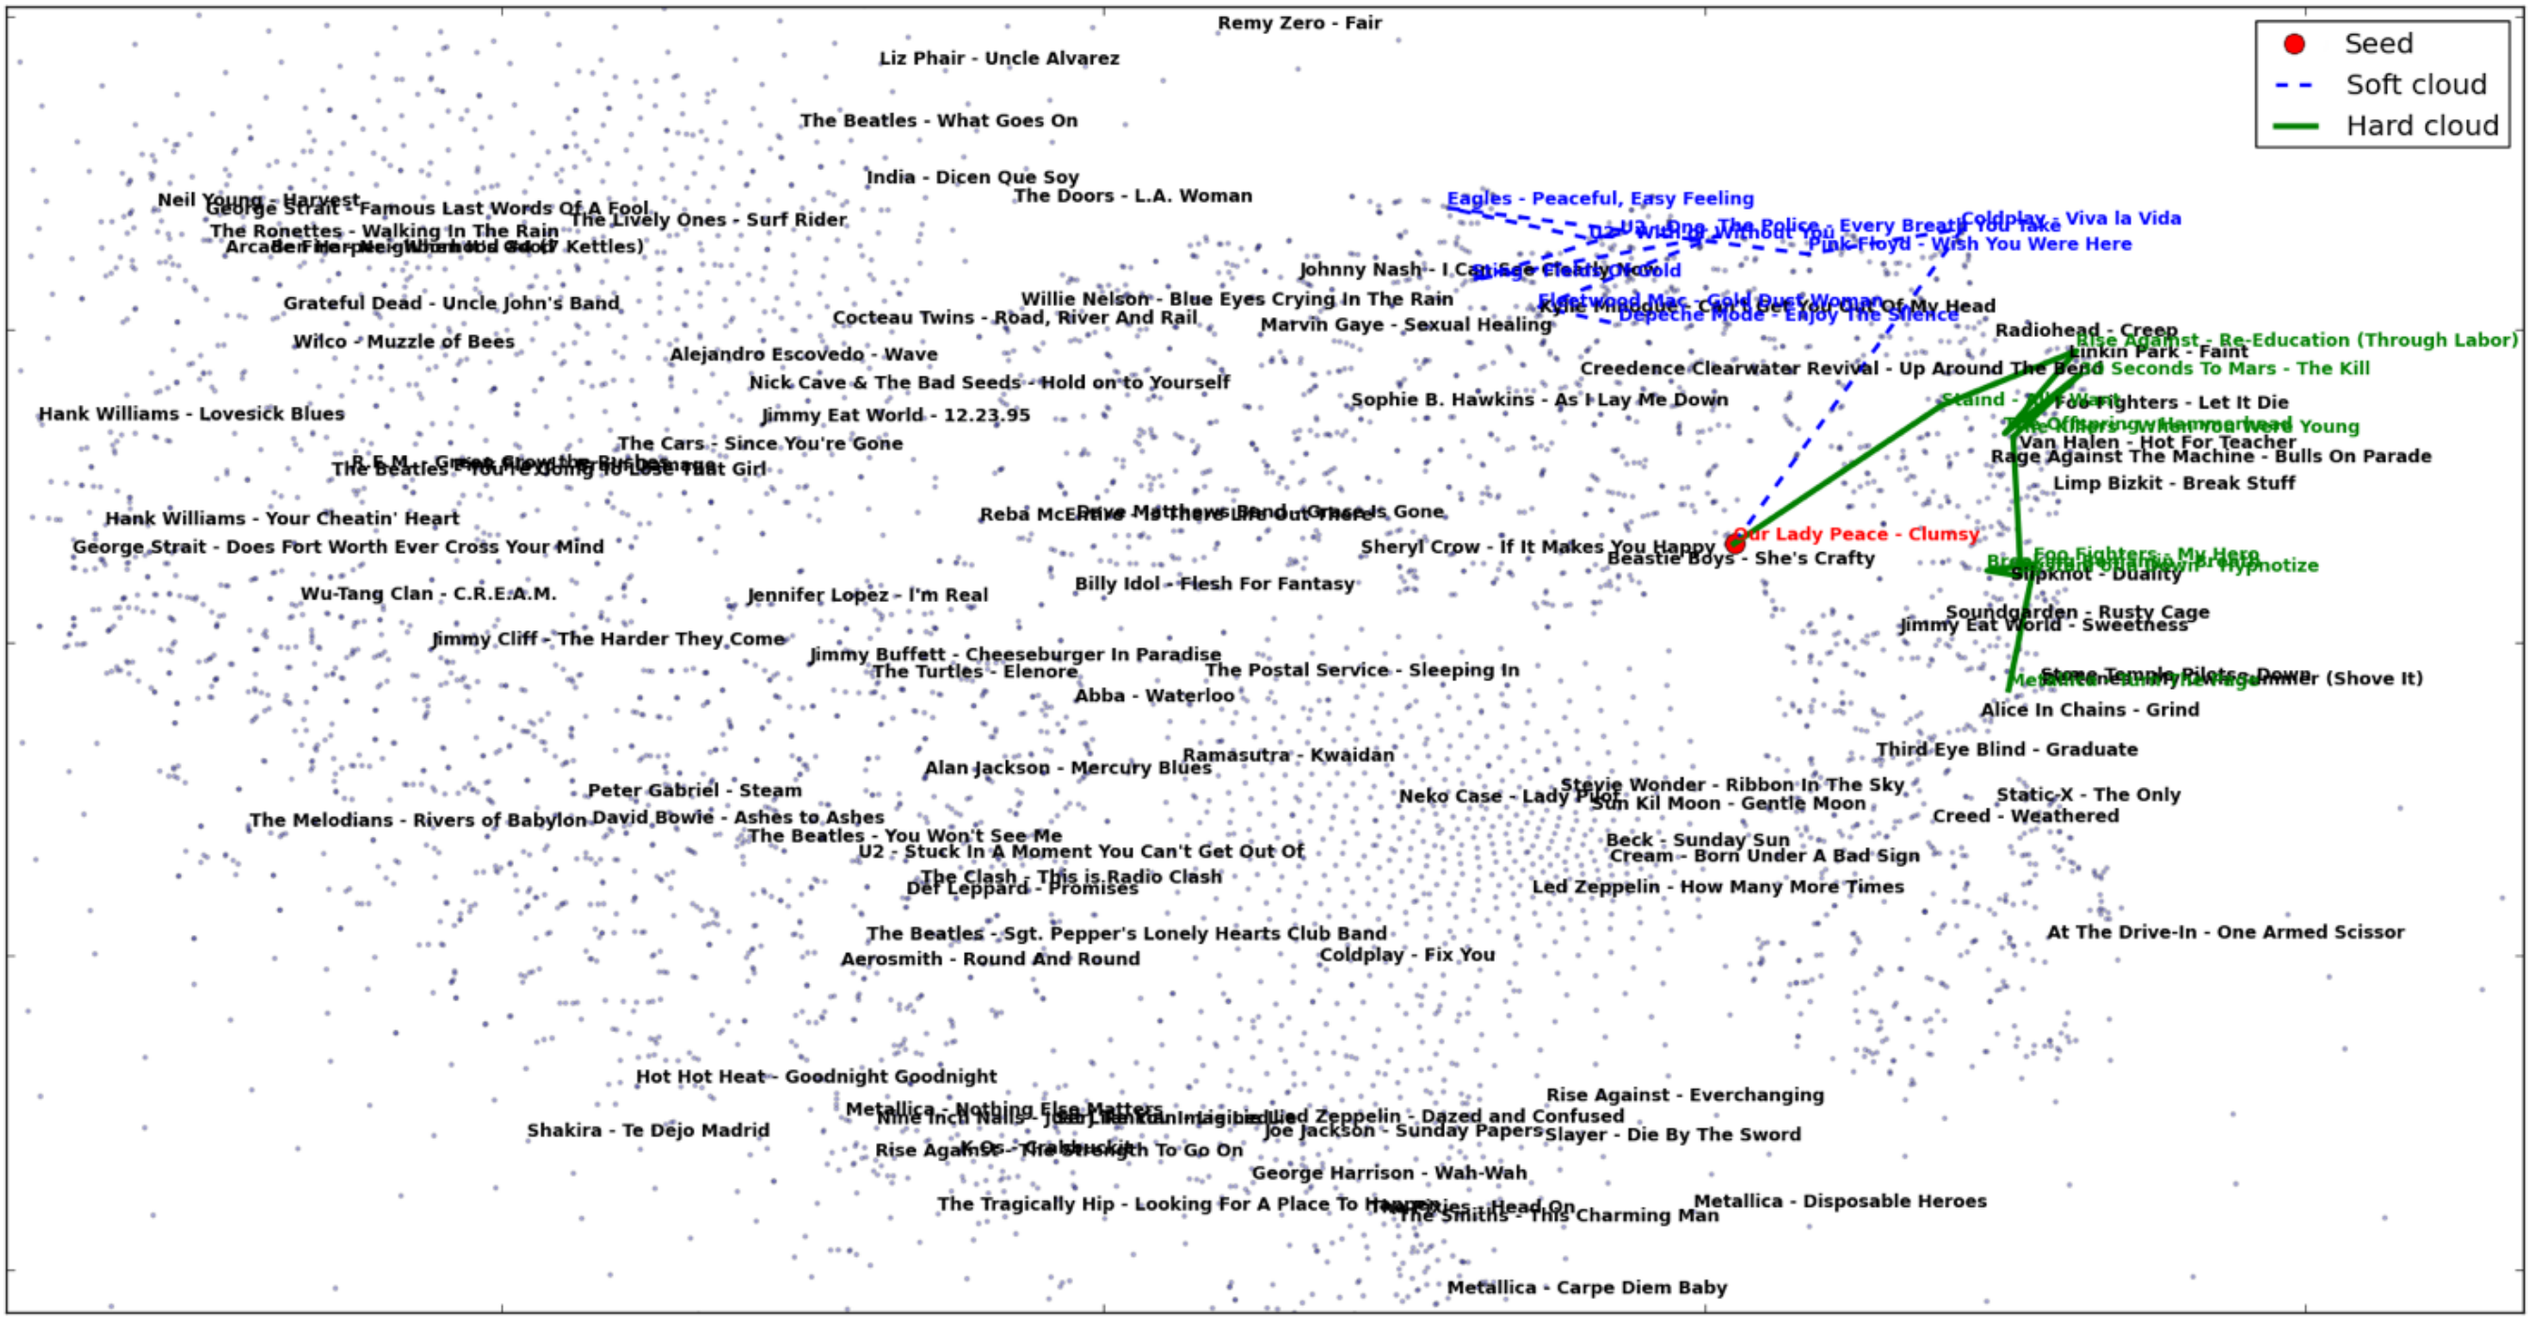

Visualization Name

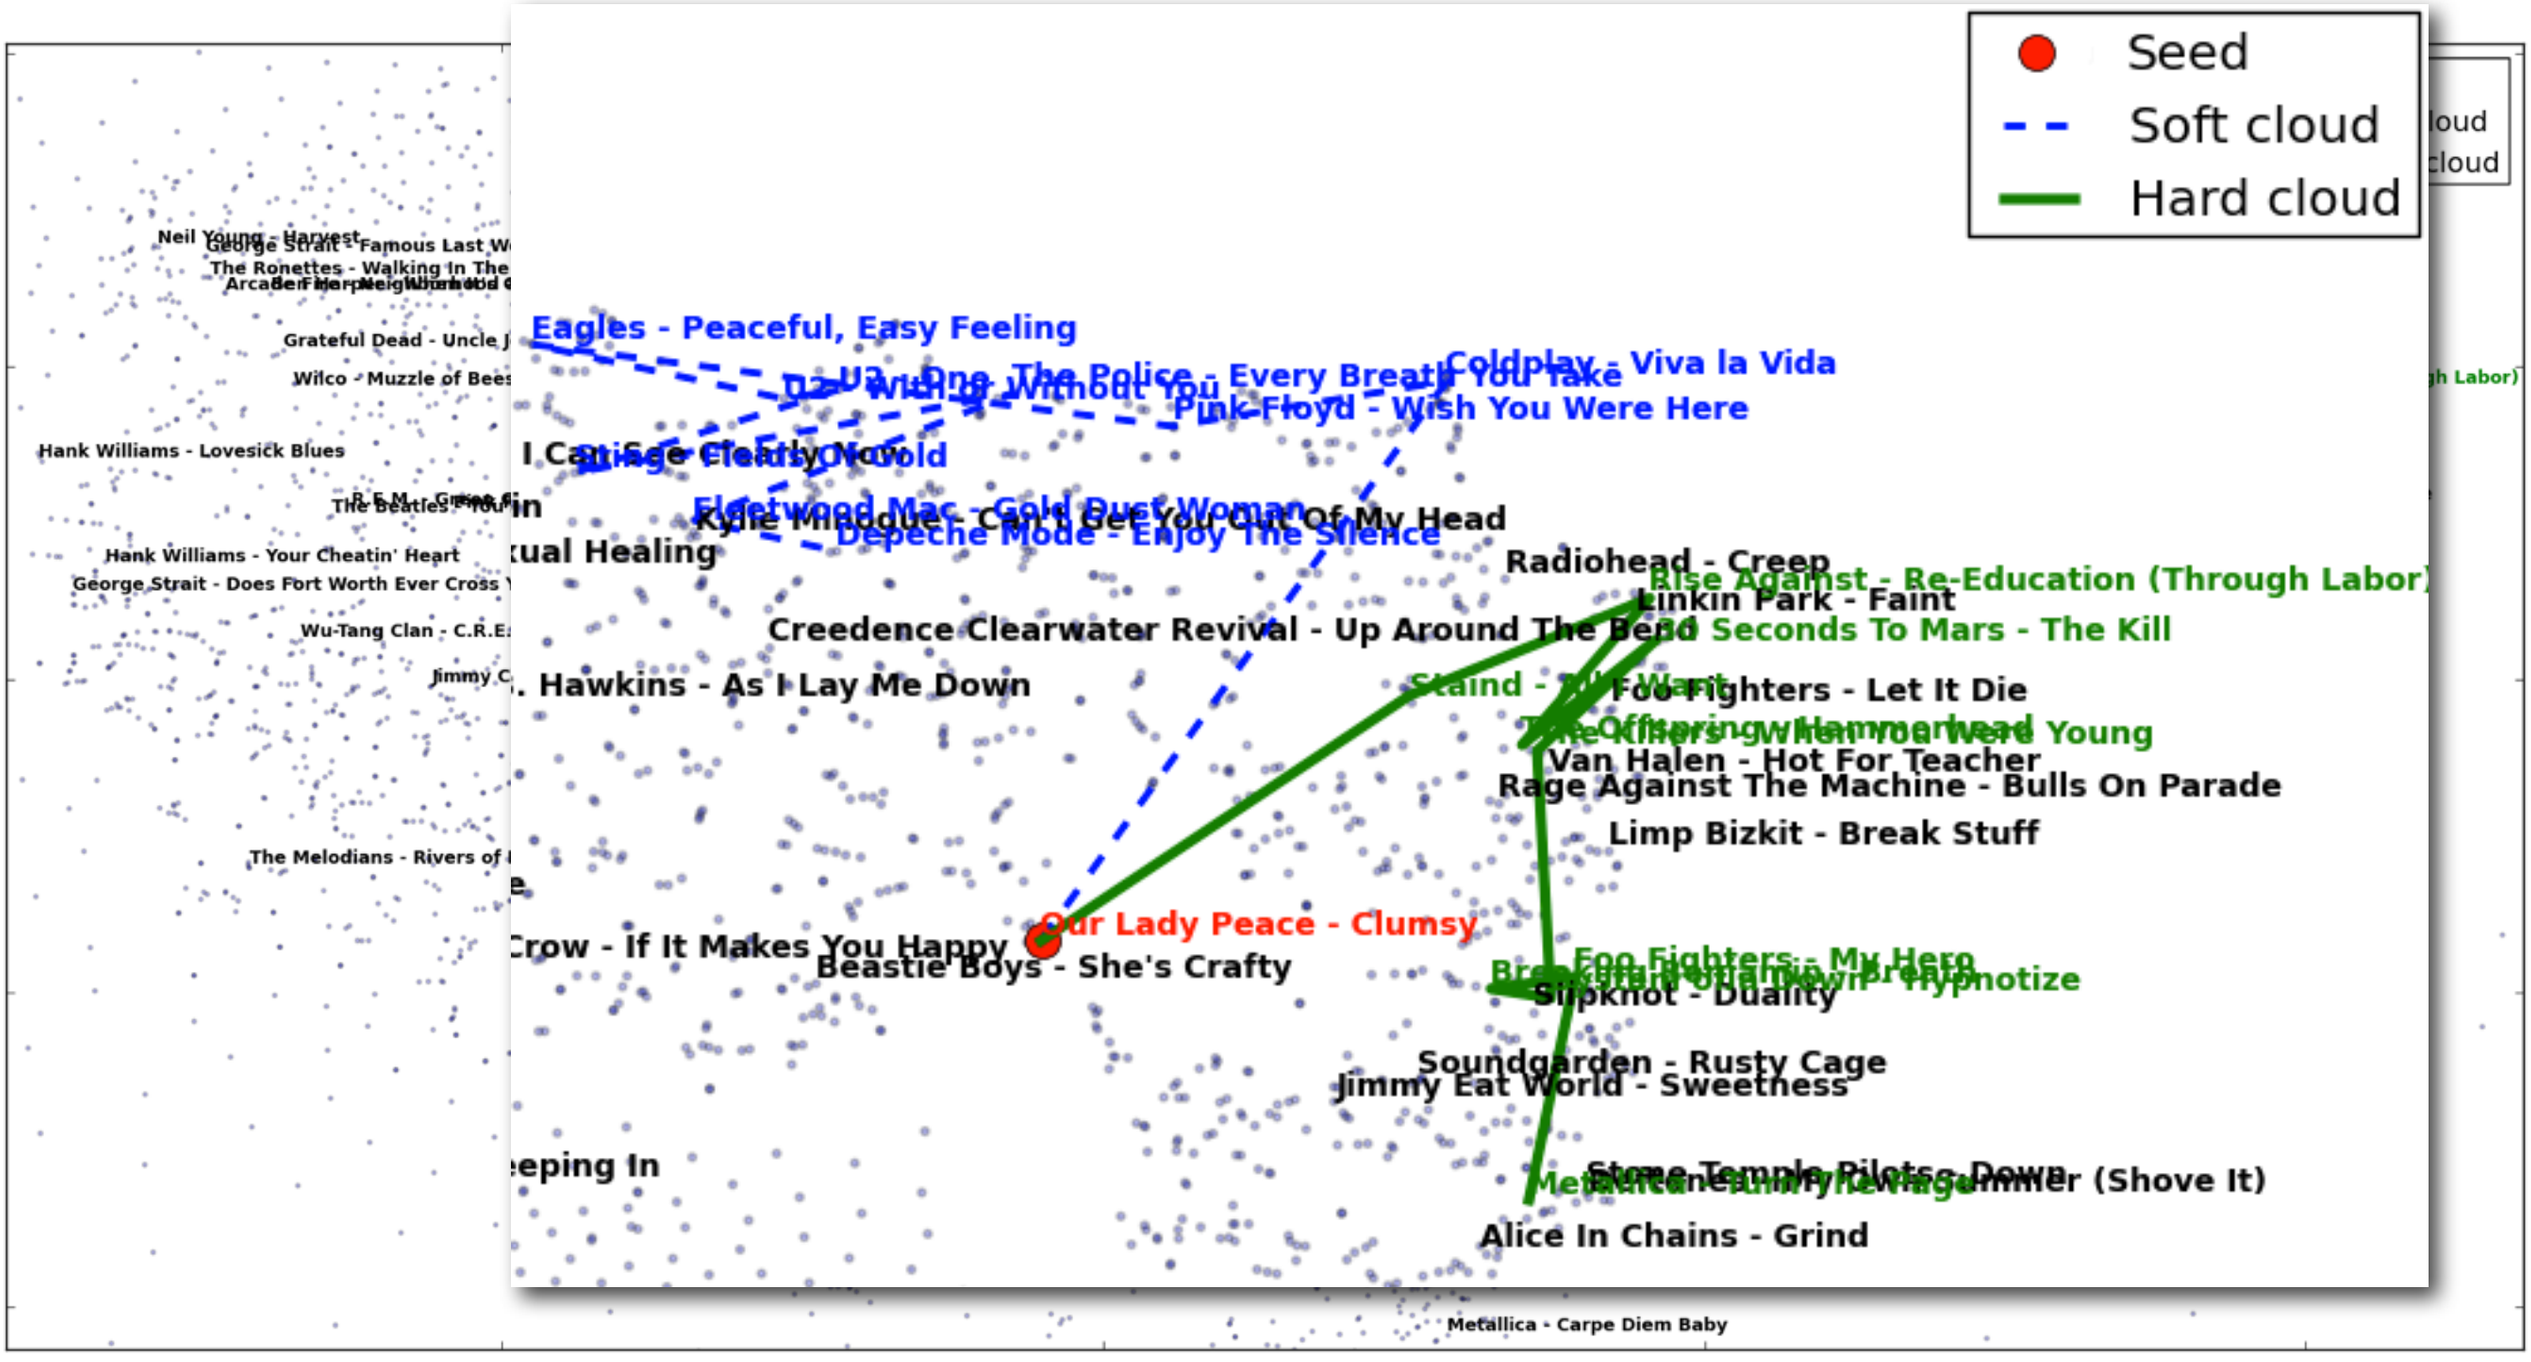

Visualization Name

Using map for playlist navigation

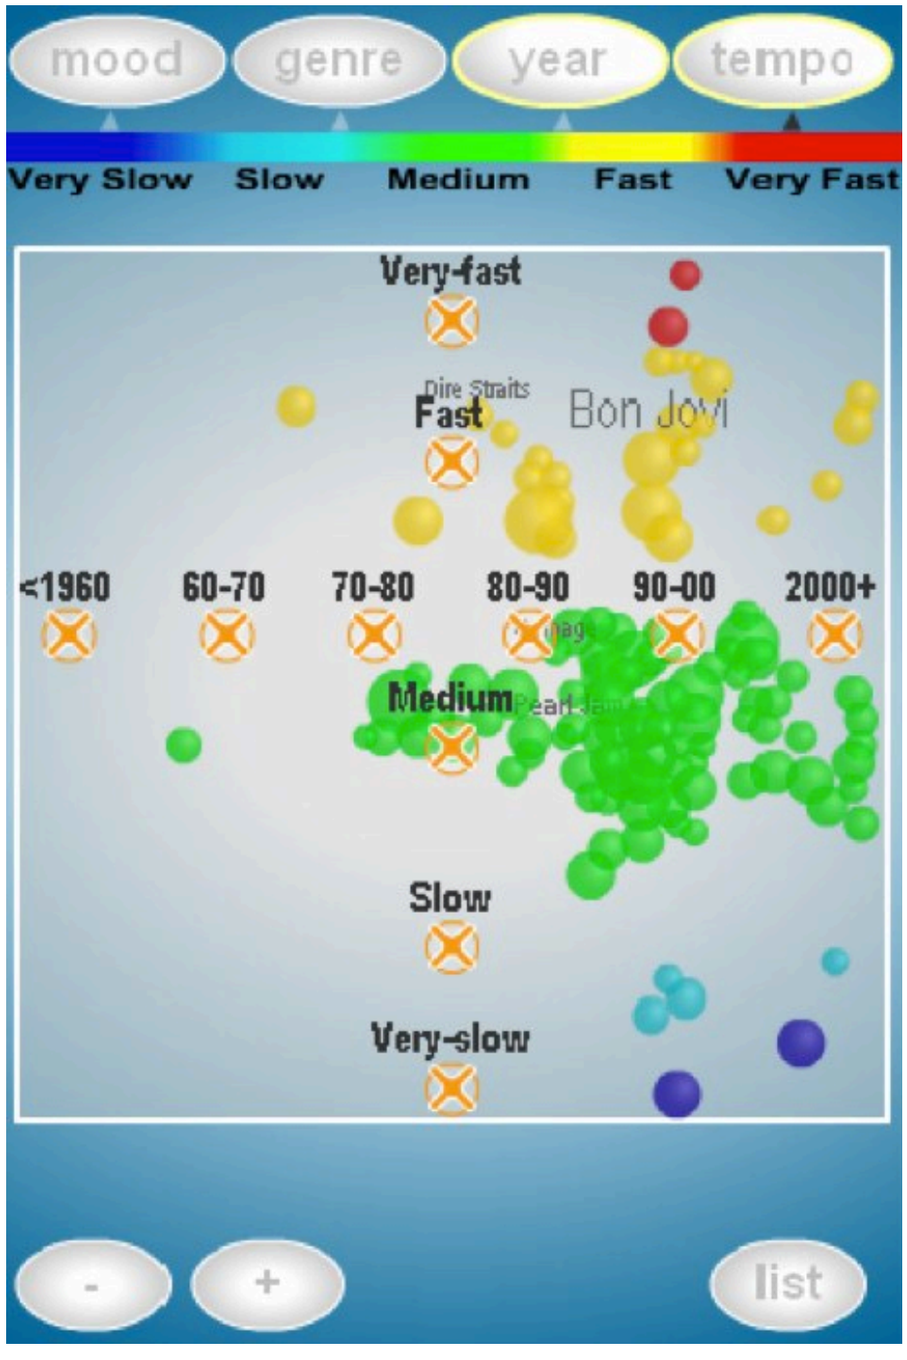

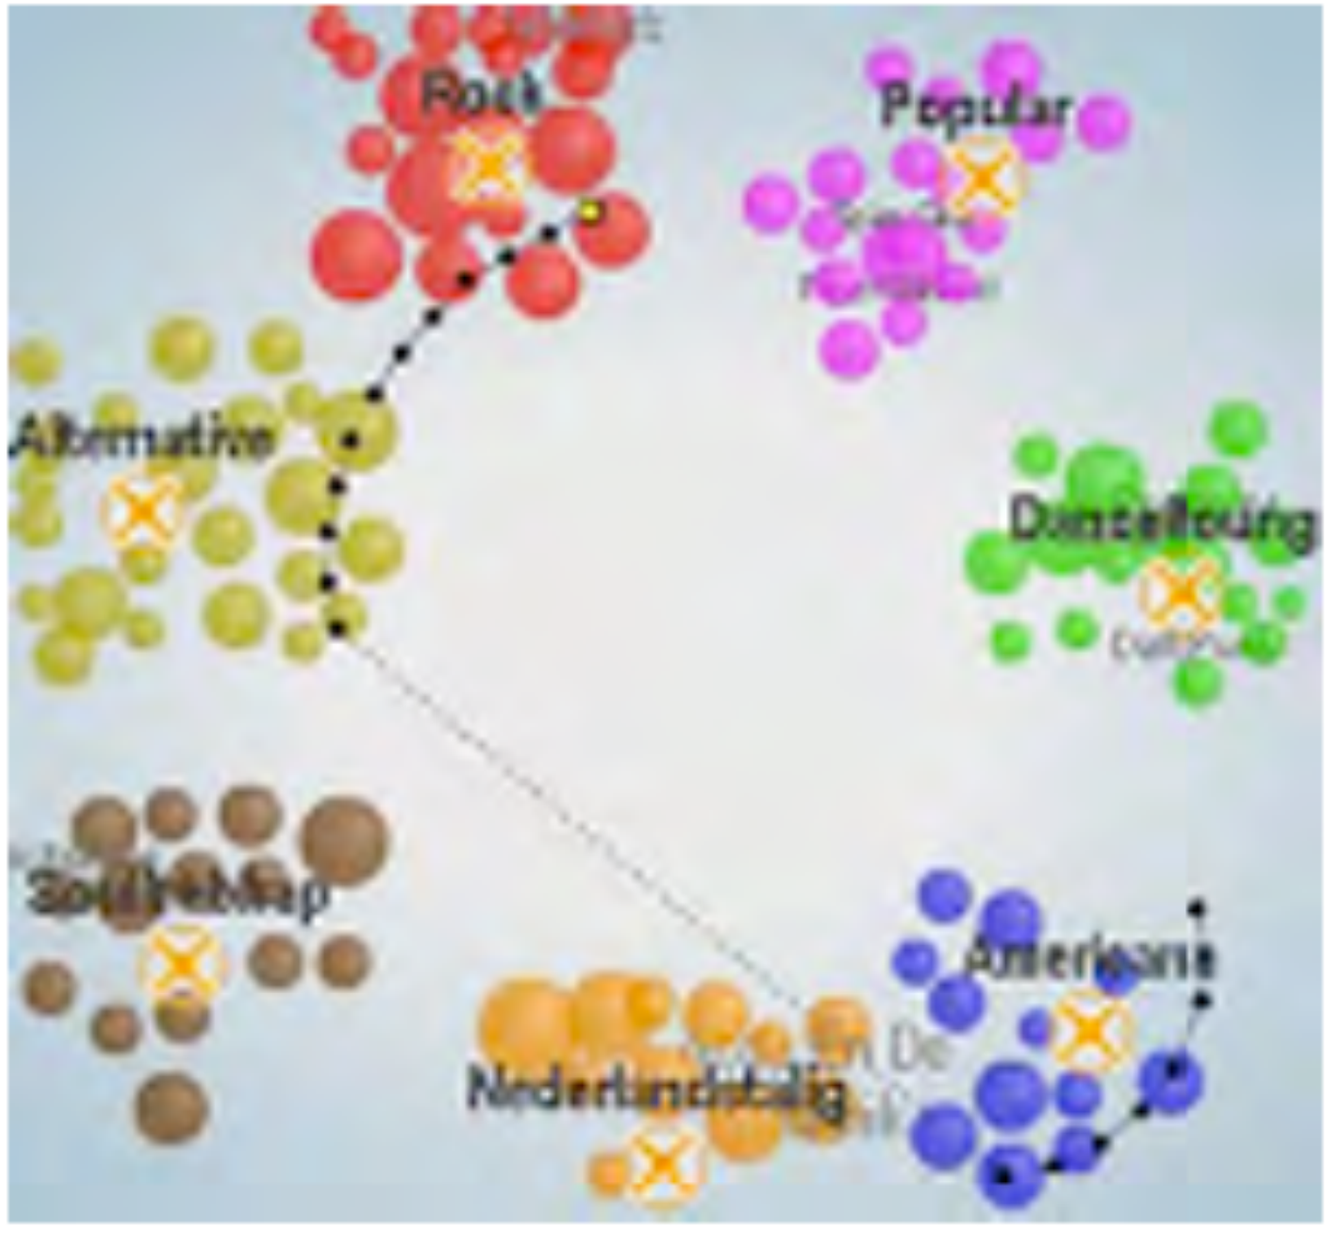

Visual Playlist Generation on the Artist Map - *Van Gulick, Vignoli*

Mr. Emo: Music Retrieval in the Emotional Plane

http://mpac.ee.ntu.edu.tw/~yihsuan/pub/MM08_MrEmo.pdf

Mr. Emo: Music Retrieval in the Emotional Plane

http://mpac.ee.ntu.edu.tw/~yihsuan/pub/MM08_MrEmo.pdf

Burst Labs

http://www.burstlabs.com/

Search Inside the Music

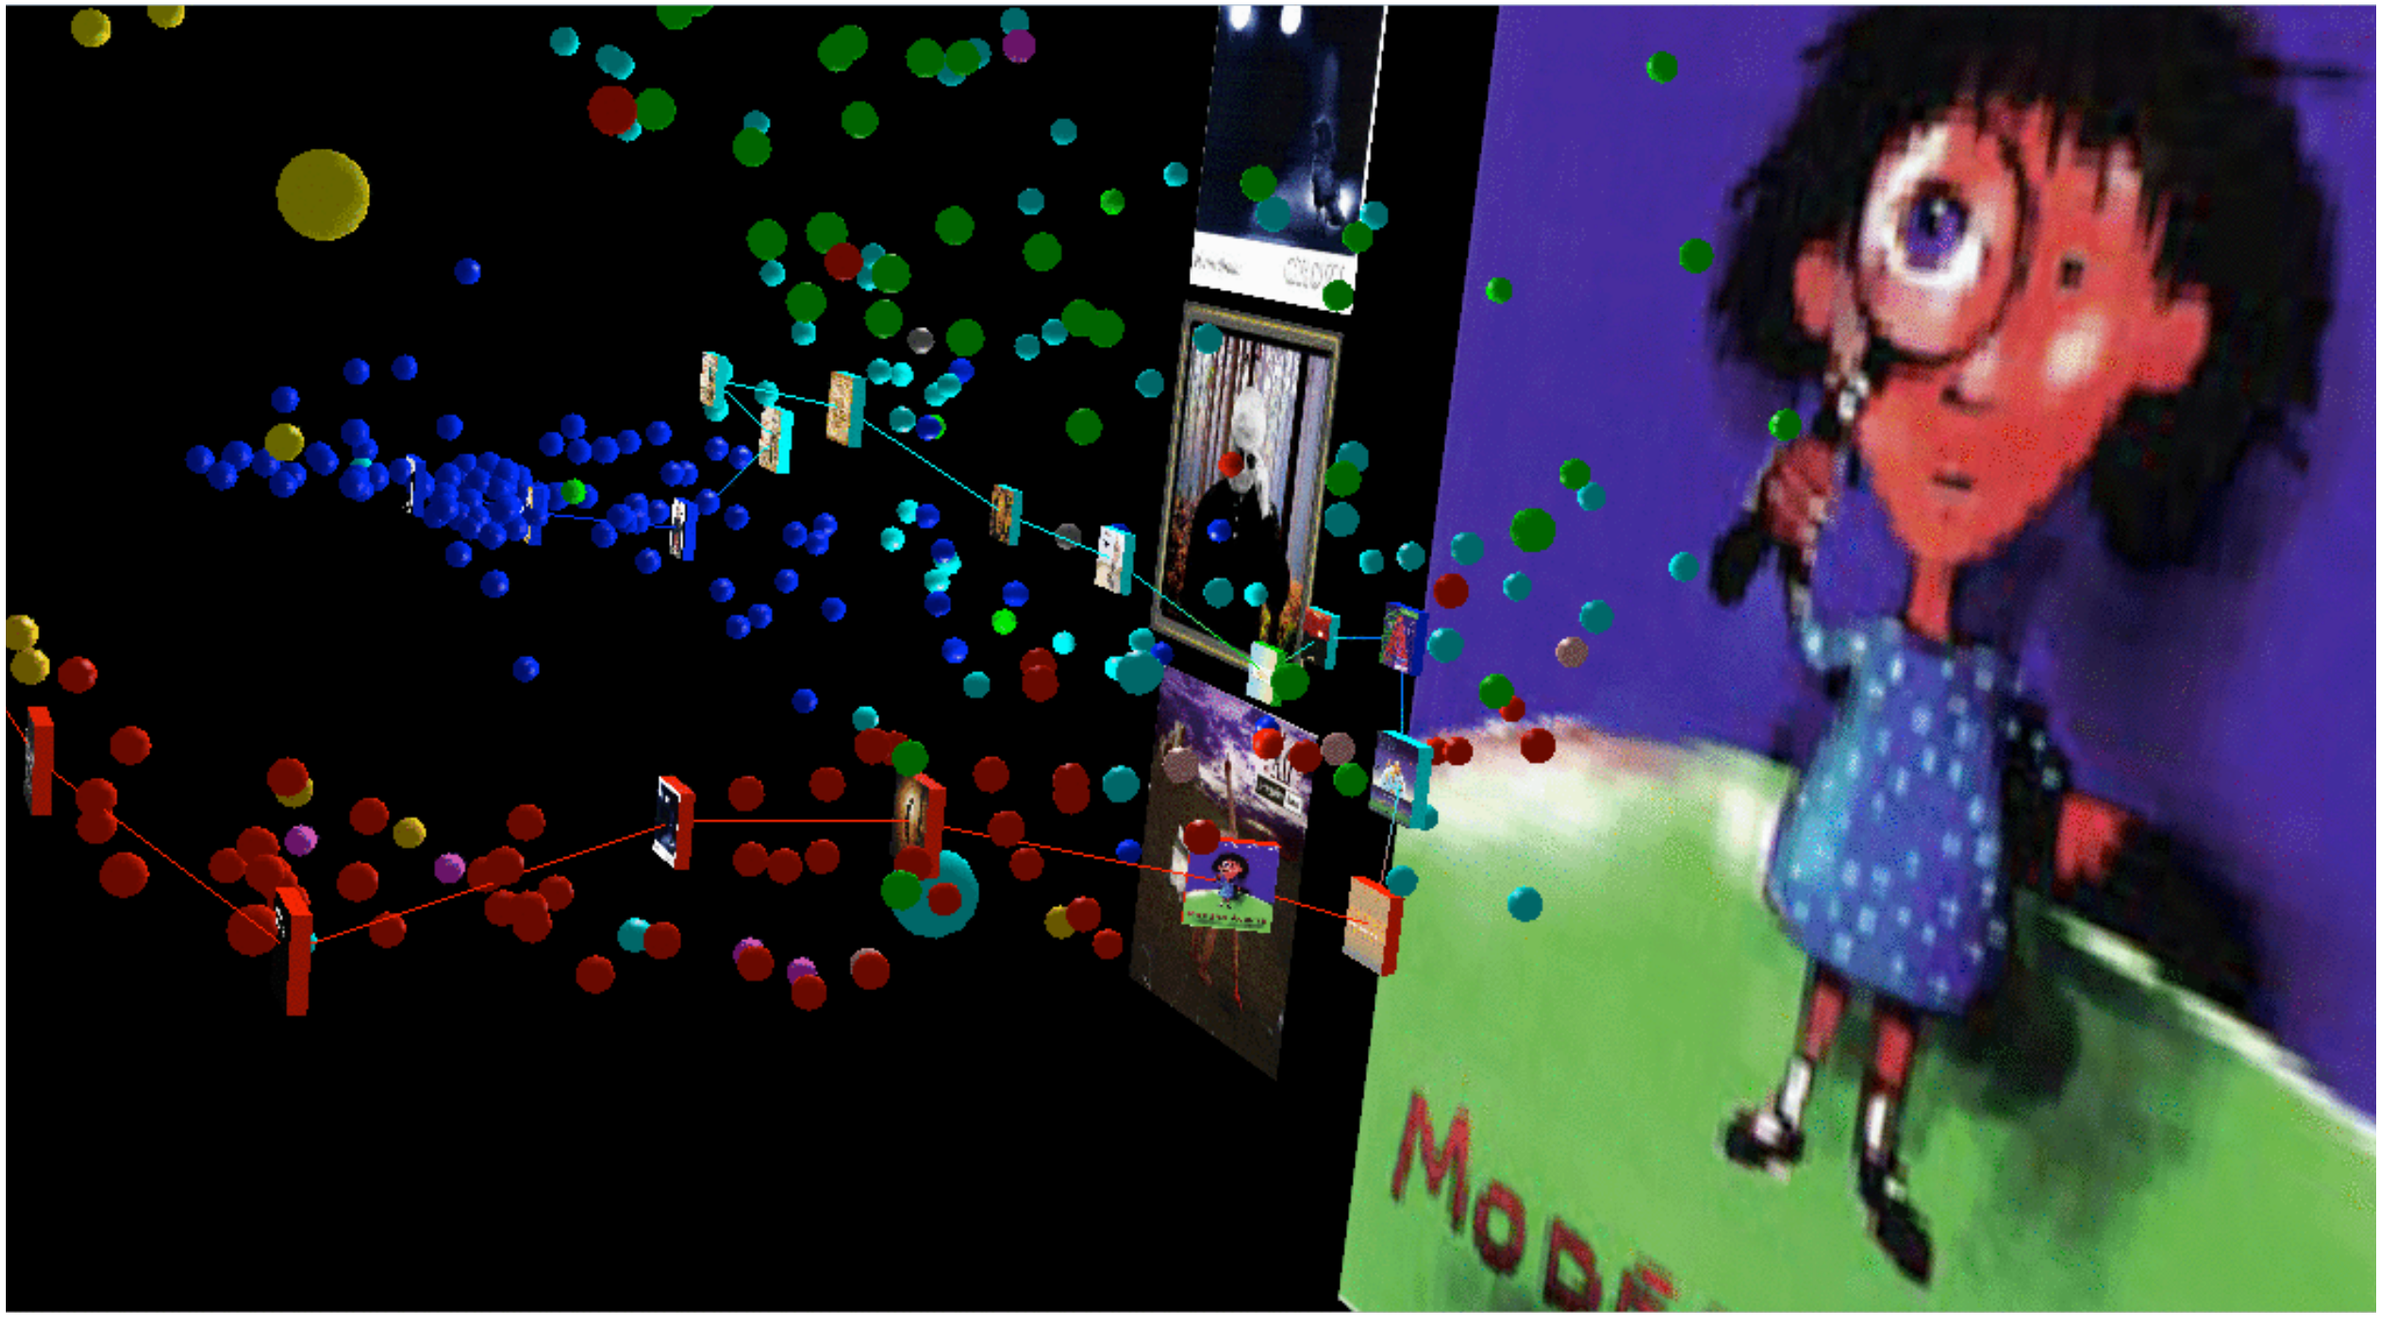

http://mediacast.sun.com/share/plamere/sitm_final.pdf - Lamere, Eck

3D Track similarity

The Mufin Player

3D Track similarity

The Mufin Player

Music Box

http://thesis.flyingpudding.com/ - Anita Lillie

Music Box

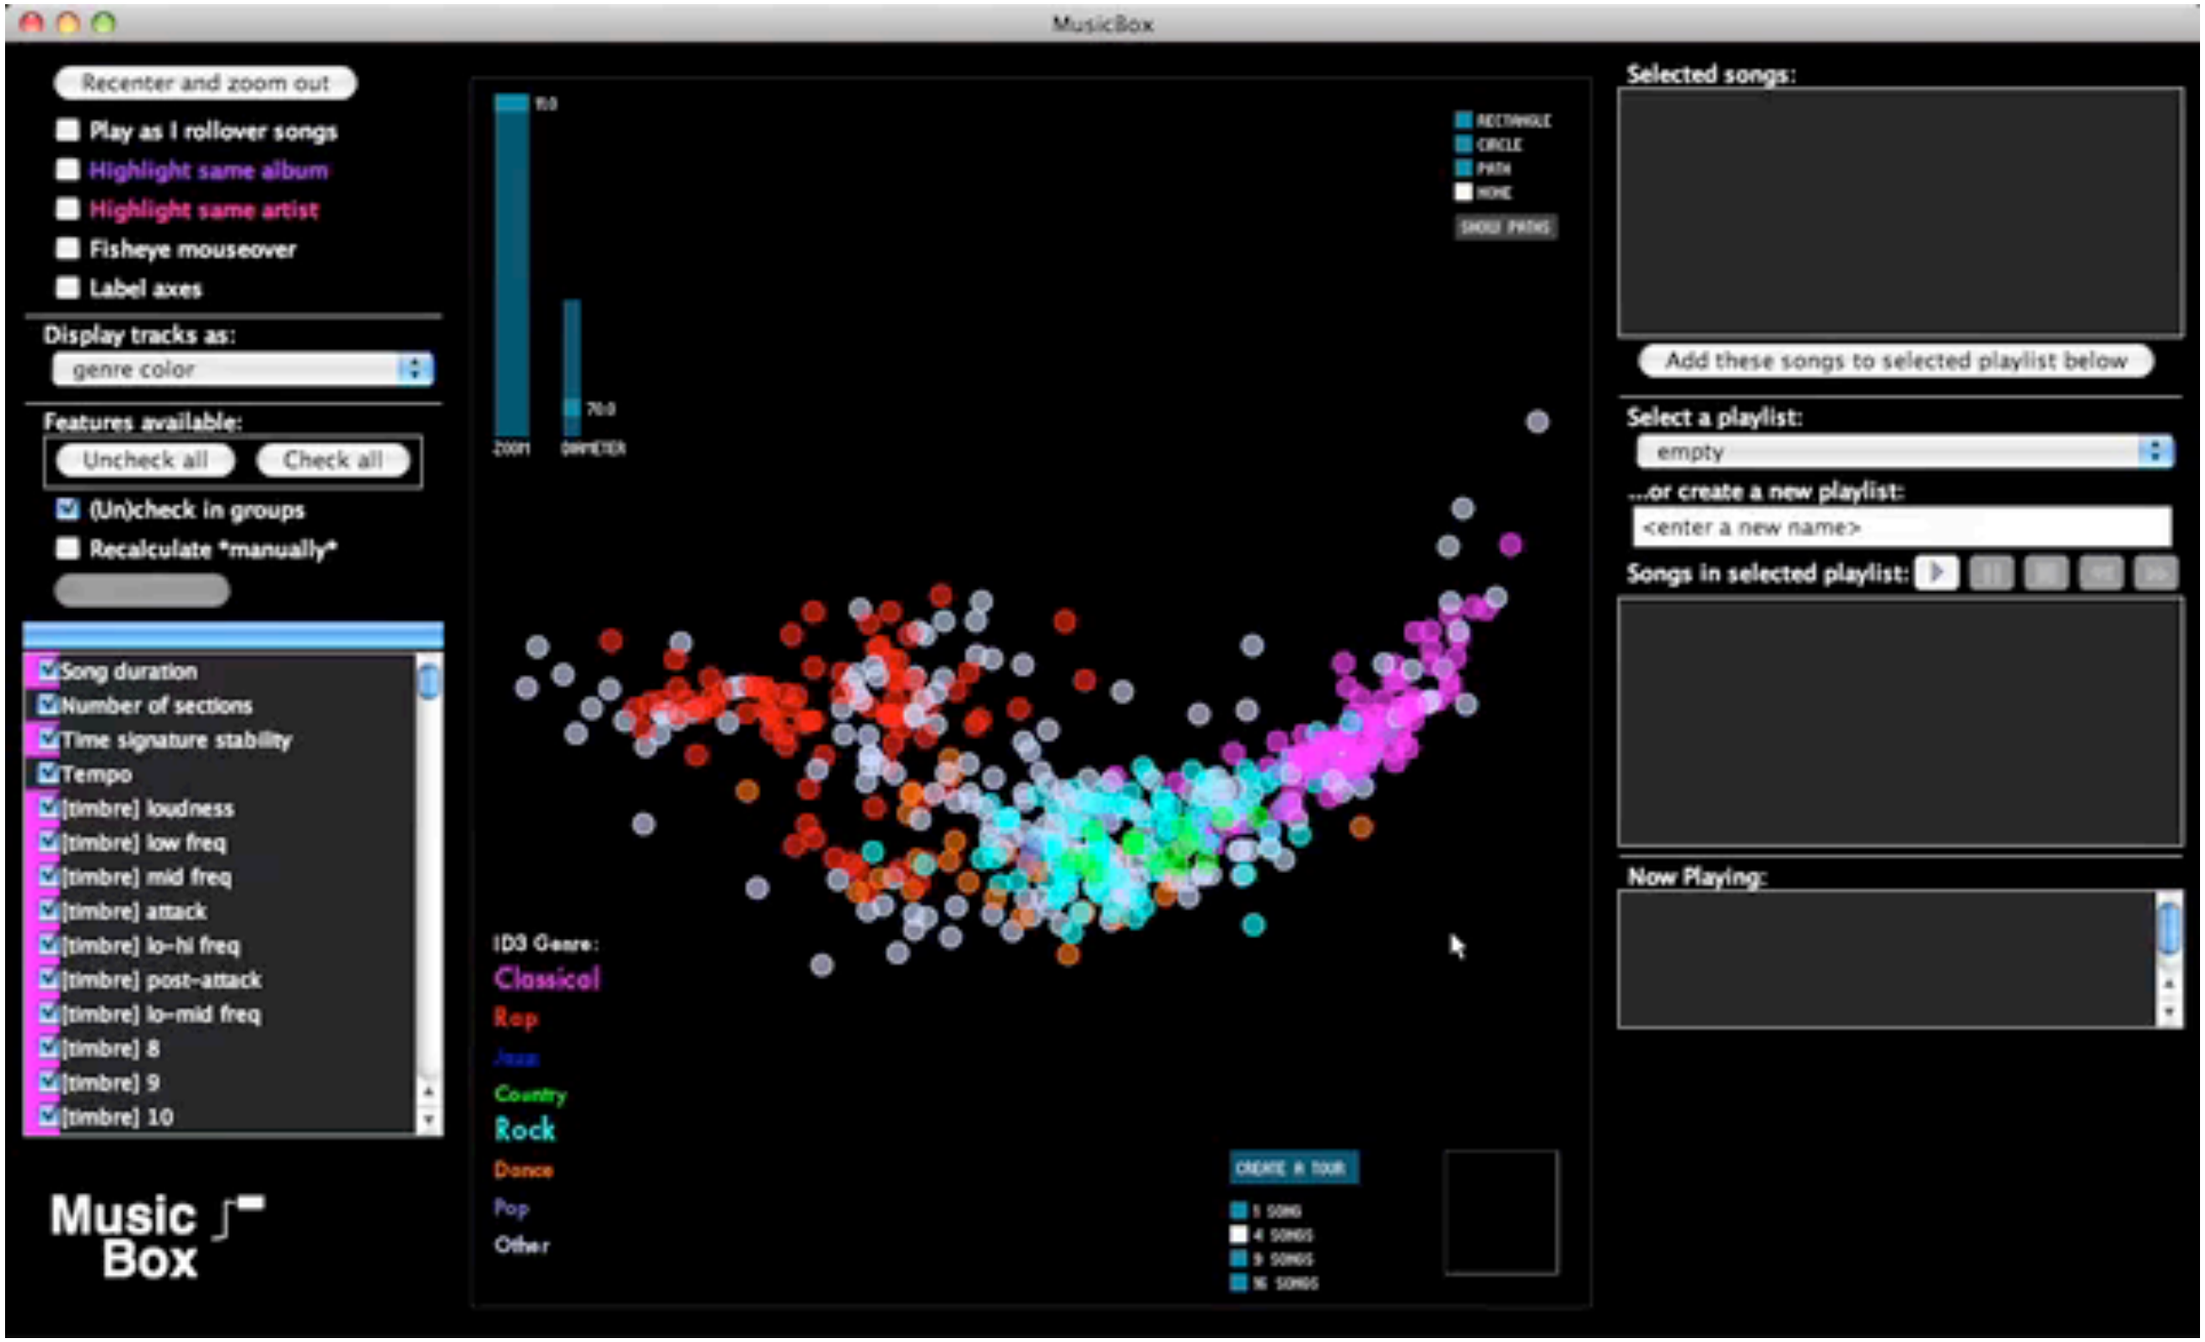

http://thesis.flyingpudding.com/ - Anita Lillie

Pictures of Songs

Beethoven, Fur Elise:

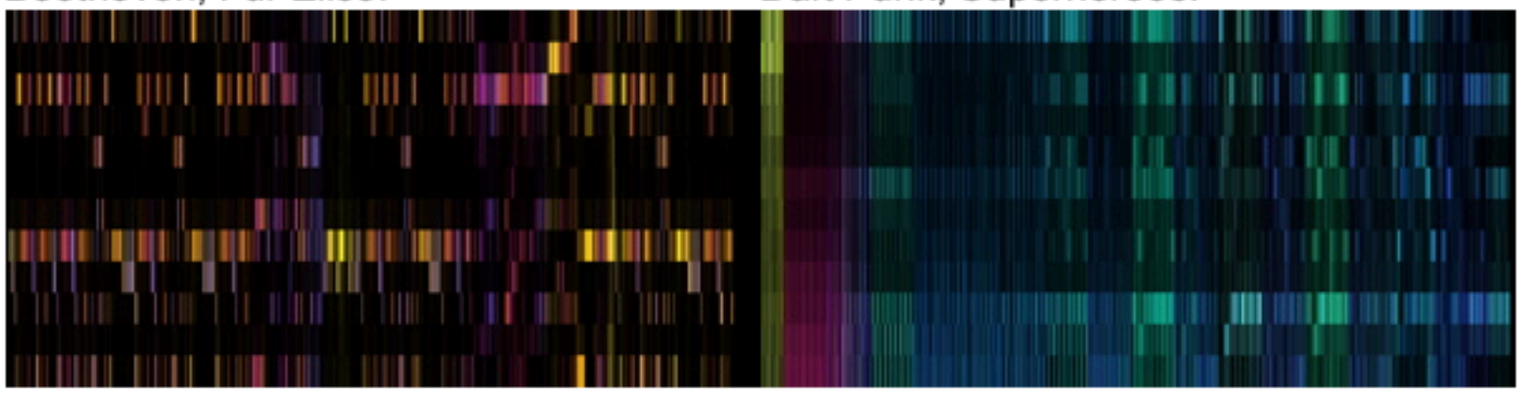

Daft Punk, Superheroes:

Beethoven, "Moonlight" Sonata:

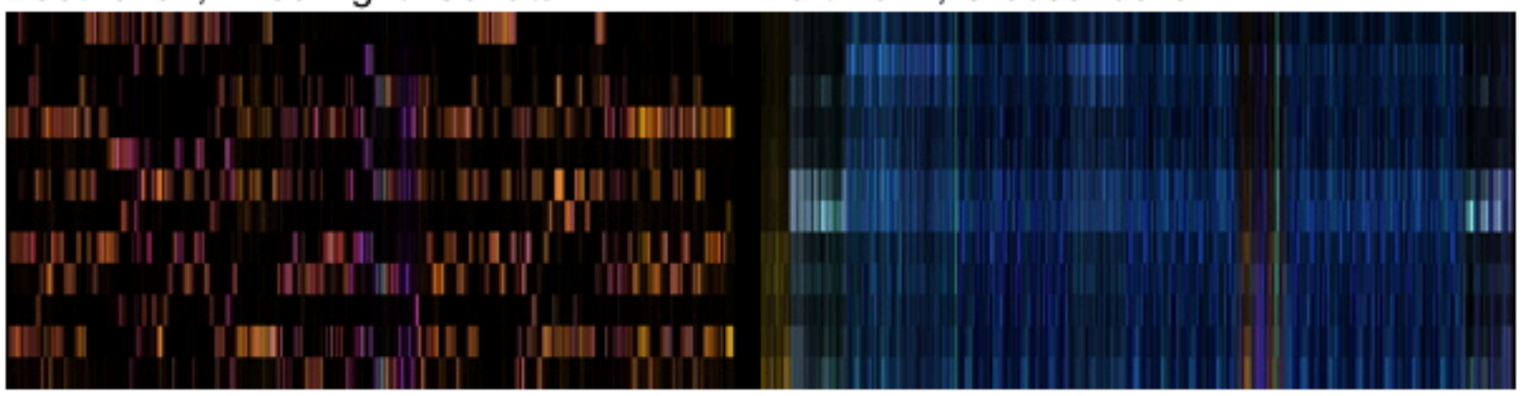

Daft Punk, Crescendolls:

Nirvana, Come As You Are:

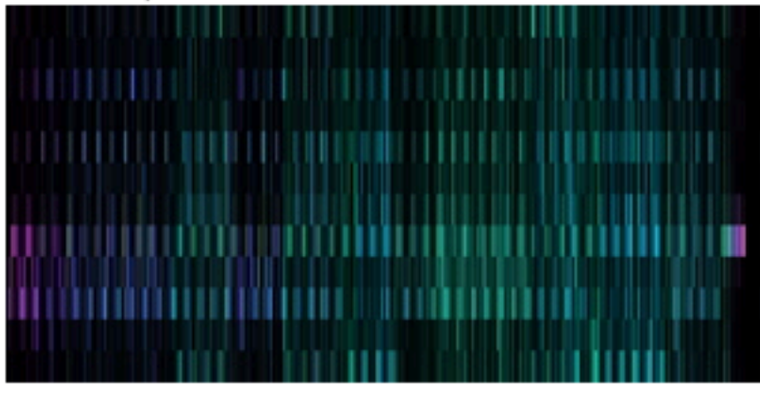

Sonic Youth, 'Cross the Breeze:

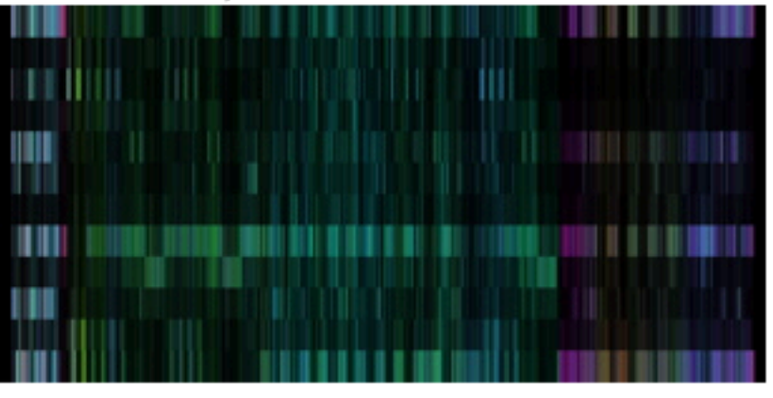

http://www.flyingpudding.com/projects/viz_music/

Visualizing Emotion in Lyrics

http://www.cs.kuleuven.be/~jorisk/thrillerMovie.mov

Sampling Connections

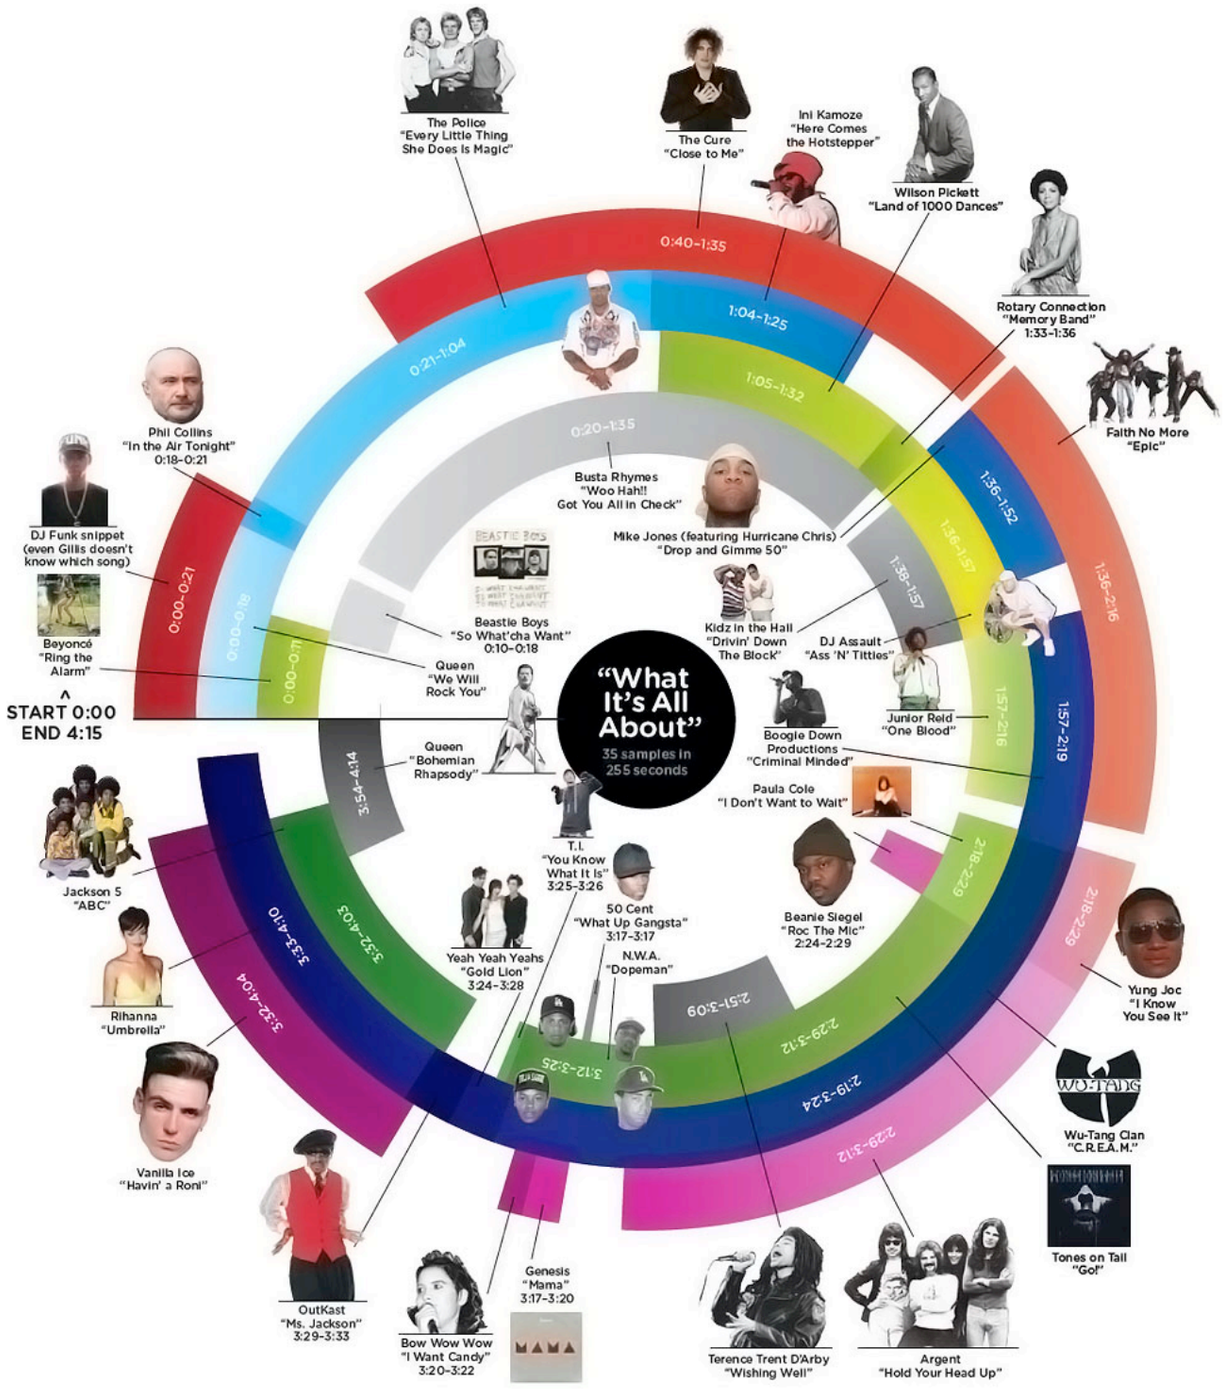

The History of Sampling

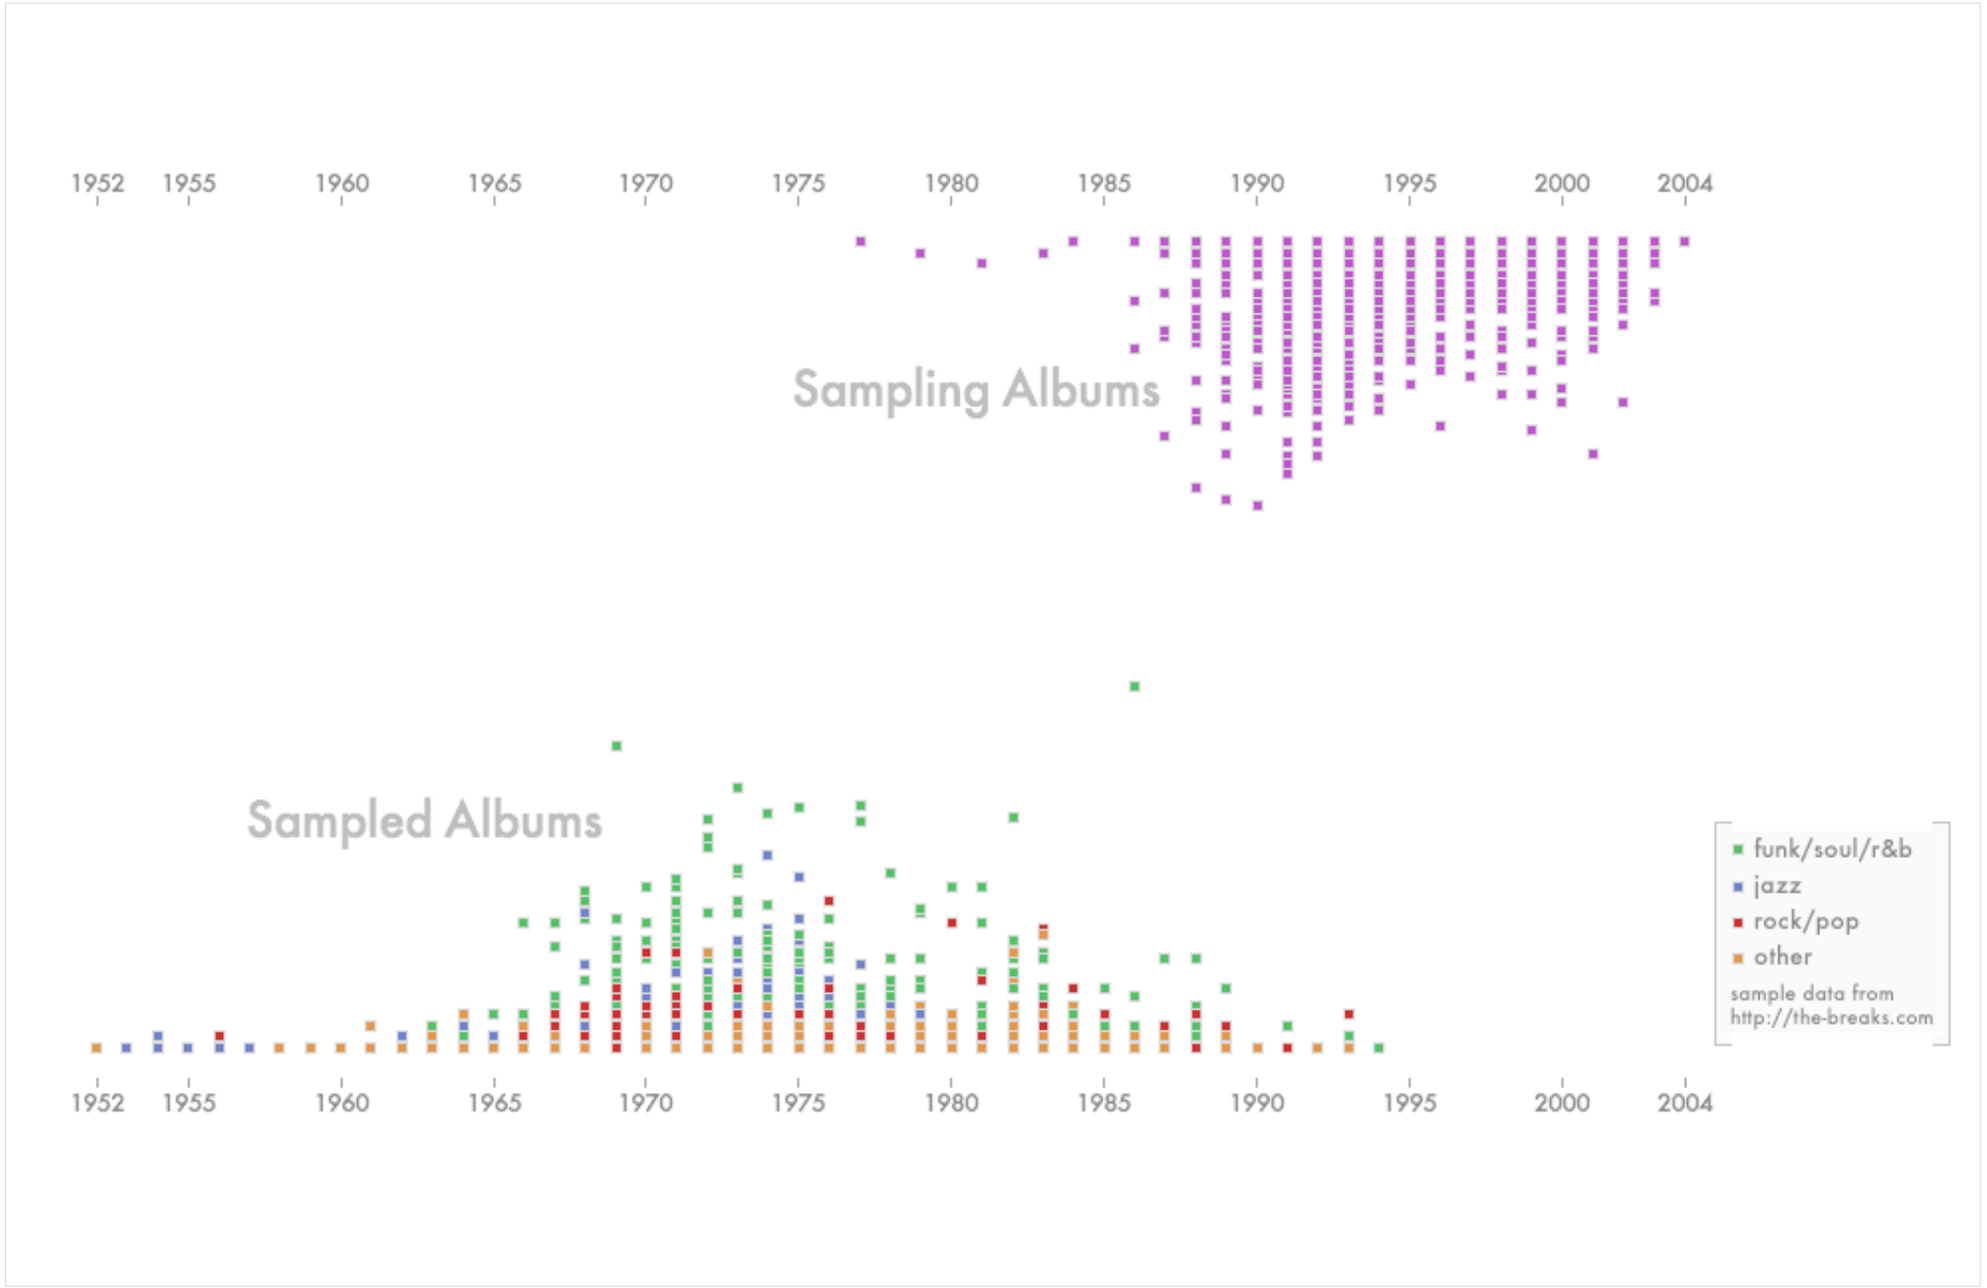

http://jessekriss.com/projects/samplinghistory/

The History of Sampling

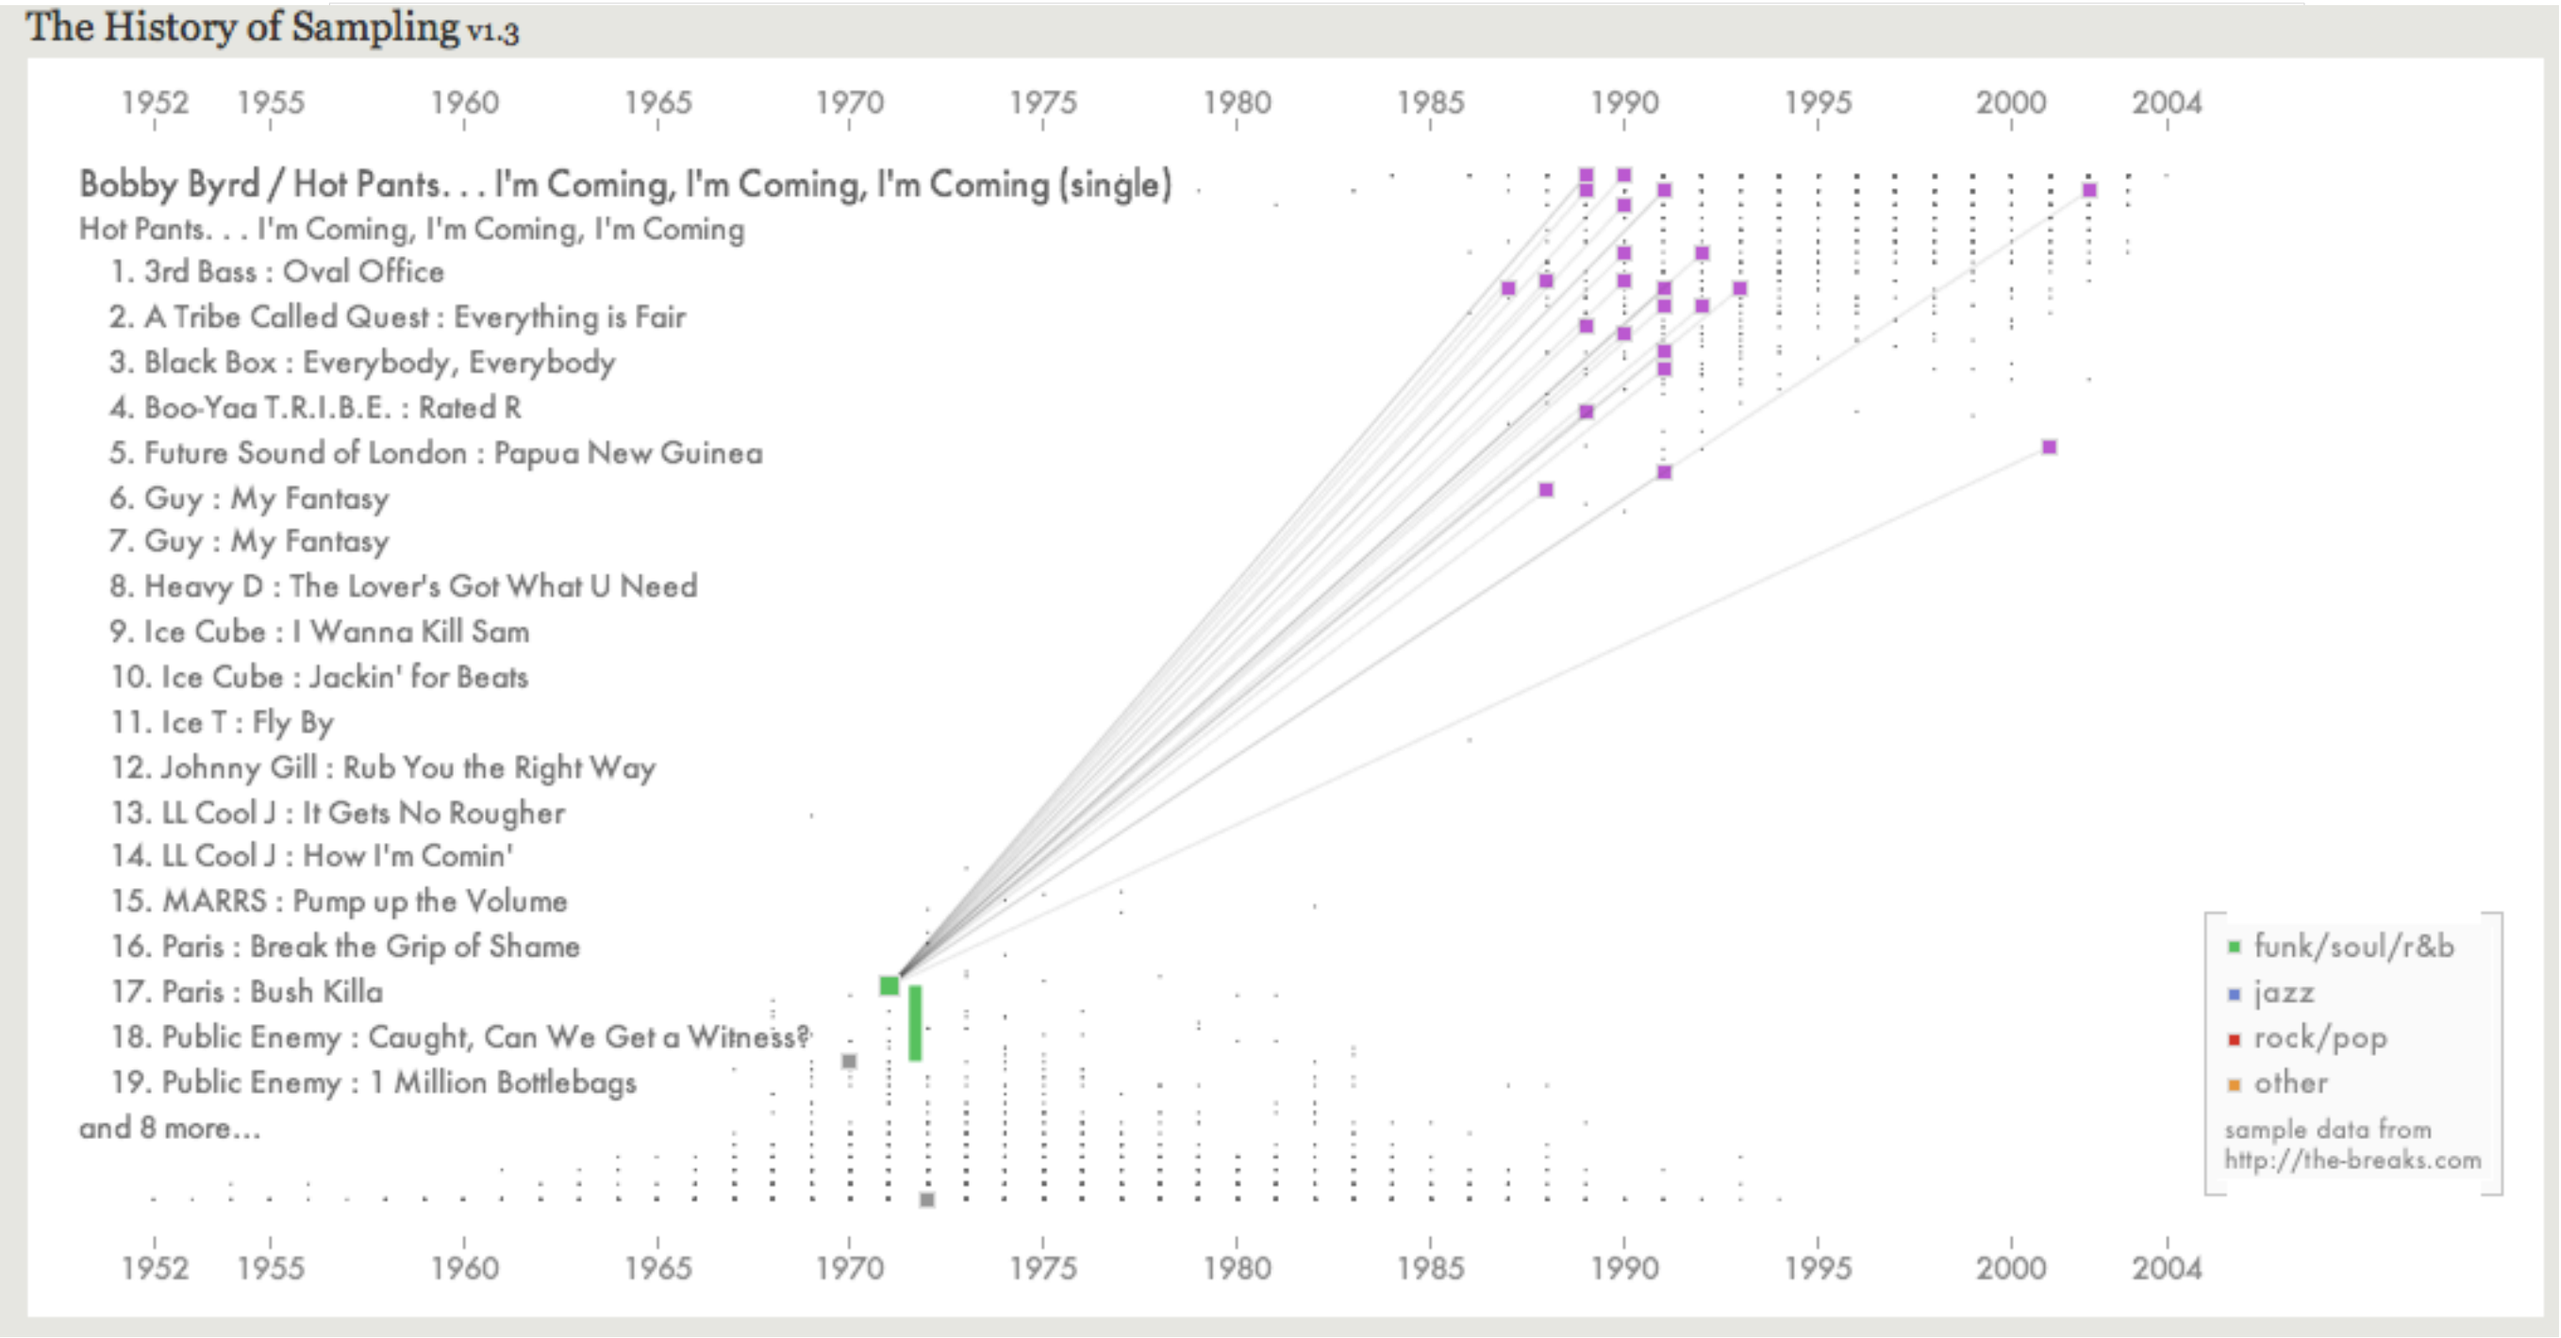

http://jessekriss.com/projects/samplinghistory/

Mused: Navigating the personal sample library

Graham Coleman - http://www.iua.upf.edu/~gcoleman/pubs/mused07.pdf

Mused: Navigating the personal sample library

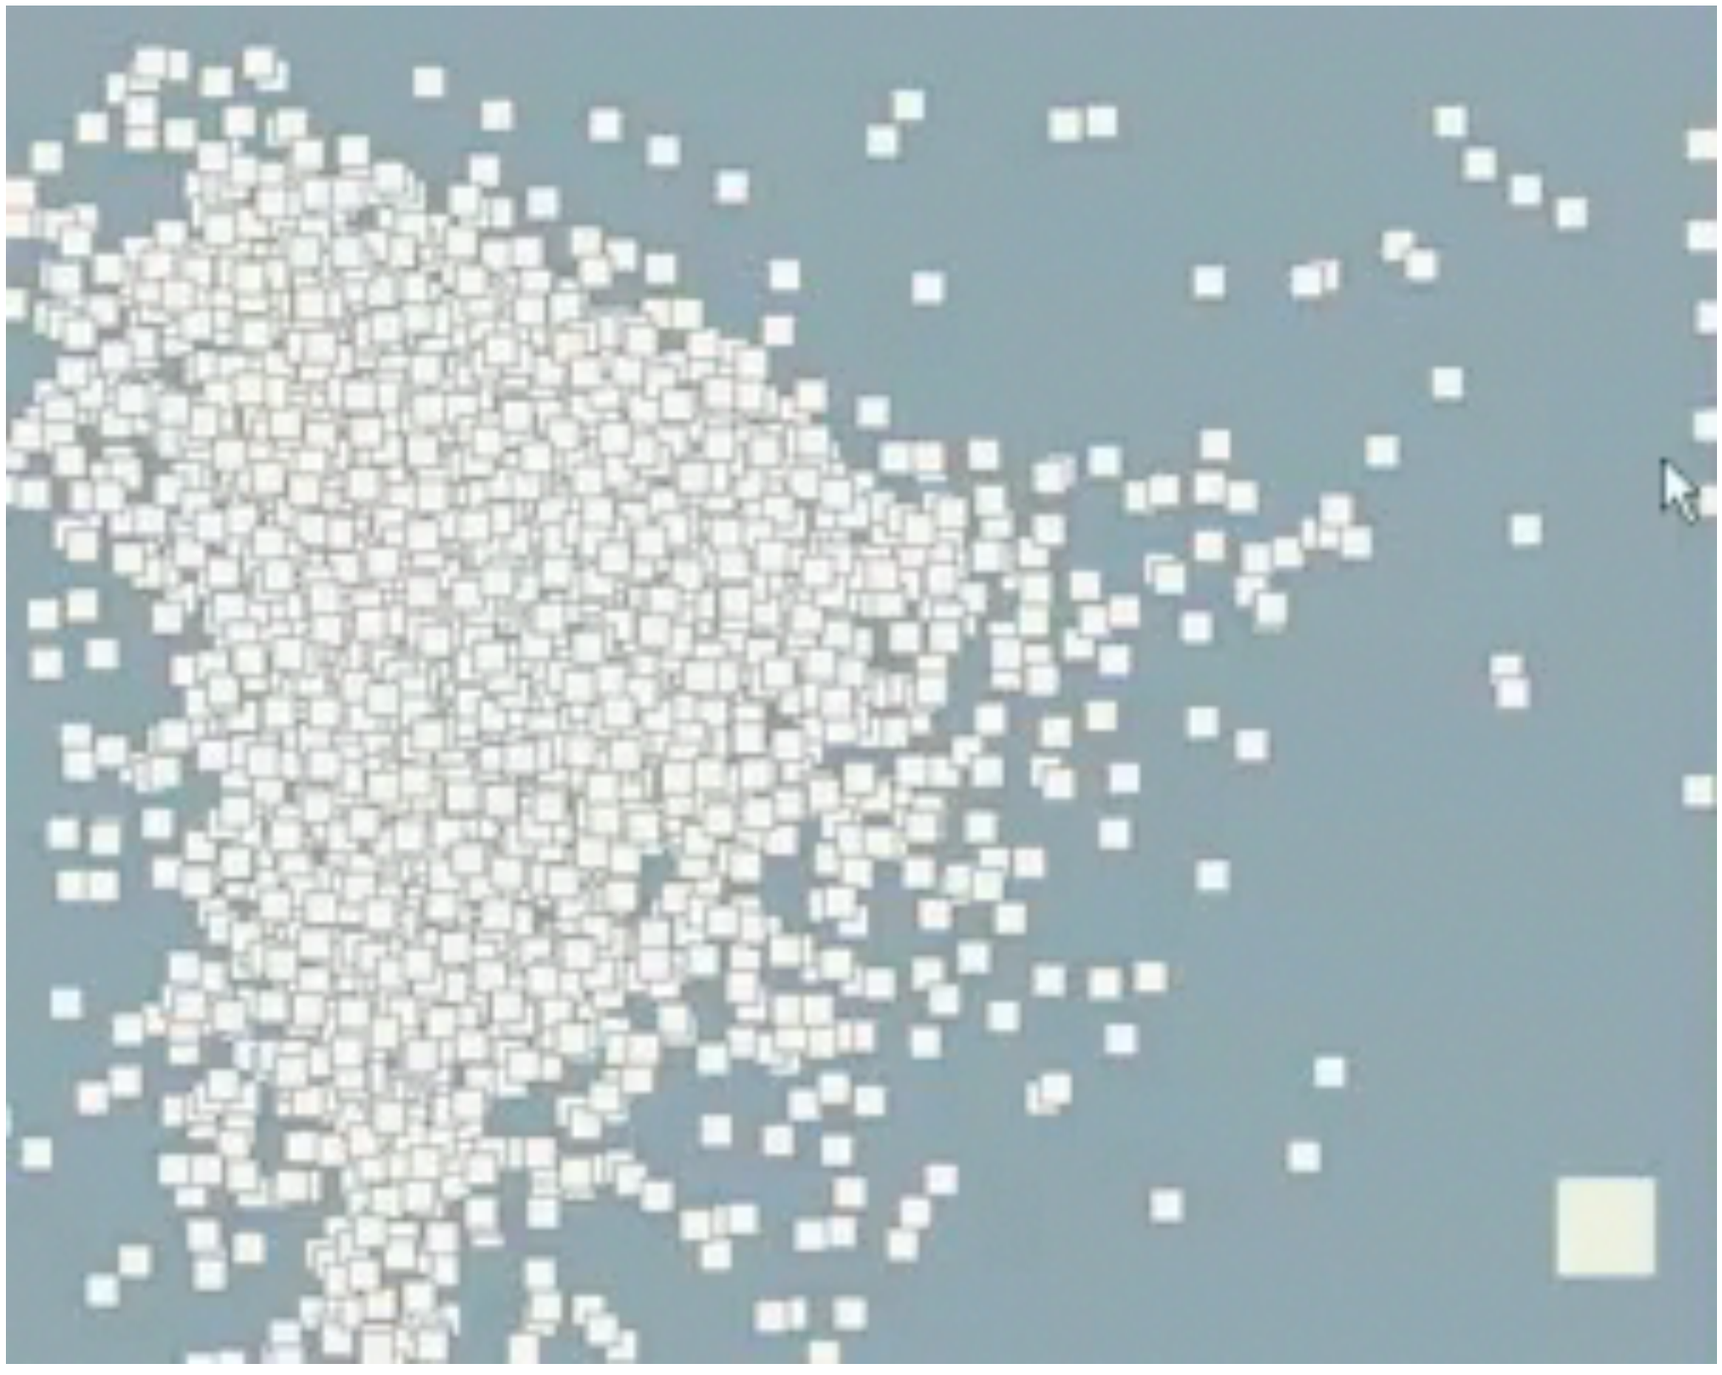

Graham Coleman - http://www.iua.upf.edu/~gcoleman/pubs/mused07.pdf

Time Series

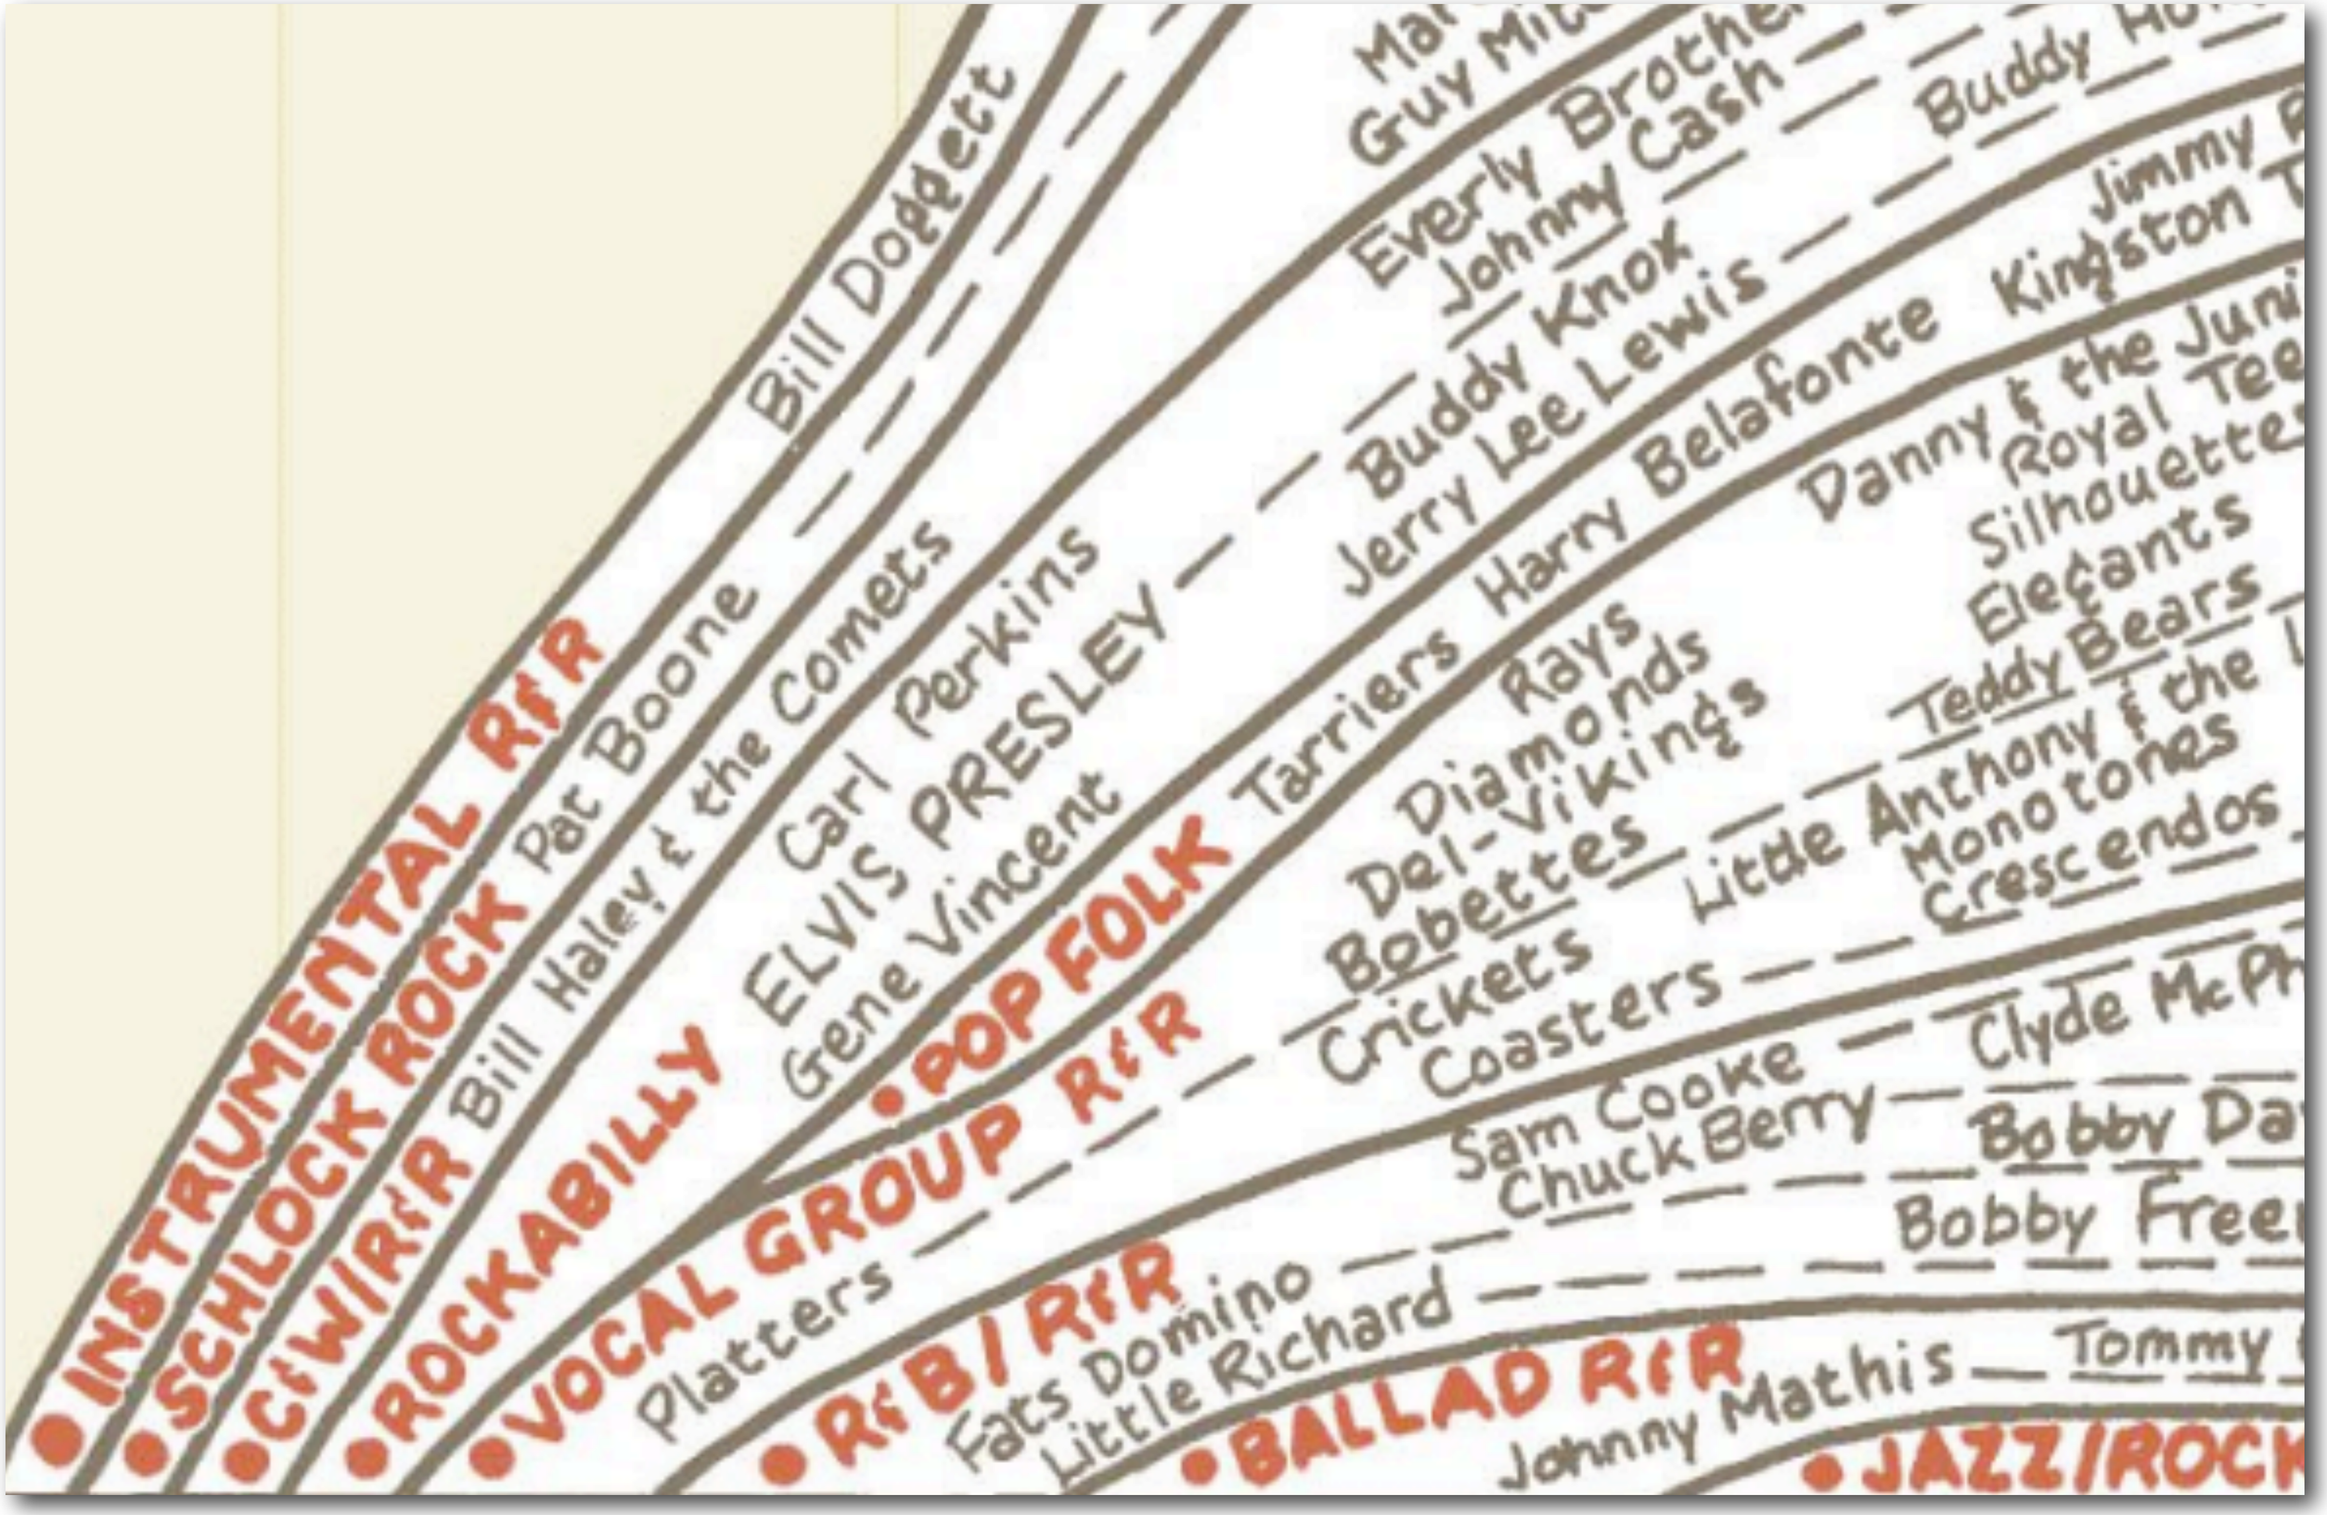

INSTRUMENTAL R&R
Bill Doggett
SCHLOCK ROCK
Pat Boone
C&W/R&R
Bill Haley & the Comets
ROCKABILLY
Carl Perkins
ELVIS PRESLEY
Gene Vincent
Guy Mitc
Everly Brother
Johnny Cash
Buddy Knox
Jerry Lee Lewis
Buddy Ho
POP FOLK
Tarriers
Harry Belafonte
Kingston T
Jimmy
VOCAL GROUP R&R
Platters
Rays
Diamonds
Del-Vikings
Bobettes
Crickets
Coasters
Danny & the Juni
Royal Tee
Silhouettes
Elegants
Teddy Bears
Little Anthony & the
Monotones
Crescendos
R&B/R&R
Fats Domino
Little Richard
Sam Cooke
Chuck Berry
Clyde McPh
Bobby Da
Bobby Free
BALLAD R&R
Johnny Mathis
Tommy
JAZZ/ROCK

last.fm spiral

http://www.diametunim.com/muse/

Visualizing Listening History

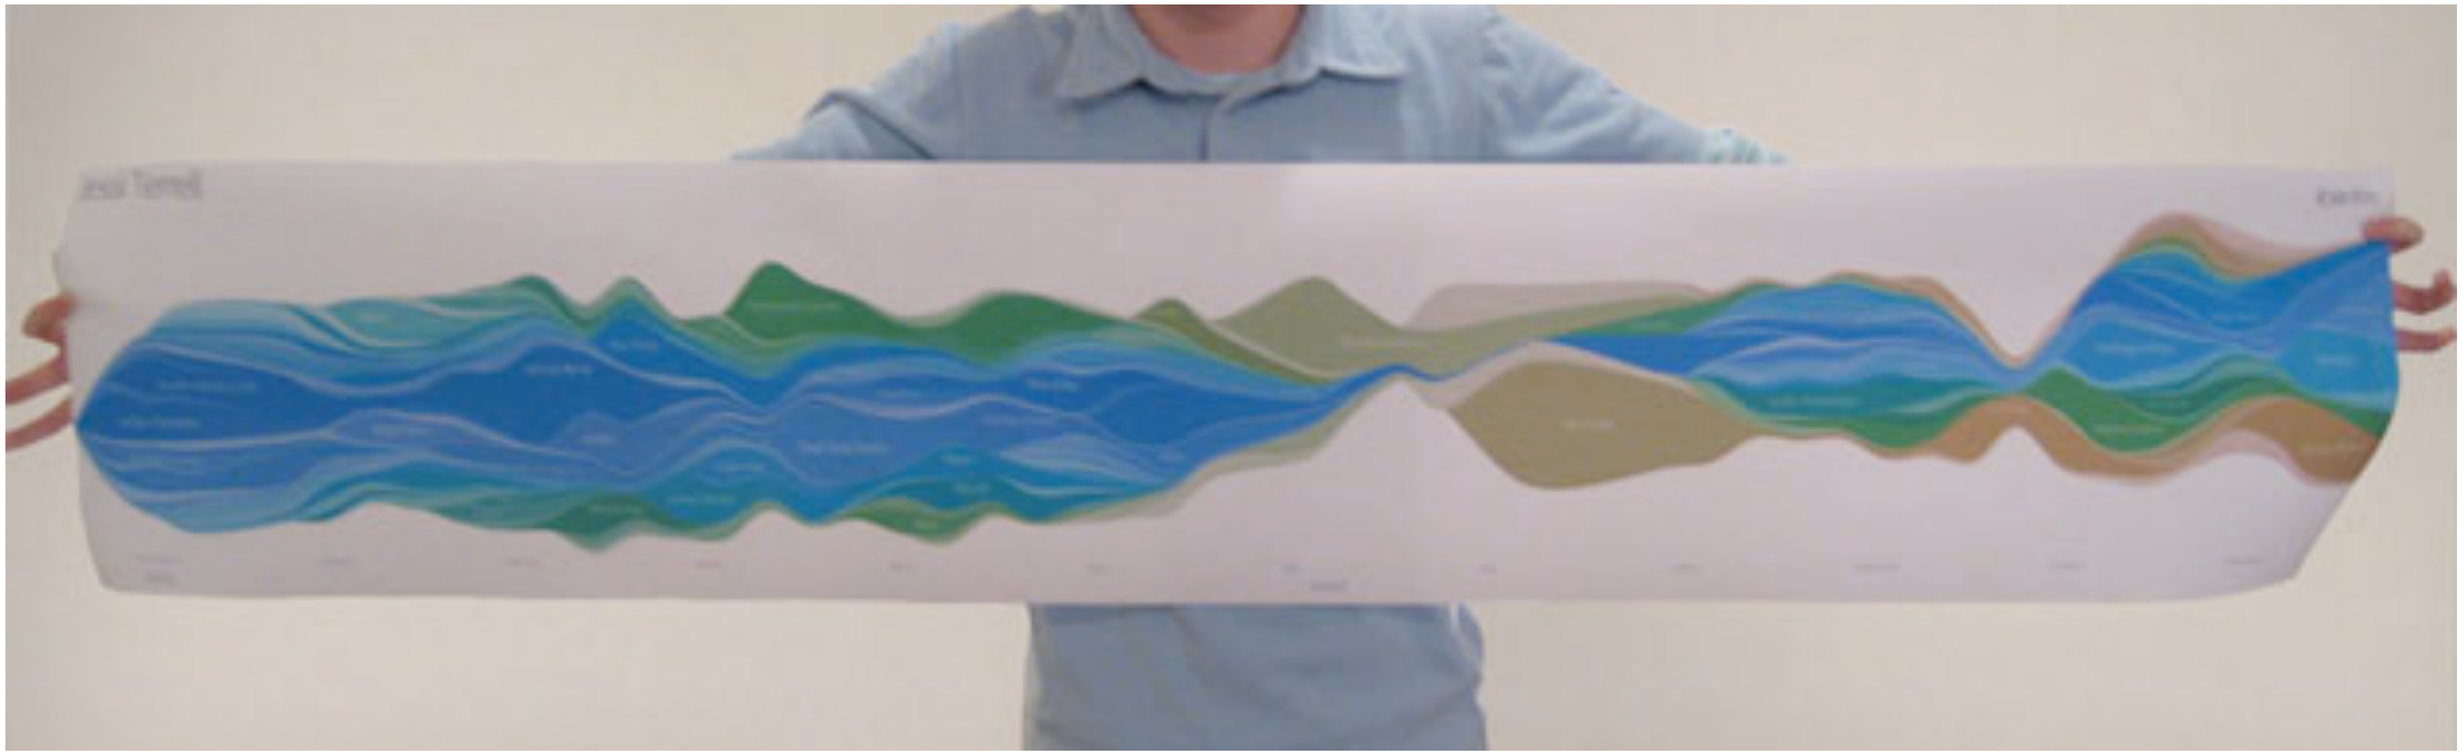

http://www.leebyron.com/what/lastfm/ - Lee Byron

Visualizing Listening History

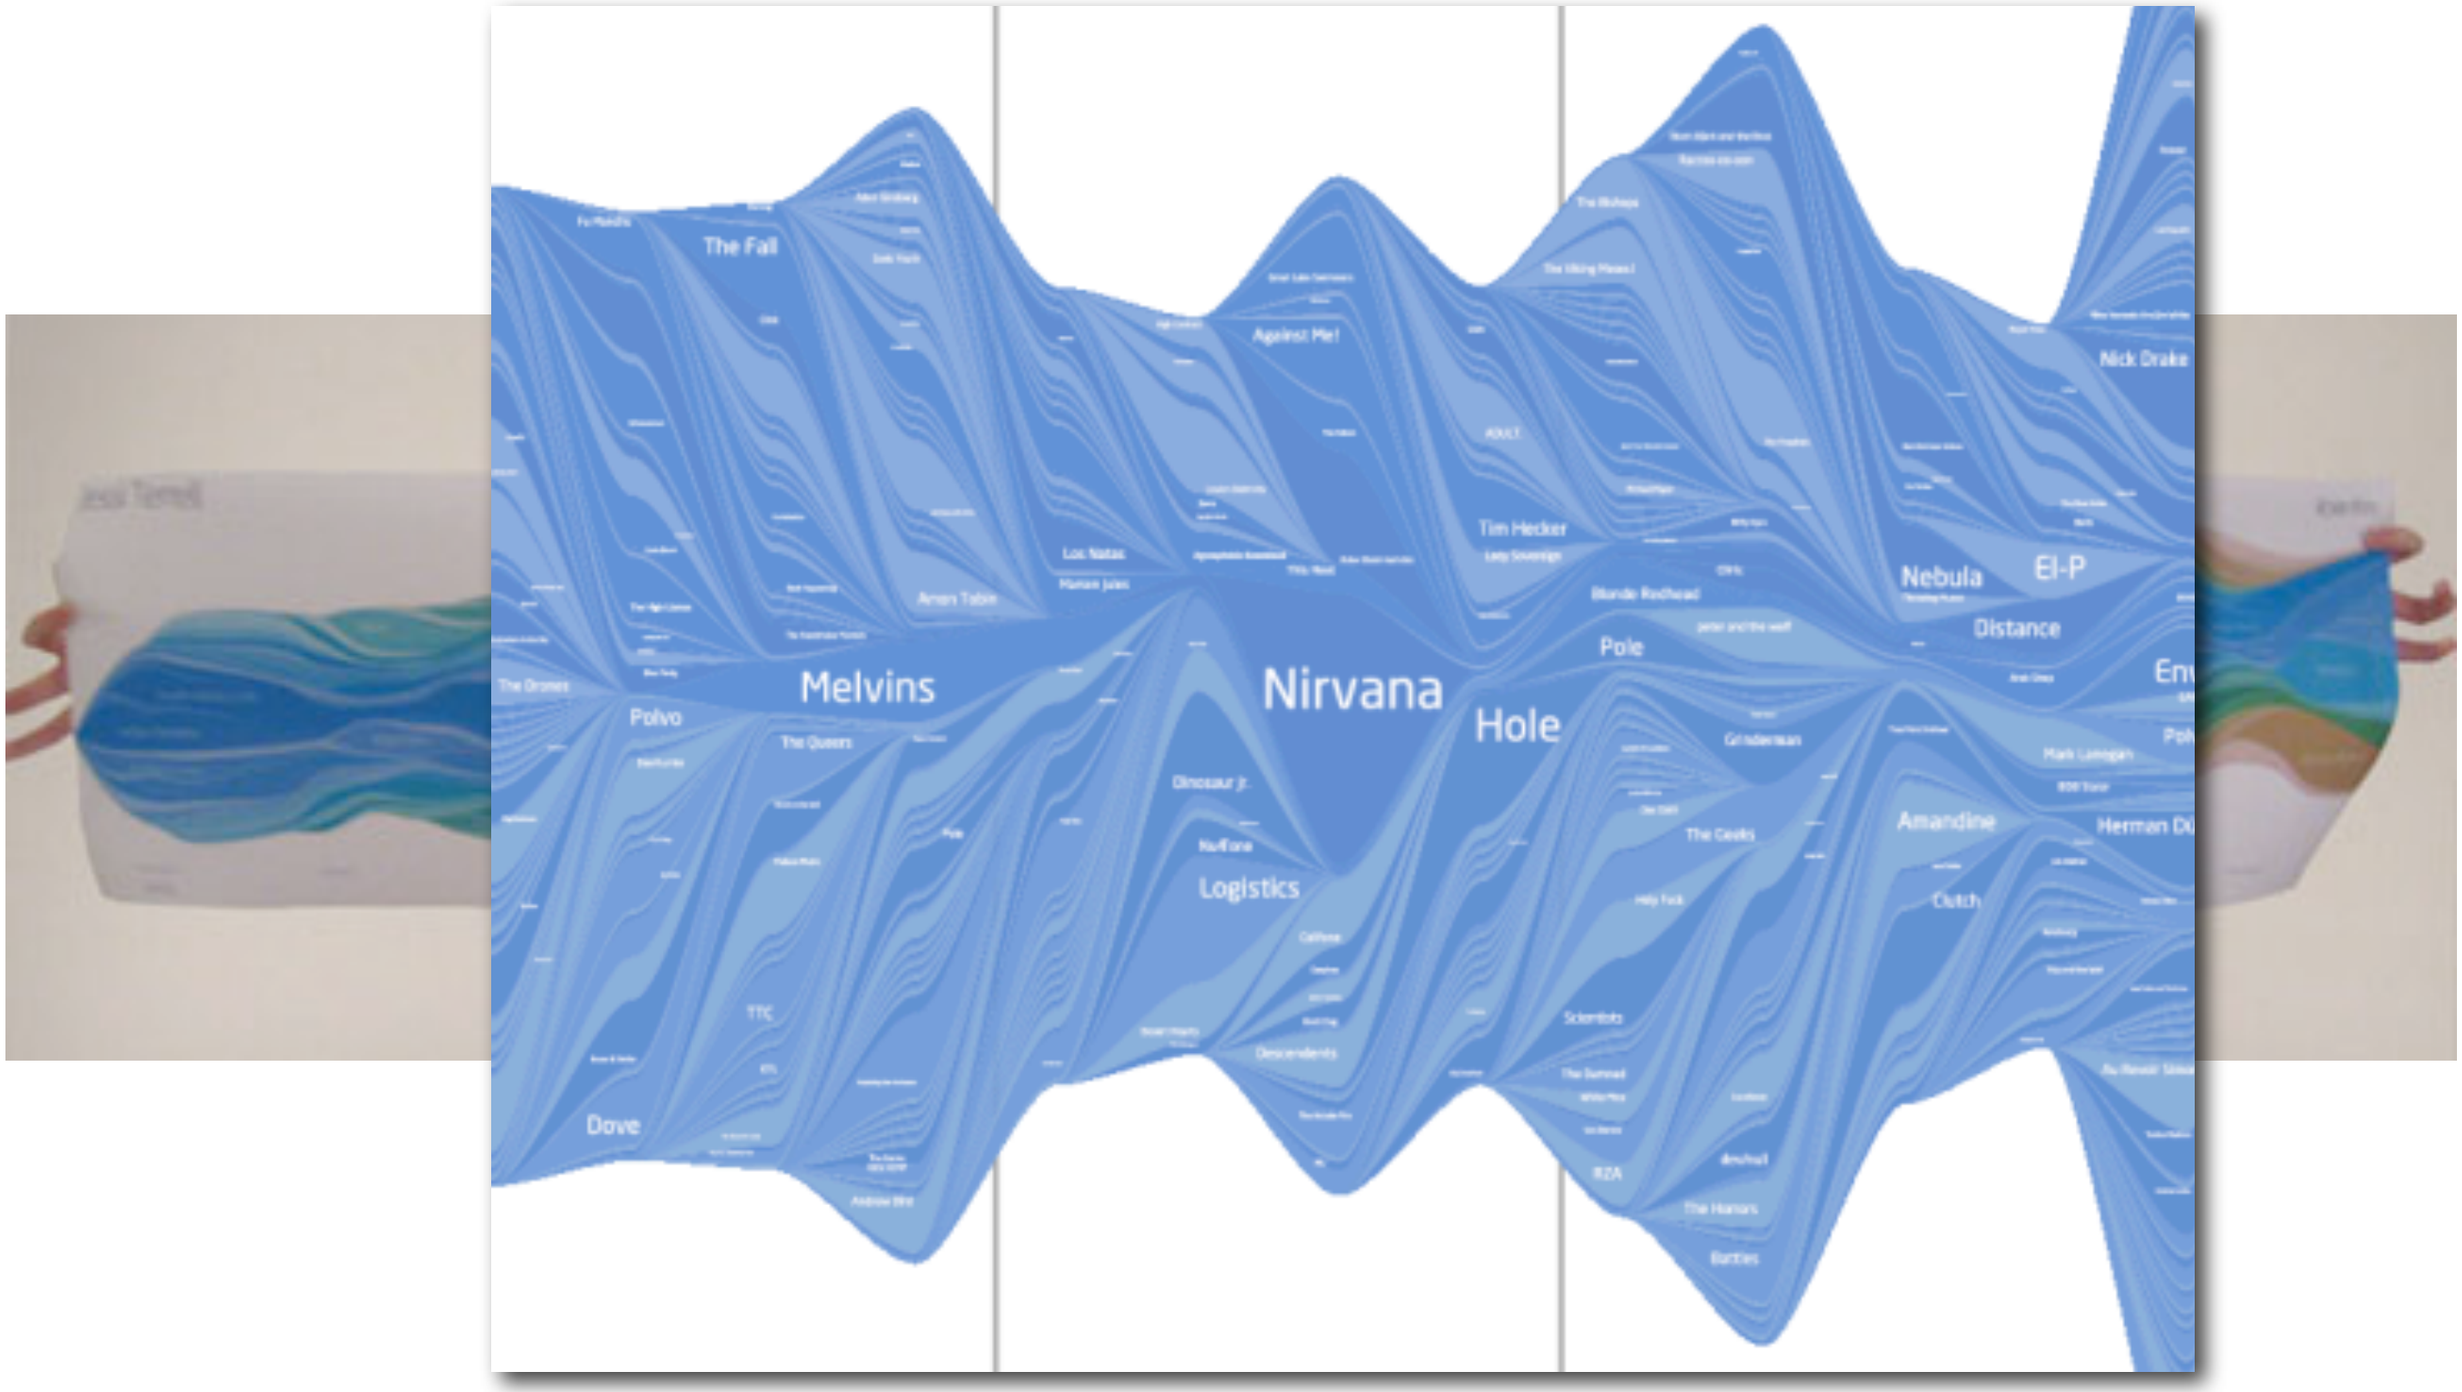

http://www.leebyron.com/what/lastfm/ - Lee Byron

Radio Waves

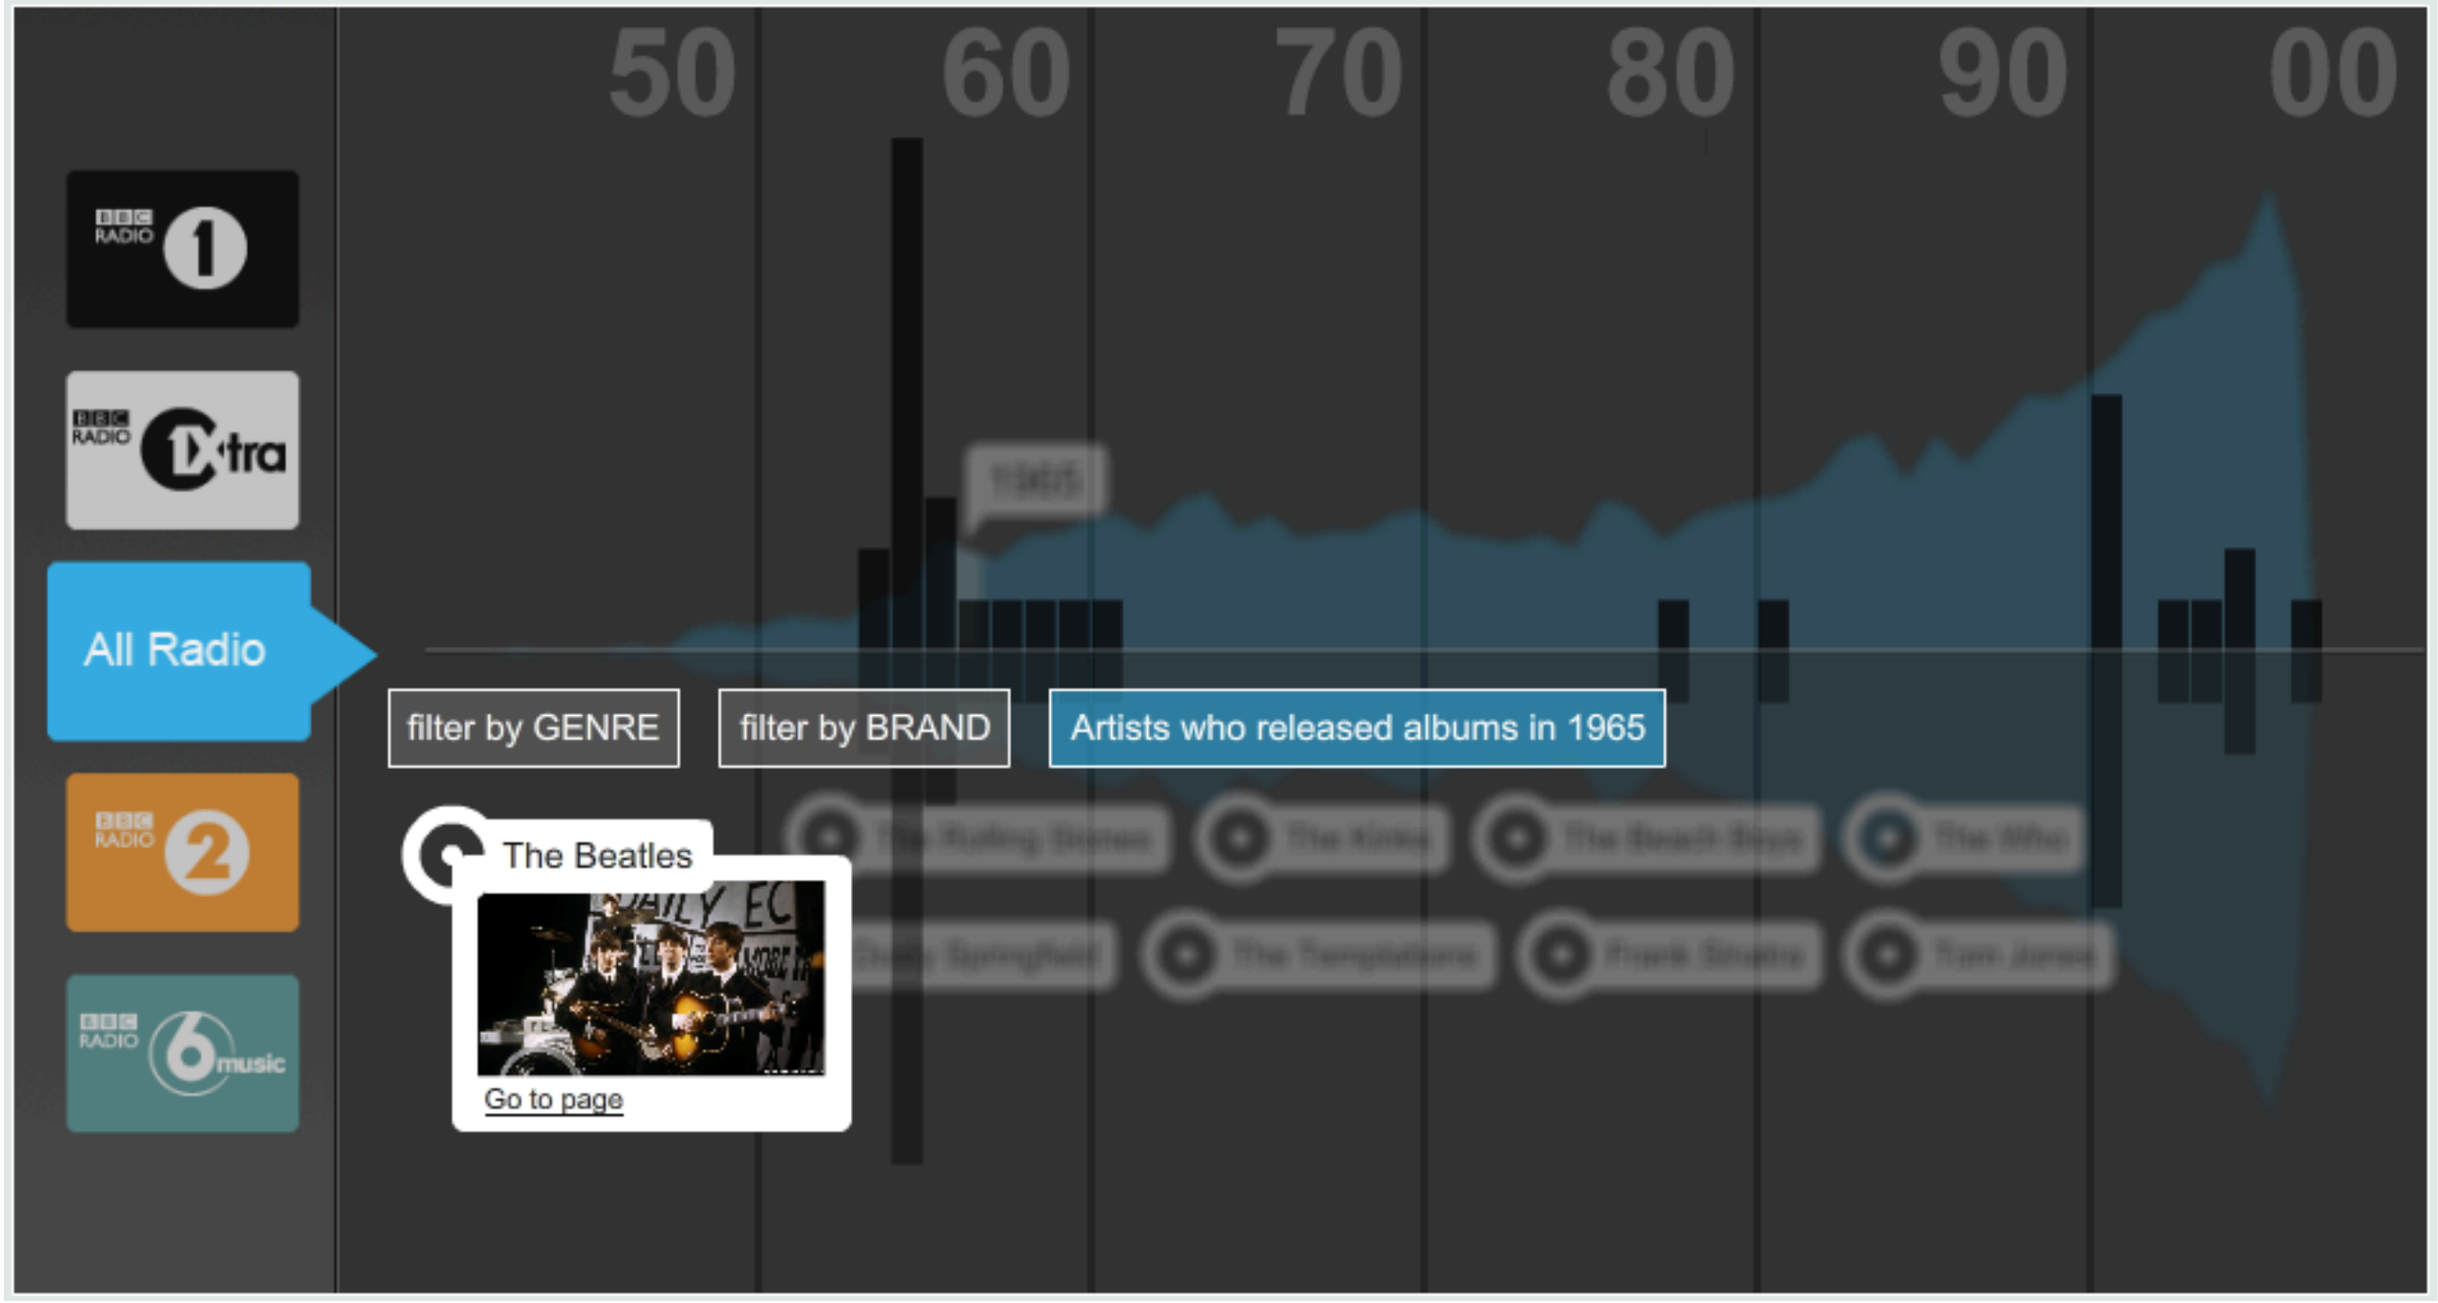

http://www.bbc.co.uk/radio/labs/radiowaves/

Radio Waves

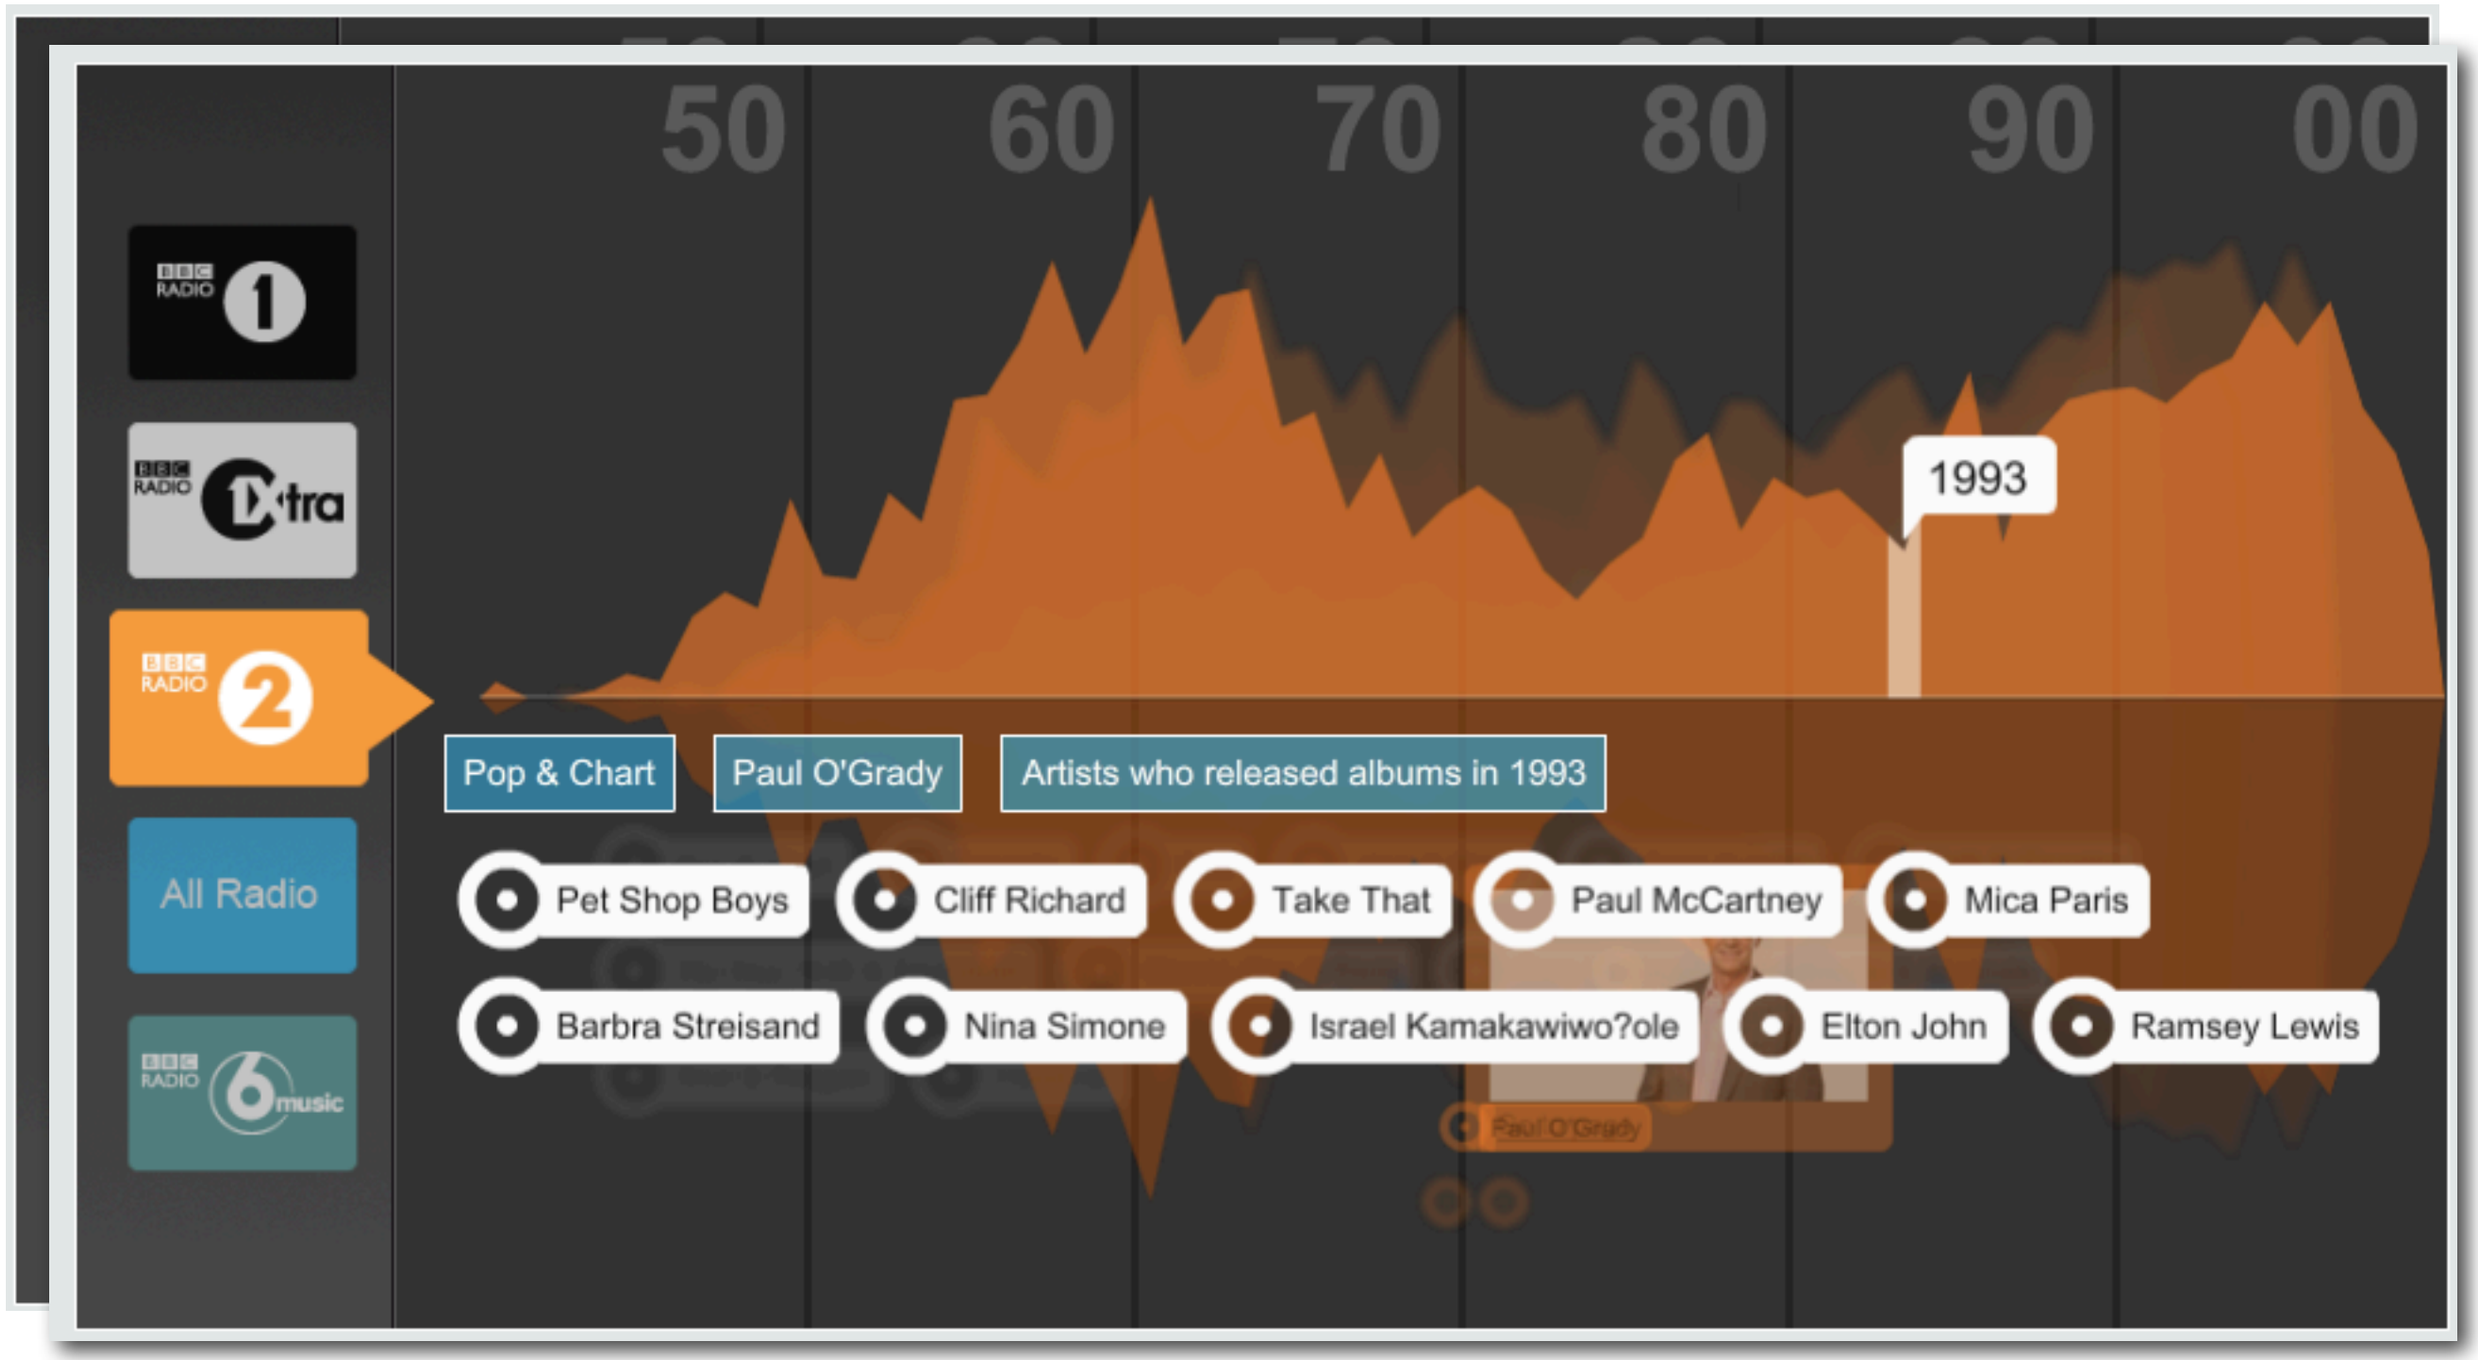

http://www.bbc.co.uk/radio/labs/radiowaves/

Pulling strings from a tangle

http://bowr.de/blog/?p=49
Dominikus Baur University of Munich

Pulling strings from a tangle

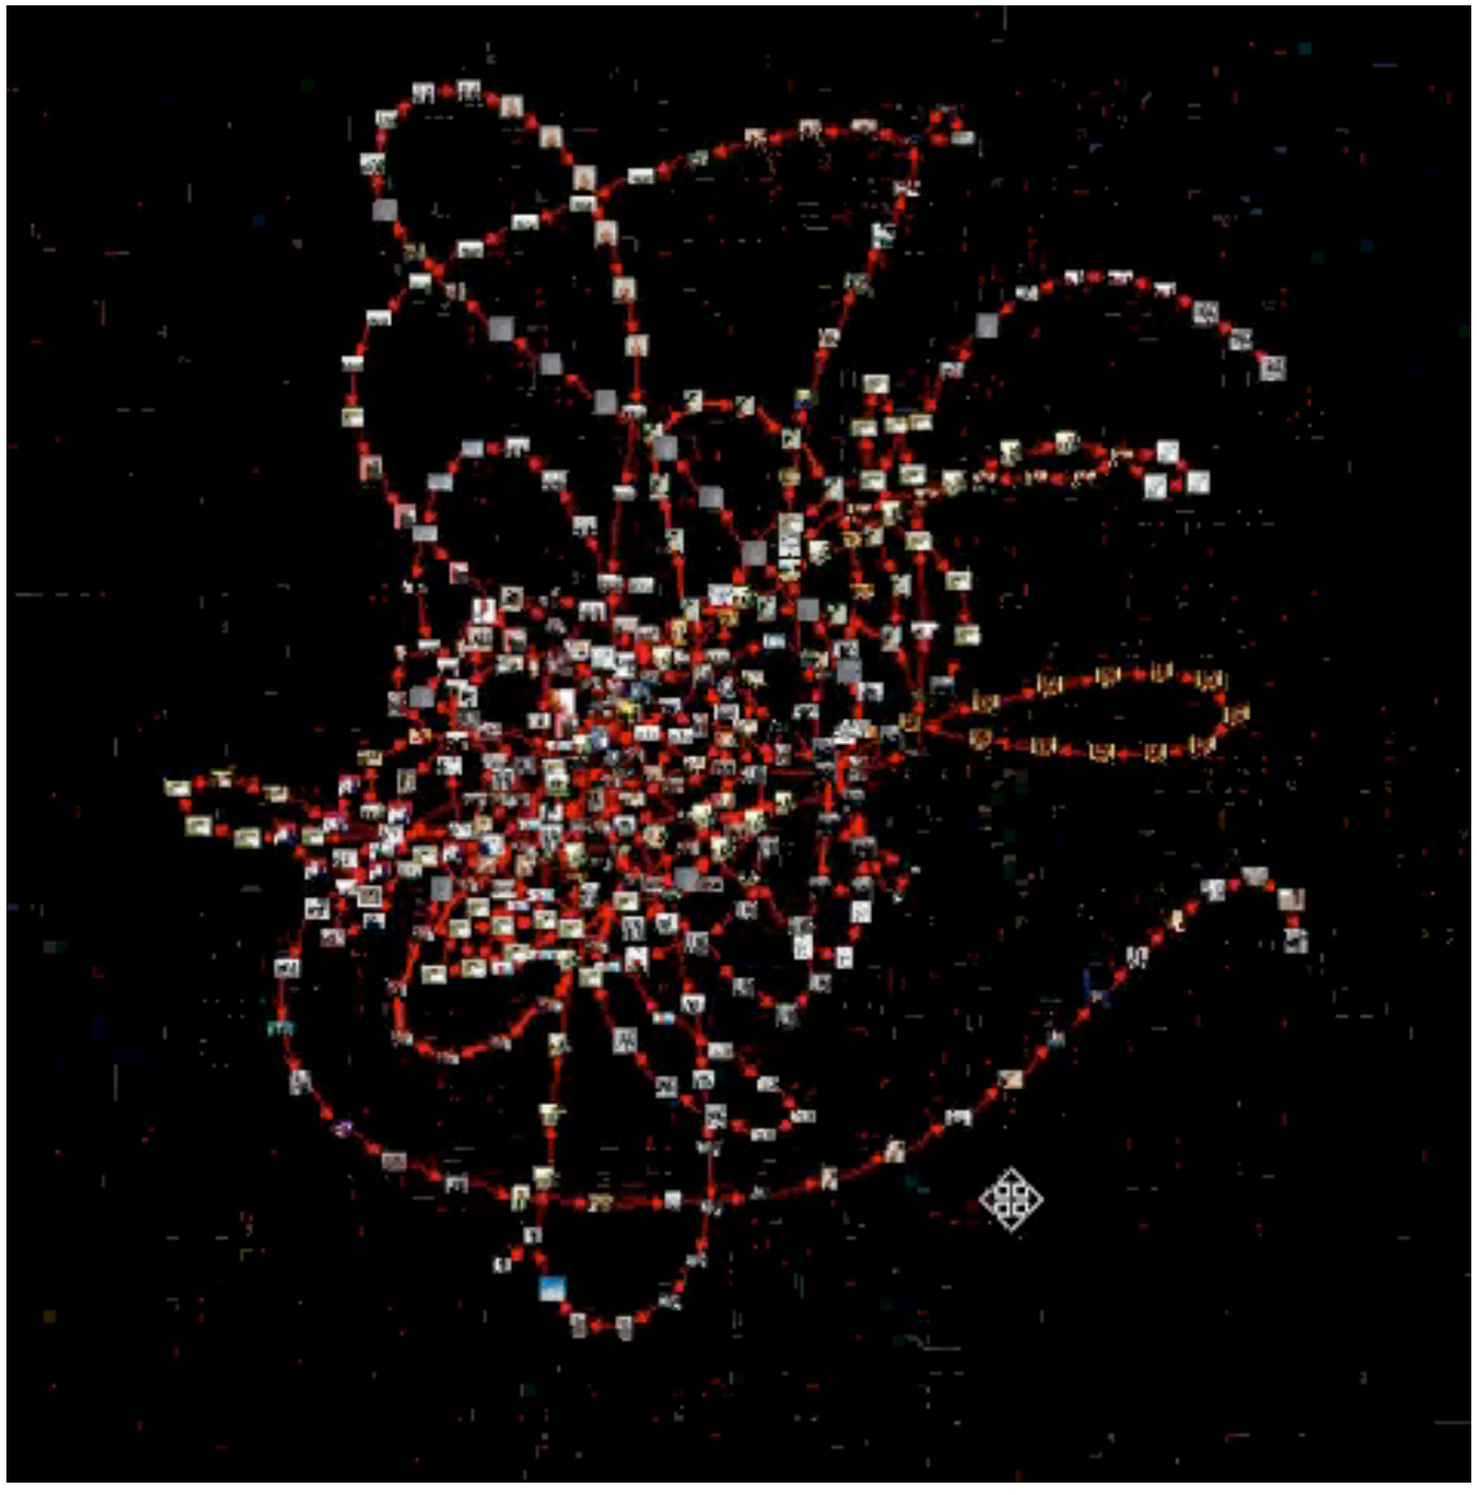

http://bowr.de/blog/?p=49
Dominikus Baur University of Munich

Trees

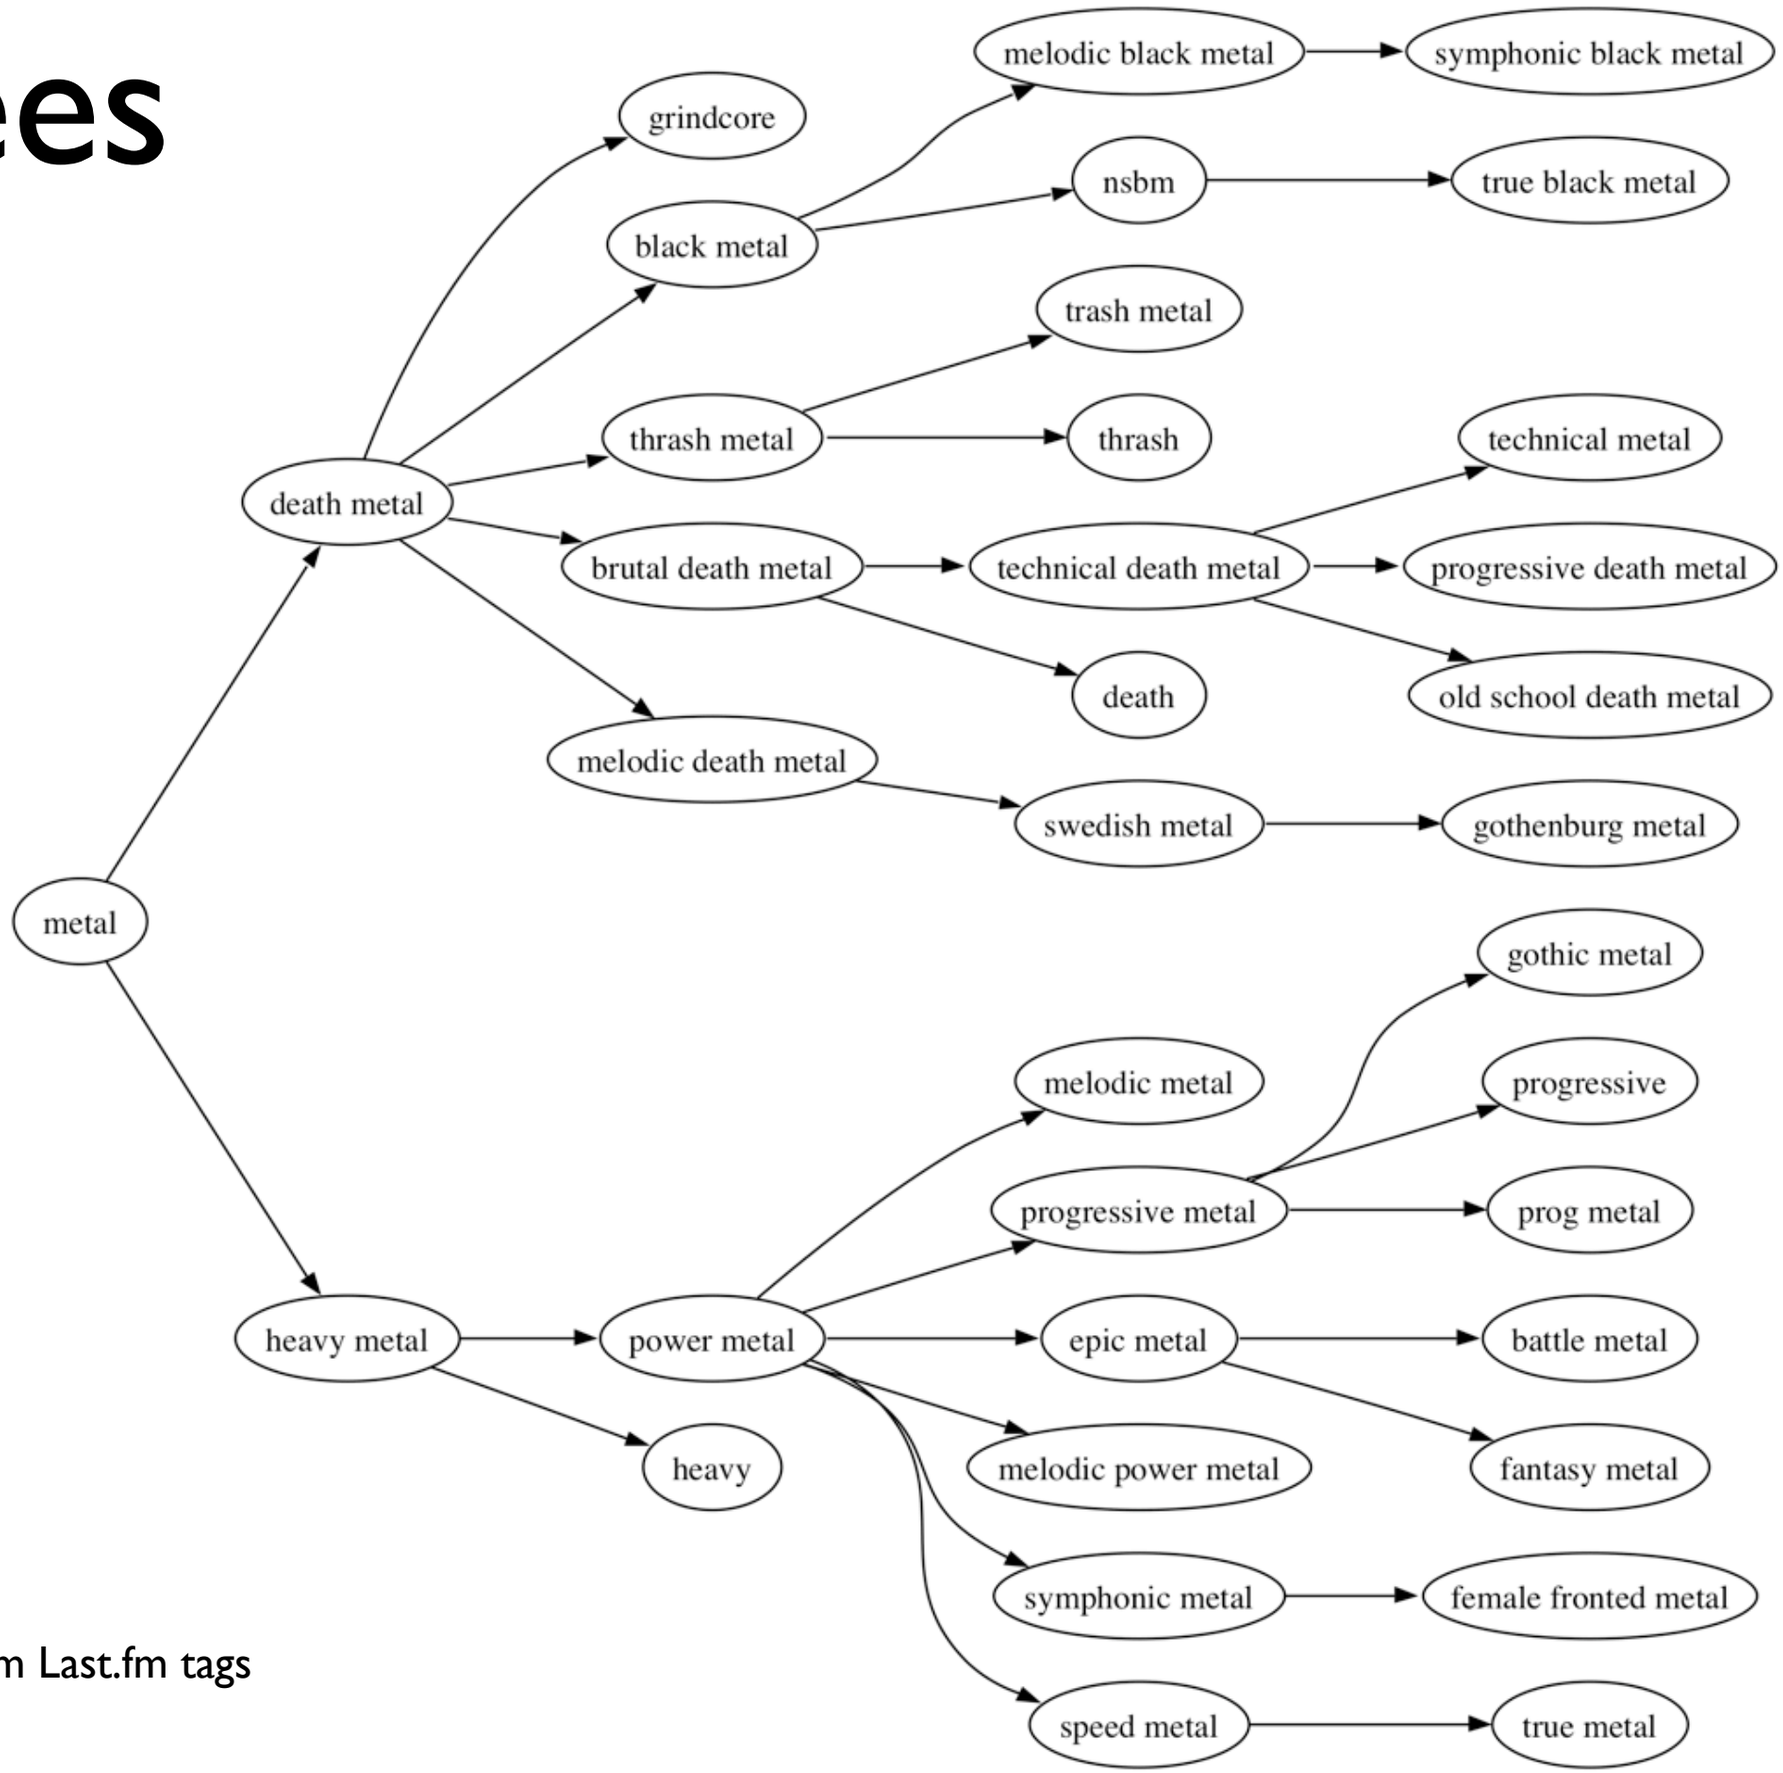

Genre Map derived from Last.fm tags
Lamere

Trees

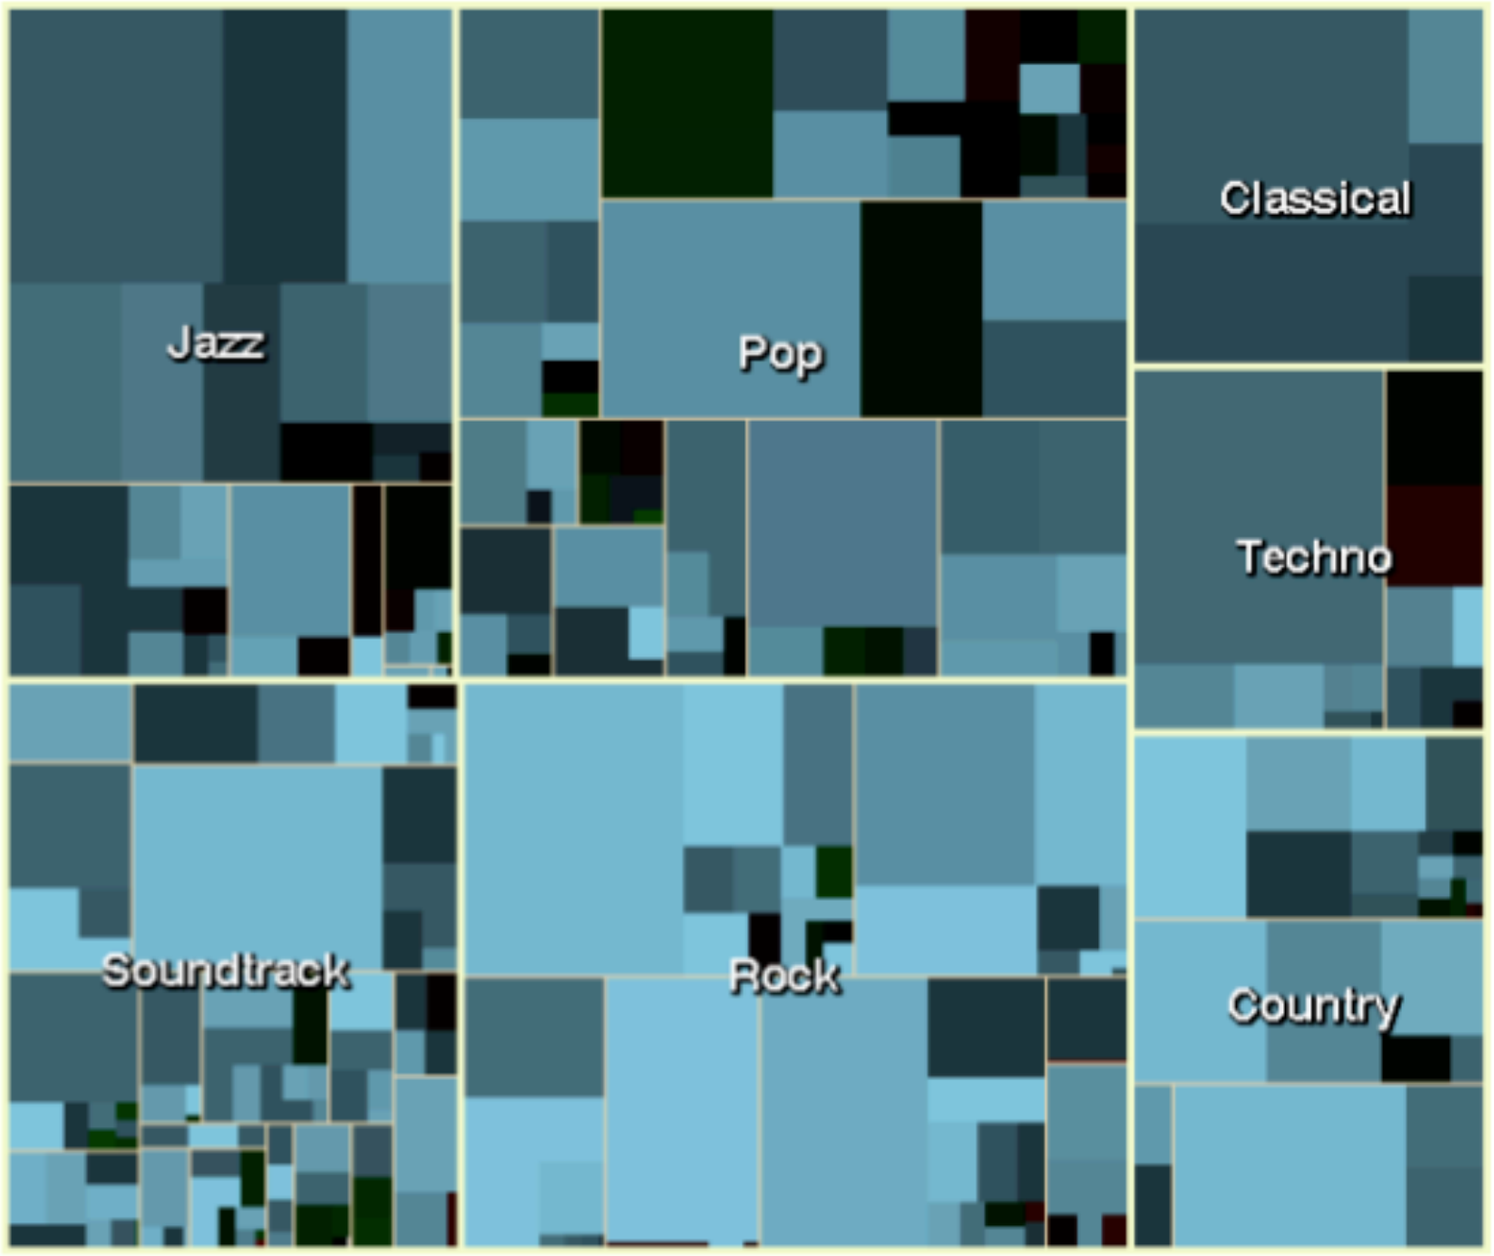

Visualizing and exploring personal music libraries
Torrens, Hertzog, Arcos

Treemap of a personal music collection

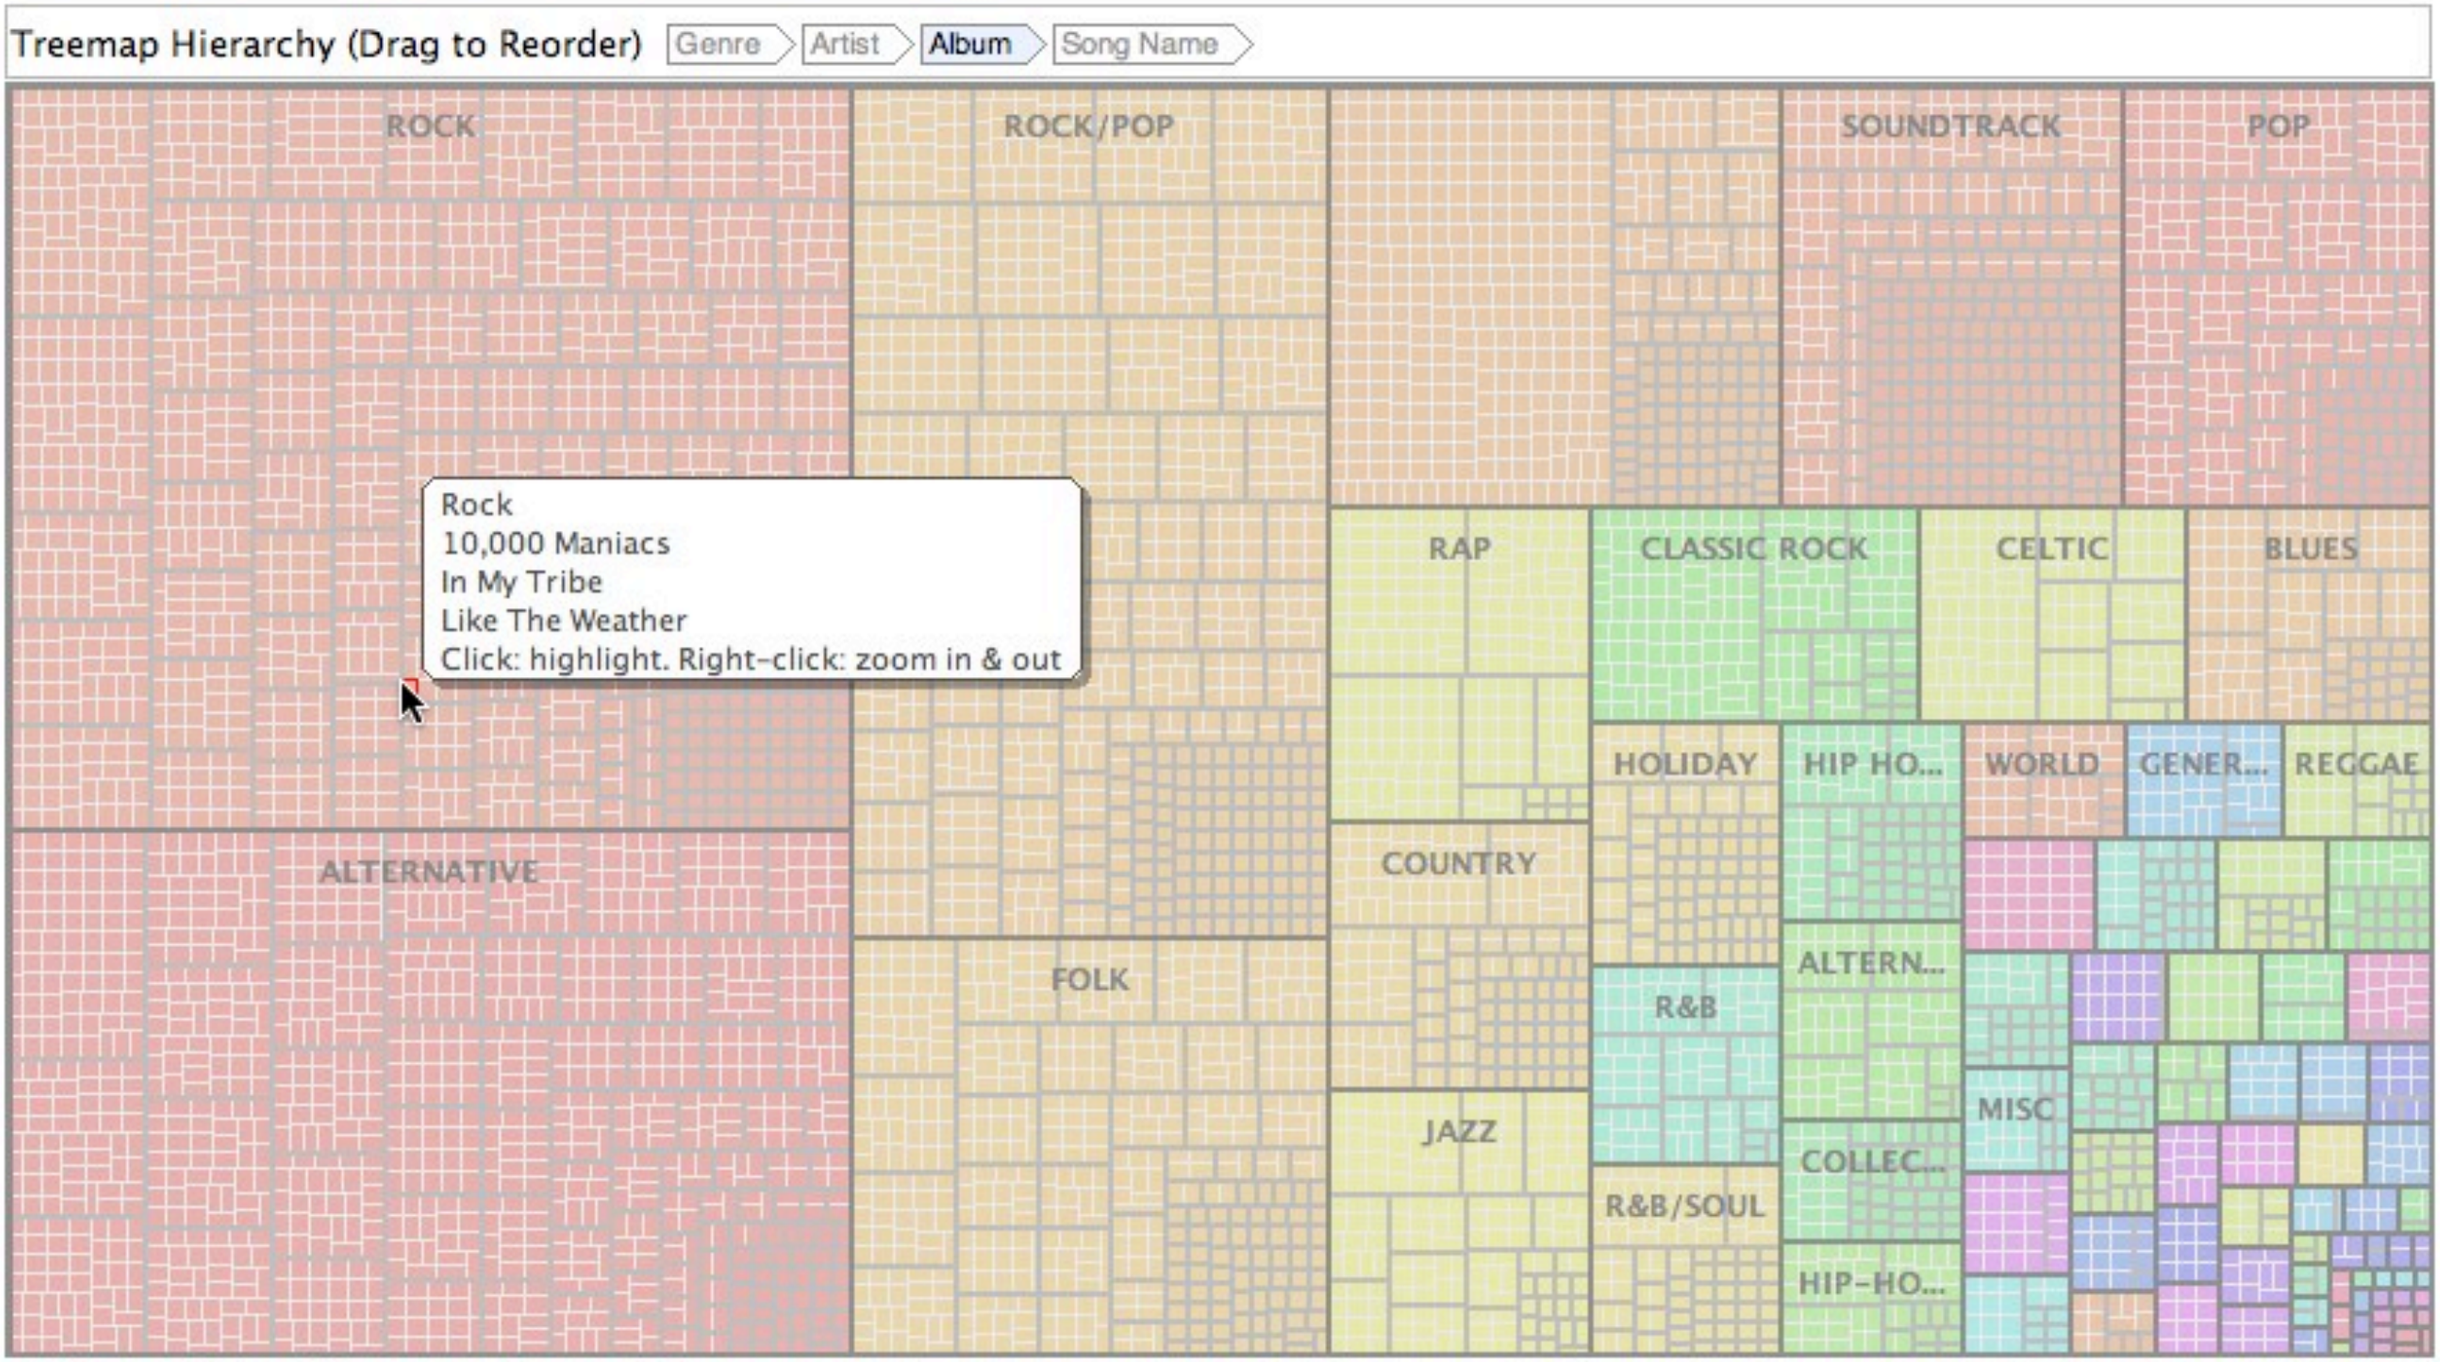

Visualizations : Dave's Music Library Contents

Interactive

Mirrored Radio dial - by thomwatson

SongExplorer

SongExplorer: Exploring Large Musical Databases Using a Tabletop Interface
Carles F. Julià - Music Technology Group - Universitat Pompeu Fabra

SongExplorer

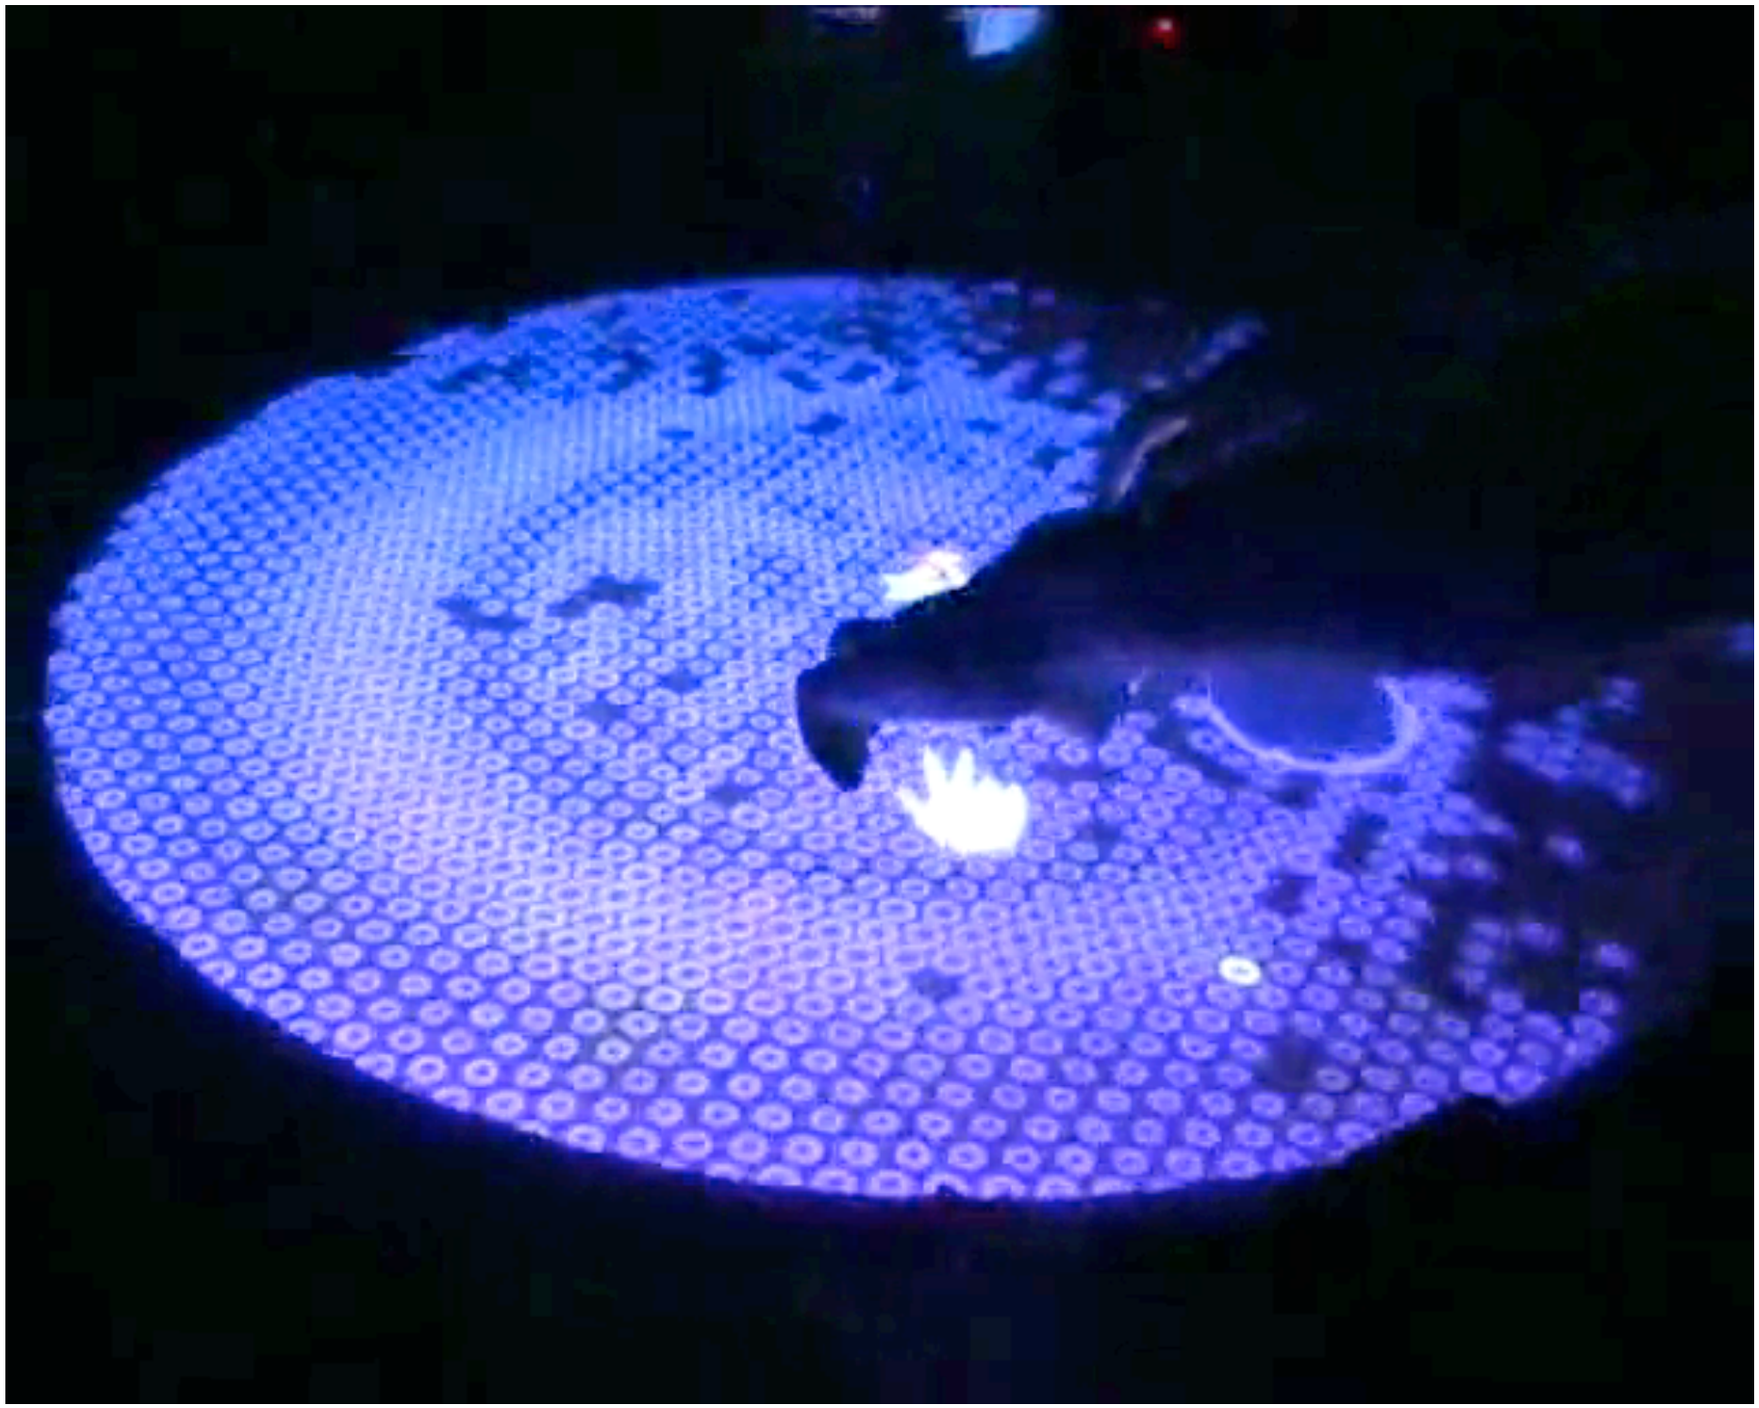

SongExplorer: Exploring Large Musical Databases Using a Tabletop Interface
Carles F. Julià - Music Technology Group - Universitat Pompeu Fabra

Fidgt: Visualize

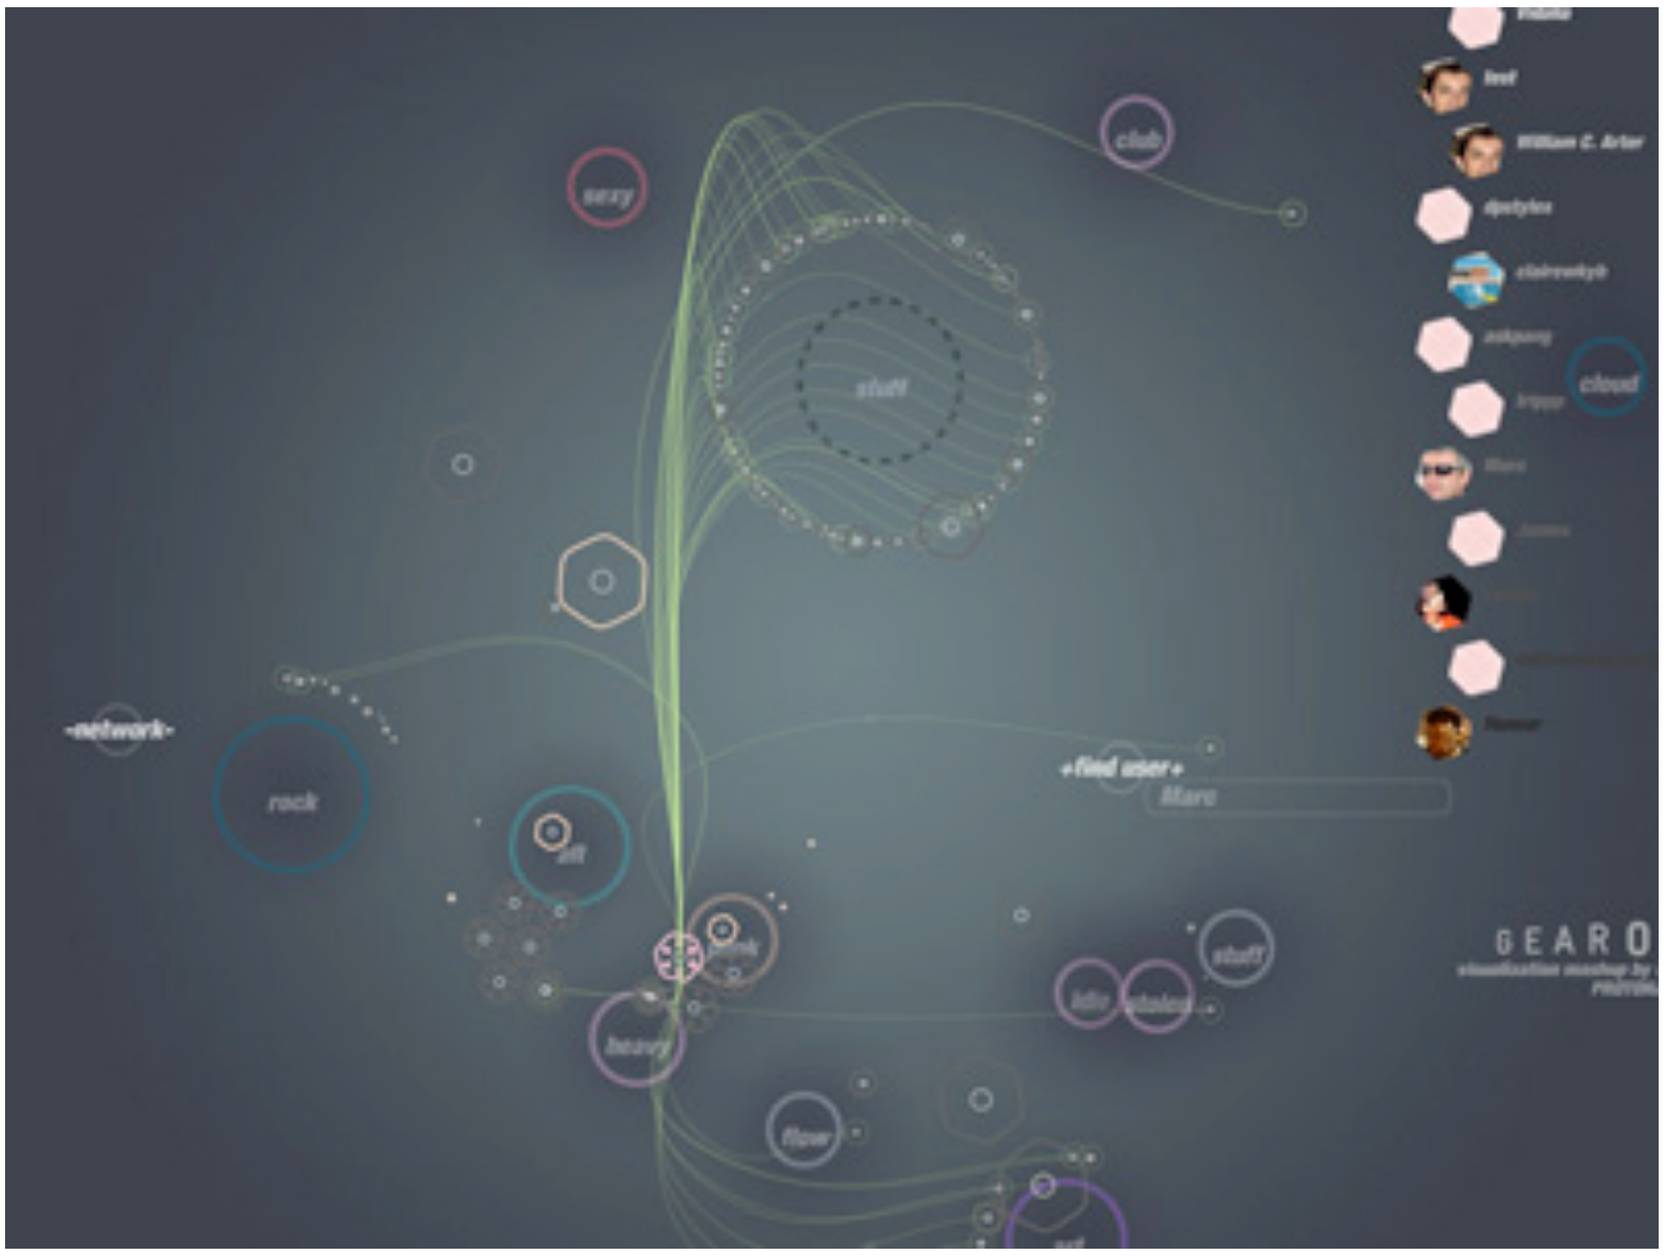

http://fidgt.com/visualize

Fidgt: Visualize

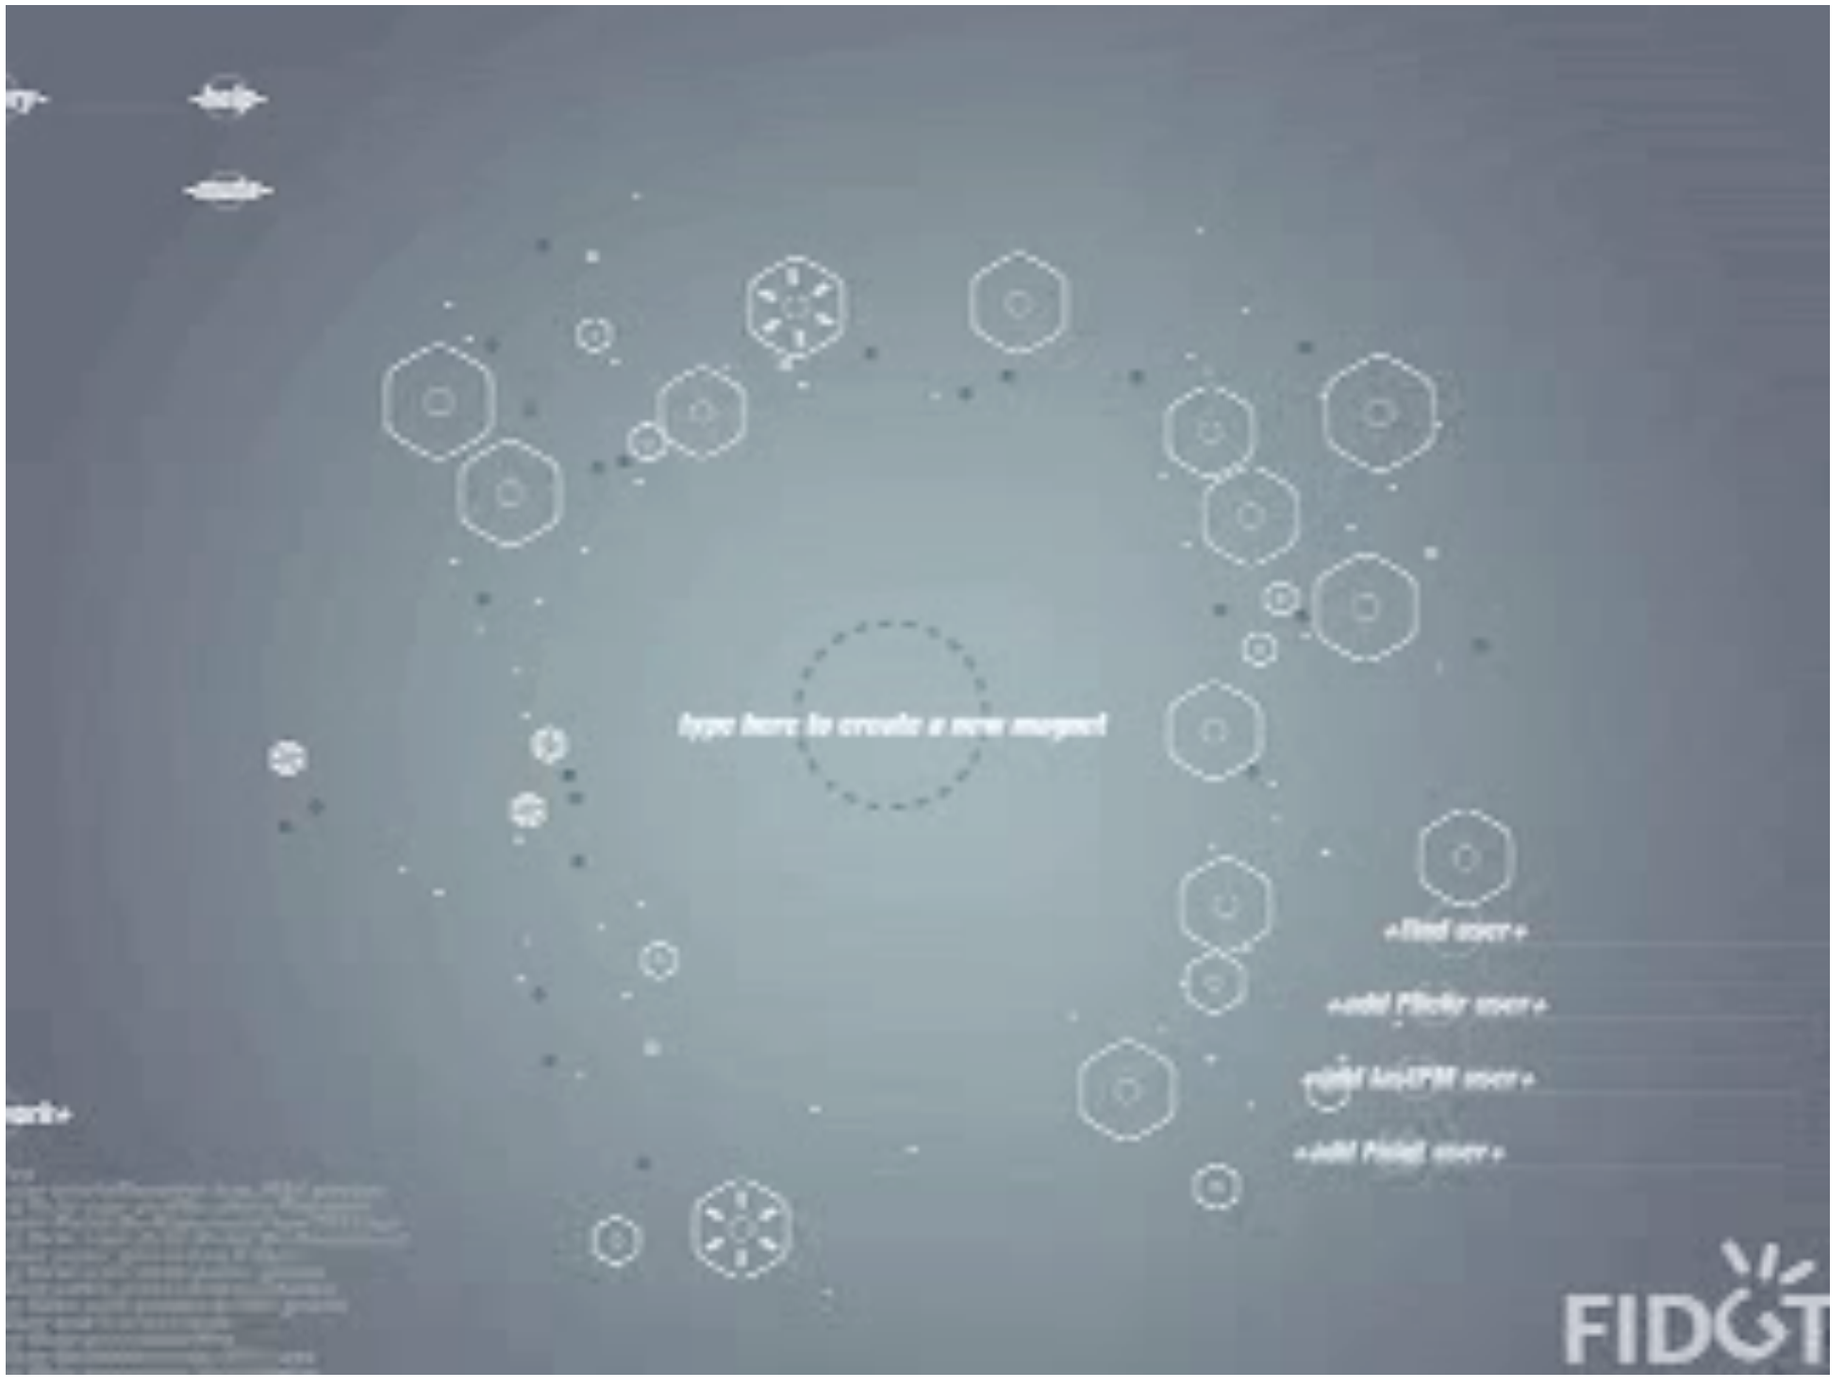

http://fidgt.com/visualize

MusicNodes

http://waks.nl/nl/musicalnodes

Lisa Dalhuijsen, Lieven van Velthoven -LIACS, Leiden University

MusicNodes

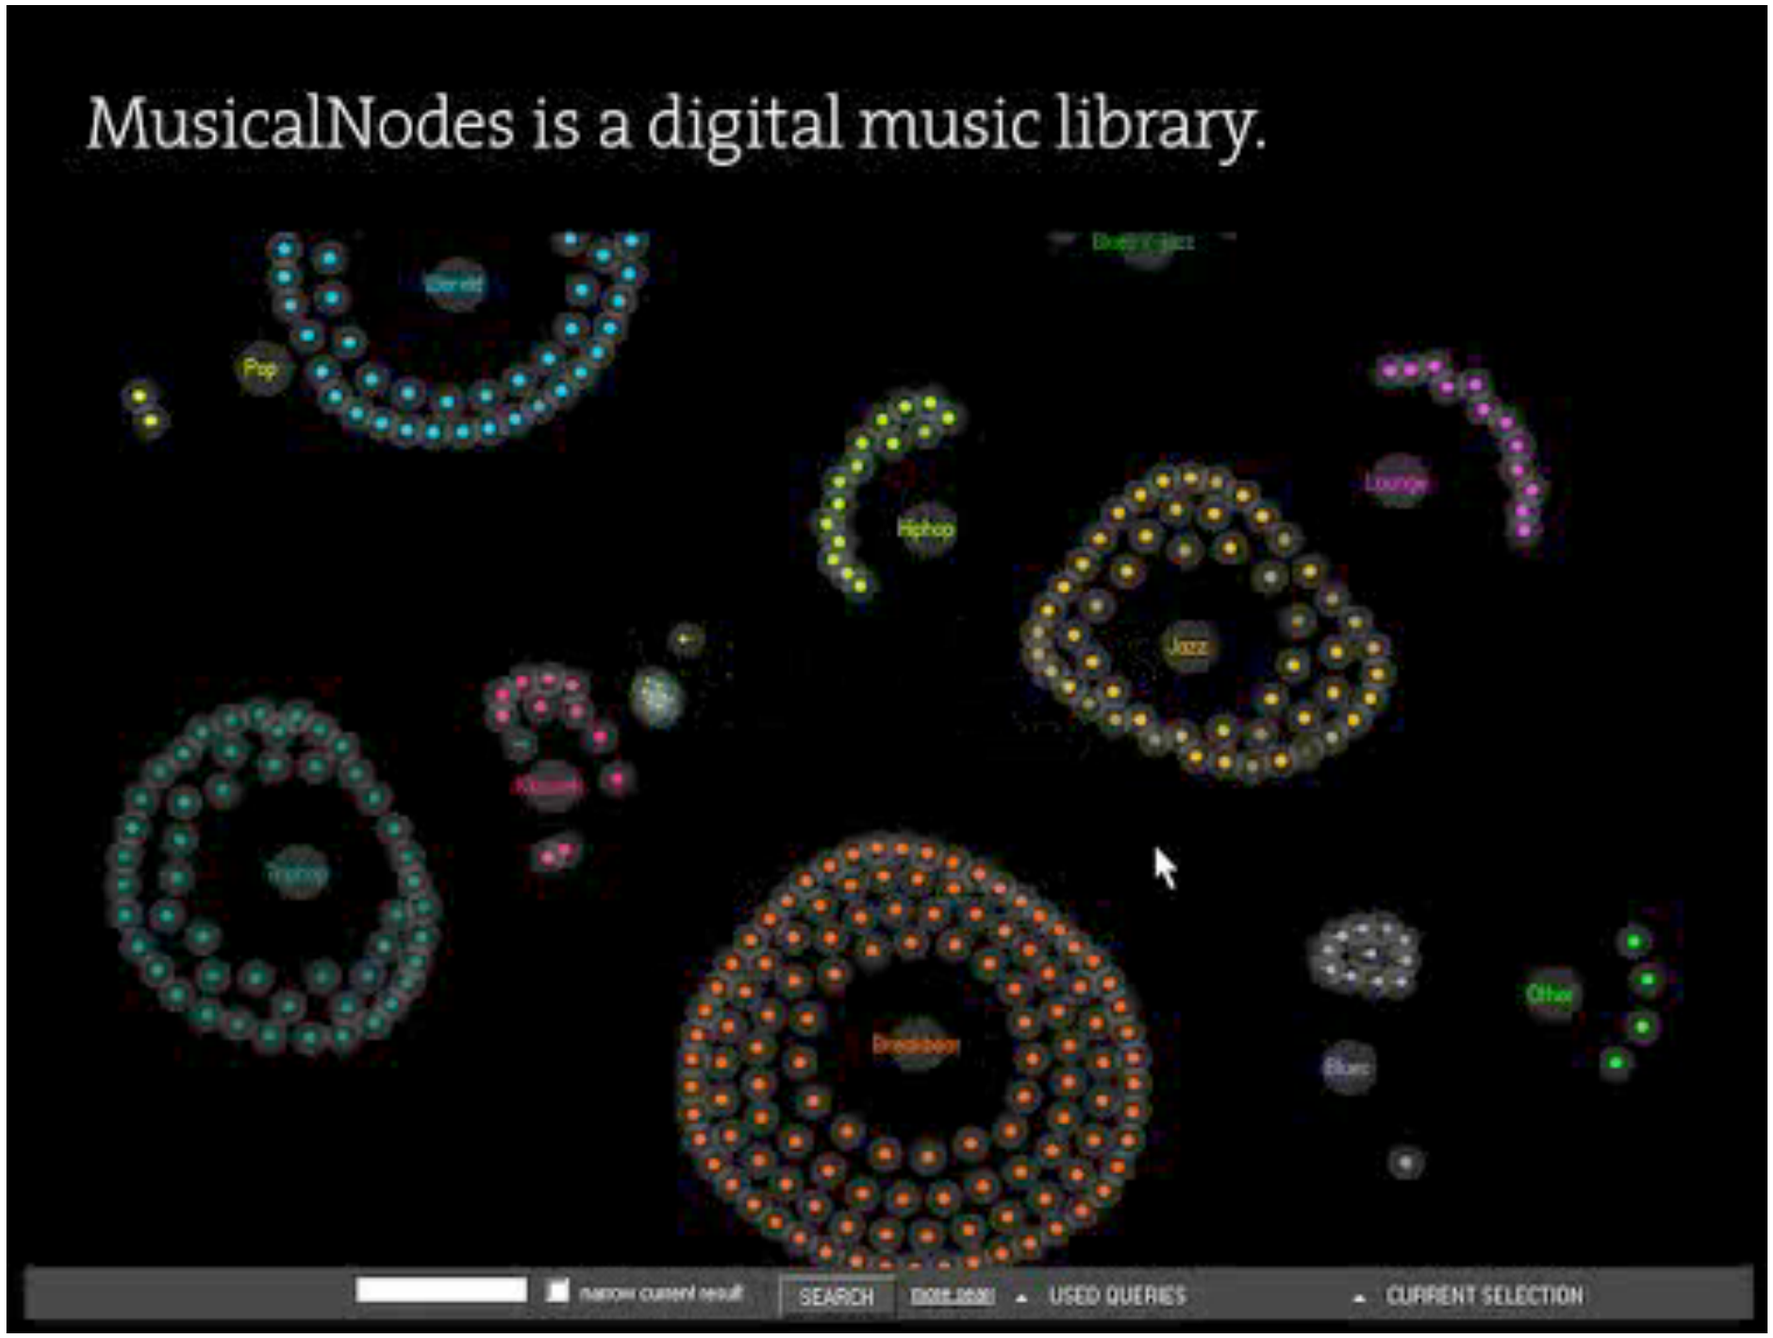

http://waks.nl/nl/musicalnodes

Lisa Dalhuijsen, Lieven van Velthoven -LIACS, Leiden University

Zune MixView

social.zune.net

Zune MixView

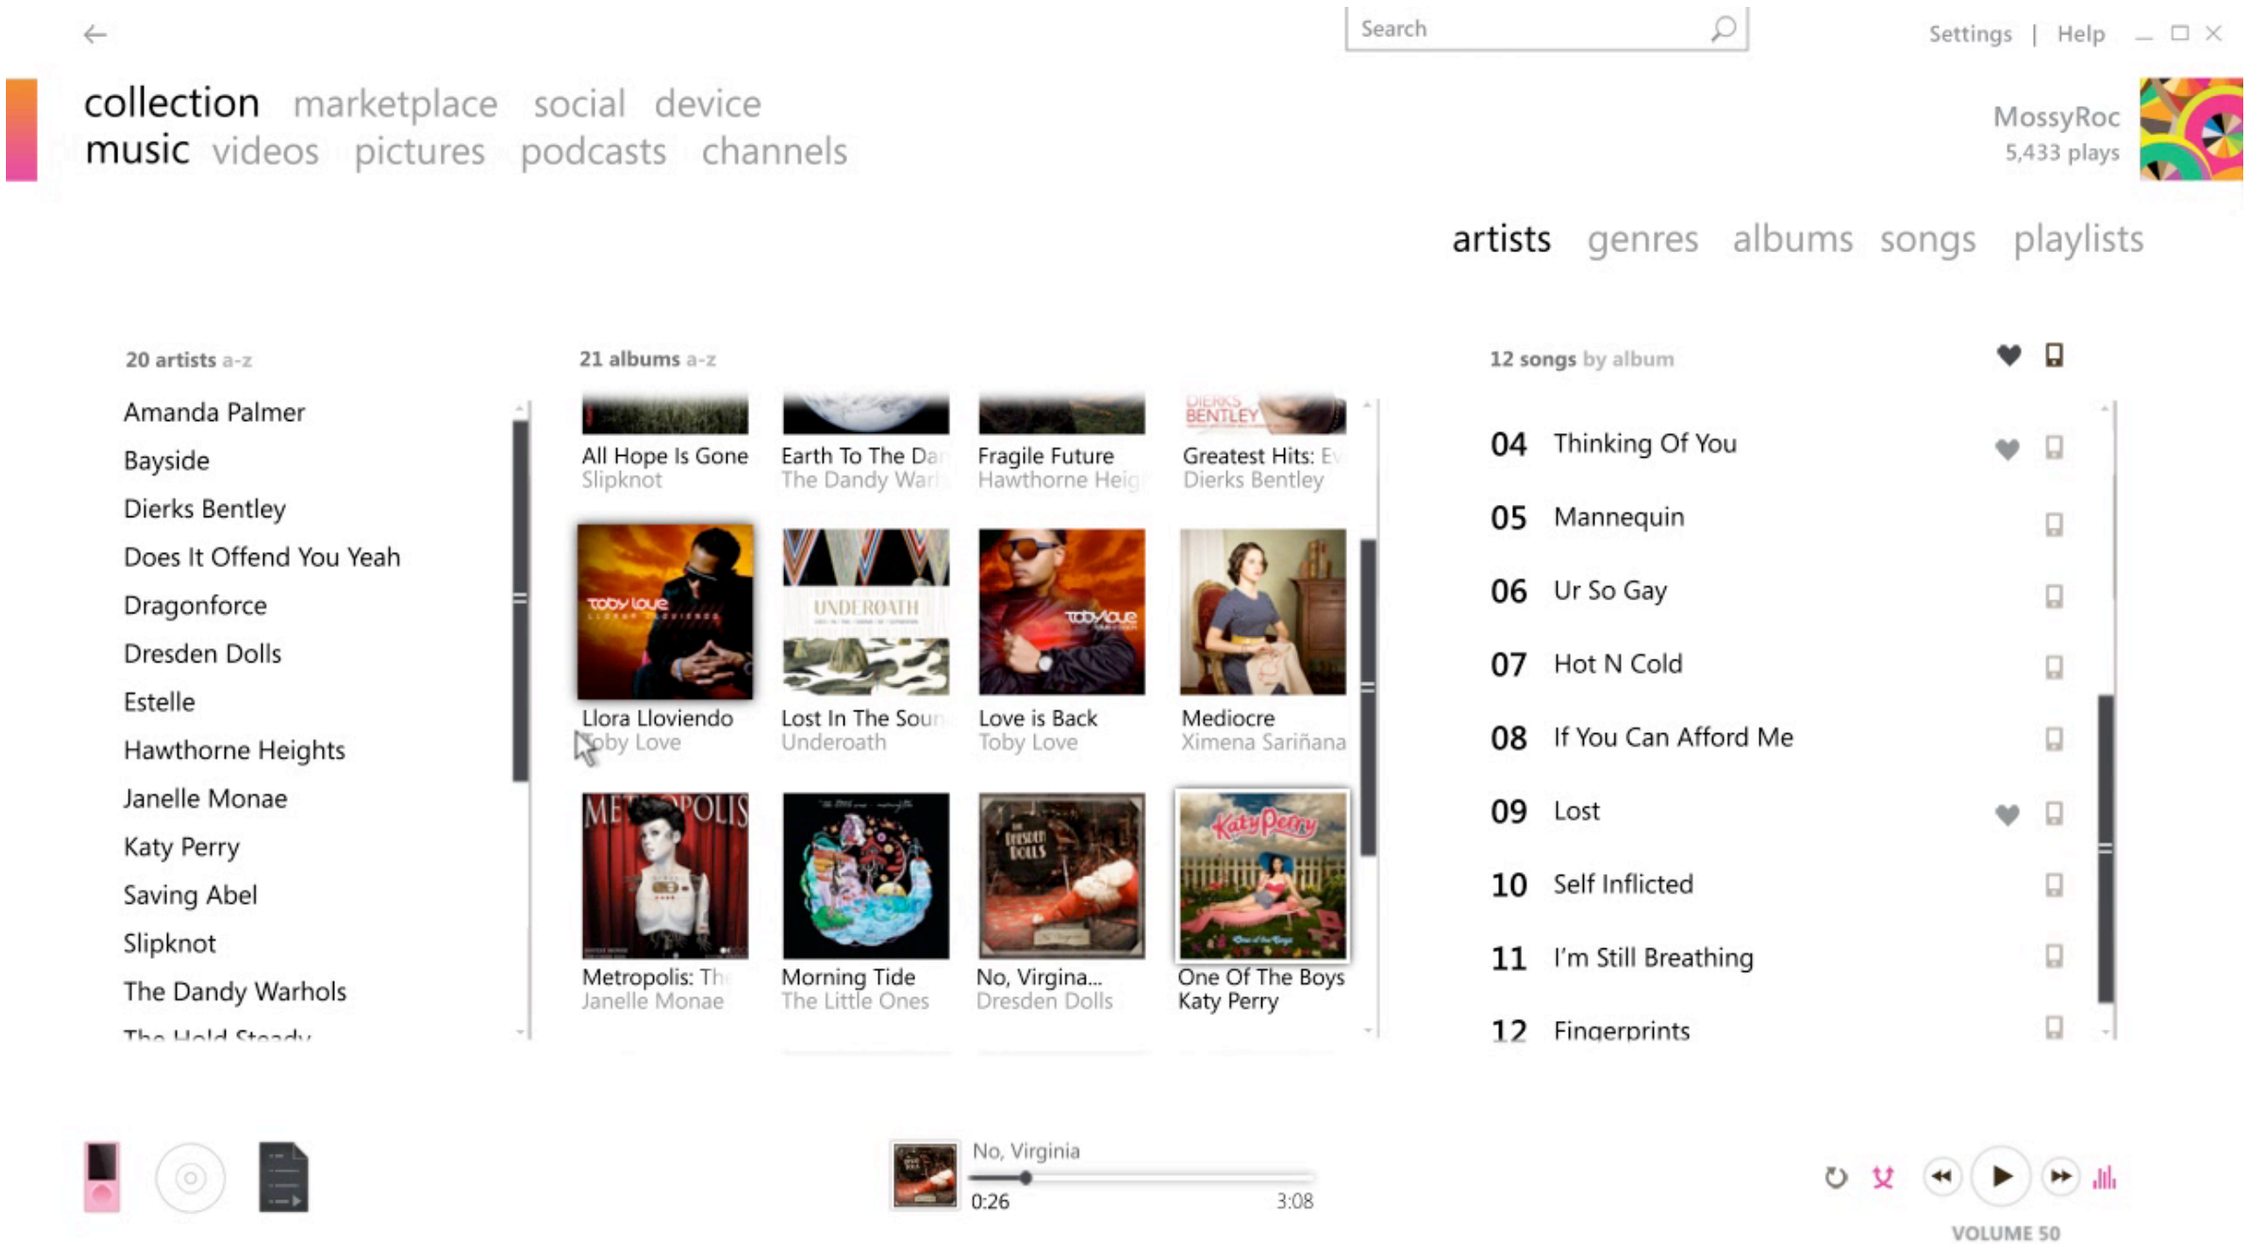

social.zune.net

Musicream

http://staff.aist.go.jp/m.goto/Musicream/ Goto and Masataka Goto.

Musicream

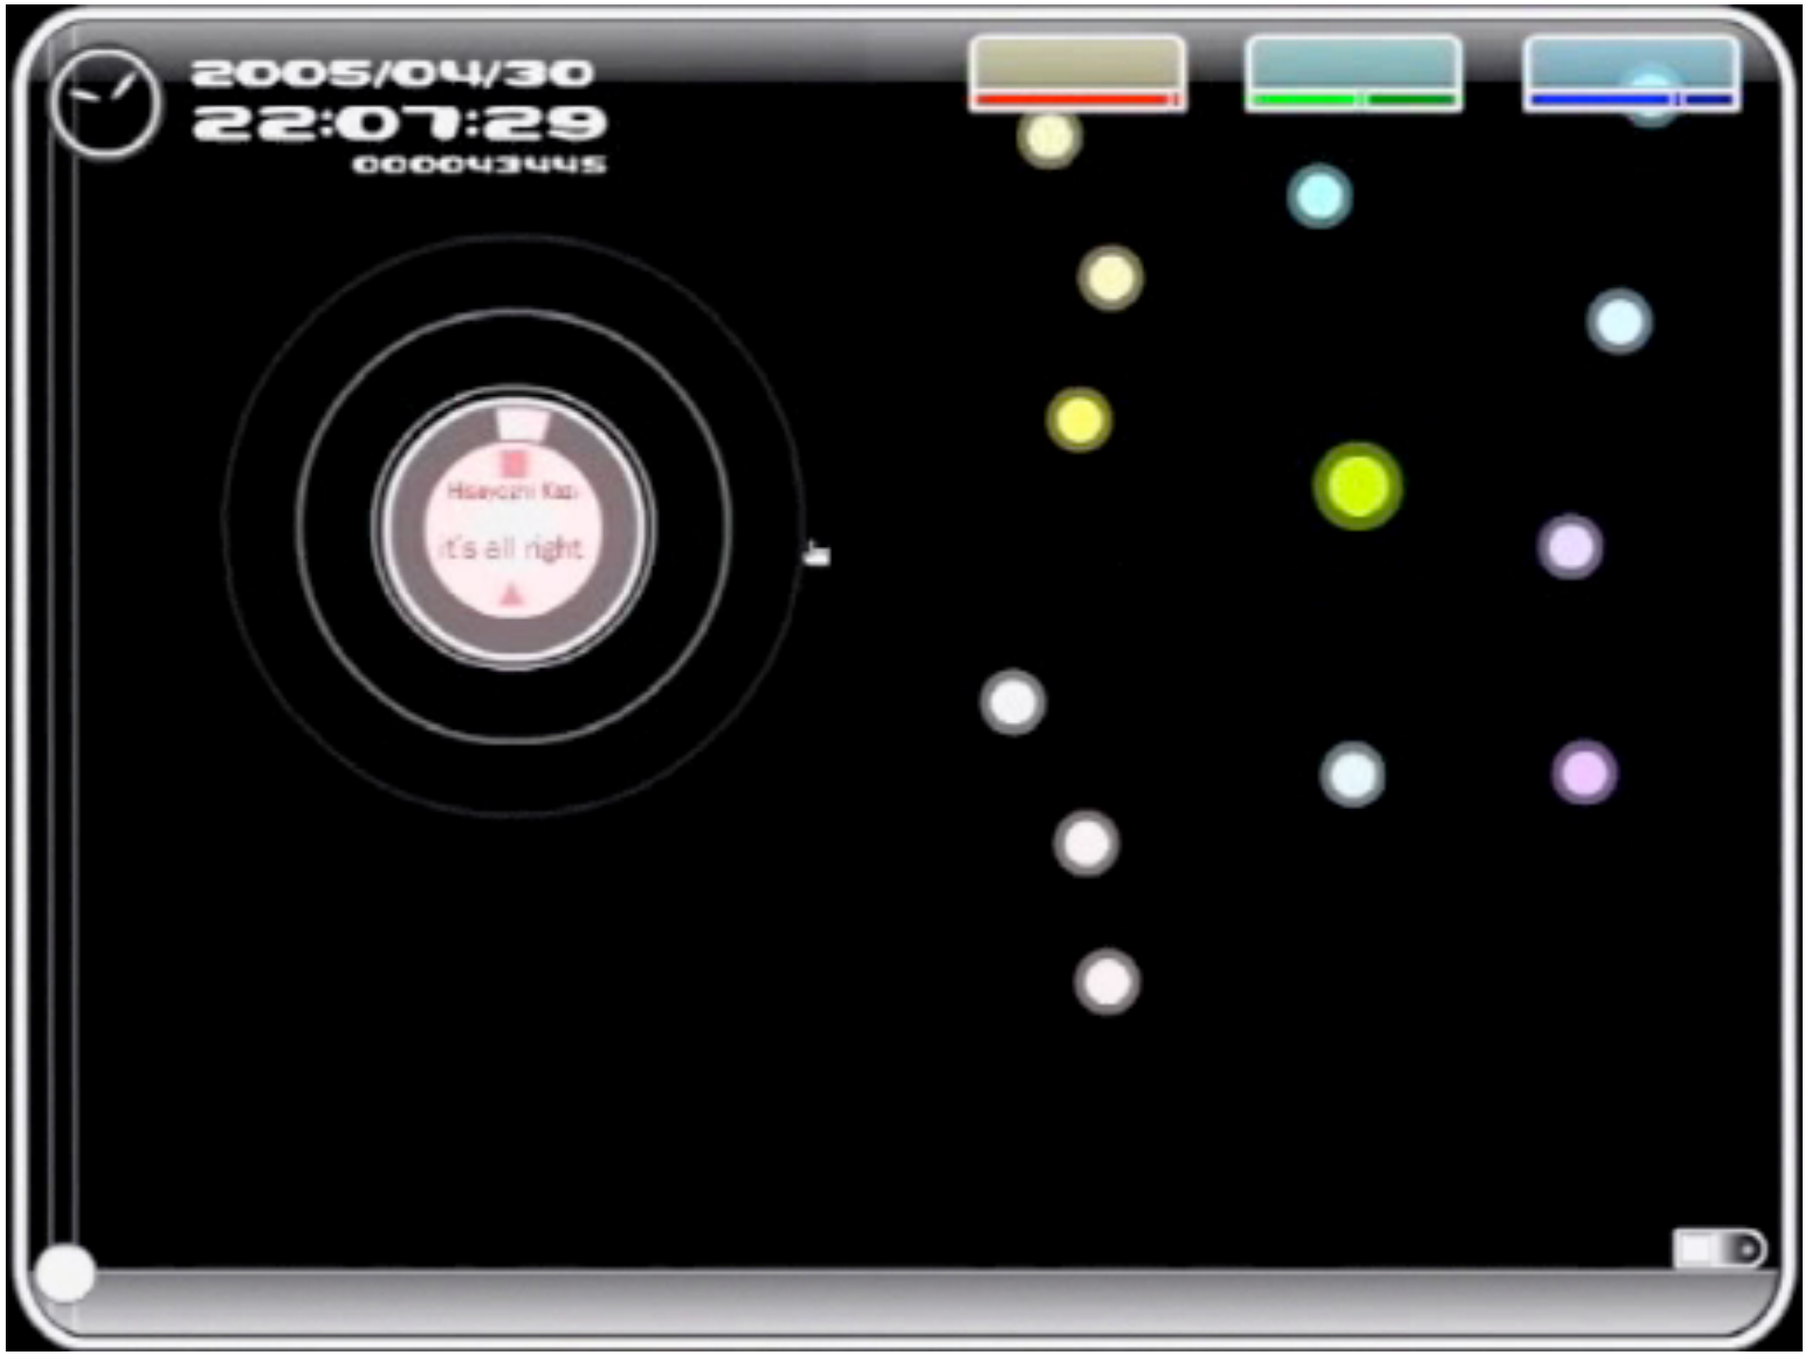

http://staff.aist.go.jp/m.goto/Musicream/ Goto and Masataka Goto.

MusicRainbow

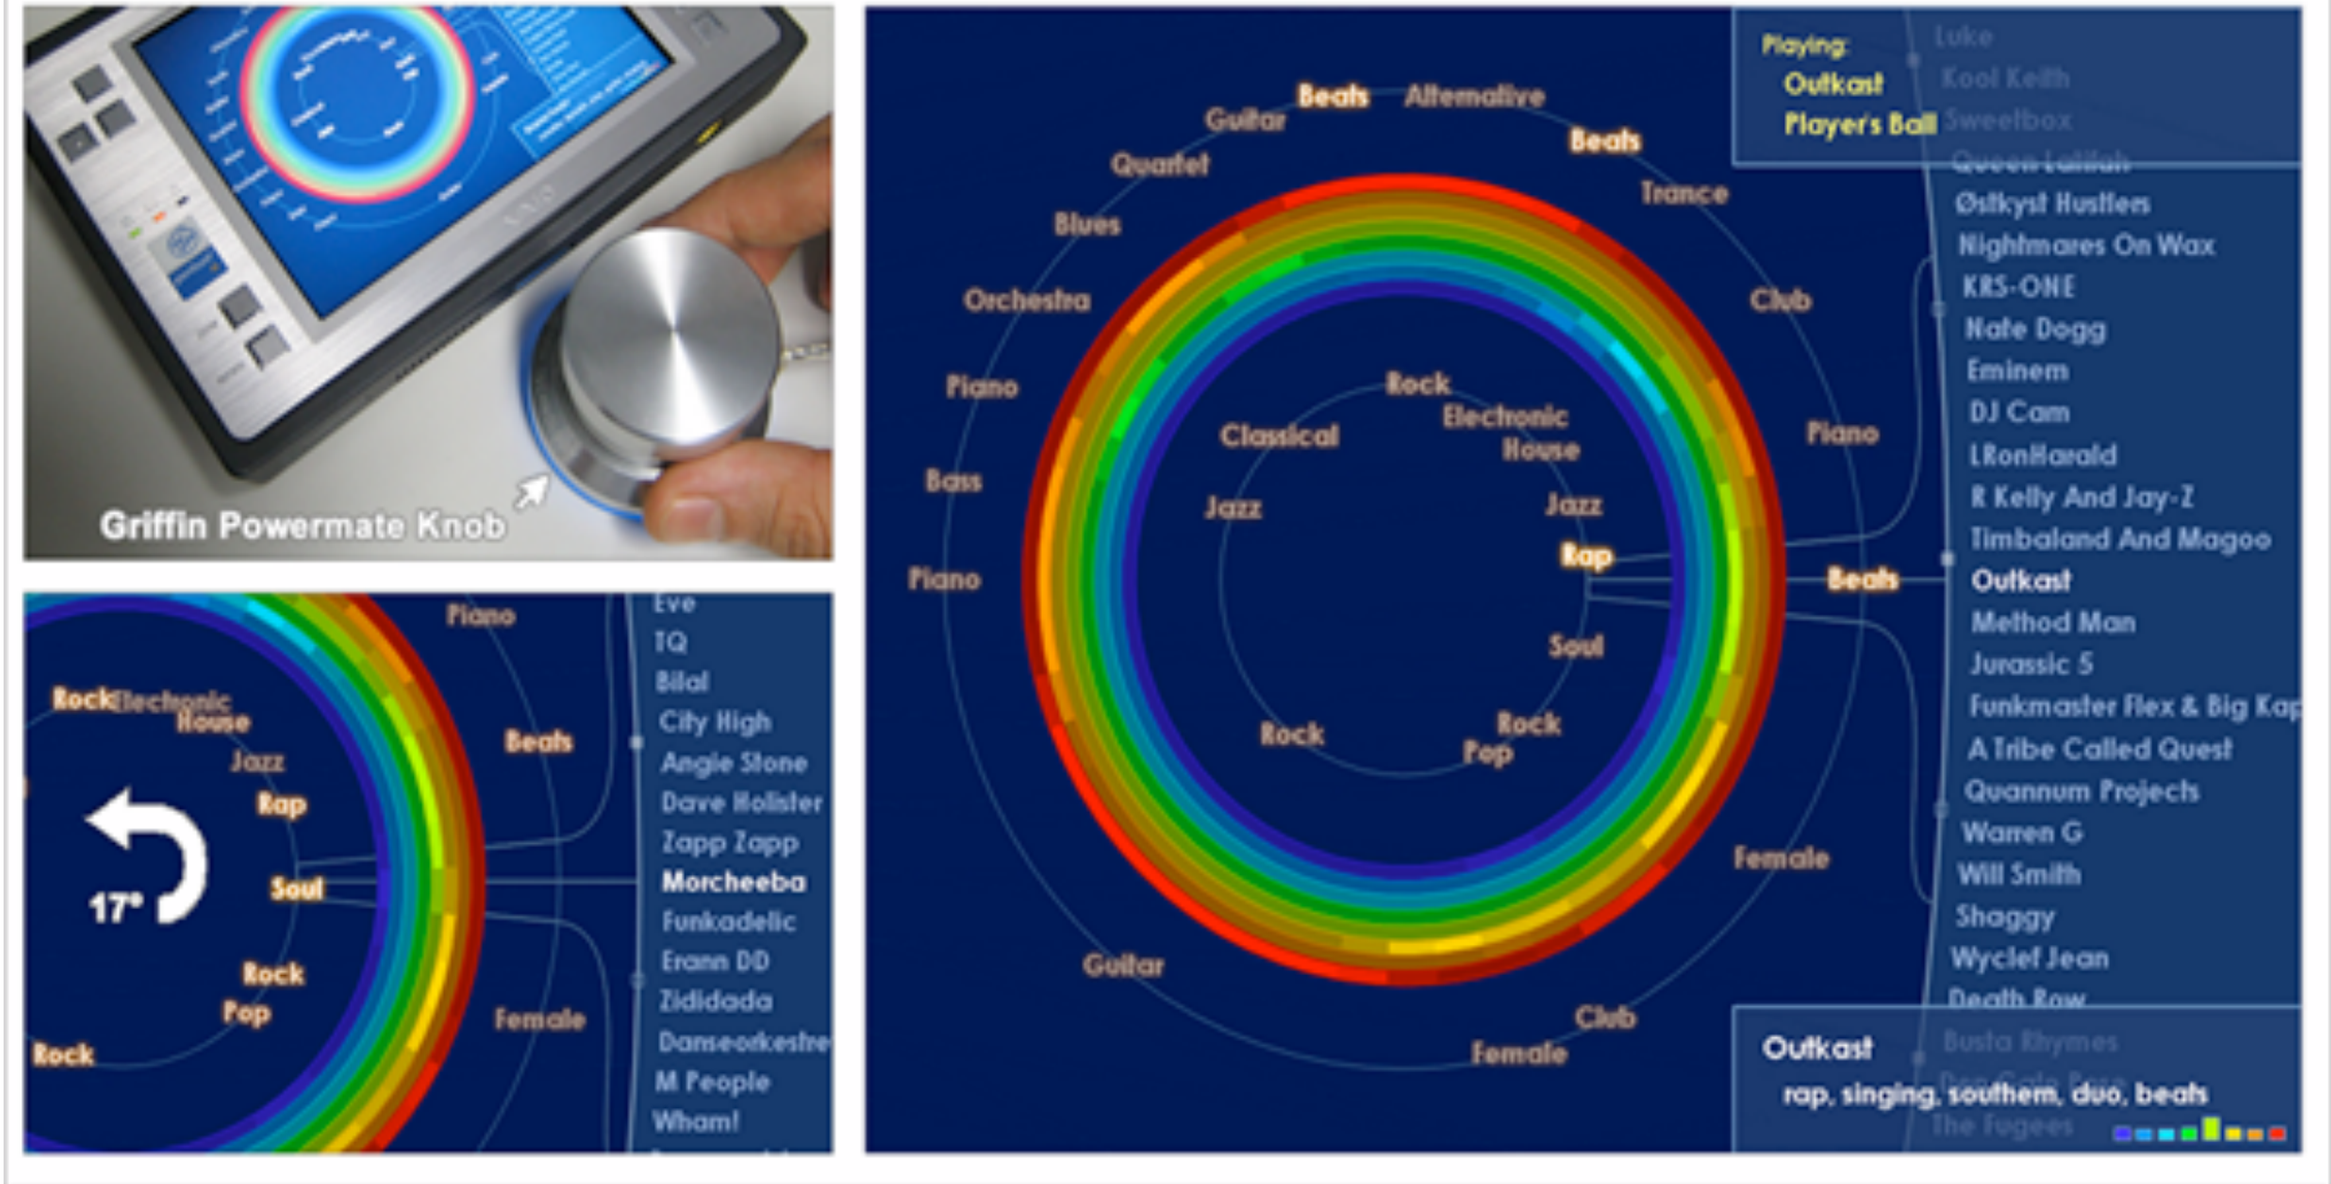

Elias Pampalk & Masataka Goto, "MusicRainbow: A New User Interface to Discover Artists Using Audio-based Similarity and Web-based Labeling

MusicSun

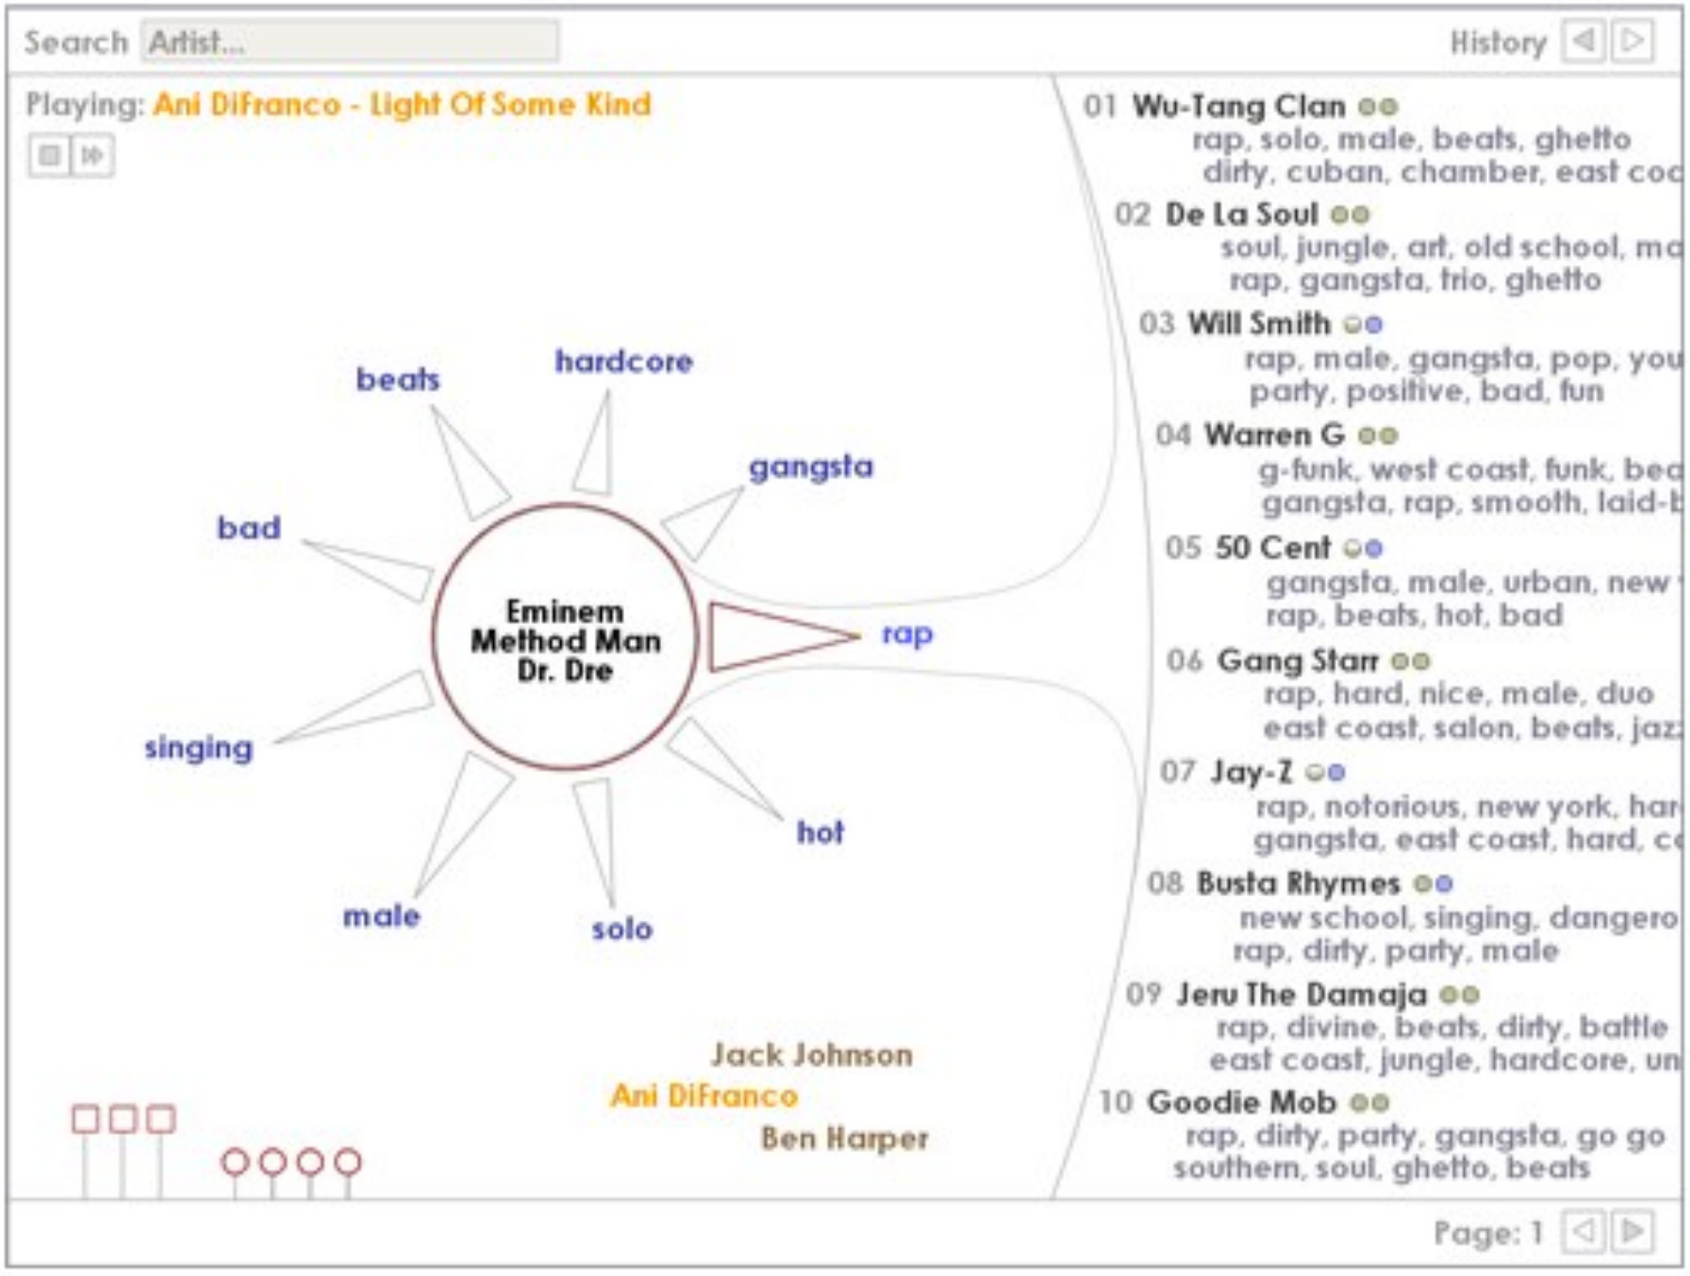

Elias Pampalk & Masataka Goto,
"MusicSun: A New Approach to Artist Recommendation"

EN Music Explorer

EN Music Explorer

Here be Dragons

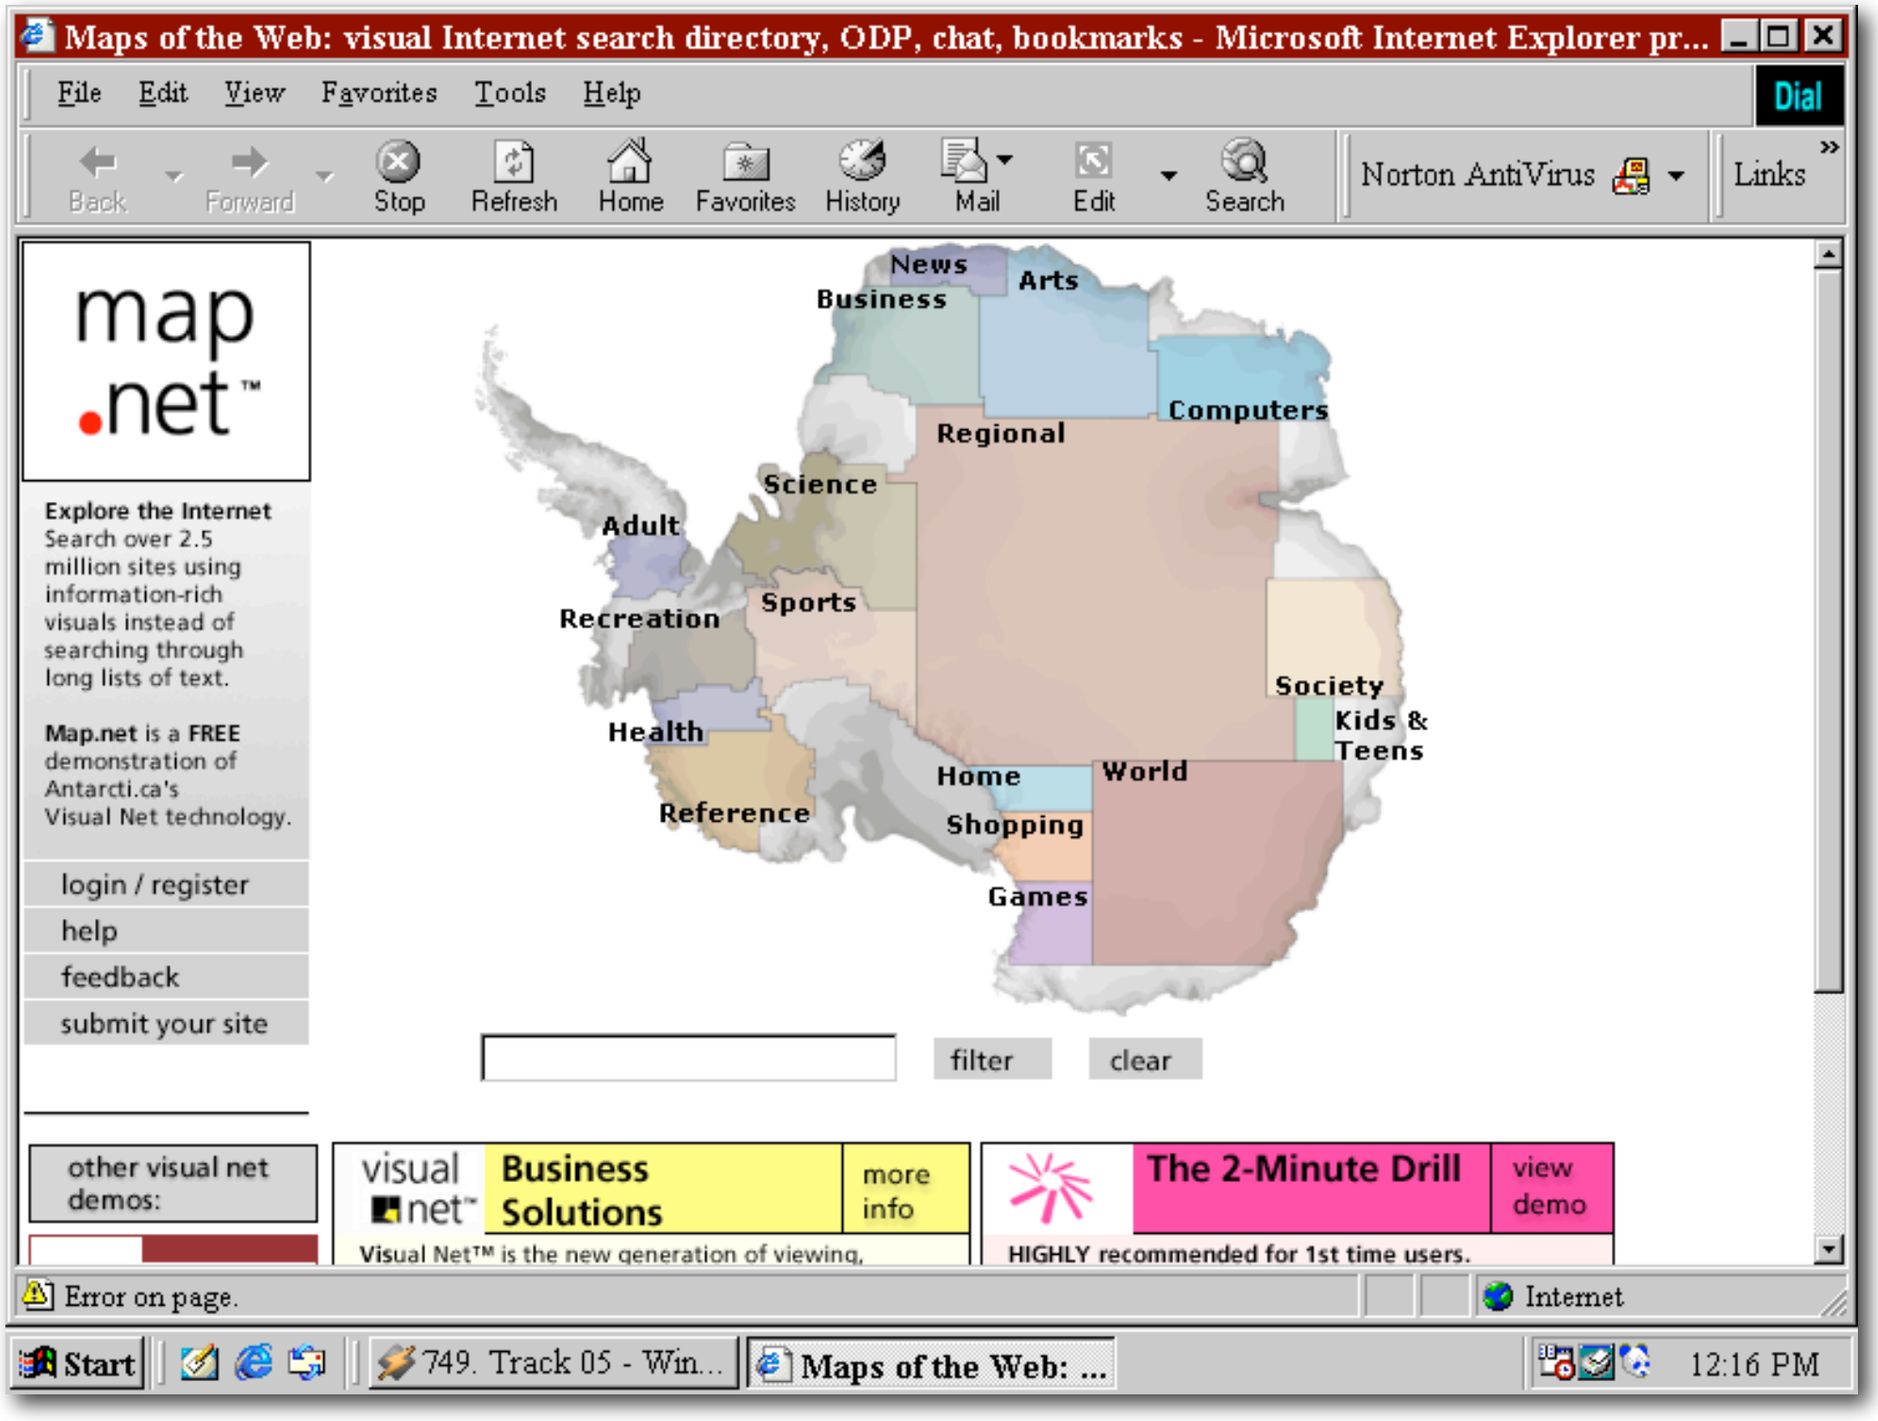

Here be Dragons

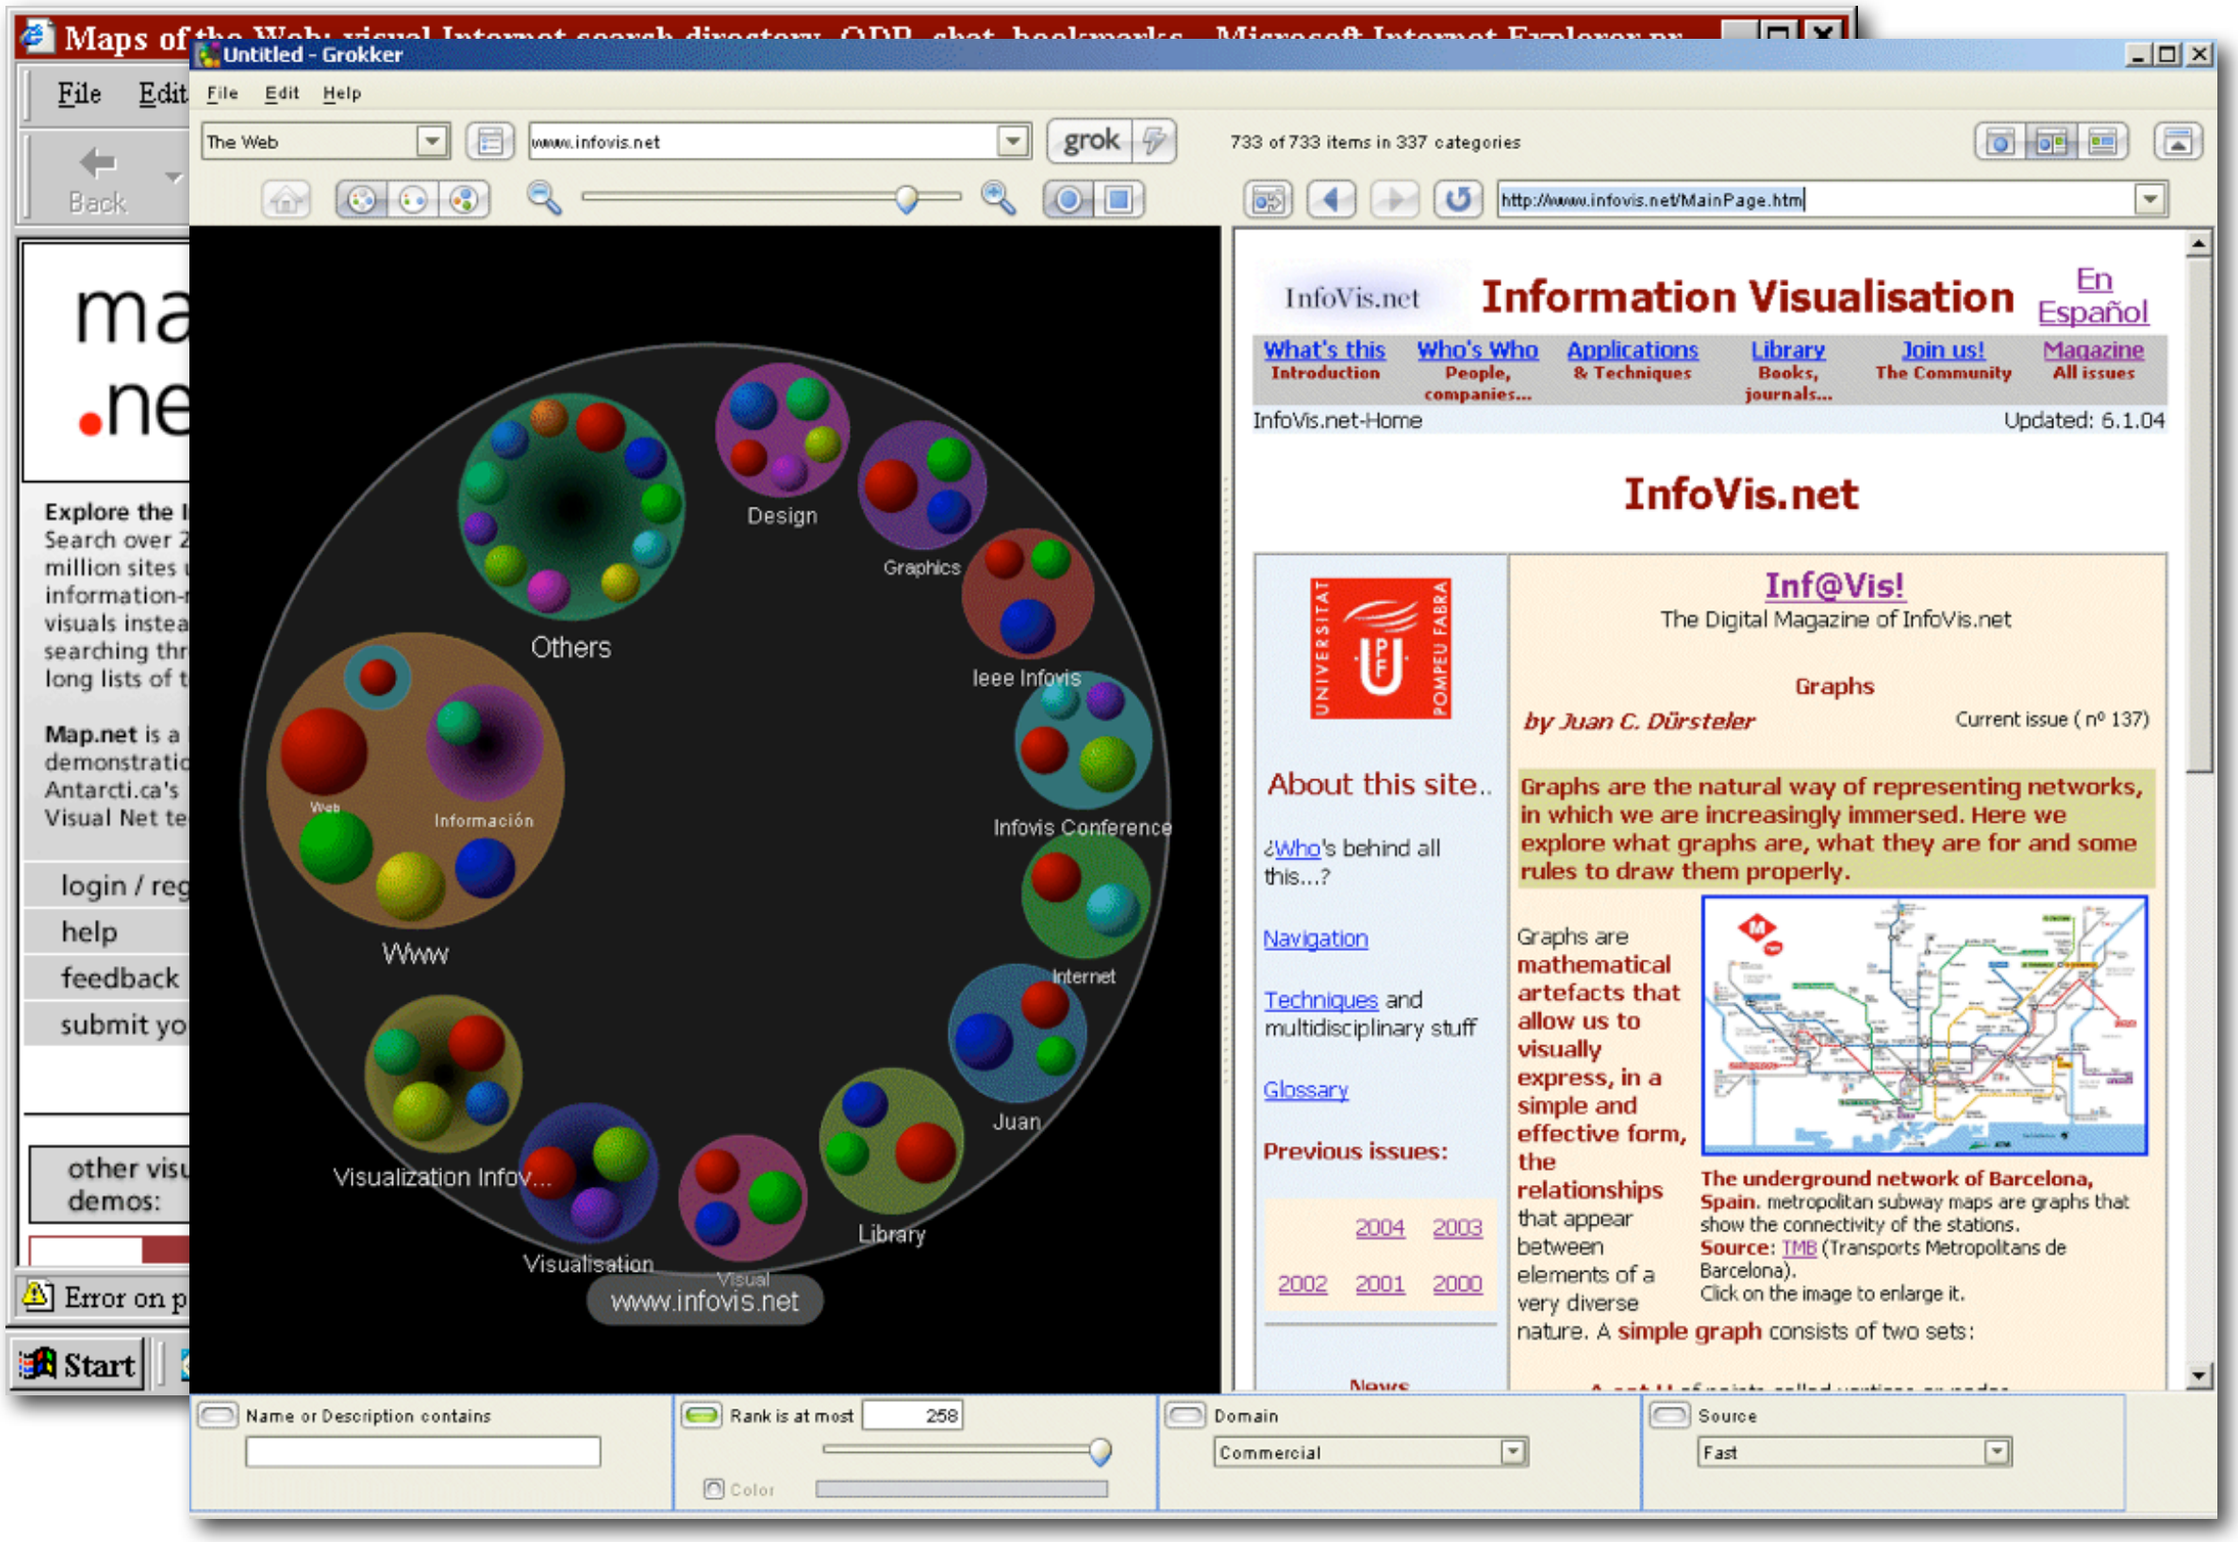

Here be Dragons

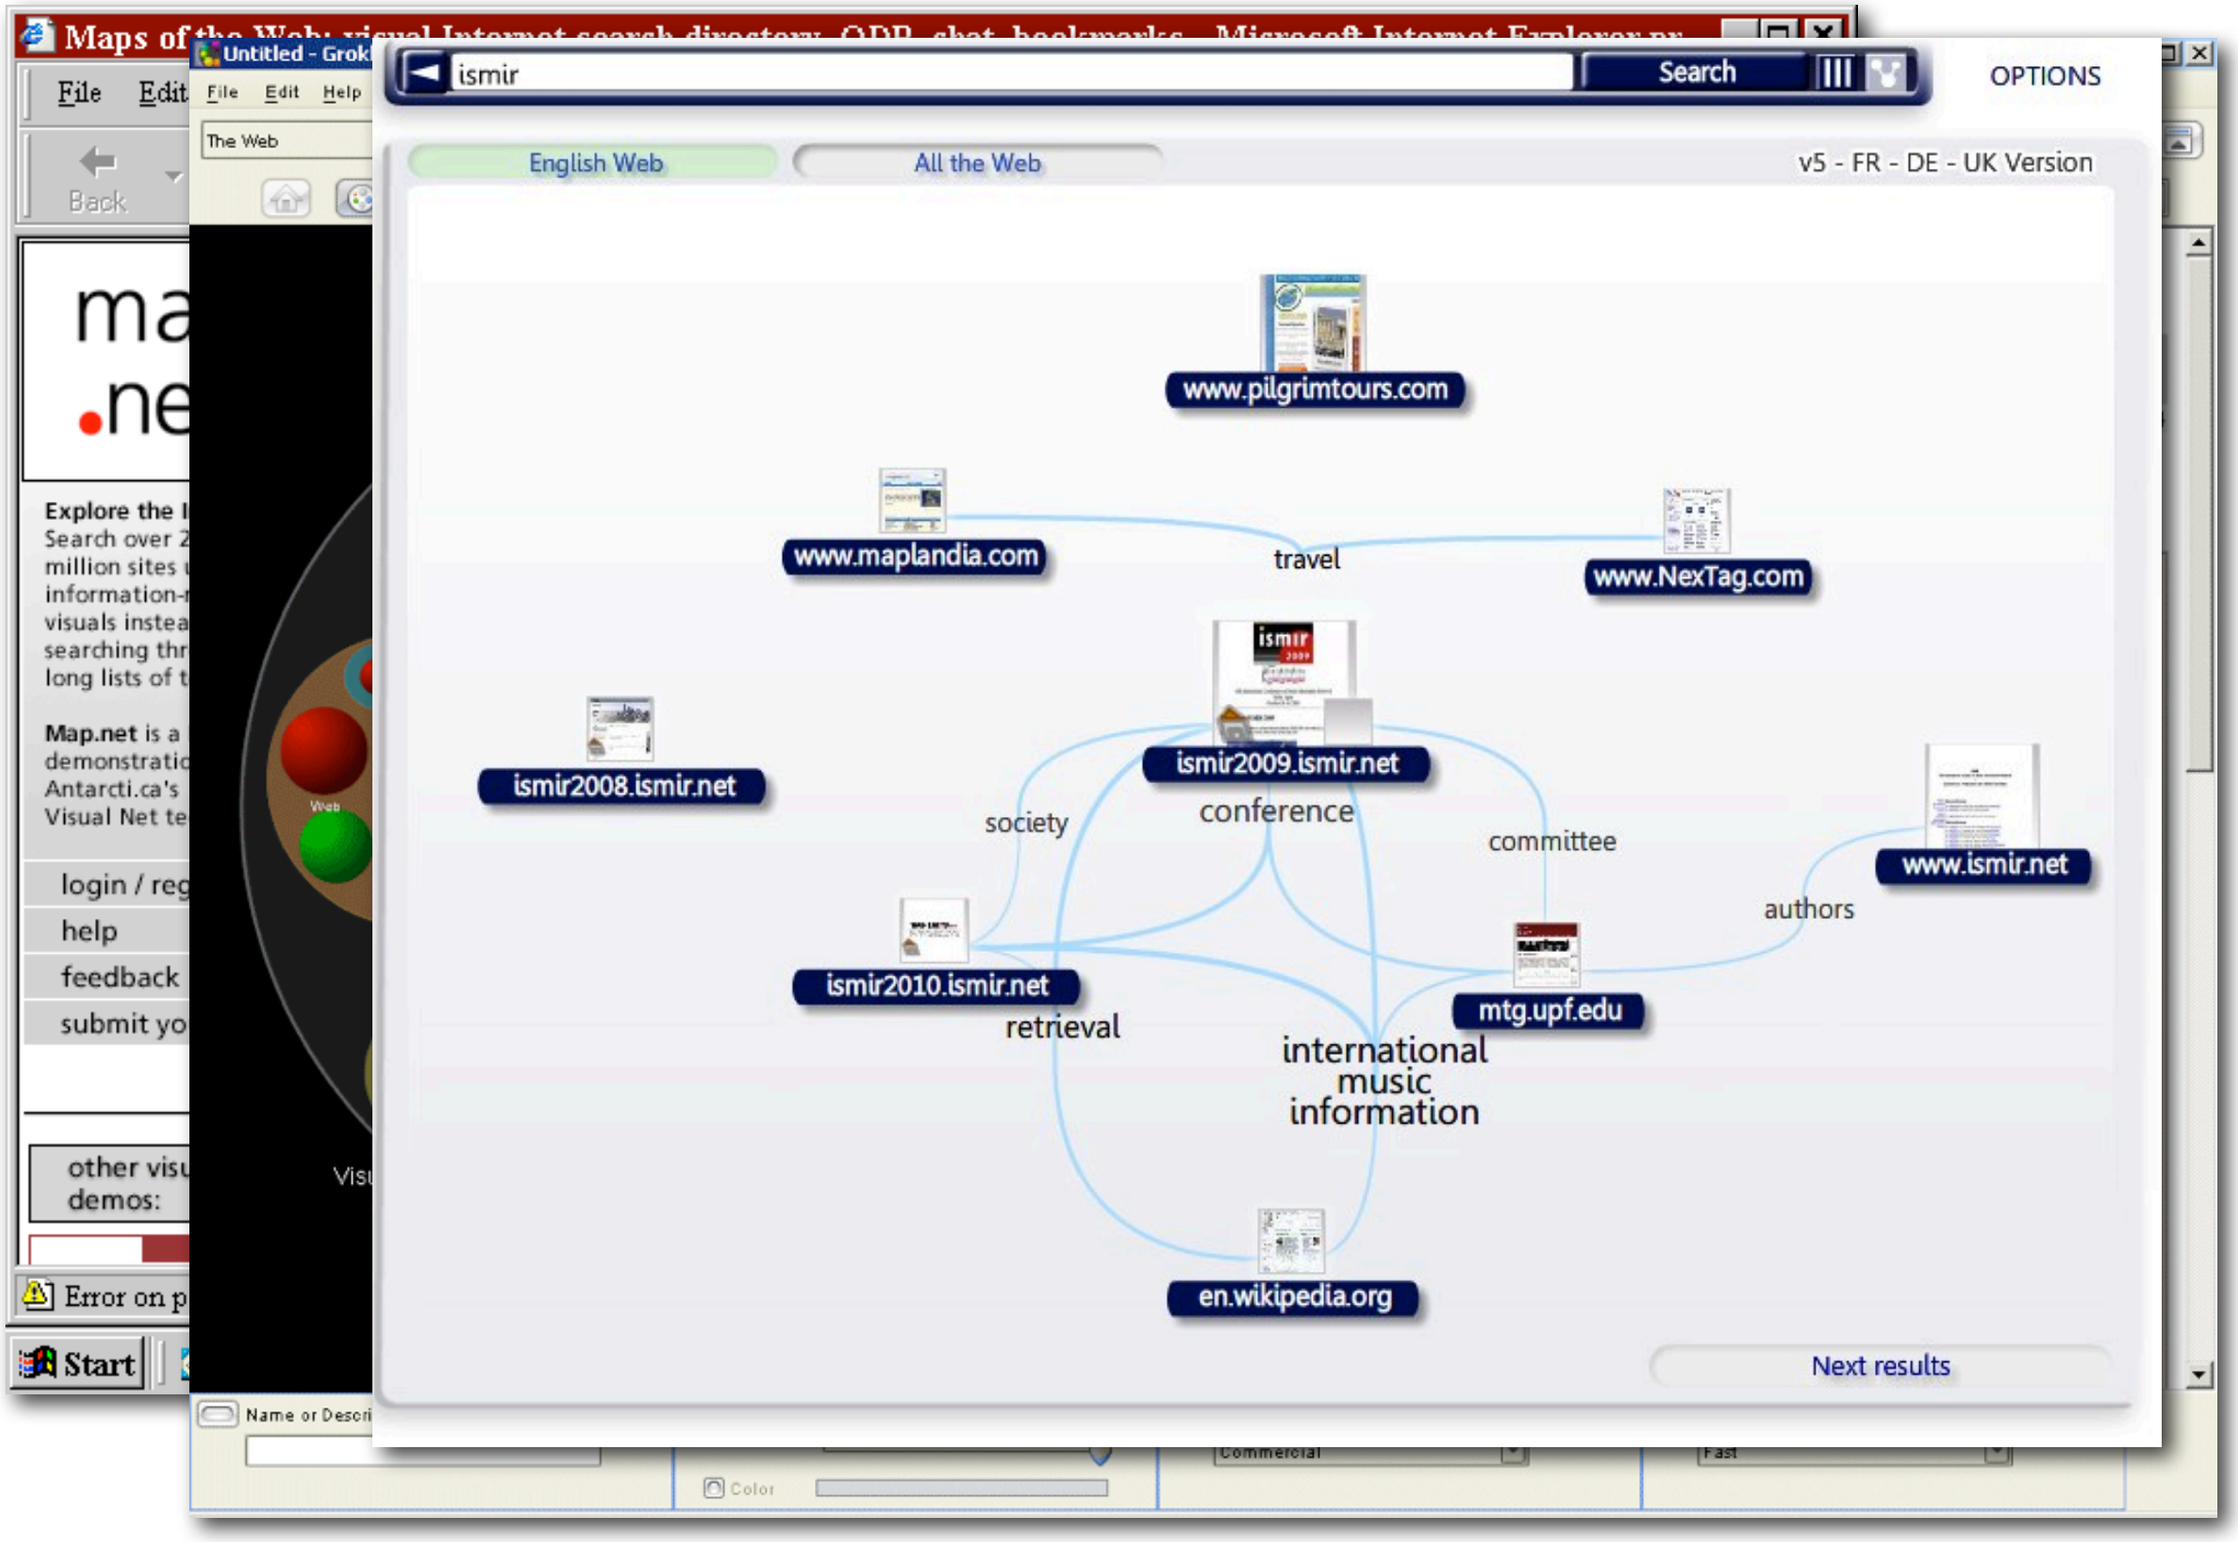

Visualization Proponents (The Yale Club)

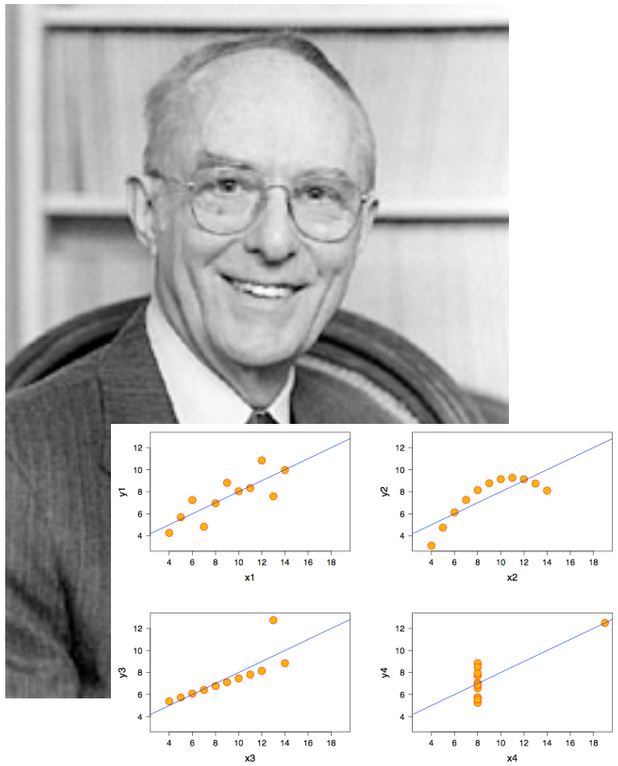

F.J. Anscombe
(Anscombe's Quartet)

0 | 0000111111222338
2 | 07
4 | 5
6 | 8
8 | 4
10 | 5
12 |
14 |
16 | 0

John Tukey
(stem/leaf plots)

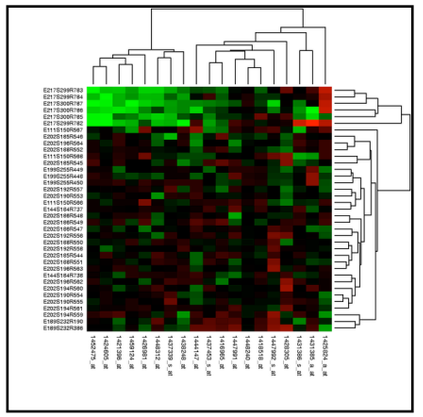

Leland Wilkinsen
(Cluster heat maps)

Edward Tufte
(Sparklines)

http://en.wikipedia.org/wiki/F.J._Anscombe
http://en.wikipedia.org/wiki/John_Tukey
http://en.wikipedia.org/wiki/Leland_Wilkinson
http://en.wikipedia.org/wiki/Edward_Tufte

Interactive Visualization Proponents

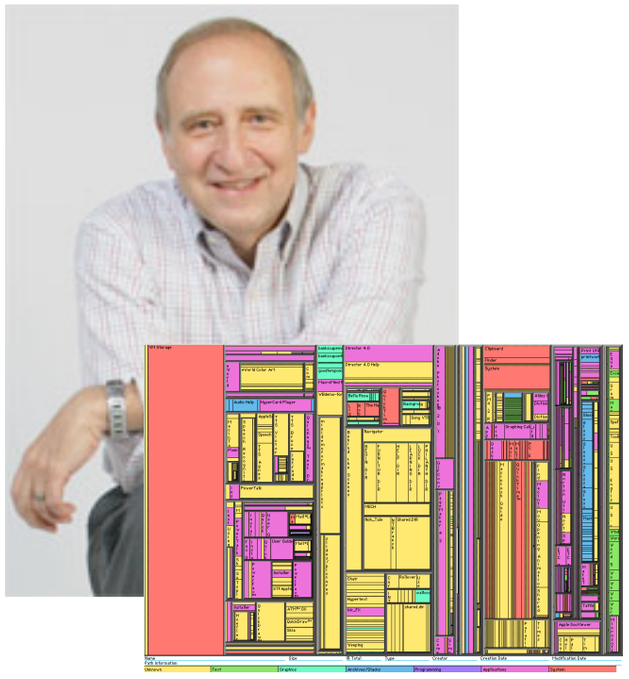

Ben Shneiderman (Treemaps, NSD diagrams)

Steve Deunes (& entire NYT Graphics Department)

Ben Fry & Casey Reas (Processing)

Resources

Resources

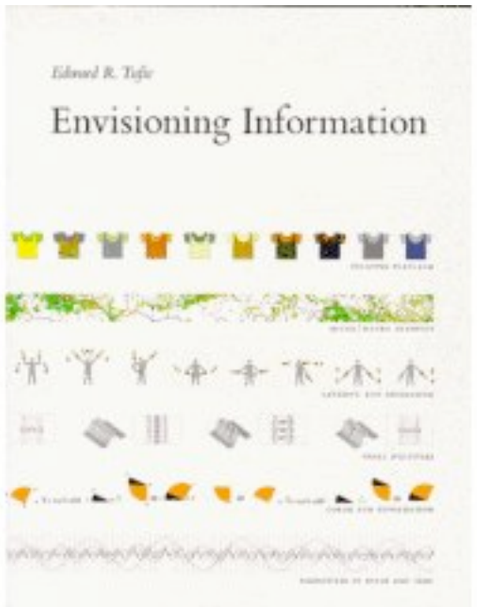

Resources

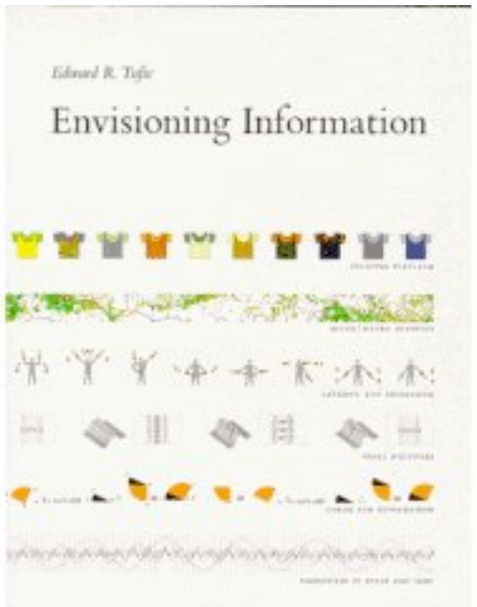
Edward R. Tufte
Envisioning Information

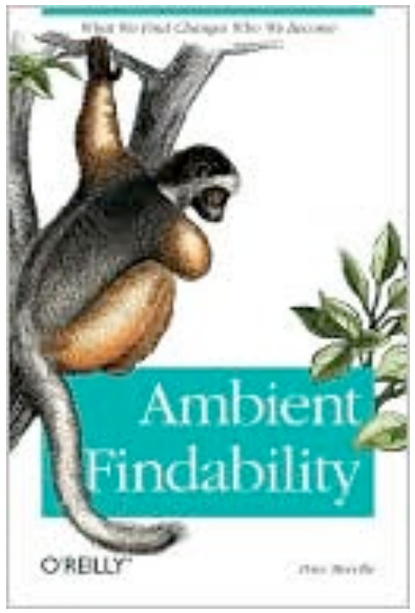
Ambient Findability
O'REILLY

Resources

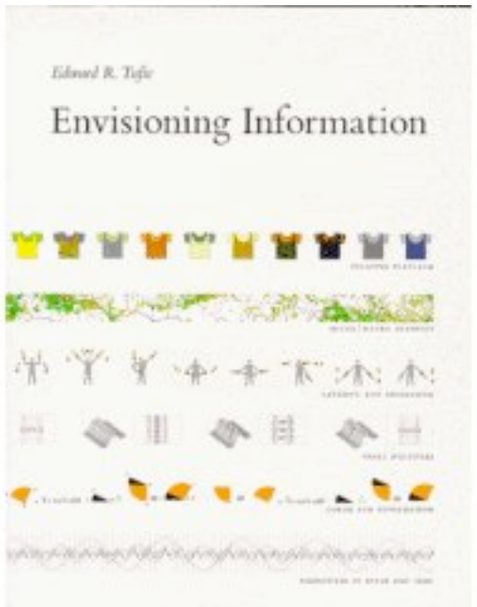
Edward R. Tufte
Envisioning Information

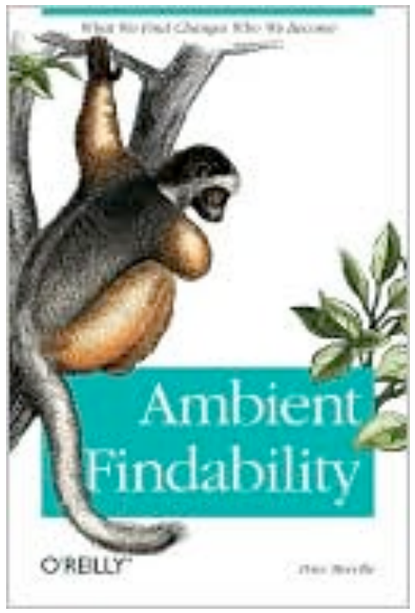
Ambient Findability
O'REILLY

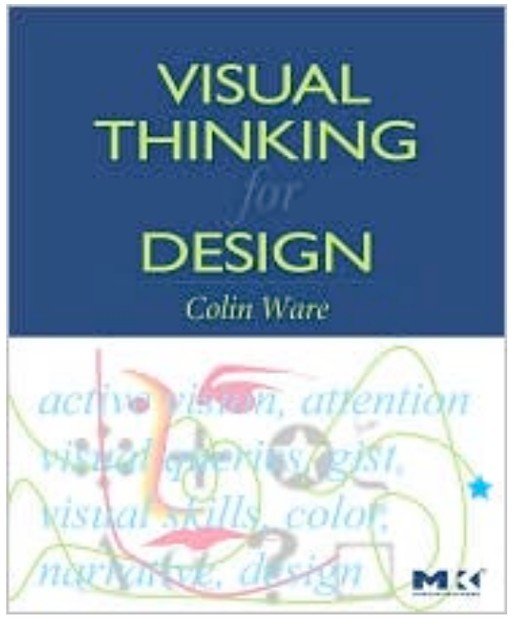
VISUAL
THINKING
for
DESIGN
Colin Ware

Resources

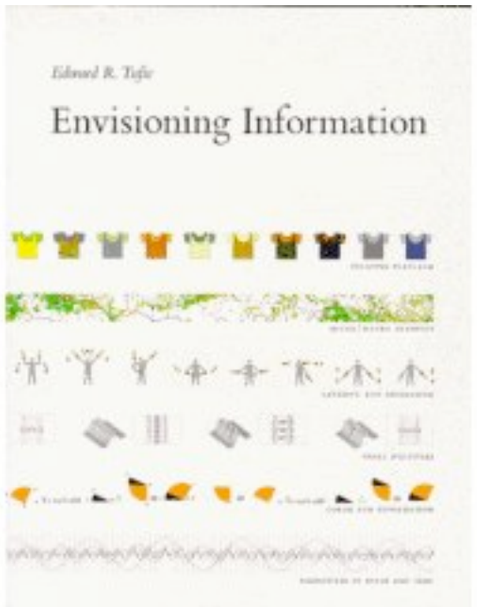
Edward R. Tufte
Envisioning Information

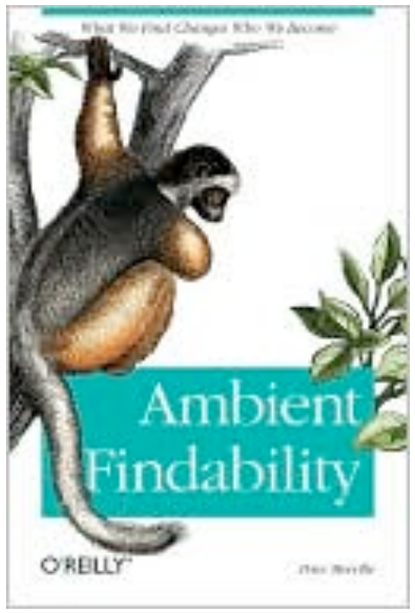
Ambient Findability
O'REILLY

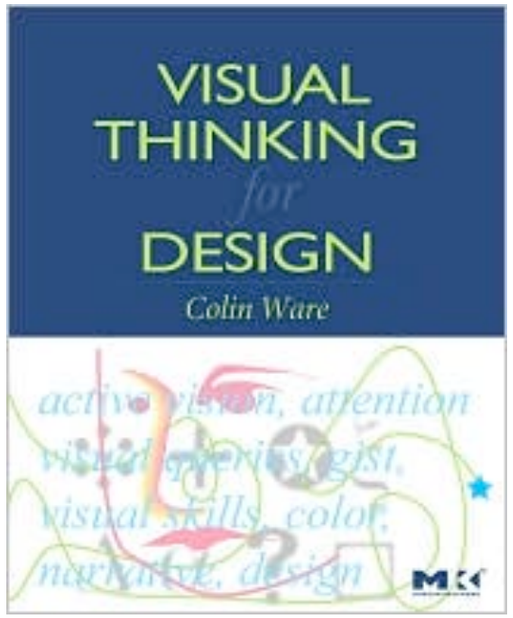
VISUAL
THINKING
for
DESIGN
Colin Ware

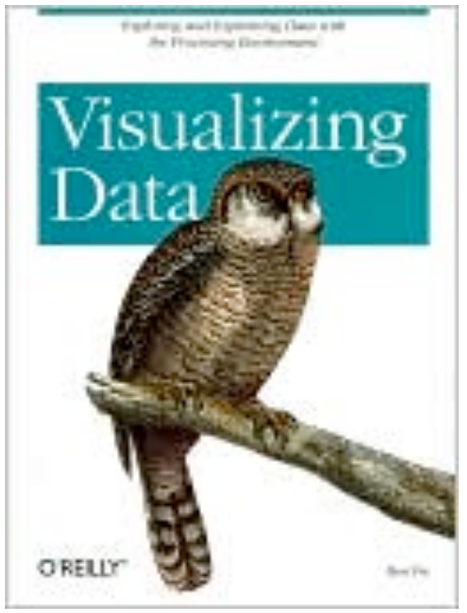
Visualizing Data
O'REILLY

Resources

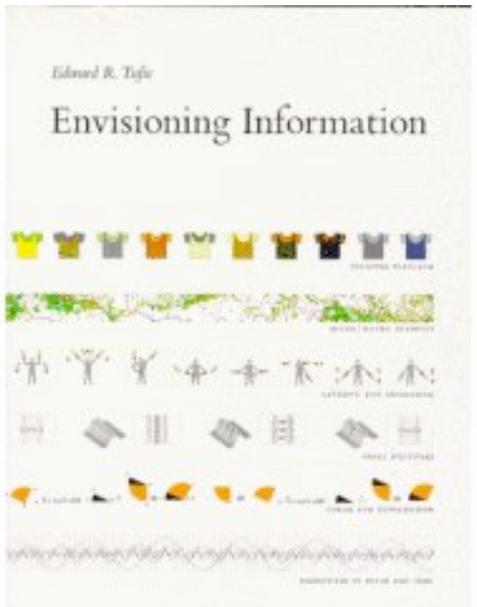
Edward R. Tufte
Envisioning Information

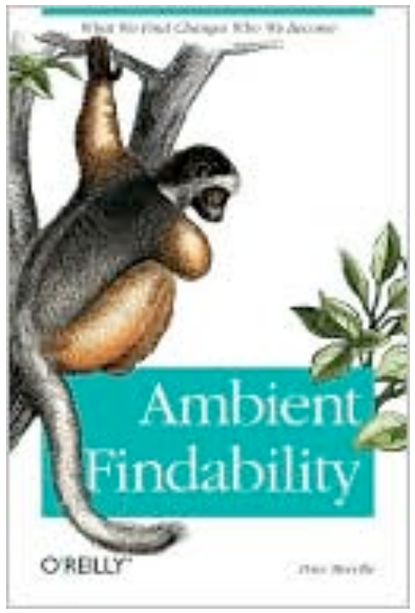
Ambient Findability
O'REILLY

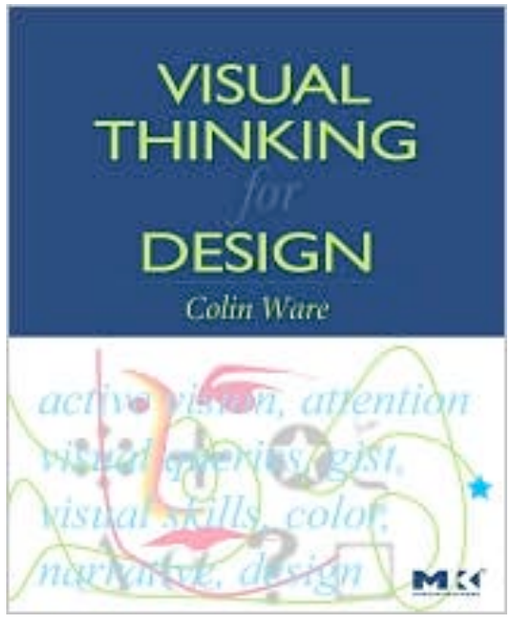
VISUAL
THINKING
for
DESIGN
Colin Ware

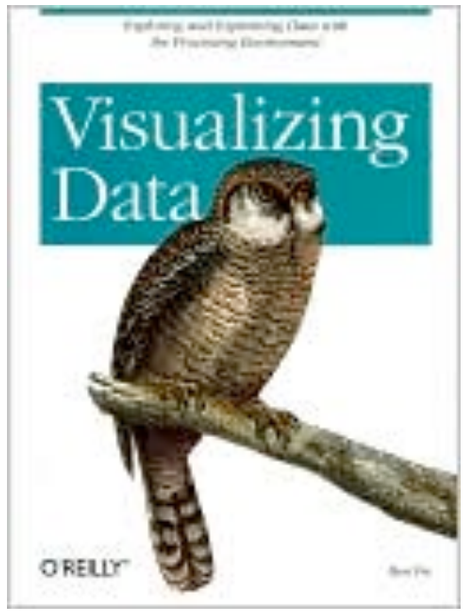
Visualizing Data
O'REILLY

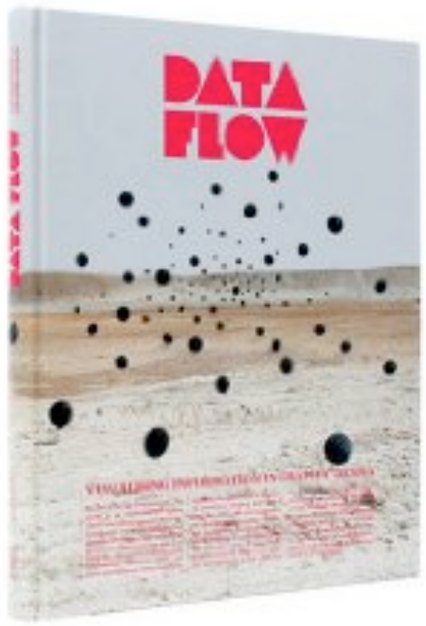
DATA FLOW

Resources

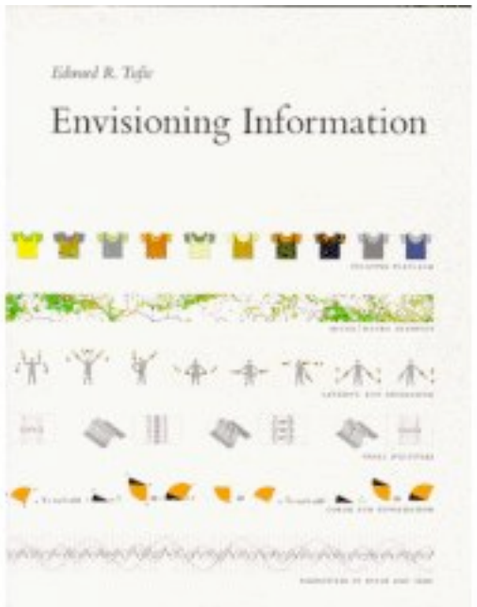
Edward R. Tufte
Envisioning Information

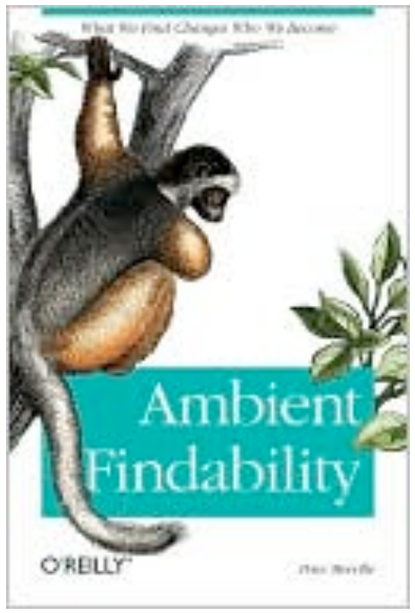
Ambient Findability
O'REILLY

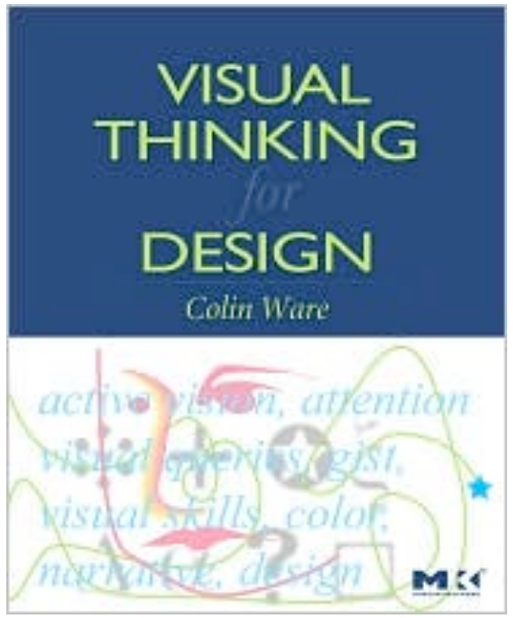
VISUAL THINKING for DESIGN
Colin Ware

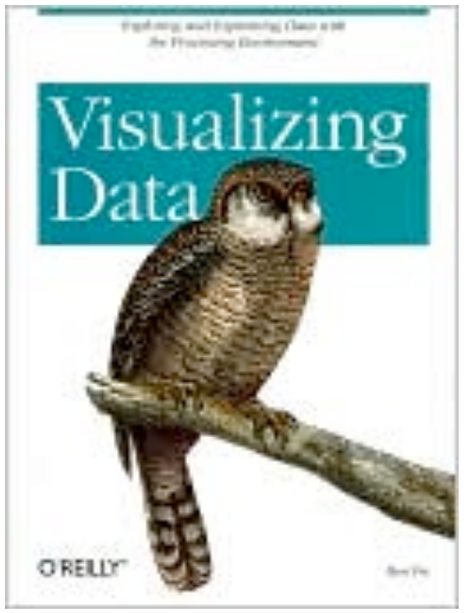
Visualizing Data
O'REILLY

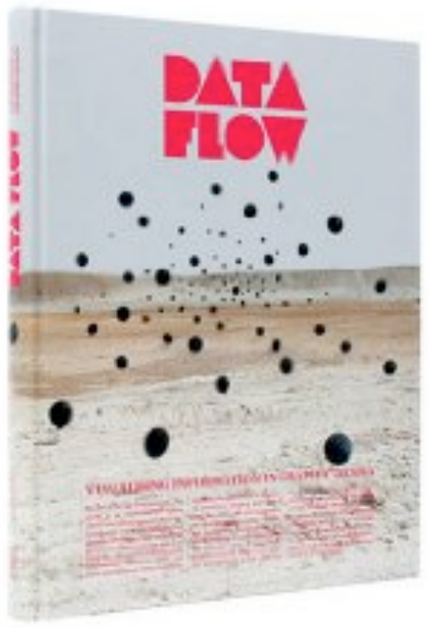
DATA FLOW

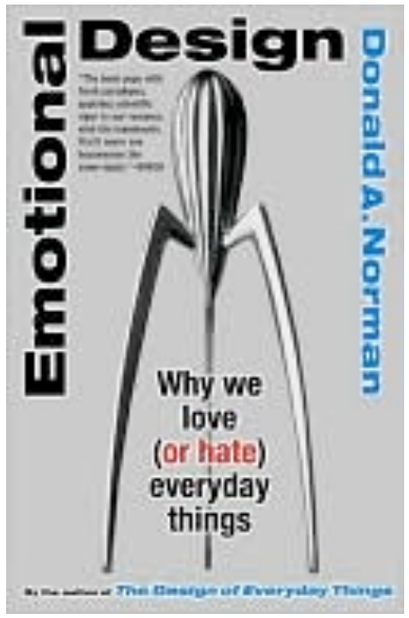
Emotional Design
Donald A. Norman
Why we love (or hate) everyday things

Resources

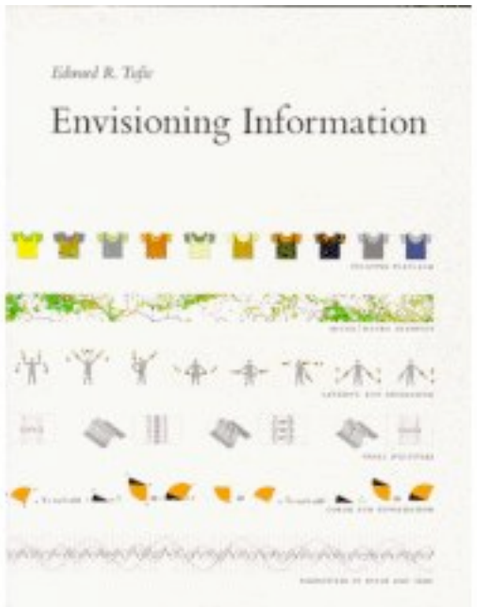
Edward R. Tufte
Envisioning Information

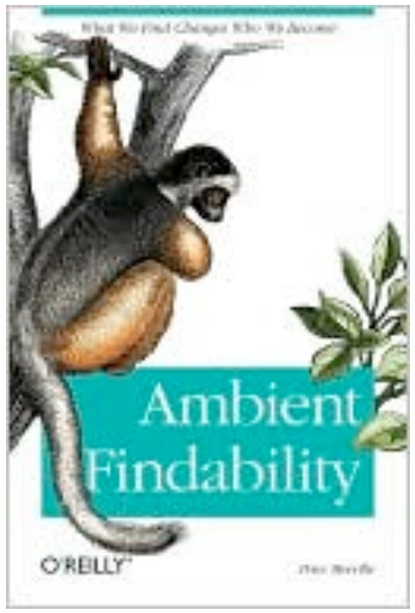
Ambient Findability
O'REILLY

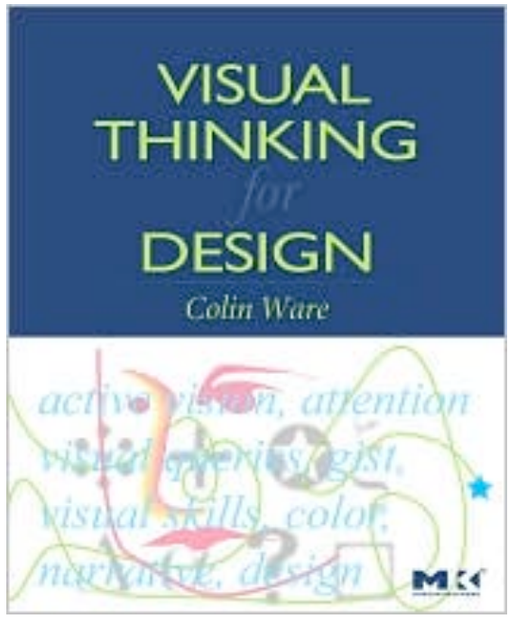
VISUAL THINKING for DESIGN
Colin Ware

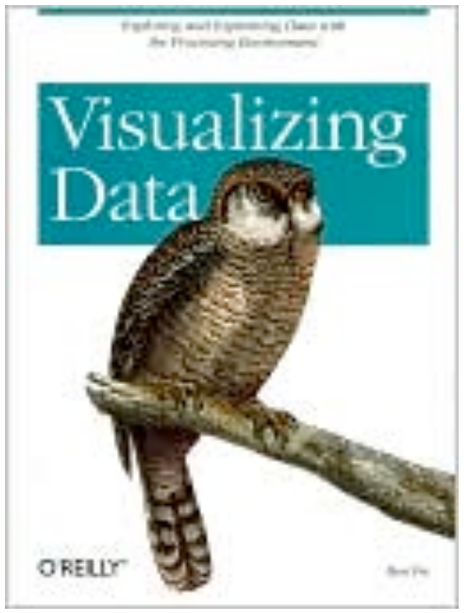
Visualizing Data
O'REILLY

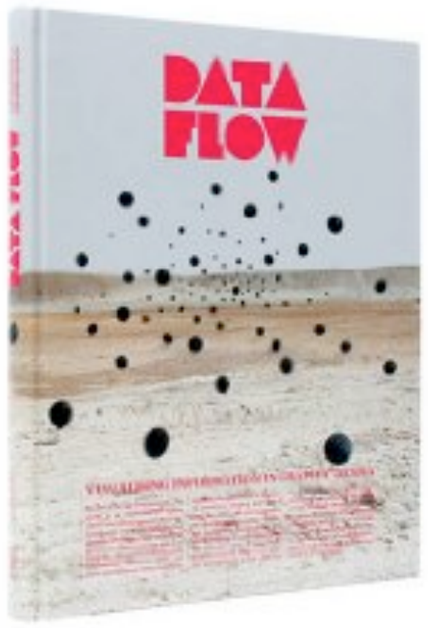
DATA FLOW

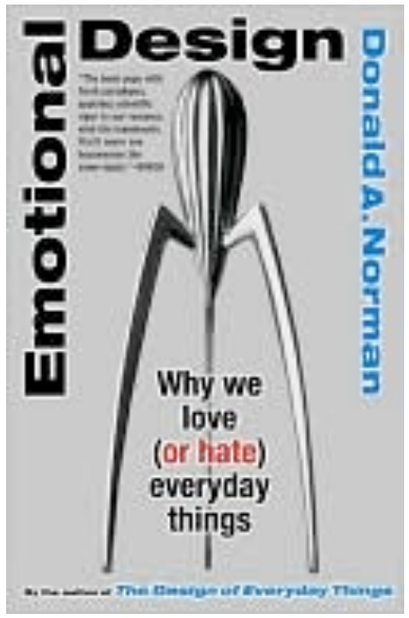
Emotional Design
Donald A. Norman
Why we love (or hate) everyday things

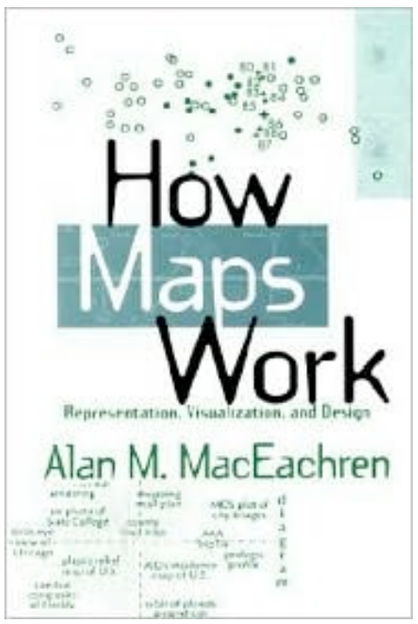
How Maps Work
Representation, Visualization, and Design
Alan M. MacEachren

Resources

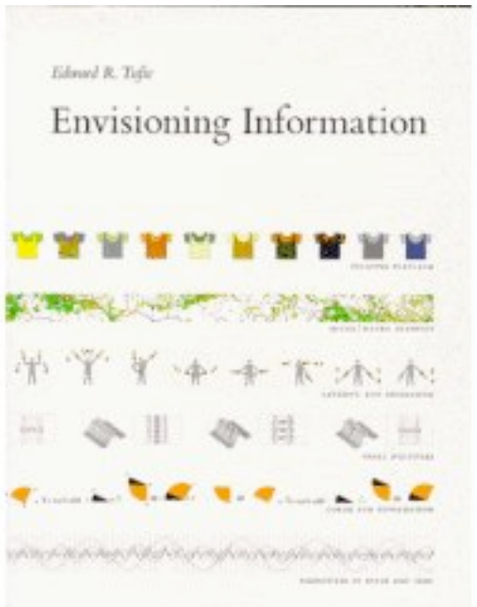
Edward R. Tufte
Envisioning Information

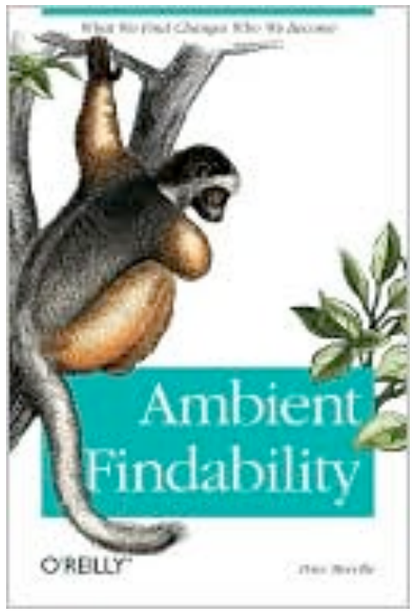
Ambient Findability
O'REILLY

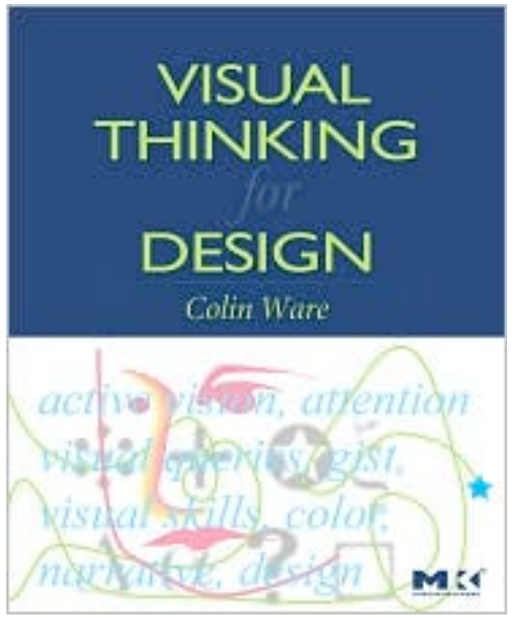
VISUAL
THINKING
for
DESIGN
Colin Ware

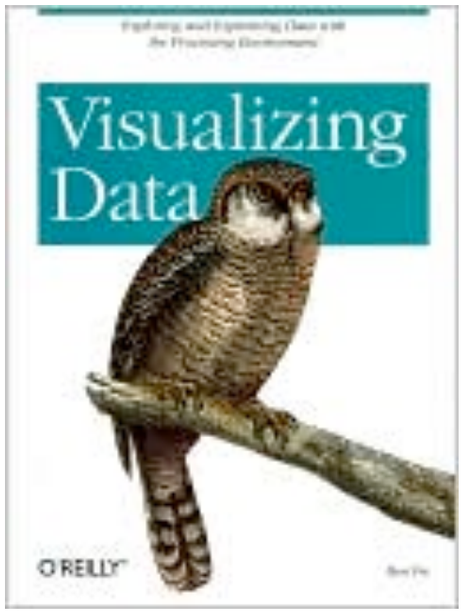
Visualizing Data
O'REILLY

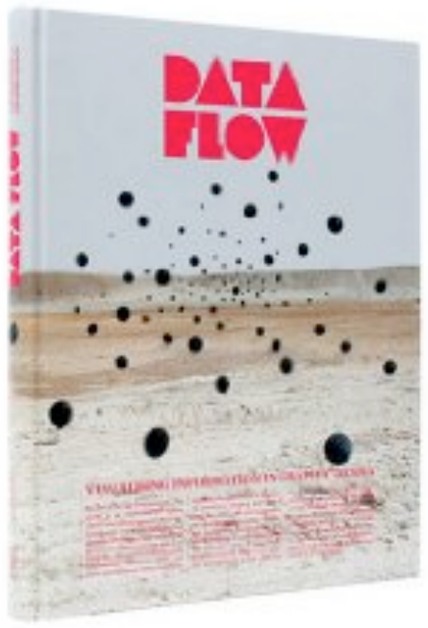
DATA FLOW

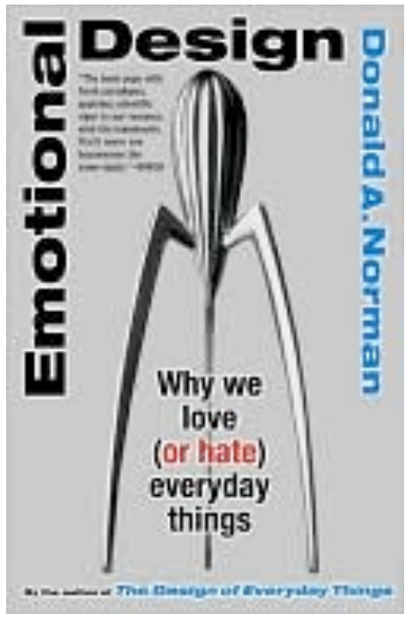
Emotional Design
Donald A. Norman
Why we love (or hate) everyday things

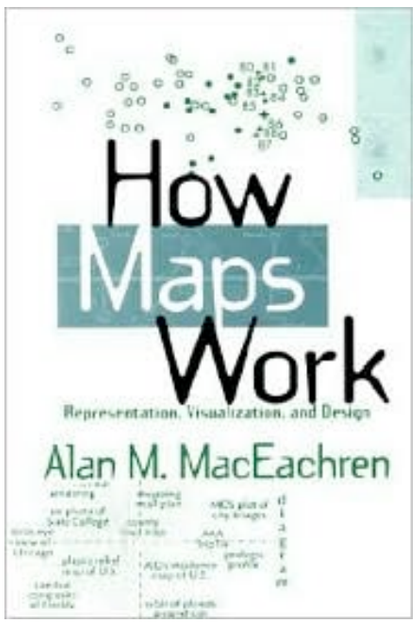
How Maps Work
Representation, Visualization, and Design
Alan M. MacEachren

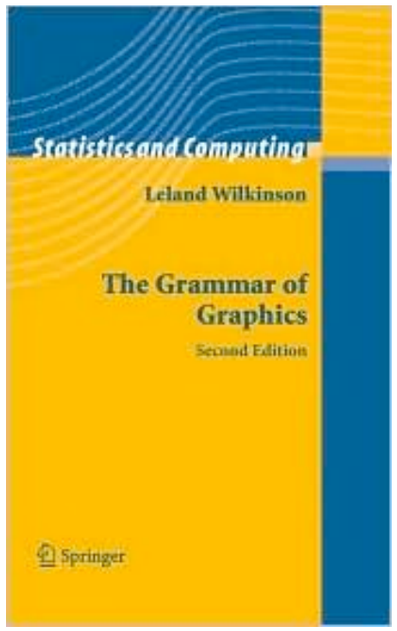
Statistics and Computing
Leland Wilkinson
The Grammar of Graphics
Second Edition
Springer

Visualizing Music

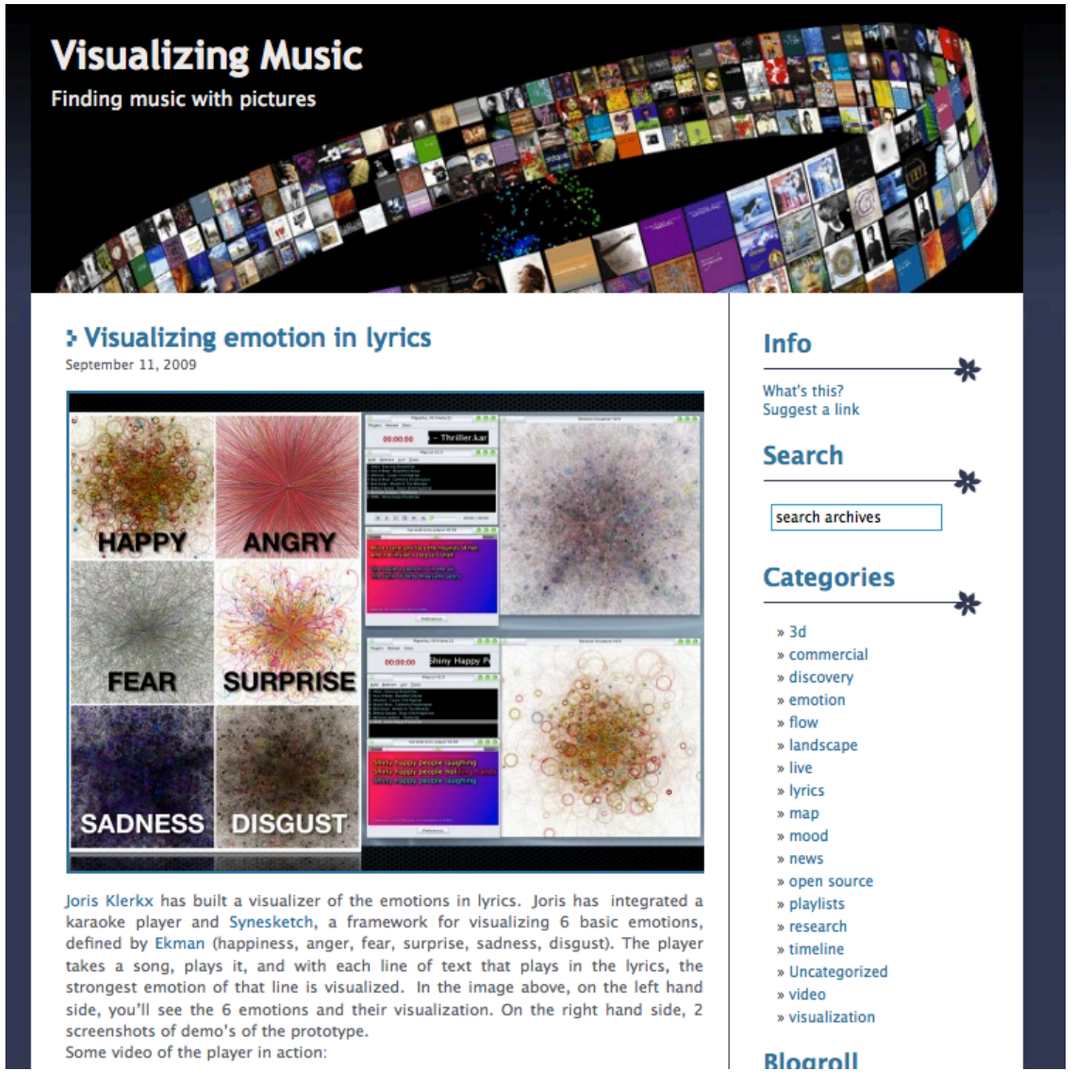

VisualizingMusic.com

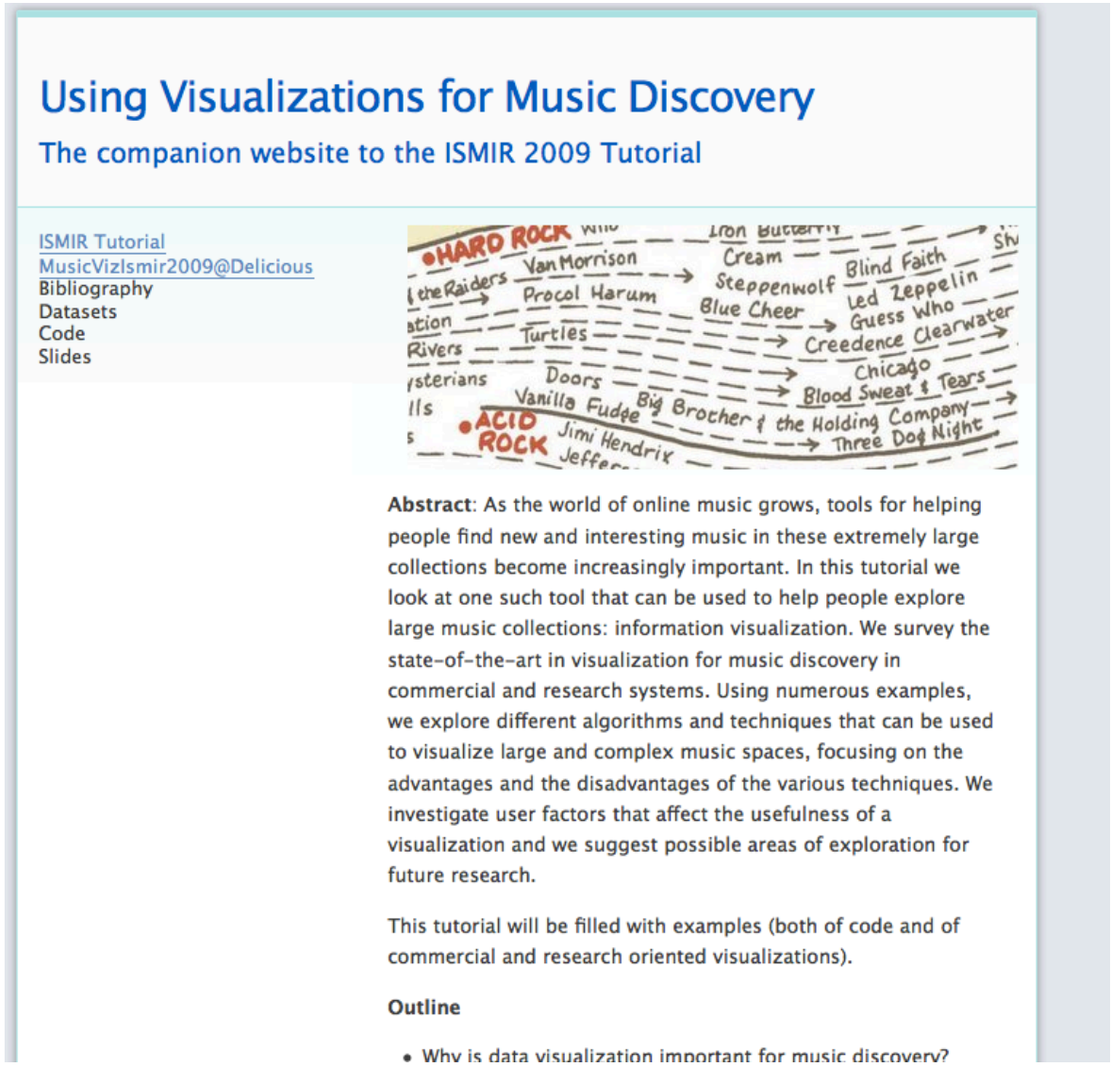

musicviz.googlepages.com/home

Other Tools

Other Analytic* Tools

- Matlab
  - Dimensionality Reduction Toolbox
  - (missing good network viz/analysis option)
  - Extensive Interactive Pub-quality plotting environment

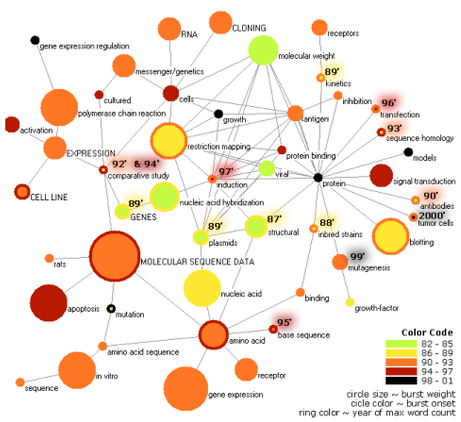

- Pajek (standalone Network Visualization/Analysis Platform)

- IVC Software Framework (research platform for analysis and viz).
- R
  - Excellent network viz/analysis package "network"
  - PCA/MDS (base library), ISOMap {vegan}
    (no LLE, T-SNE yet)
  - Pub quality plotting environments (default, ggplot2, iplots)

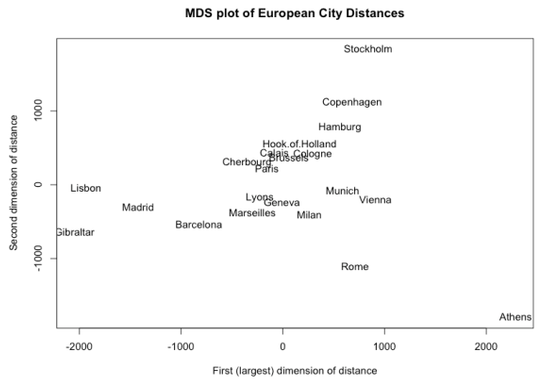

http://ict.ewi.tudelft.nl/~lvandermaaten/Matlab_Toolbox_for_Dimensionality_Reduction.html
http://cran.r-project.org/web/packages/network/index.html
http://cran.r-project.org/web/packages/vegan/index.html
http://vlado.fmf.uni-lj.si/pub/networks/pajek/
http://iv.slis.indiana.edu/sw/index.html

*for publications

Other "Generic" Viz Tools

- Prefuse (java/as3)
- Processing*
- Gephi
- TouchGraph
- Gnuplot
- matplotlib (python)

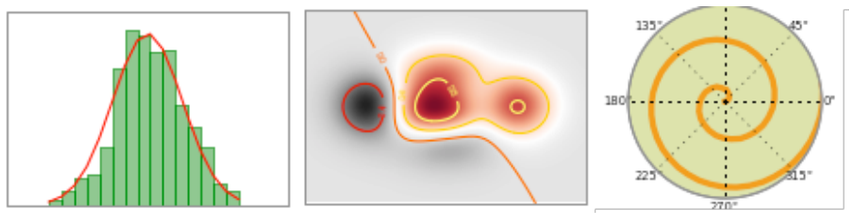

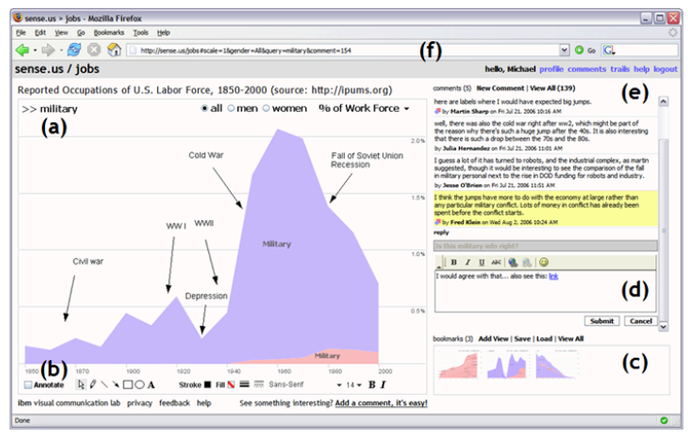

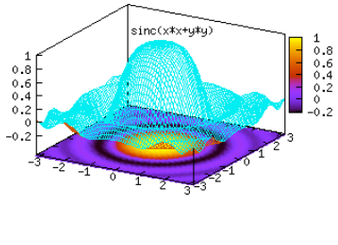

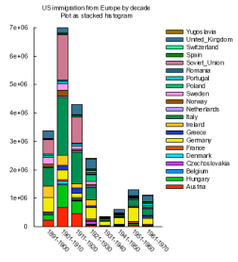

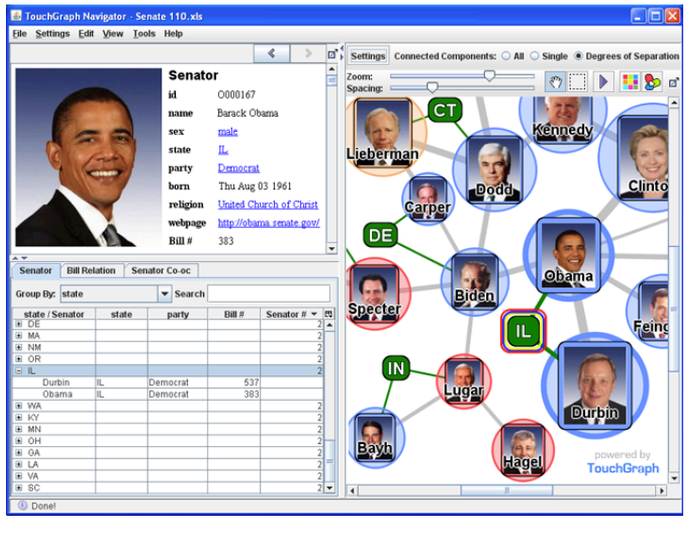

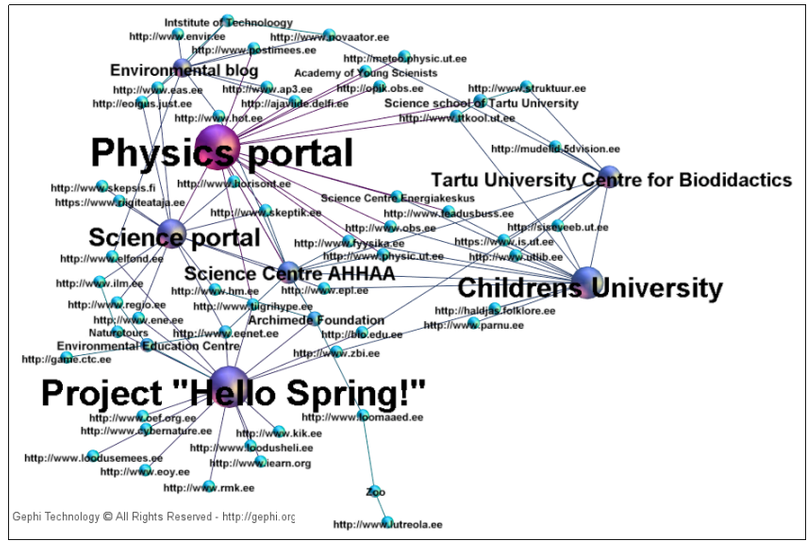

http://www.prefuse.org/
http://processing.org/
http://gephi.org/
http://www.touchgraph.com/navigator.html
http://www.gnuplot.info/
http://matplotlib.sourceforge.net/

Conclusions

- Music visualizations reflect *ways of understanding music*.
- Music understanding draws from many different *type*s of information (musicological, acoustic, social, bibliographic, etc.)
- Music visualizations can help to express this understanding in a way that can help aid navigation and discovery.

Questions?

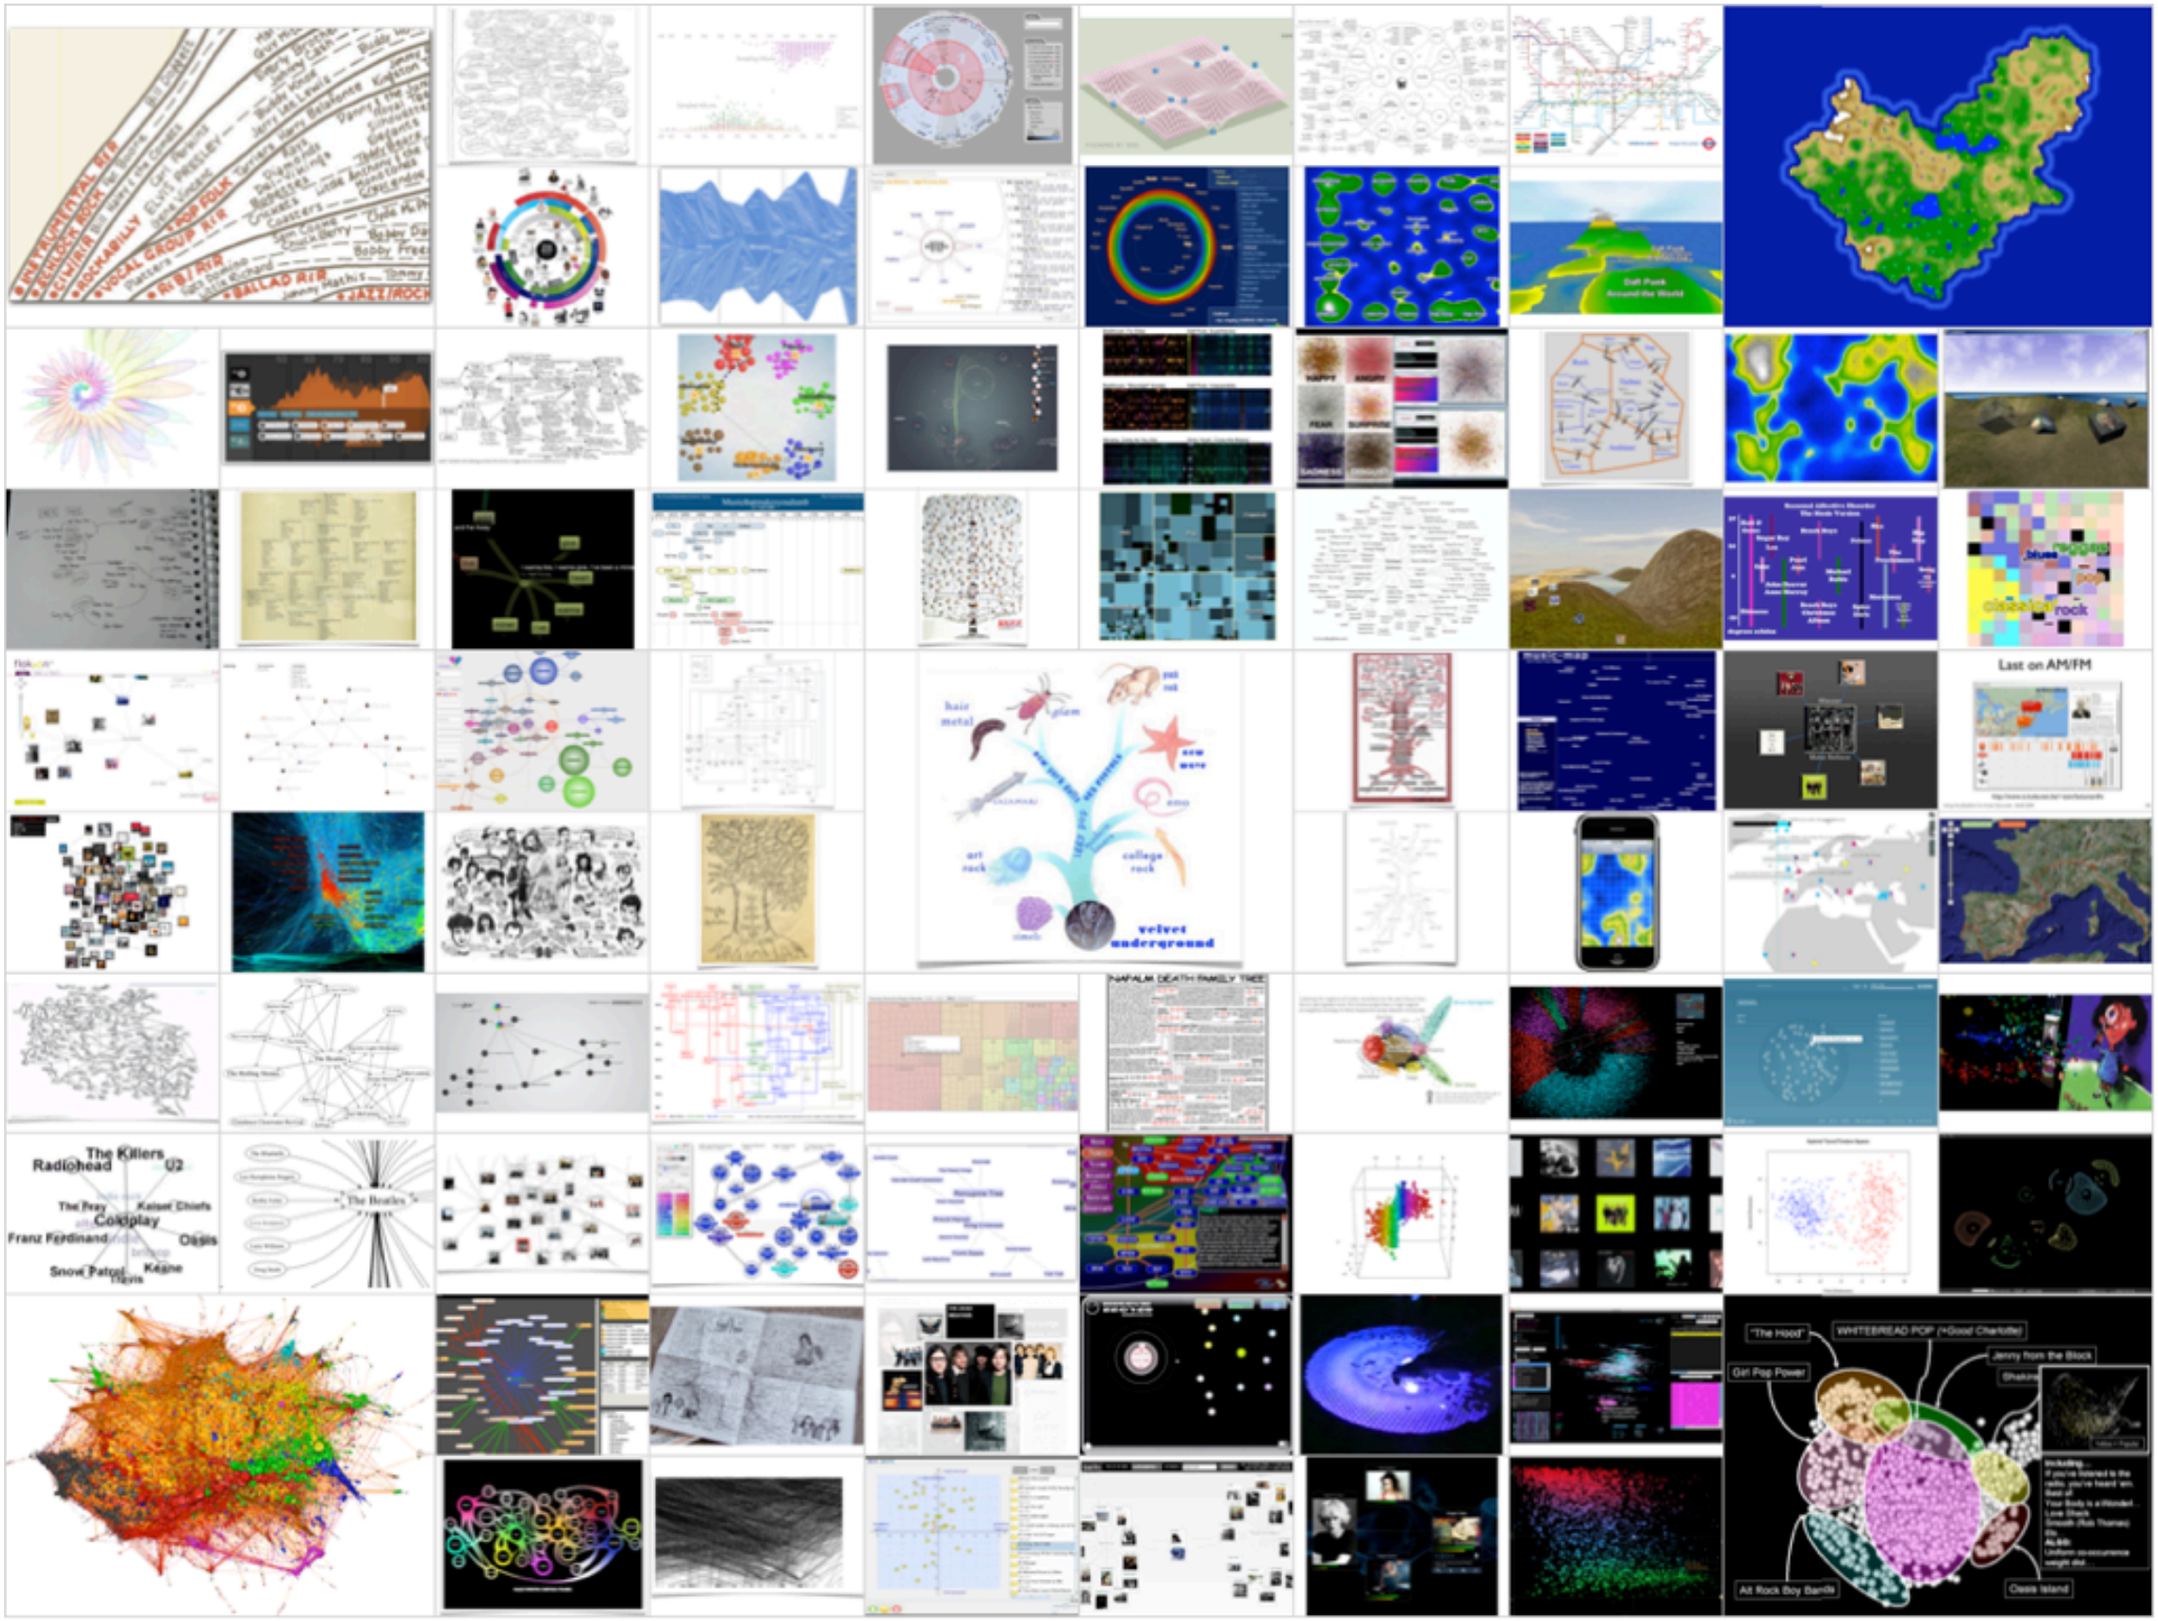
hair metal
art rock
college rock
velvet underground
The Killers
Radiohead
U2
Coldplay
Oasis
Franz Ferdinand
Snow Patrol
Keane
The Beatles
"The Hood"
WHITEBREAD POP' (+Good Charlotte)
Girl Pop Power
Jenny from the Block
Alt Rock Boy Bands
Oasis Island